

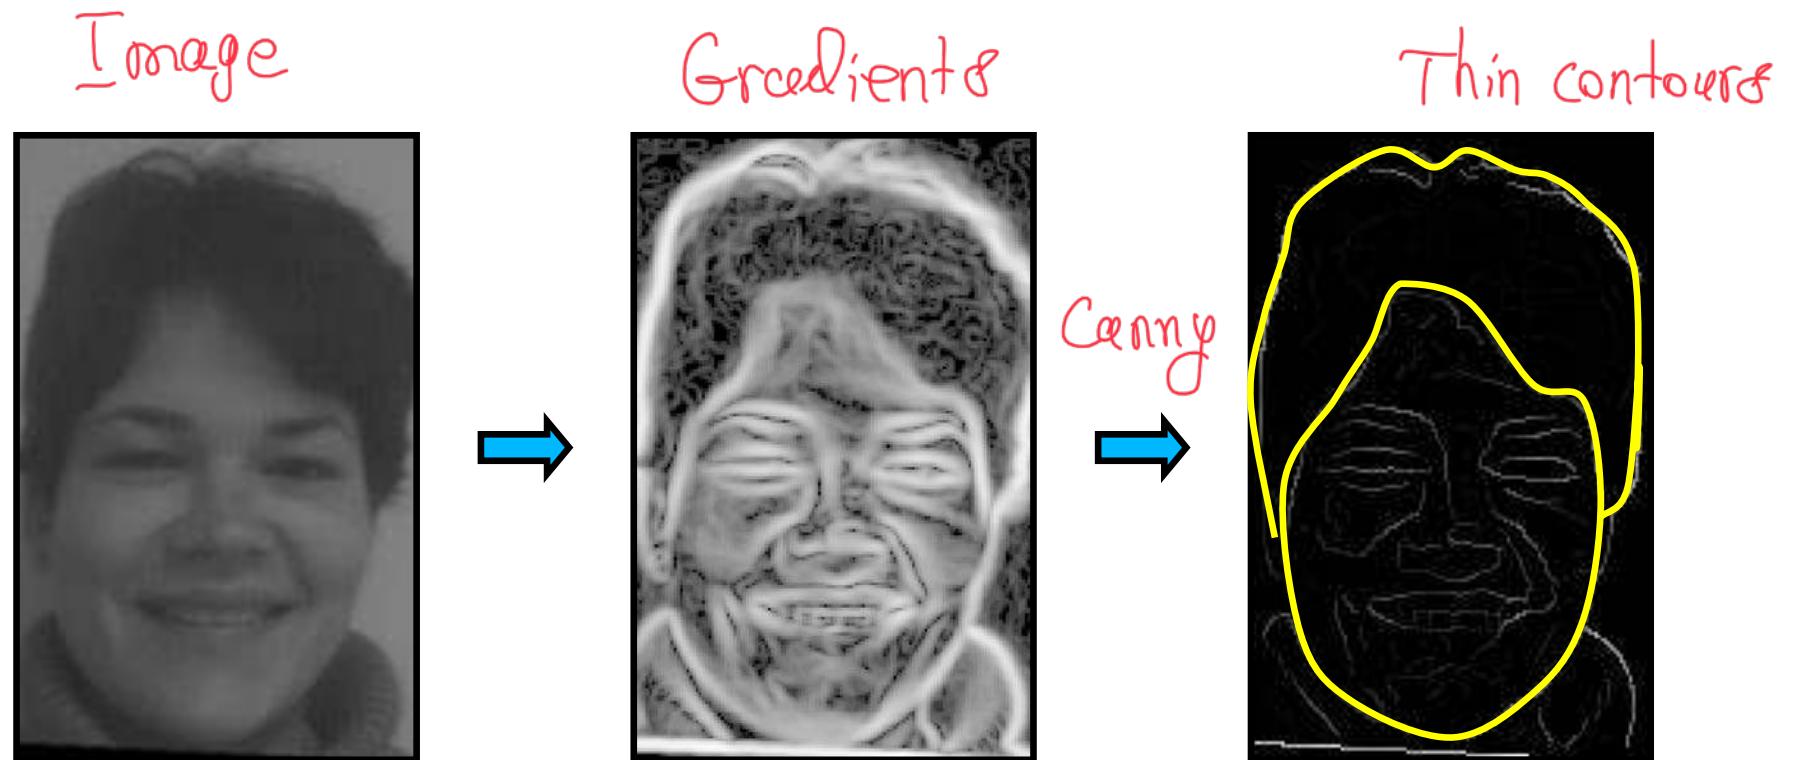
Delineation

Edge map \Rightarrow How they
useable

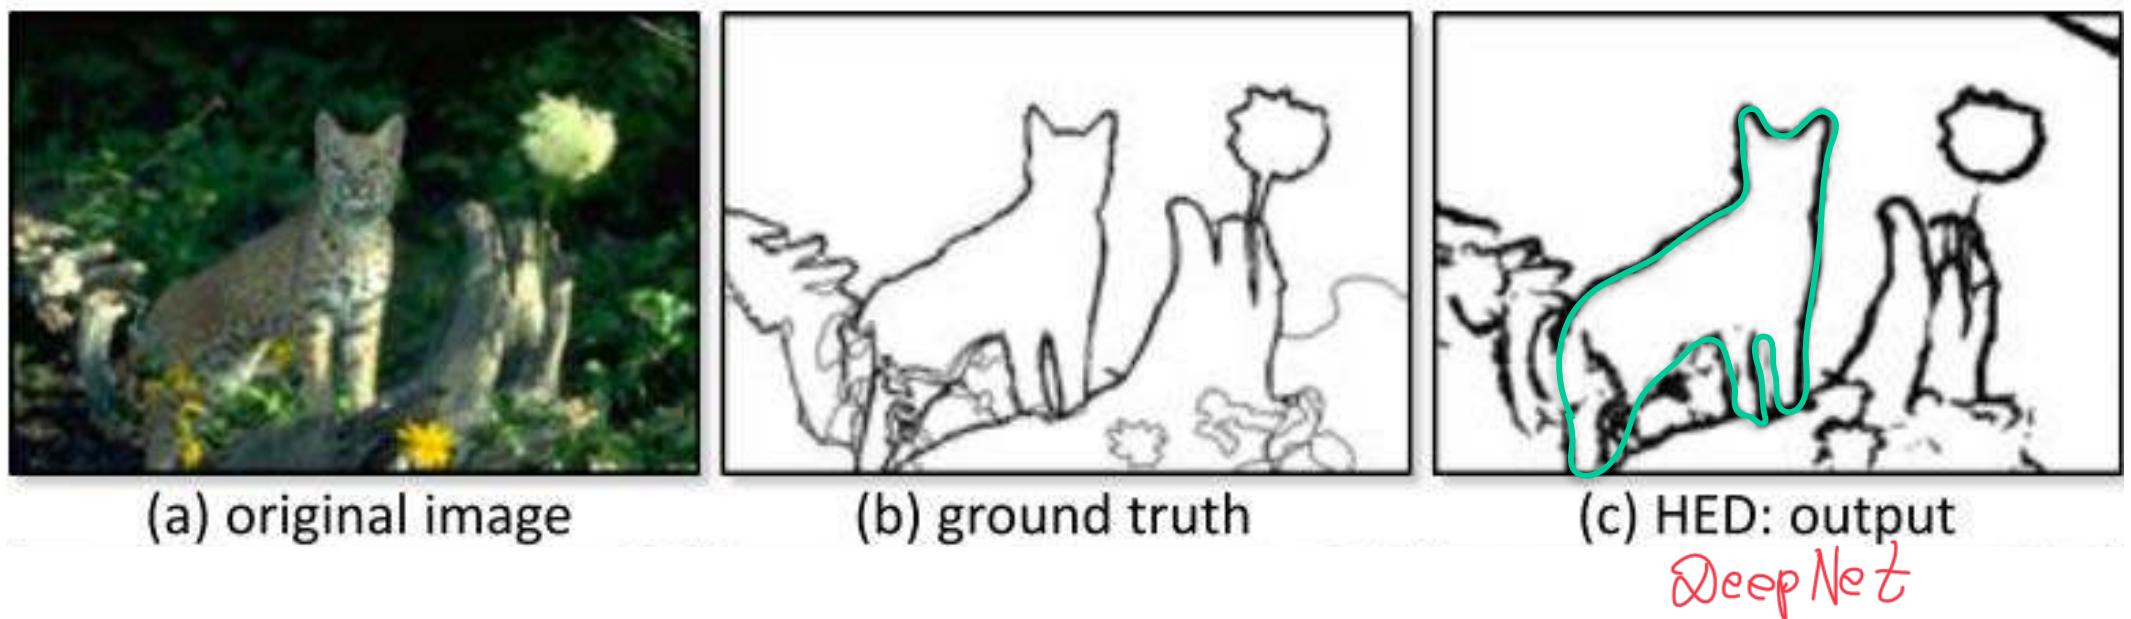
- Dynamic Programming
- Deformable Models
- Hough Transform
- Graph Based Approaches



From Gradients to Outlines

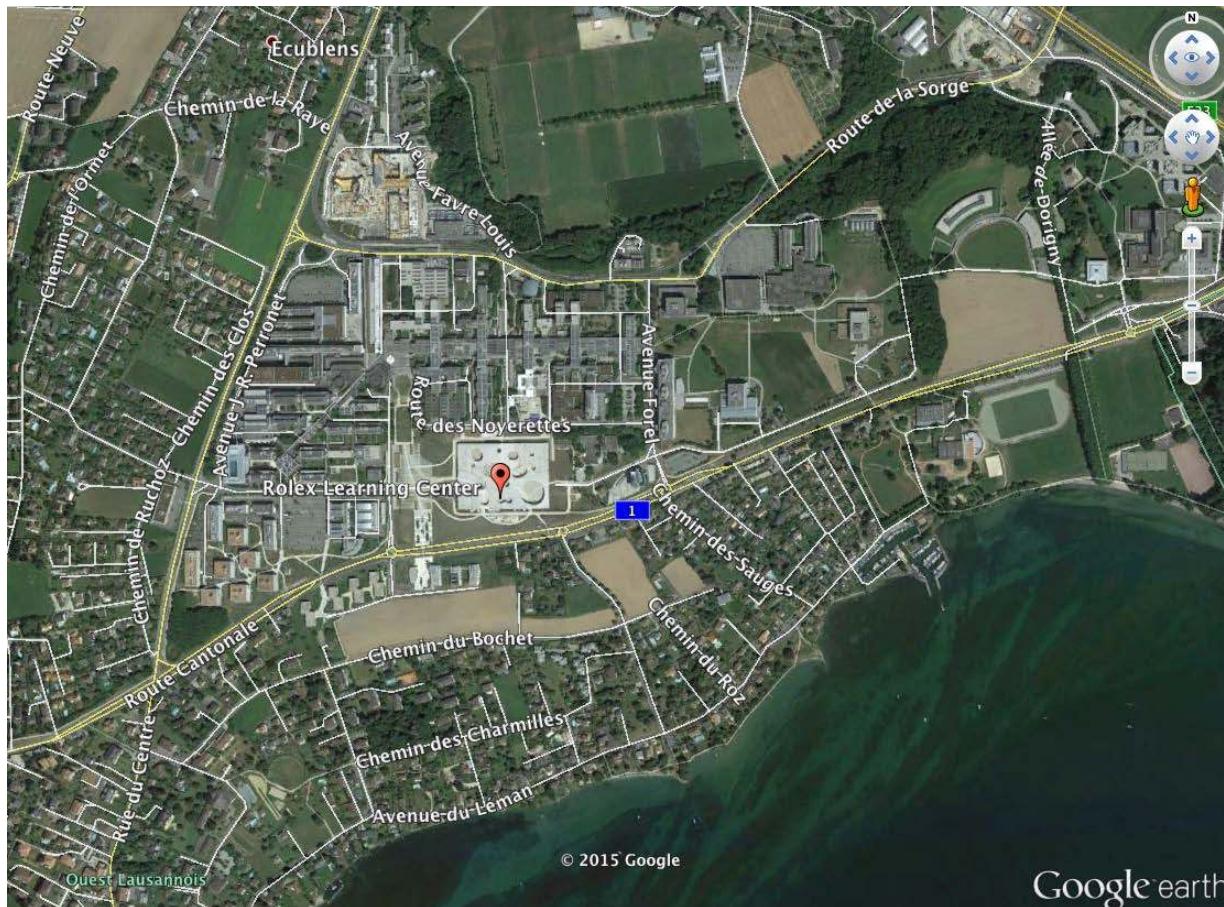


From Deep-Nets to Outlines



→ Still work to do!
not complete contour

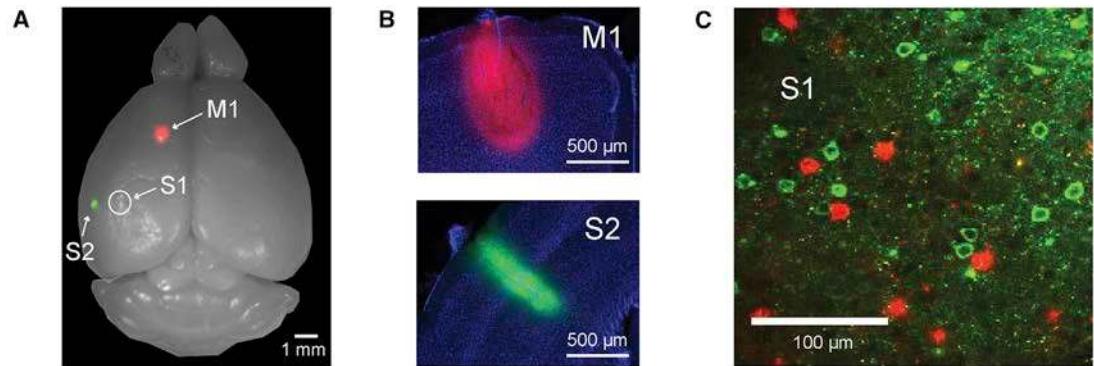
Mapping and Overlays



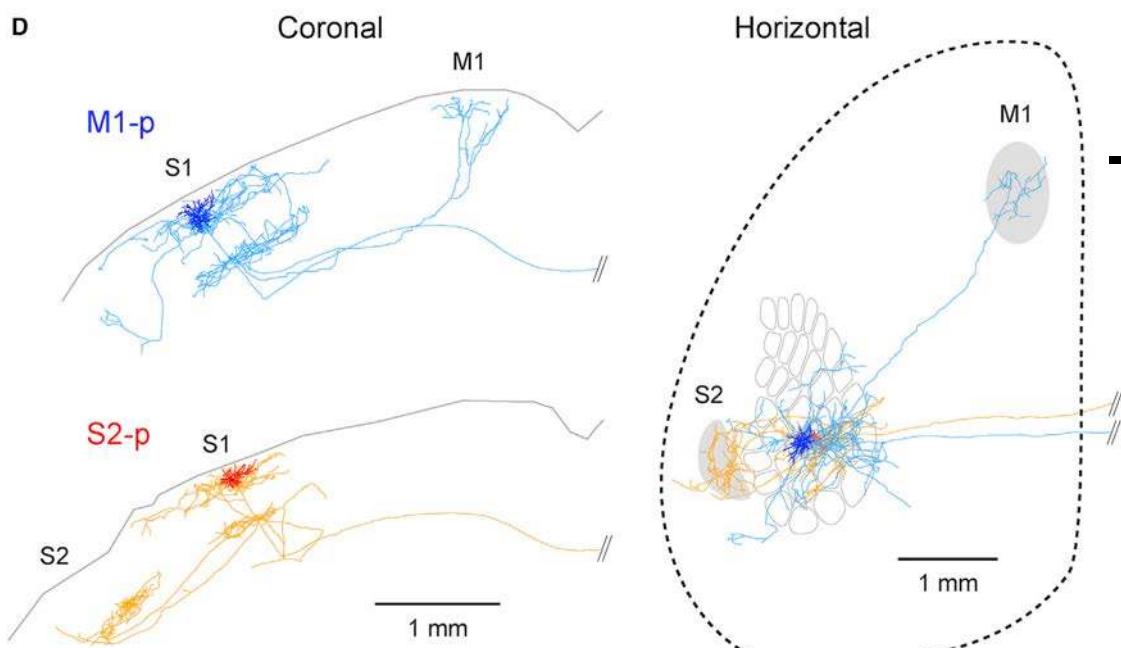
Google earth

Connection
Connectivity matters!

Connectomics



Connectivity matters



→ Topology needed

Analogy



Low level processing

- Uses Deep Nets to find the most promising locations to focus on.

High level processing

- Performs tree-based search when possible.
- Relies on reinforcement learning and other ML techniques to train.

Techniques

Semi-Automated Techniques:

(We'll have human
in the loop)

- Dynamic Programming
- Deformable Models

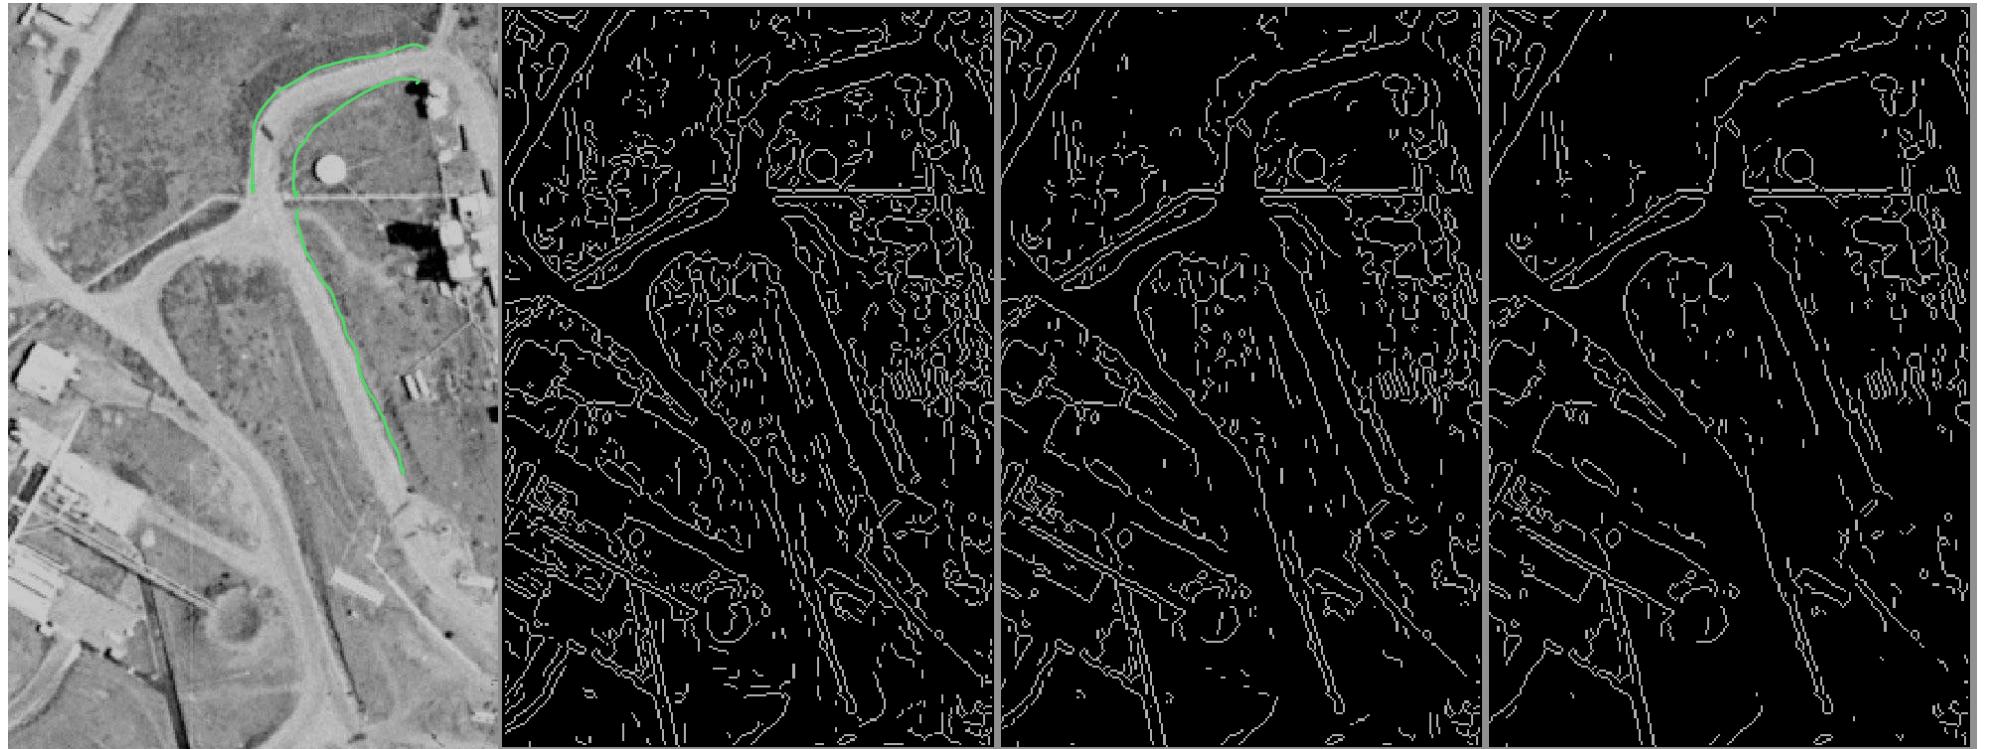
Fully Automated Techniques:

- Hough Transform
- Graph Based Approaches

Reminder: Canny Limitations

of the road
We want edges

not usable for current



- There is no ideal value of σ !
- Deep nets can help but do not solve the problem.

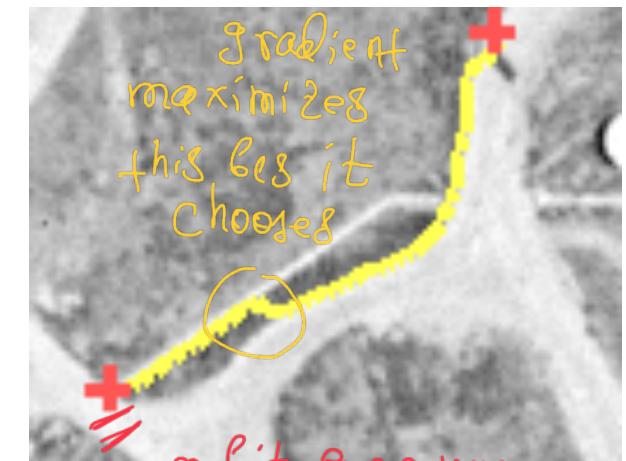
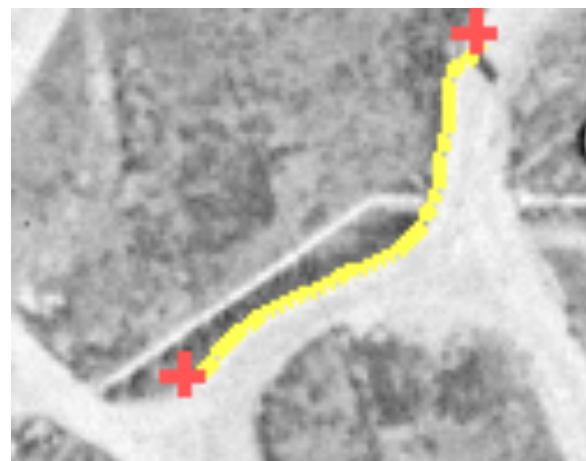
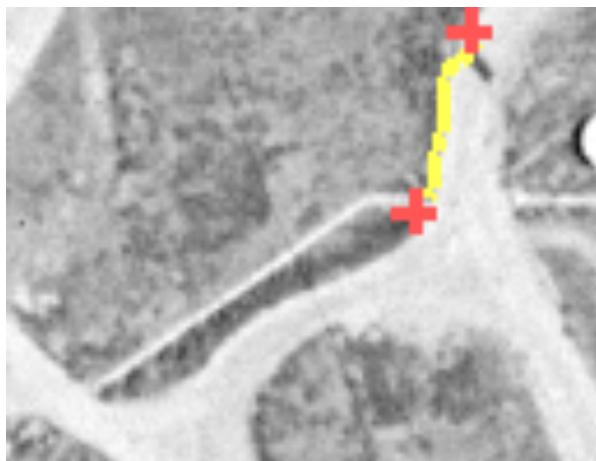
Interactive Delineation



Image



Gradient



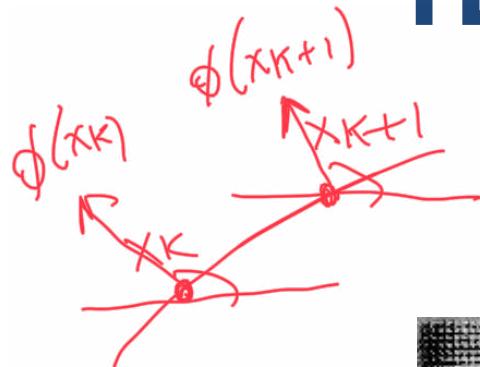
- The user provides the start and end points (**red x**).
- The algorithm does the rest (**yellow line**).

Live Wire in Action

✓ Human Poop

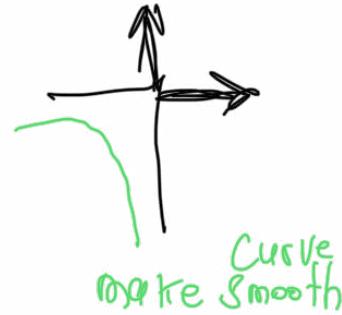


1D Dynamic Programming

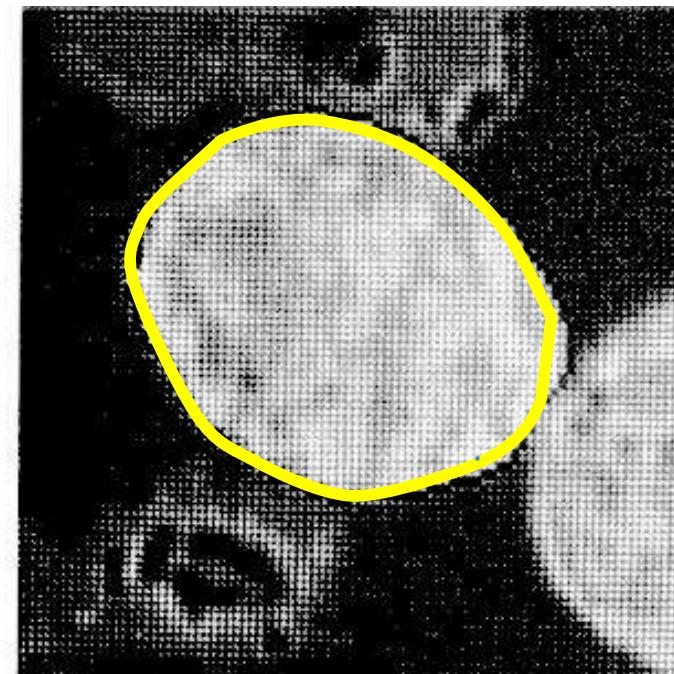
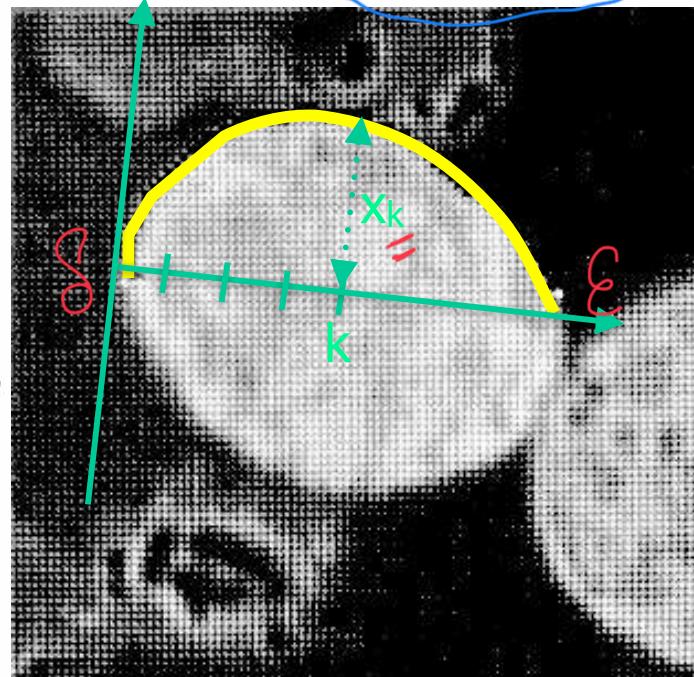


Markov order 1

$$g < t < e \Rightarrow x = f(t)$$



maximize



$$h(x_1, x_2, \dots, x_n) = - \sum_{k=1}^n g(x_k) + \sum_{k=1}^{n-1} r(x_k, x_{k+1})$$

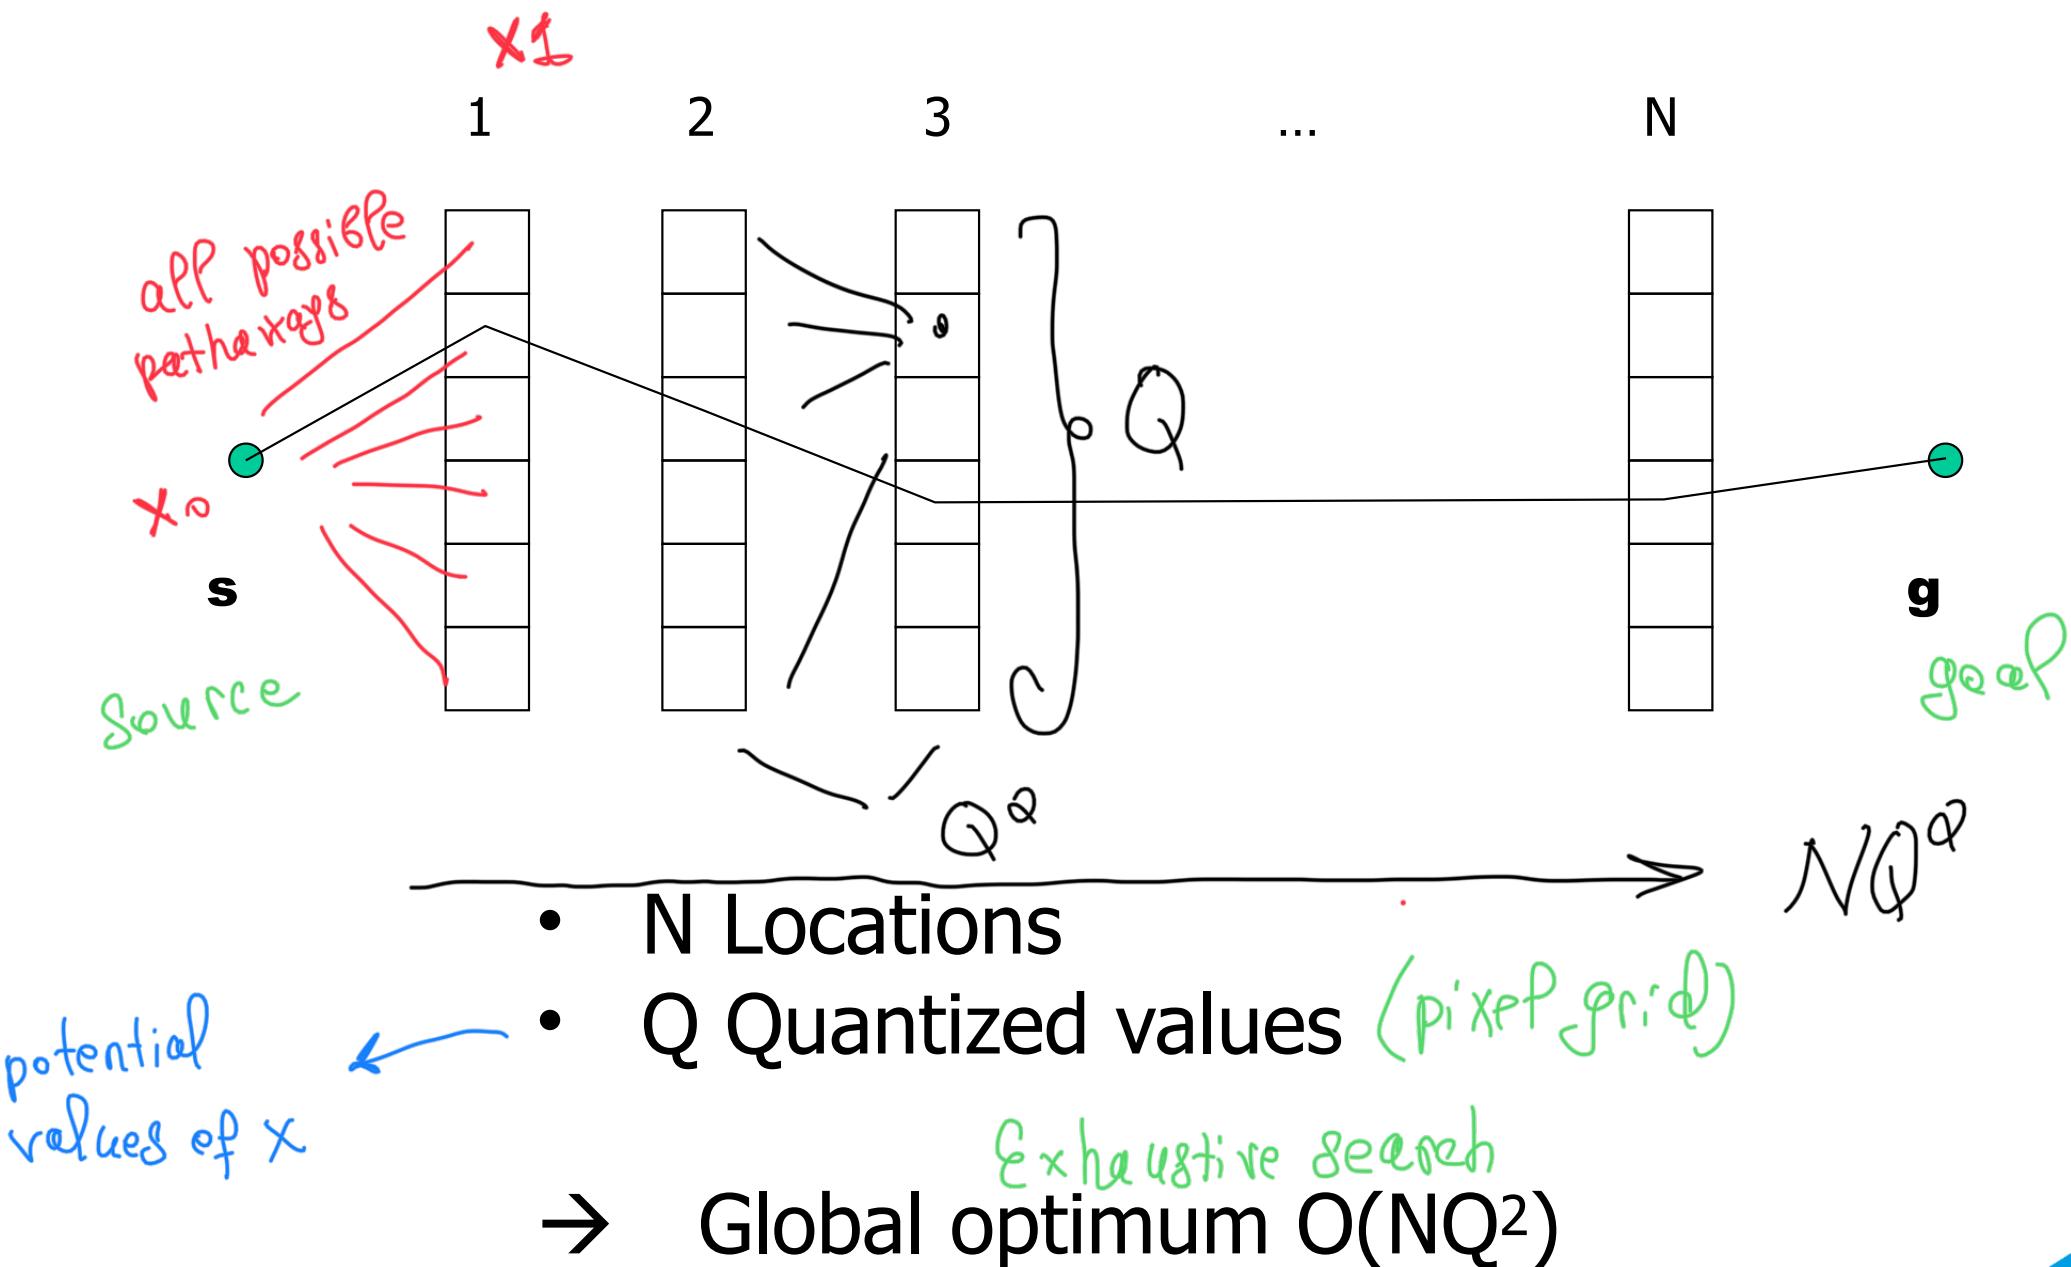
$$r(x_k, x_{k+1}) = \text{diff}(\phi(x_k), \phi(x_{k+1}))$$

gradient defined at x_k
difference between 2 angles

where ϕ denotes the gradient orientation.

1D Dynamic Programming

maximizing gradients $\sum g(x_i)$



1D Dynamic Programming

To find

$$\min_{x_i} h(x_1, x_2, \dots, x_n)$$

where

$$h(x_1, x_2, \dots, x_n) = r(s, x_1) + \sum_{i=1}^{n-1} r(x_i, x_{i+1}) + r(x_n, g)$$

define

$$f_1(x_2) = \min_{x_1} (r(s, x_1) + r(x_1, x_2))$$

$$f_2(x_3) = \min_{x_2} (r(x_2, x_3) + f_1(x_2))$$

$$\vdots \quad \vdots \quad \vdots$$

$$f_{n-1}(x_n) = \min_{x_{n-1}} (r(x_{n-1}, x_n) + f_{n-2}(x_{n-1}))$$

$$\Rightarrow \min h(x_1, x_2, \dots, x_n) = \min_{x_n} (r(x_n, g) + f_{n-1}(x_n))$$

2D Dynamic Programming

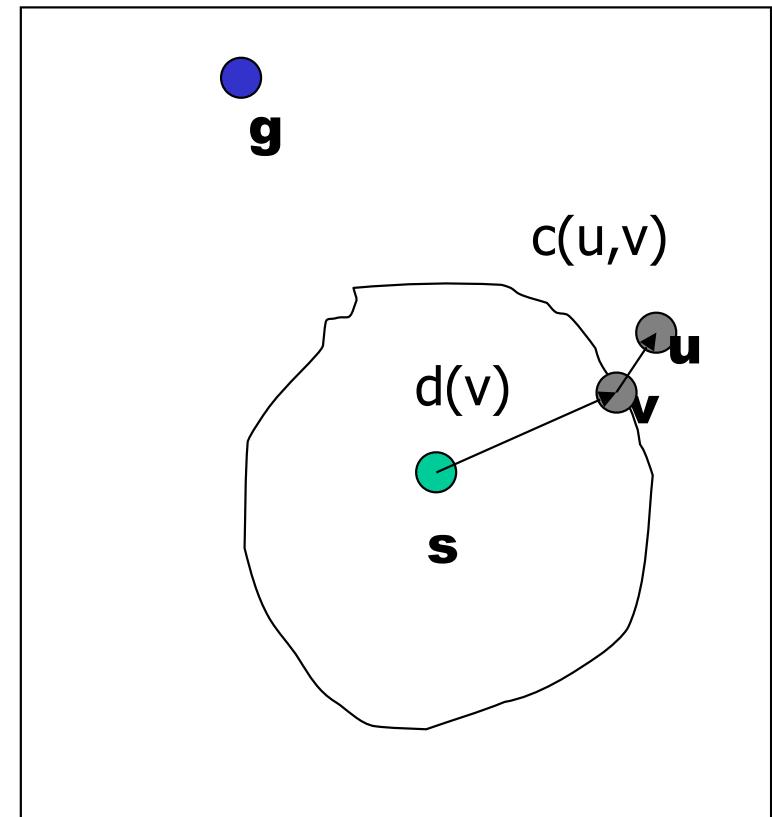
Notations:

s Start point

L List of active nodes
*to nodes which we know
Best paths*

$c(u,v)$ Local costs for link $u \rightarrow v$

$d(v)$ Total cost from s to v



Dijkstra Path Expansion



Open nodes represent "unvisited" nodes. Filled nodes are visited ones, with color representing the distance: The greener, the shorter the path. Nodes in all the different directions are explored uniformly, appearing more-or-less as a circular **wavefront**.

Dijkstra's Algorithm

Initialization :

$$d(s) \leftarrow 0 \text{ and } d(u) \leftarrow \infty \text{ for } u \neq s$$

$$T = \emptyset$$

$$v = s$$

not connected yet

List of paths

progressively create fronts until

end of currently shortest path

Loop until goal is reached :

$$T \leftarrow T \cup \{v\}$$

for all $v \rightarrow u$ edges such that $u \notin T$

$$\text{if } d(v) + c(v,u) \stackrel{\text{cost}}{<} d(u)$$

$$d(u) \leftarrow d(v) + c(v,u)$$

end

end

$$v = \operatorname{argmin}_{w \notin T} d(w)$$

Maintain a sorted list of paths

Live Wire

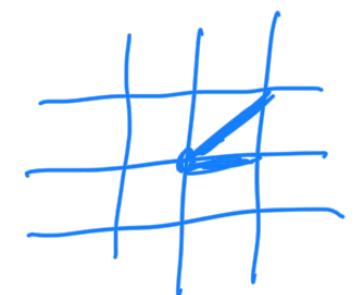
- Sorting is the expensive operation. Normally $n \log(n)$, but can be reduced to $\log(n)$ if all costs are integer costs

- Local costs computed using gradient:

edge along which gradient
is high \Rightarrow low cost
(high)
(low)

$$c(u,v) = 255 - \frac{1}{2} (g(u) + g(v))$$

gradient



- Diagonal penalized by multiplying cost of non diagonal edges by:

$$\frac{1}{\sqrt{2}} = \frac{5}{7}$$

for longer edges, you need to
(like diag)
reduce gradient

more edges \Rightarrow more cost

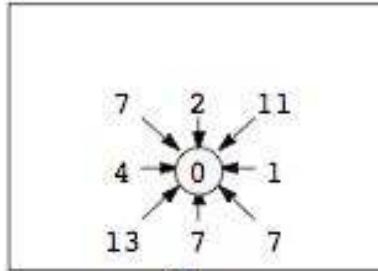
- Add a constant cost for each edge.

maybe gradient was
small ≈ 0
(all edges would be equivalent)

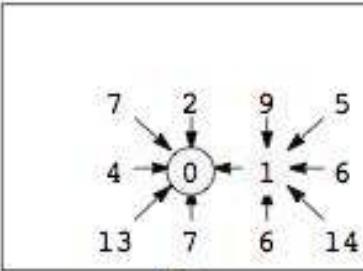
Cost Expansion

11	13	12	9	5	8	3	1	2	4	10
14	11	7	4	2	5	8	4	6	3	8
11	6	3	5	7	9	12	11	10	7	4
7	4	6	11	13	18	17	14	8	5	2
6	2	7	10	15	15	21	19	8	3	5
8	3	4	7	9	13	14	15	9	5	6
11	5	2	8	3	4	5	7	2	5	9
12	4	(2)	1	5	6	3	2	4	8	12
10	9	7	5	9	8	5	3	7	8	15

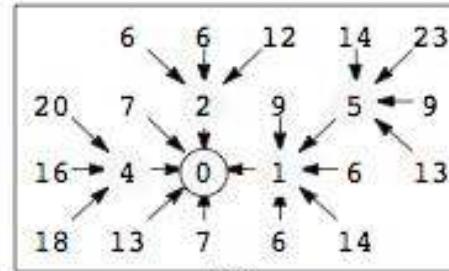
(a)



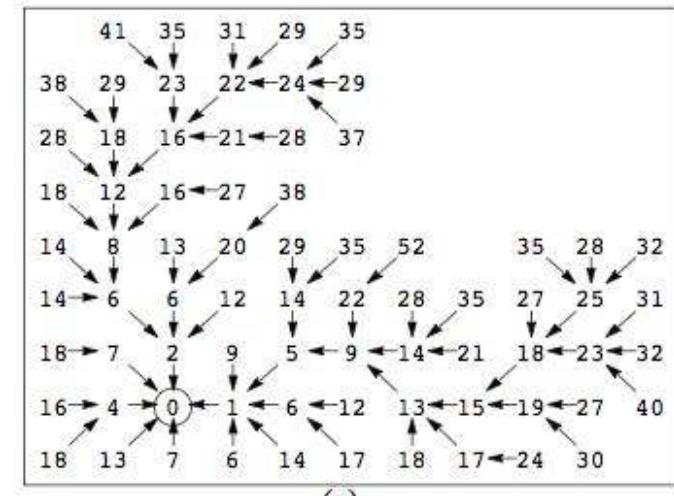
(b)



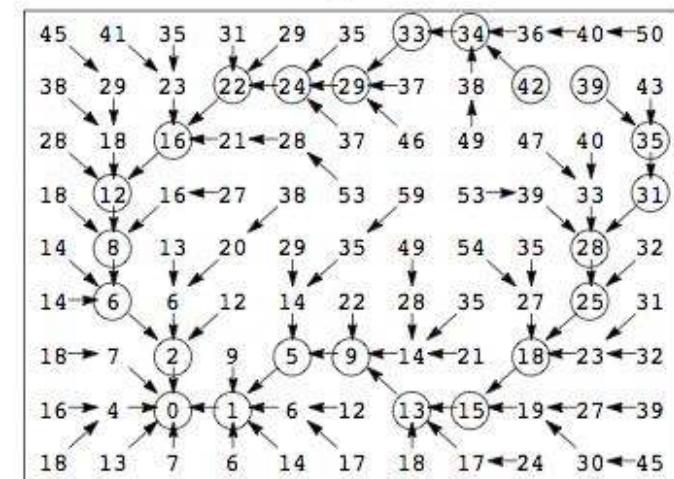
(c)



(d)



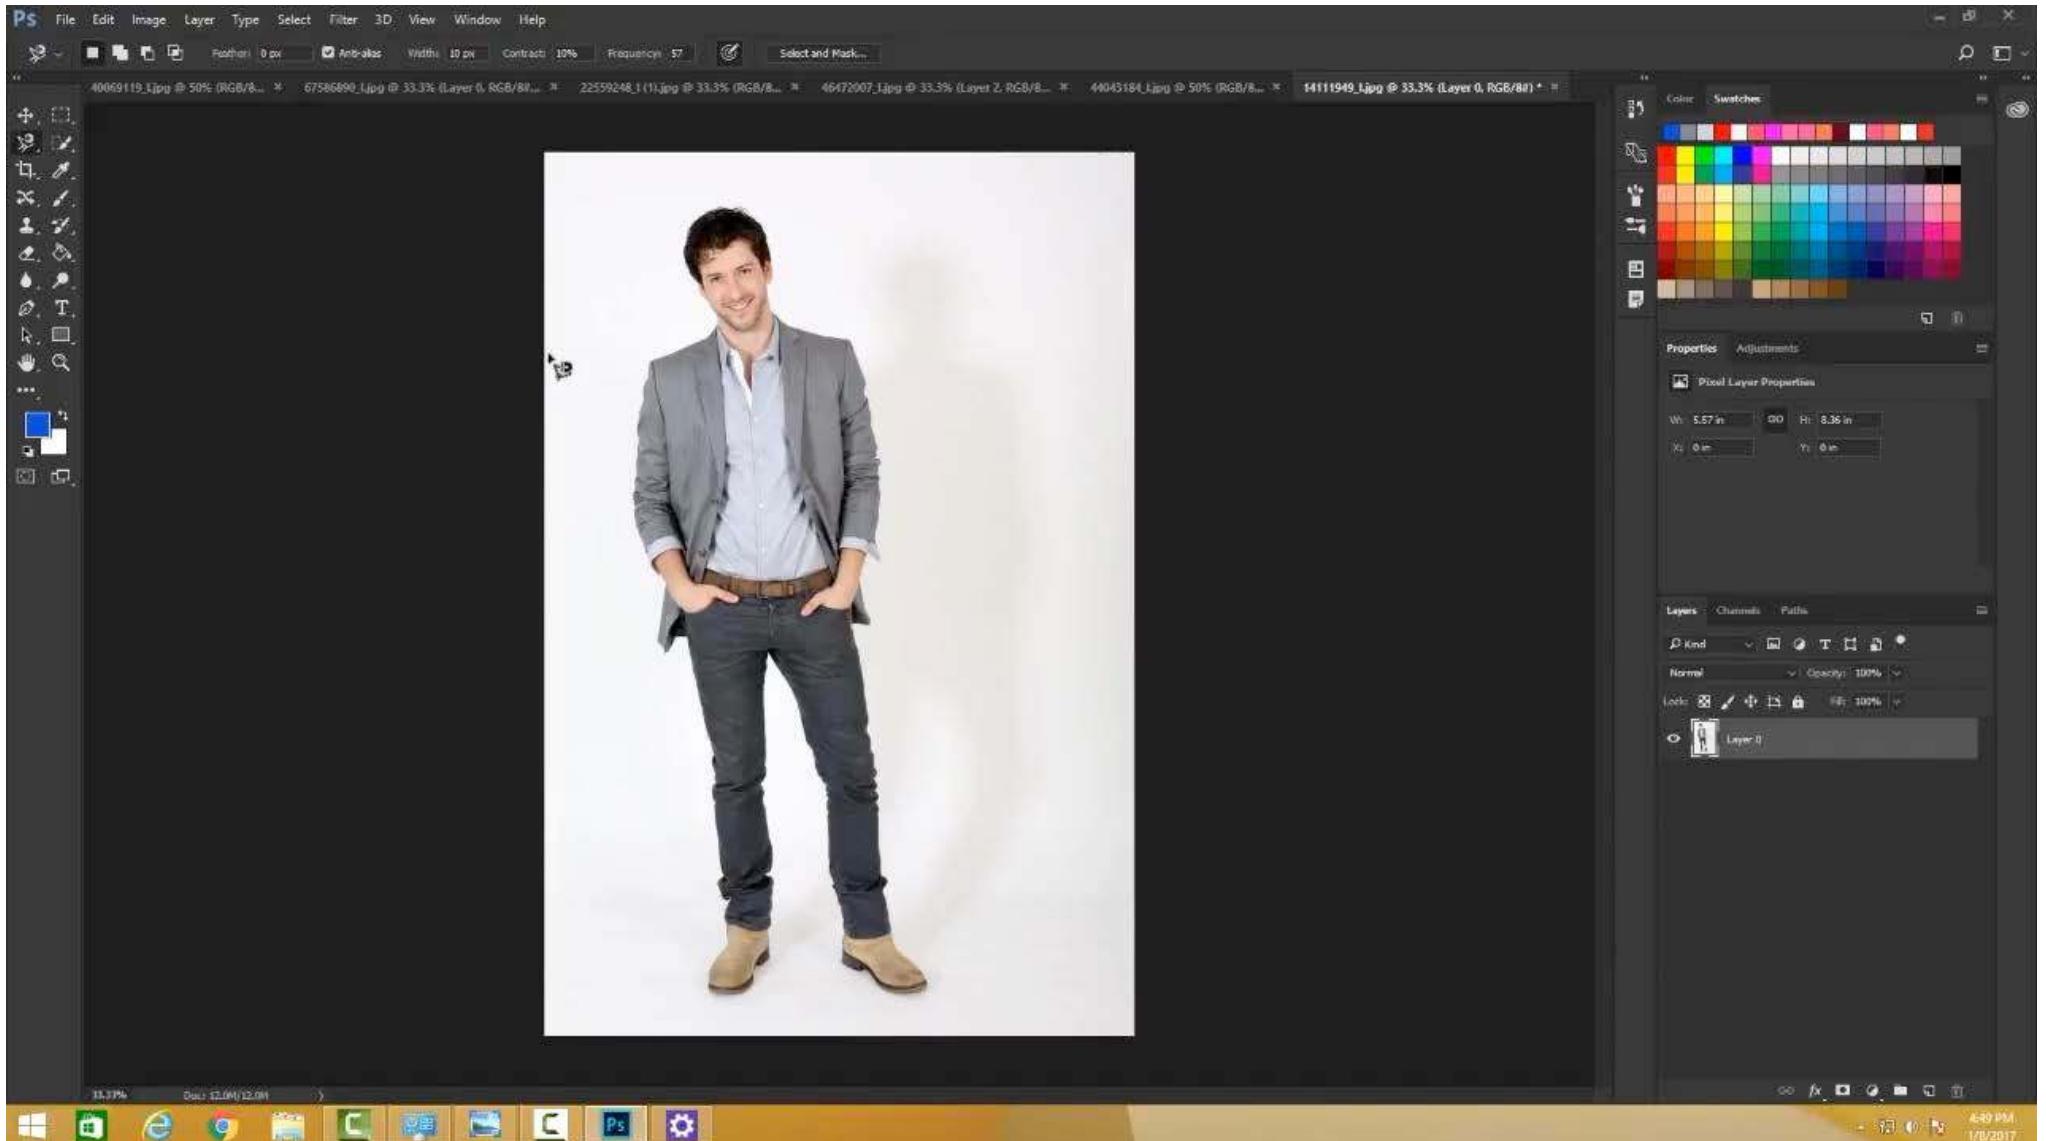
(e)



(f)

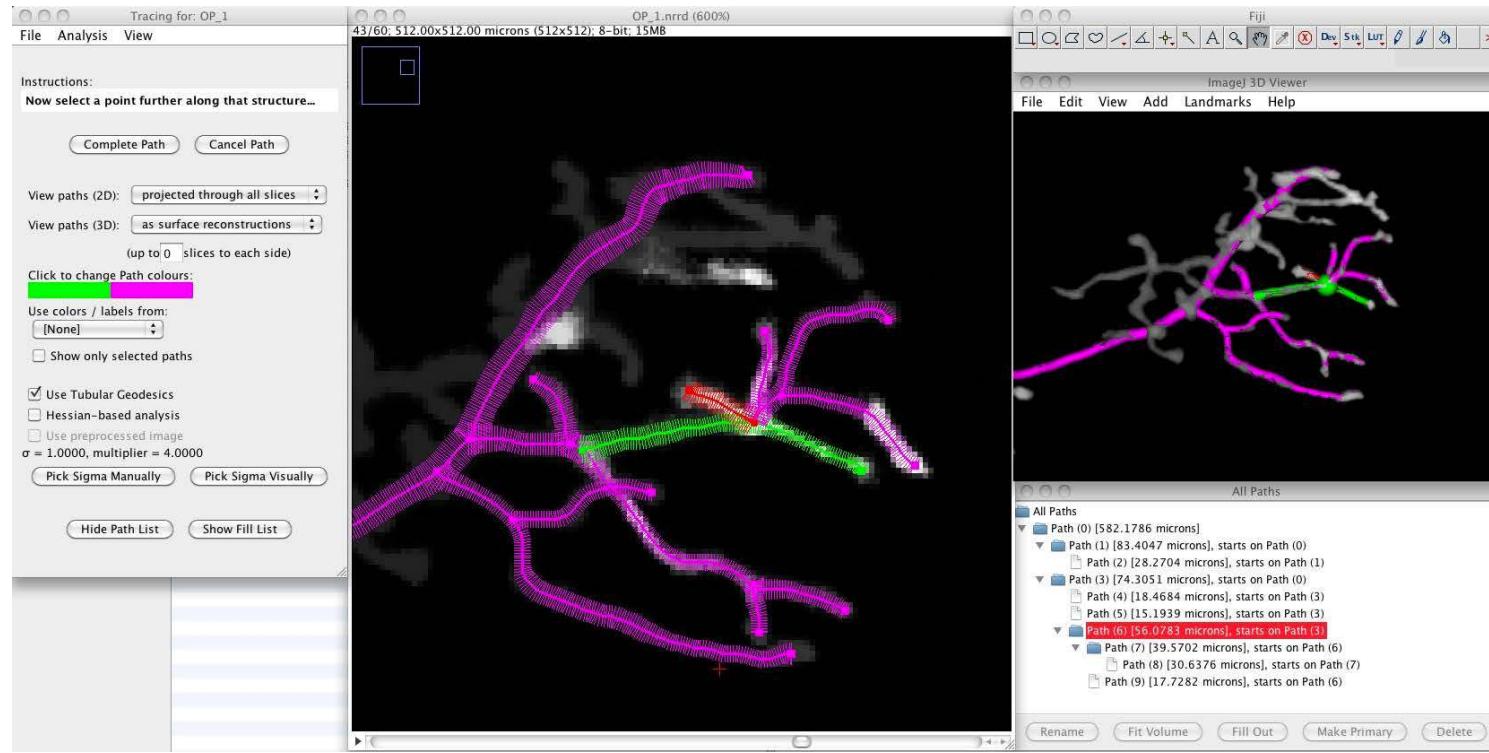
(a) Local cost map. (b) Seed point expanded. (c) 2 points expanded. (d) 5 points expanded. (e) 47 points expanded. (f) Completed cost path-pointer map with optimal paths shown from nodes with total costs 42 and 39.

Magnetic Lasso in Photoshop



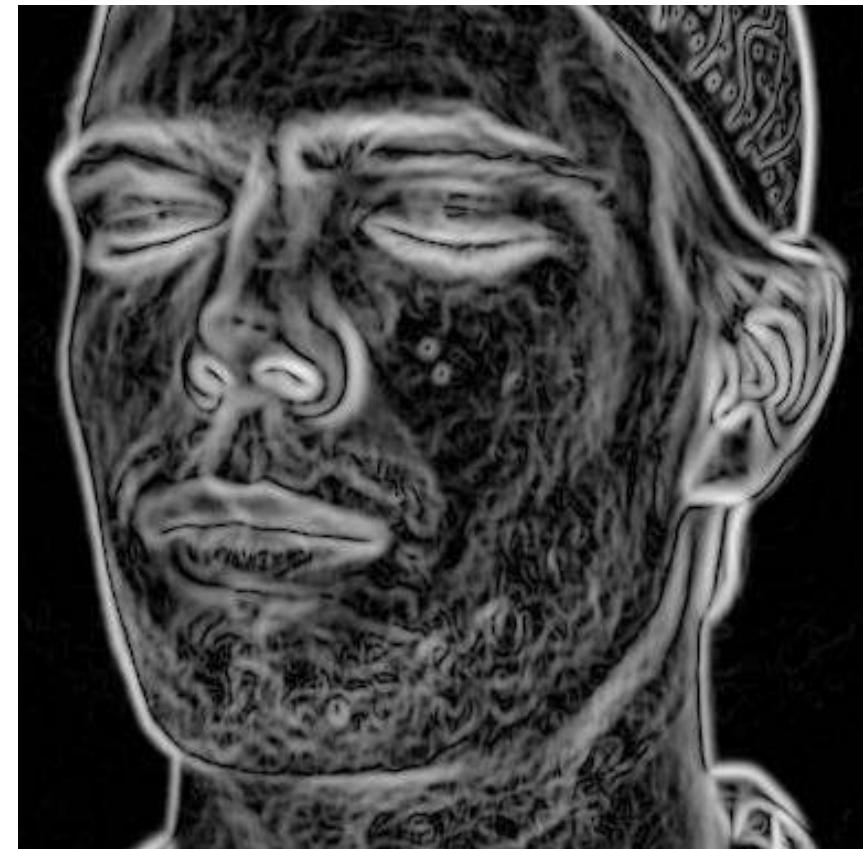
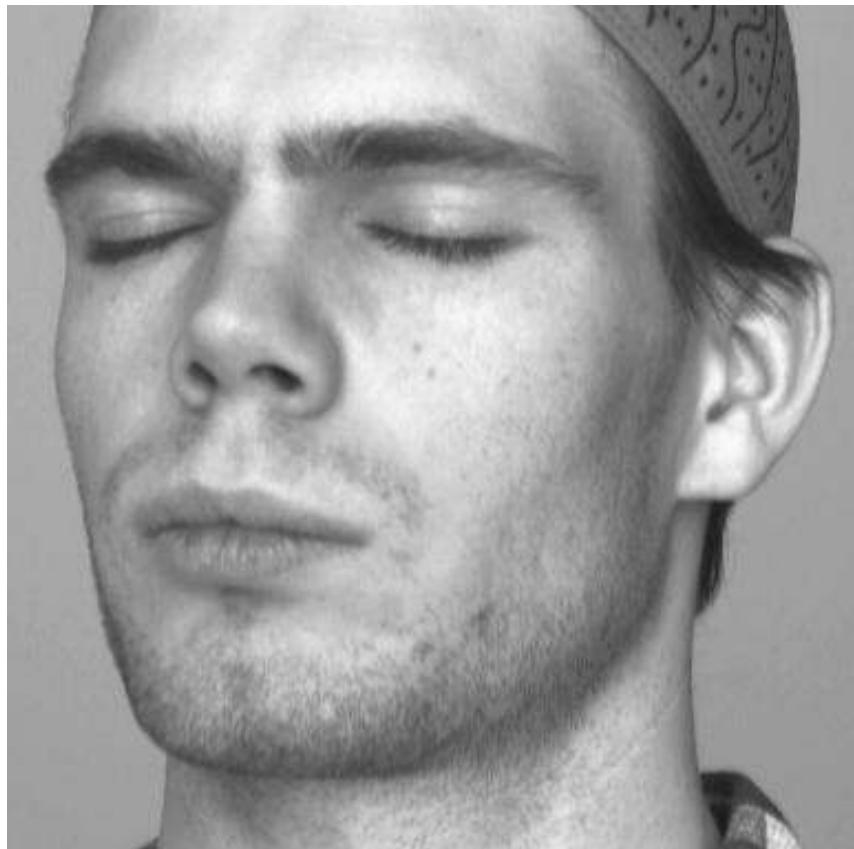
Integrating the LiveWire into a powerful interface that allows a user to correct mistakes yields a useful tool.

Tracing Neurons



- In the biomedical world, images are 3D cubes of data.
- The approach extends naturally to tracking of 3D structures such as dendritic trees in the brain, blood vessels, etc ...

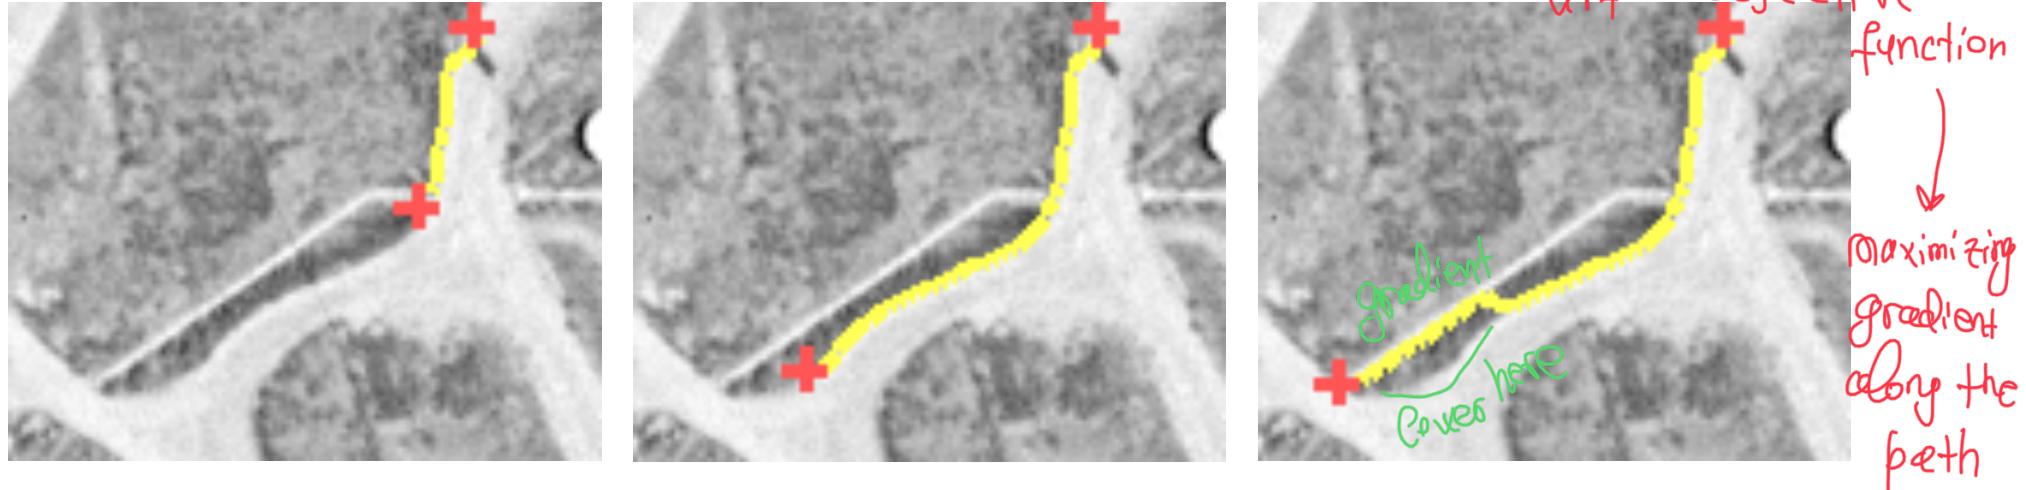
Face Image



Live Wire



Limitations



- The “optimal” path is not always the “best” one.
 - Difficult to impose global constraints. (road is globally smooth)
Markov states \mapsto points and their immediate neighbors
 - The cost grows exponentially with the dimension of the space in which we work.
Dynamic programming \Rightarrow global optimum (not as expressive as you want)
- > Must often look for local, as opposed to global, optimum using gradient descent techniques.

Techniques

Semi-Automated Techniques:

- Dynamic programming
- Deformable Models

Fully Automated Techniques:

- Hough transform
- Graph Based Approaches

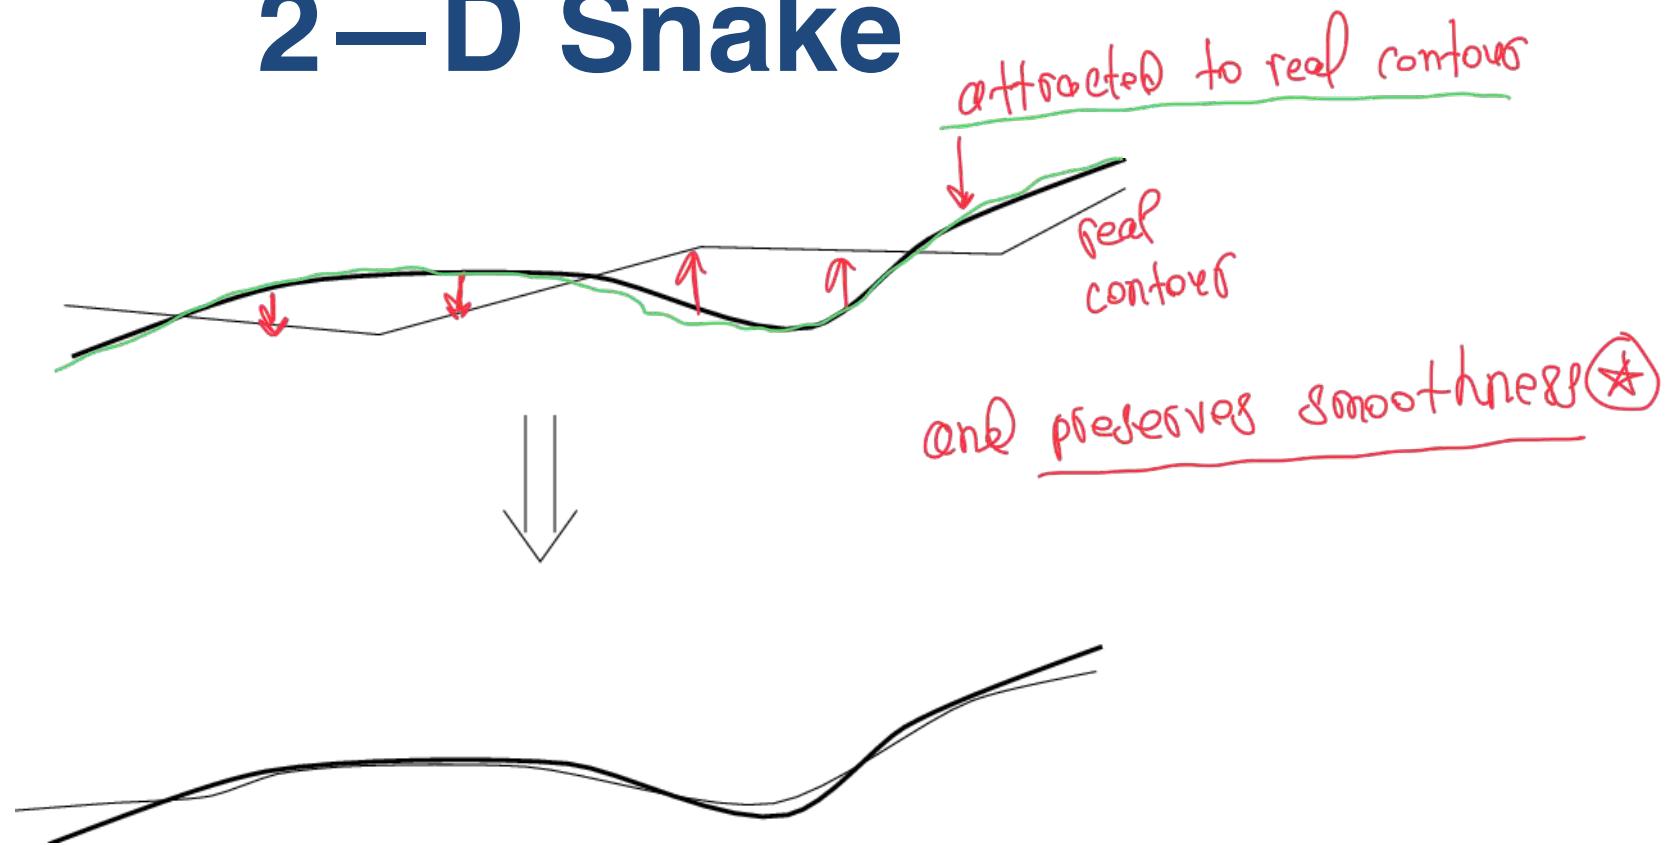
Snakes



Prop: Rough approximation
(contour being deformed)
Contour is attracted
to gradient &

When it Deforms:
i) it remains smooth
ii) maximize gradient on
the curve

2-D Snake



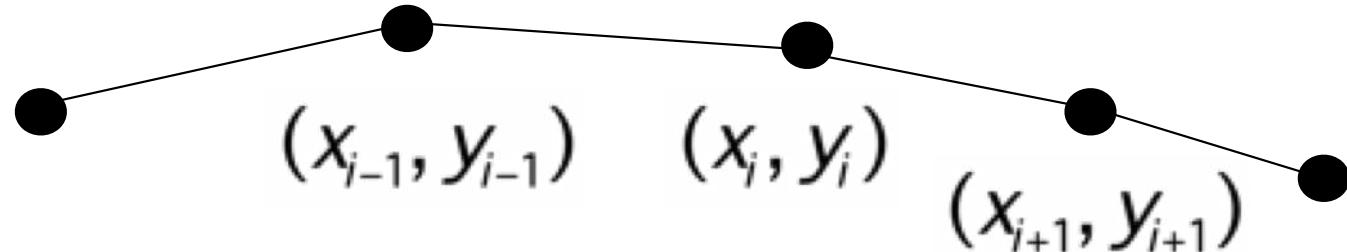
Deformable contours that

- Maximize the gradient along the curve;
 - Minimize their deformation energy. \rightarrow remains smooth
- > Interactive tools for contour detection that can be generalized to handle sophisticated models

Polygonal Approximation

defined by its vertices

Compromise between largest possible gradient and smallest possible curvature



Weighting coefficient

$$E = -\frac{\lambda}{N+1} \sum_{i=0}^N G(x_i, y_i) + -\frac{1}{N} \sum_{i=1}^N ((2x_i - x_{i-1} - x_{i+1})^2 + (2y_i - y_{i-1} - y_{i+1})^2)$$

Average gradient

of 2nd derivatives
finite difference approximation
 $\frac{\partial x}{\partial s}$ $\frac{\partial y}{\partial s}$

Average sum of squared 2nd derivatives

\approx

Average sum of square curvature

Matrix Notation

along the curve
 $\nabla G(x,y)$ no gradients

$$E = E_G + 1/2X^t K X + 1/2Y^t K Y$$

$$X = [x_1, \dots, x_N]^t$$

$$Y = [y_1, \dots, y_N]^t$$

Sparse matrix

Local Optimum

Γ

$$\frac{\delta E}{\delta X} = \begin{bmatrix} \frac{\partial E}{\partial x_1} & \dots & \frac{\partial E}{\partial x_n} \end{bmatrix}$$

We are looking for minimum of E

$$E = E_G + \frac{1}{2} X^T K X + \frac{1}{2} Y^T K Y$$

$$\frac{\delta E}{\delta X} = \frac{\delta E_G}{\delta X} + KX = 0$$

$$\frac{\delta E}{\delta Y} = \frac{\delta E_G}{\delta Y} + KY = 0$$

positive but not invertible

But K is not invertible!

Dynamics

compute $\{X_0, X_1, \dots, X_T\}$
 $\{Y_0, Y_1, \dots, Y_T\}$

rubber band

Embed curve in a viscous medium and solve at each step:

*faster it goes,
more slow down*

$$0 = \frac{\partial E}{\partial X} + \alpha \frac{dX}{dt} = \frac{\partial E_G}{\partial X} + KX + \alpha \frac{dX}{dt}$$

$$0 = \frac{\partial E}{\partial Y} + \alpha \frac{dY}{dt} = \frac{\partial E_G}{\partial Y} + KY + \alpha \frac{dY}{dt}$$

Iterating

Reformed curves \Rightarrow computing X_t, Y_t

At every step:

Solving 2 differential equations

$$0 = \frac{\delta E_G}{\delta X} + KX_t + \alpha(X_t - X_{t-1}) \quad \xleftarrow{\text{finite-difference expression of derivative}} \quad (K + \alpha I)X_t = \alpha X_{t-1} - \frac{\delta E_G}{\delta X}$$

$$0 = \frac{\delta E_G}{\delta Y} + KY_t + \alpha(Y_t - Y_{t-1}) \quad \Rightarrow \quad (K + \alpha I)Y_t = \alpha Y_{t-1} - \frac{\delta E_G}{\delta Y}.$$

derivative
of gradient
terms

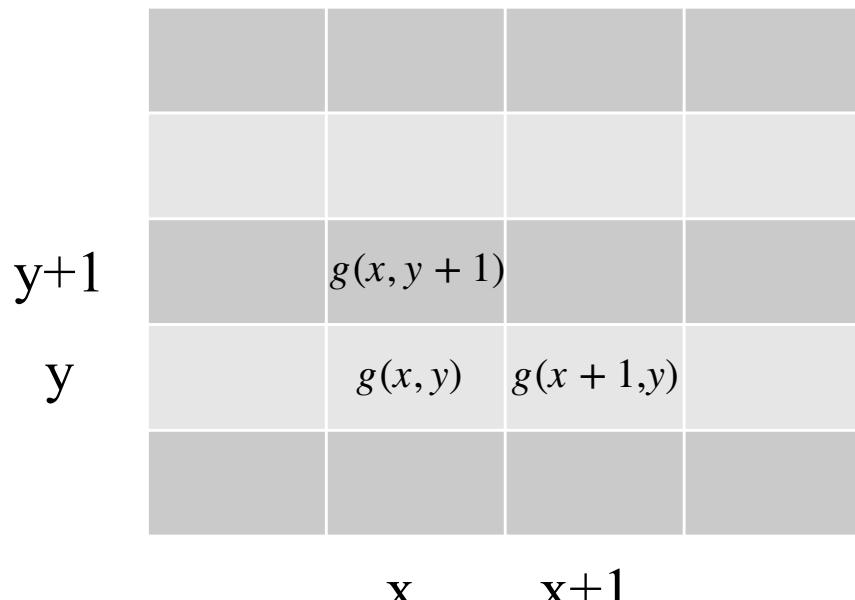
→ Solve two linear equations at each iteration.

Derivatives of the Image Gradient

$$E_G = -\frac{I}{N} \sum_{i=1}^N G(x_i, y_i)$$

$$\frac{\partial E_G}{\partial X} = \begin{bmatrix} \frac{\partial E_G}{\partial x_1} & \dots & \frac{\partial E_G}{\partial x_N} \end{bmatrix}, \quad \frac{\partial E_G}{\partial Y} = \begin{bmatrix} \frac{\partial E_G}{\partial y_1} & \dots & \frac{\partial E_G}{\partial y_N} \end{bmatrix}$$

$$\frac{\partial E_G}{\partial x_i} = -\frac{I}{N} \frac{\partial G}{\partial x_i}(x_i, y_i), \quad \frac{\partial E_G}{\partial y_i} = -\frac{I}{N} \frac{\partial G}{\partial y_i}(x_i, y_i)$$

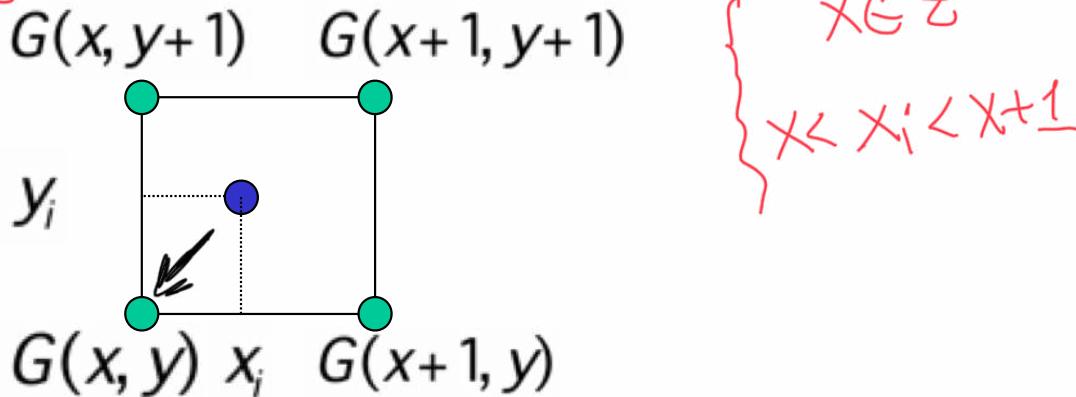


- We have values of g for integer values of x and y .
- But x_i and y_i are not integers.

—> We need to interpolate.

Bilinear Interpolation

You have gradients at 4 pts



$$0 < p, q \leq 1$$

The more closer it's
to one of these 4 pts,
the more influence it gets
from weight

$$p = x_i - x$$

$$q = y_i - y$$

Bilinear interpolation

$$G(x_i, y_i) = \underbrace{(1-p)(1-q)}_{\frac{\partial G}{\partial x_i}} G(x, y) + (1-p)q G(x, y+1) + p(1-q) G(x+1, y) + pq G(x+1, y+1)$$

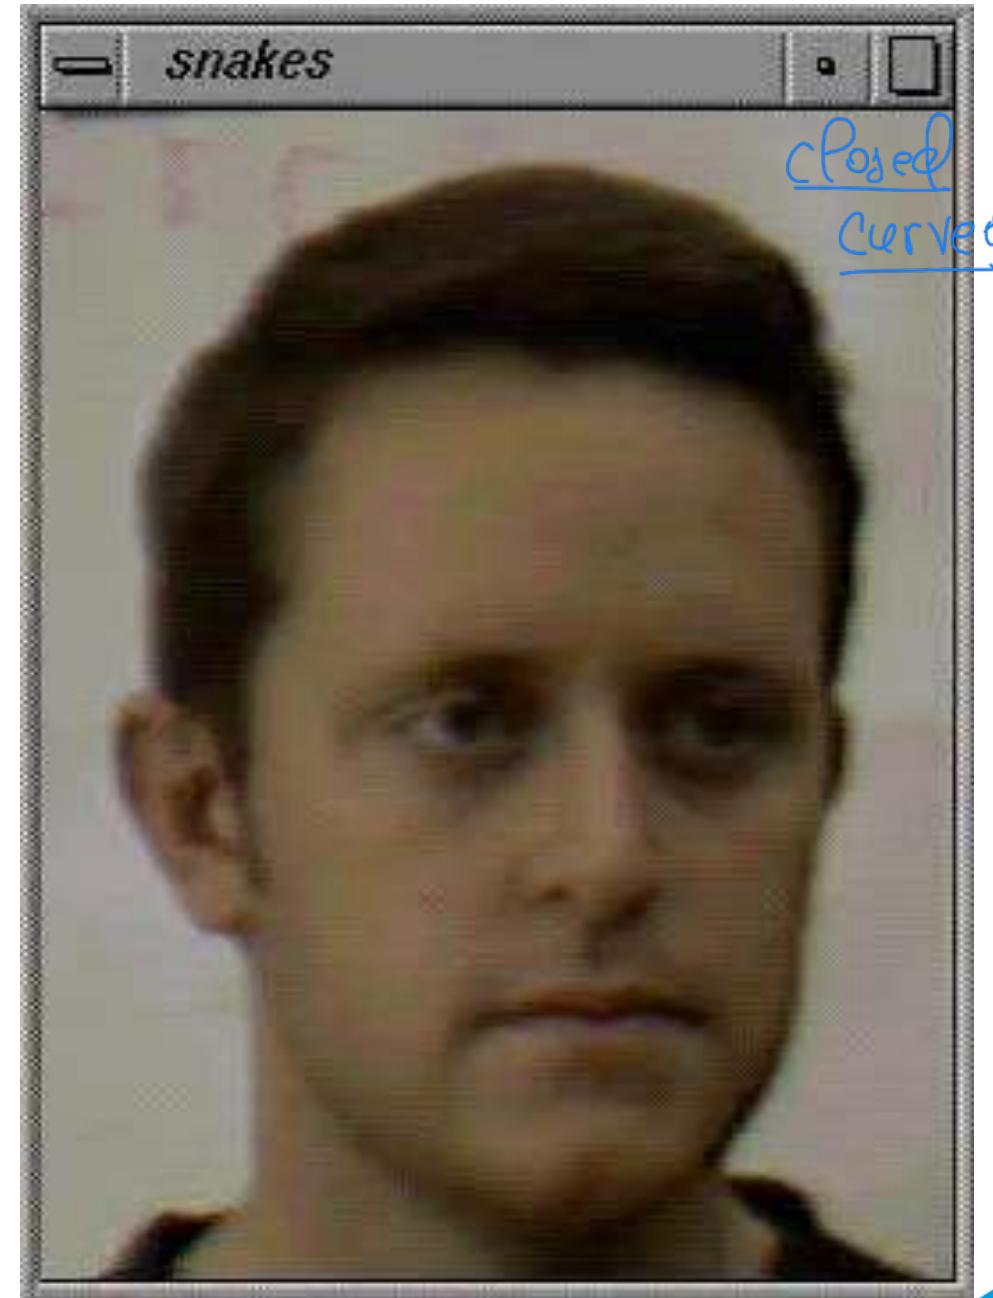
$$\frac{\partial G}{\partial x_i} = (1-q)(G(x+1, y) - G(x, y)) + q(G(x+1, y+1) - G(x, y+1))$$

$$\frac{\partial G}{\partial y_i} = (1-p)(G(x, y+1) - G(x, y)) + p(G(x+1, y+1) - G(x+1, y))$$

We only have values of G
for integers, we can estimate
for non-integers and compute derivatives

$$\frac{\underline{\partial G}}{\underline{\partial x_i}} = \frac{\underline{\partial G}}{\underline{\partial p}}$$

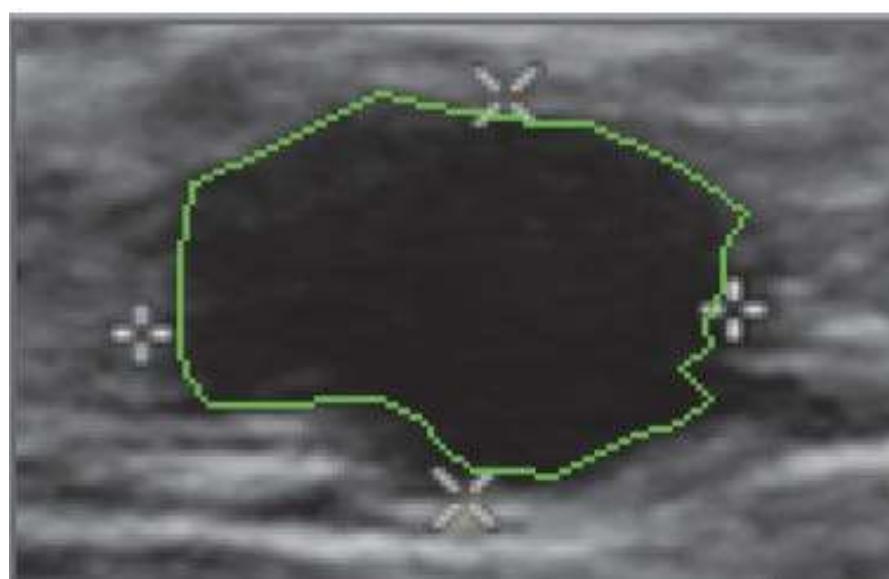
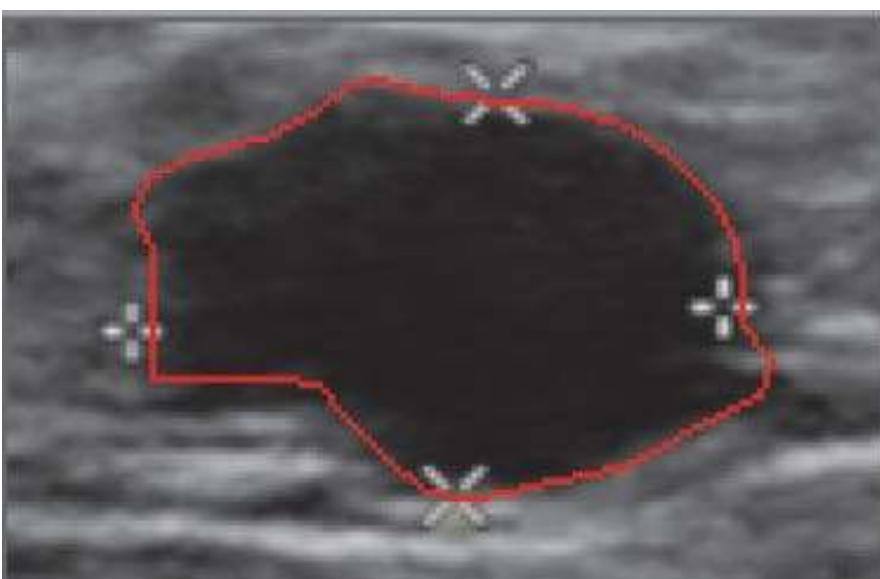
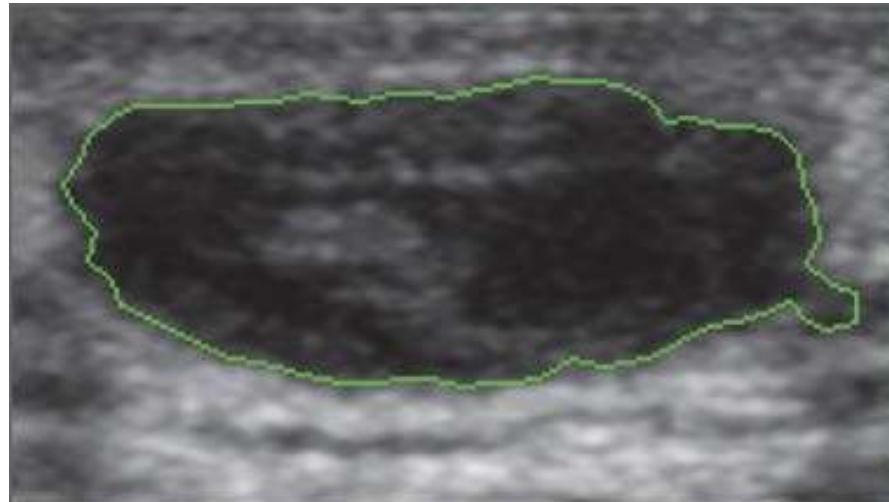
Open and Closed Snakes



Handle boundaries differently

Cysts Tumors in Ultrasound Images

↳ extremely noisy



Drawn by the physician.

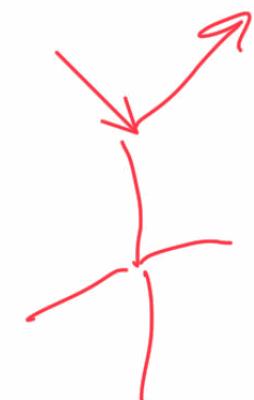
Refined by the Computer.

Network Snakes

Cadastral
maps



You have network
of curves → deformed
by themselves

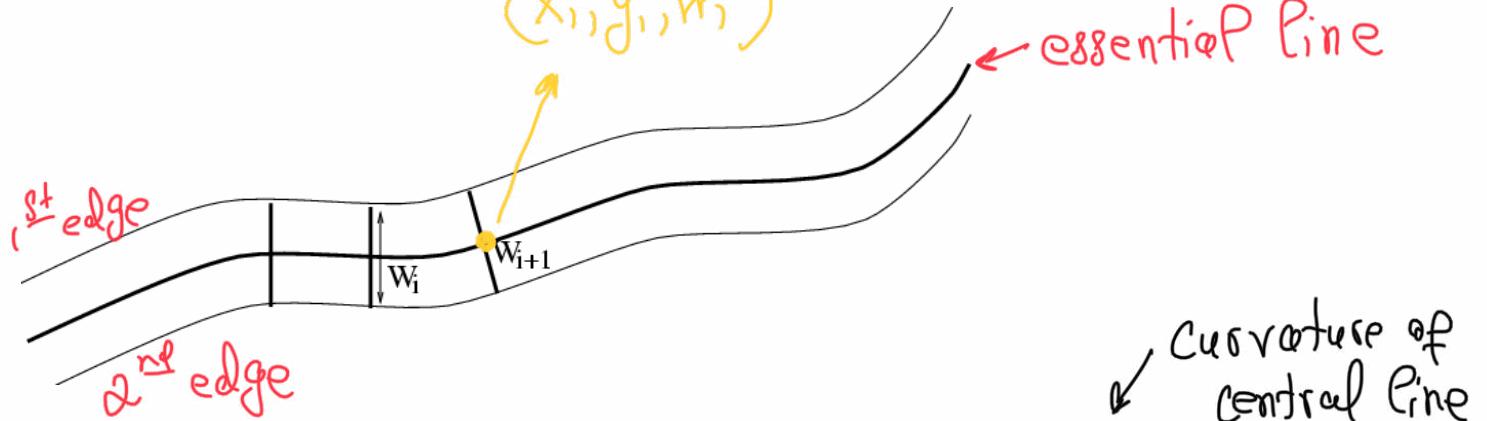


--> Updated field boundaries. → Who owns what?

Ribbon Snakes

How to represent road?

particular point
position of centre of line and width at that
 (x_i, y_i, w_i)



Curvature of
central Line

$$E = E_G + 1/2X^t K X + 1/2Y^t K Y + 1/2W^t K_W W$$

$$W = [w_1, \dots, w_N]^t$$

$$K_W = \begin{bmatrix} \cdot & \cdot \\ \cdot & \cdot & \cdot & -1 & 2 & -1 & \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & -1 & 2 & -1 & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & -1 & 2 & -1 & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & -1 & 2 & -1 & \cdot & \cdot \\ \cdot & -1 & 2 & -1 & \cdot \\ \cdot & -1 & 2 & -1 \\ \cdot & -1 & 2 \\ \cdot & -1 \end{bmatrix}$$

Dynamics Equations

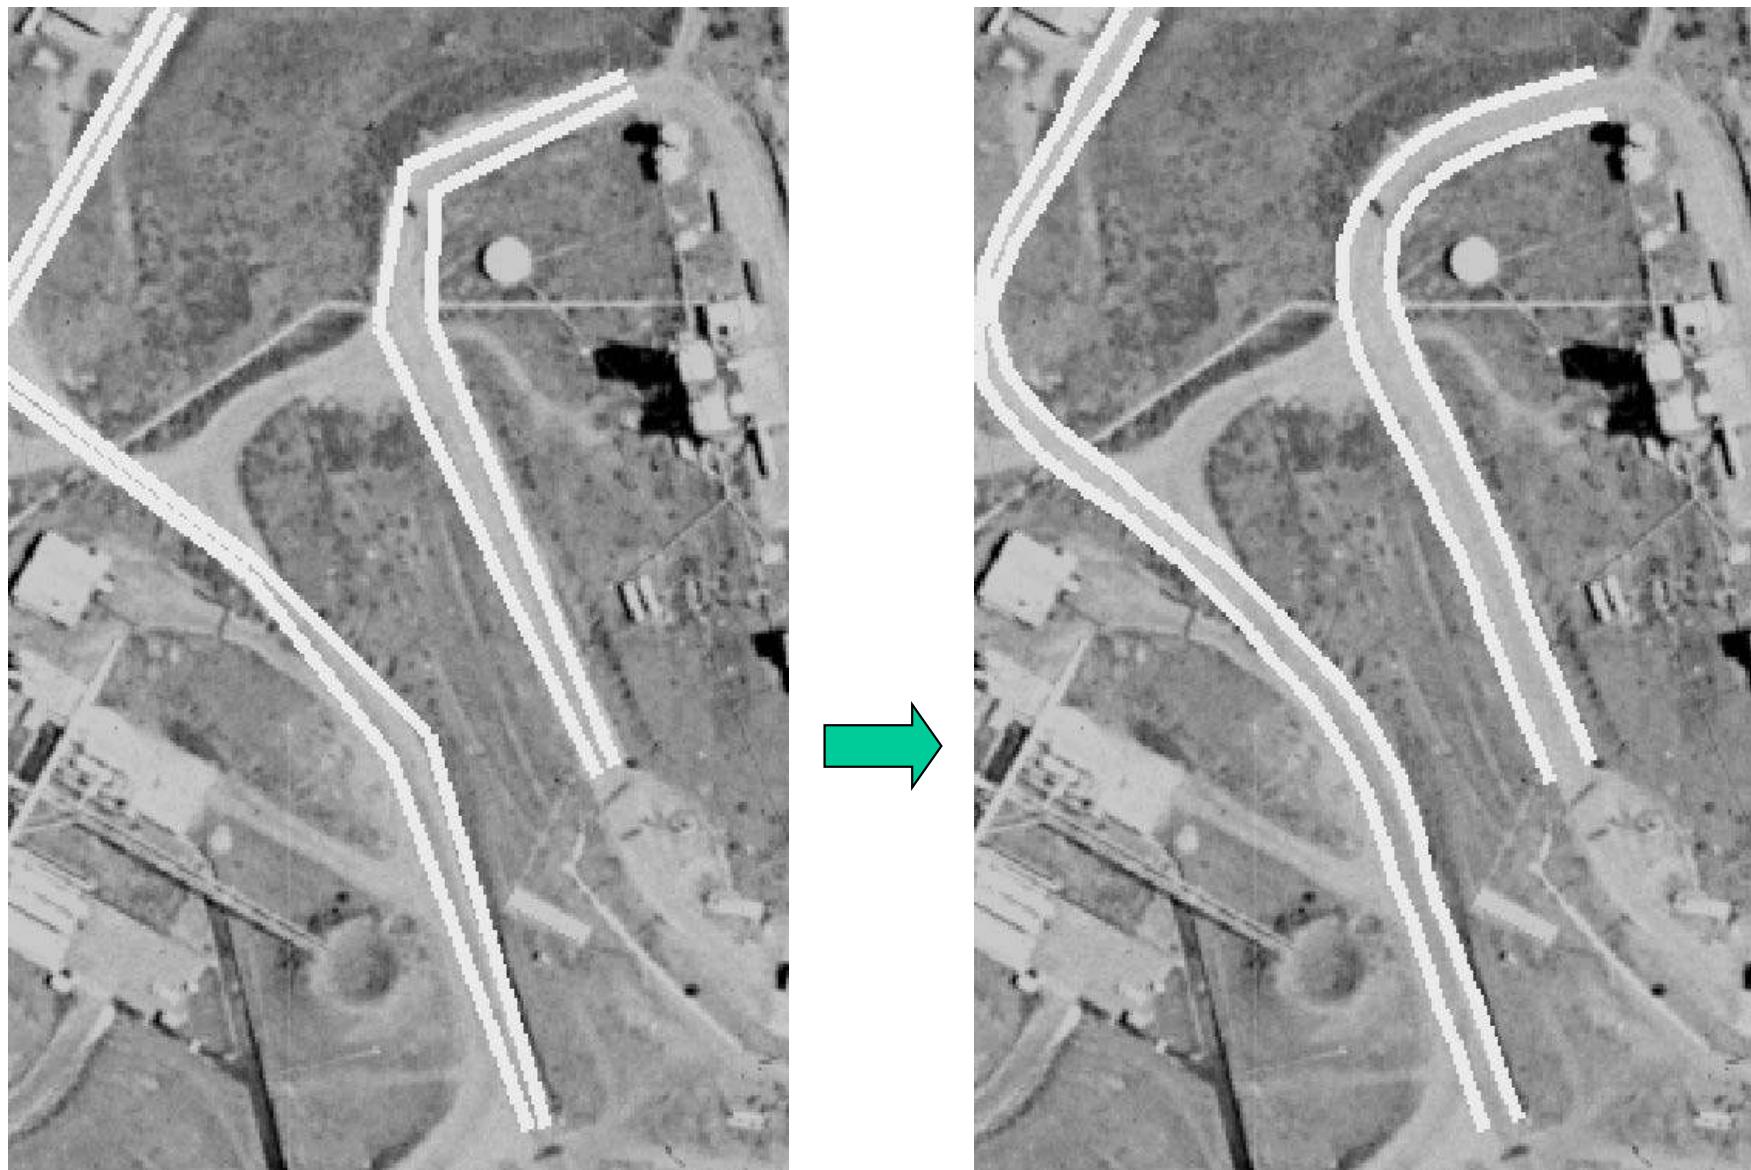
3 eq's
to solve
at each
iteration

$$\left\{ \begin{array}{l} (K + \alpha I)X_t = \alpha X_{t-1} - \frac{\delta E_G}{\delta X} \\ (K + \alpha I)Y_t = \alpha Y_{t-1} - \frac{\delta E_G}{\delta Y} \\ (K + \alpha I)W_t = \alpha W_{t-1} - \frac{\delta E_G}{\delta W} \end{array} \right.$$

~~\equiv~~

→ Solve three linear equations at each iteration.

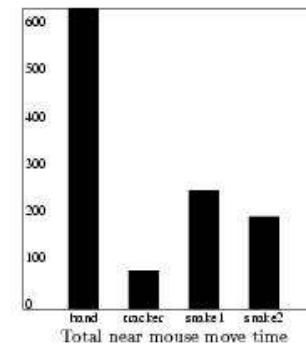
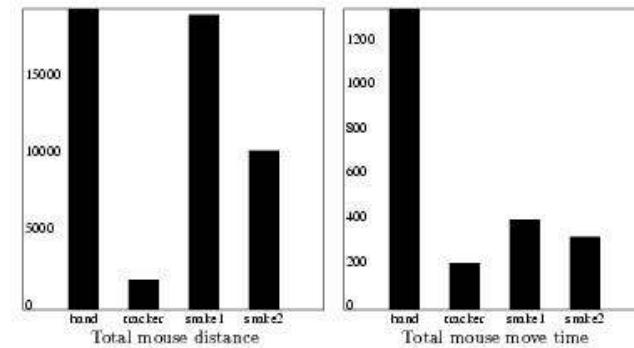
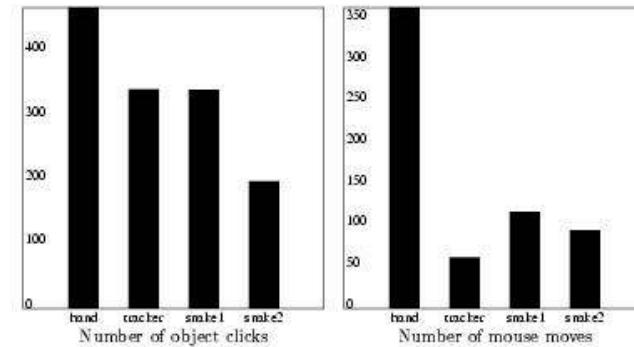
Delineating Roads



Delineating Roads

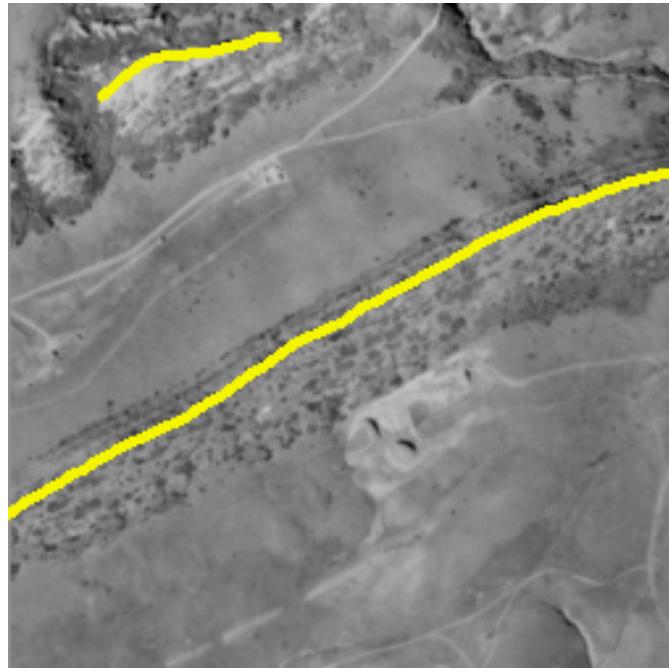


Evaluation



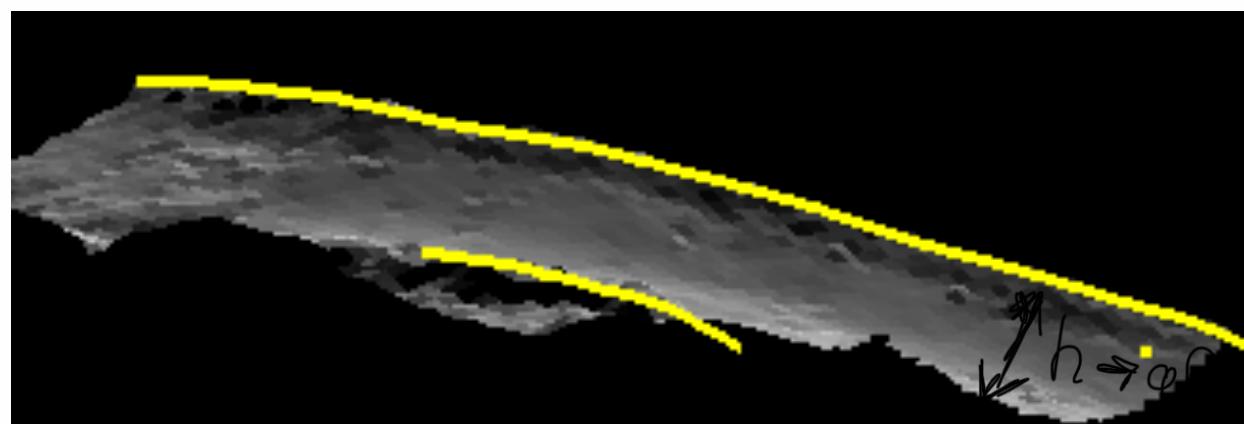
It takes far fewer clicks to trace the roads using semi-automated tools than doing entirely by hand.

Modeling a Ridge Line in 3D

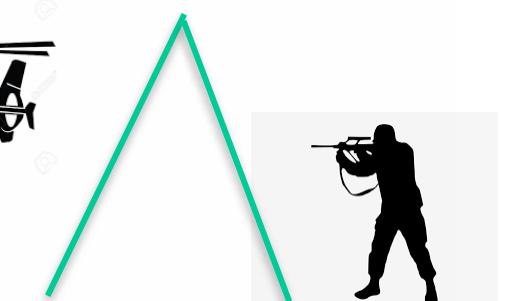


γ 2 views to model
(Stereo vision)

Three different views of the same landscape



Synthetic side view.



Modeling a Building in 3D

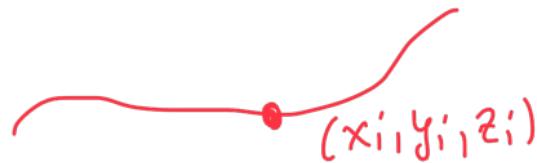


Sketch roughly the contours

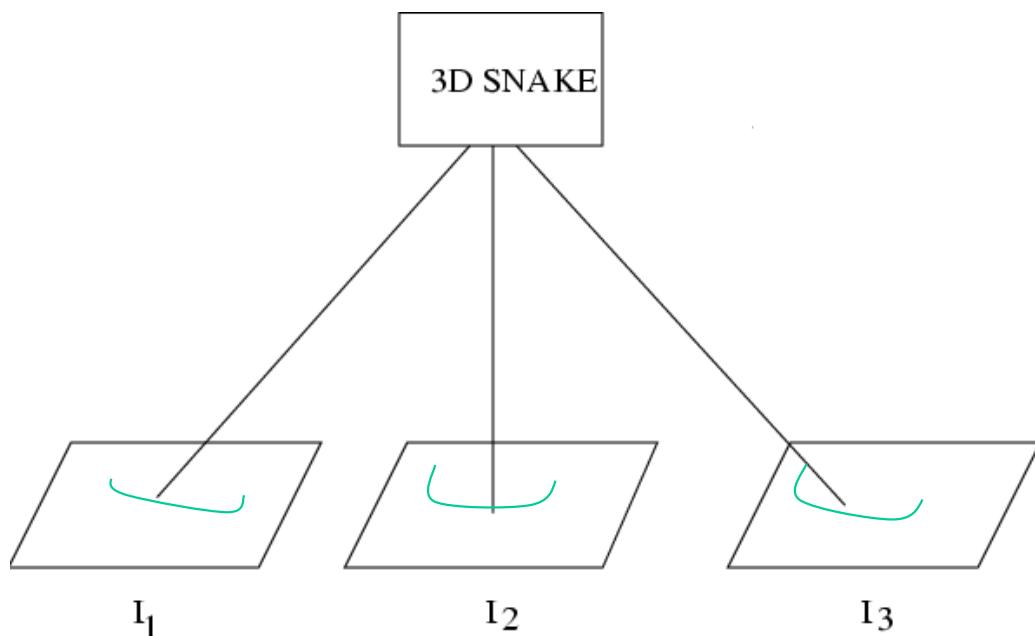
Ribbon
Snake optimizing it

⇒ finding exact
heights
of the building

3D Snakes

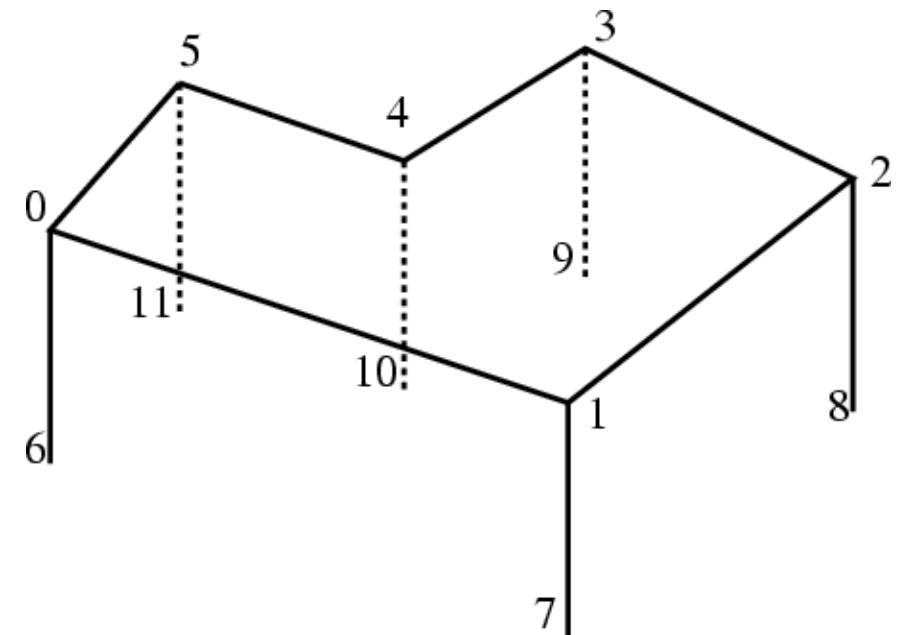


For rigid line



Smooth 3—D snake

Using the camera models, you project into 2D



Rectilinear 3—D snake

Dynamics Equations

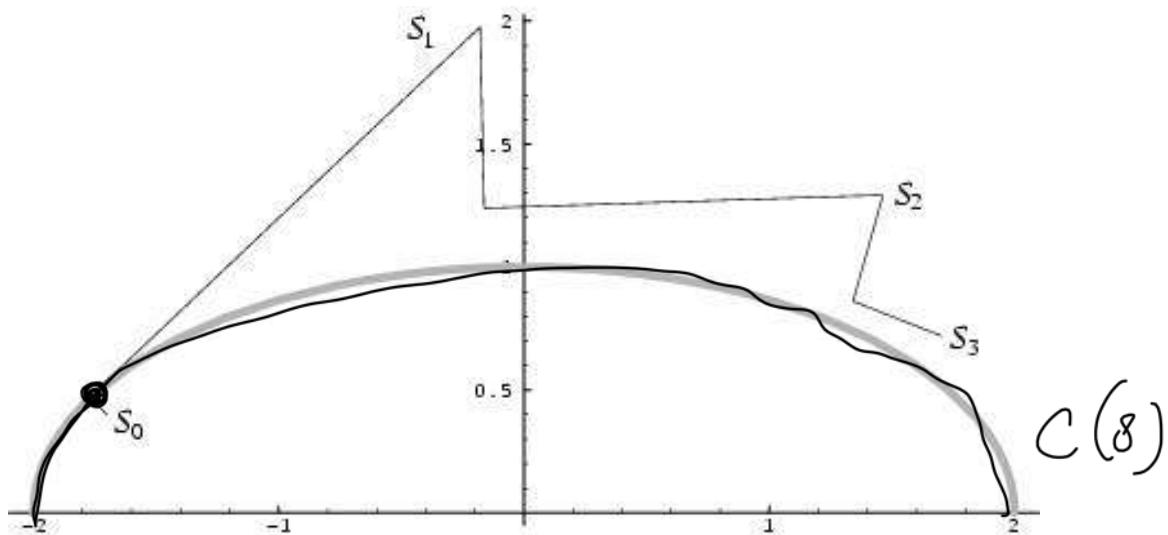
$$(K + \alpha I)X_t = \alpha X_{t-1} - \frac{\delta E_G}{\delta X}$$

$$(K + \alpha I)Y_t = \alpha Y_{t-1} - \frac{\delta E_G}{\delta Y}$$

$$(K + \alpha I)Z_t = \alpha Z_{t-1} - \frac{\delta E_G}{\delta Z}$$

→ Solve three linear equations at each iteration.

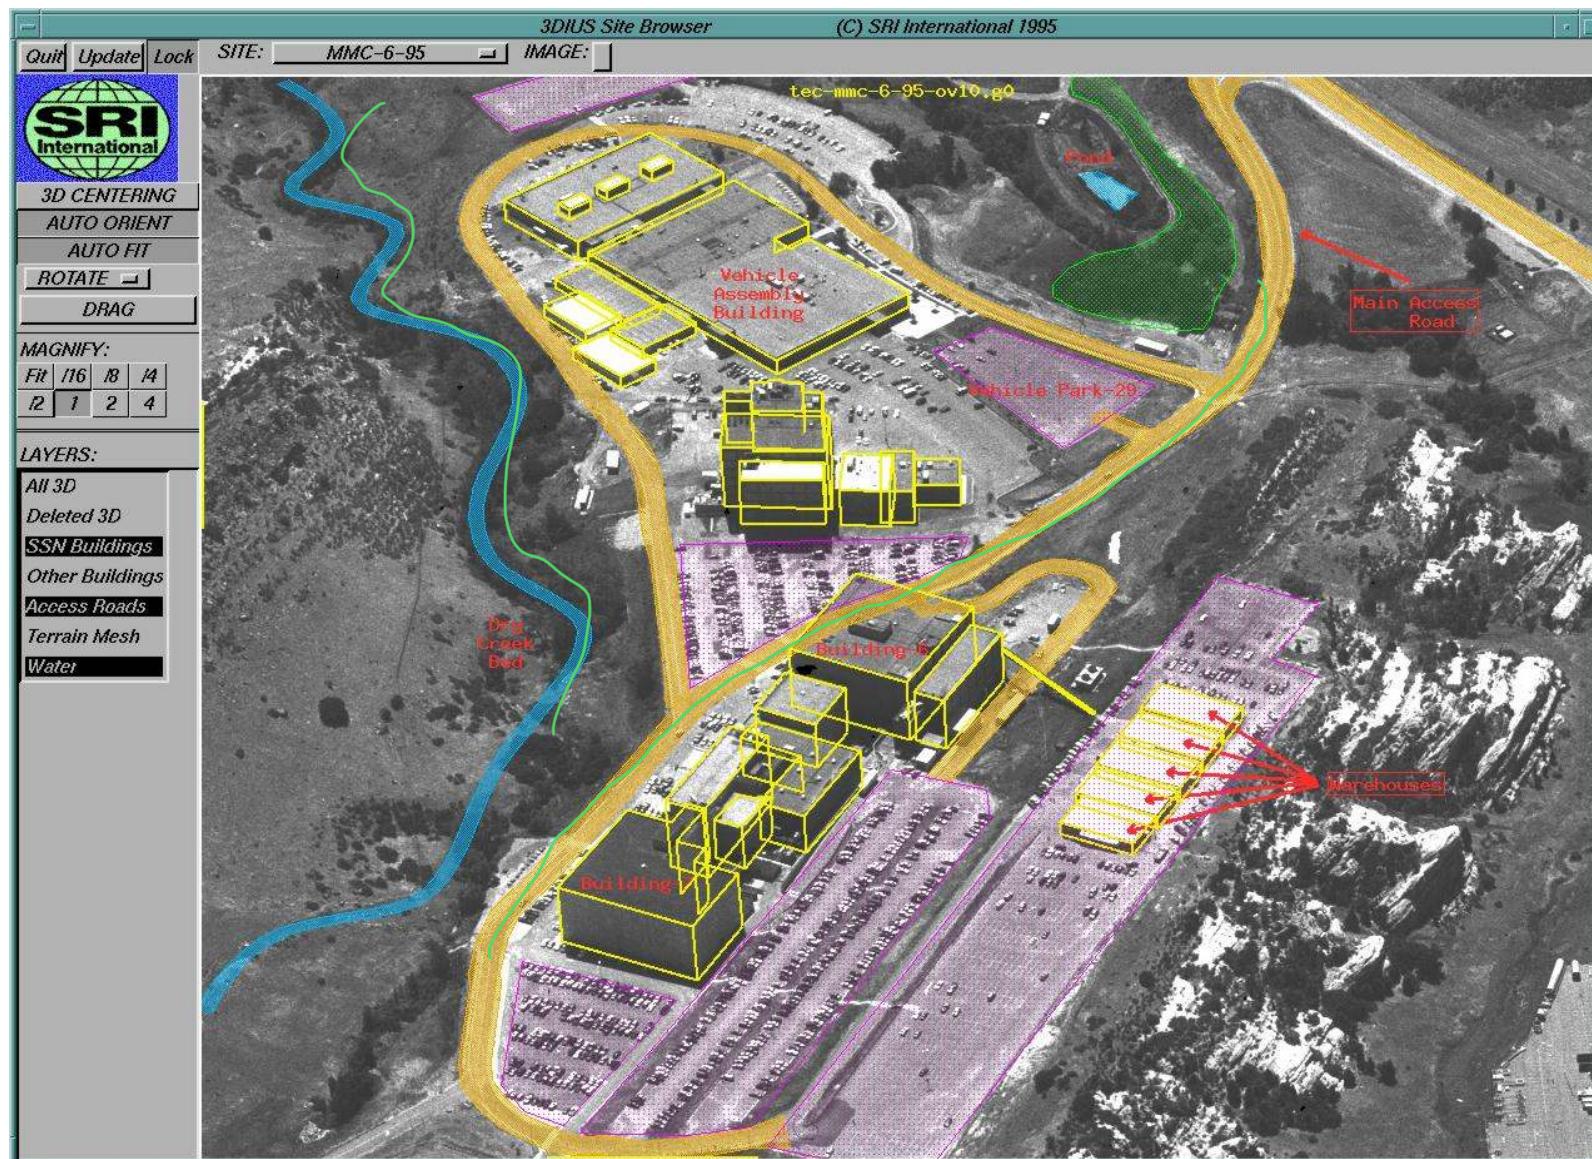
Constrained Optimization



- Minimize $F(S)$ subject to $C(S) = 0$

Site Modeling (1996)

Constraints => Roads on the terrain, roads flow downhill, and consistent



Site Modeling (2019)

Reconstructed images/shapes



Level Sets



Implicit vs Explicit

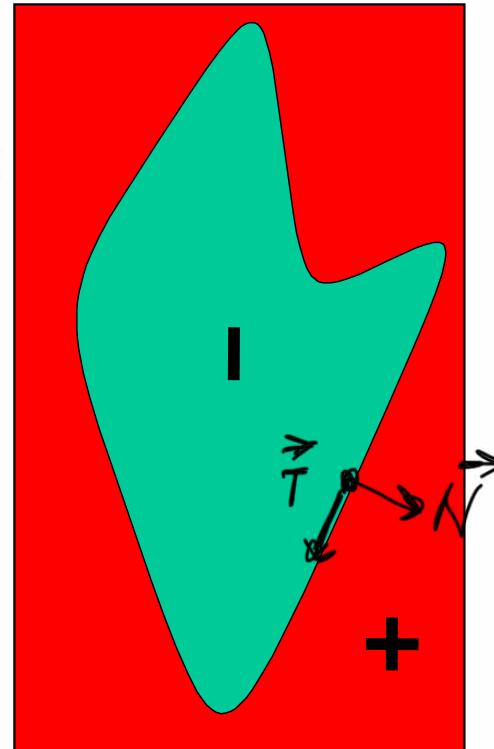
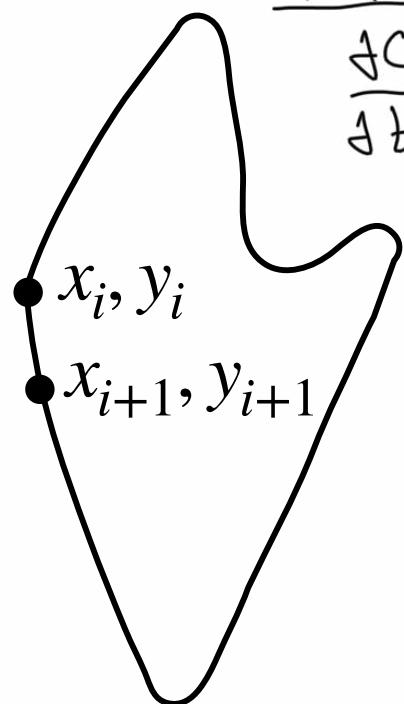
Generic formulation:

$$\frac{dC}{dt} = \alpha(s, t) \vec{T} + \beta(s, t) \vec{N}$$

Reparameterization:

$$\frac{dC}{dt} = \beta(s, t) \vec{N}$$

displacement along the normal



Special cases:

$\beta = \pm 1$ ← curve will grow/shrink

$$\beta = \beta_0 - \beta_1 K$$

curvature

$$z = \Phi(x, y),$$

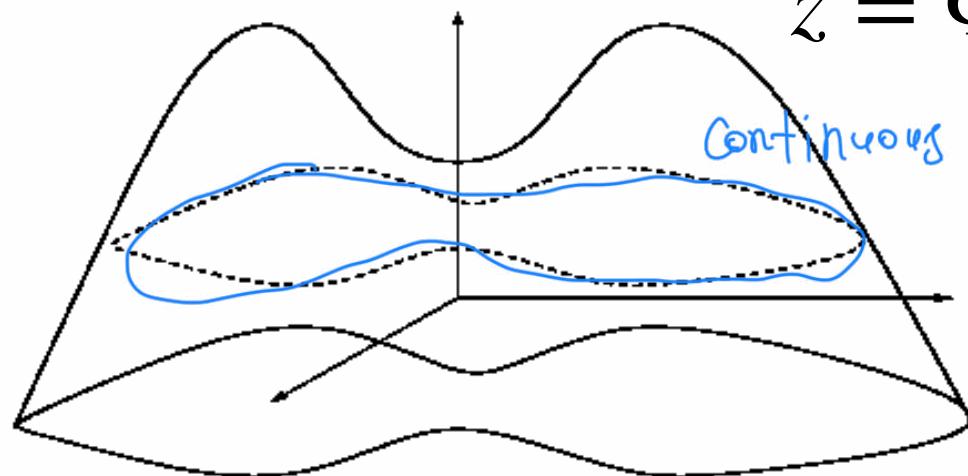
$z > 0$ outside,

$z < 0$ inside,

→ Consider the curve as the zero level set of a surface.

Topology Changes are Possible

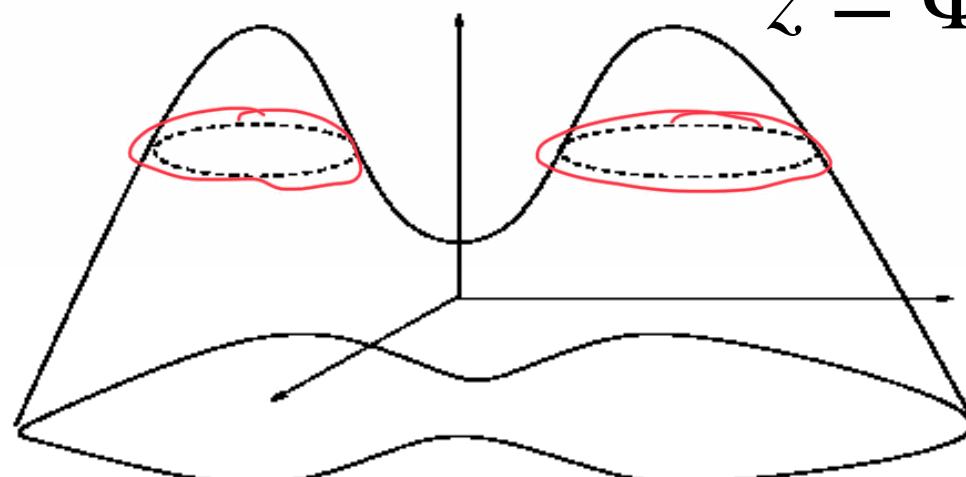
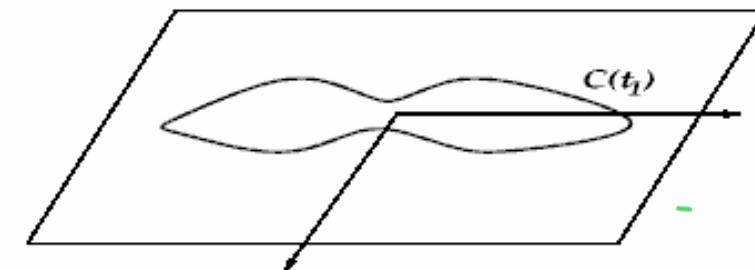
Surface is where $\underline{\Phi}(x, y) = 0$



$$z = \Phi(x, y, t_1)$$

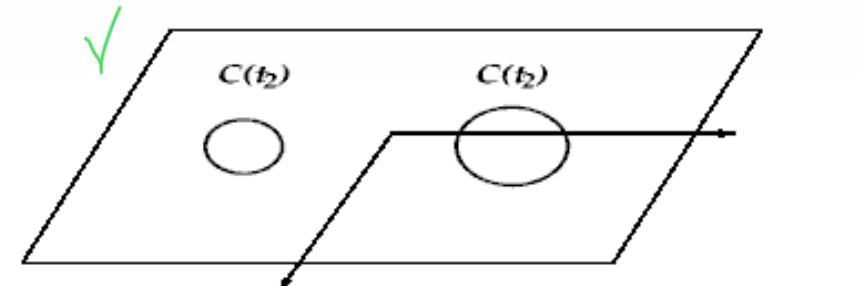
Continuous Surface

Both of them
are zero-crossings of the
✓ implicit field



$$z = \Phi(x, y, t_2)$$

When $\underline{\Phi}$ is deformed
(with some field \underline{F})



Curve Evolution

Level Sets



$$\underline{\Phi}(x, y) = 0$$



Consider the curve as the zero level set of the surface:

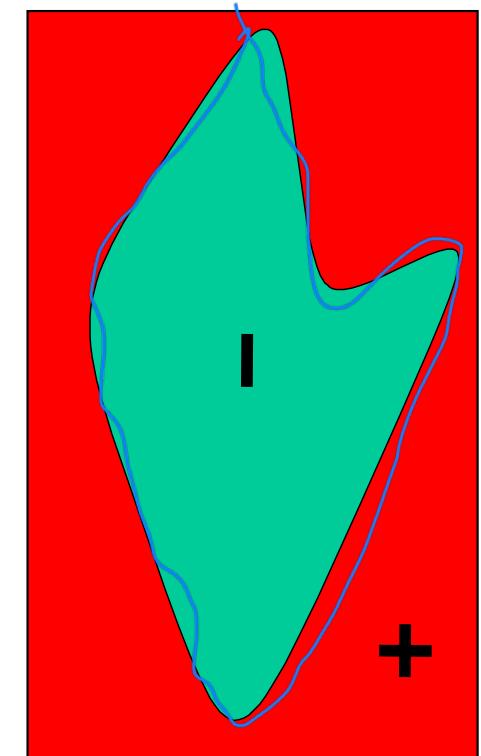
$$z = \Phi(x, y, t)$$

Evolution equation:

$$0 = \Phi_t + \beta(\kappa) |\nabla \Phi|$$

$$\text{where } \kappa = \frac{\Phi_{xx}\Phi_y^2 - 2\Phi_{xy}\Phi_x\Phi_y + \Phi_{yy}\Phi_x^2}{\Phi_x^2 + \Phi_y^2}$$

↑
curvature



$\beta(\kappa)$ is the speed at which the surface deforms.

Level Set Smoothing

Smoothing occurs when $\beta(\kappa) = -\kappa$

Desirable properties:

- Converges towards circles.
- Total curvature decreases.
- Number of curvature extrema and zeros of curvature decreases.

Relationship with Gaussian smoothing:

- Analogous to Gaussian smoothing of boundary over the short run, but does not cause self-intersections or overemphasize elongated parts.
- Can be implemented by Gaussian smoothing the characteristic function of a region.

Shape Recovery

If you want snake-like behaviour

Evolution equation: $0 = \Phi_t + \underbrace{\beta(\kappa)}_{\approx} |\nabla \Phi|$

where: $\beta(\kappa) = k_I(1 - \epsilon\kappa)$

If there are a lot of

gradients $\Rightarrow k_I$ -small \Rightarrow
 $\beta(\kappa)$ -small \Rightarrow curve will
almost not move (Φ_t)

$$k_I = \frac{1}{1 + \nabla I}$$

gradient of the
image

no gradients \Rightarrow curve move
more

→ Expansion stops at the boundaries.

Level Sets

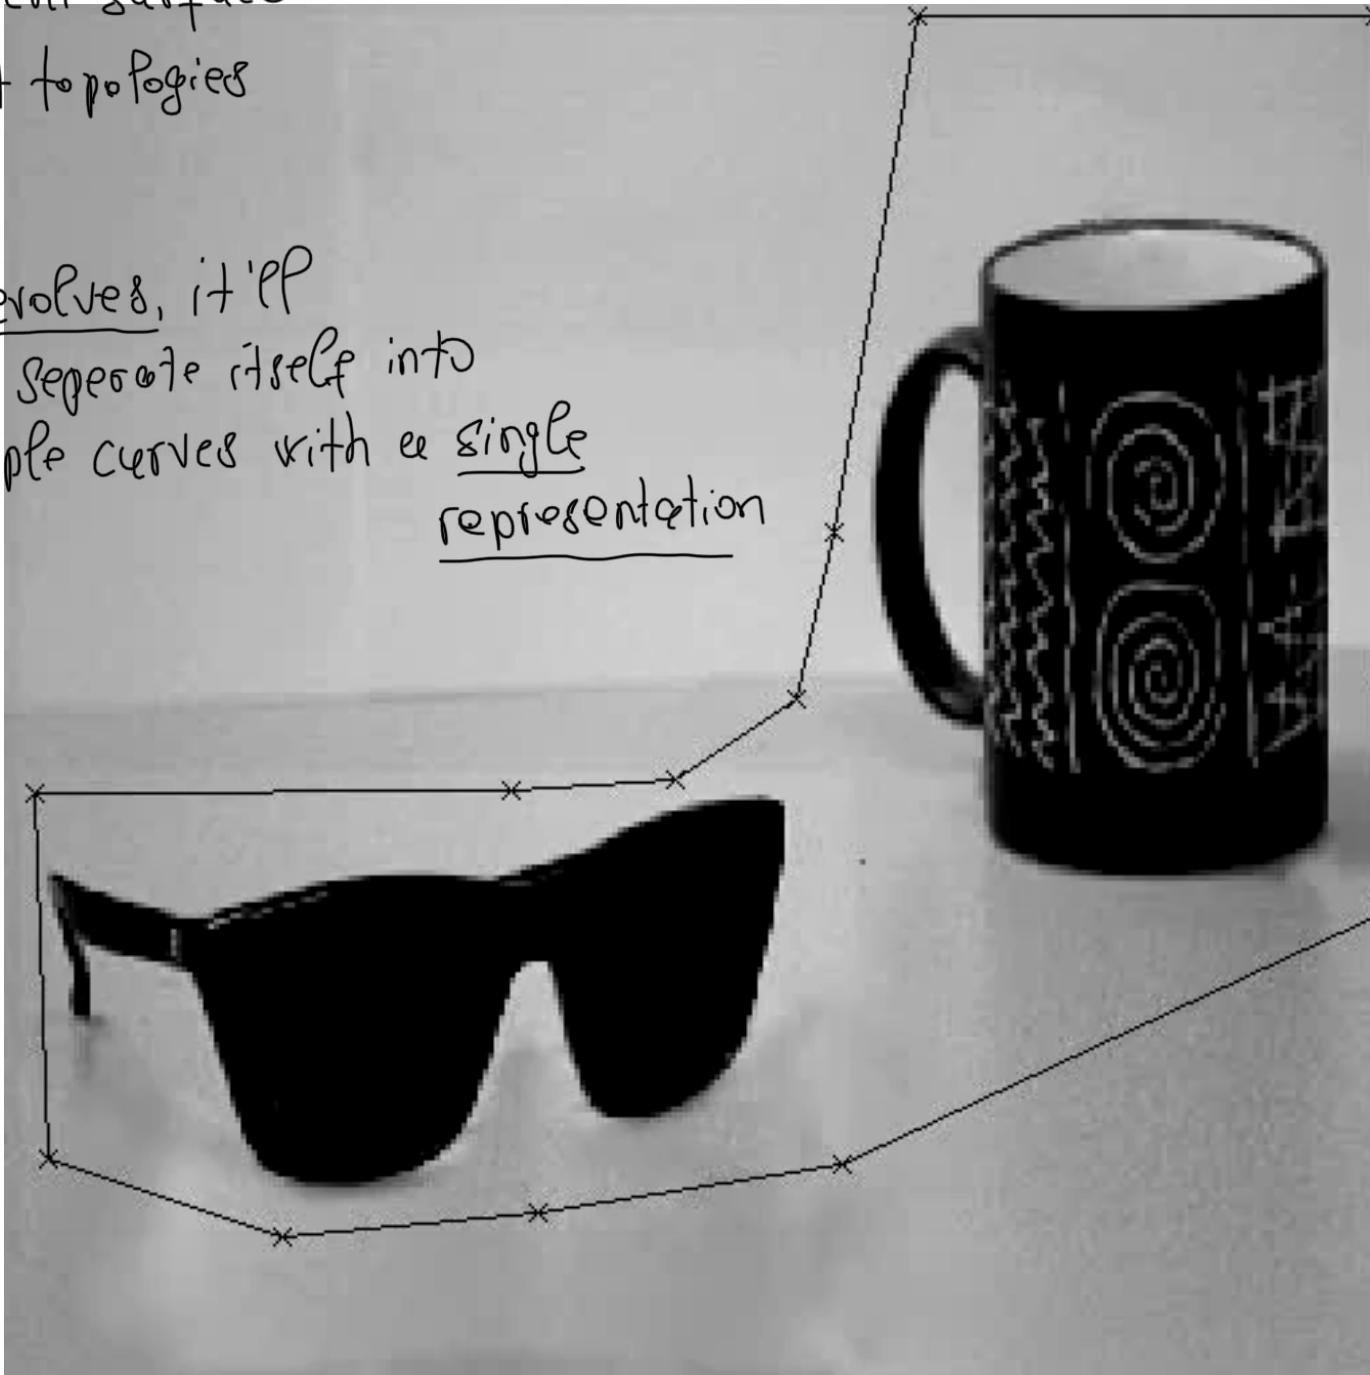


Stops by the
boundaries

with implicit surface represent,
you can represent surface
with different topologies

Level Sets

As the field evolves, it'll
separate itself into
multiple curves with a single
representation



Techniques

Semi-Automated Techniques:

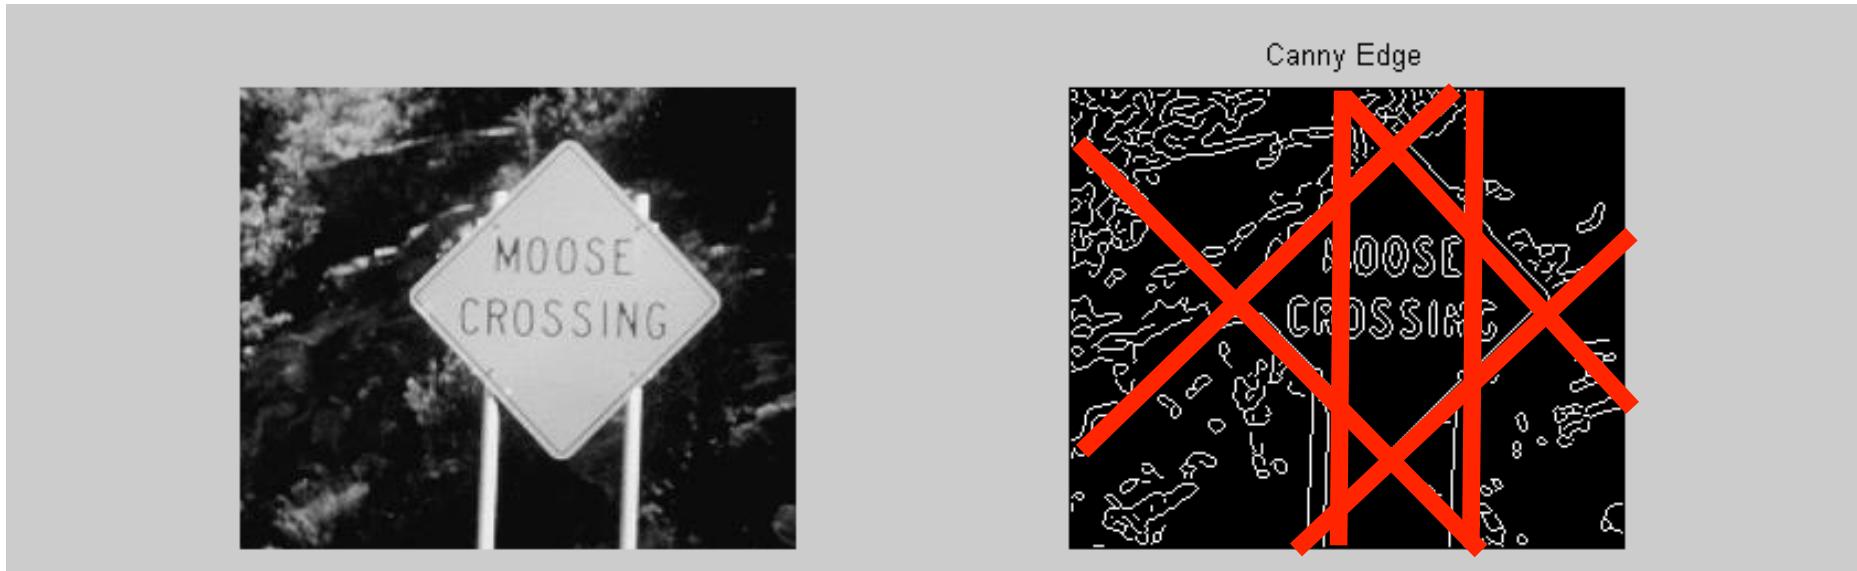
- Dynamic programming
- Deformable Models

Fully Automated Techniques:

- Hough transform
- Graph Based Approaches

User not intervened

Finding Lines



Input:

- Canny edge points.
- Gradient magnitude and orientation.

Output: *Find*

- All straight lines in image.

Hough Transform

Simple

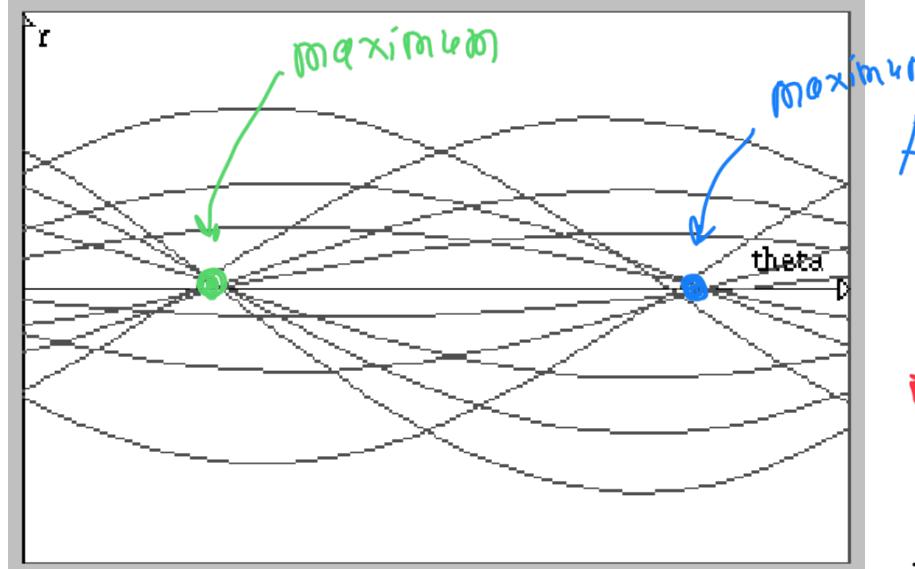
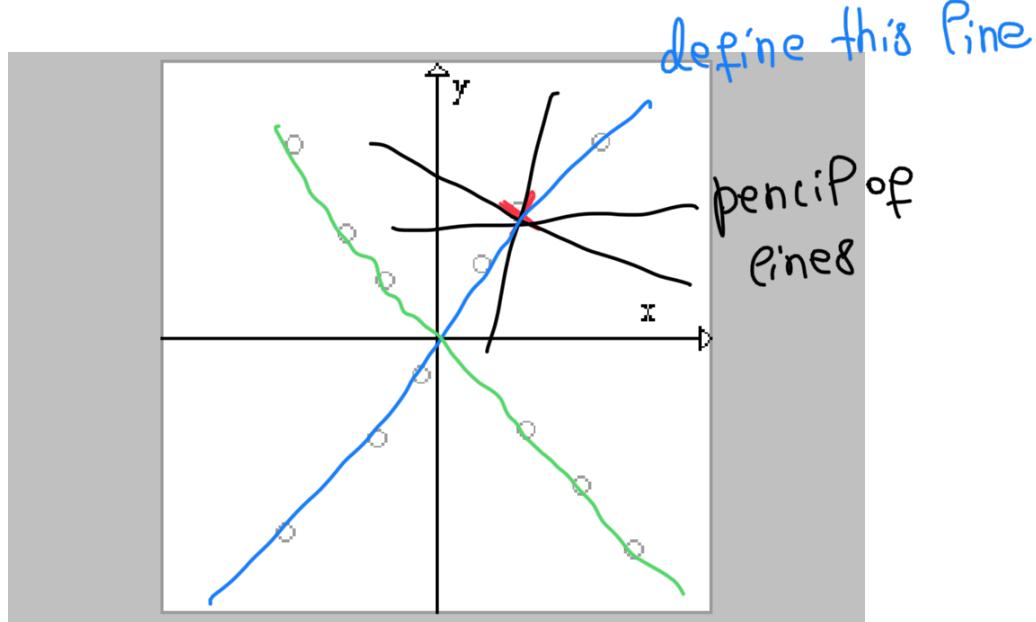
Given a parametric model of a curve:

(Line defined by
2 parameters)

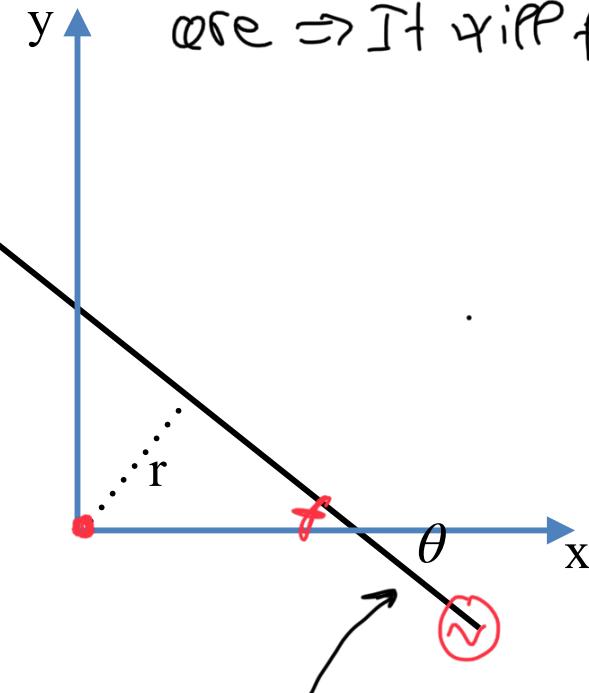
Voting scheme

- Map each contour point onto the set of parameter values for which the curves passes through it.
voting for curves
- Find the intersection for all parameter sets thus mapped.
Accumulate votes

Voting Scheme



Strength: You don't need to specify how many lines there are \Rightarrow It will find by itself

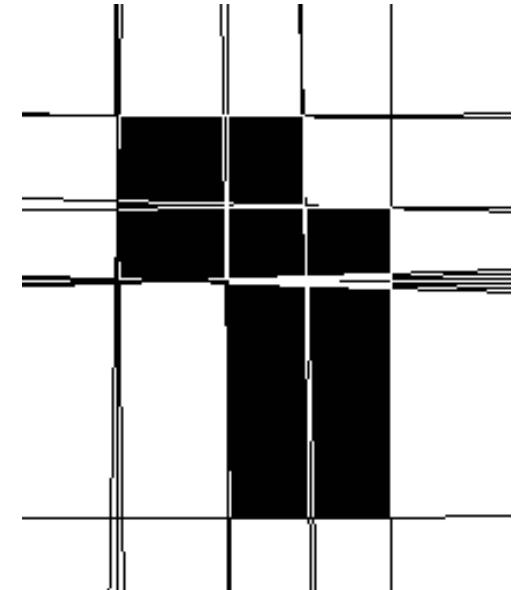
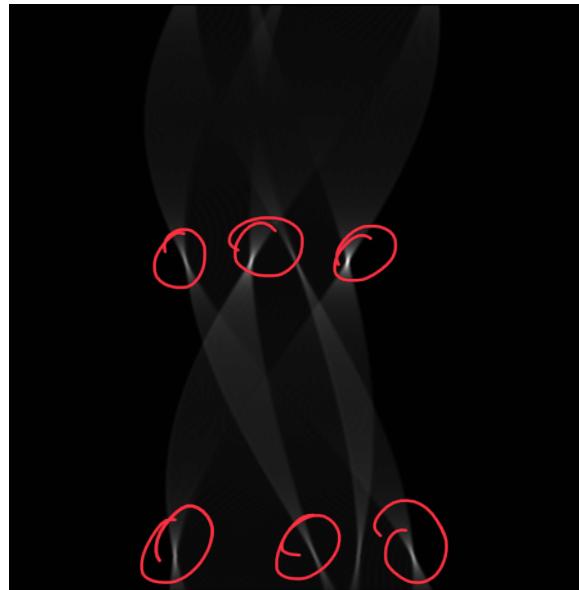
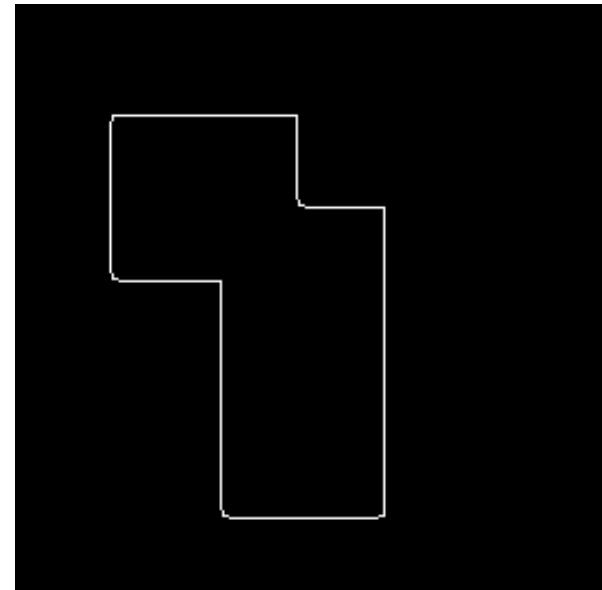


(x, y) points are on the pts that lie on that line parameterized by r, θ

Reverse: Given (x, y) with pencil of lines, draw figure on the left to find on (r, θ)

$$x \cos(\theta) + y \sin(\theta) = r, \quad 0 \leq \theta \leq \pi$$

Synthetic Lines



Image

Contours
from Canny

Accumulator

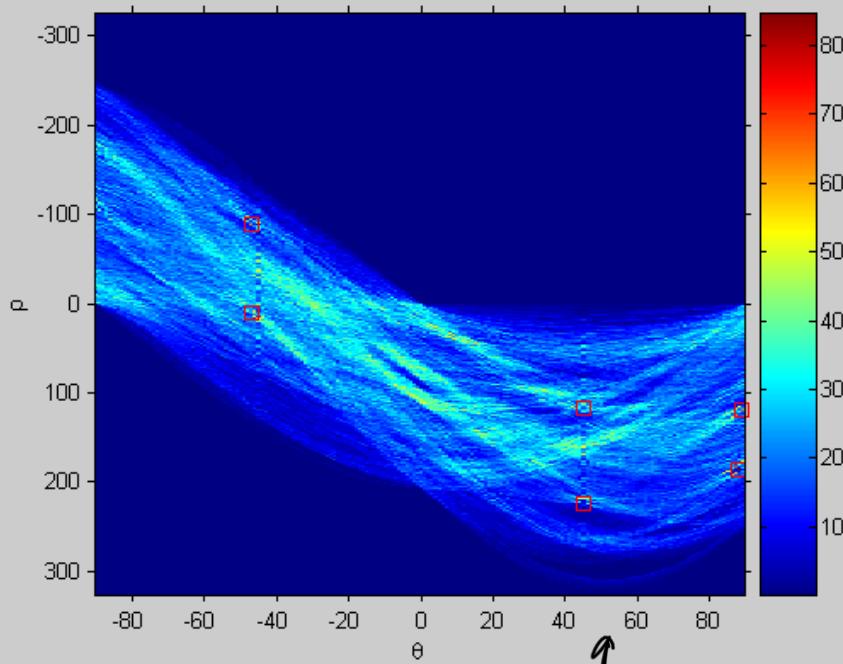
Lines

array

Once the contour points are associated to individual lines, you can perform least squares fitting.

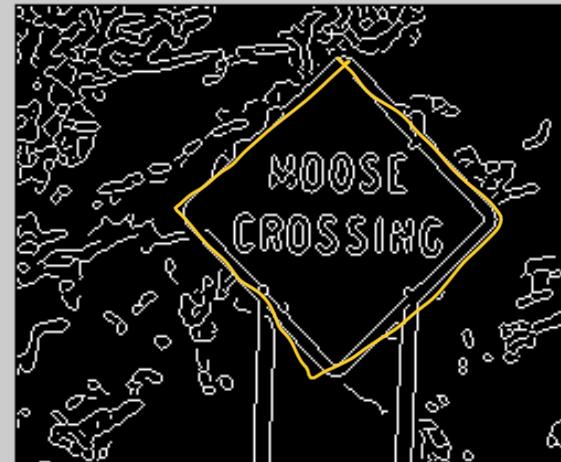
Real Lines

Line Segments



Votes in the accumulator array

Canny Edge



A maximum associates with a straight line (infinite length)

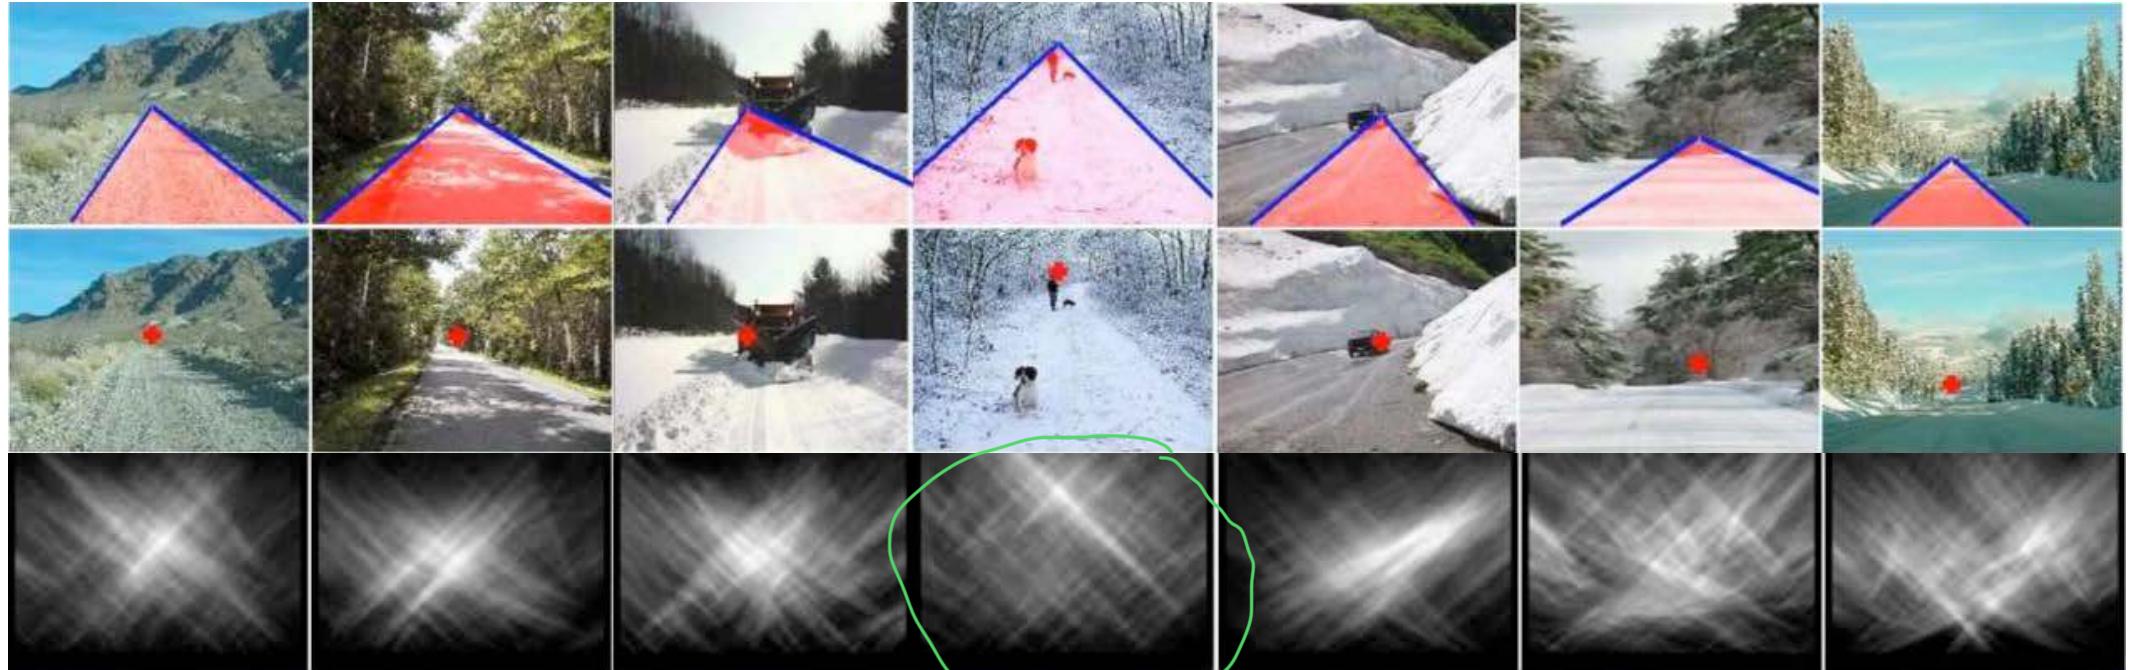
(they vote for all possible lines)
that go through themselves

Road Lines



Road Edges

\Rightarrow you will poor for 2 maxima



↑
contours
vote &

\Rightarrow reveal edges
of the road

accumulator
array on the scene

In projection, they meet
at vanishing pt

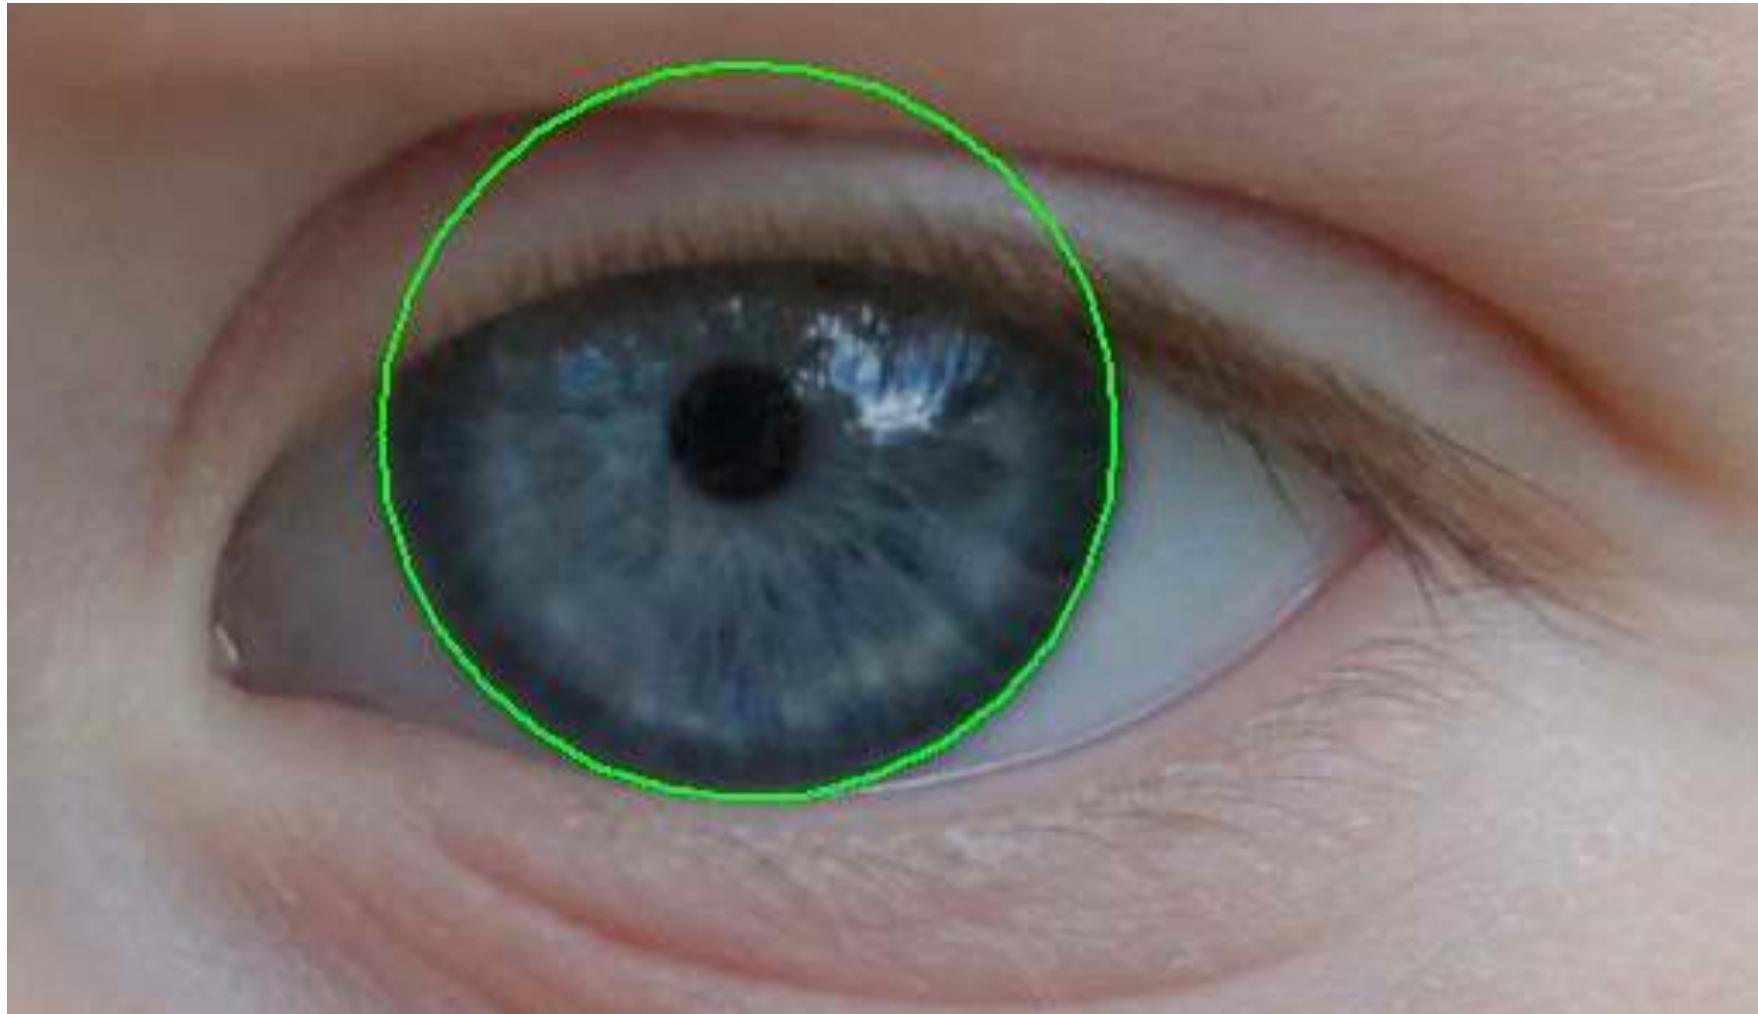
→ We are looking for the peak
which corresponds to (r, θ)

Generic Algorithm

- Quantize parameter space with 1 dimension per parameter.
- Form an accumulator array. *(grid, some dimension → 1 per parameter)*
- For each point in the gradient image such that the gradient strength exceeds a threshold, increment appropriate element of the accumulator.
(Canny image)
- Find local maxima in the accumulator.

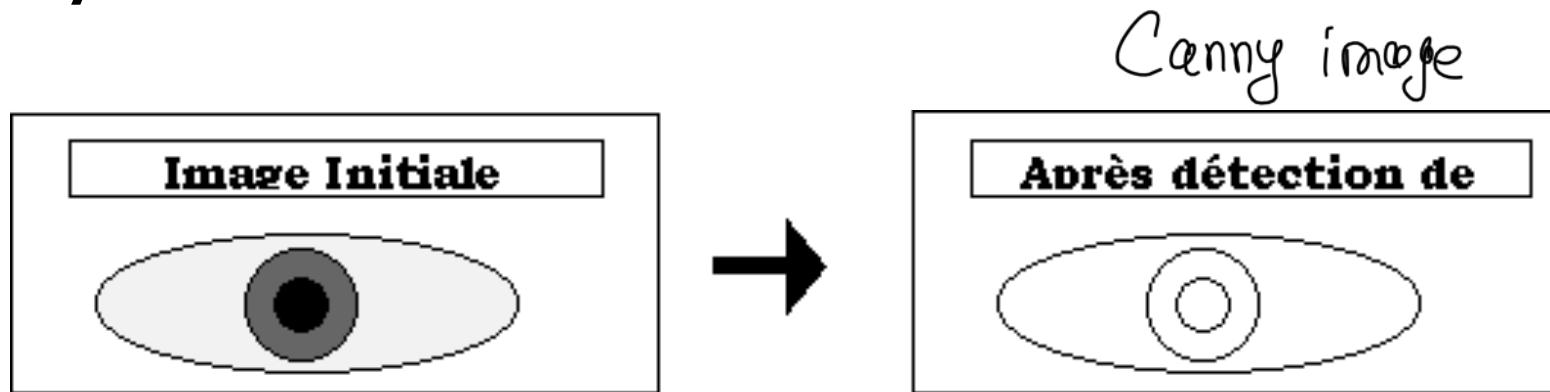
they vote for the
curve they can be on

Iris Detection



Occlusions

In theory:



In practice:



Circle Detection

$\forall \theta$ with $0 < \theta < 2\pi$

Circle of equation:

$$x = x_0 + r \cos(\theta)$$

$$y = y_0 + r \sin(\theta)$$

Therefore:

If you know point (x, y) (belonging to an edge)

center of circle

$$x_0 = x - r \cos(\theta)$$

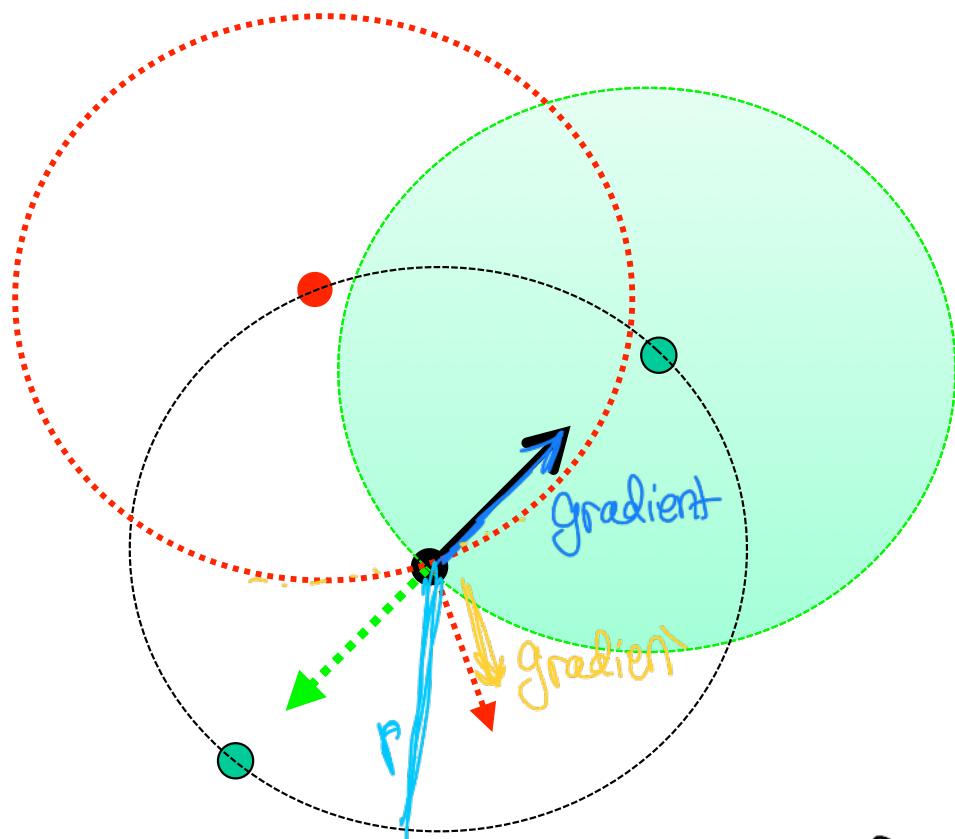
$$y_0 = y - r \sin(\theta)$$

for fixed r

for different r 's, you get
different centers

$\Rightarrow (x, y)$ votes on the circle as well

Gradient Orientation



In practice, we don't know where circle is \Rightarrow only gradient

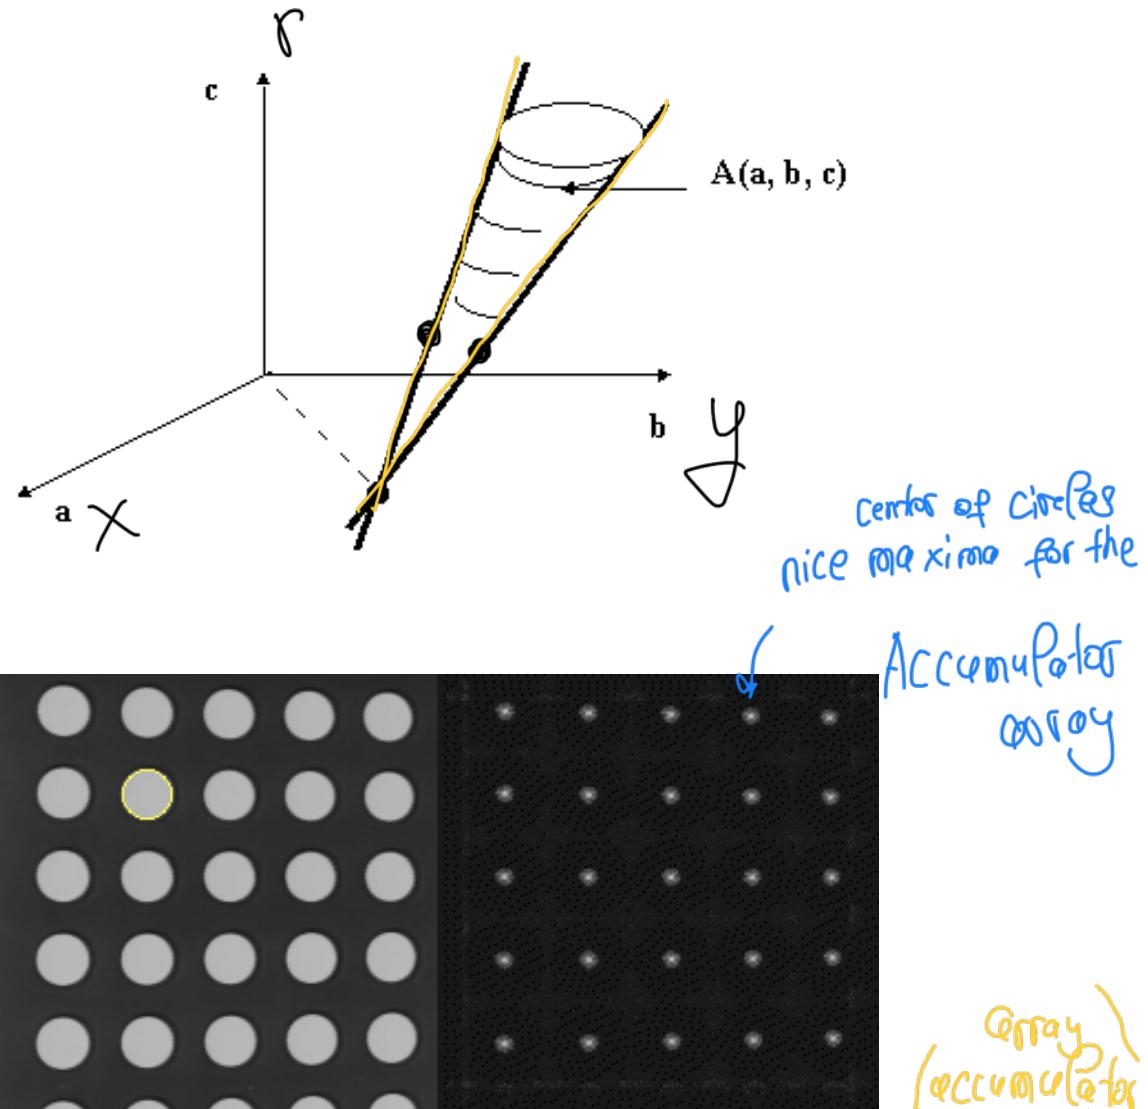
black-dotted

Can vote either along the entire circle or only at two points per value of the radius.

Because you take into consideration the gradient direction 69

Simple Image

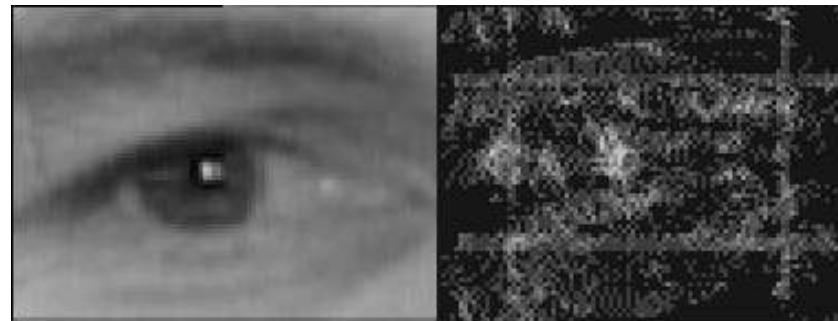
Voting scheme:



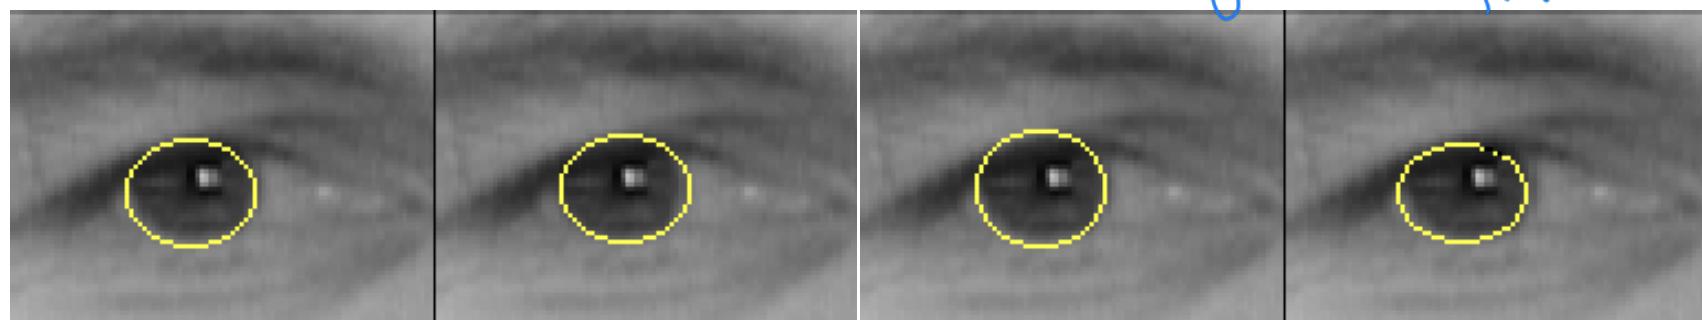
Result:

Eye Image

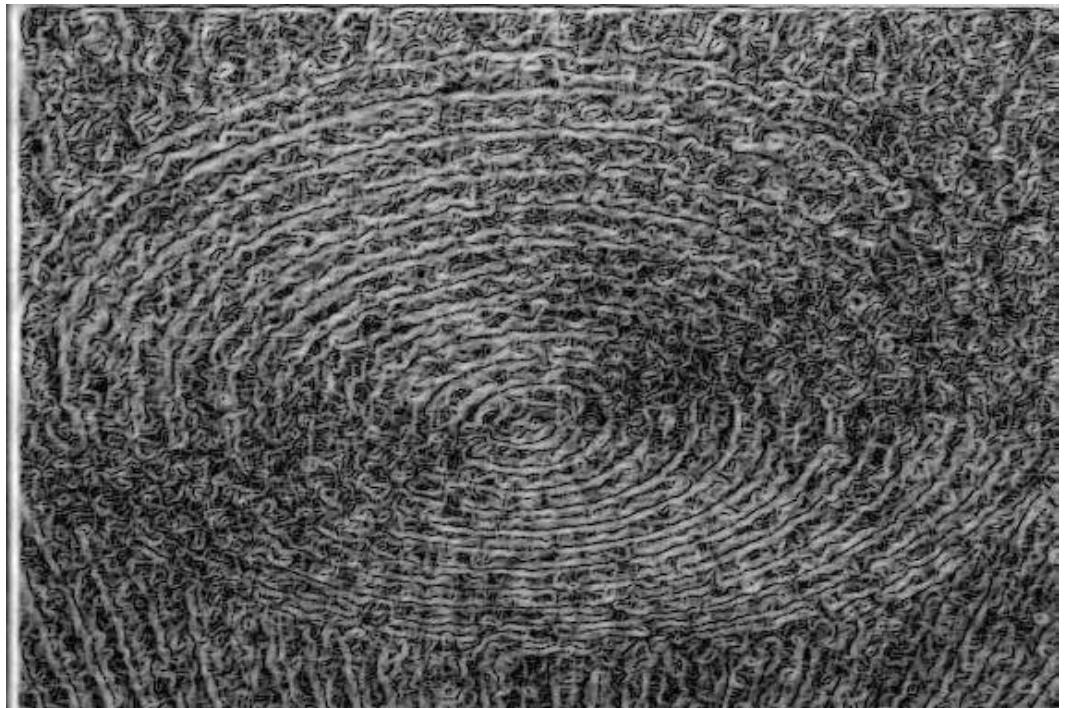
Image and accumulator:



Best four candidates:



Ellipses



Gradient of the image → Ellipses

Delineation does not work well

Ellipse Detection

Ellipse of equation:

$$x = x_0 + a \cos(\theta)$$

$$y = y_0 + b \sin(\theta)$$

Therefore:

$$x_0 = x - a \cos(\theta)$$

$$y_0 = y - b \sin(\theta)$$

2D-accumulator array

If you have a point (x, y) , for a given (a, b)
⇒ it will vote for
the center (x_0, y_0)

Gradient Orientation

For each ellipse point:

$$\frac{dx}{d\theta} = -a \sin(\theta)$$

$$\frac{dy}{d\theta} = b \cos(\theta)$$

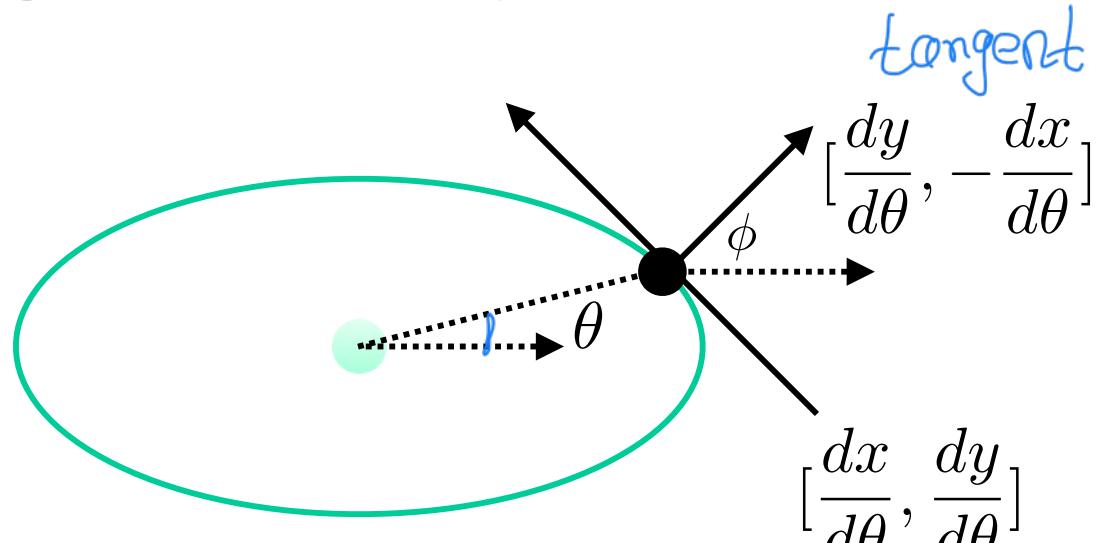
$$\phi = \text{atan}\left(-\frac{dx}{d\theta}, \frac{dy}{d\theta}\right)$$

$$= \text{atan}(a \sin(\theta), b \cos(\theta))$$

$$= \text{atan}\left(\frac{a}{b} \tan(\theta)\right)$$

$$\tan(\phi) = \frac{a}{b} \tan(\theta)$$

$$\Rightarrow \theta = \text{atan}\left(\frac{b}{a} \tan(\phi)\right)$$

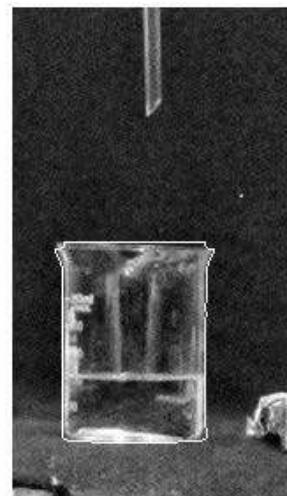
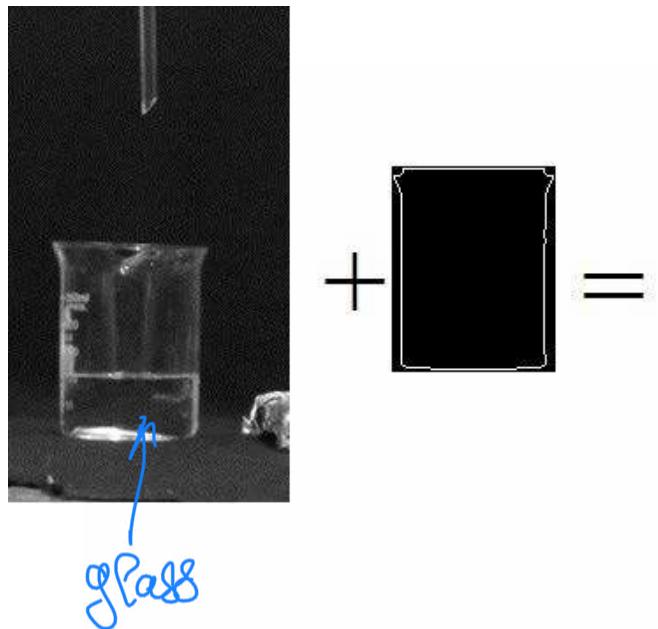


vote for the center where tangent will be
competitive

If ellipse can be rotated \Rightarrow 5D accumulator array

The accumulator need only be incremented for this θ .

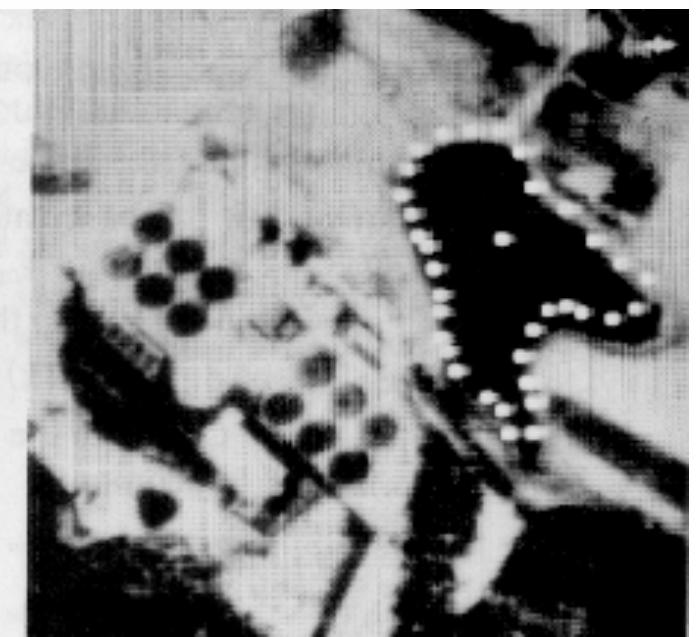
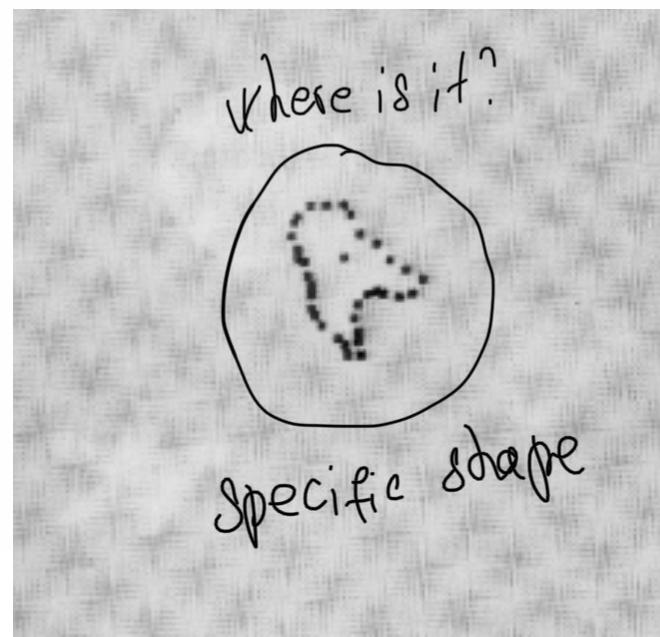
Generalized Hough



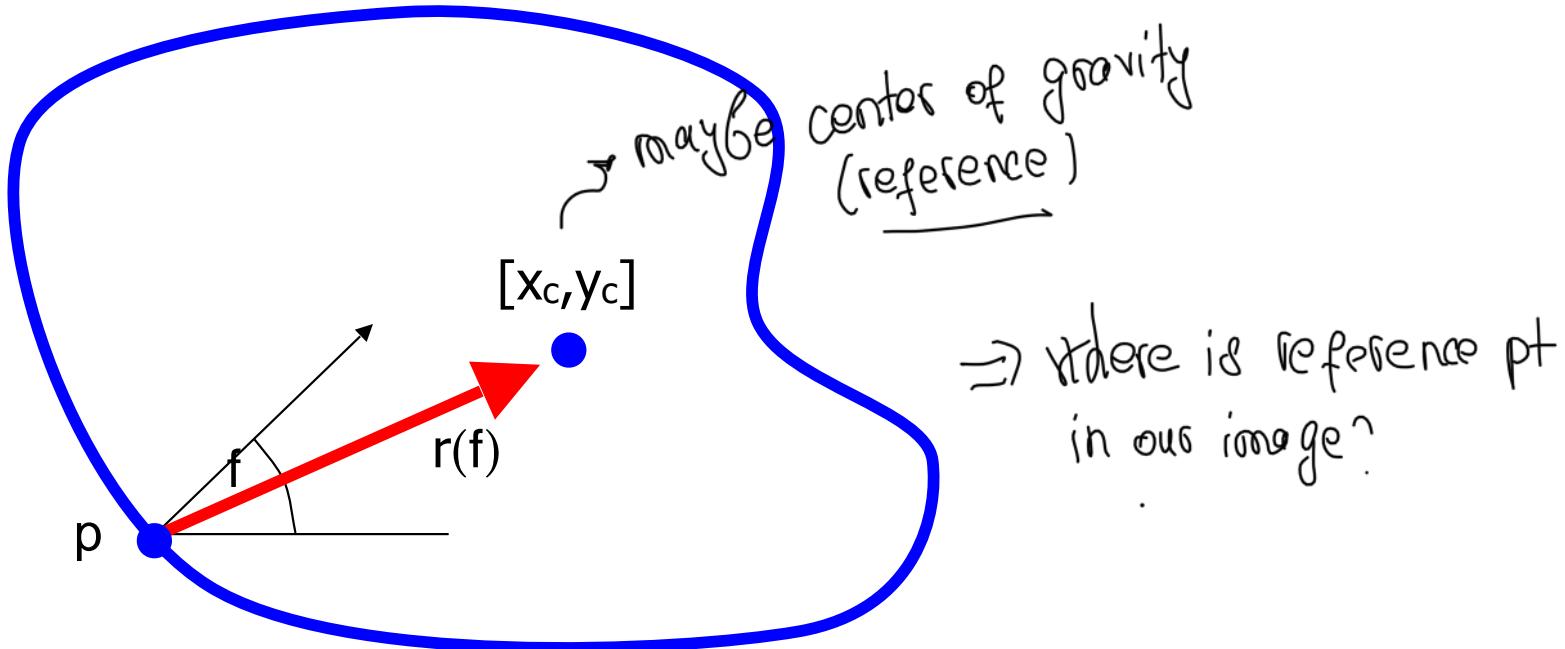
... or a lake.

Draw
Contours of a becher and then
Finding a becher ...

Paint it



Generalized Hough

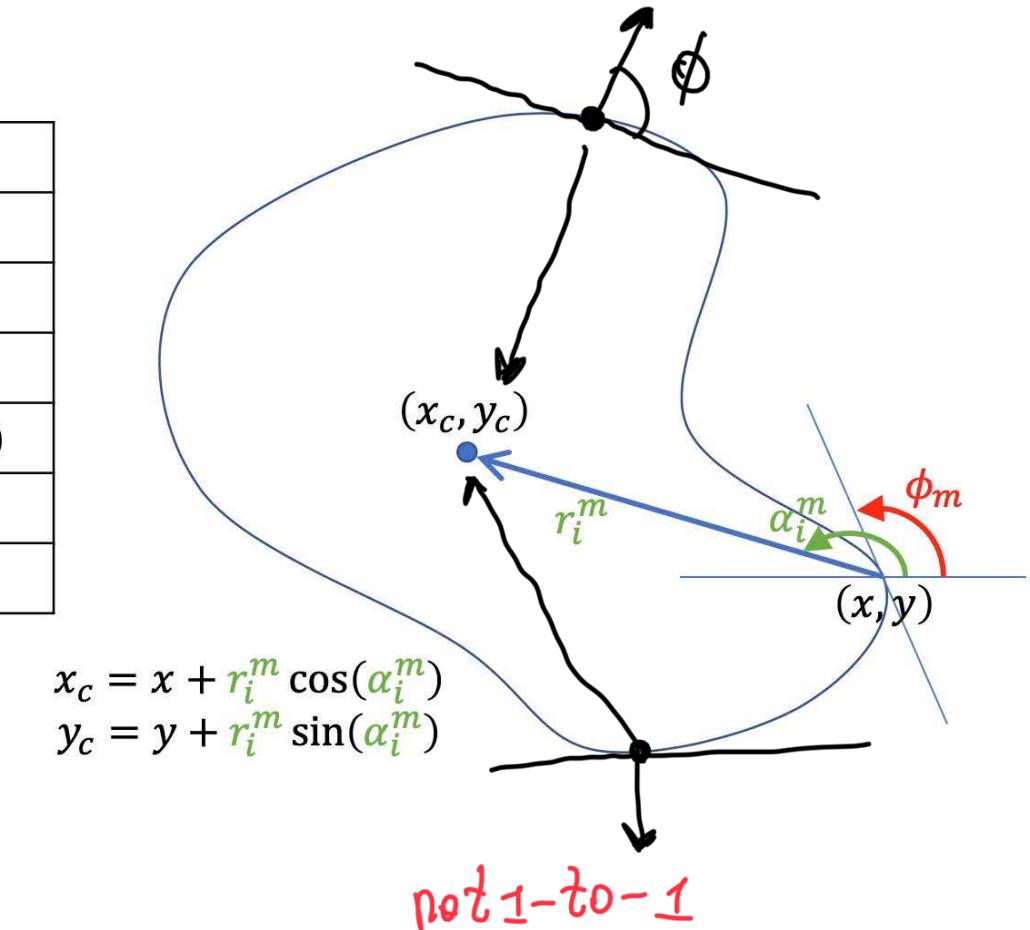


- We want to find a shape defined by its boundary points in terms of the location of a reference point $[x_c, y_c]$.
- For every boundary point p , we can compute the displacement vector $r = [x_c, y_c] - p$ as a function of local gradient orientation f .

R-Table

ϕ	$R(\phi_i)$
ϕ_1	$(r_1^1, \alpha_1^1), (r_2^1, \alpha_2^1), \dots, (r_{n1}^1, \alpha_{n1}^1)$
ϕ_2	$(r_1^2, \alpha_1^2), (r_2^2, \alpha_2^2), \dots, (r_{n2}^2, \alpha_{n2}^2)$
..
ϕ_m	$(r_1^m, \alpha_1^m), (r_2^m, \alpha_2^m), \dots, (r_i^m, \alpha_i^m), \dots, (r_{nm}^m, \alpha_{nm}^m)$
..
ϕ_M	$(r_1^M, \alpha_1^M), (r_2^M, \alpha_2^M), \dots, (r_{nM}^M, \alpha_{nM}^M)$

(x, y) which has gradient
 \Rightarrow it'll vote for "center" (x_c, y_c)
 for all elements in R-table



Set of potential displacement vectors r, a
 given the boundary orientation f . ϕ_i 's

--> Generalized template matching.

Algorithm

1. Make an R-table for the shape to be located.
2. Form an accumulator array of possible reference points initialized to zero.
3. For each edge point,
 - Compute the possible centers, that is, for each table entry, compute
$$x = x_e + r_\phi \cos(\theta(\phi))$$
$$y = y_e + r_\phi \sin(\theta(\phi))$$
 - Increment the accumulator array

Voting scheme \Rightarrow All edge point votes for all places the center could be wrt to it given its gradient direction

Robust, reliable,
fast

Real-Time Hough

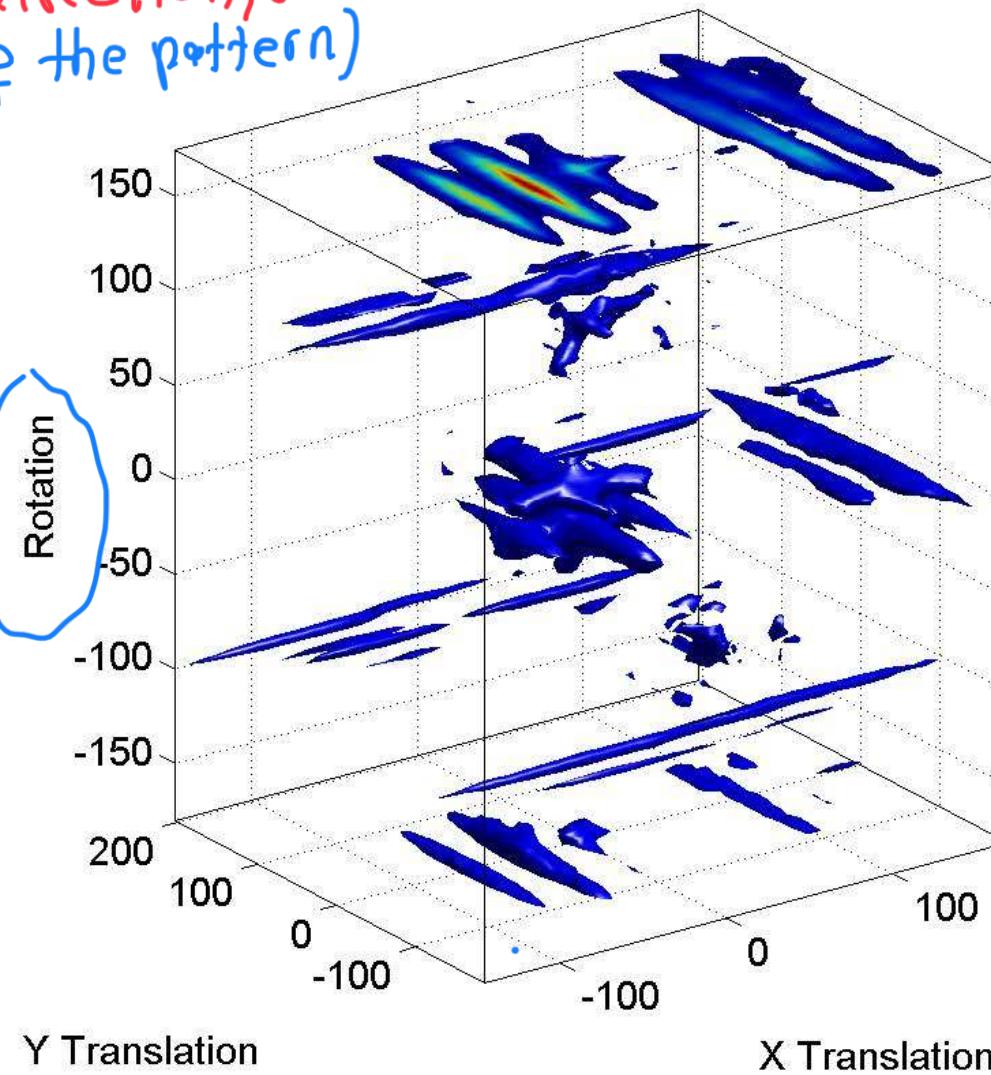


Accumulator

2D-accumulator

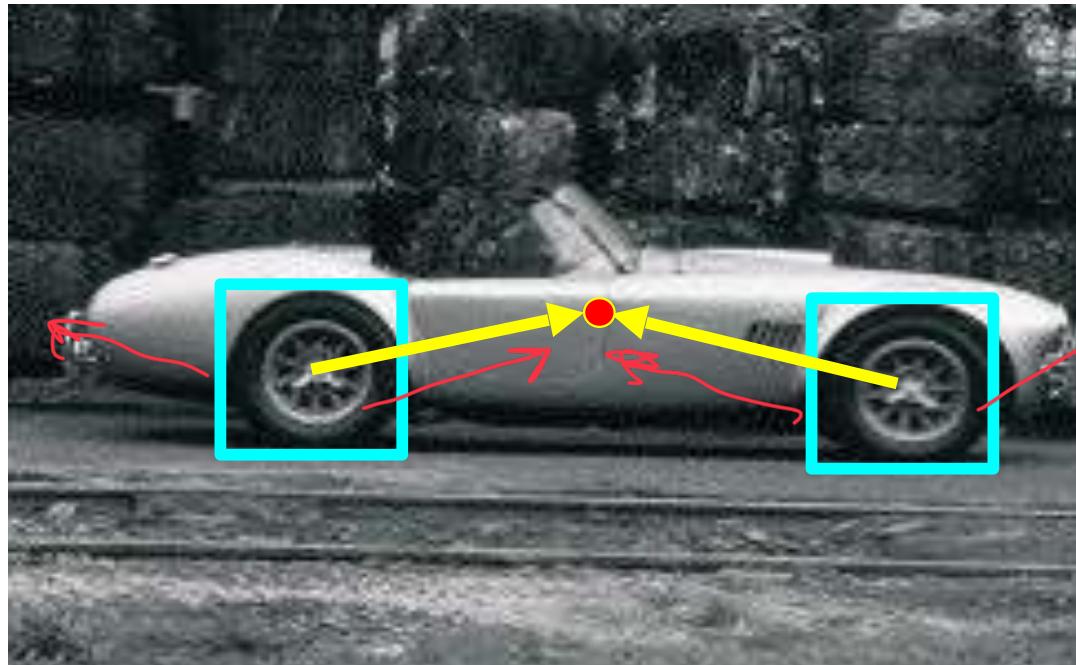
Previously, we were assuming to know gradient direction, but in practice \rightarrow no
(orientation of the pattern)

A possible orientation
of the pattern



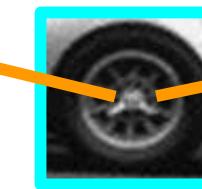
From Delineation To Detection

Detecting car \Rightarrow detecting wheels



Training image

they can vote for the position of the car and get estimate for the center of the car



Visual codeword with displacement vectors

Instead of indexing displacements by gradient orientation, index by “visual codeword”.

Hough
(from generalized)

Limitations

Fairly, effectively detects curves → including complicated ones
↳ defined relatively by small # of parameters
↳ knowing shape priori, only looking for orientations

Computational cost grows exponentially with the number of model parameters:

→ Only works for objects whose shape can be defined by a small number of parameters.

→ when assumptions met

→ Approach is robust but lacks flexibility.

From Delineation To Detection

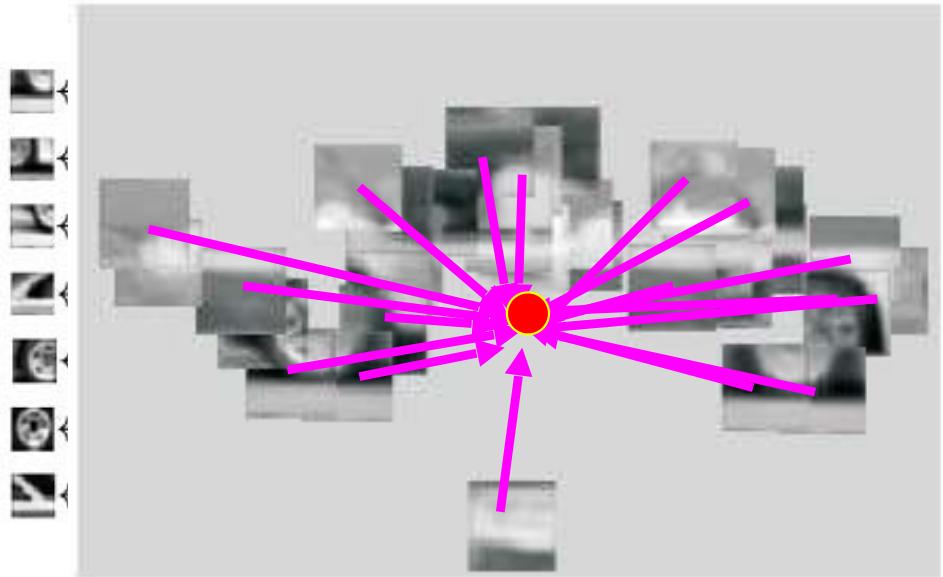
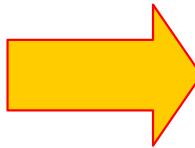
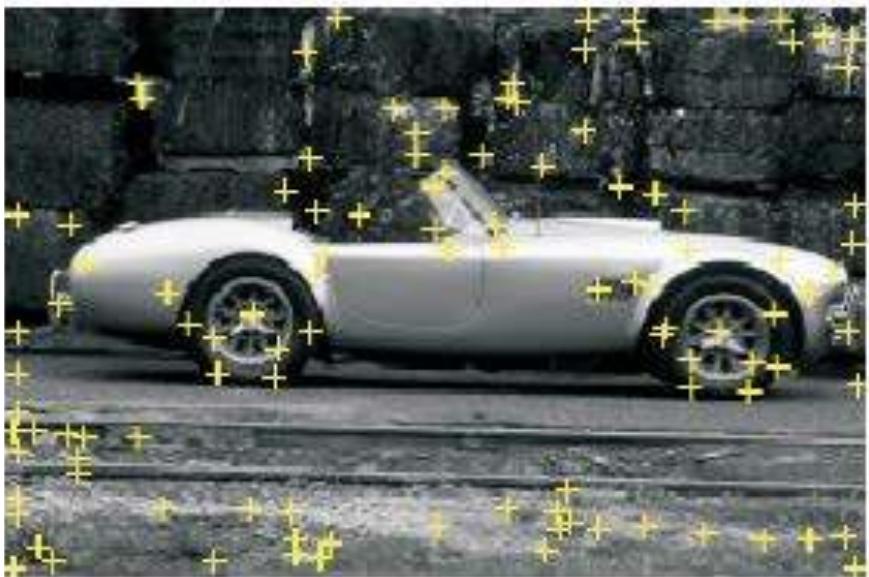


Test image

Instead of indexing displacements by gradient orientation, index by “visual codeword”.

Optional: Training

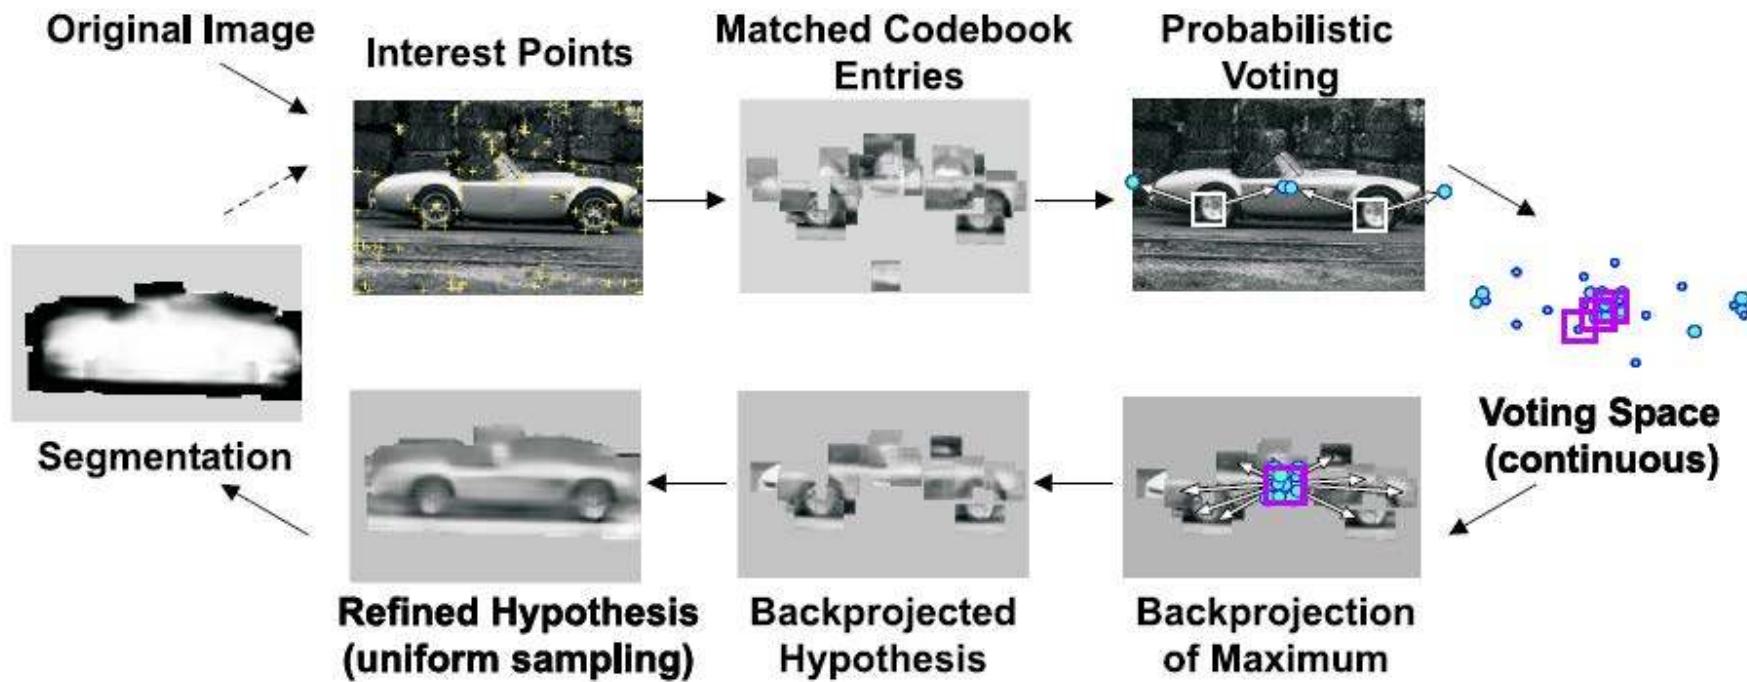
(ML, not DL)



1. Use clustering to build codebook of patches around extracted interest points.
≡ of images (repeatedly)
2. Map the patch around each interest point to closest codebook entry.
3. For each codebook entry, store all positions it was found, relative to object center.

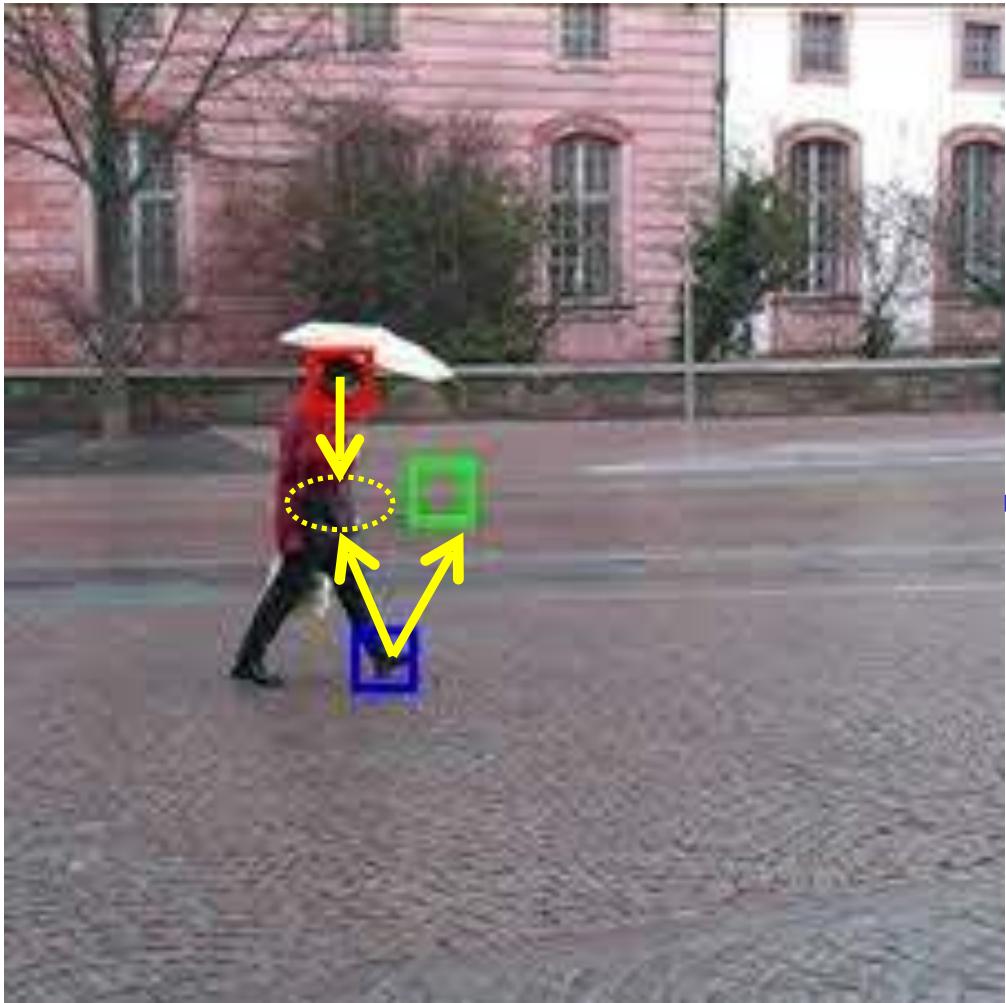
--> Build an R table.

Optional: Testing



1. Given test image, extract patches, match to codebook entry.
2. Cast votes for possible positions of object center.
3. Search for maxima in voting space.
4. Extract weighted segmentation mask based on stored masks for the codebook occurrences.

Pedestrian Detection



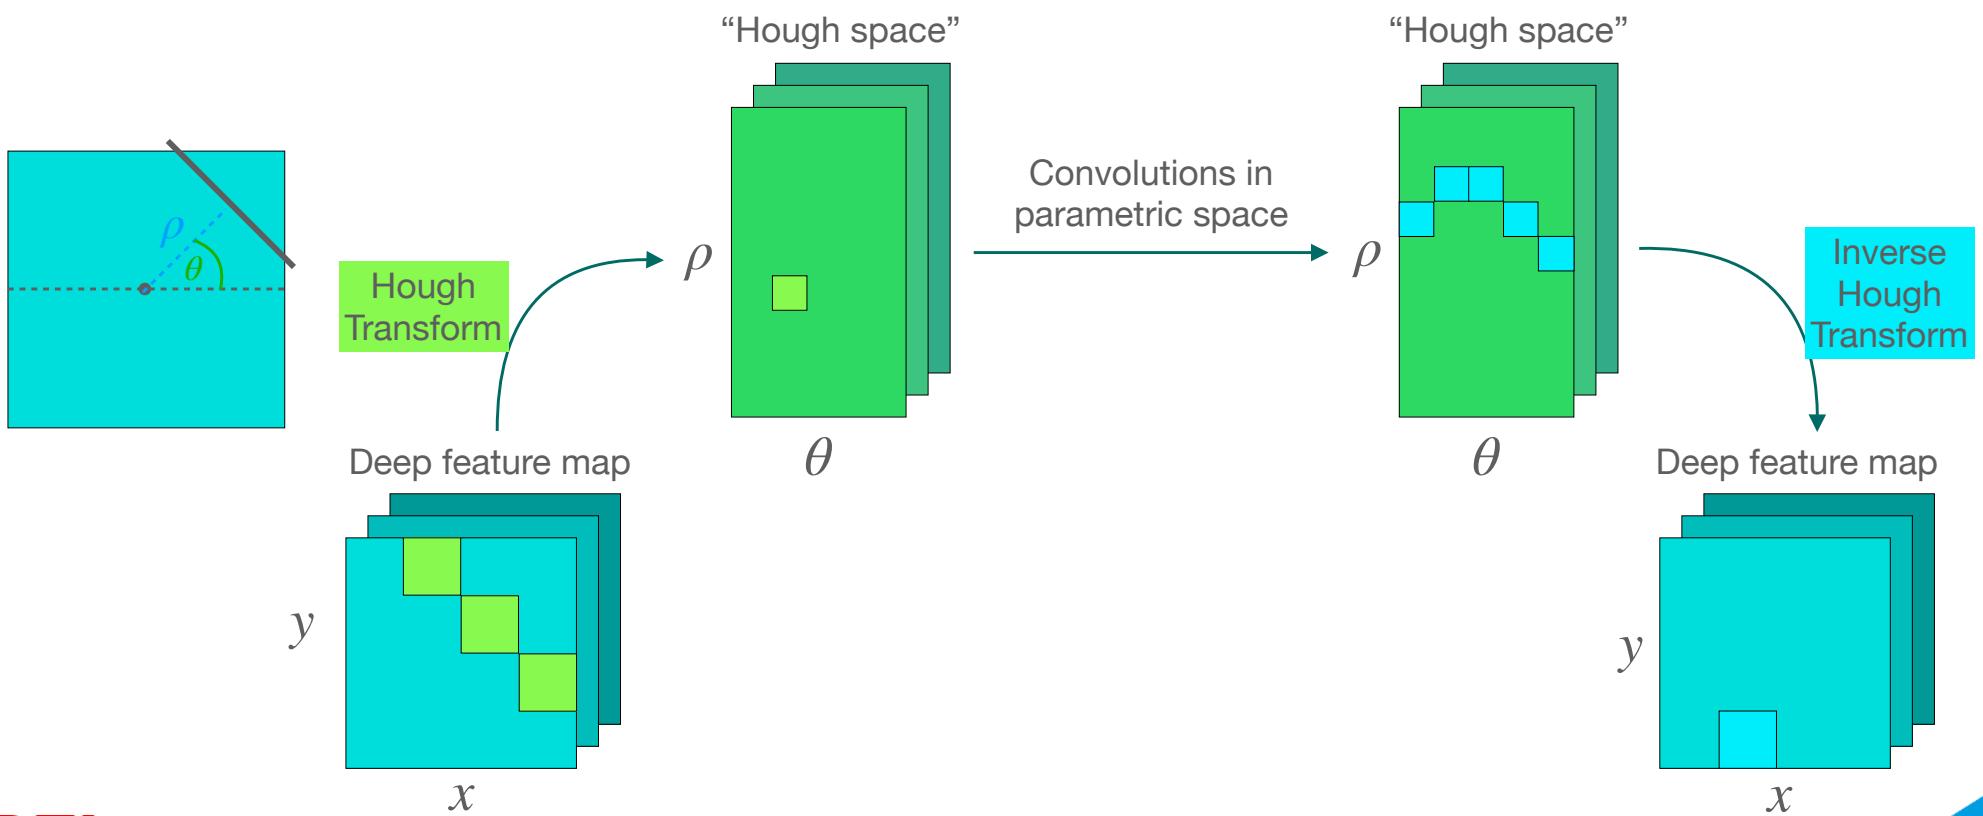
Find patches of images for
head, feet, ...

$\Rightarrow \underline{\text{Vote}}$ for the center of
gravity

Occlusion Handling



Deep Hough Transform



Techniques

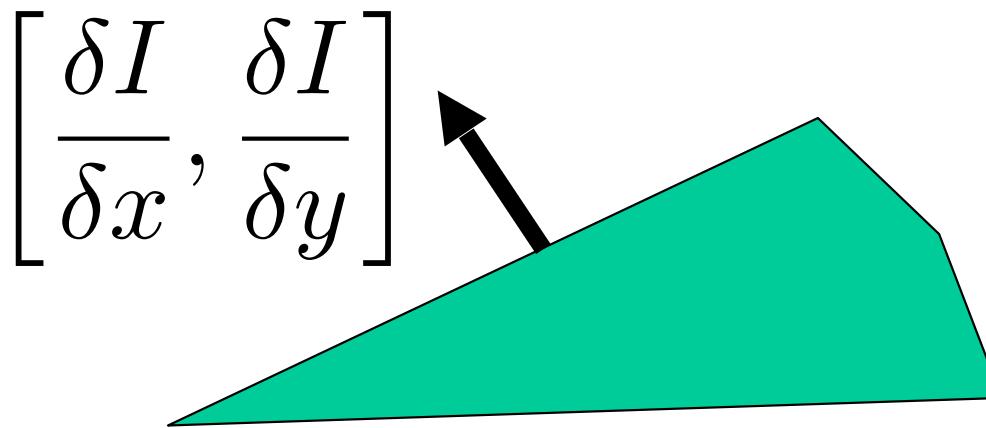
Semi-Automated Techniques:

- Dynamic programming
- Deformable Models

Fully Automated Techniques:

- Hough transform
- Graph Based Approaches

Magnitude and Orientation



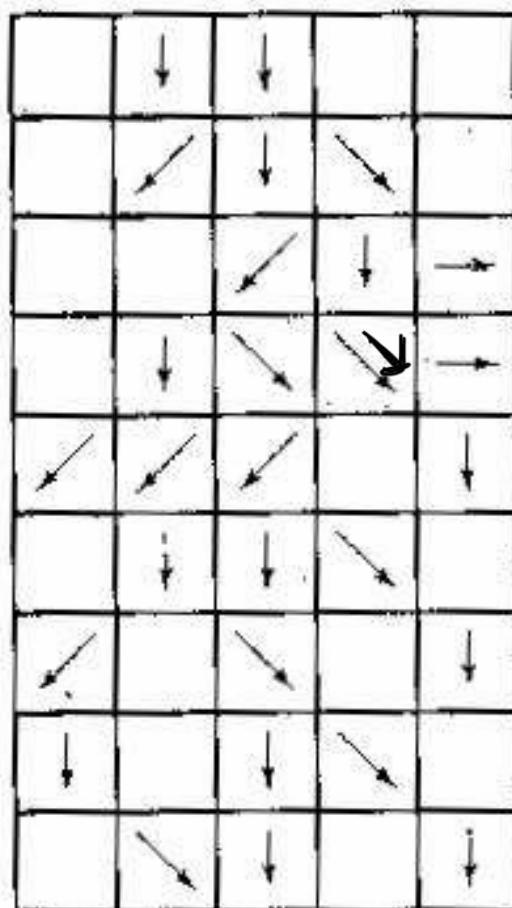
$$\text{Contrast: } G = \sqrt{\frac{\delta I^2}{\delta x} + \frac{\delta I^2}{\delta y}}$$

$$\text{Orientation: } \Theta = \arctan\left(\frac{\delta I}{\delta y}, \frac{\delta I}{\delta x}\right)$$

Minimum Spanning Tree

Image modeled as a graph:

Image represented
as a grid



gradient

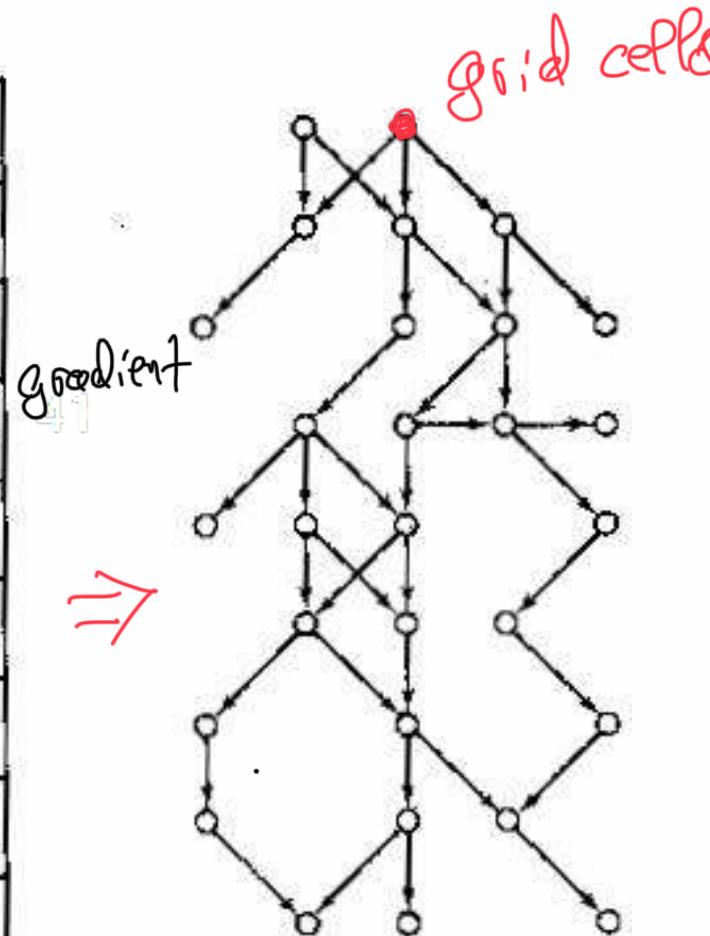


Image as a
graph

-> Generate minimal distance graph $O(N \log(N))$ algorithm)

↑ not connected
↔ they are likely
to be connected

Delineation 1998



Detecting contours
Detect road centerlines

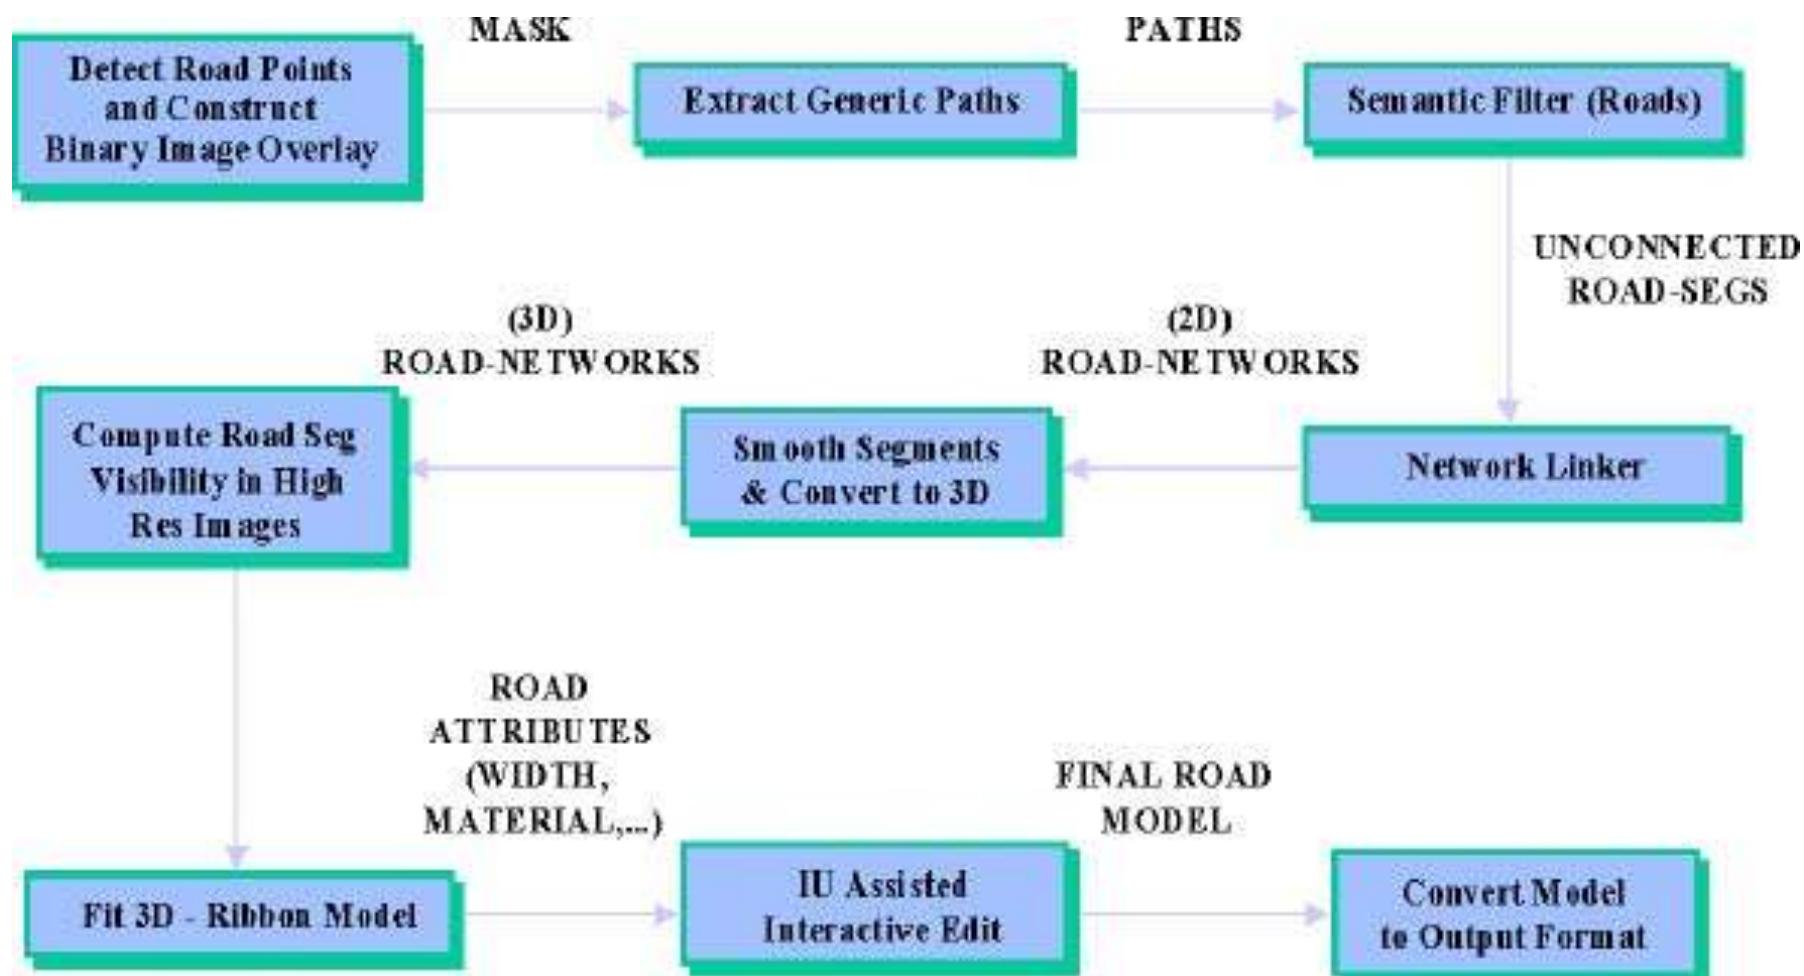
Pong, not too
Build ST and find paths^{ferry}
Find generic paths

Get rid of short, wiggly, not
informative
Apply semantic filter

Find road widths

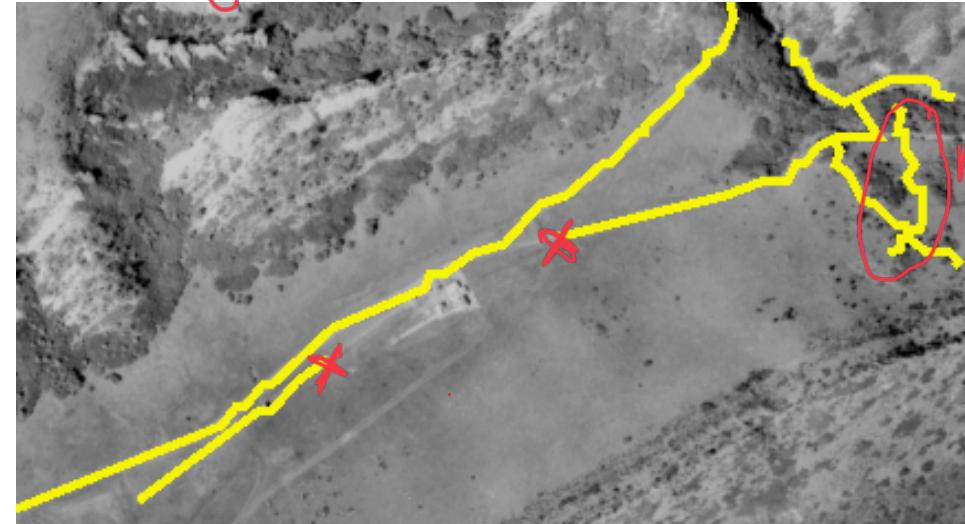
From Image To Roads

A lot of params need to be tuned
by human



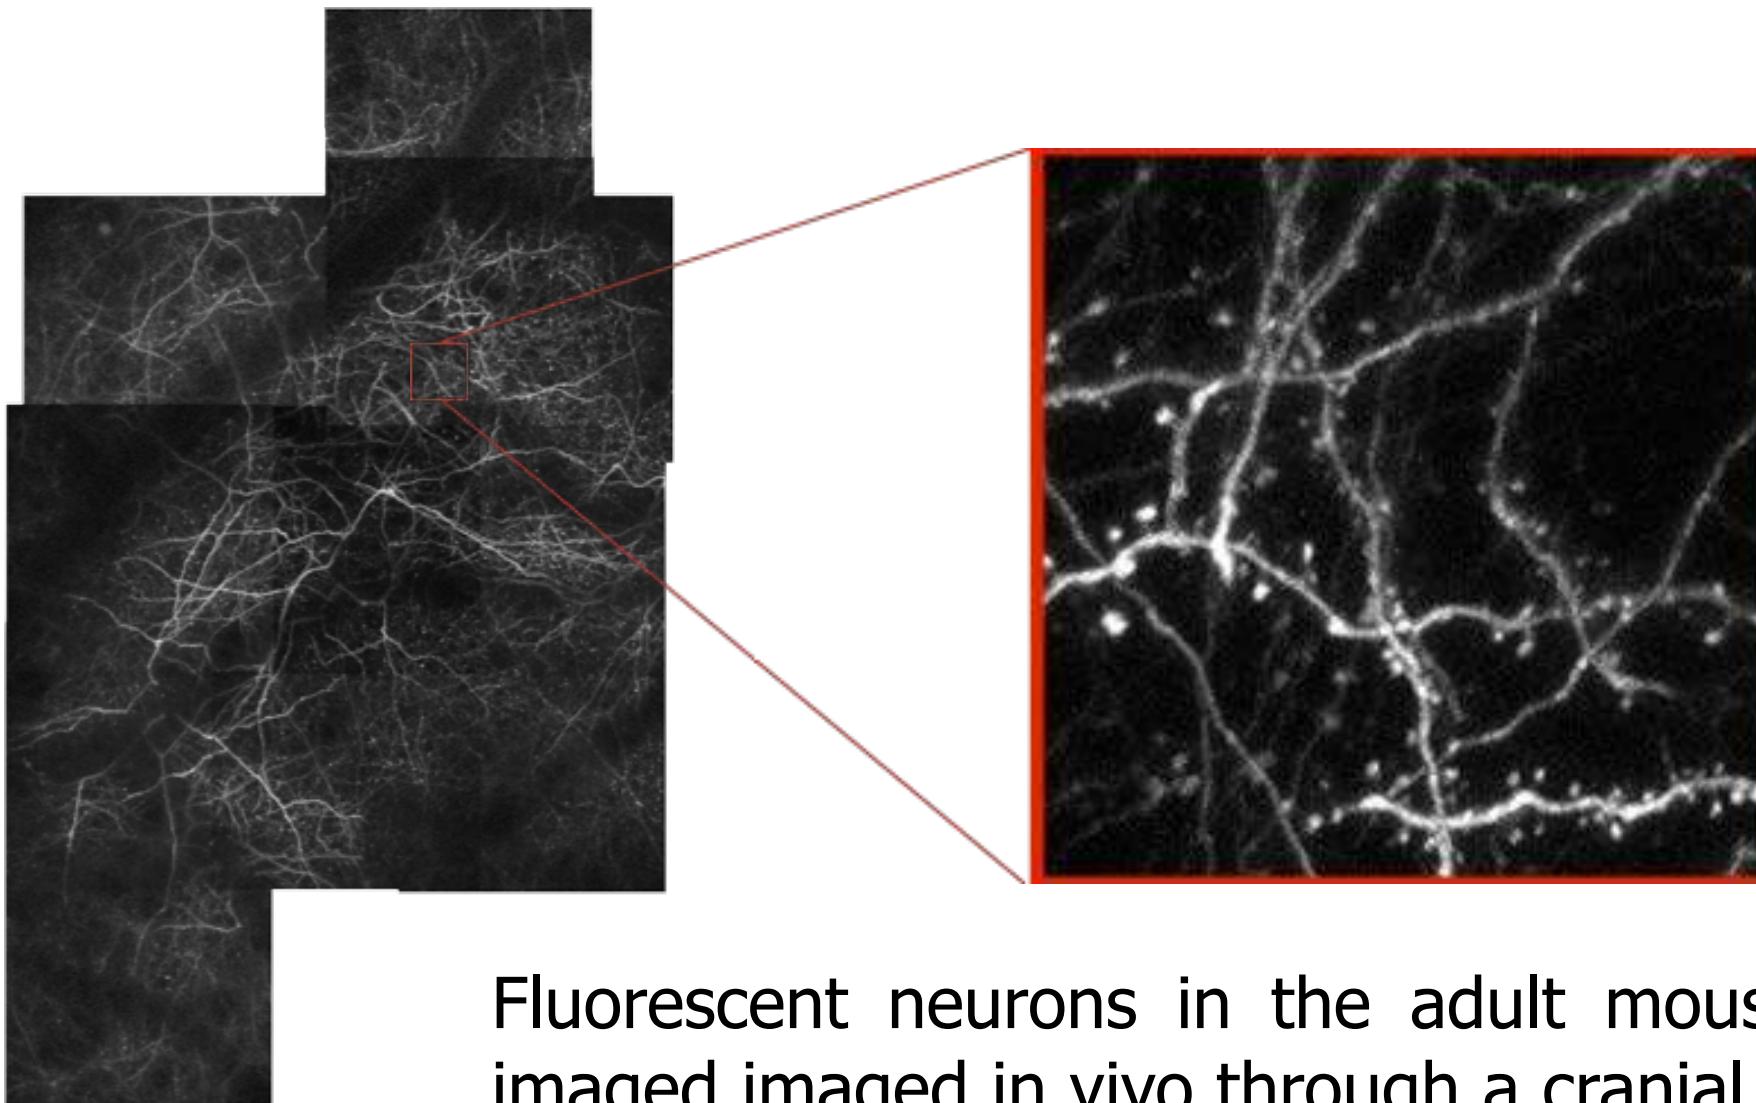
Road Editing

Fully automated



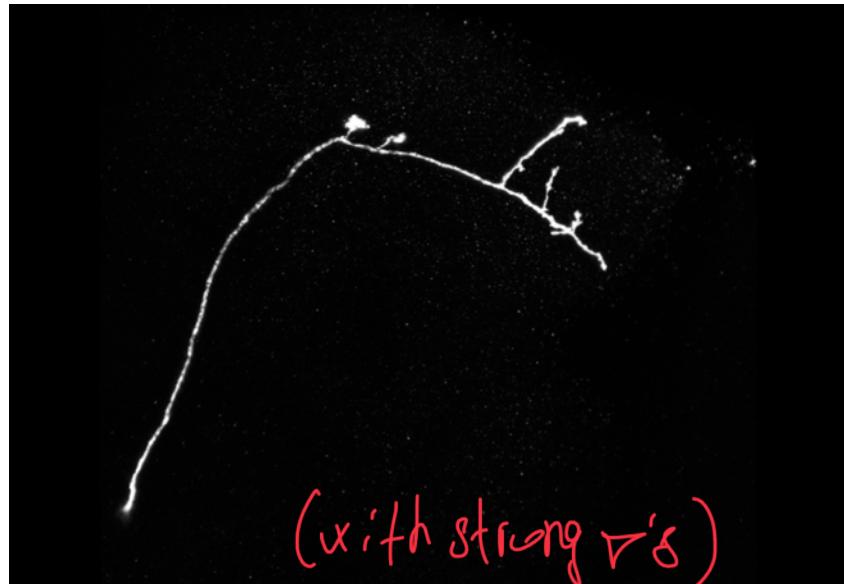
Use DP to finish
connecting

Dendrites And Axons

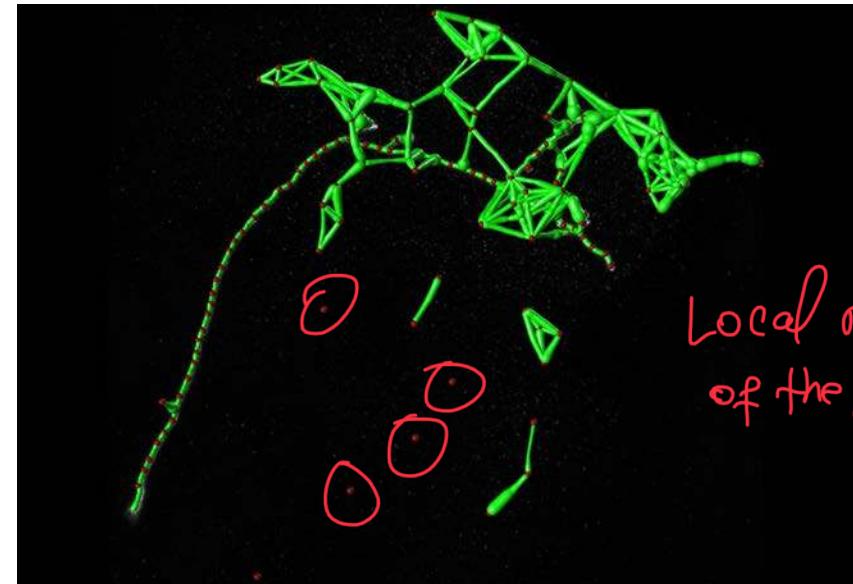


Fluorescent neurons in the adult mouse brain imaged *in vivo* through a cranial window using a 2-photon microscope.

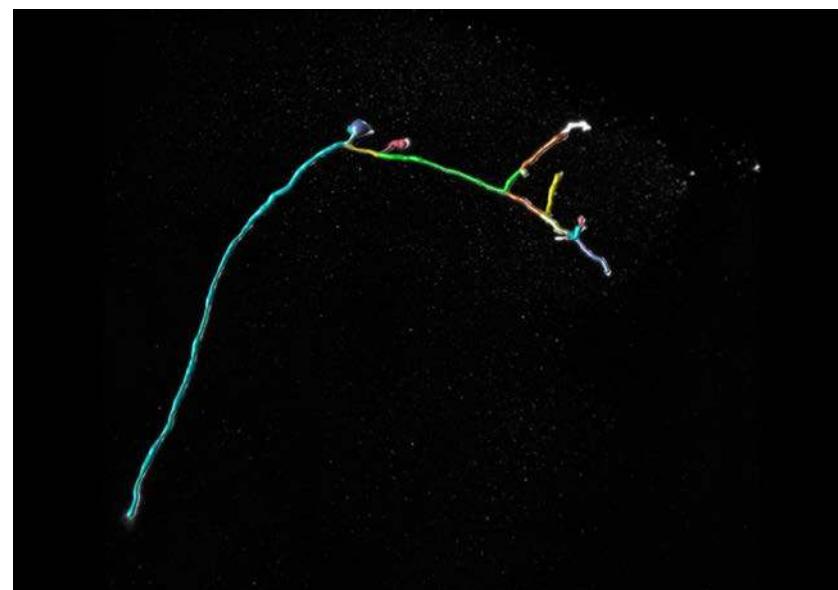
Delineation 2012: Neurites ...



Filtered Image



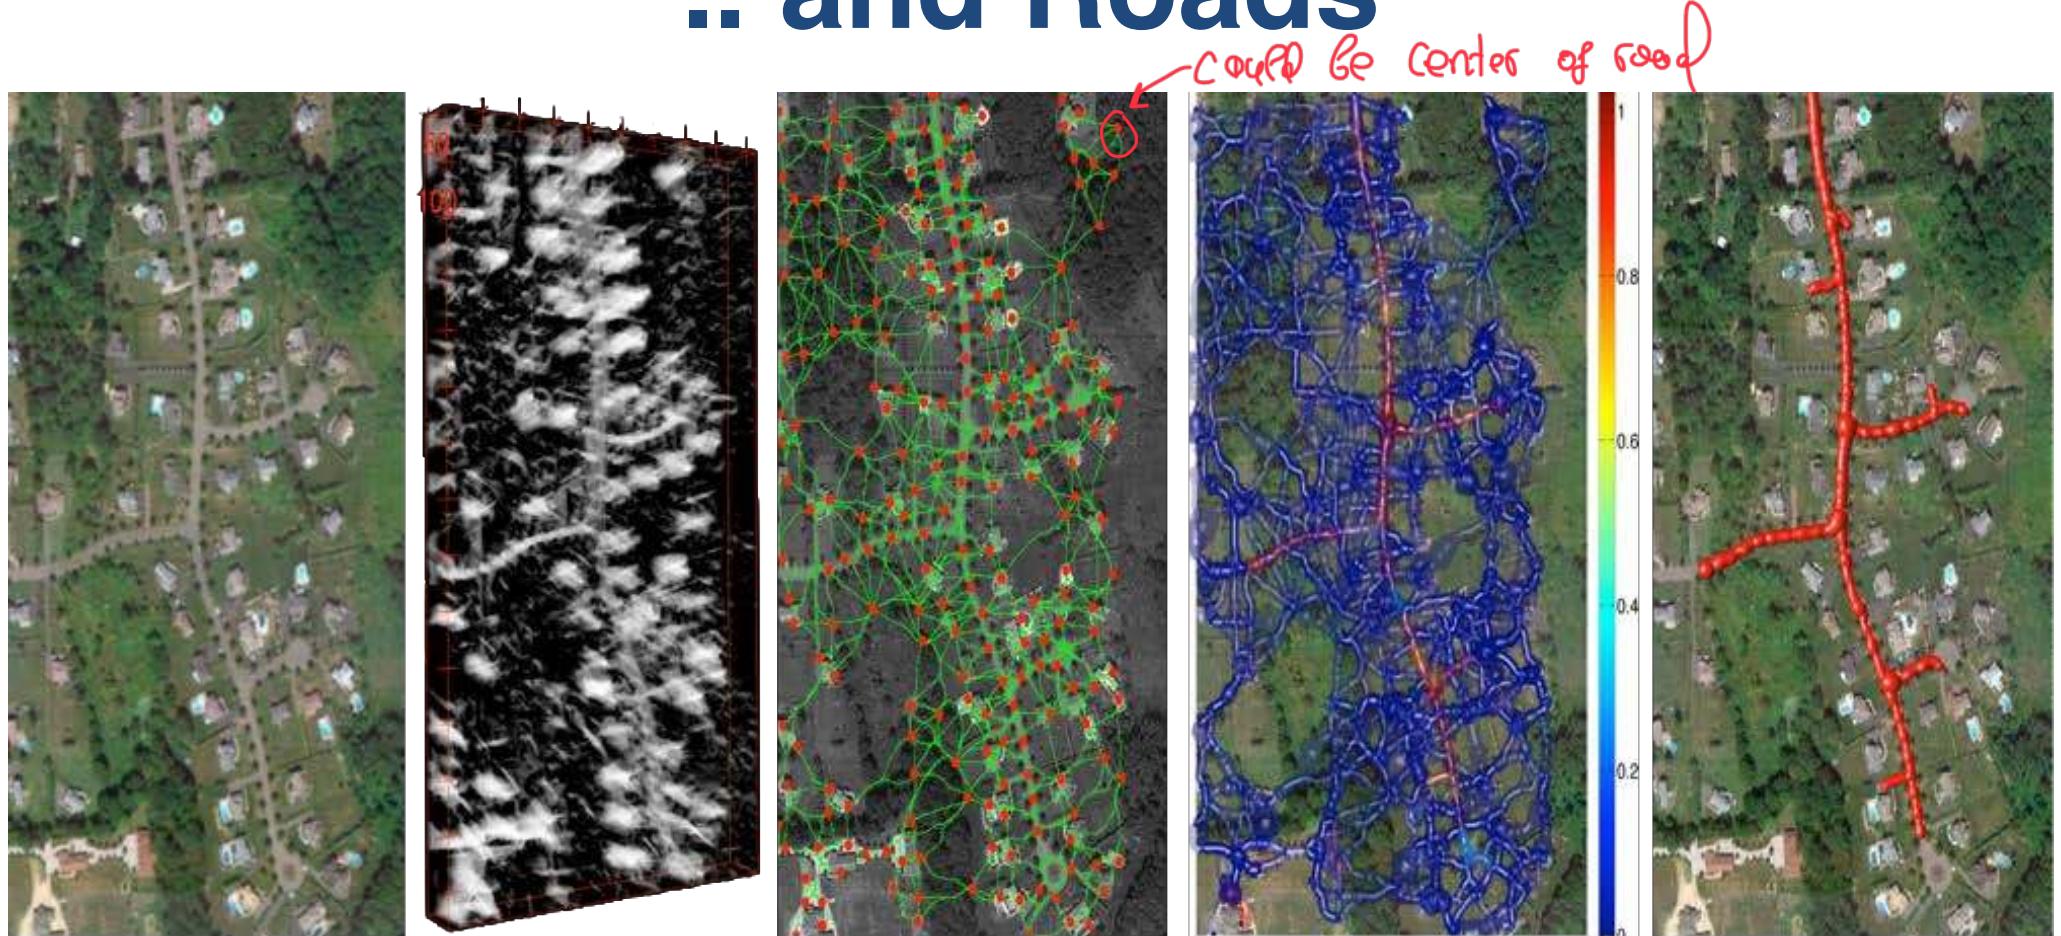
Graph



Maximum Likelihood Subtree

Keep paths that're
most probable and
give final delineation
of the axon

.. and Roads



Image

Filtered image

2D at several scales

Graph

Local maxima

Weighted graph

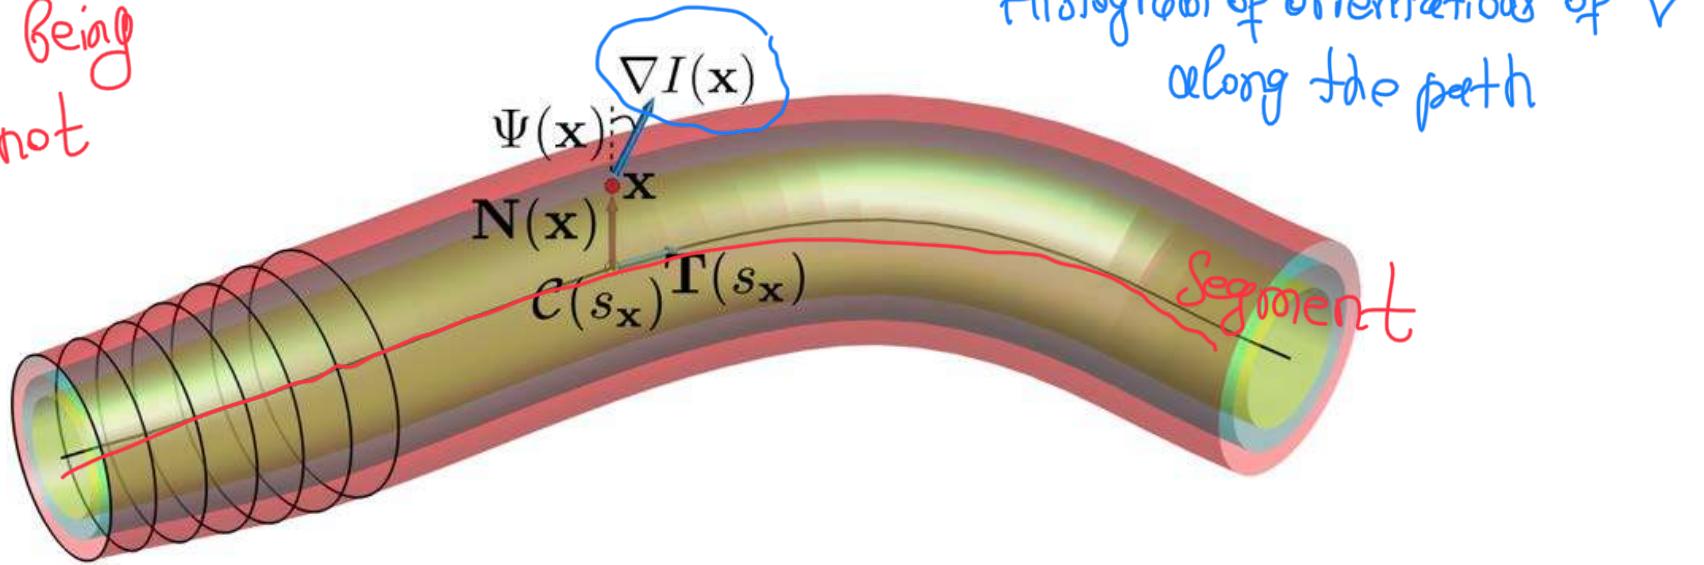
(ML \Rightarrow Likelihood of being part of road)

Subtree

—> Machine plays a crucial role to ensure that the **same algorithm works in different situations.**

Histogram of Gradient Deviations

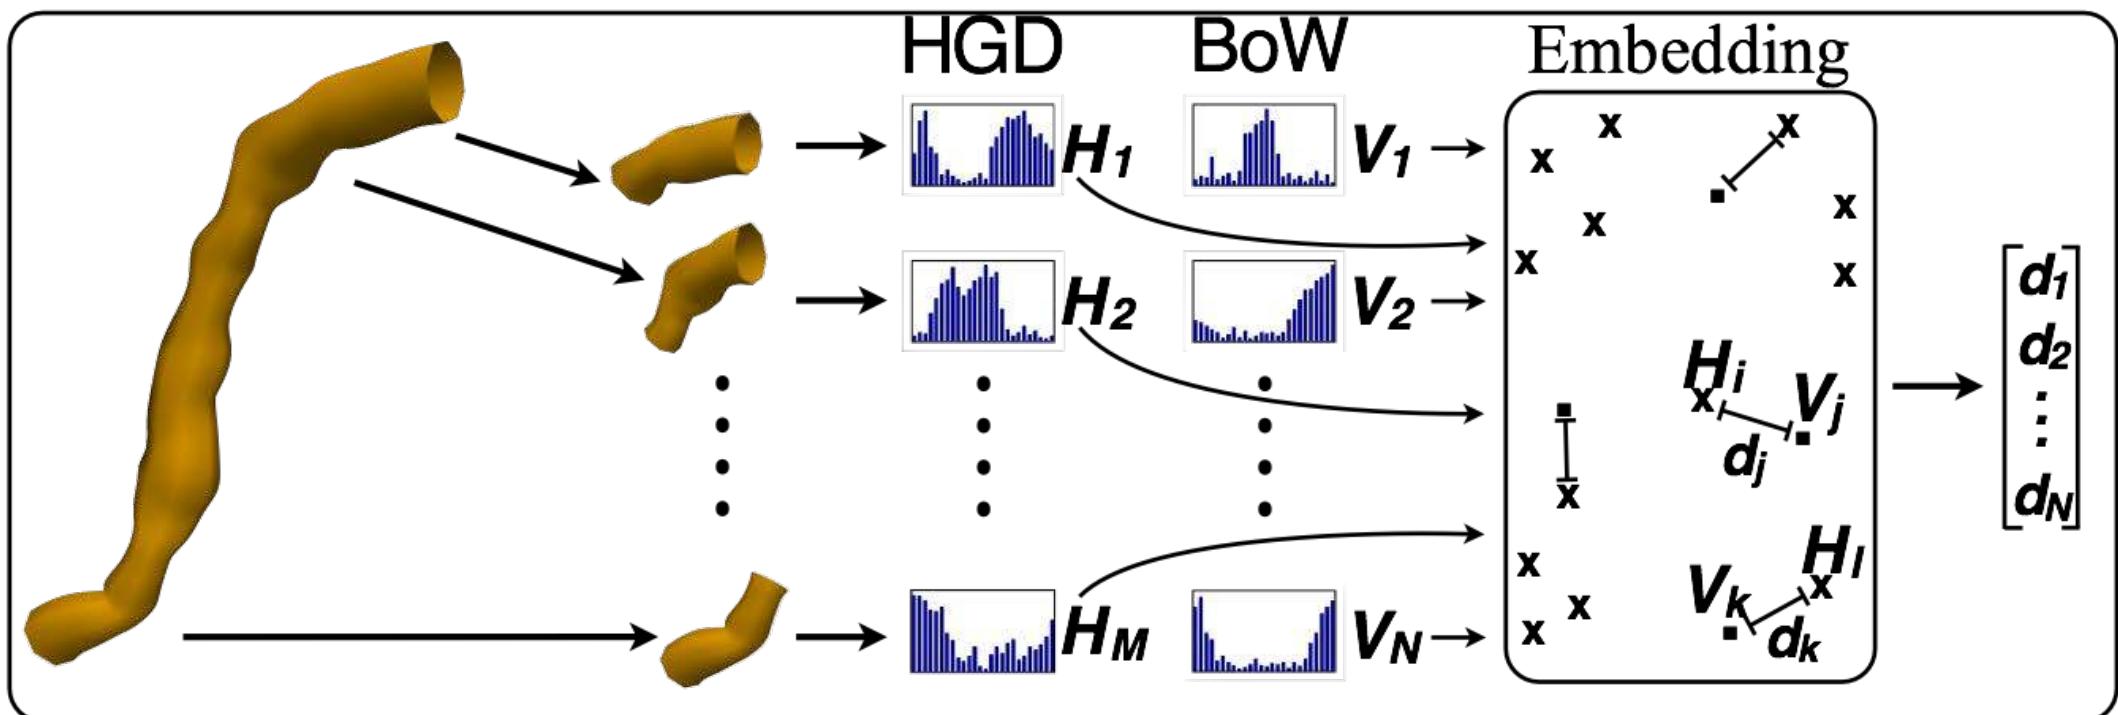
Classify segment being
a road line or not



$$\Psi(\mathbf{x}) = \begin{cases} \text{angle}(\nabla I(\mathbf{x}), \mathbf{N}(\mathbf{x})) , & \text{if } \|\mathbf{x} - \mathcal{C}(s_{\mathbf{x}})\| > \varepsilon \\ \text{angle}(\nabla I(\mathbf{x}), \Pi(\mathbf{x})) , & \text{otherwise,} \end{cases}$$

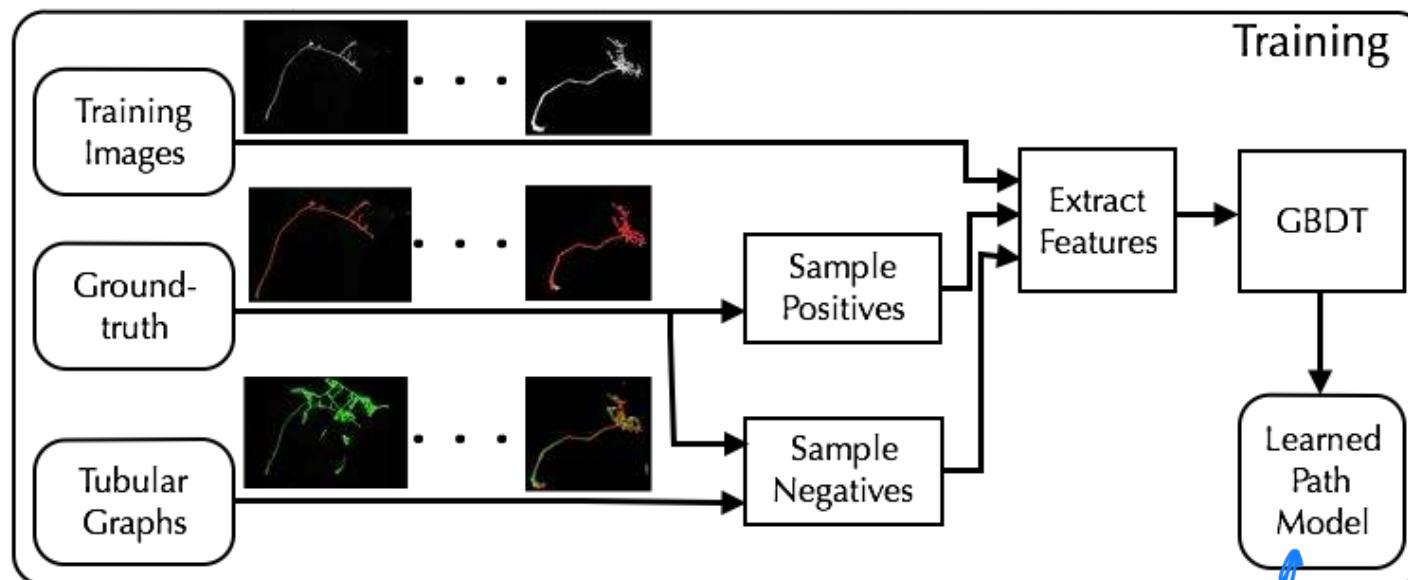
→ One histogram per radius interval plus four geometric features (curvature, tortuosity,).

Optional: Embedding

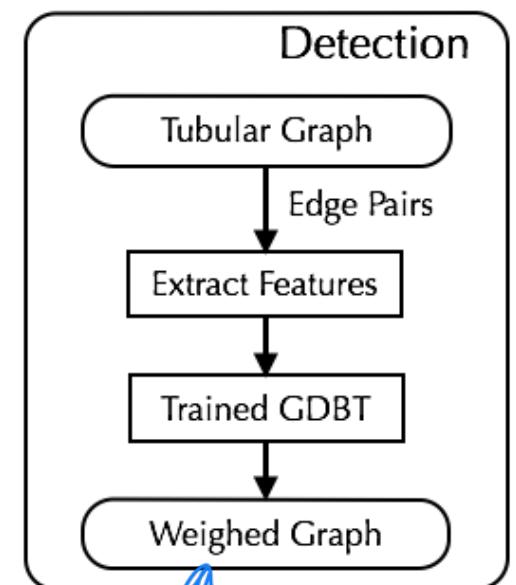


→ Same length feature vectors whatever the actual length of the path.

Optional: Path Classification

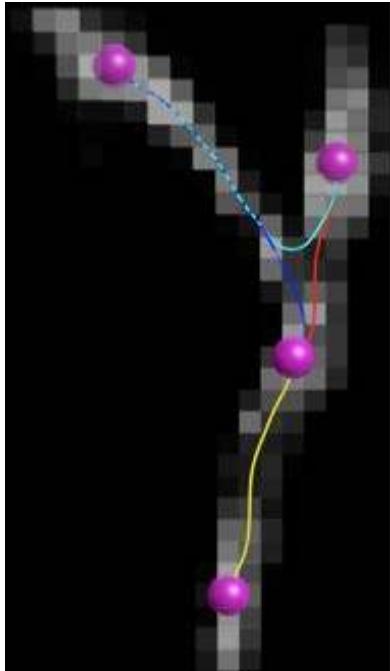


SYM's

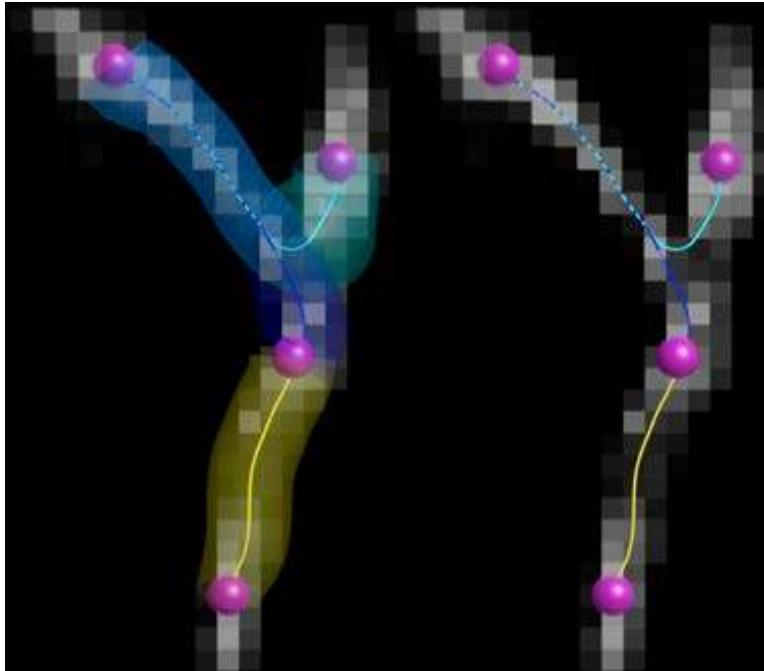


assign weights to
various edges

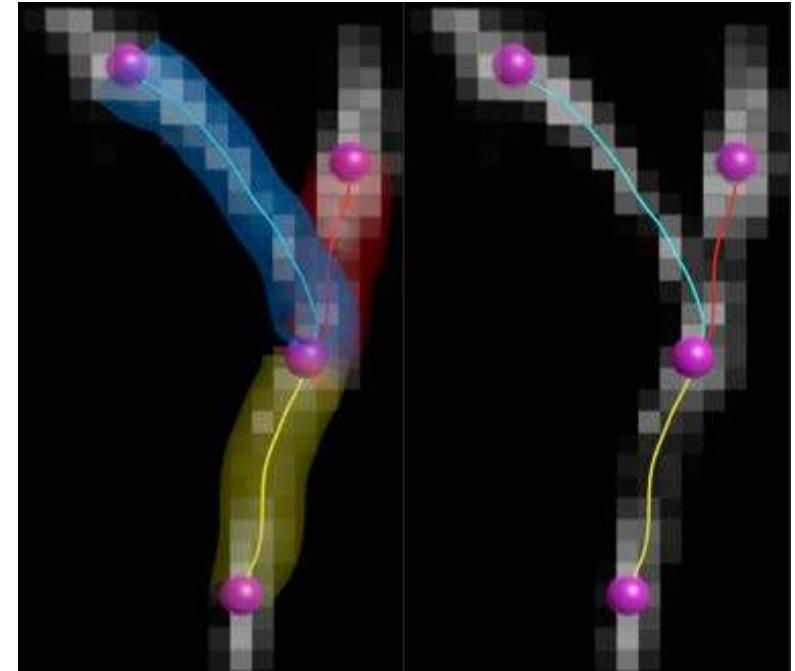
Optional: Finding the Best Tree



Tubularity graph



Without edge pair term



With edge pair term

$$\begin{aligned} \mathbf{t}^* &= \underset{\mathbf{t} \in \mathcal{T}(G)}{\operatorname{argmax}} P(\mathbf{T} = \mathbf{t} | I) , \\ &= \underset{\mathbf{t} \in \mathcal{T}(G)}{\operatorname{argmax}} P(I | \mathbf{T} = \mathbf{t}) P(\mathbf{T} = \mathbf{t}) , \\ &= \underset{\mathbf{t} \in \mathcal{T}(G)}{\operatorname{argmin}} \sum_{e_{ij} \in G} c_{ij}^d t_{ij} \end{aligned}$$

maximize objective function

Optional: QMIP Formulation

$$\begin{aligned} \min \quad & \sum_{e_{ij} \in E, e_{jk} \in E} c_{ijk} t_{ij} t_{jk} \\ \text{s.t.} \quad & \sum_{v_j \in V \setminus \{v_r\}} y_{rj}^l \leq 1, \quad \forall v_l \in V \setminus \{v_r\}, \\ & \sum_{v_j \in V \setminus \{v_k\}} y_{jk}^l \leq 1, \quad \forall v_l \in V \setminus \{v_r\}, \\ & \sum_{v_j \in V \setminus \{v_i, v_r\}} y_{ij}^l - \sum_{v_j \in V \setminus \{v_i, v_l\}} y_{ji}^l = 1, \quad \forall v_k \in V \setminus \{v_r\}, \\ & \forall v_i \in V \setminus \{v_r, v_k\}, \\ & y_{ij}^l \leq t_{ij}, \quad \forall e_{ij} \in E, v_l \in V \setminus \{v_r, v_i, v_j\}, \\ & y_{il}^l = t_{il}, \quad \forall e_{il} \in E, \\ & y_{ij}^l \geq 0, \quad \forall e_{ij} \in E, v_l \in V \setminus \{v_r, v_i\}, \\ & t_{ij} \in \{0, 1\}, \quad \forall e_{ij} \in E. \end{aligned}$$

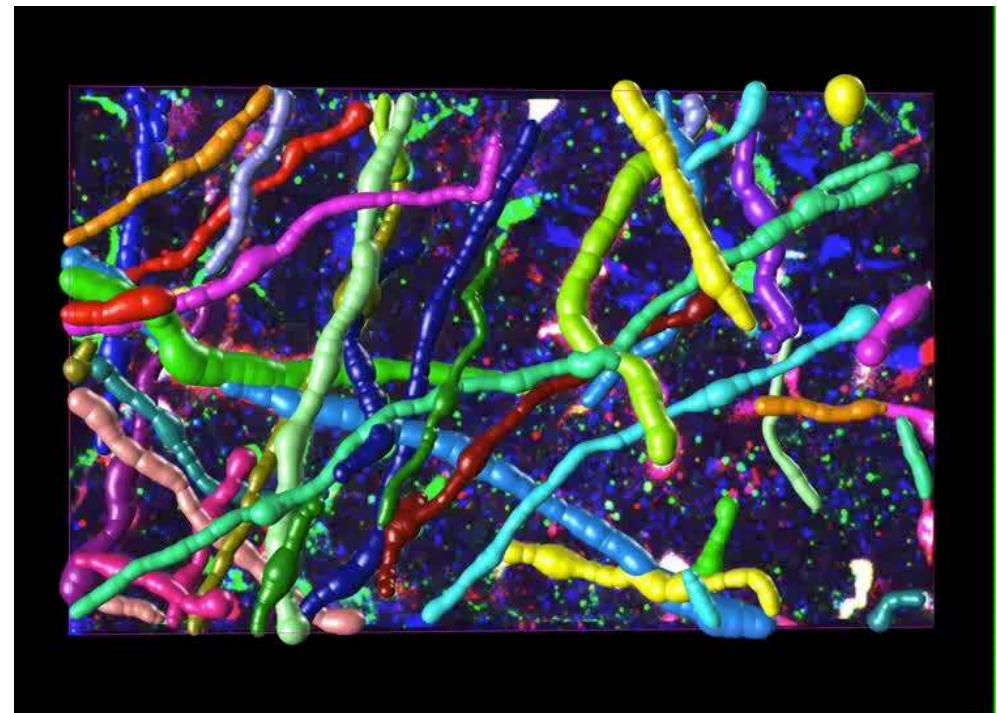
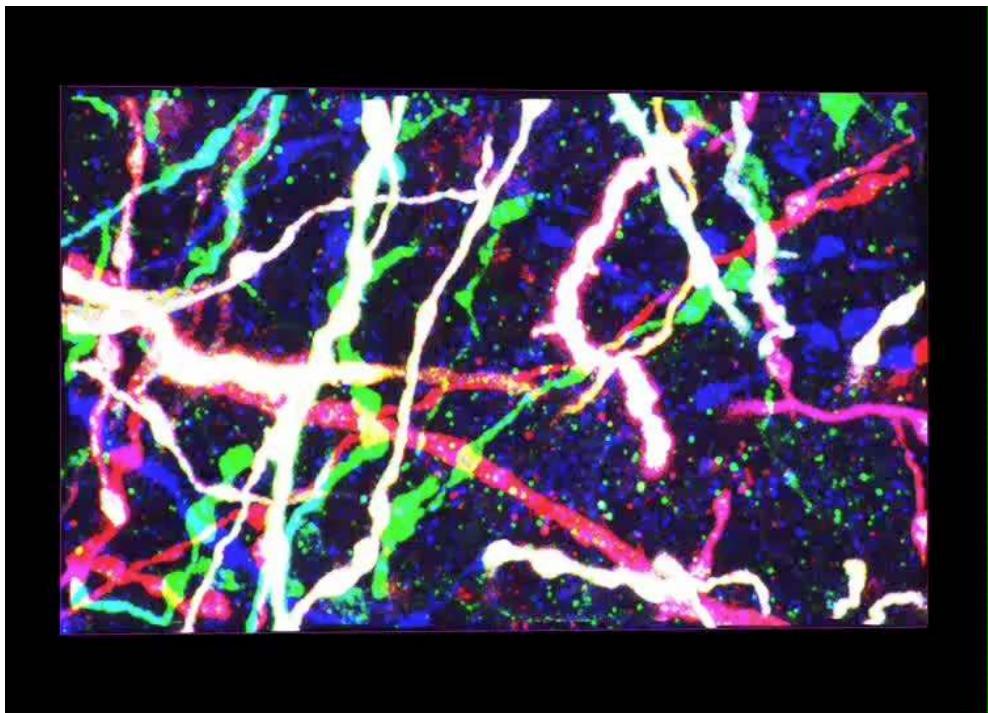
given the root note v_r .

Roads

Very good delineation

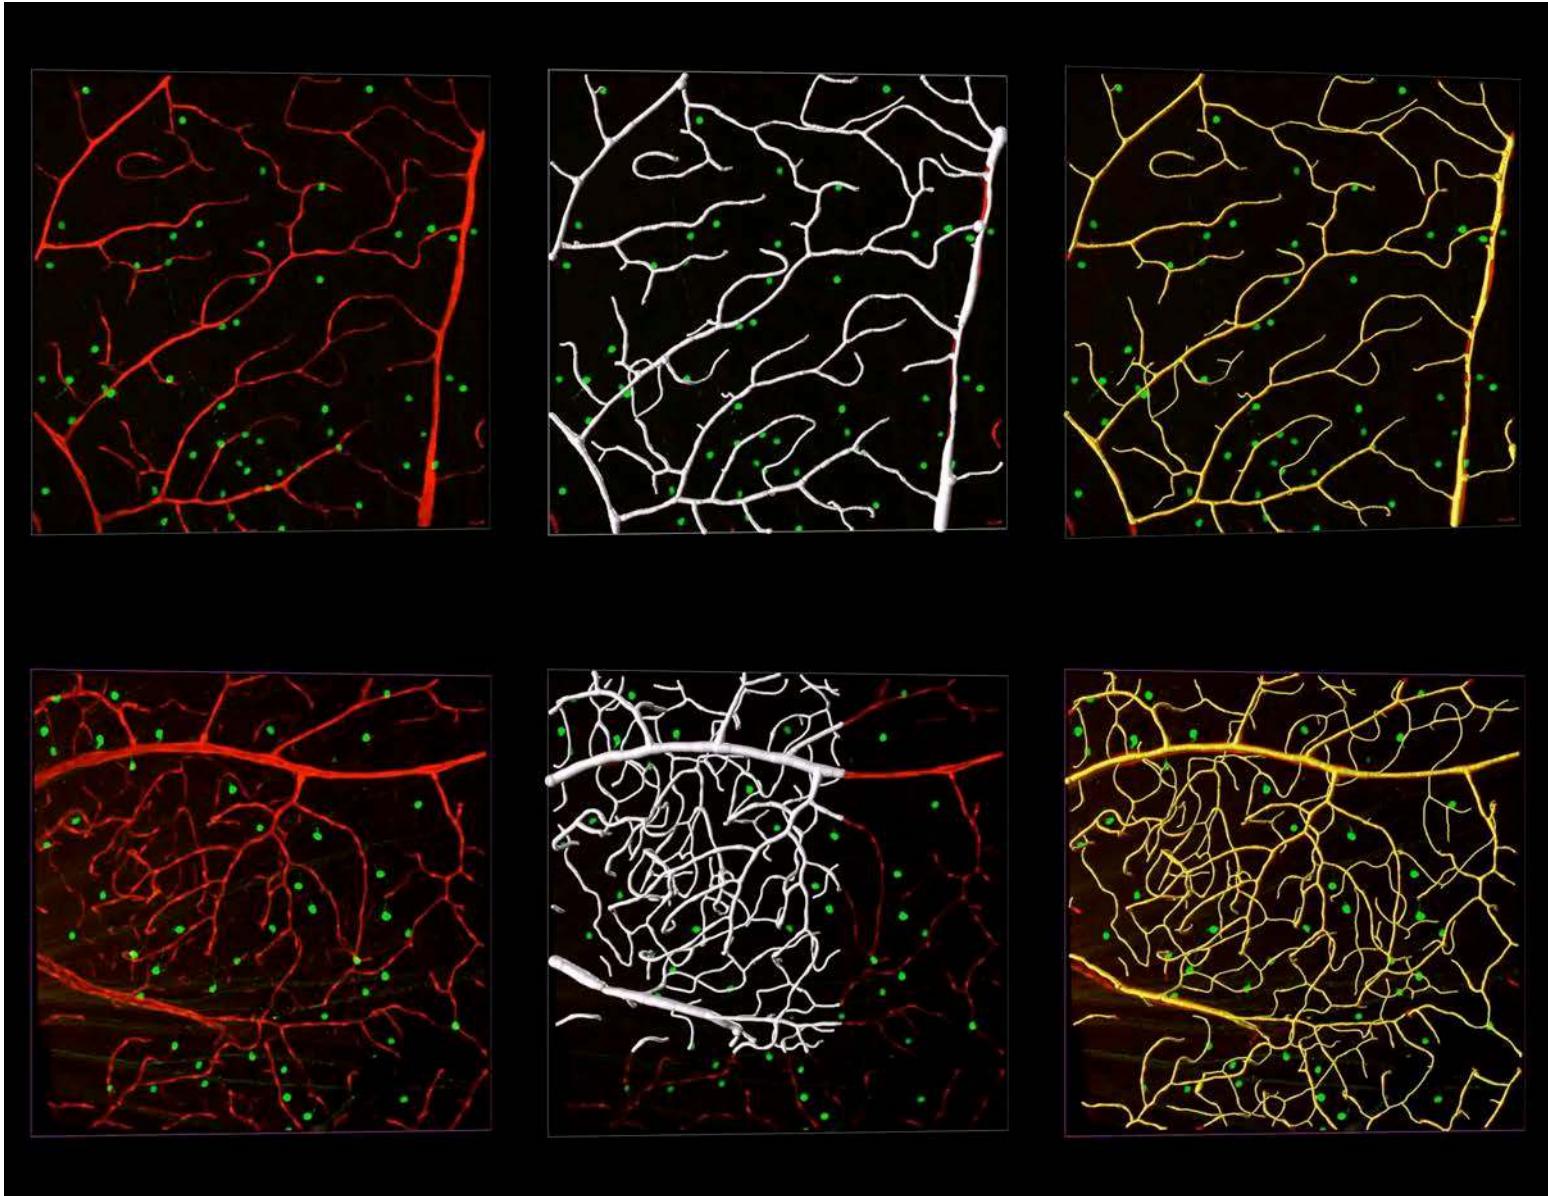


Brainbow Images



Blood Vessels

Only changing classifier \Rightarrow which paths are good/bad?



Deep Tsunami

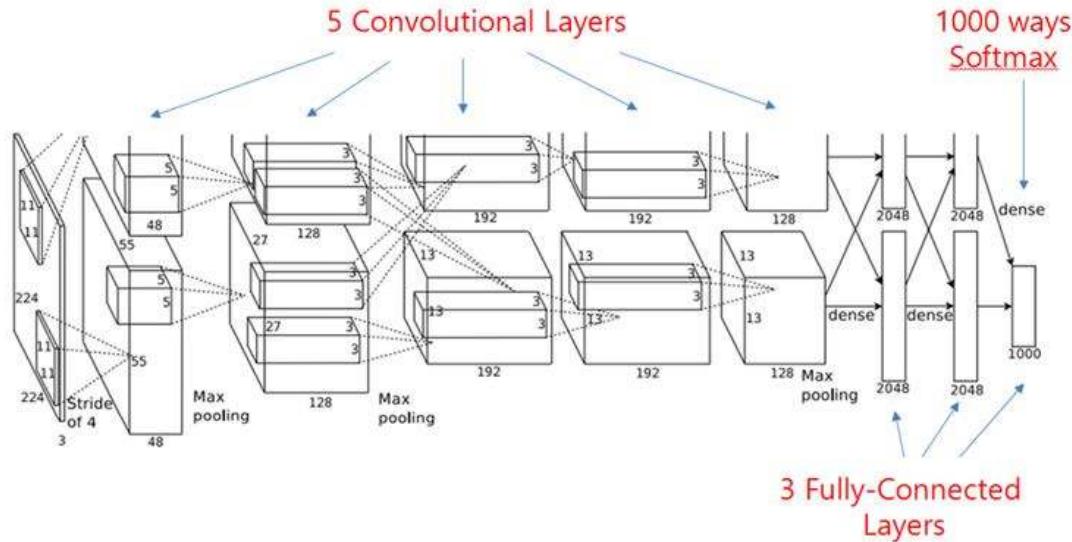
The New York Times

*Turing Award Won by Three
Pioneers in Artificial Intelligence*



From left, Yann LeCun, Geoffrey Hinton and Yoshua Bengio. The researchers worked on key developments for neural networks, which are reshaping how computer systems are built.

Reminder: AlexNet (2012)



Task: Image classification

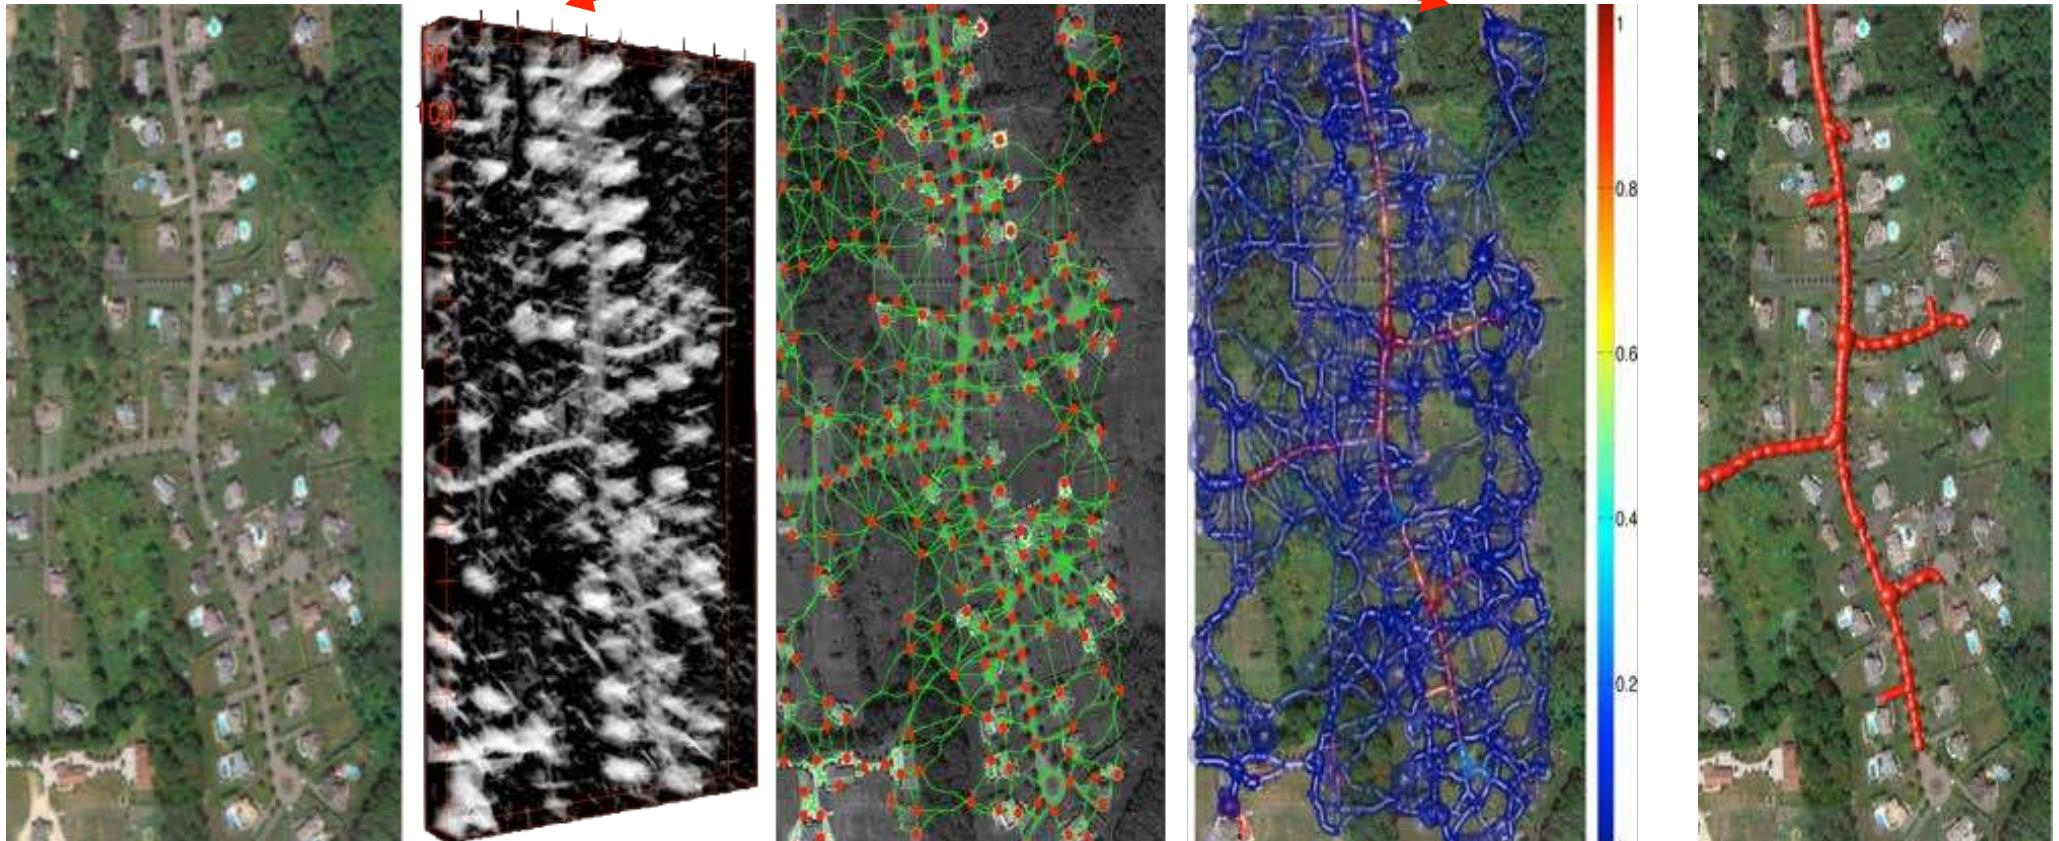
Training images: Large Scale Visual Recognition Challenge 2010

Training time: 2 weeks on 2 GPUs

Major Breakthrough: Training large networks has now been shown to be practical!!

Delineation 2012

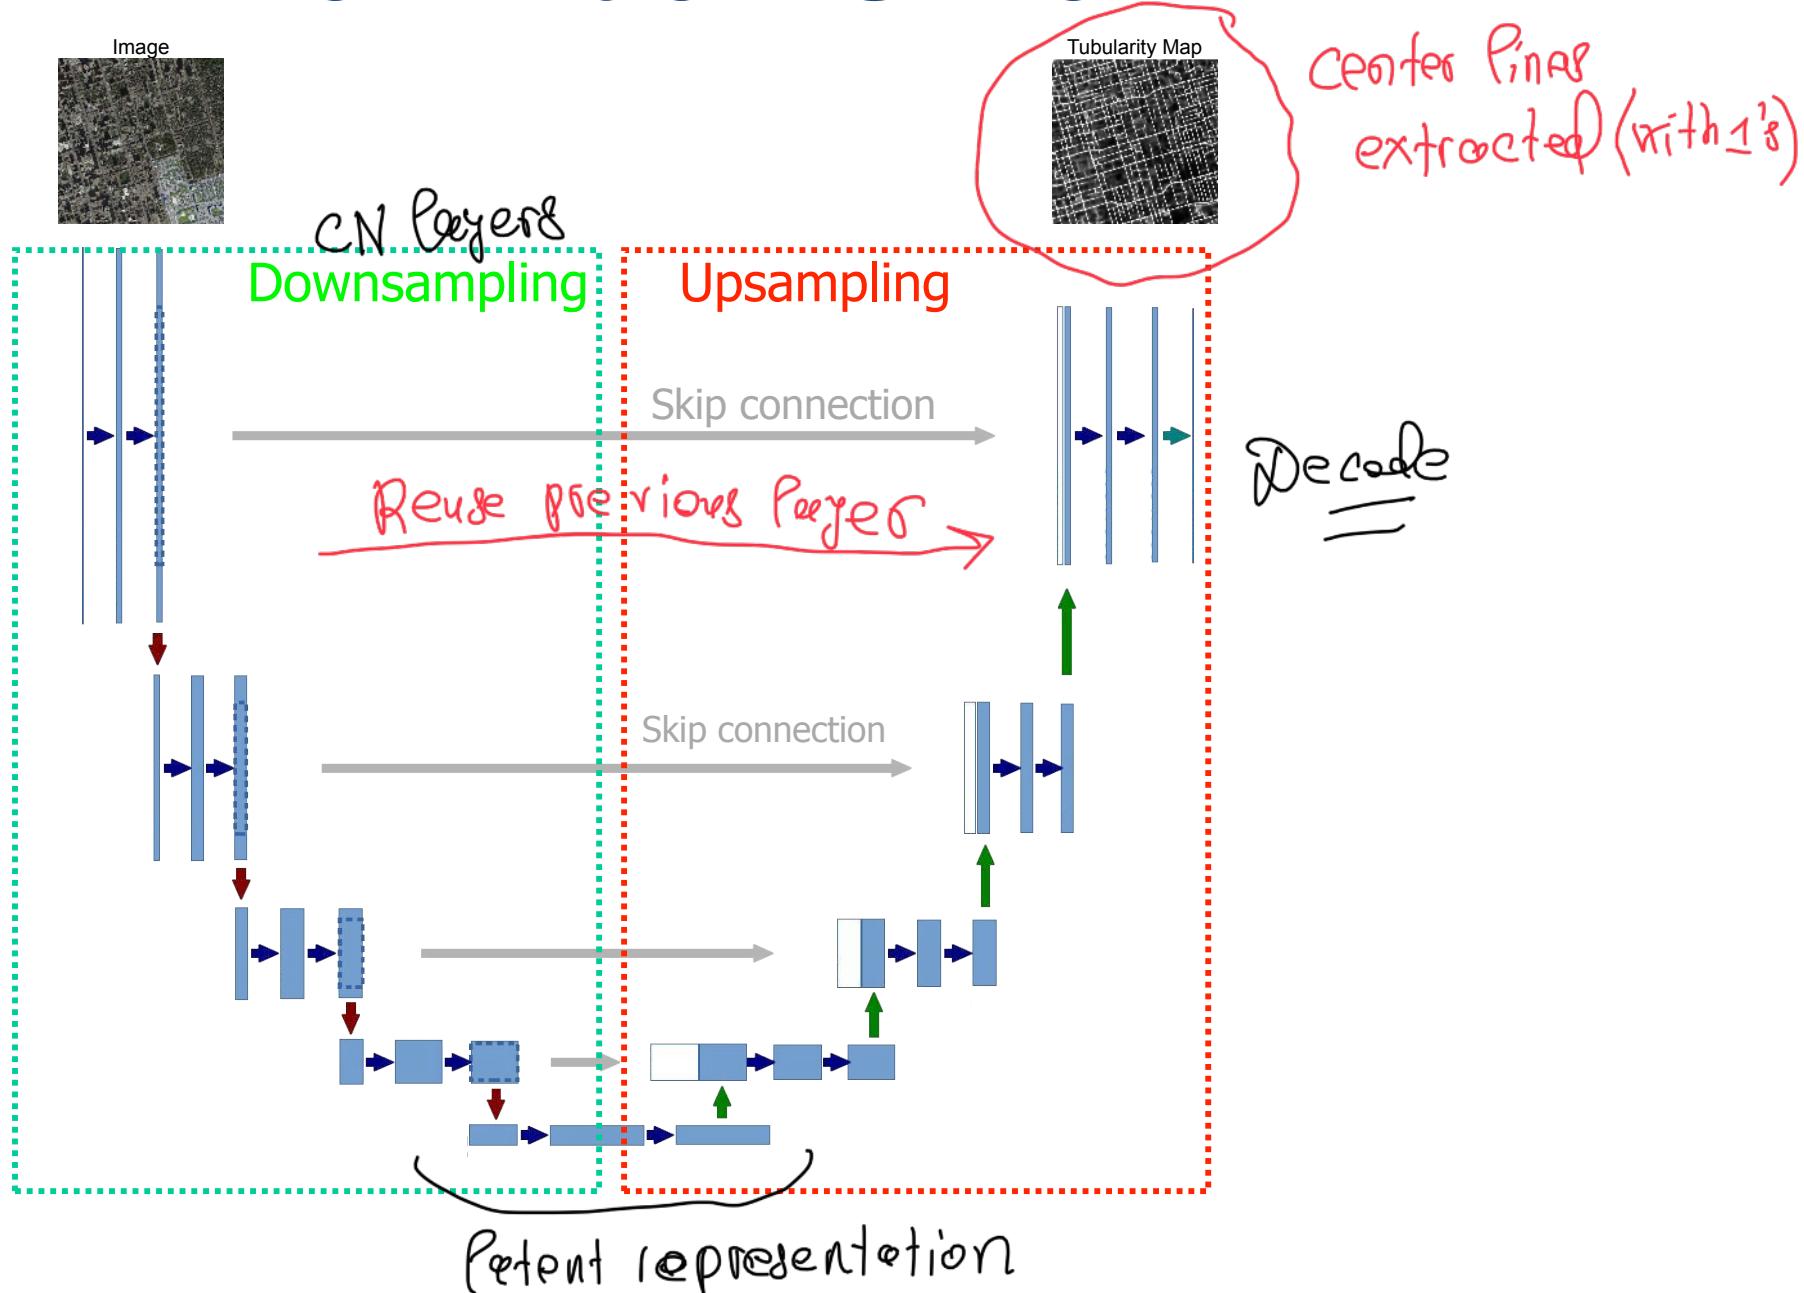
These two steps are closely related!



↑
Detecting
centerlines

↑
raiding whether
path is road/not

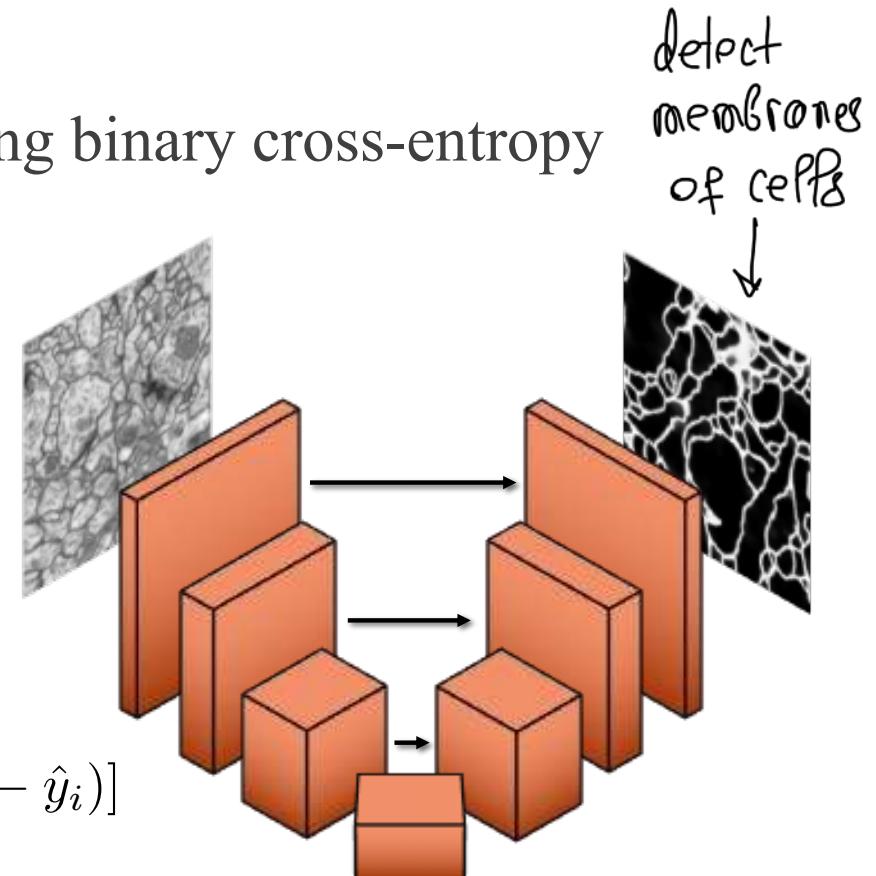
Reminder: U-Net



→ Train a U-Net to output a tubularity map.

Training a U-Net

Train Encoder-decoder U-Net architecture using binary cross-entropy



Minimize

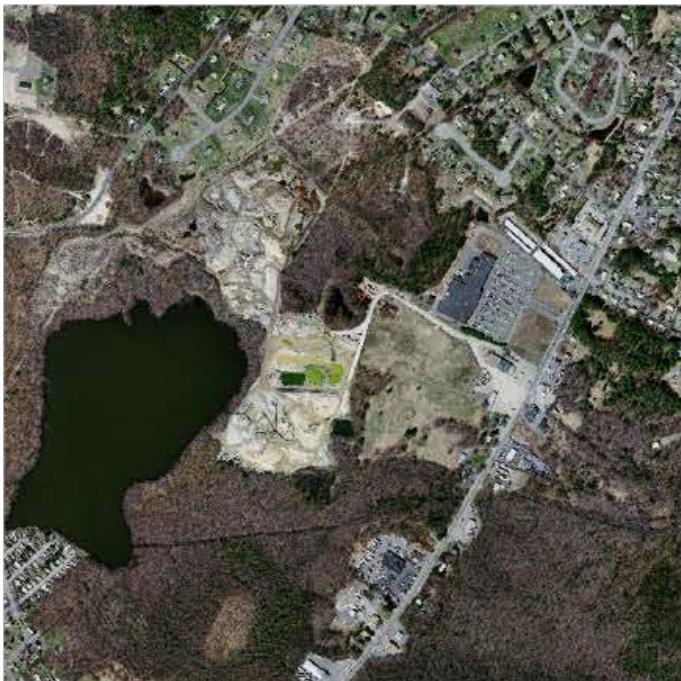
$$L_{bce}(\mathbf{x}, \mathbf{y}; \mathbf{w}) = -\frac{1}{i} \sum_1^P [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

where

- $\hat{\mathbf{y}} = f_{\mathbf{w}}(\mathbf{x})$,
- \mathbf{x} is an input image,
- \mathbf{y} the corresponding ground truth.

If you have centerline $\rightarrow 1$
(not) $\rightarrow 0$

Network Output



Image



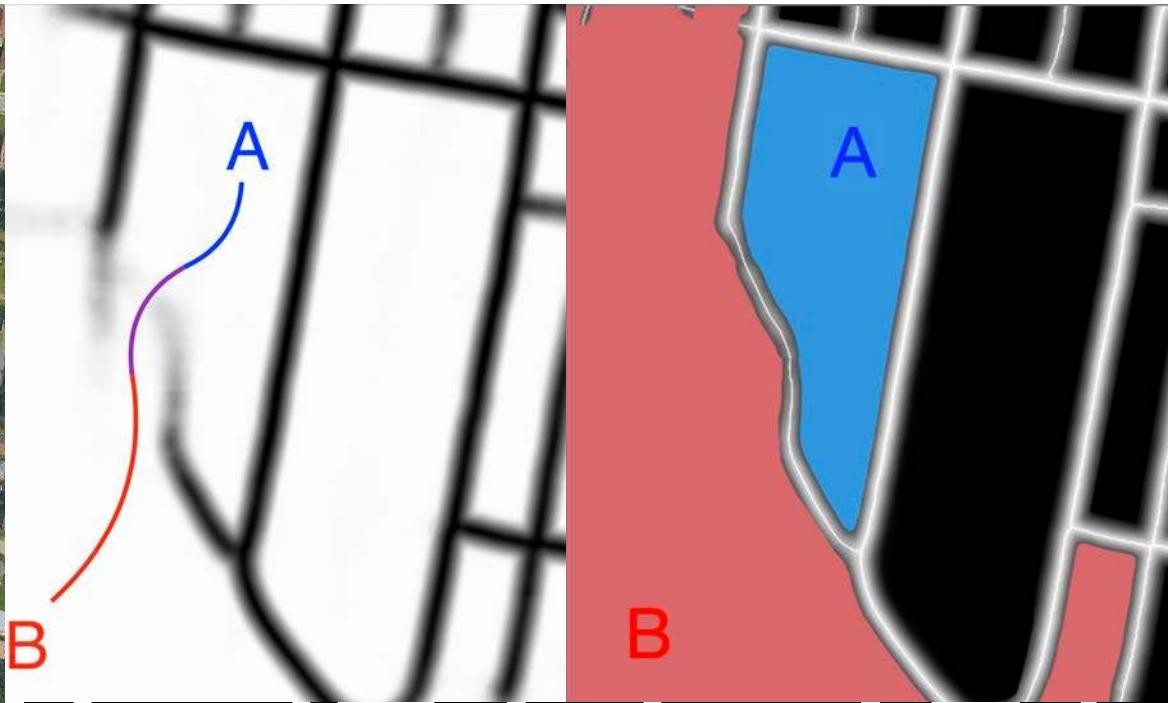
BCE Loss



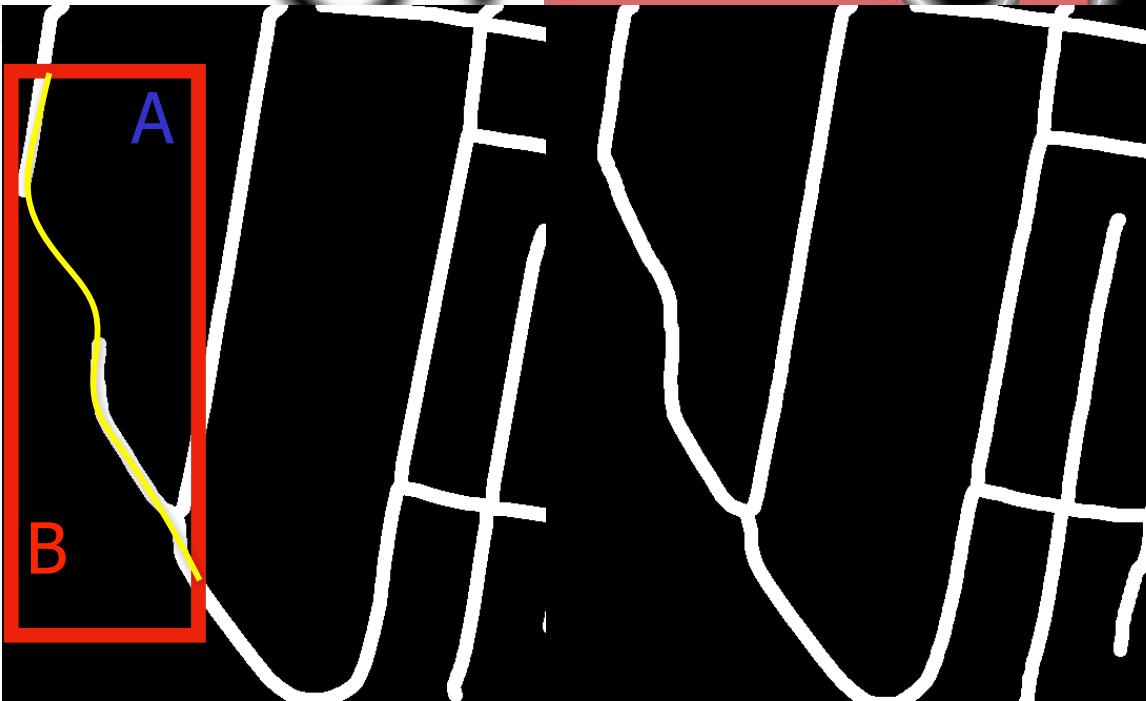
Ground truth

Accounting for Topology

→ Looking for specific pixels that have some form of Π



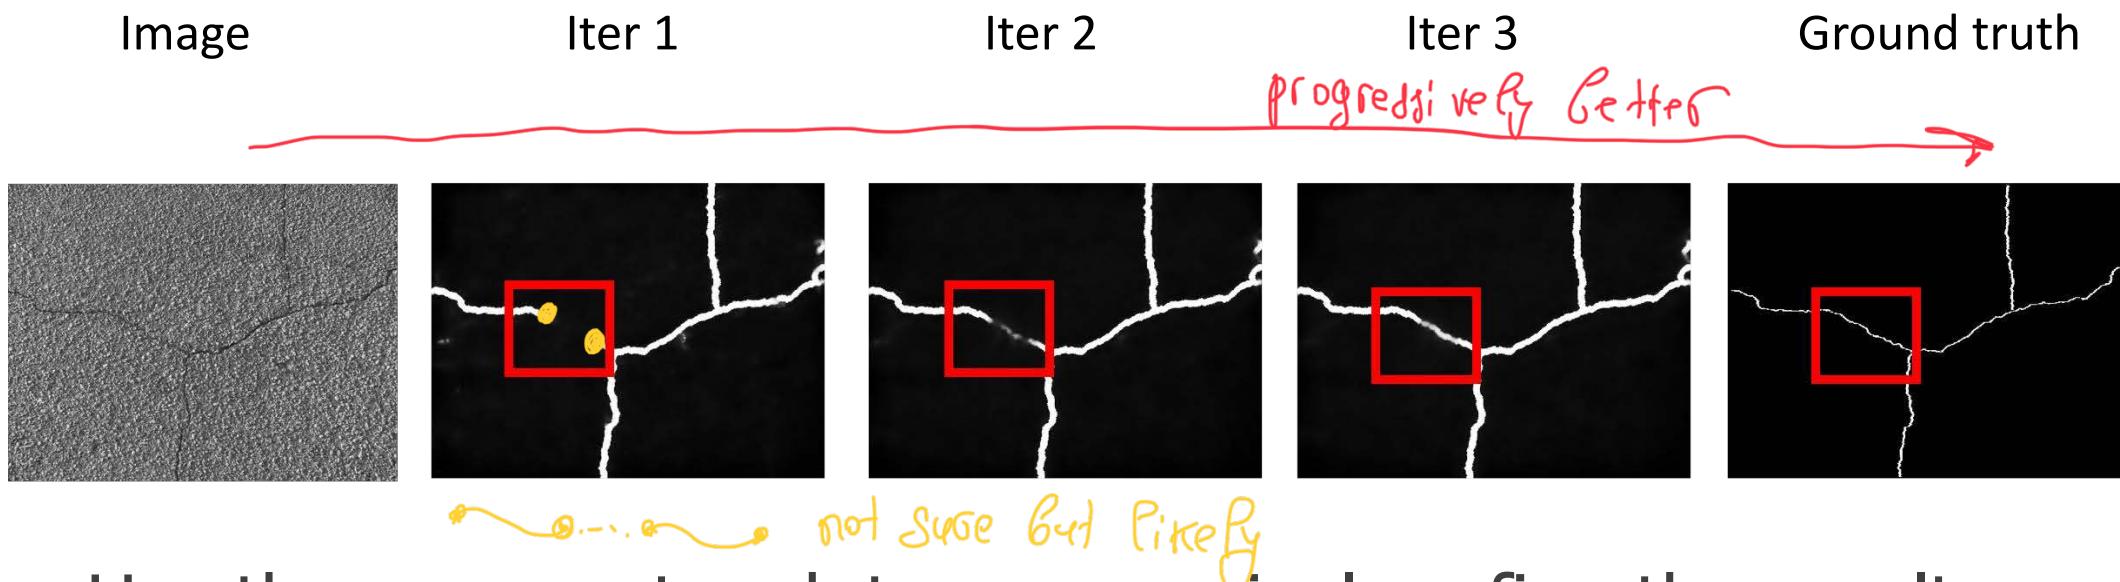
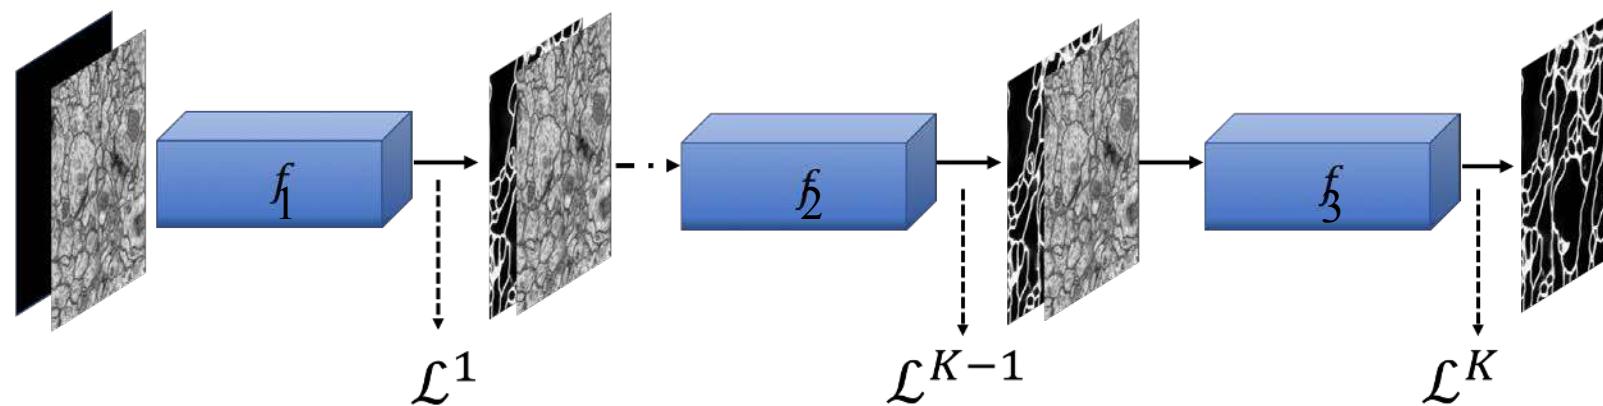
Onur et al., PAMI'21



- The yellow road is partially hidden by trees.
- A standard U-Net misses the hidden portion.
- We add to the loss function used to train the network a term that encourages points such as A and B to be separated.
- The re-trained U-Net now finds the complete road.

BCE and
Topology loss

Iterative Refinement



Use the same network to progressively refine the results keeping the number of parameters constant

Delineation Steps

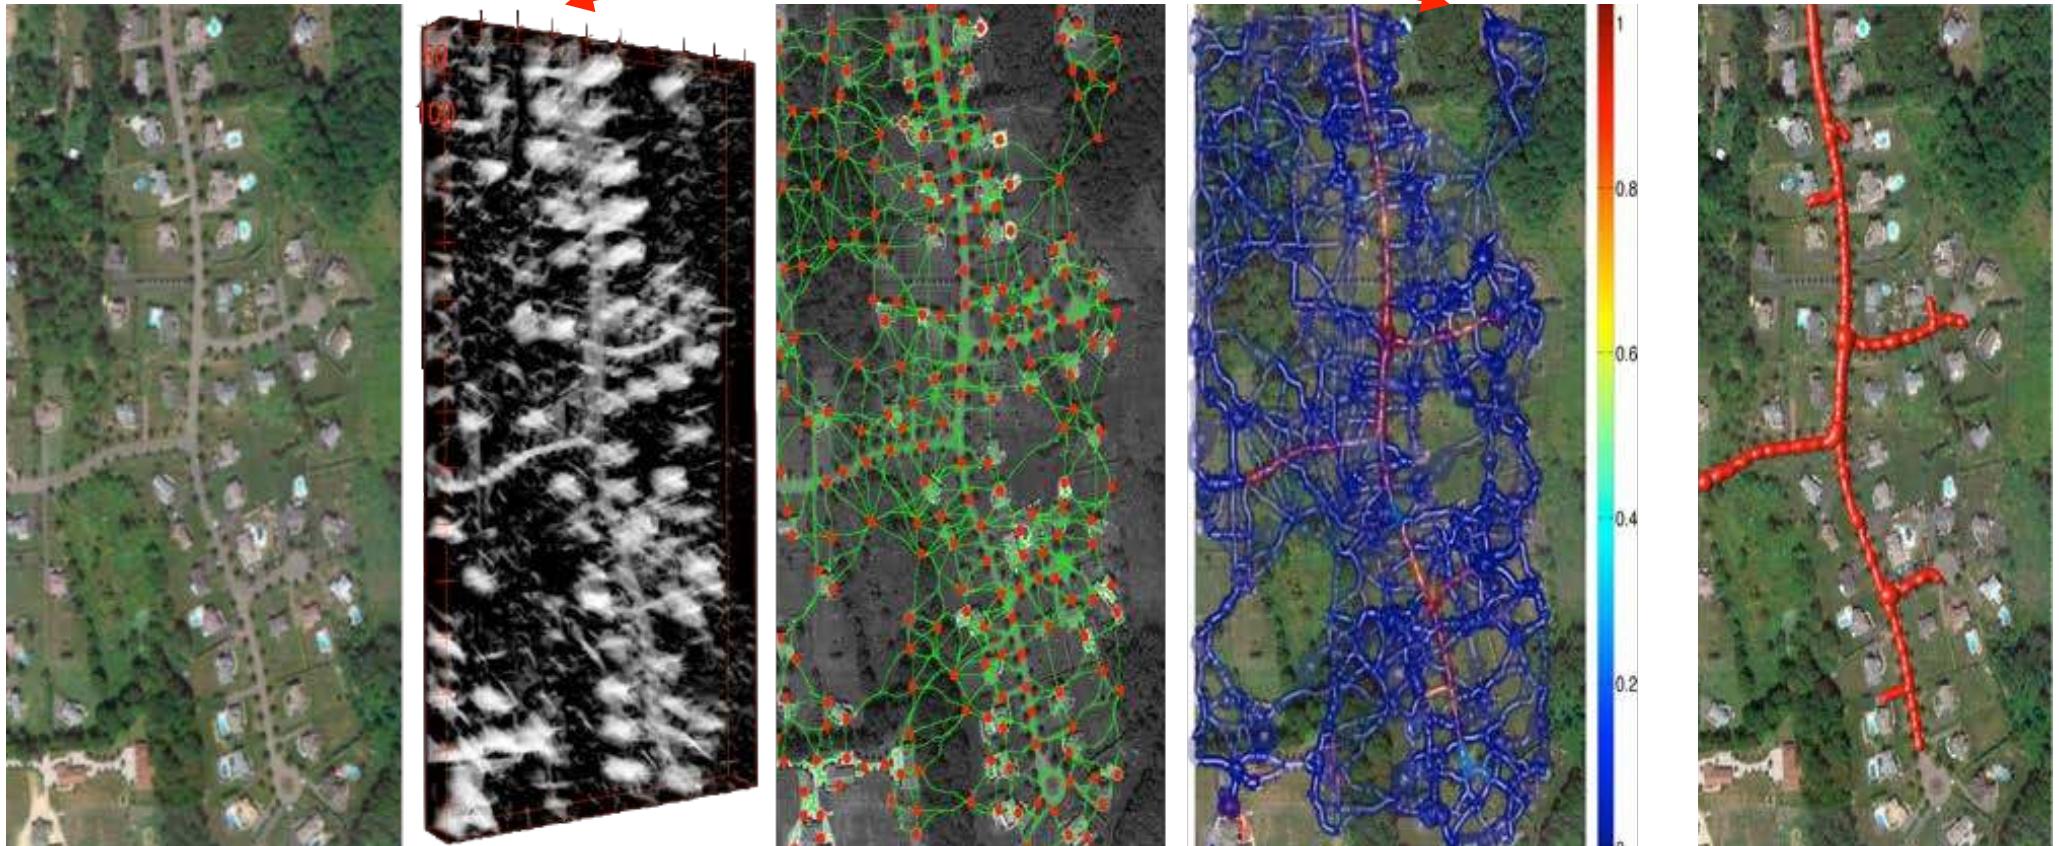


1. Compute a probability map.
2. Sample and connect the samples.
3. Assign a weight to the paths.
4. Retain the best paths.

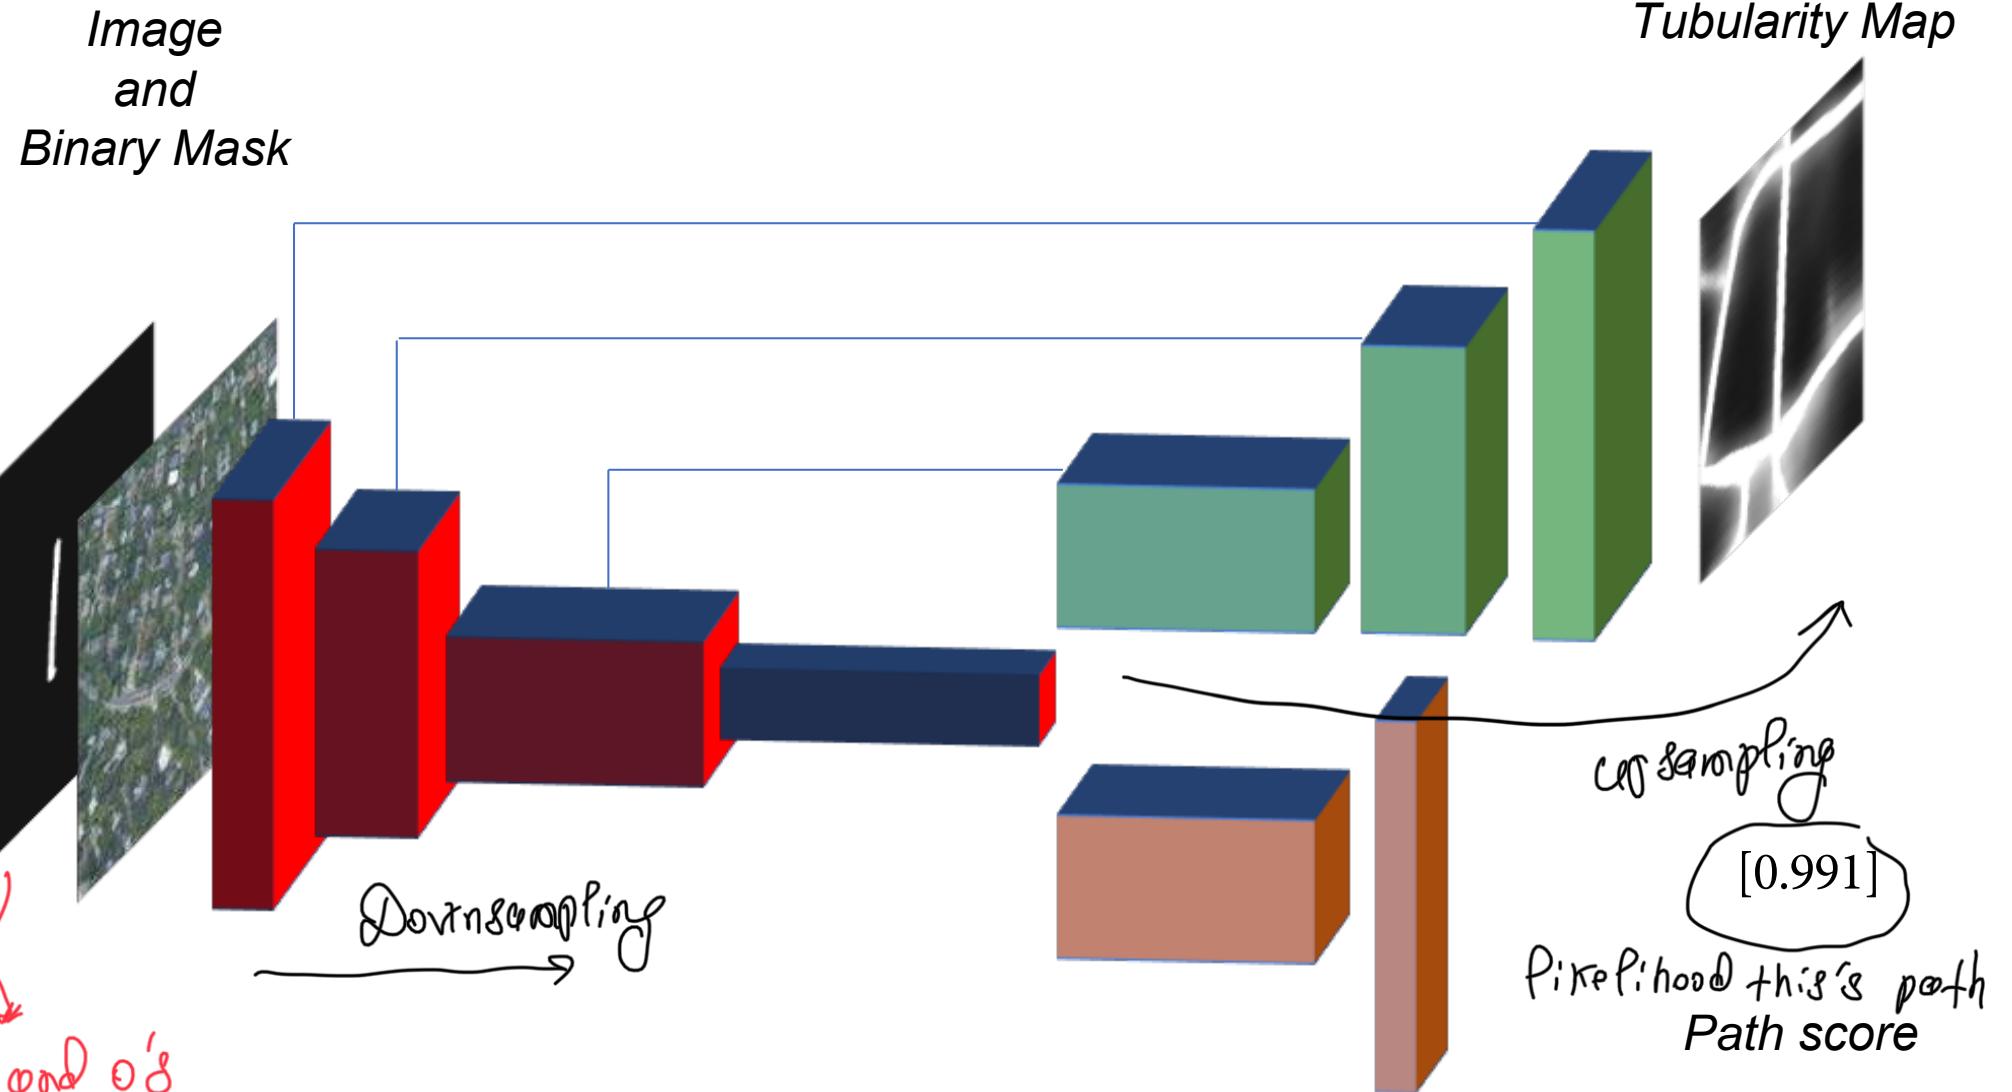
Iterative refinement:
→ Done both at training and testing

Delineation 2019

These two steps are performed by the same network



Dual Use U-Net



Streets Of Toronto

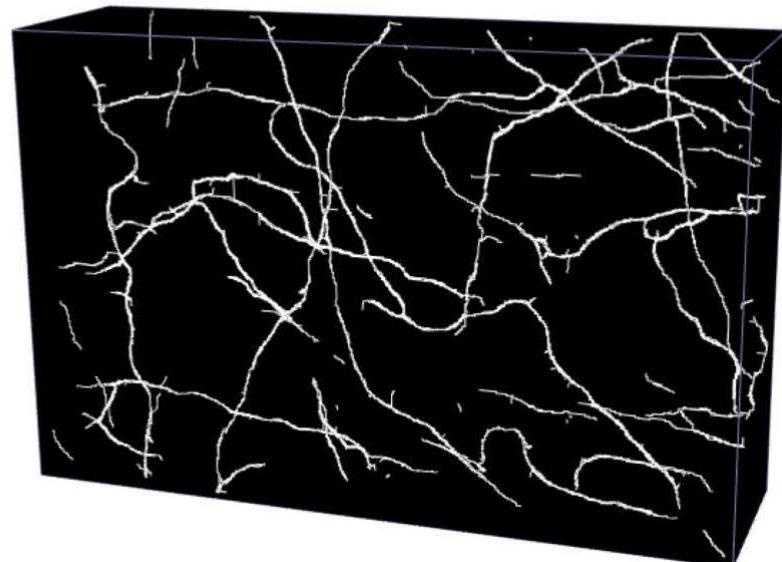
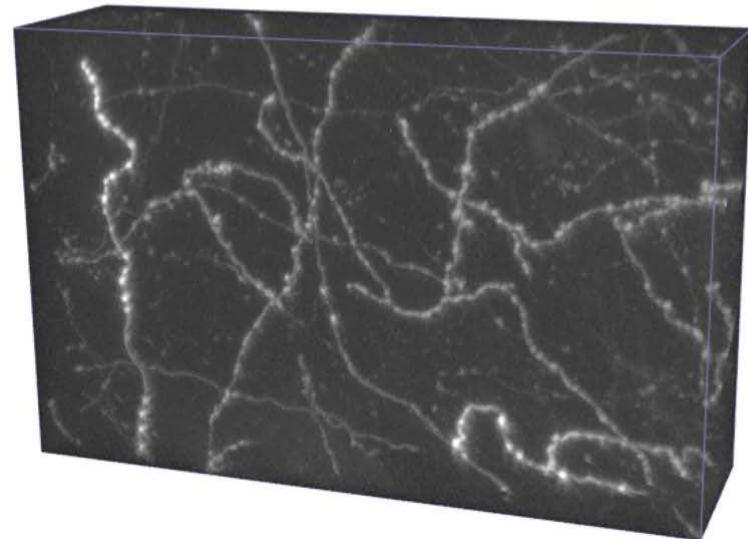
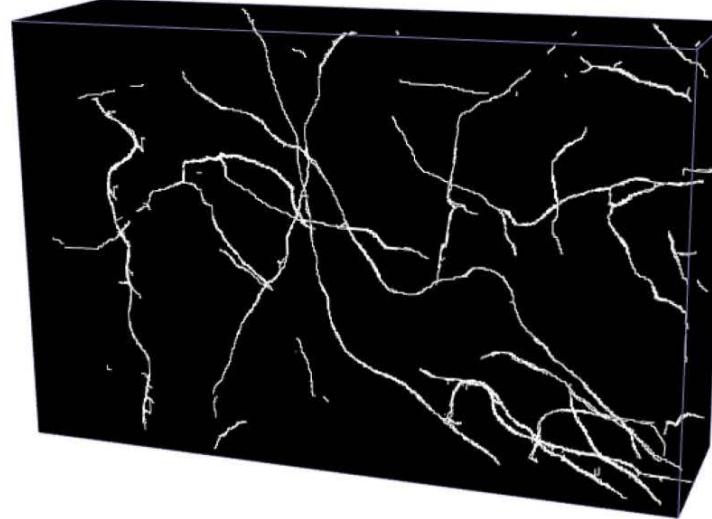
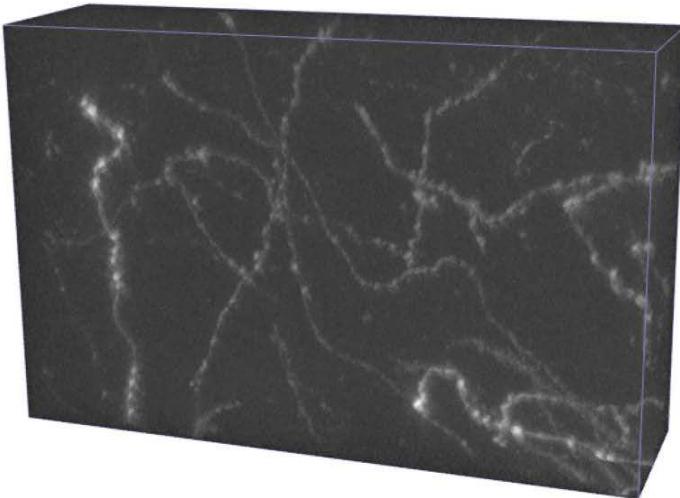


green → good
others

- False negatives
- False positives

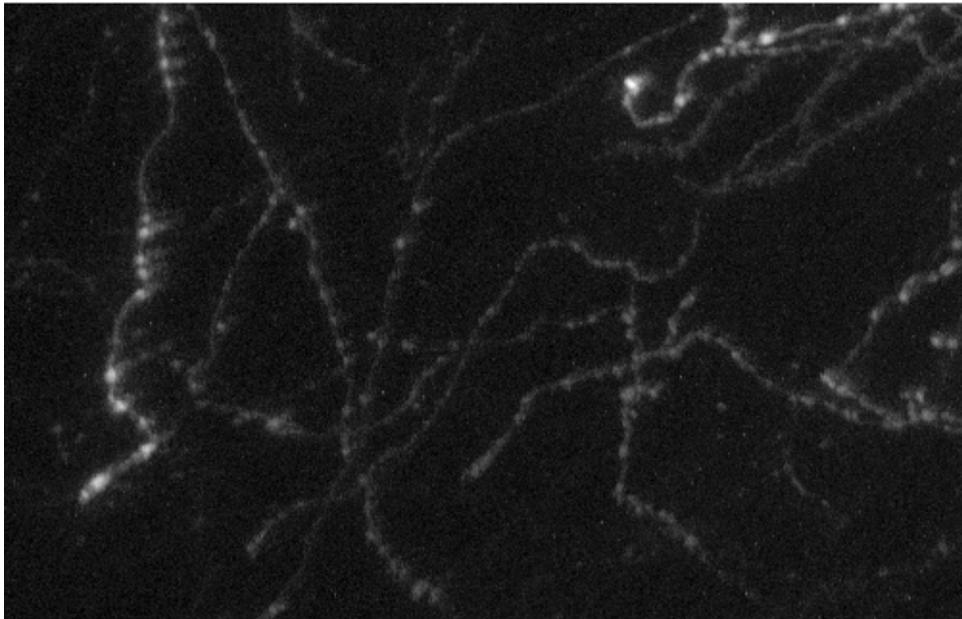
Dendrites And Axons

Generic
==

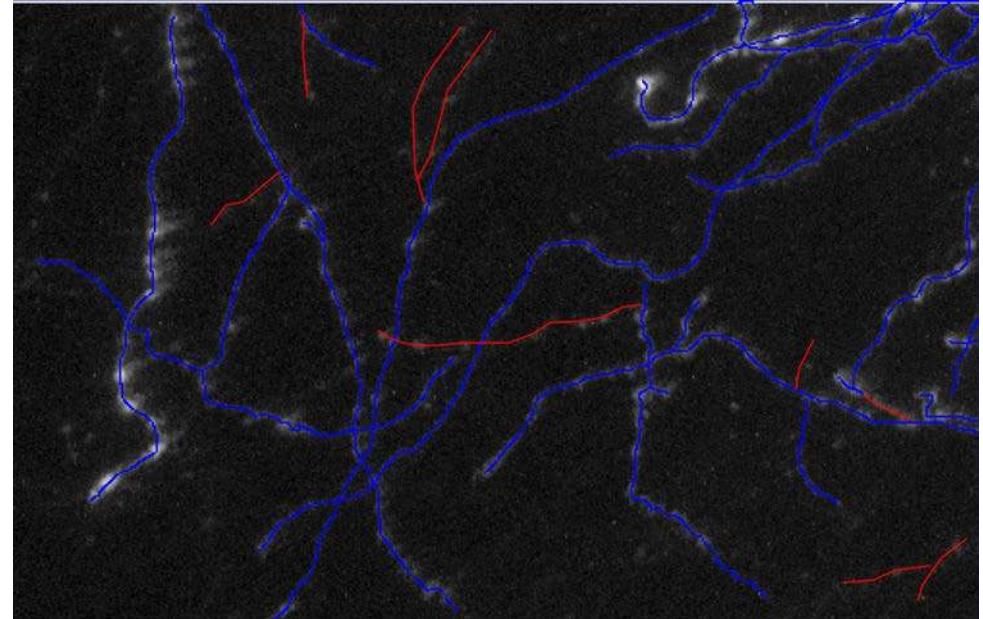


Just train differently

Typical Annotations



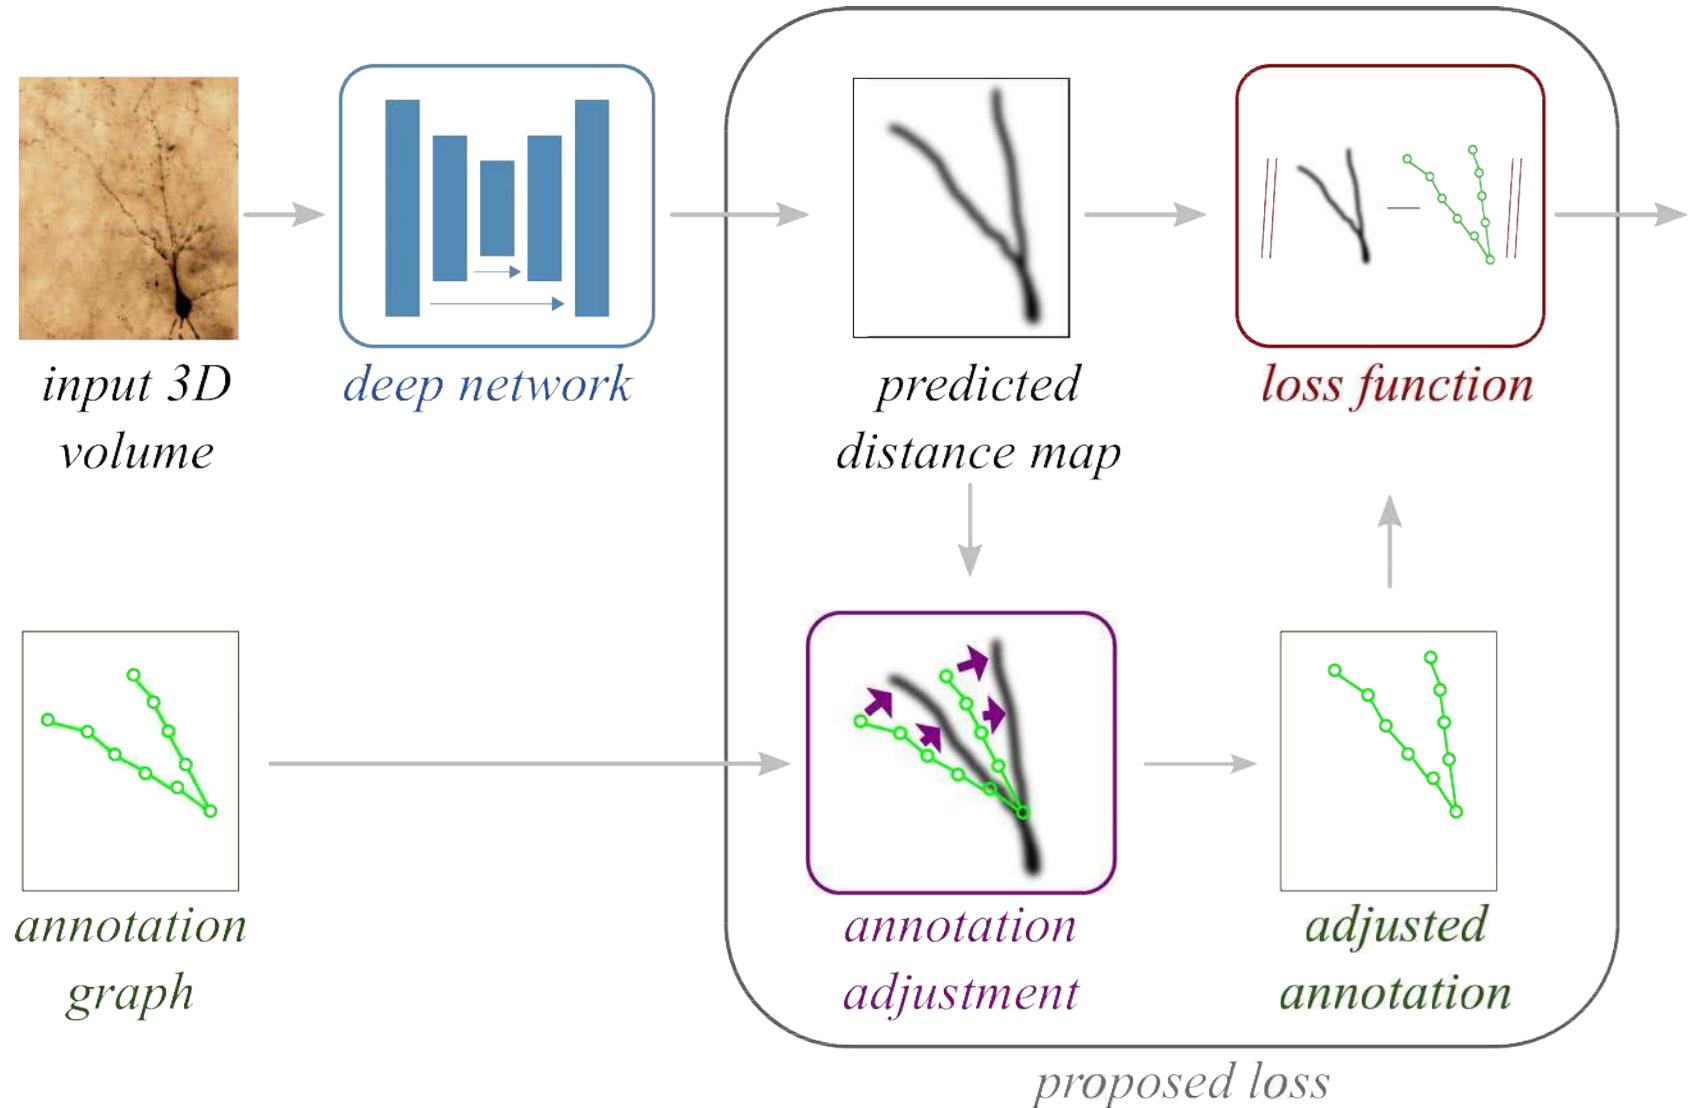
Original Image



“Ground truth” + **Mistakes**

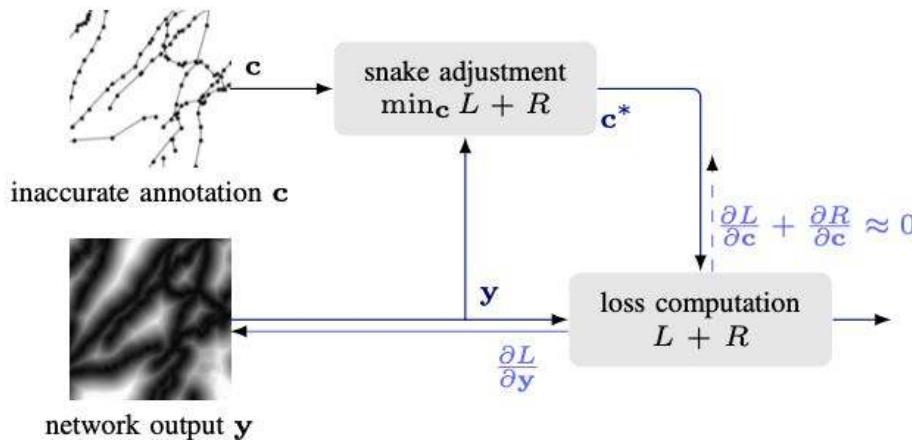
→ Human annotations are often imprecise.

Correcting the Annotations



To account for annotation inaccuracies during training, we jointly train the network and adjust the annotations while preserving their topology.

Annotations as Network Snakes



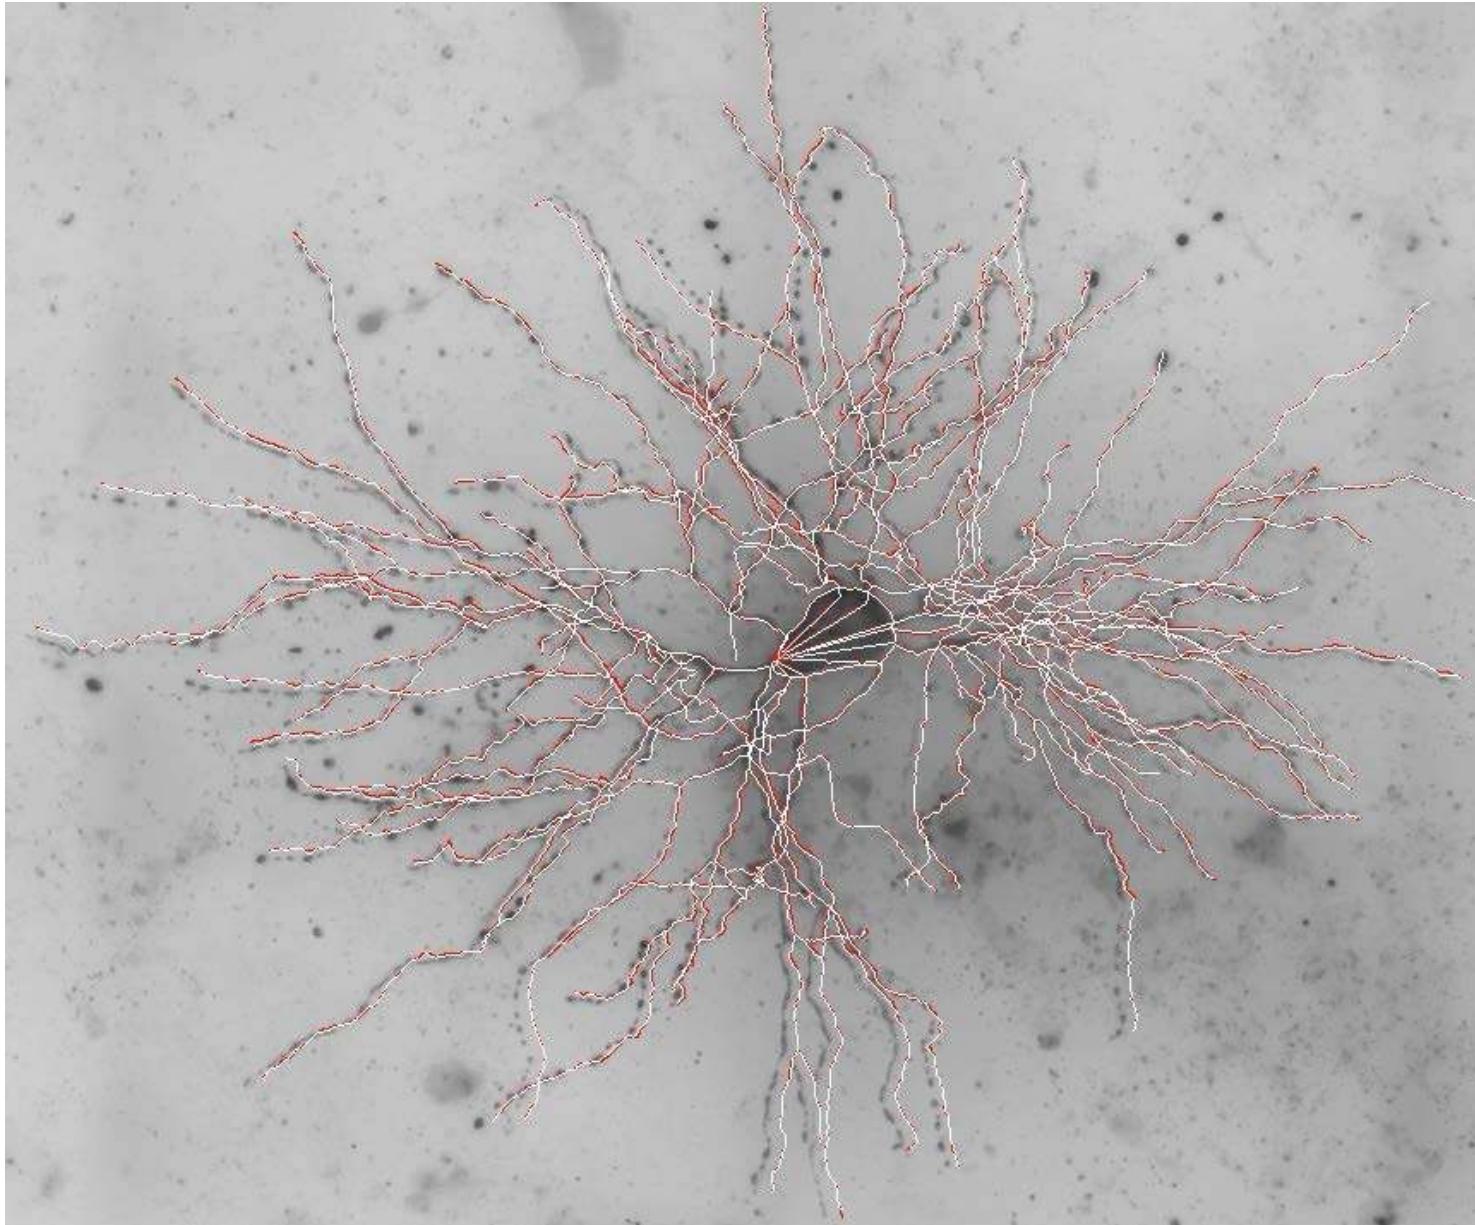
$$\Theta^*, \mathbf{C}^* = \operatorname{argmin}_{\Theta, \mathbf{C}} \sum_{i=1}^N \mathcal{L}(D(\mathbf{c}_i), \mathbf{y}_i) + R(\mathbf{c}_i)$$

where

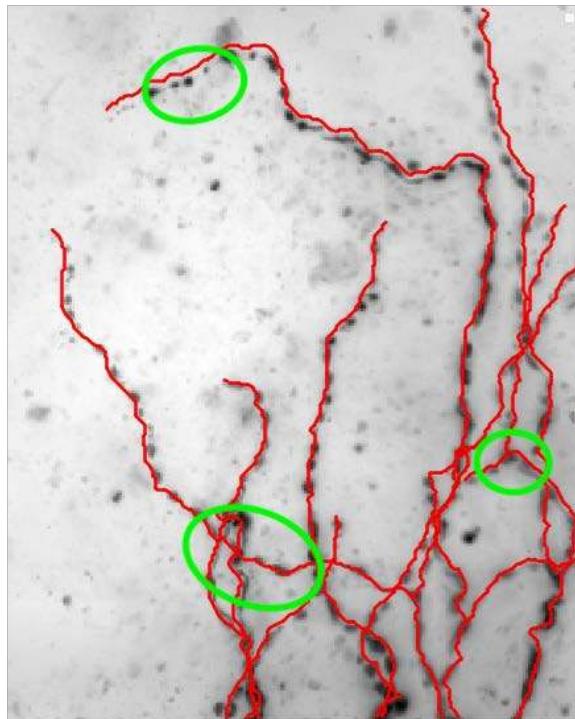
Distance between network
output and annotations.

- The y_i are the network outputs;
- The \mathbf{c}_i are the annotation vertices;
- \mathbf{C} is the vector obtained by concatenating all the \mathbf{c}_i ;
- D is a distance transform ;
- \mathcal{L} is the MSE loss;
- R is a regularization term.

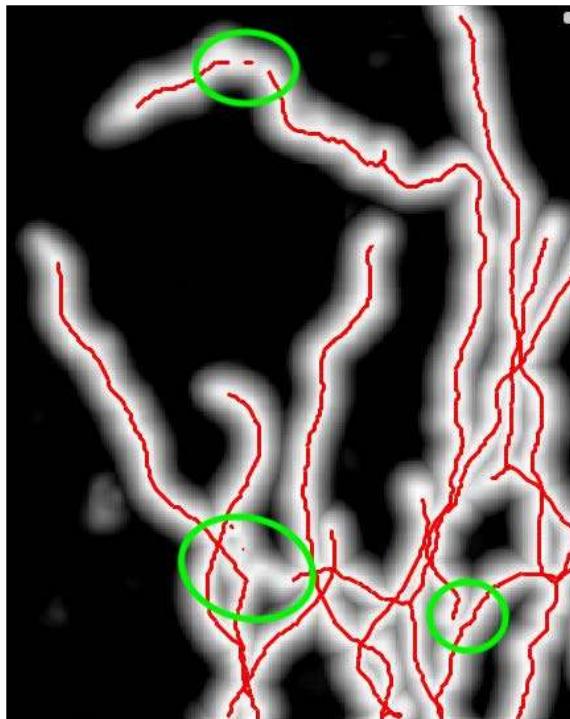
Snake Optimization



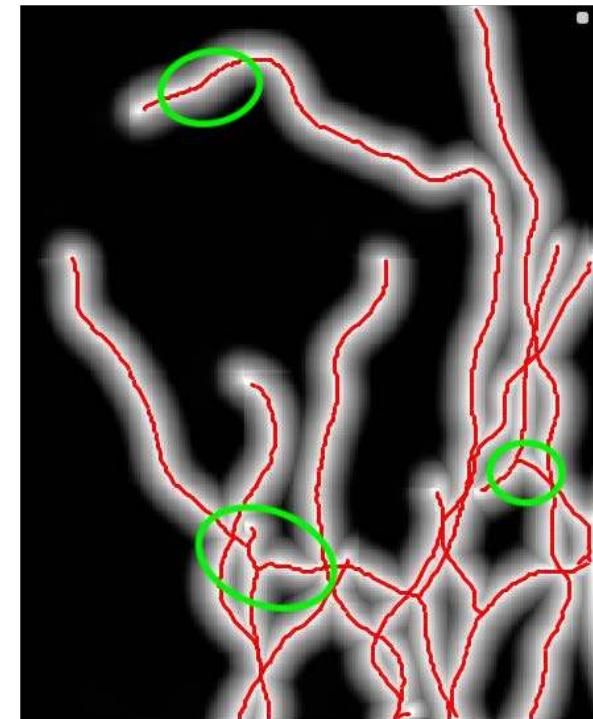
Improved Results



Annotated image



Vanilla U-Net



Network snakes

On the Job

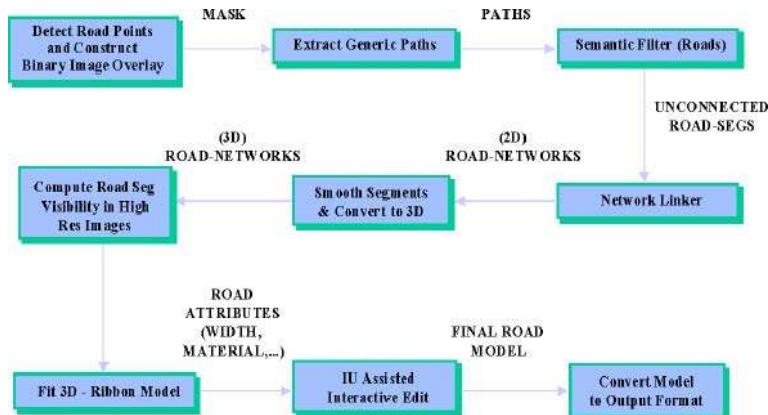


The new job is good, we do of course a lot of Deep Learning, but **also some good old-school computer vision** e.g. registration 😊 So the material from Computer Vision class is definitely helpful and I wouldn't change it to another all-Deep Learning class (even in the light of today's Turing Award).

Best,

Agata

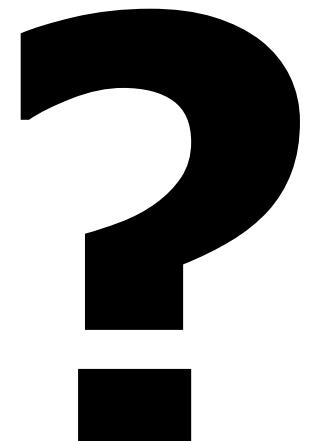
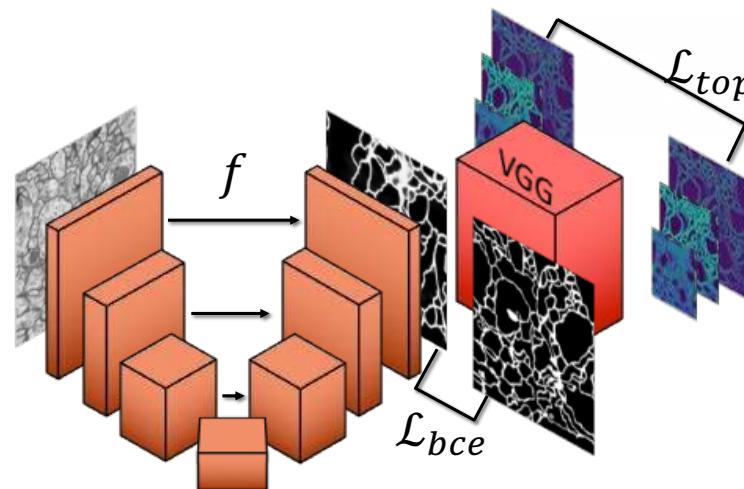
1998 - 2038



1998

2018

2038



It is difficult to make predictions, especially about the future.
Sometimes attributed to Niels Bohr.

In Short

- Edge and image information is noisy.
 - Models are required to make sense of it.
out of it
- An appropriate combination of graph-based techniques, machine learning, and semi-automated tools is required.