

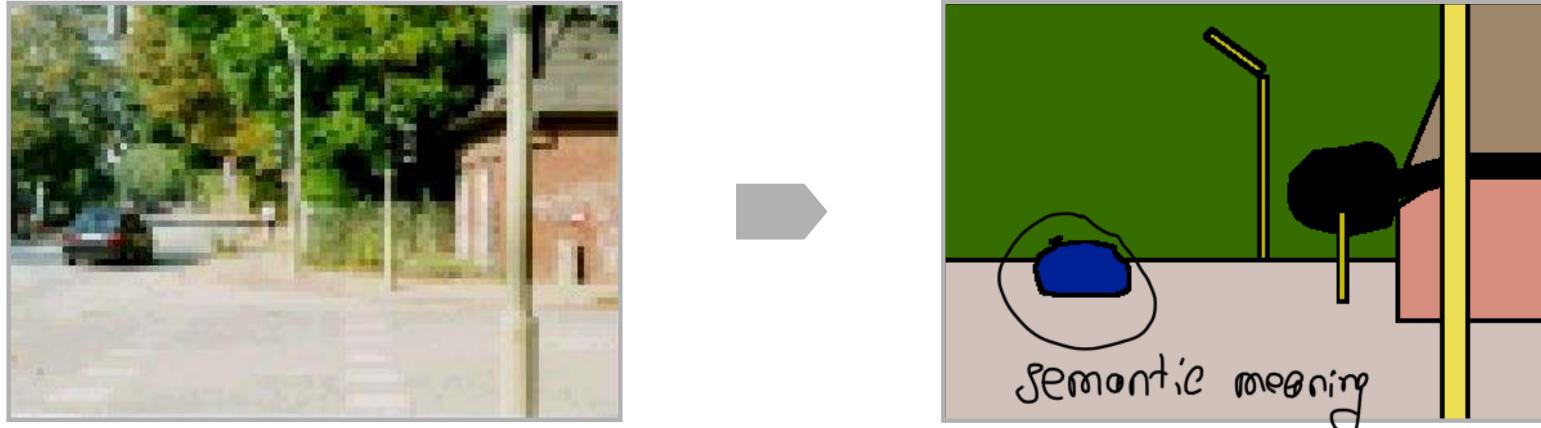
Regions

→ correspond to
meaningful
semantic object



Pascal Fua
IC-CVLab

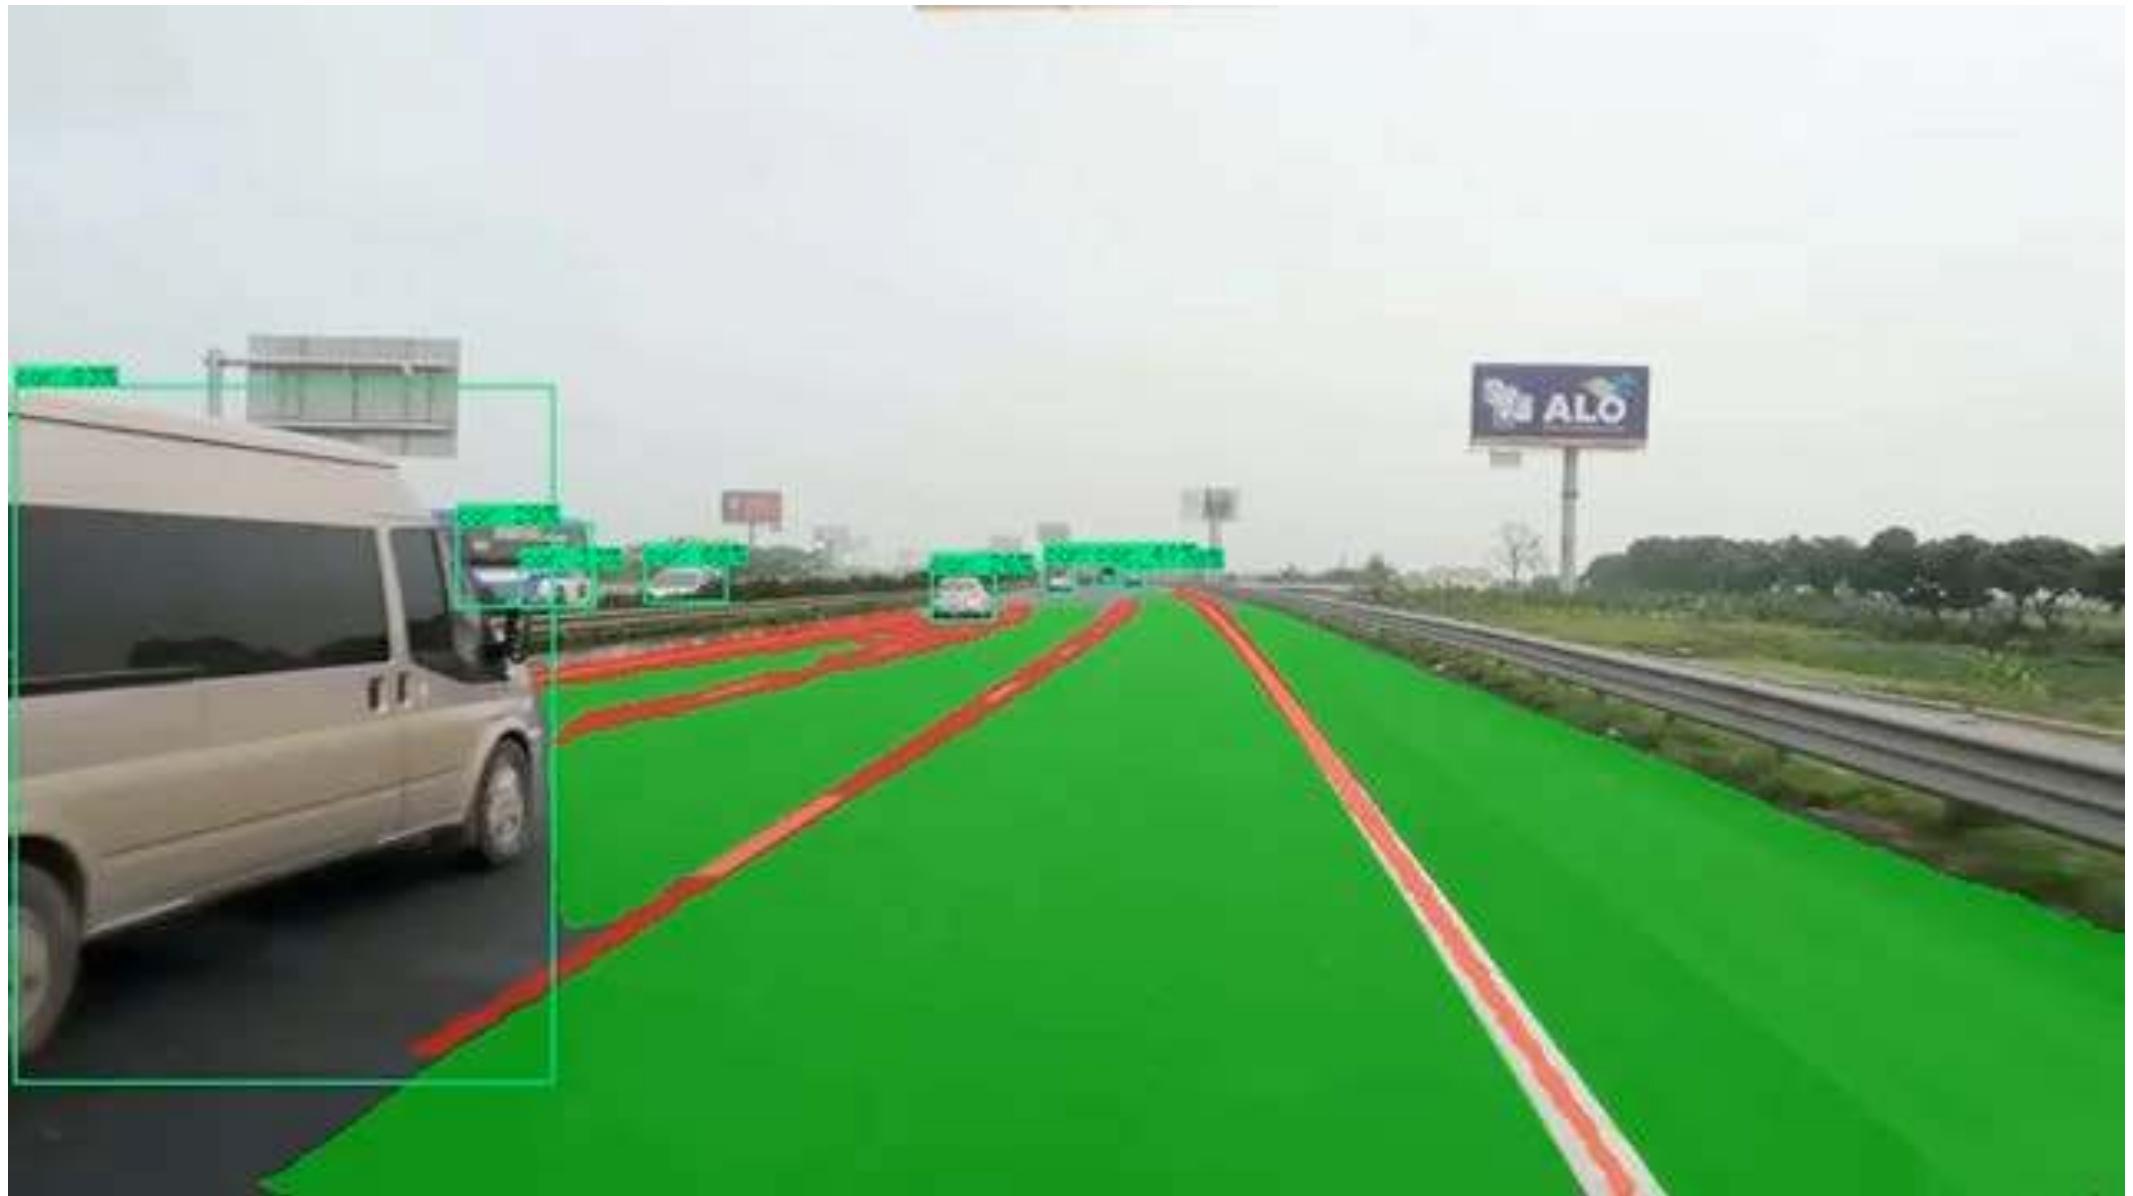
Region Segmentation



Ideal region: Set of pixels with the same statistical properties and corresponding to the same object.

Purpose: Should help with recognition, tracking, image database retrieval, and image compression among other high-level vision tasks.

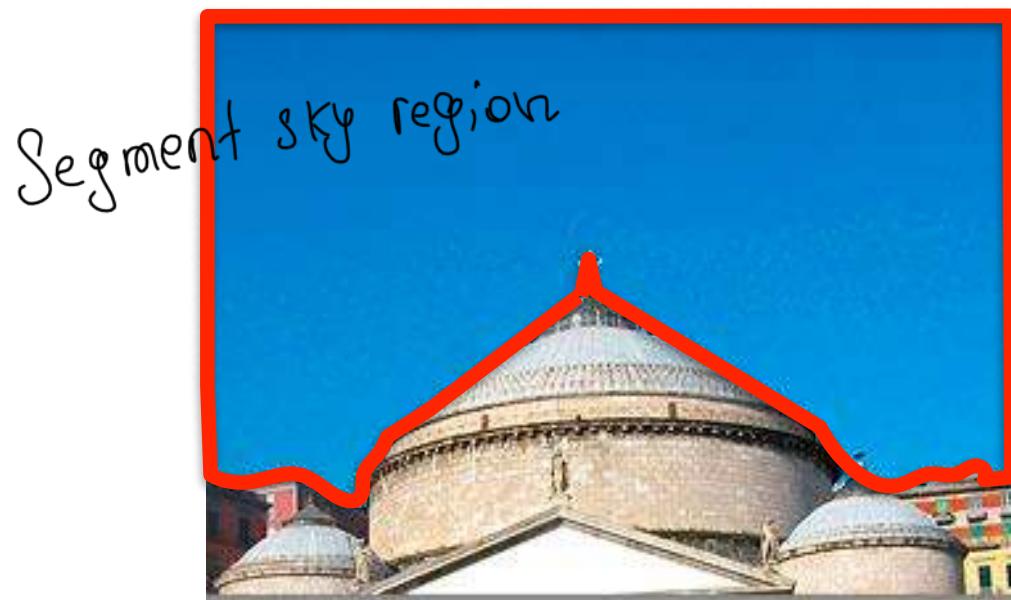
Application: Automated Driving



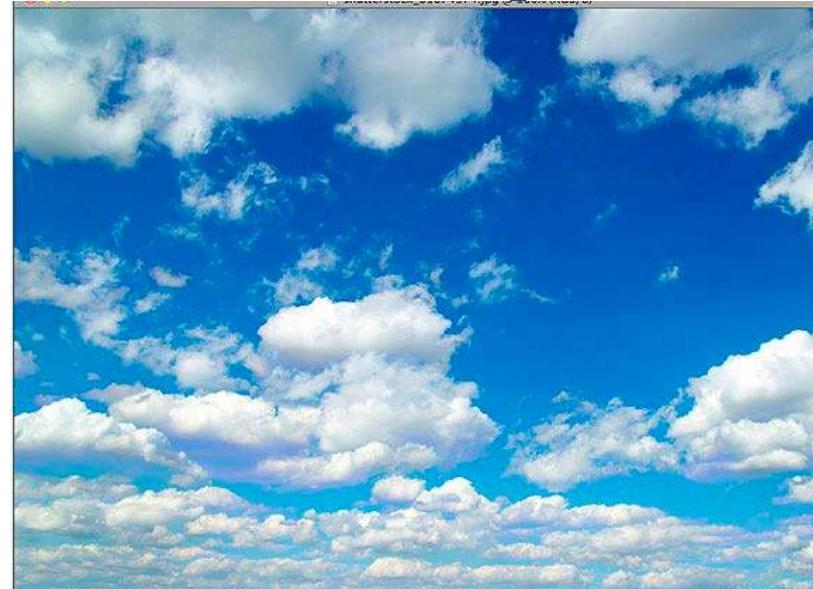
Applications: Photoshop



I find this blue sky too bland!



Find the sky region.



I prefer this one.



Replace it.

In Theory

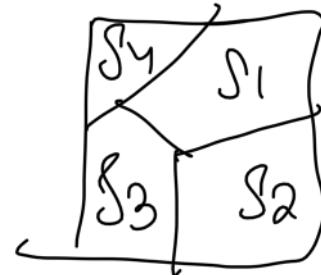
Look for an image partition such that:

$$I = \bigcup_{i=1}^m S_i$$

$$S_i \cap S_j = 0, \forall i \neq j$$

$$H(S_i) = \text{True}, \forall i \leftarrow \begin{array}{l} \text{homogeneously uniform} \\ \text{in some sense} \end{array}$$

$$H(S_i \cup S_j) = \text{False}, \text{ if } S_i \text{ and } S_j \text{ are adjacent.}$$



where H measures homogeneity.

Complex Gray Level Variations



do not correspond
to free

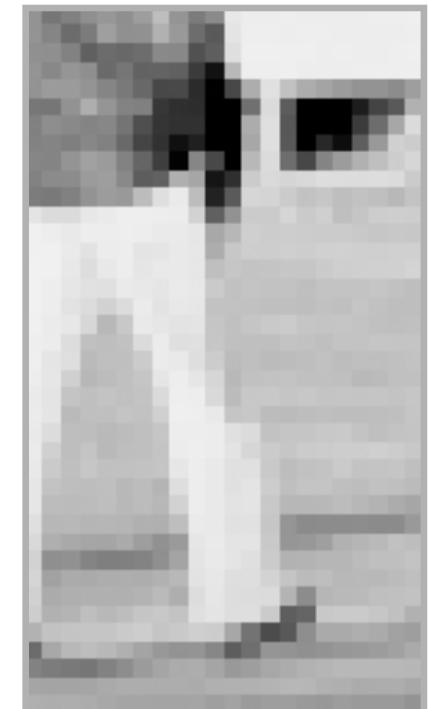
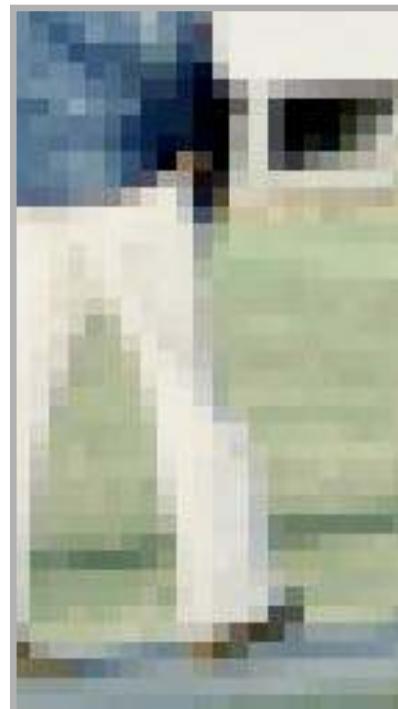
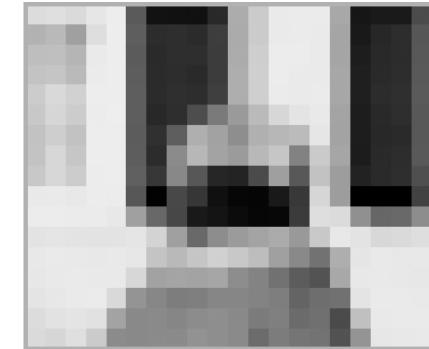
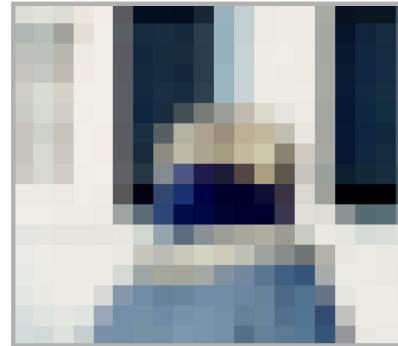
- Simple thresholding and other basic image operations do not suffice.
- The H predicate is difficult to define.

does not work



Context is Essential

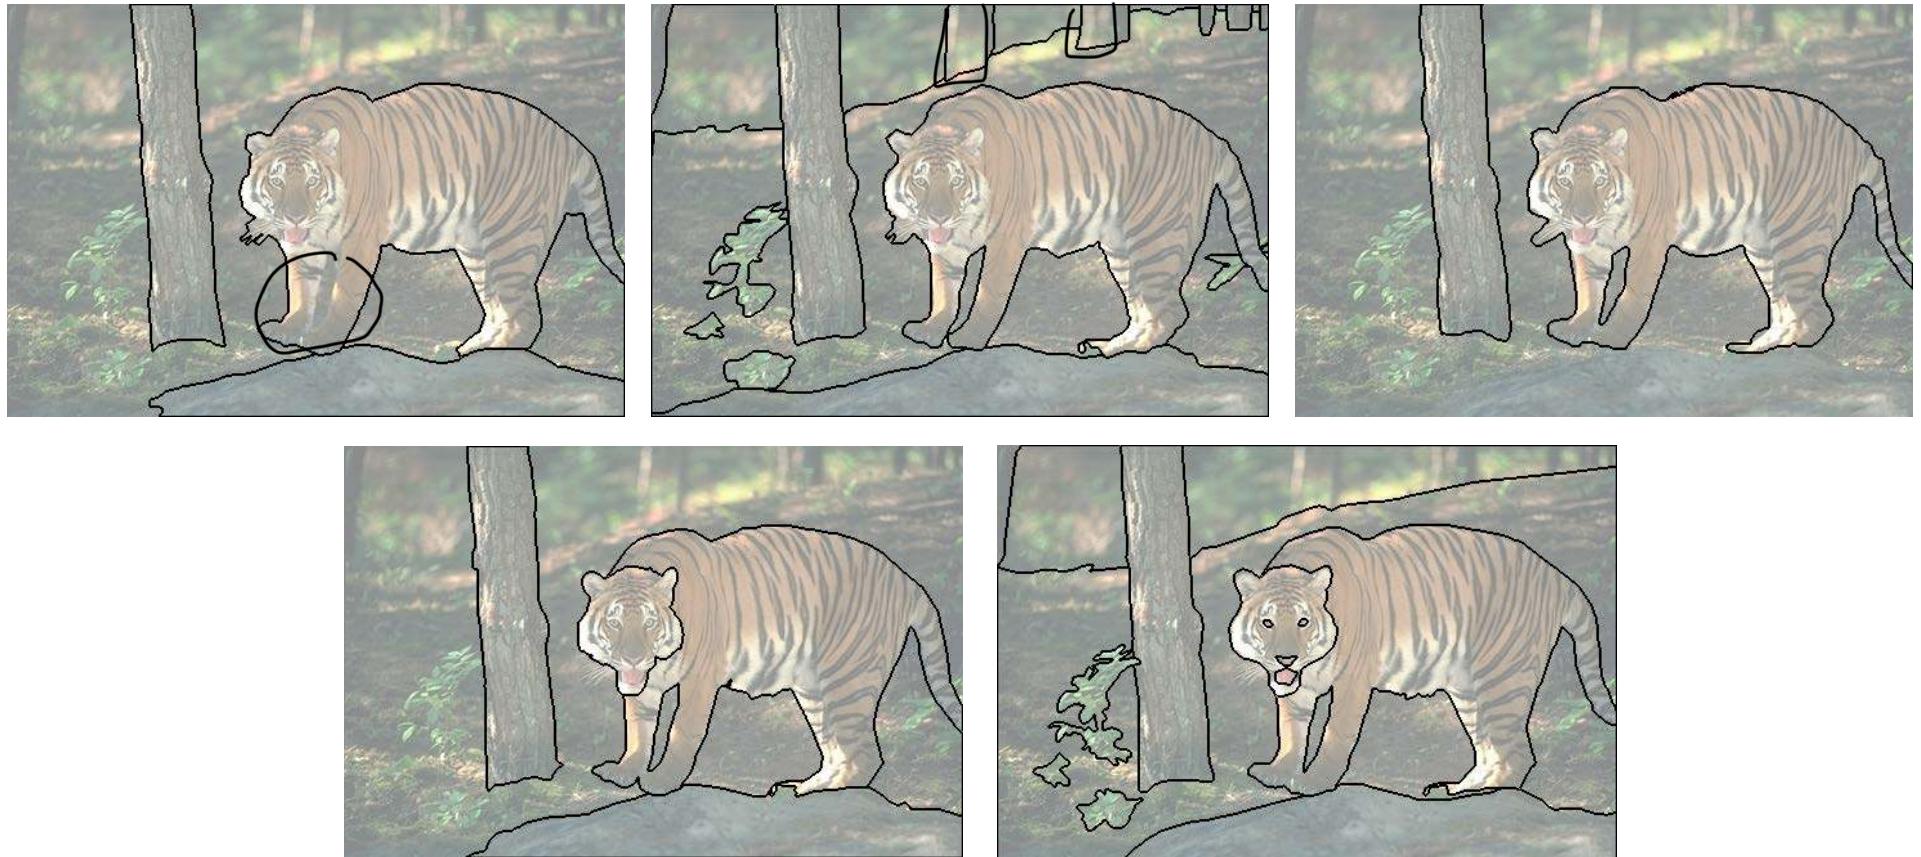
What is H?



Without context, it's hard to ^{segment}

Without the whole image it is hard to make sense of small image windows.

There is not always a Single Answer



- Segmentations hand-drawn by 5 different people.
purpose?
- We cannot say that one is right and the others wrong.
No absolute definition on segmentation

Homogeneous or Not?



In terms of color → ✗
Statistics → ✓
(orange with
stripes)

texture



What is homogeneous in some parts of these images are the statistical properties, not the actual pixel values.

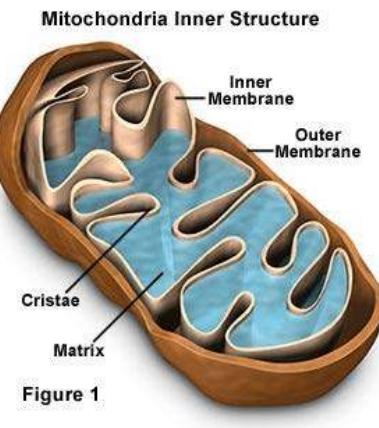
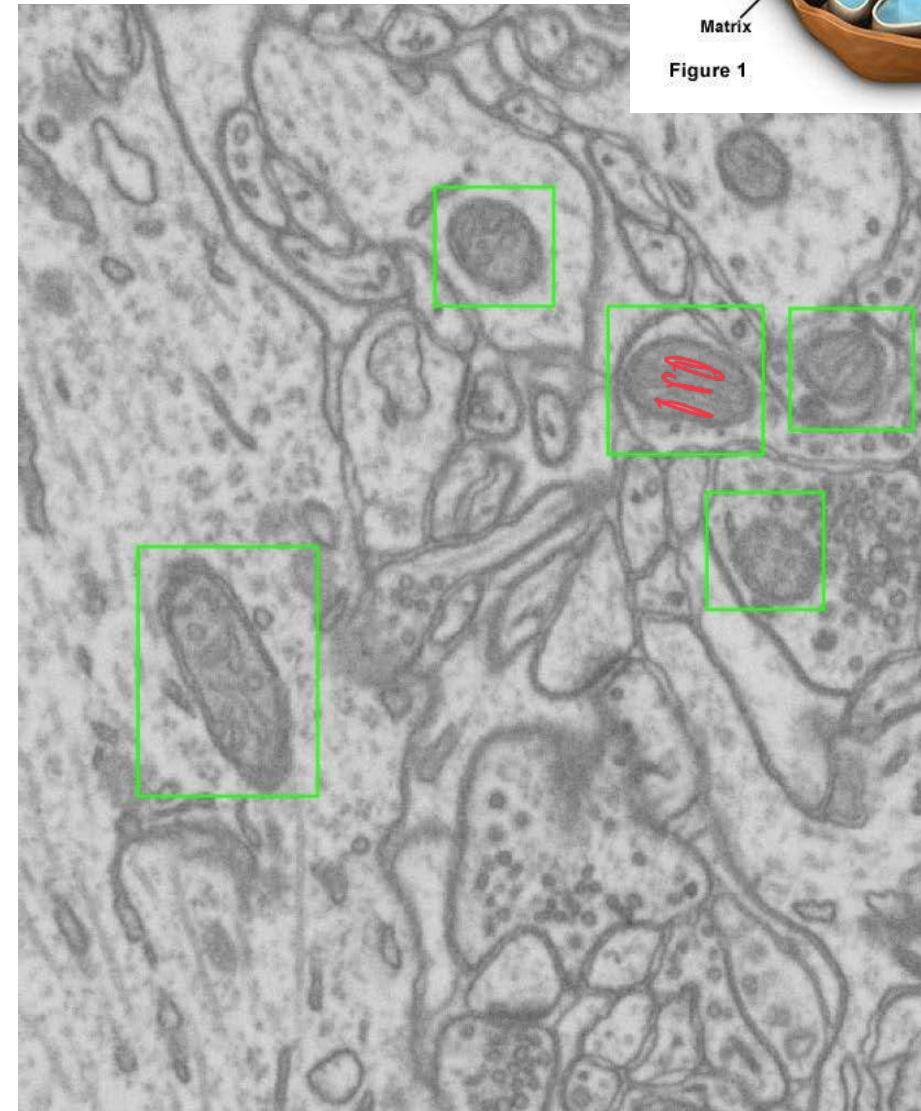
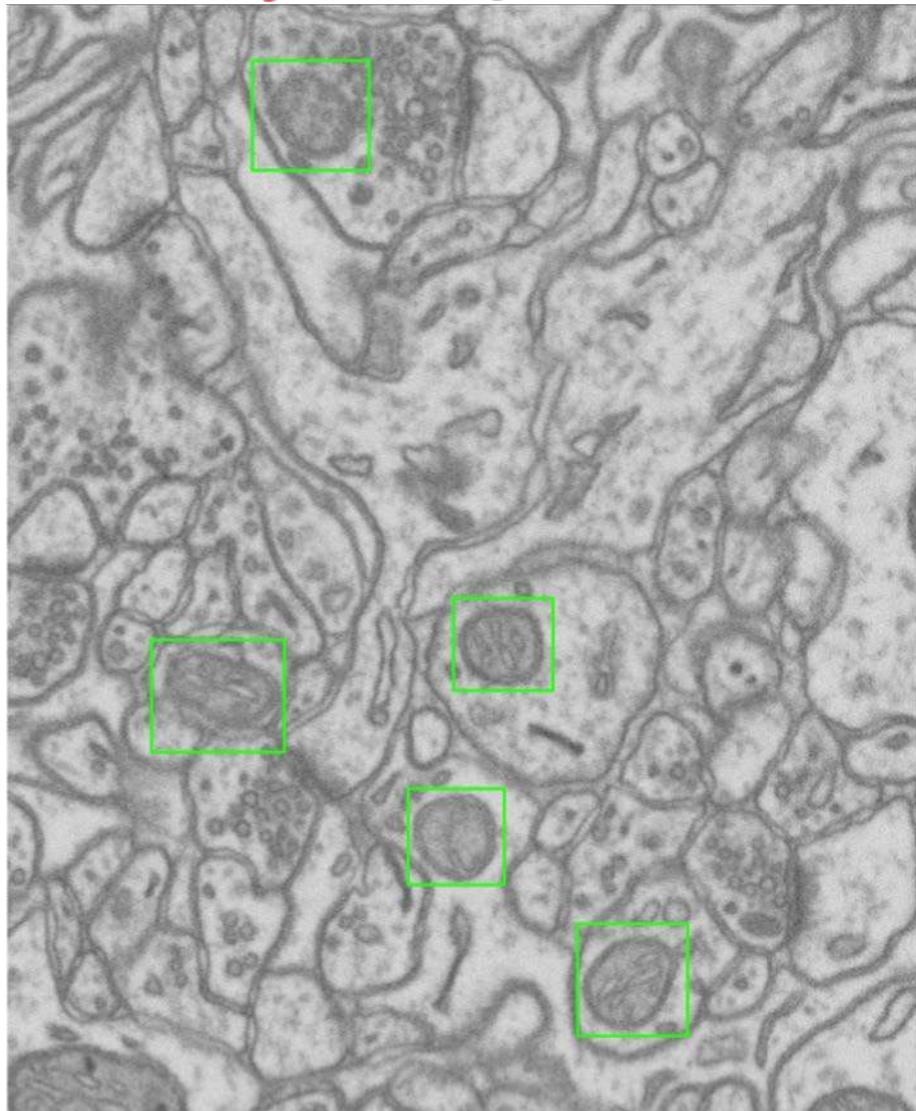


orientation of gradients

Mitochondria

homogenous, segment ✓

textual information



In Theory

Merge:

- Start with a partition that satisfies Eq. 3.
- Satisfy Eq. 4 by merging regions.

Split:

- Start with a partition that satisfies Eq. 4.
- Split regions until they all satisfy Eq. 3.

Homogeneity:

- Uniform gray-level or color statistics.
- Regions to which a parametric surface can be fitted.

From Simple to Complex Algorithms

- Region Growing.
- Histogram Splitting.
- K-Means.
- Graph Theoretic Methods.
- Convolutional Neural Nets.

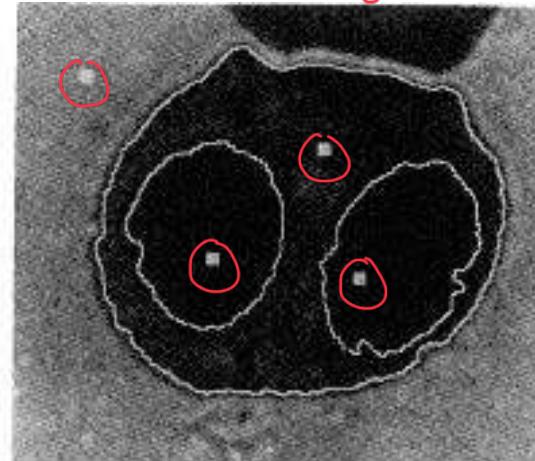
Region Growing

slightly darker

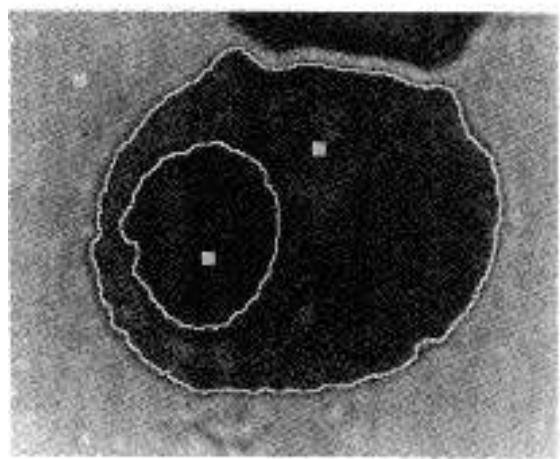


(a)

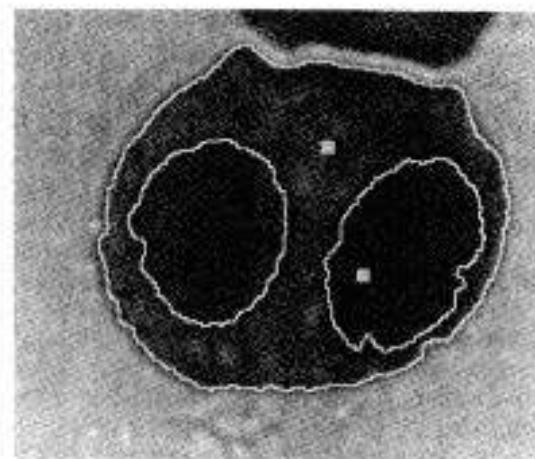
give seeds



(b)



(c)

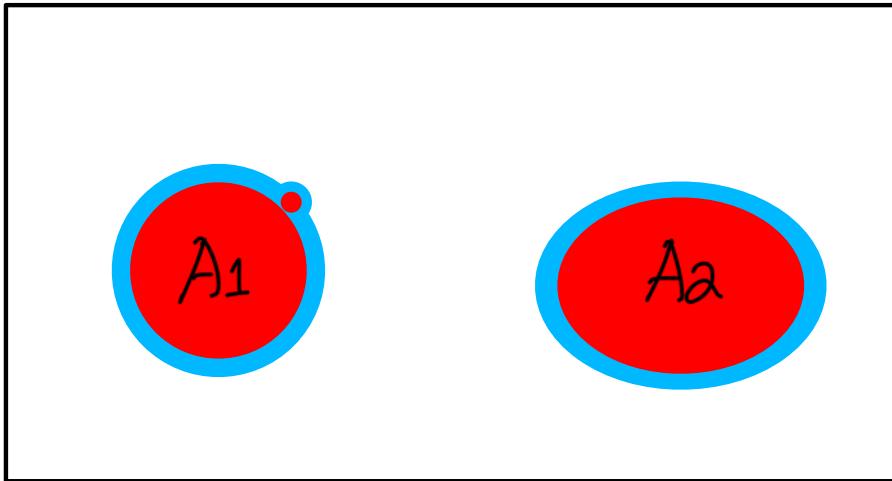


(d)

Interactive Segmentation of a Cell

Region Growing

progressively
adding neighbors



- Labeled pixels.
- Unlabeled pixels.

Given a set of regions $\underbrace{A_1, \dots, A_n}$, let
 $\underbrace{\text{seed regions}}$

$$T = \{x \notin (\cup_{i=1}^n A_i), N(x) \cap (\cup_{i=1}^n A_i) \neq \emptyset\}, \text{ unlabeled pixels}$$

the set of unlabeled pixels that are neighbors of already labeled ones and d be a metric, such as
immediate

$$\delta(x) = |g(x) - \text{mean}_{y \in A_{i(x)}}[g(y)]|.$$

Until all pixel are labeled: \rightarrow gray-level of x

1. Represent T as a sorted list, the SSL, according to this metric.
2. Label the **first point** in T . $\not\equiv T$ sorted list
3. Add its the neighbors to the SSL.

\rightsquigarrow fast

Region Growing

While SSL is not empty do

 Remove first pixel y from SSL.

If all already labeled neighbors of y, other than boundary pixels, have the same label

then

 Set y to this label.

 Update running mean of corresponding region.

 Add neighbors of y that are neither already set nor already in the SSL to the SSL according to their distance value.

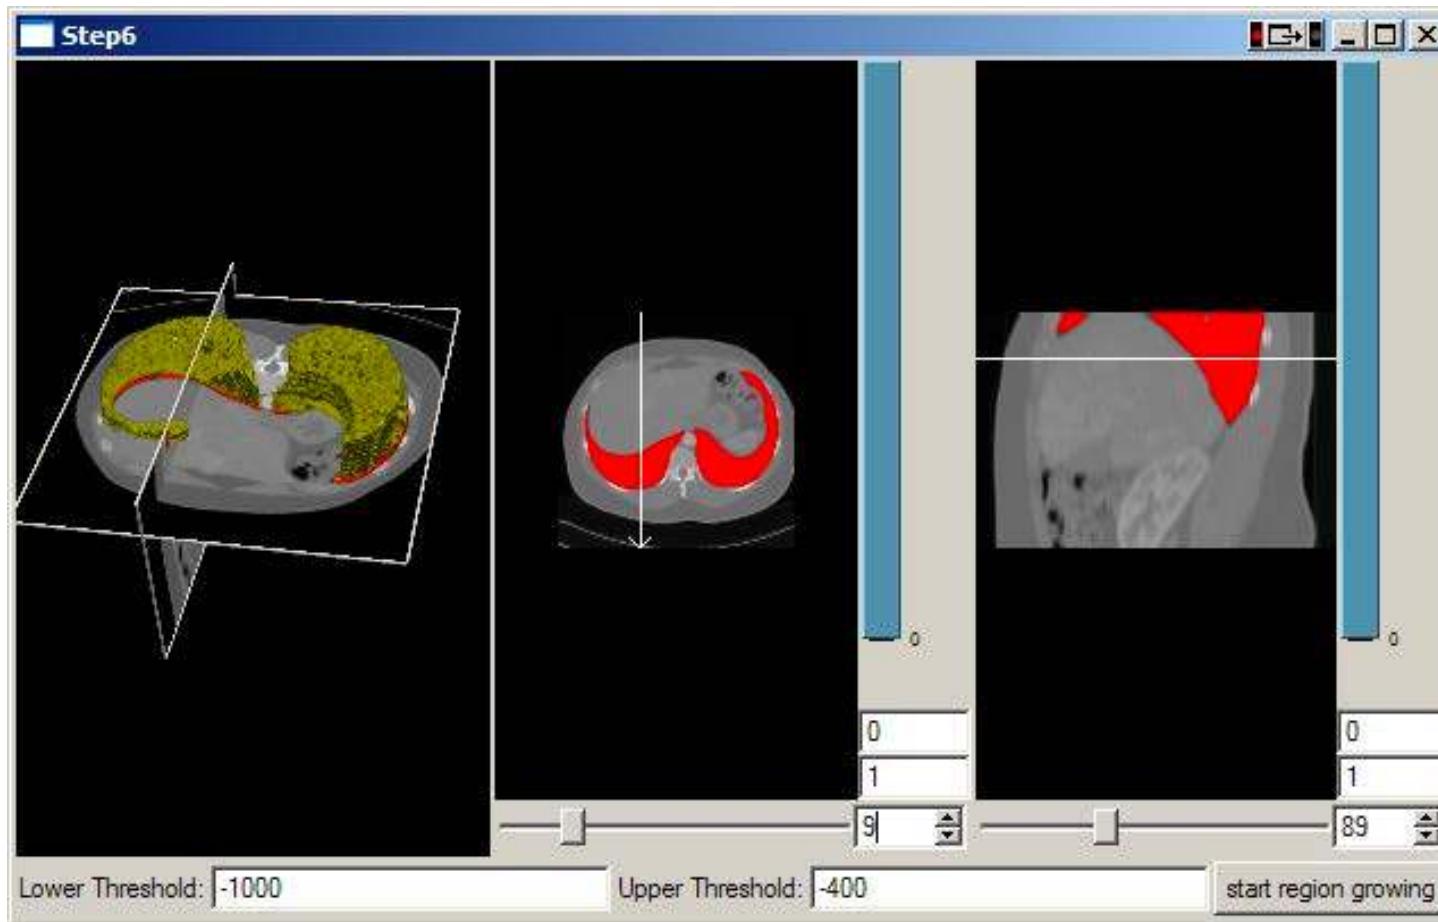
else

 Flag y as a boundary pixel.

fi

od

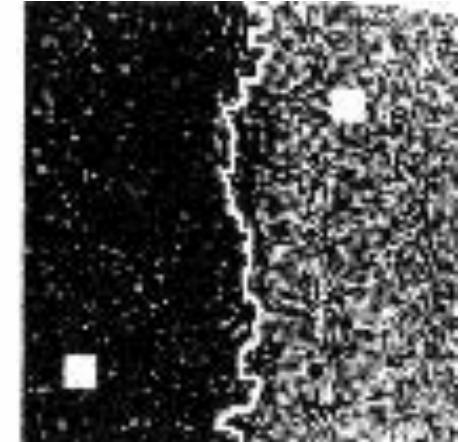
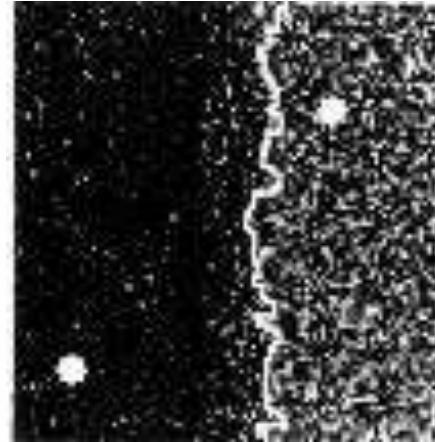
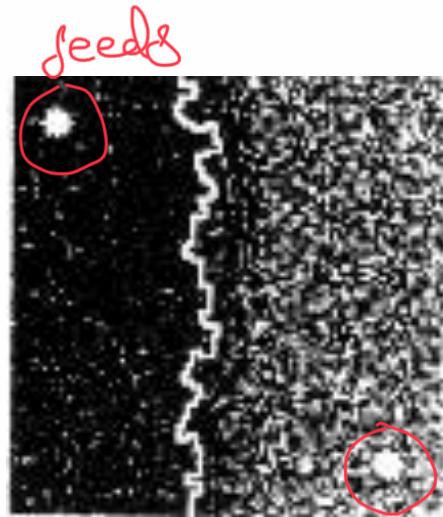
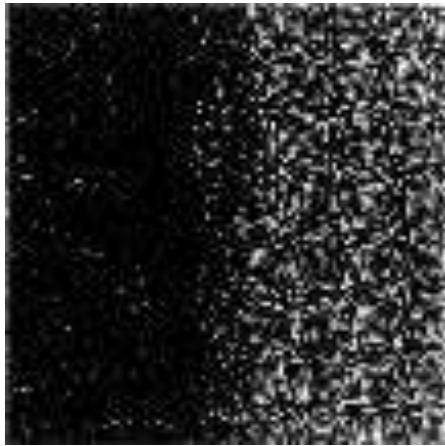
Interactive Region Grower



Medical Imaging Interaction Toolkit

Limitations

Synthetic image



Original image

Result given
two seeds.

Result given
two different seeds.

Result given
larger seeds.

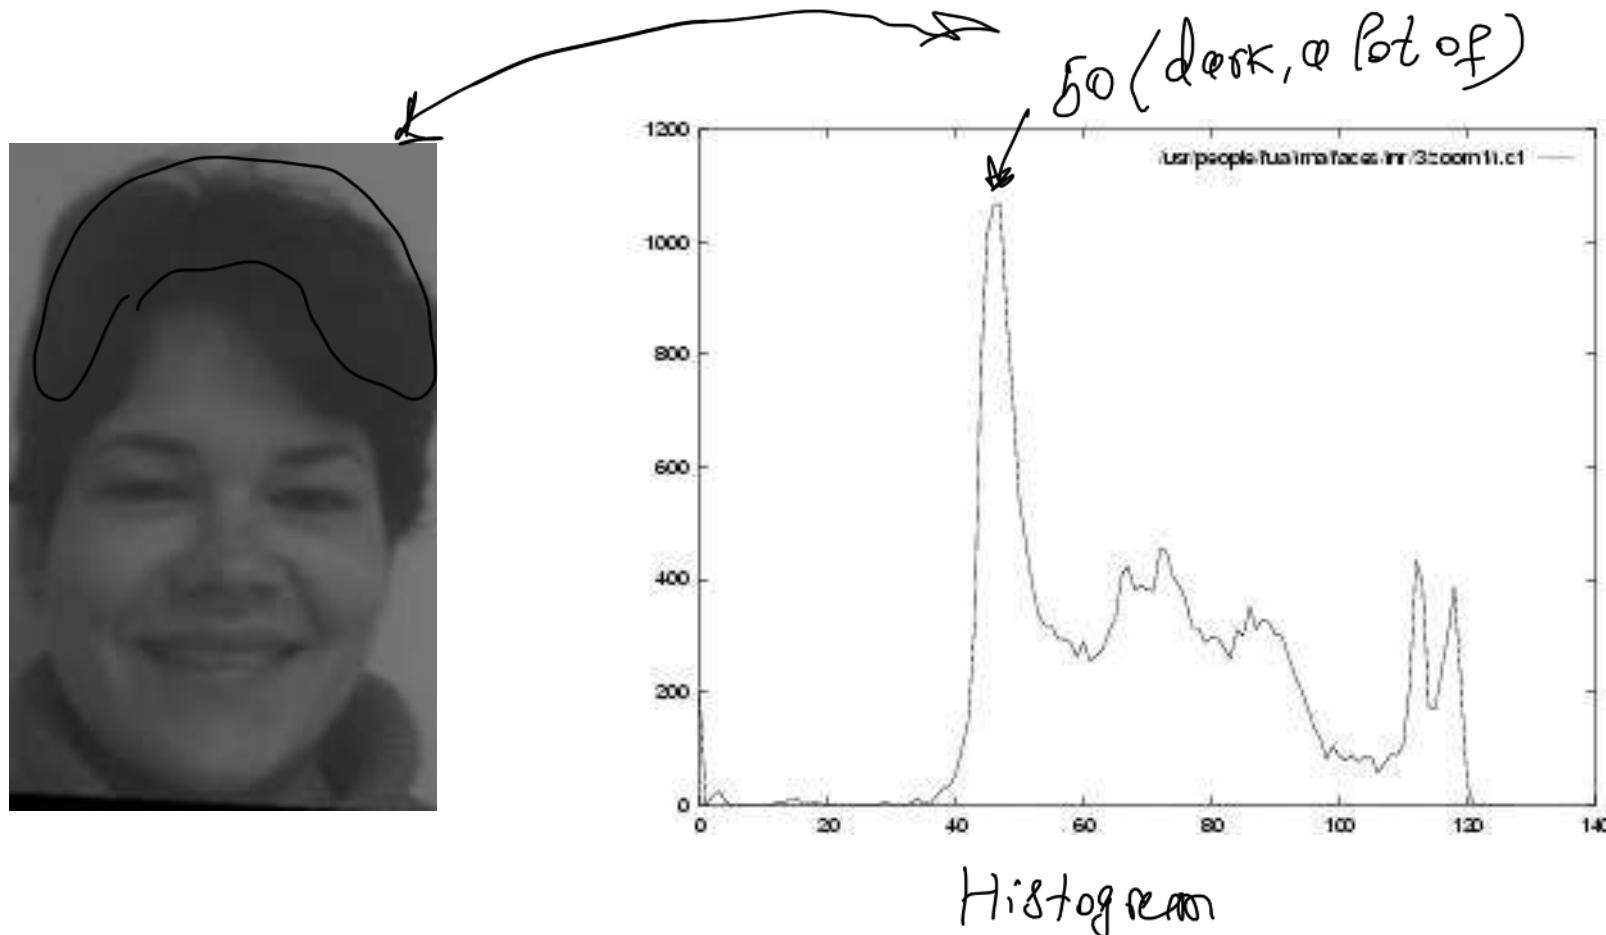
Groedy Algorithm

Statistics become more reliable

- The result depends on the order in which the pixels are taken into consideration.
- The homogeneity measure is noise sensitive.

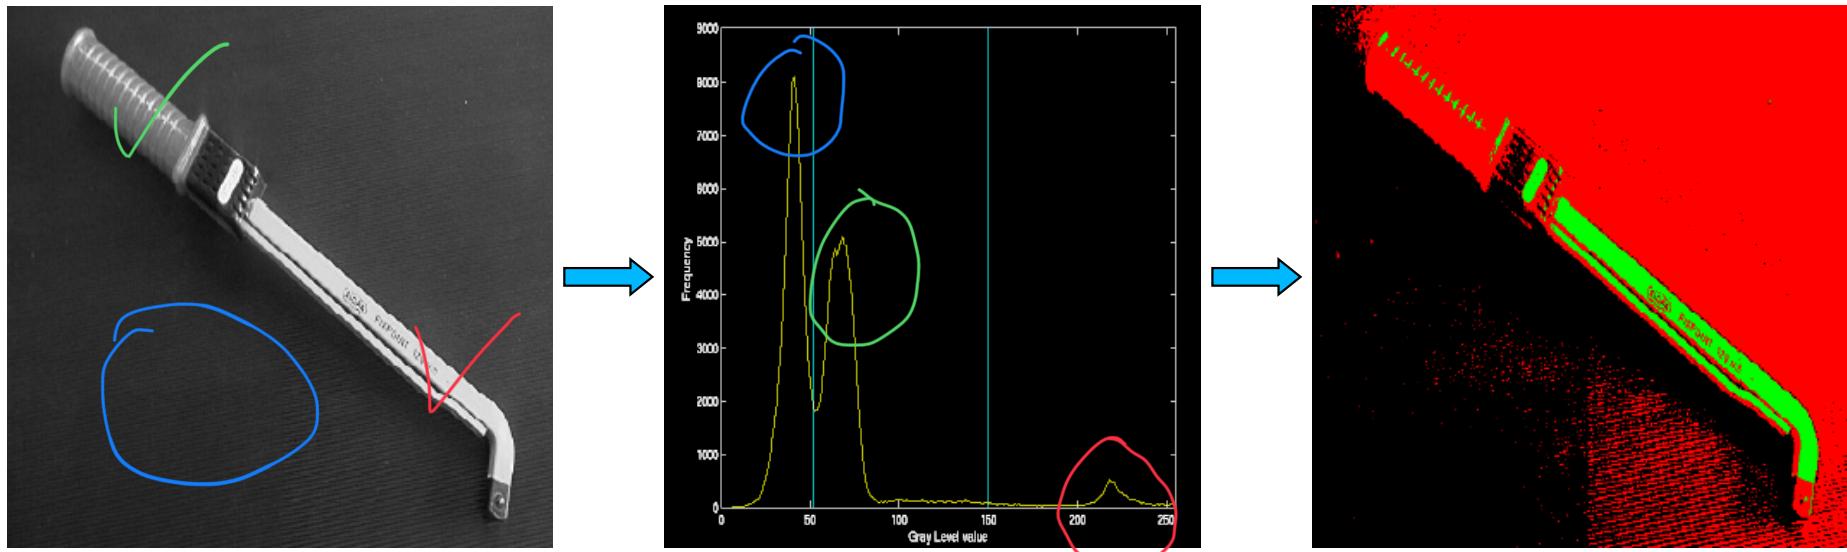
↳ not very robust

Image Histogram



Number of pixels that have a given gray level.

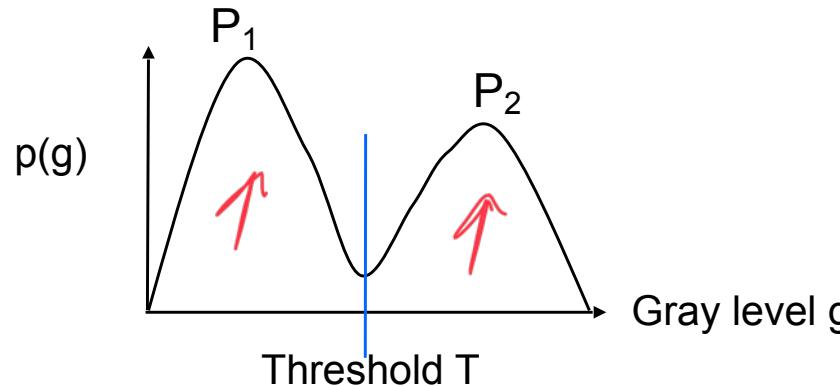
Histogram Splitting



- Groups of similar pixels appear as **bumps** in the brightness histogram
(Blue vertical lines)
 - Split the histogram at local-minima
 - Label pixels according to which bump they belong to
- connected components give regions
- Cannot stop there, must go on.

not considering geometric proximity (Pixels or Pot) .

Recursive Splitting

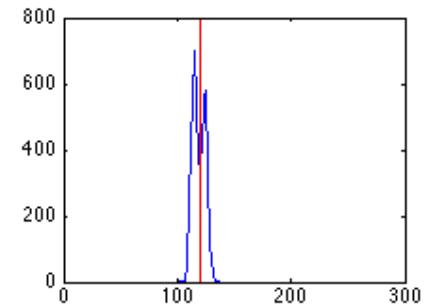
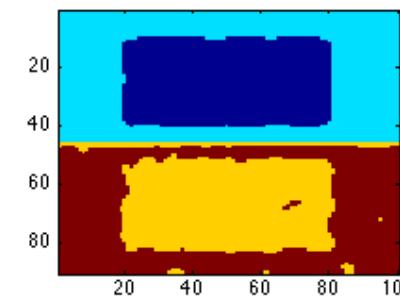
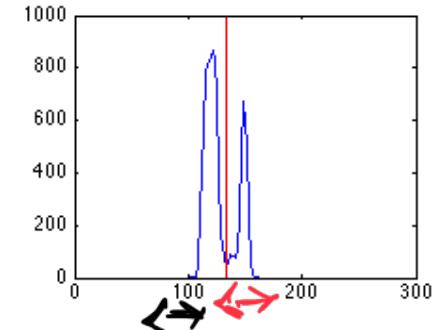
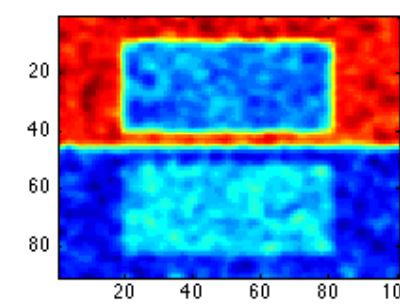
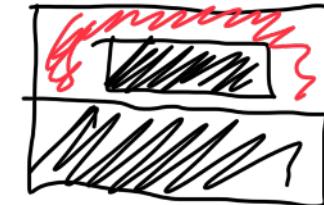


- Compute image histogram.
- Smooth histogram. *so that you don't have too many local minima*
- Look for peaks separated by deep valleys.
- Group pixels into connected regions. ✓
- Smooth these regions. *(within separate images)*
- Iterate.

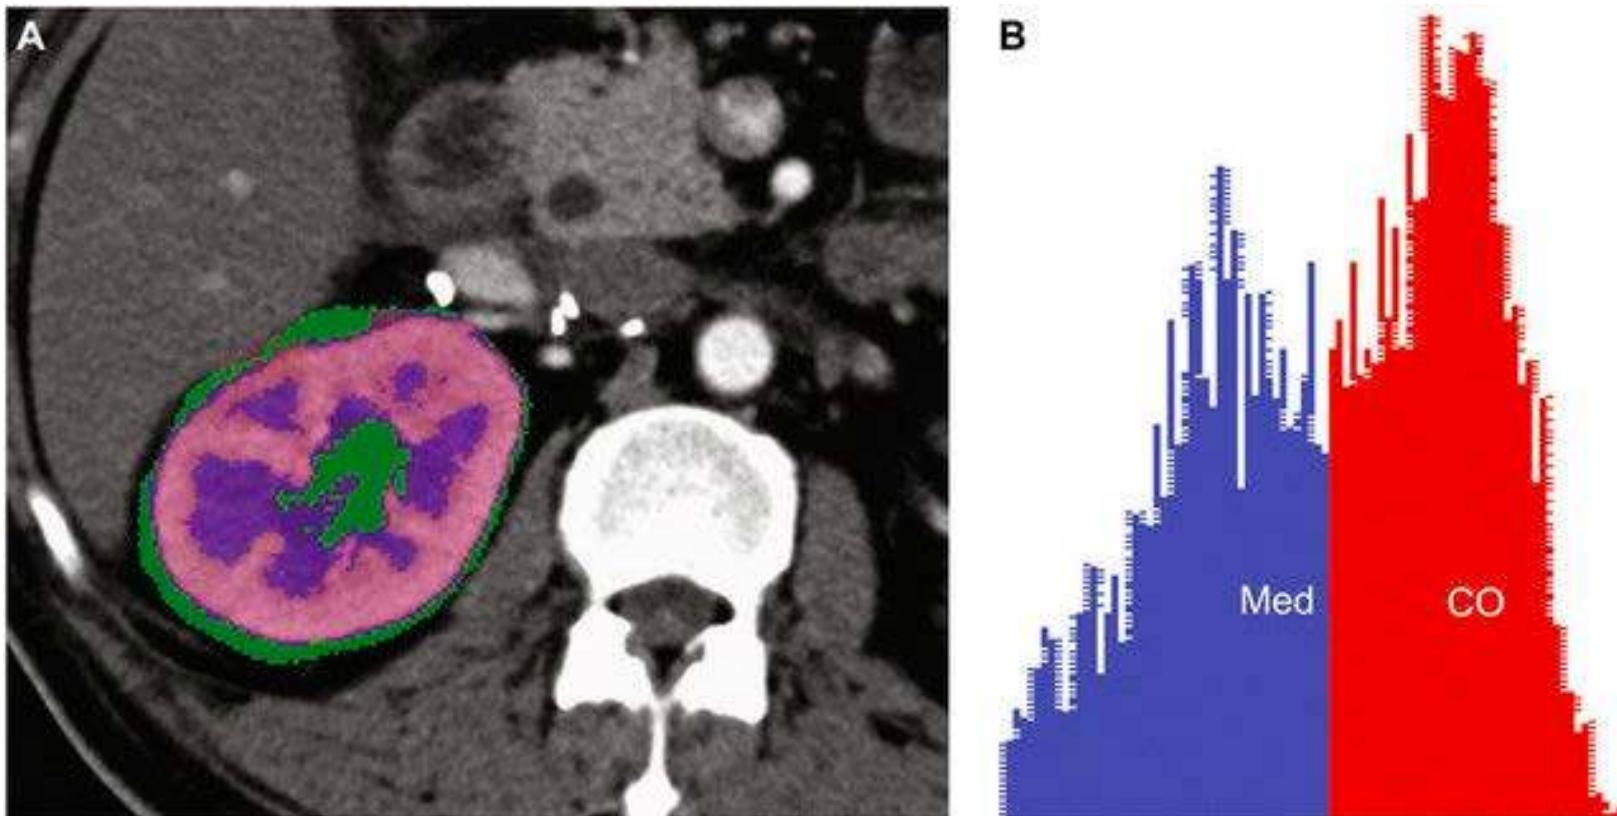
Recursive Splitting

- A first threshold is used to segment the dark pixels.
 - This yields two regions, the bottom half of the picture and the dark rectangle at the top.
 - The bottom half of the picture can now be more easily segmented into two regions.
- > Decisions can be deferred until enough information becomes available.

At first =>
Step

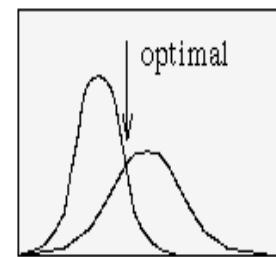
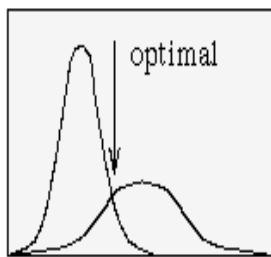
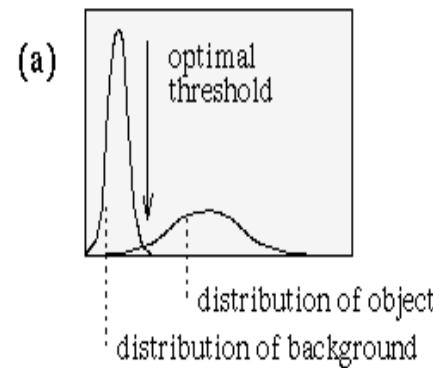


Medical Application



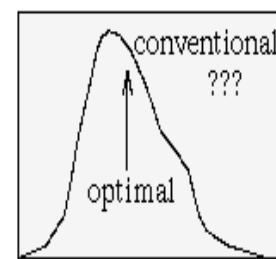
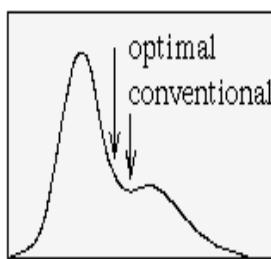
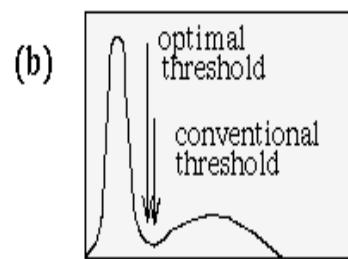
It has its applications but

Finding Thresholds is Hard



Probability distributions

distributions getting more and more similar



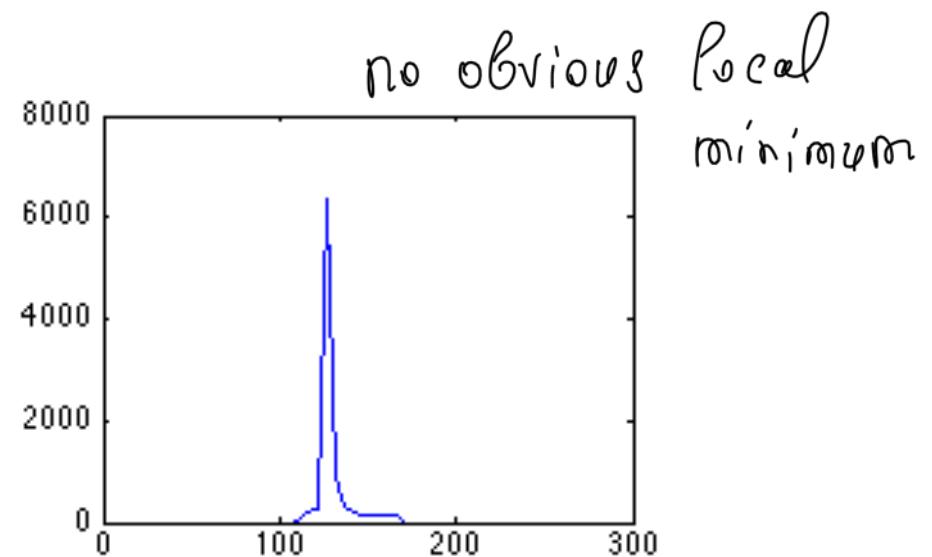
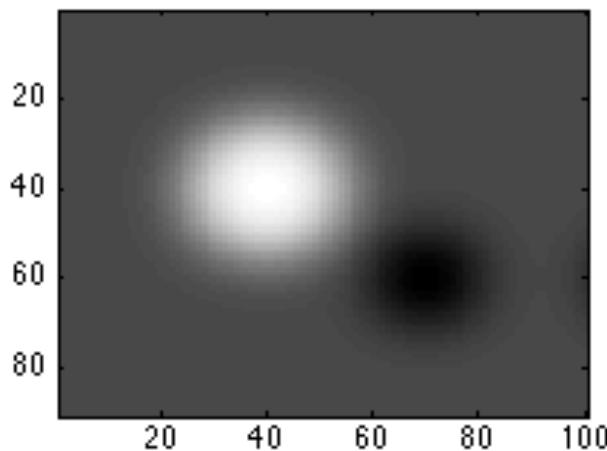
Corresponding histograms

harder and harder to cut

→ Choosing optimal thresholds is a difficult optimization problem.

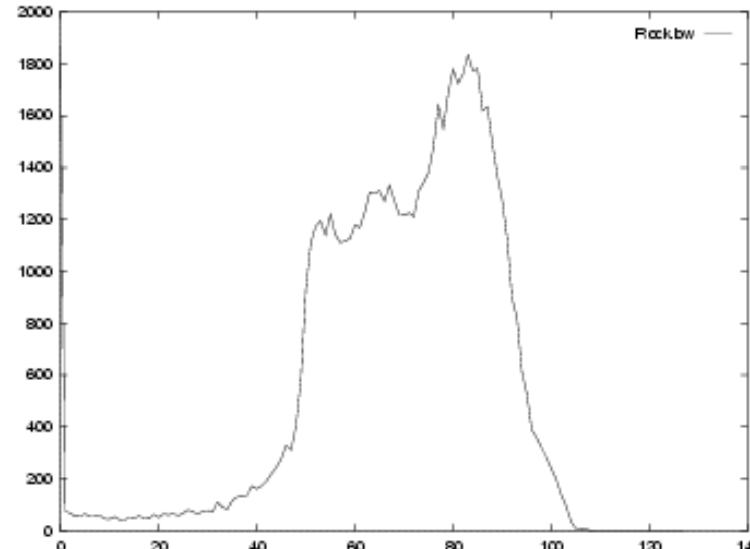
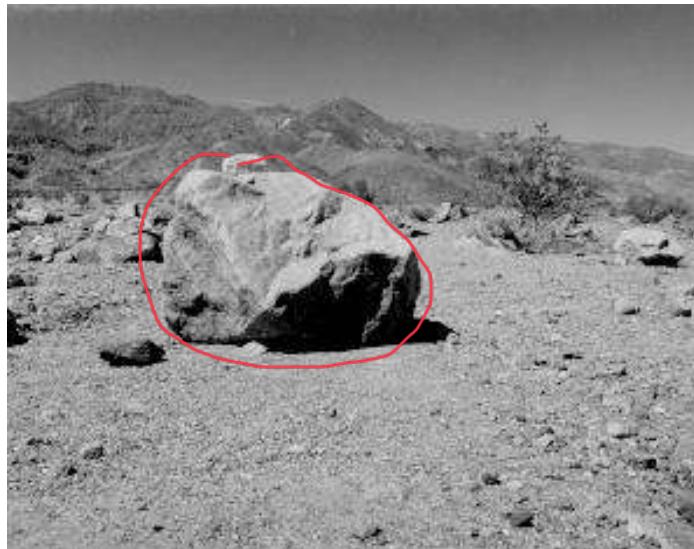
Stop split histogram → when you don't have local minima

No Obvious Threshold



There are thresholds that would work but you can't find them from the histogram only.

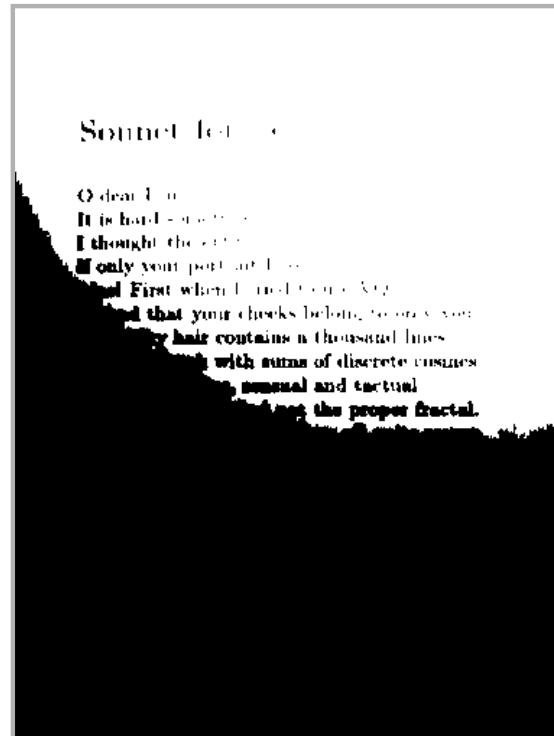
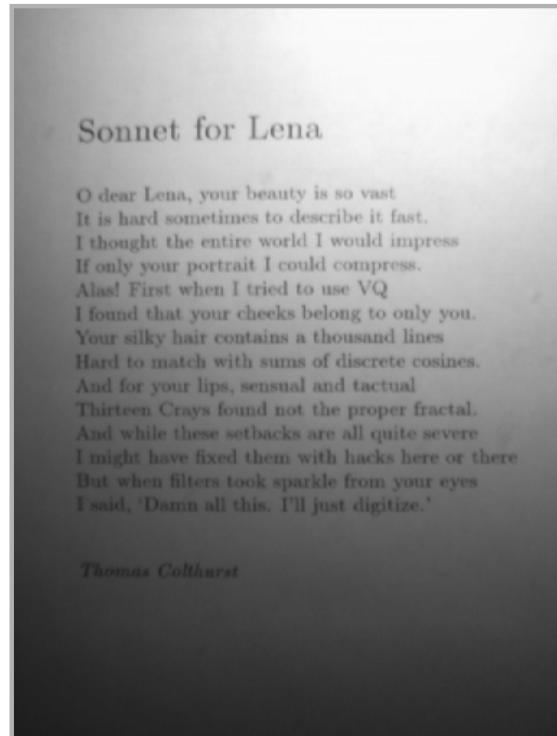
No Valid Threshold at All



hard to find
threshold that
you'd find the
stone's boundary

In this image, no threshold can generate meaningful results.

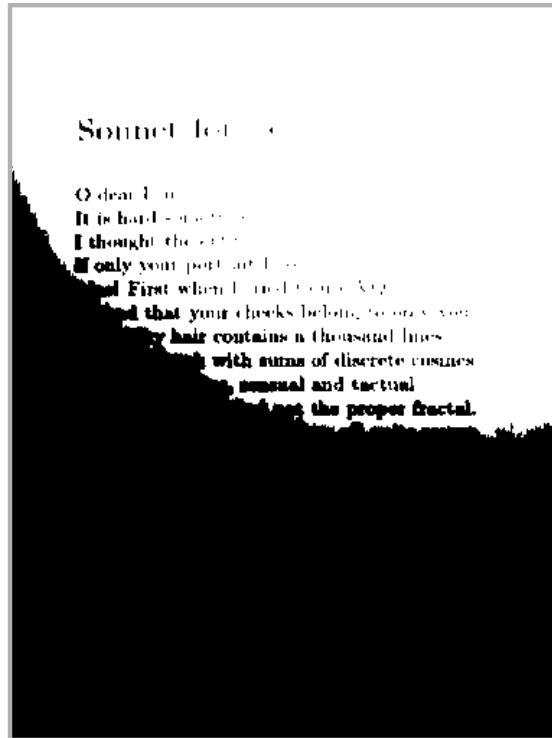
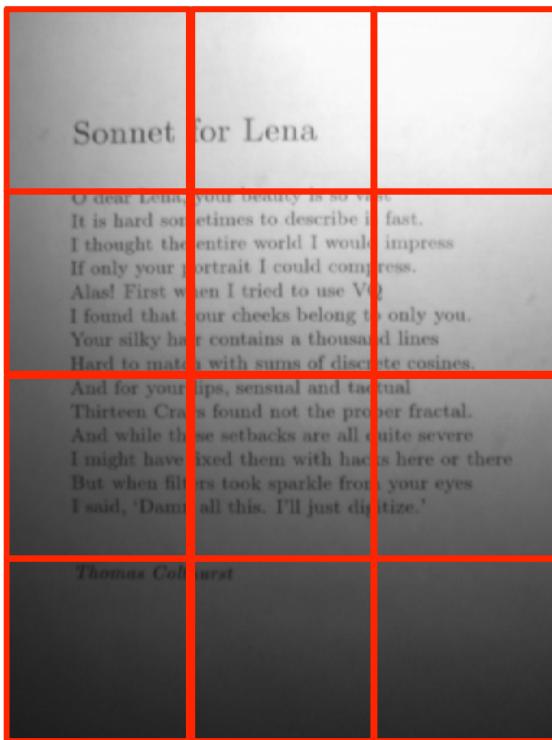
No Global Threshold



Because of the complex illumination:

- The top of the page is much lighter than the bottom.
- No global threshold can be found.
- Local thresholds can account for this.

Use Local Histograms



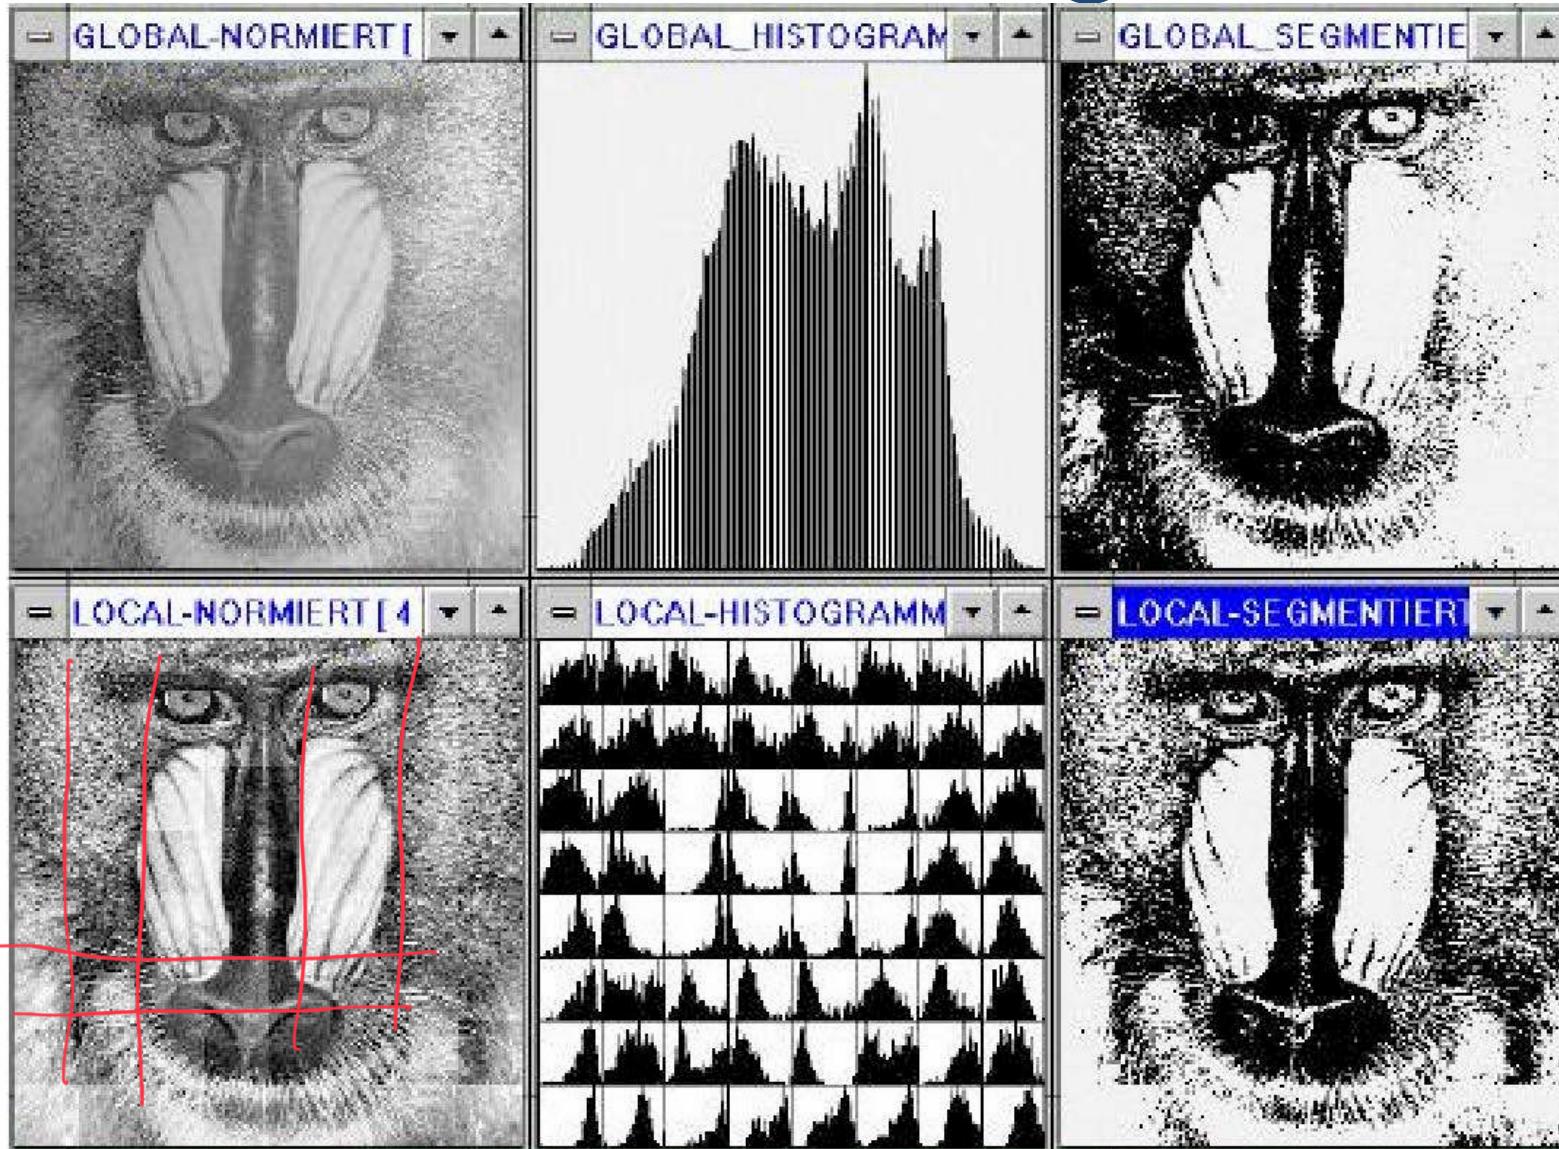
Global

Local

- Compute a histogram in each red square.
- Find a local threshold.
 - Segment independently

Local information
⇒ proximity,
geometry
matter!

Use Local Histograms

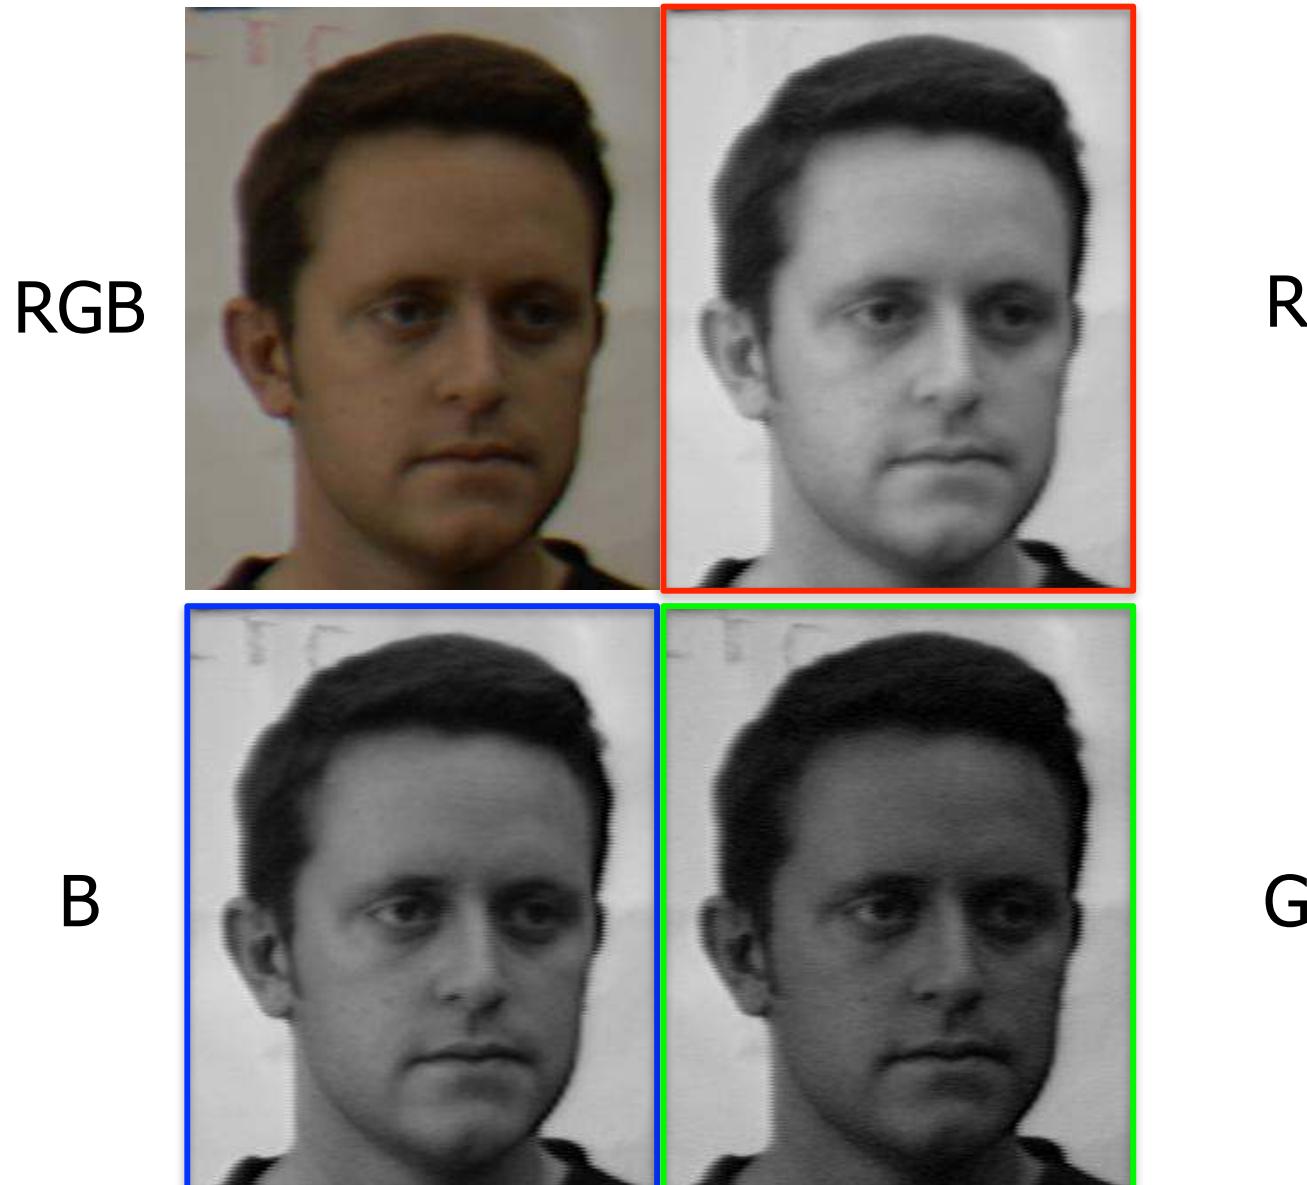


- Compute local histograms on a coarse grid.
- Use them instead of a global one.

Limitations of Global Histograms

- Histograms do not account for neighborhood relationships.
↳ spatial information between regions ⇒ not at all!
- Thresholds are hard to find.
- Some boundaries will not be found because the gray levels on both side belong to the same histogram peak.
(rock example)

Using Color



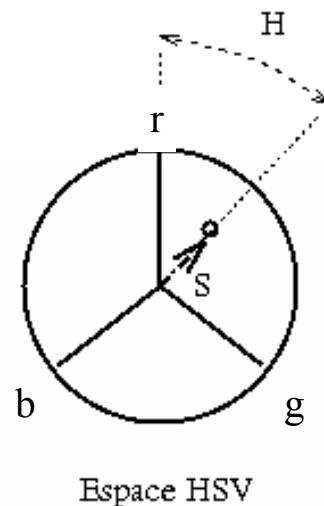
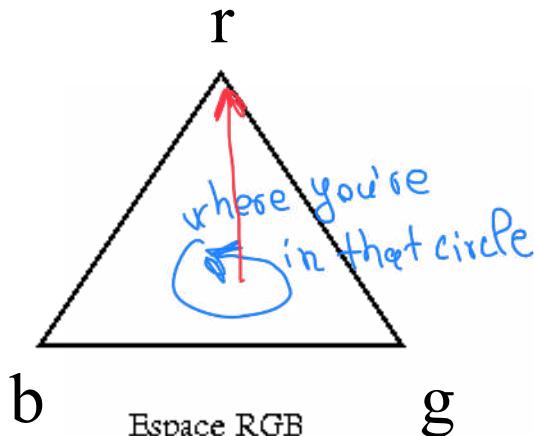
Color Space

$$V = R + G + B$$

$$r = \frac{R}{V}$$

$$g = \frac{G}{V}$$

$$b = \frac{B}{V}$$



- Value



- Hue

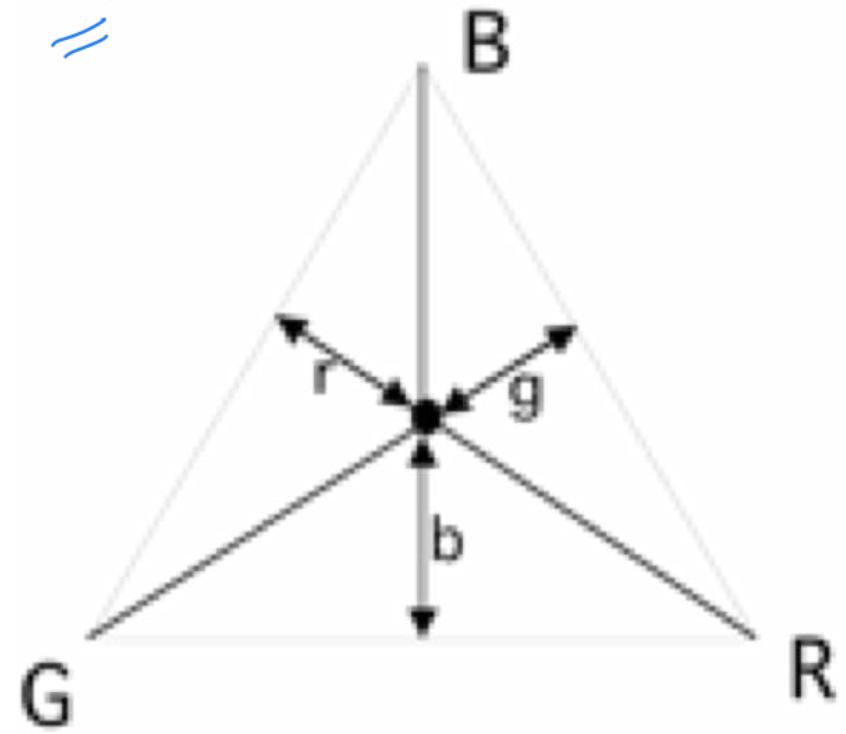
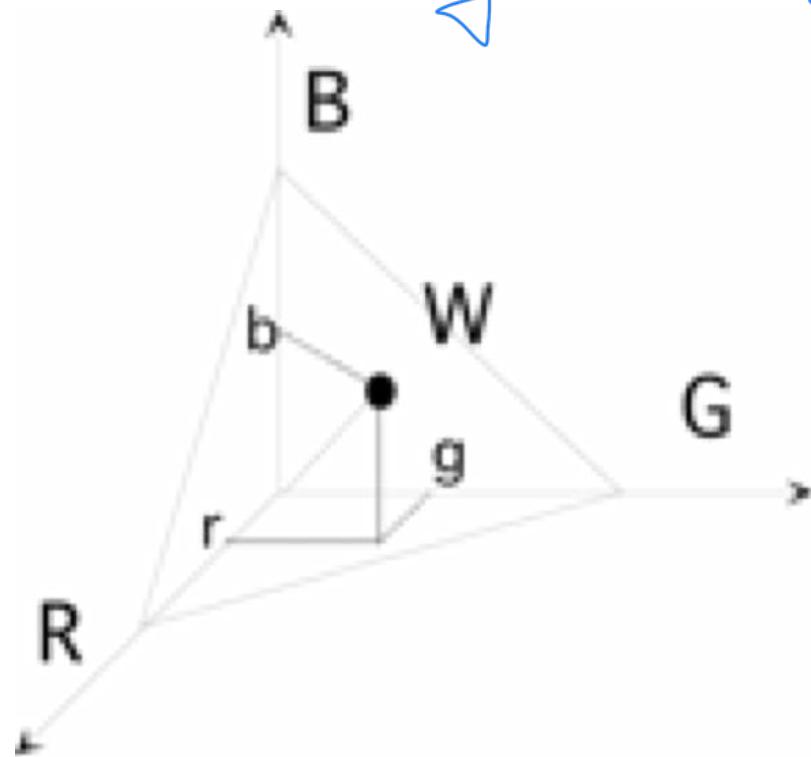


- Saturation



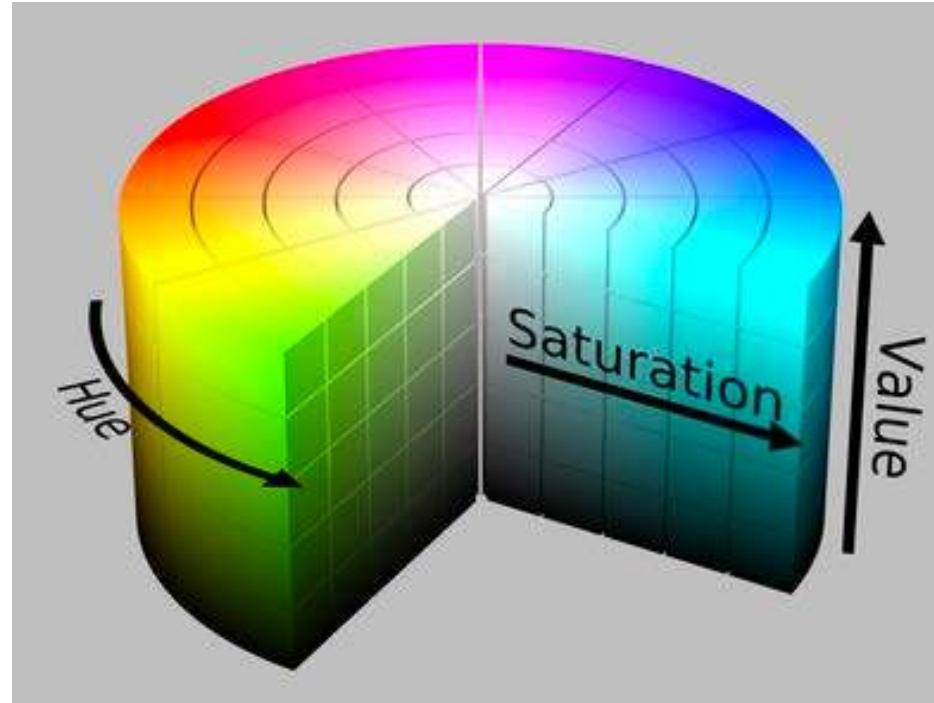
RGB Chromaticity Diagram

Chromaticity of the pixel $P \rightarrow$ color



The Maxwell triangle involves projecting the colors in RGB space onto $R+G+B=1$ plane.
→ Chromaticity becomes independent from luminance.

HSV Space



Hue / Saturation / Value



angle

$$H = \arccos\left(\frac{0.5(2r - g - b)}{\sqrt{(r - g)^2 + (r - g)(g - b)}}\right) \text{ if } b < g$$

$$1 - \frac{3 \min(r, g, b)}{T} \quad S = \max(r, g, b) - \min(r, g, b) \rightsquigarrow \begin{array}{l} \text{How close you're to the} \\ \text{boundary of triangle} \end{array}$$

$V = R + G + B \rightsquigarrow \text{how bright you are}$

HSV Images

RGB



saturation's low

How close we're
to the pure color?

Saturation

- Skin is saturated
- Lips is even more saturated



$R + G + B$

gray-level image

Value

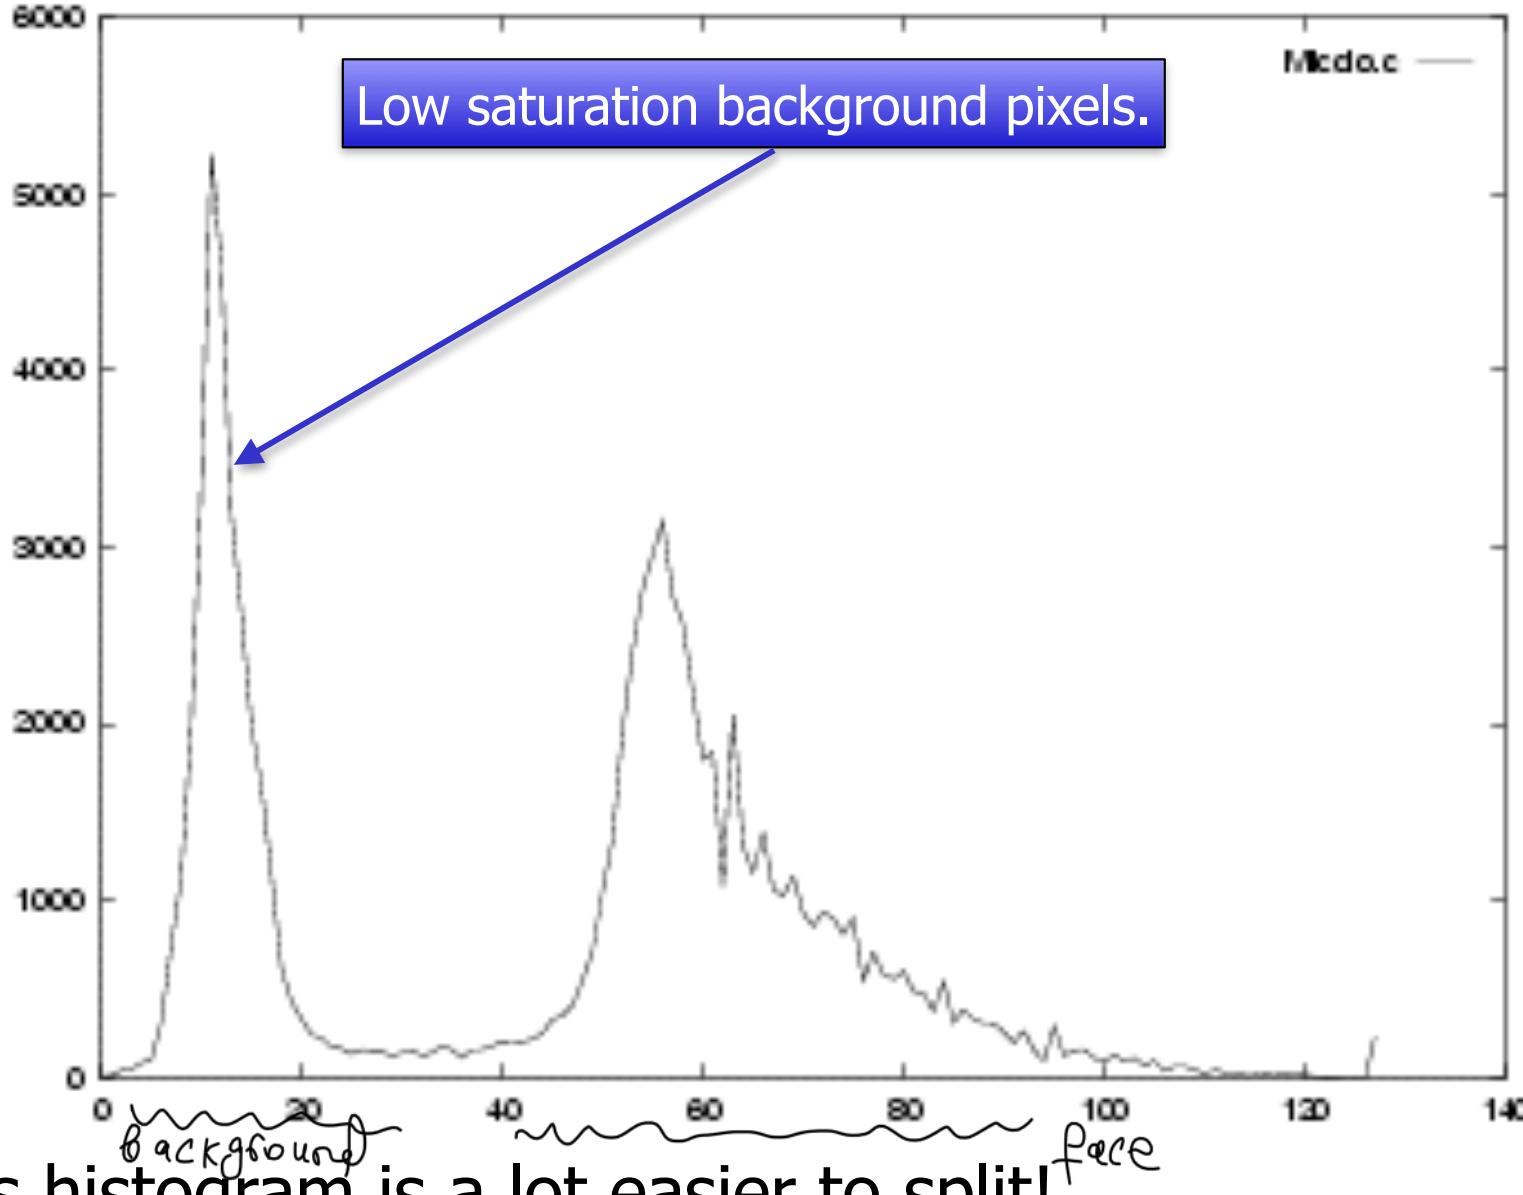


What kind of color
we're dealing with?

Hue



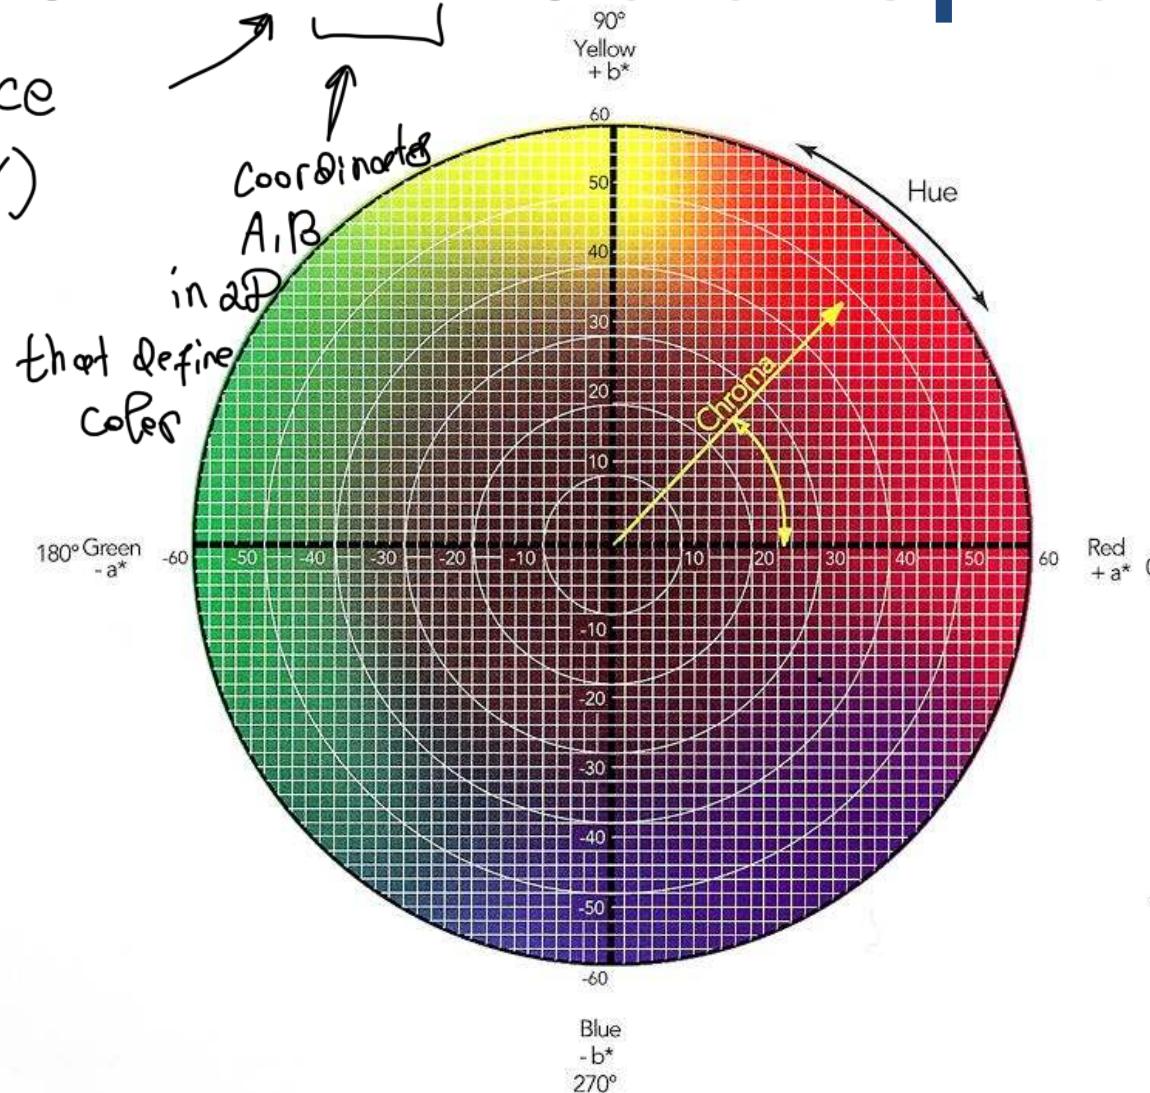
Saturation Histogram



- This histogram is a lot easier to split!
- It makes it easy to segment the head from the background.

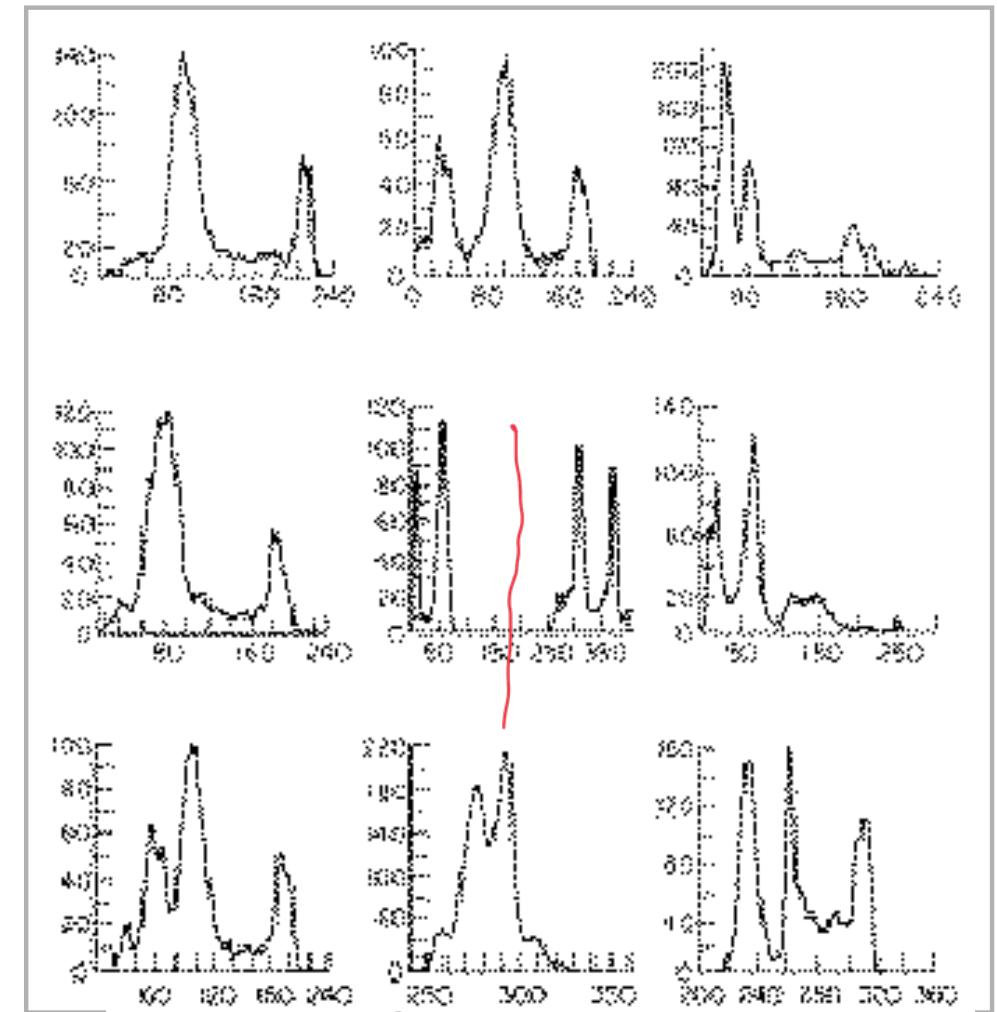
CIE LAB Color Space

Luminance
(similar to H_{8V})



- Another way to represent color in terms of three coordinates, one for “brightness” and the other two for chromaticity.
- Designed so that the same amount of numerical change in these values corresponds to roughly the same amount of visually perceived change.

Using Multiple Histograms

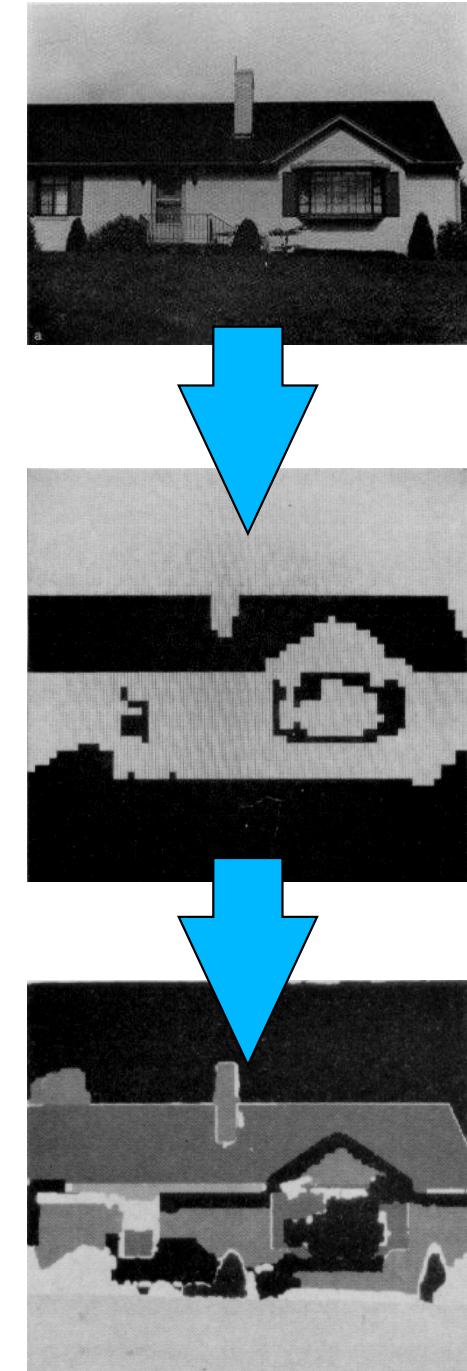
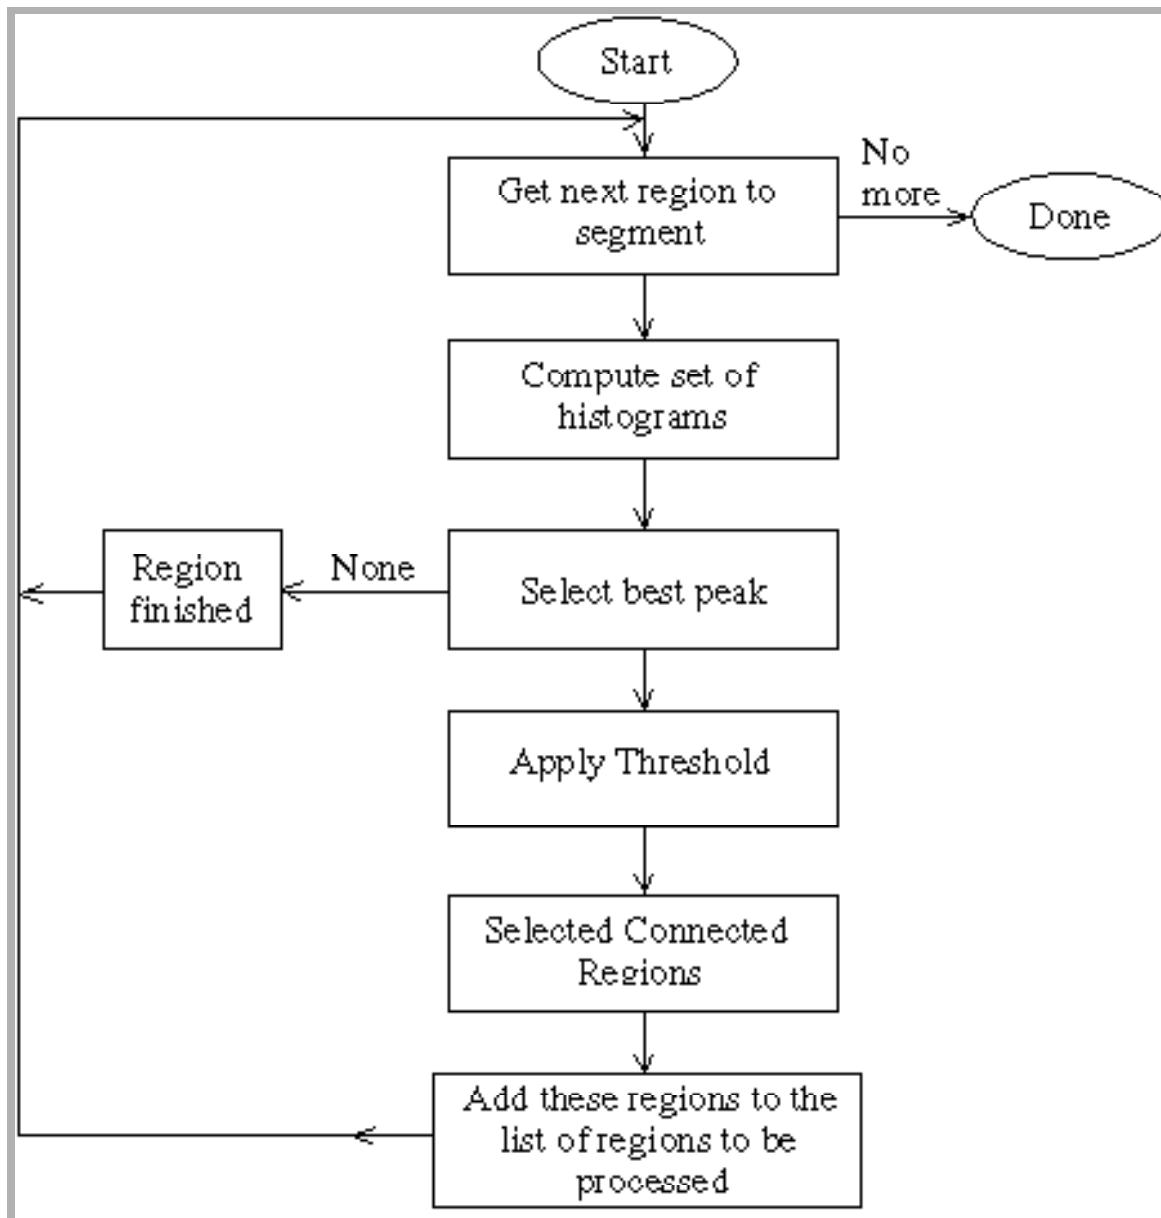


→ not only traditional but also HSV histograms
(or LAB)

- Compute multiple histograms for each segment.
- Use one of them to split each segment.
- Repeat on the resulting smaller segments.

use the one that's
easier to segment

Recursive Algorithm



Early Color Segmentation



More Recent Color Segmentation



I find this blue sky too bland!



Replace it.



puts boundaries
like this

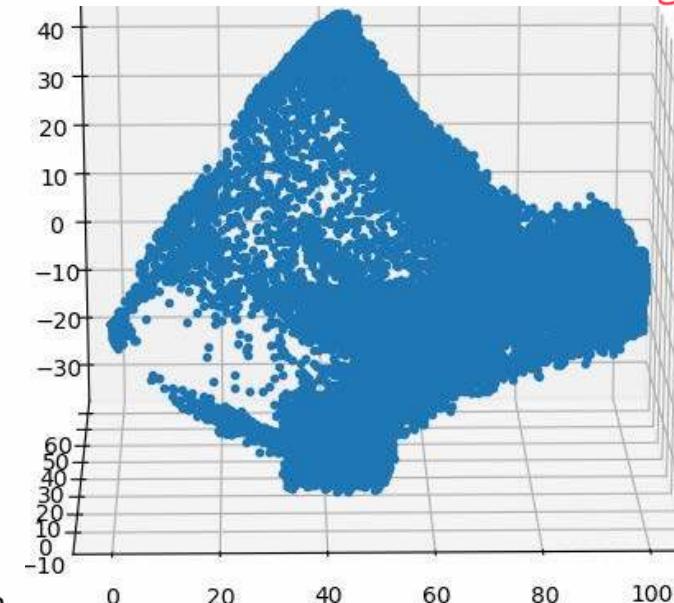
Manual intervention was still required to find the whole sky region the sky region!

Segmentation as Clustering



→ Luminance;
gray-level
 L

b



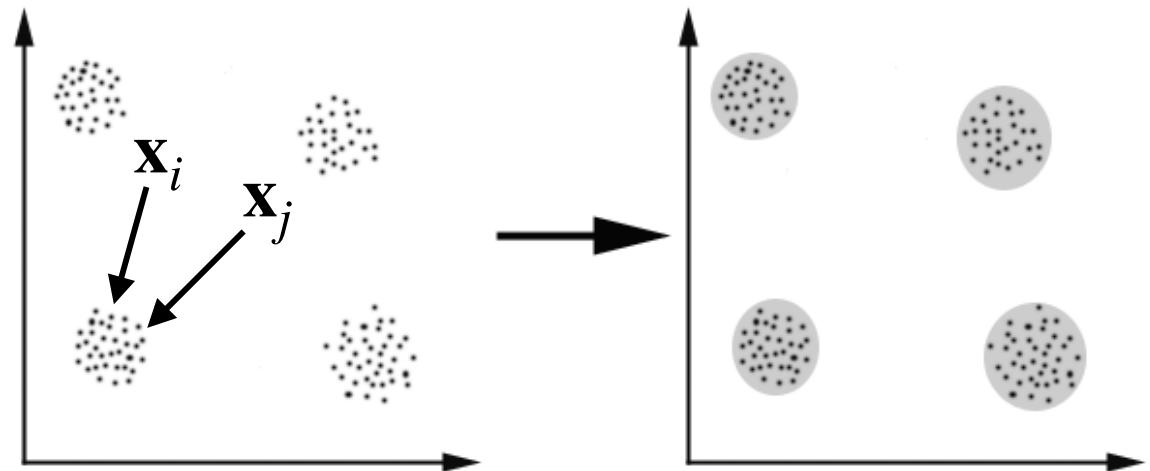
Cloud of
3D pts

channel 6 defines what kind of colors we're dealing with

- ⑤ Each pixel has 2 spatial coordinates and 1 gray level or 3 color components. The main problem with histogram splitting
- Segmentation can be understood as clustering in
 - 1D space (G);
 - 3D space (x,y,G), (H,S,V), (L,a,b);
 - 5D space (x,y,L,a,b). complete description of pixels
- Ideally each cluster should be as compact as possible.

it completely forgets
where the pixels come
from

K-Means Clustering

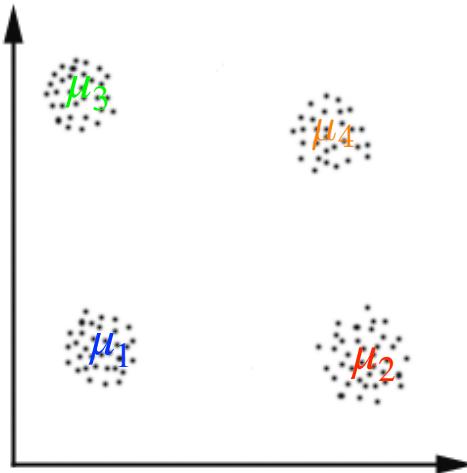


in n -dimension

Given a set of input samples:

- Group the samples into K clusters.
- K is assumed to be known/given. *priori*
- In the example above, each sample is a point in 2D.
- In images, samples can be points in 1D, 3D, or 5D.

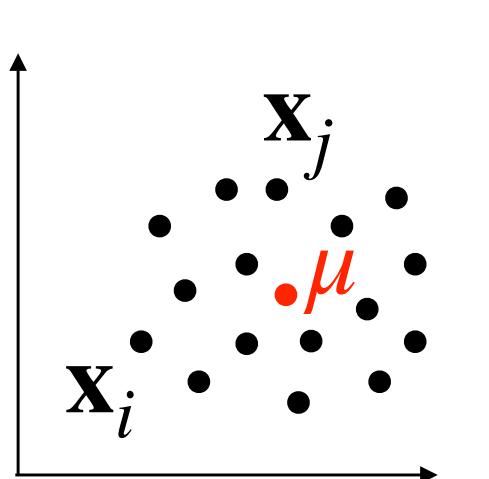
K-Means Clusters



- Cluster k is formed by the points $\{\mathbf{x}_{i_1^k}, \dots, \mathbf{x}_{i_{n^k}^k}\}$.
- μ_k is the **center of gravity** of cluster k.

Sample
Subset of

The mean of points $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, $\mathbf{x}_i \in \mathbb{R}^D$ is

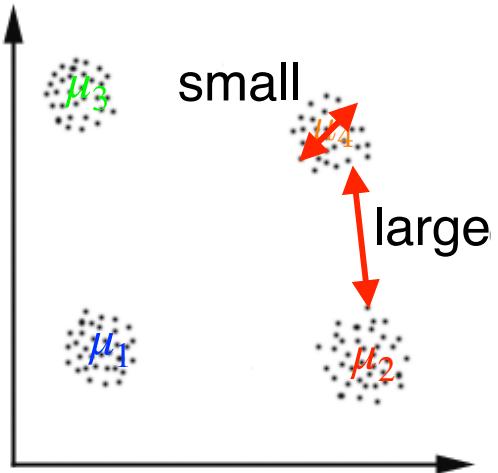


$$\mu = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i, \mu \in \mathbb{R}^D$$

In 2D

- If the \mathbf{x}_i were physical points of equal mass, μ would be their center of gravity.
- This applies in any dimension.

Formalization



- Cluster k is formed by the points $\{\mathbf{x}_{i_1^k}, \dots, \mathbf{x}_{i_{n^k}^k}\}$.
- μ_k is the **center of gravity** of cluster k .

intra distances of the cluster \rightarrow small

- The distances between the points within a cluster should be small.
- The distances across clusters should be large.
- This can be encoded via the distance to cluster centers $\{\mu_1, \dots, \mu_K\}$:

$$\xrightarrow{\text{Minimize}} \sum_{k=1}^K \sum_{j=1}^{n_k} (\mathbf{x}_{i_j^k} - \mu_k)^2$$

we want to find
sum over all clusters

where $\{\mathbf{x}_{i_1^k}, \dots, \mathbf{x}_{i_{n^k}^k}\}$ are the n^k samples that belong to cluster k .

Difficult Minimization Problem

Minimize

$$J = \sum_{k=1}^K \sum_{j=1}^{n_k} (\mathbf{x}_{i_j^k} - \underline{\mu_k})^2$$

- i) we want to find cluster centers
ii) Assignments to clusters of points to minimize J

but:

- We don't know what points belong to what cluster.
- We don't know the center of gravity of the clusters.

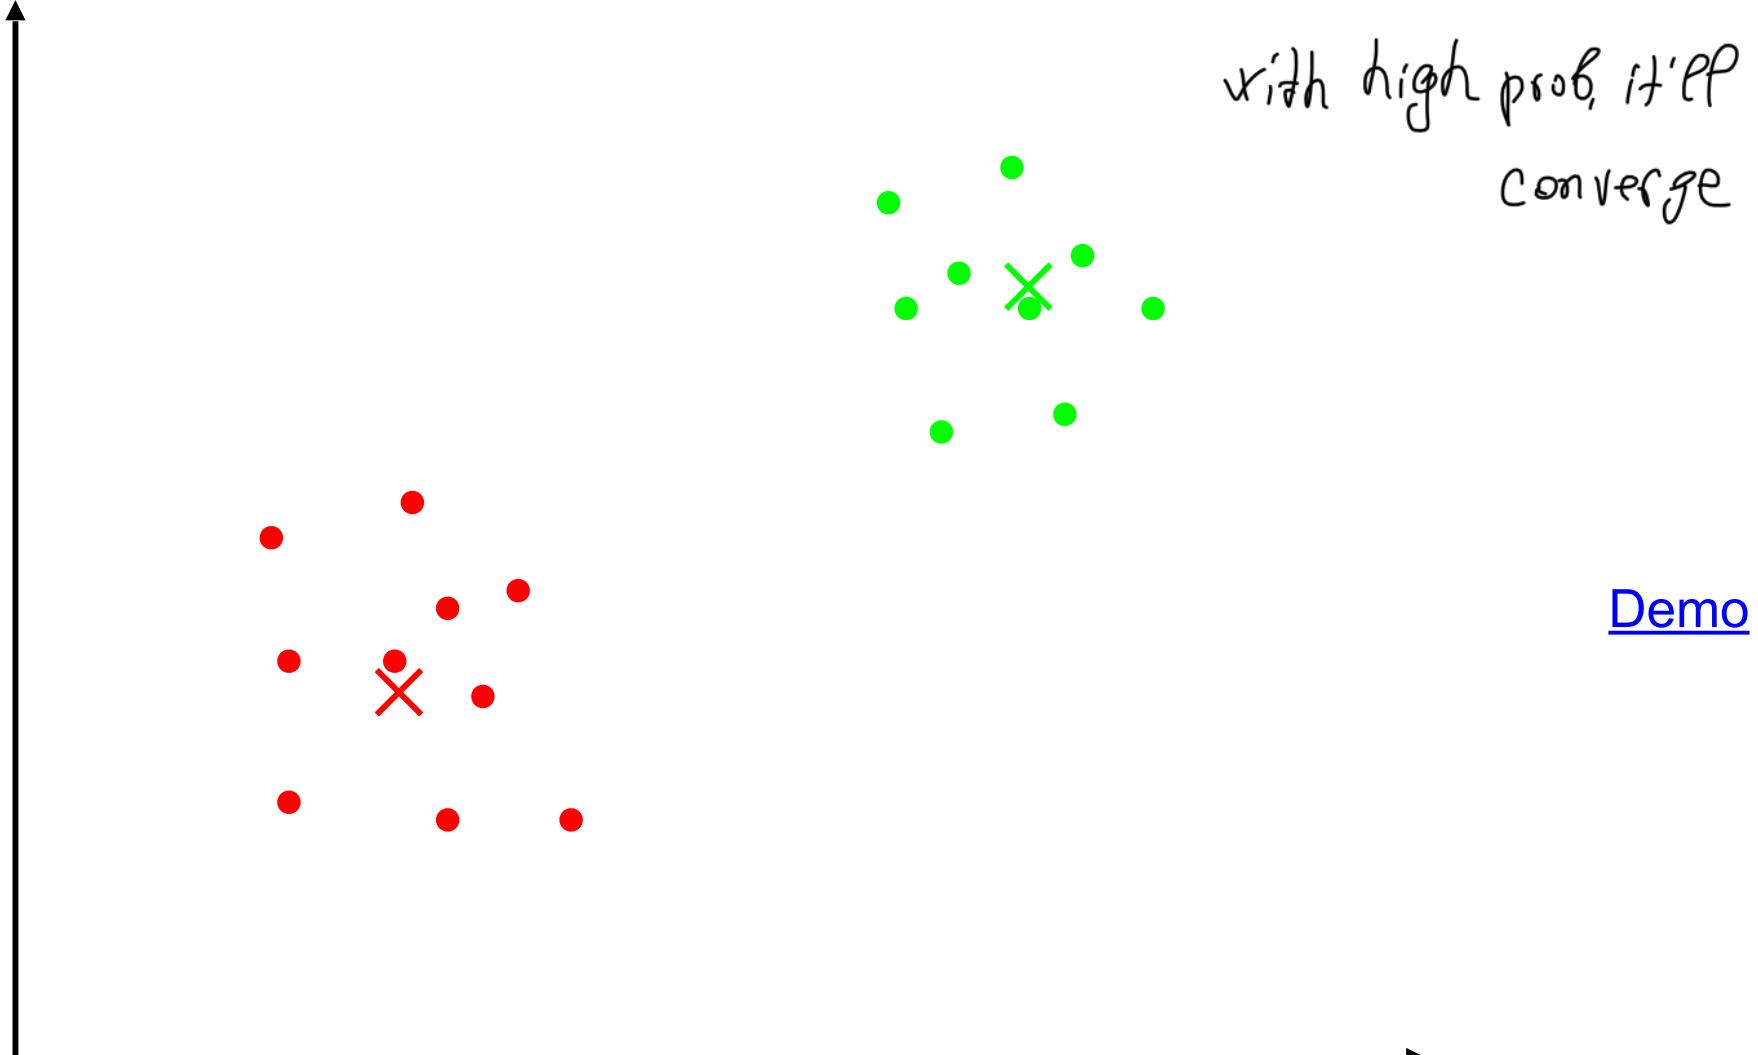
We don't know μ 's



Simple Solution to the Problem

1. Initialize $\{\mu_1, \dots, \mu_K\}$, randomly if need be.
2. Until convergence
 - 2.1. Assign each point x_i to the nearest center μ_k
 - 2.2. Update each center μ_k given the points assigned to it
By computing mean of those pts

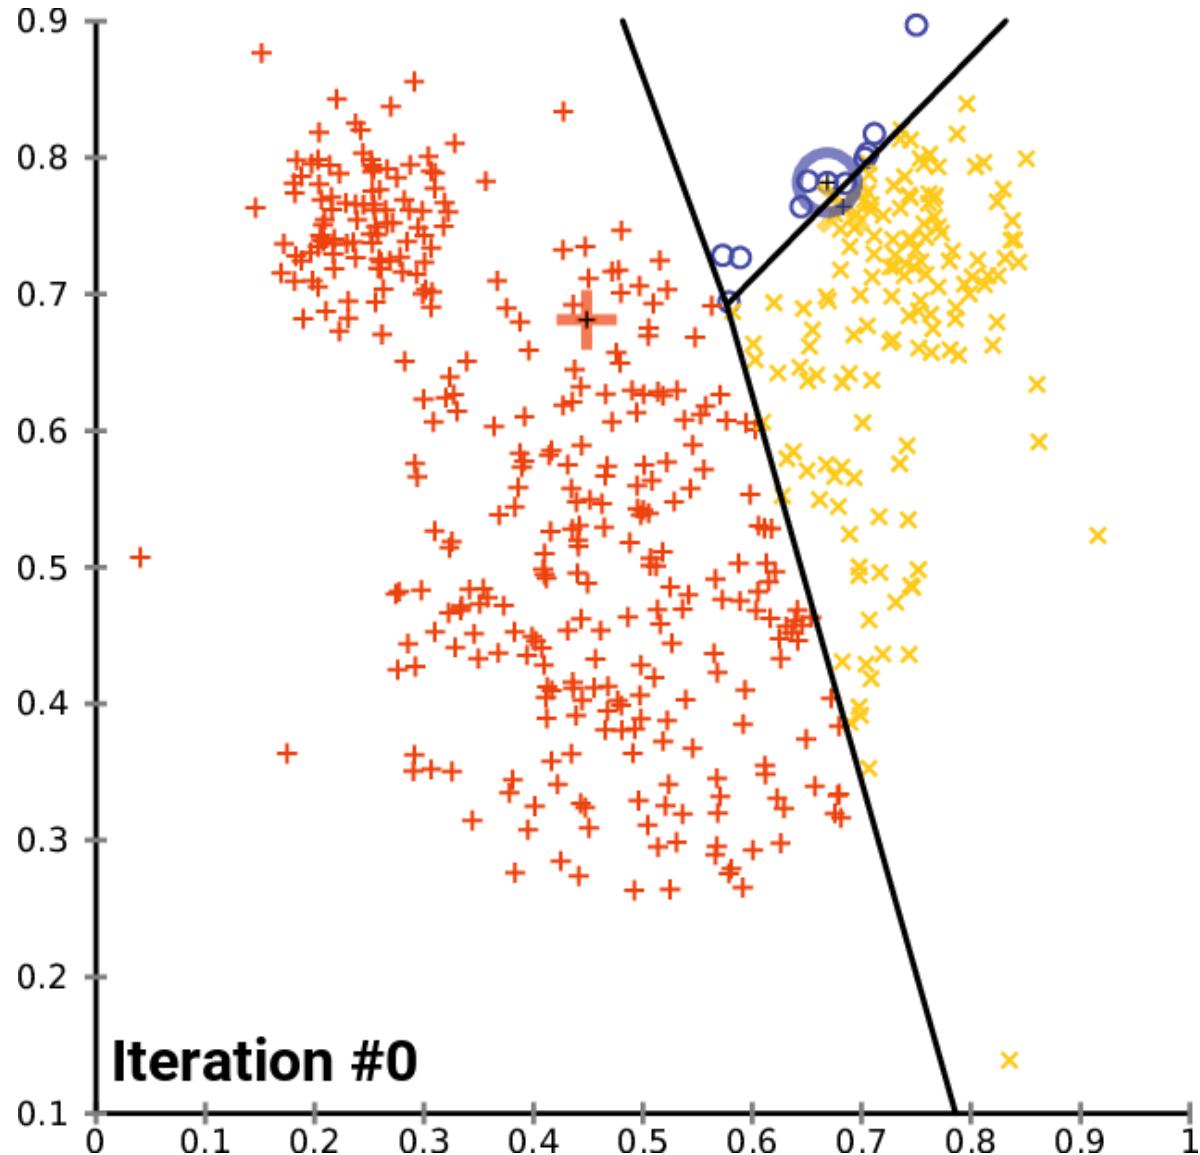
Alternating Optimization



- Initialize
- Associate point to centers
- Recompute centers

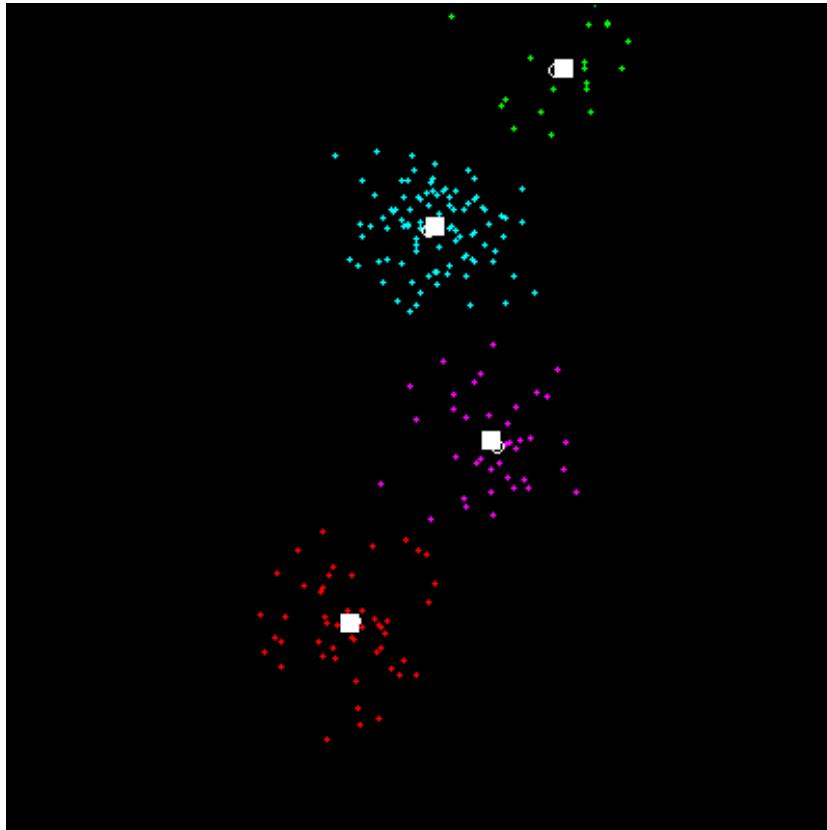


Three Classes

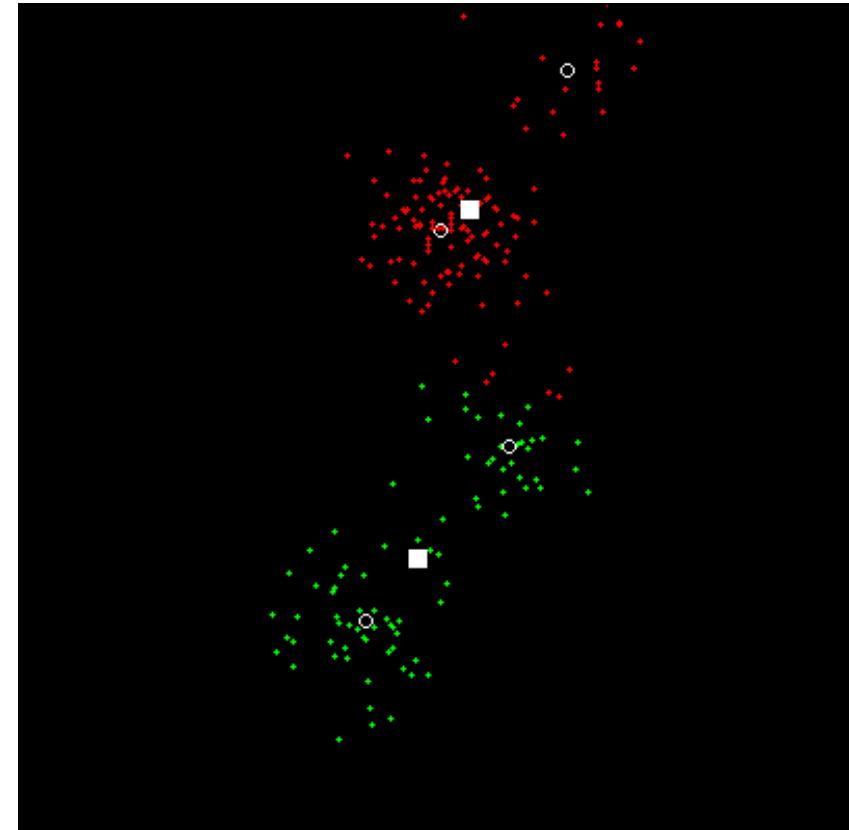


[Demo](#)

Initial Conditions Matter



Initially, the points are assigned to the clusters at random → Success.

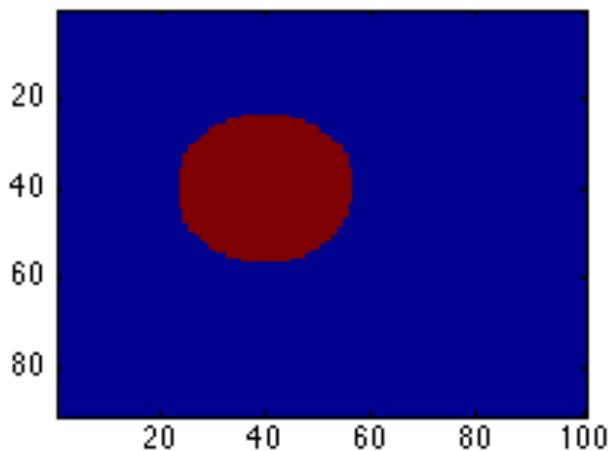
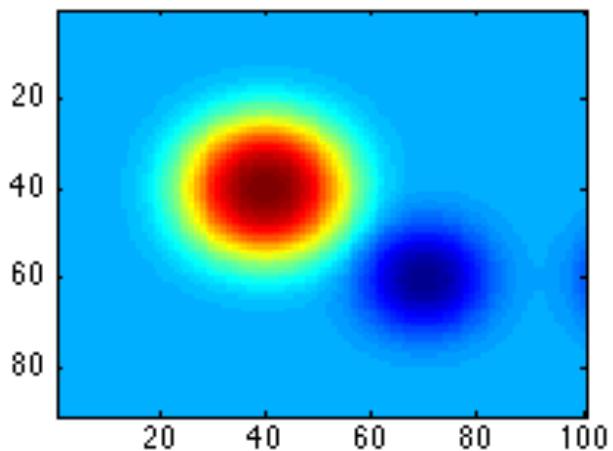


Initially, the points are assigned to the closest cluster → Failure.

→ In practice try several different random initialization and keep the one that yields the best result in term of the sum of square distances.

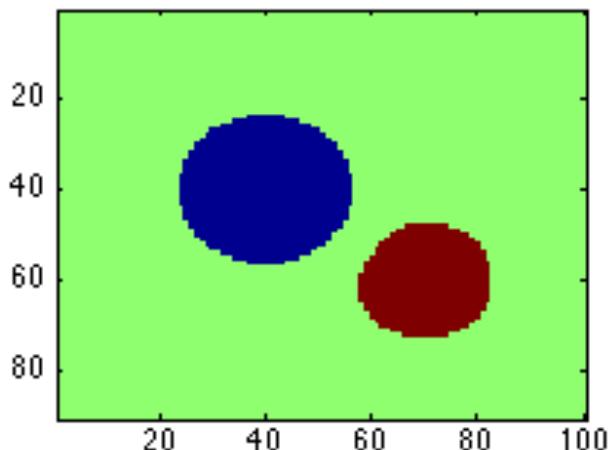
our criteria

GRAY-LEVEL ONLY (1D)

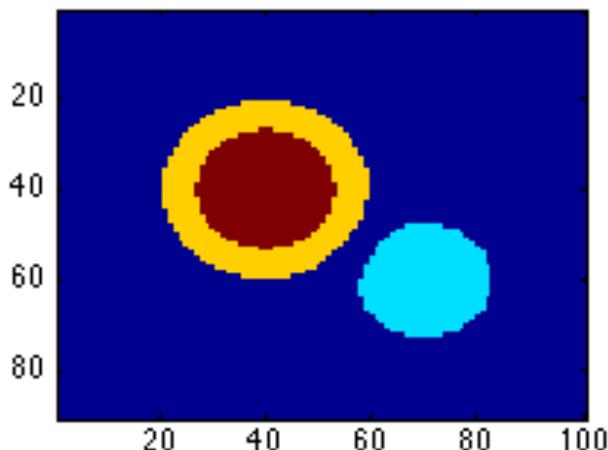


$K \rightarrow \# \text{ of } \text{Clusters}$

$K=2$



$K=3$

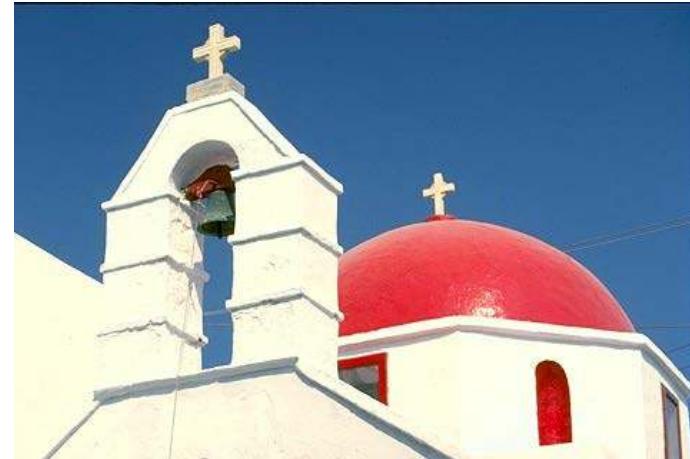


$K=4$

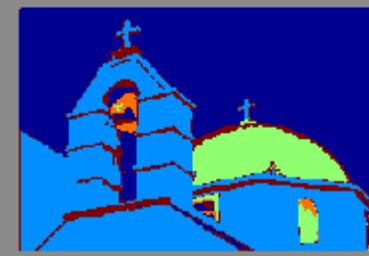
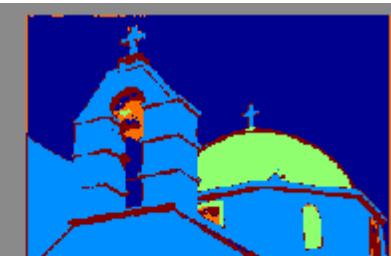
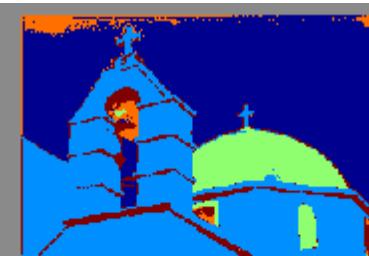
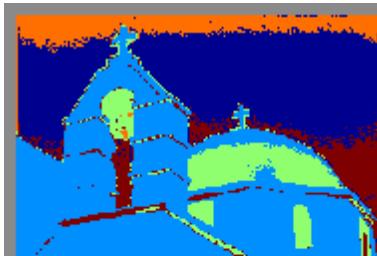
k -means worked!

For this picture, histogram splitting failed!

Color Only (3D)



LAB-space

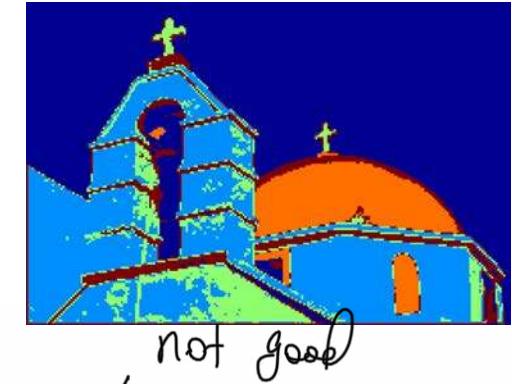
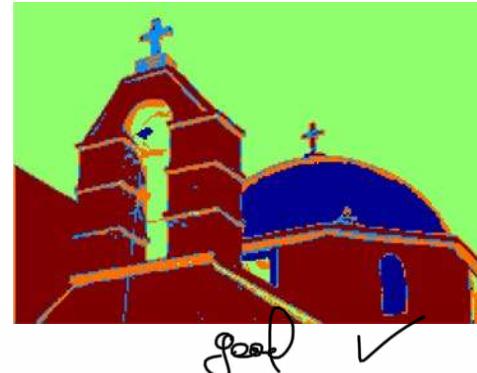
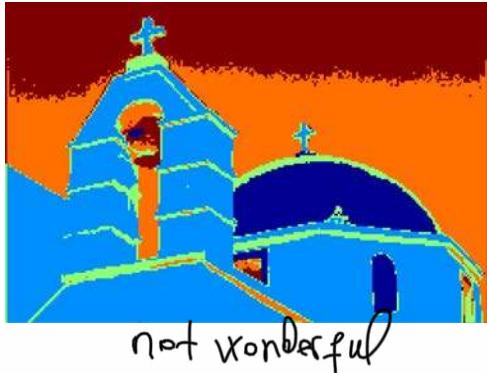
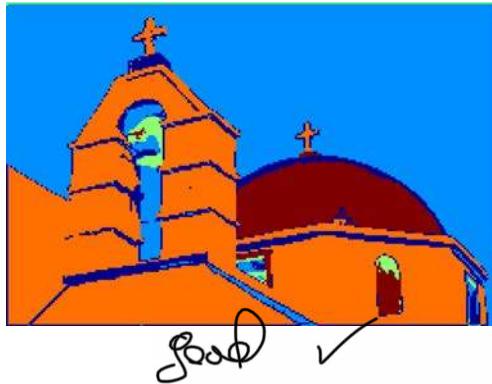


Dome, roof, walls ✓

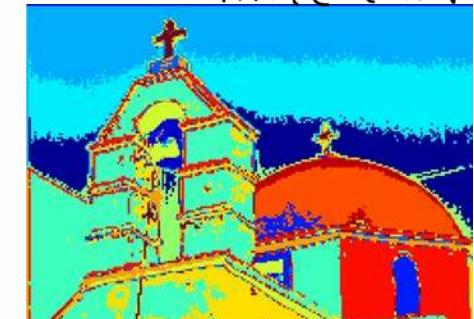
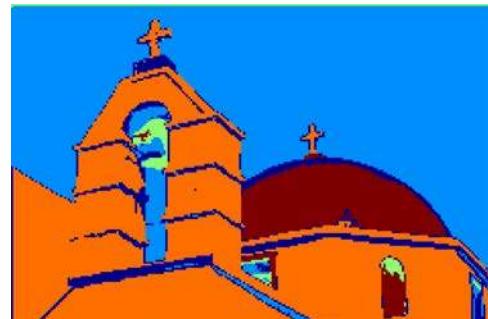
8 iterations for k=5

Color Only (3D)

Different Initializations for $k=5$

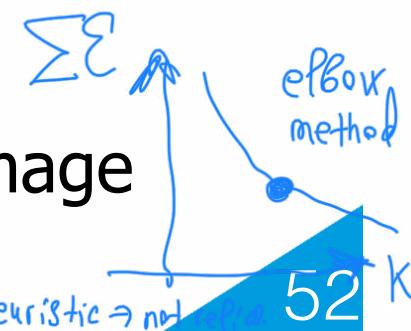


Different values of k

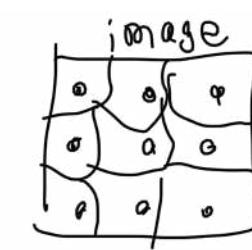
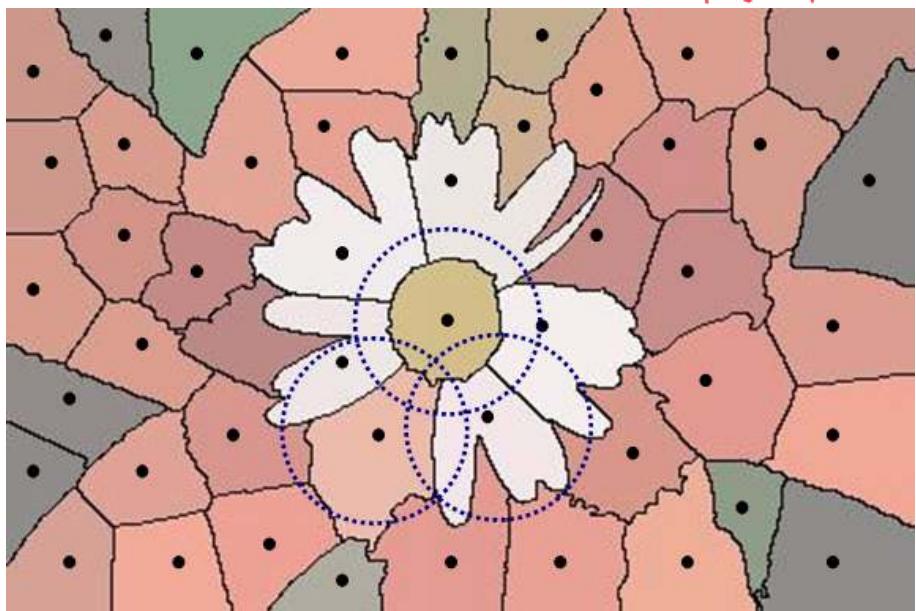


How # Clusters?

- Different results for different initializations.
- For k well chosen, the results are good for this image because it features bright and distinctive colors.



XY + Color (5D)



Start with
seeds on a
regular grid

$$E(\mathcal{C}_1, \dots, \mathcal{C}_k, \mathbf{c}_1, \dots, \mathbf{c}_k) = \sum_j \sum_{i \in \mathcal{C}_j} d(\mathbf{x}_i, \mathbf{c}_j)^2$$

$$\mathbf{x} = \begin{bmatrix} u \\ v \\ L \\ a \\ b \end{bmatrix} \text{ or } \begin{bmatrix} u \\ v \\ I \\ 0 \\ 0 \end{bmatrix}$$

L color Intensity

$$d(\mathbf{x}, \mathbf{c})^2 = \frac{(\mathbf{x}[0] - \mathbf{c}[0])^2 + (\mathbf{x}[1] - \mathbf{c}[1])^2}{h_s^2} + \frac{(\mathbf{x}[2] - \mathbf{c}[2])^2 + (\mathbf{x}[3] - \mathbf{c}[3])^2 + (\mathbf{x}[4] - \mathbf{c}[4])^2}{h_r^2}$$

cluster Centers different weights

Run K-Means algorithm with regularly spaced seeds on a grid and using a distance that is a weighted sum of distances in image space and in gray level/color space.

→ small region centered around seed

→ Superpixels

SLIC Superpixels

1024x1024

Fewer seeds



Bigger regions

Superpixel
segmentation

256x256



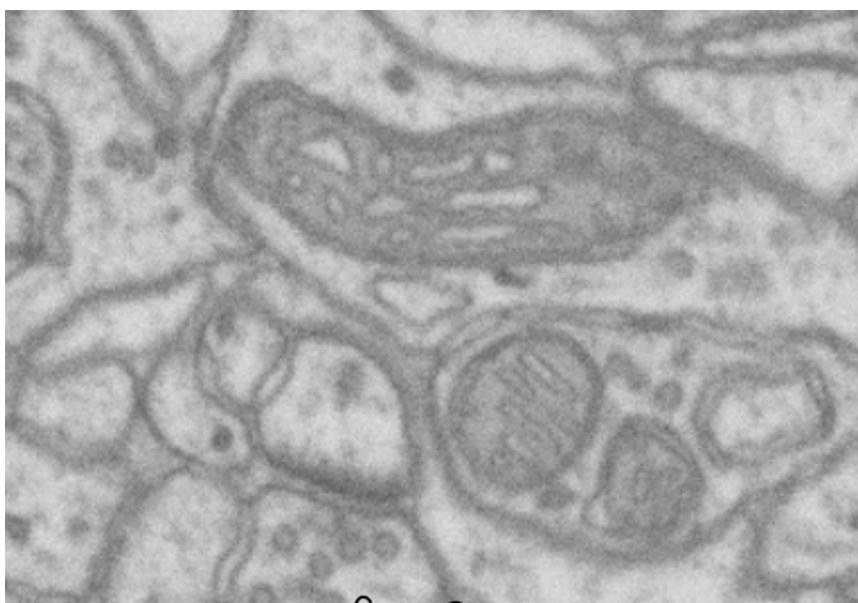
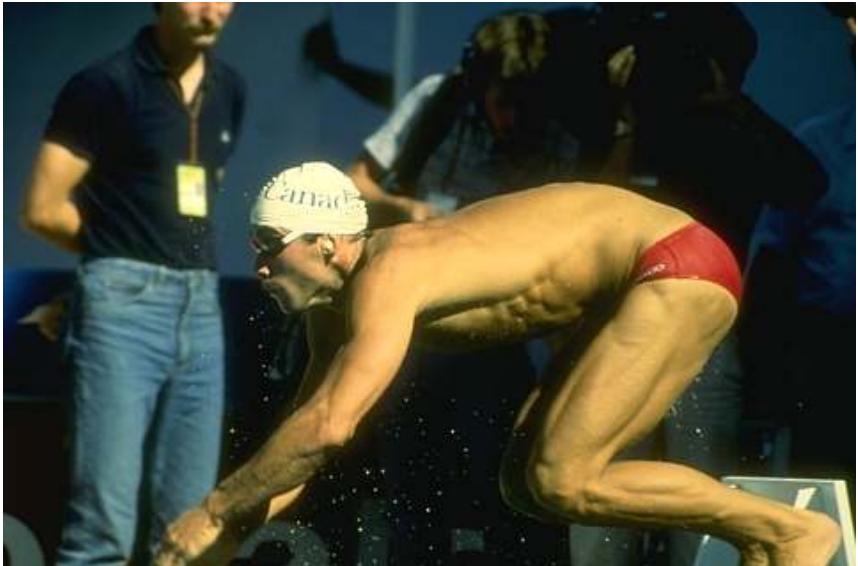
256x256

Think about
grass issue!

64x64

- Superpixel segmentations with centers on a 64x64, 256x256, and 1024x1024 grid.
- Can be used to describe the image in terms of a set of small regions. ↗ on which you can compute statistical (Local) 54

Color or Black and White

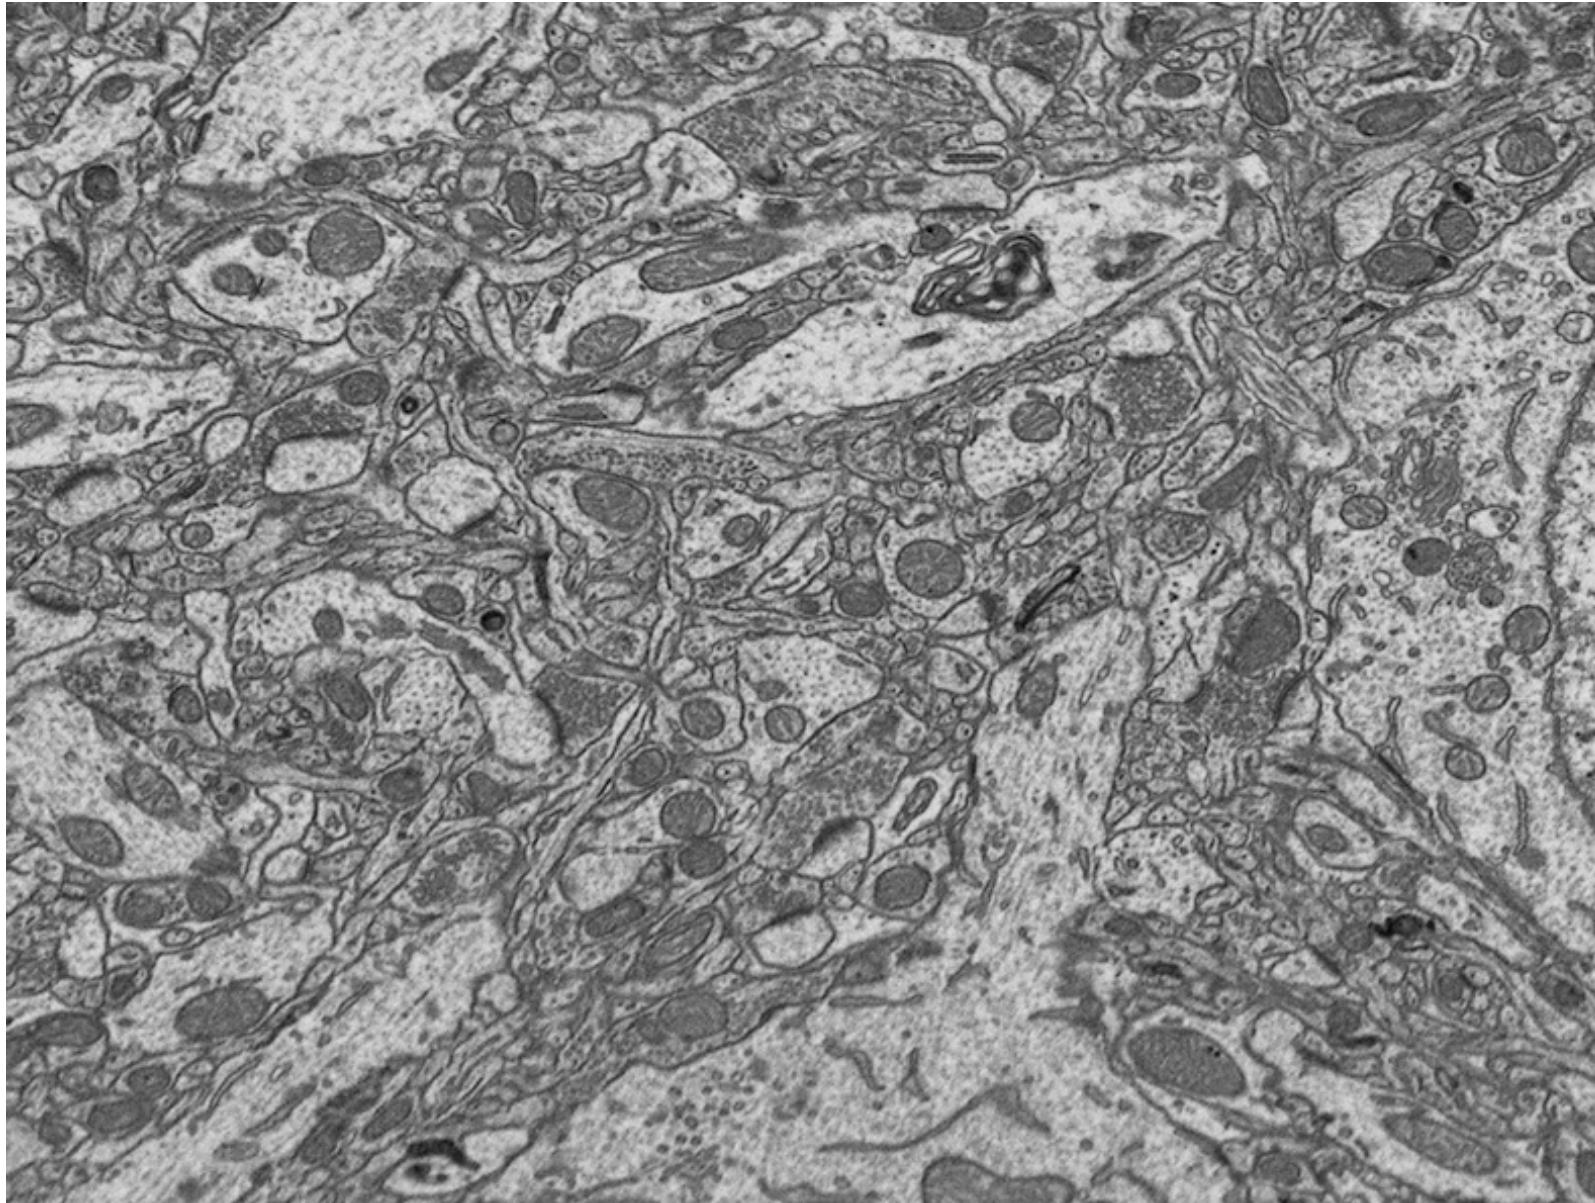


gray - prep



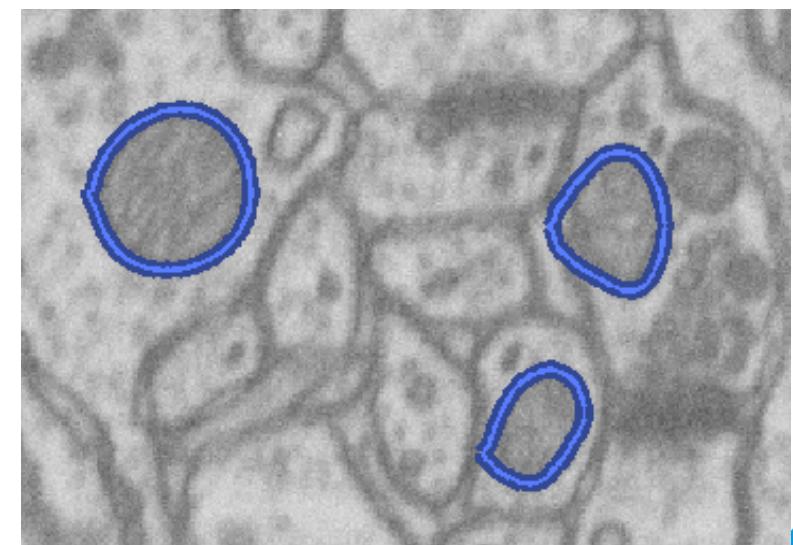
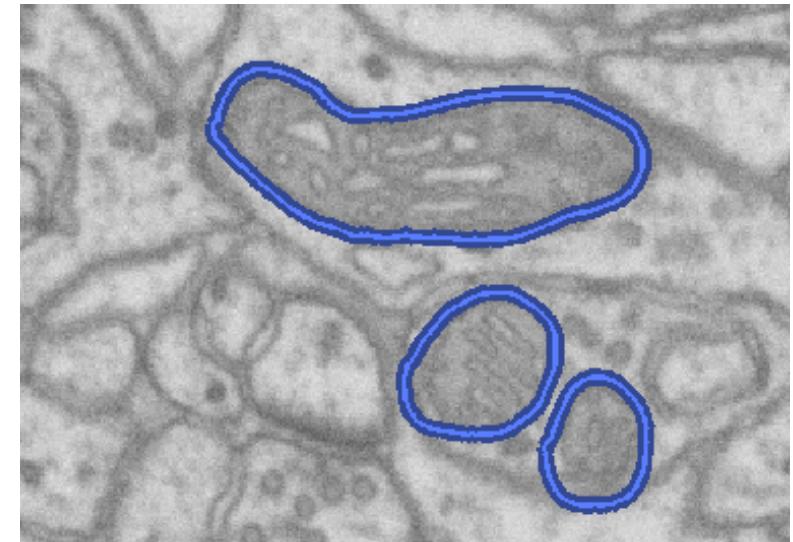
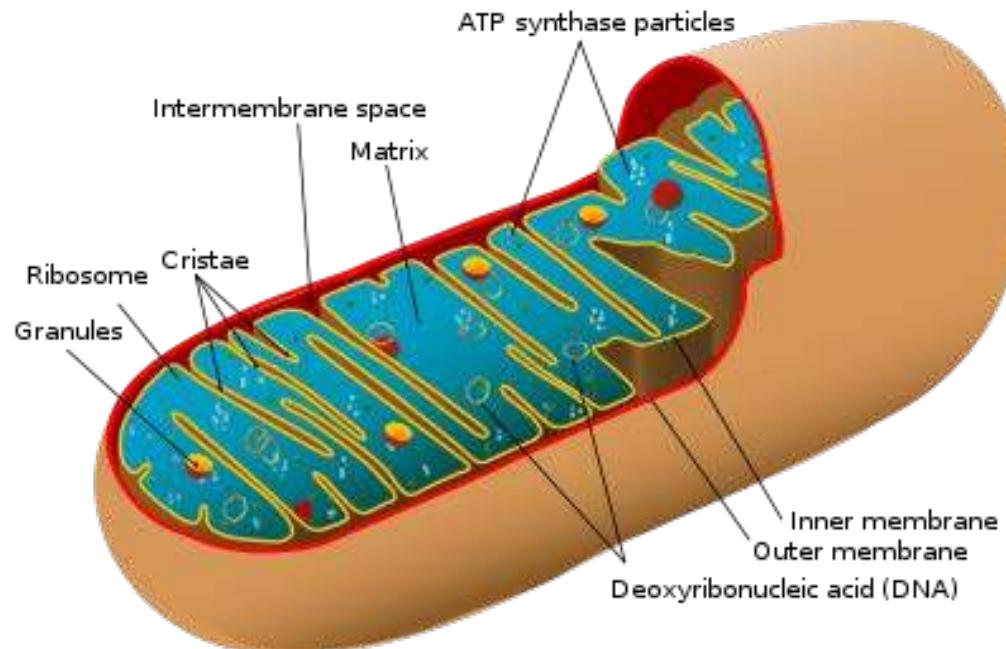
It works in both cases.

Electron Microscopy



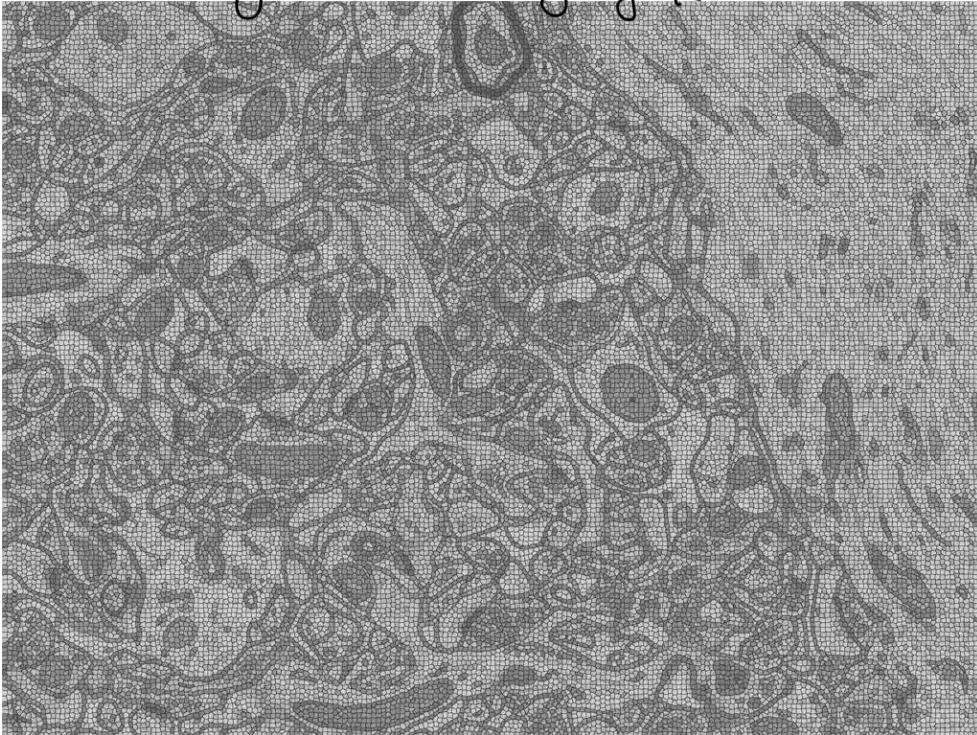
Mitochondria Segmentation

and find boundaries
Slice of mitochondria

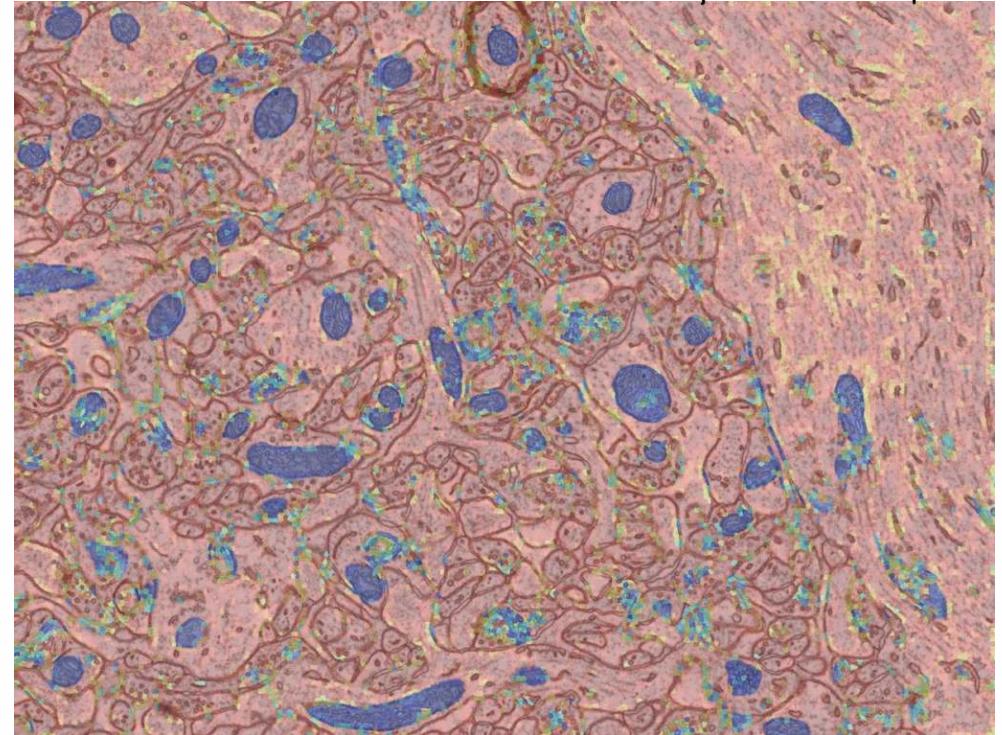


Assigning Probabilities

solely based on gray-pixels



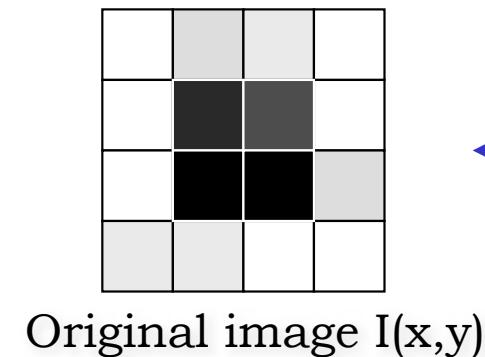
after classifier



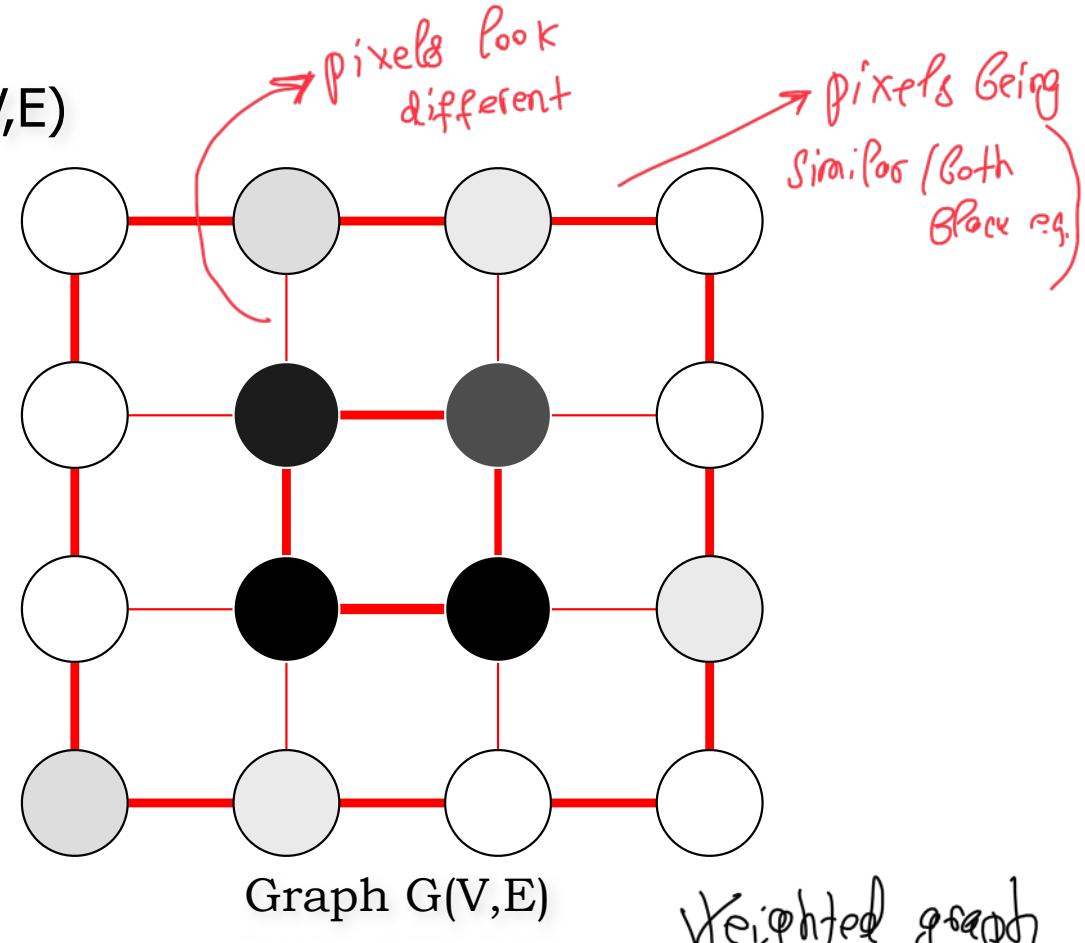
- Compute image statistics for each superpixel.
- Train a classifier to assign a probability to be within a mitochondria.
 - Based on image statistics
 - Red → unlikely to be a mitochondria
 - Blue → likely
- Can be used to produce segmentations using the graph-based techniques we will describe next.

Images as Graphs

An image $I(x,y)$ is equivalent to a graph $G(V,E)$



Original image $I(x,y)$



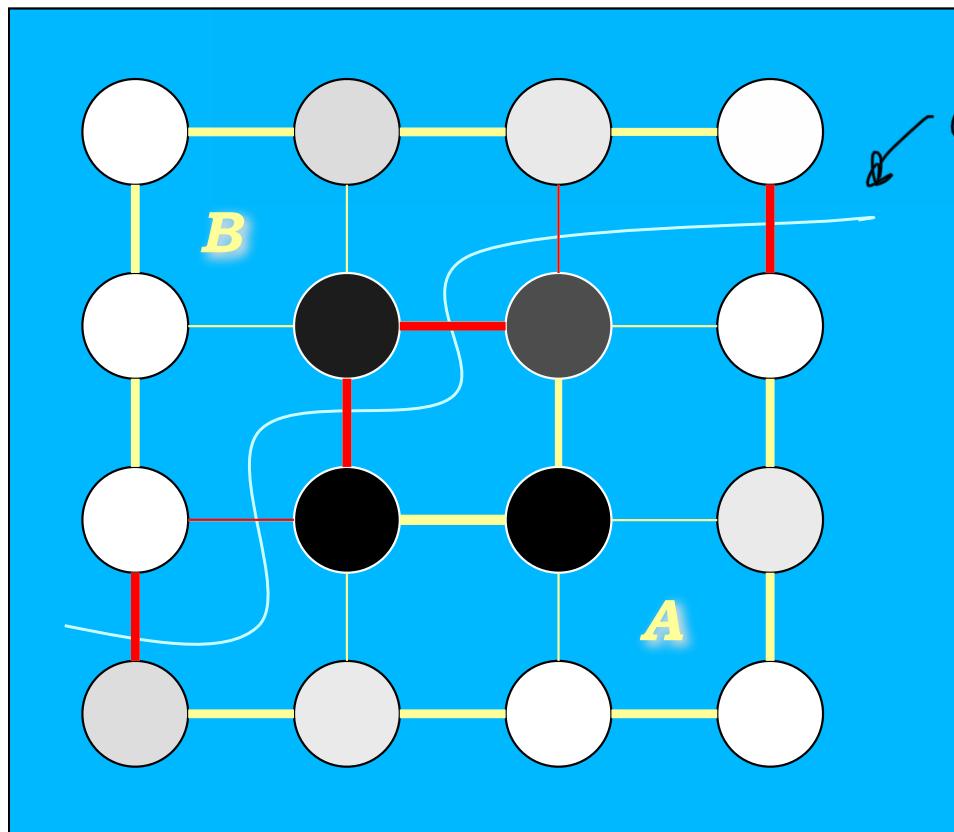
Graph $G(V,E)$

Weighted graph

- V is a set of vertices or nodes that represent individual pixels.
- E is a set of edges linking neighboring nodes together. The weight or strength of the edge is proportional to the similarity between the vertices it joins together.

Graph-Cut

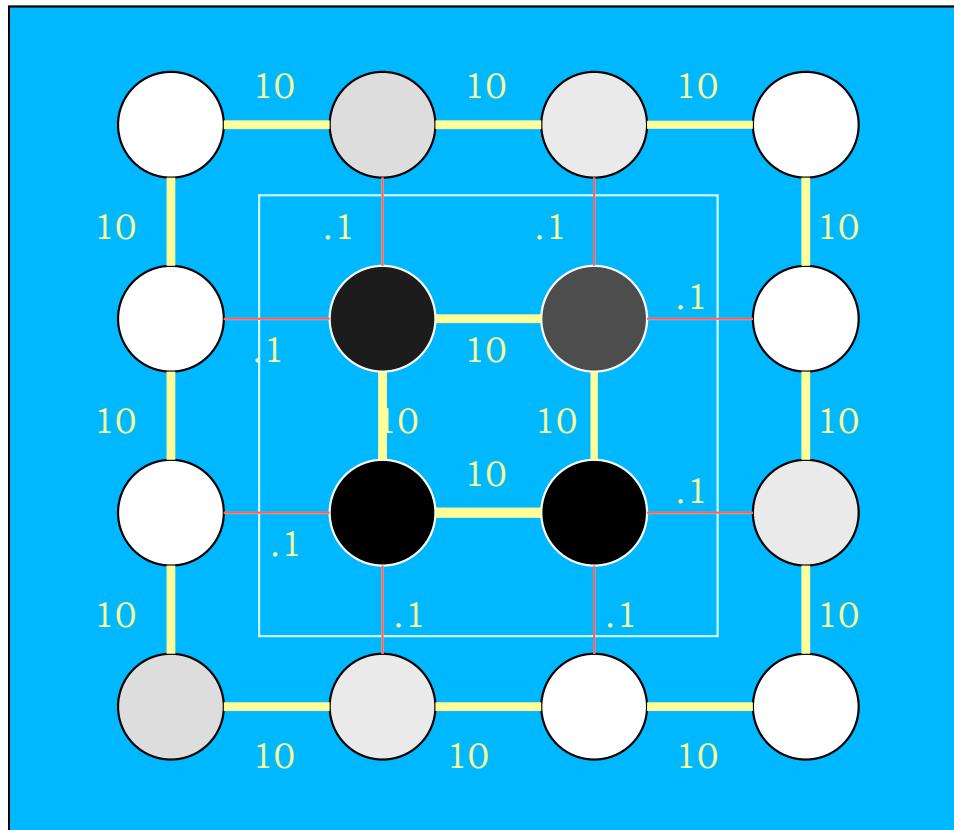
Segment given image into 2 regions



Cuts graph into 2 independent connected components

A cut through a graph is defined as the total weight of the links that must be removed to divide it into two separate components.

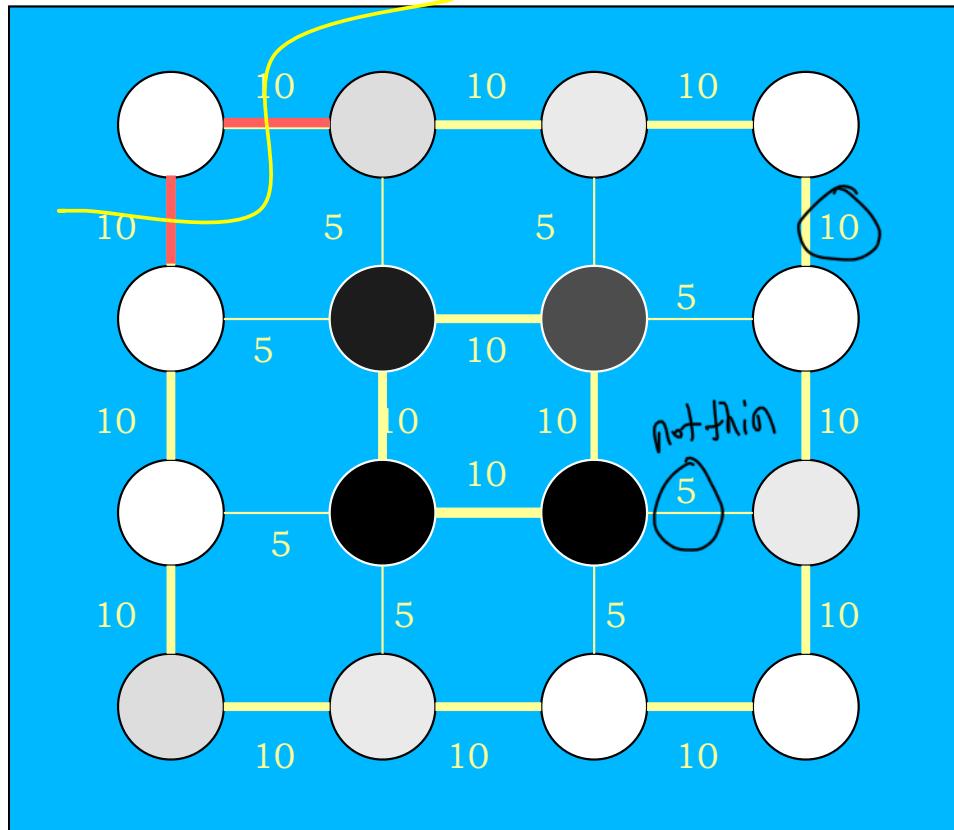
Min-Cut



- Find the cut through the graph that has the overall minimum weight, which can be done effectively. → *cut edges that are thinnest*
- It should be the subset of edges of least weight that can be removed to partition the graph.
- Since weight encodes similarity, this should be equivalent to partitioning the graph along the boundary of least similarity.
(thin edges)

Trivial Cut

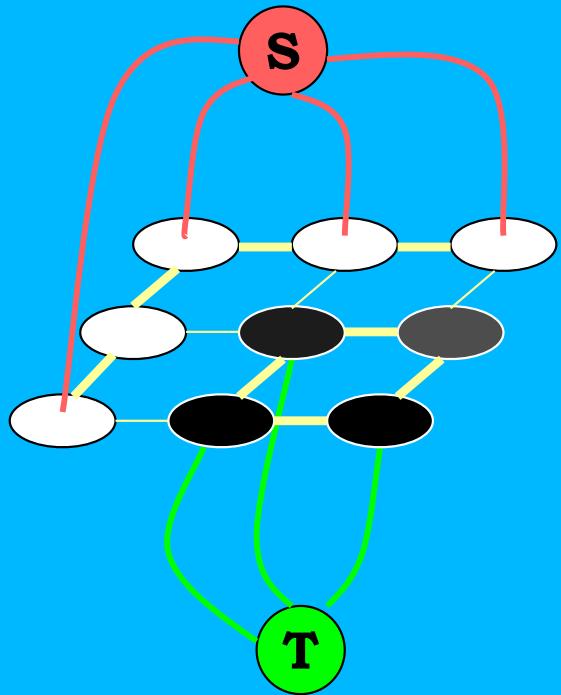
cheapest cut, hmm...



"Thin" edges are not that thin

- Has a preference for short cuts, which may sometime result in trivial solutions.
 - Must be constrained to avoid them.

ST Min-Cut



- Introduce two special nodes called source (S) and sink (T)
- S and T are linked to some image nodes by links of very large weight that will never be selected in a cut.
- Find the minimum cut separating the source from the sink.

→ will avoid trivial cut

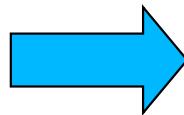
--> The problem becomes deciding how to connect S and T to the image nodes.

Interactive ST Min-Cut

Human intervention: You had to choose foreground and background pixels



Run
graph-cut



User-selected pixels connected to S (red)
and pixels connected to T (cyan)

Minimum S-T cut

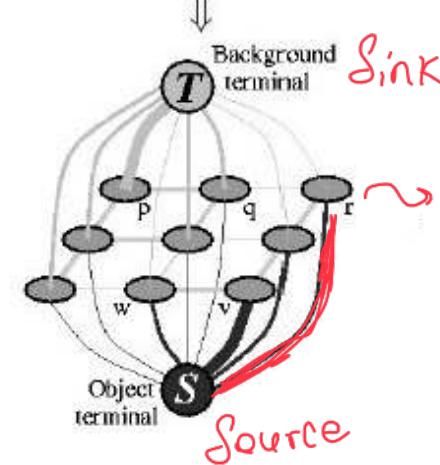
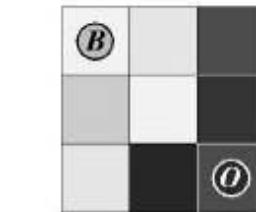
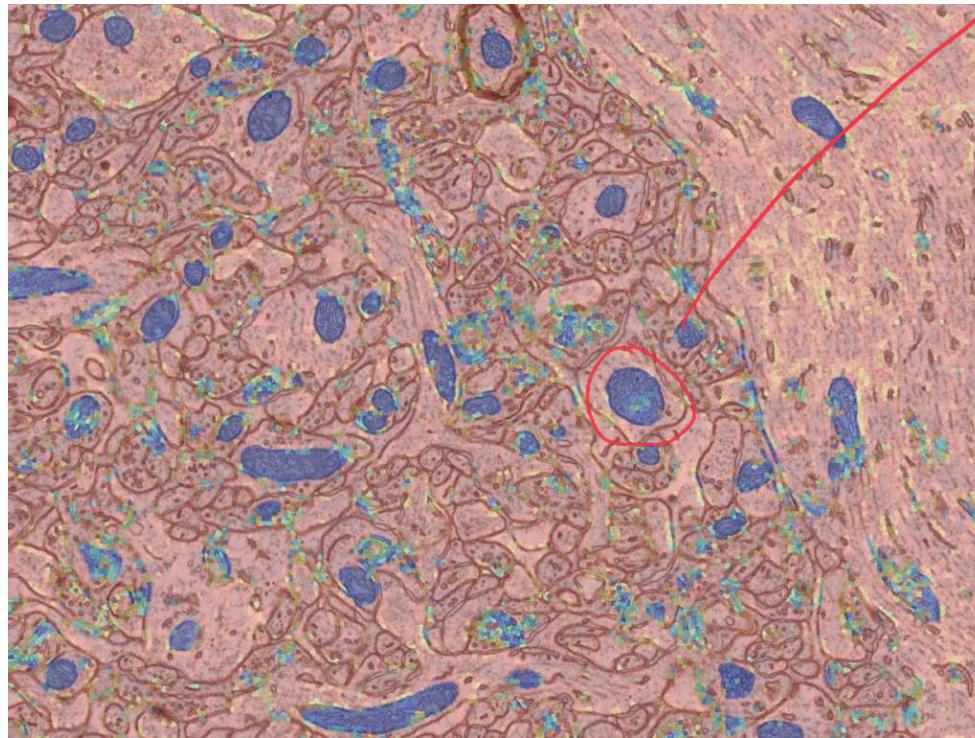
↳ background

↳ background

--> If we have a good initial 'guess' to tell us how to link the source and sink to the image, we will get an optimal segmentation.

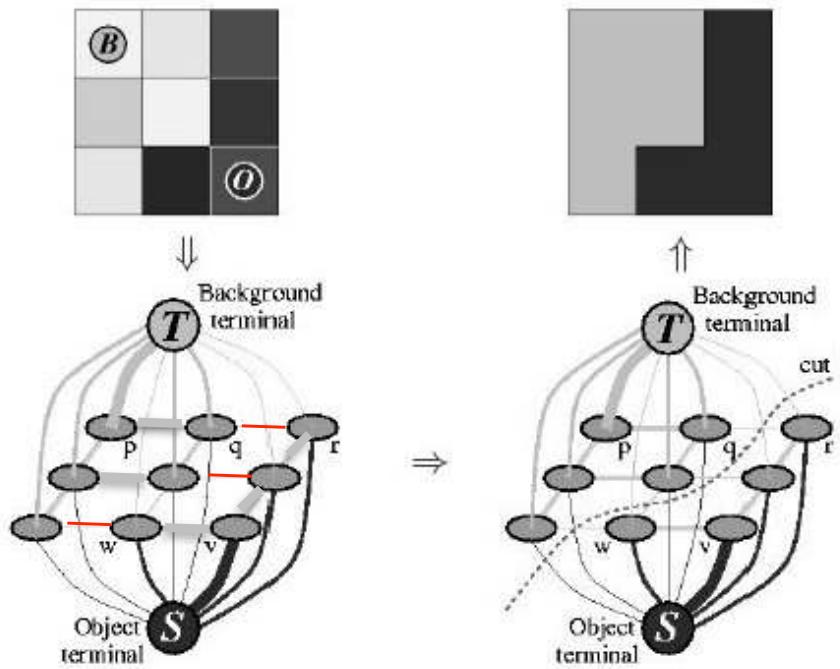
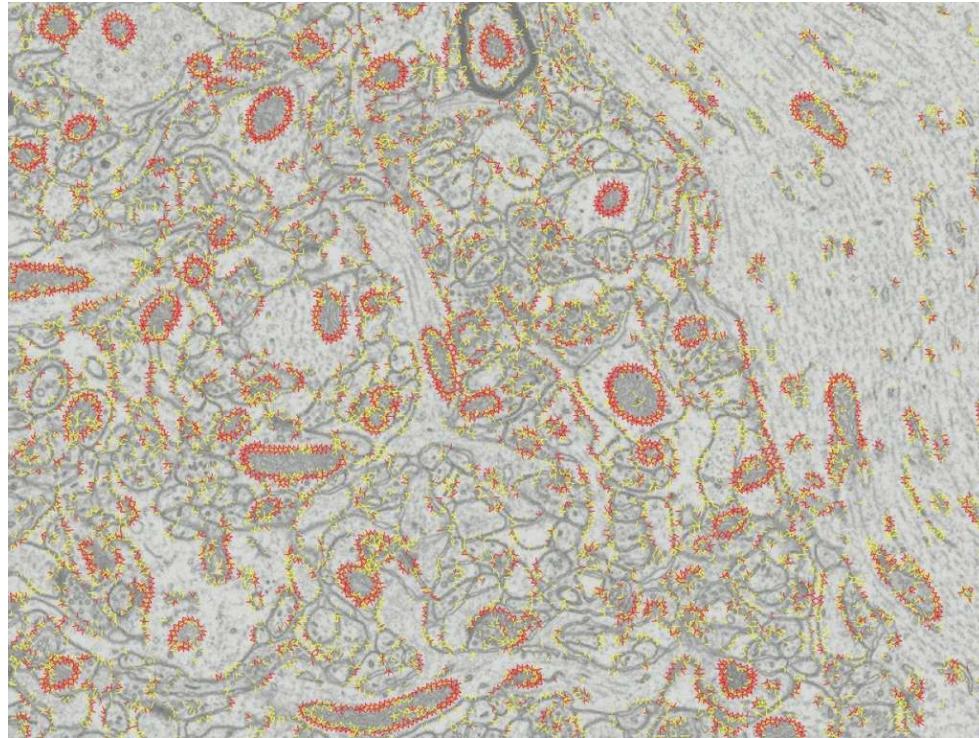
Back To Mitochondria

Probabilities those supervoxels will be within mitochondria or background



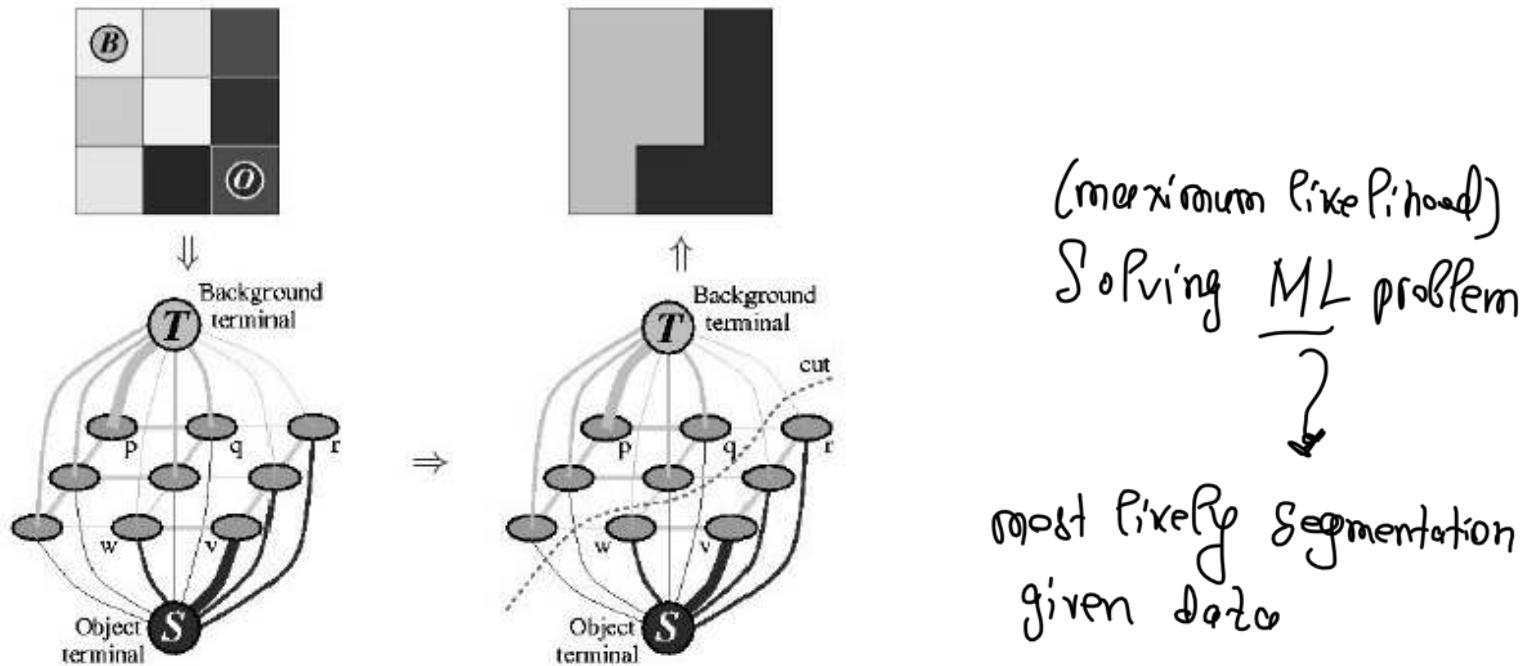
- A high probability of being a mitochondria can be represented by a strong edge connecting a supervoxel to the source and a weak one to the sink.
- And conversely for a low probability.

Back To Mitochondria



Another classifier can be trained to assign a high-weight to edges connecting supervoxels belonging to the same class and a low one to others.

Minimizing a Loss Function

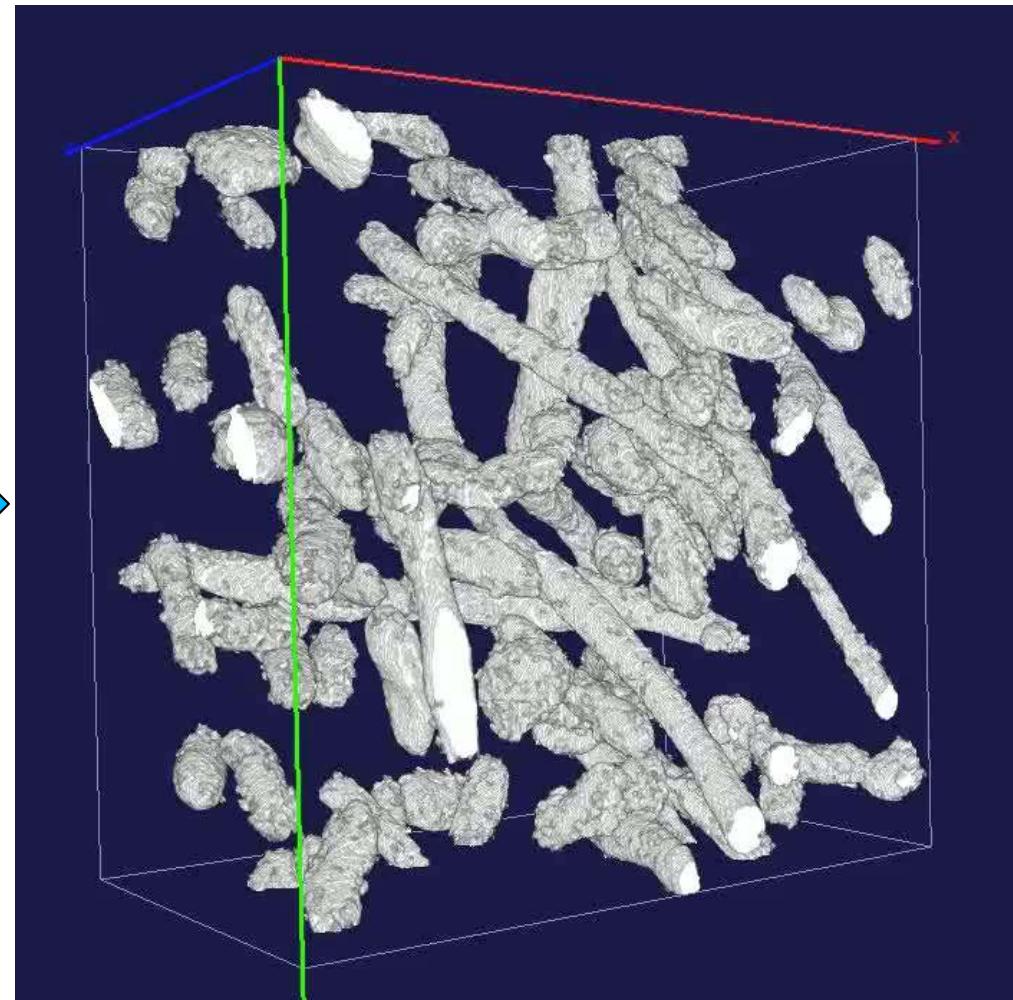
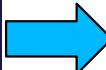
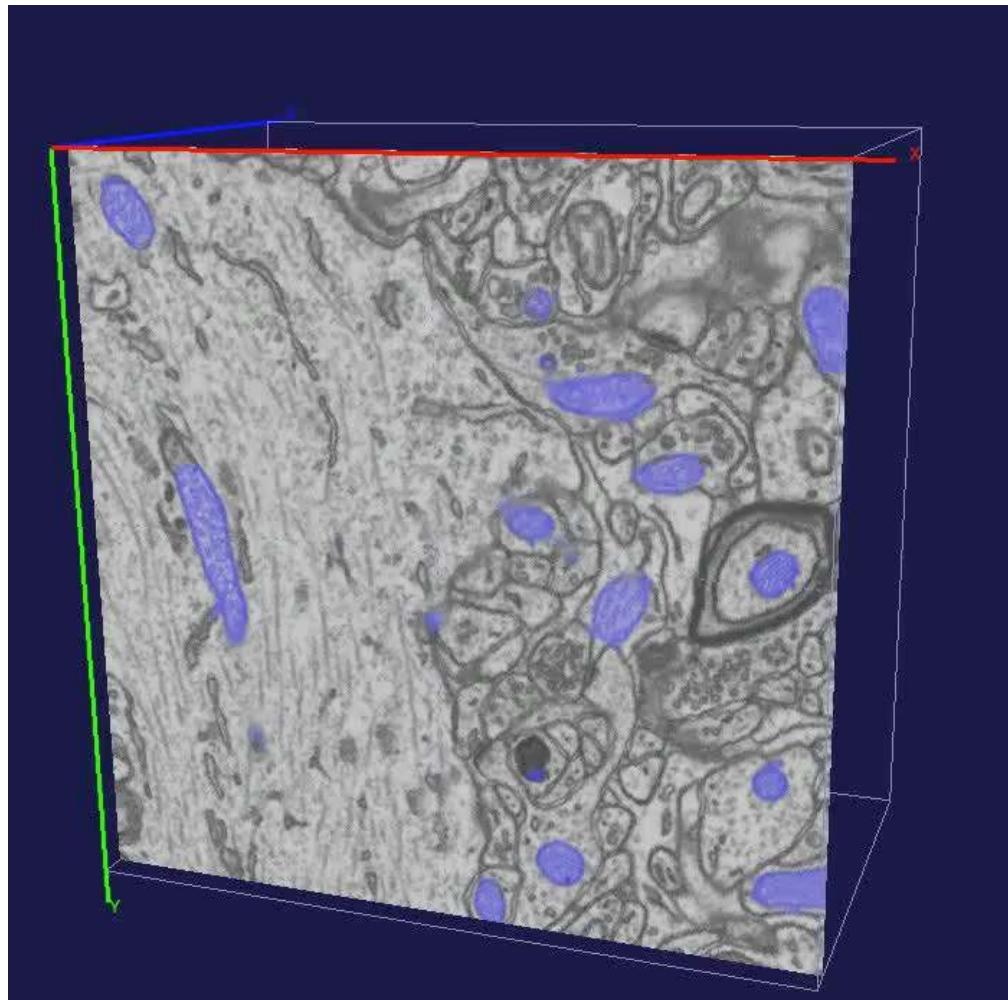


In probabilistic terms, this amounts to minimizing

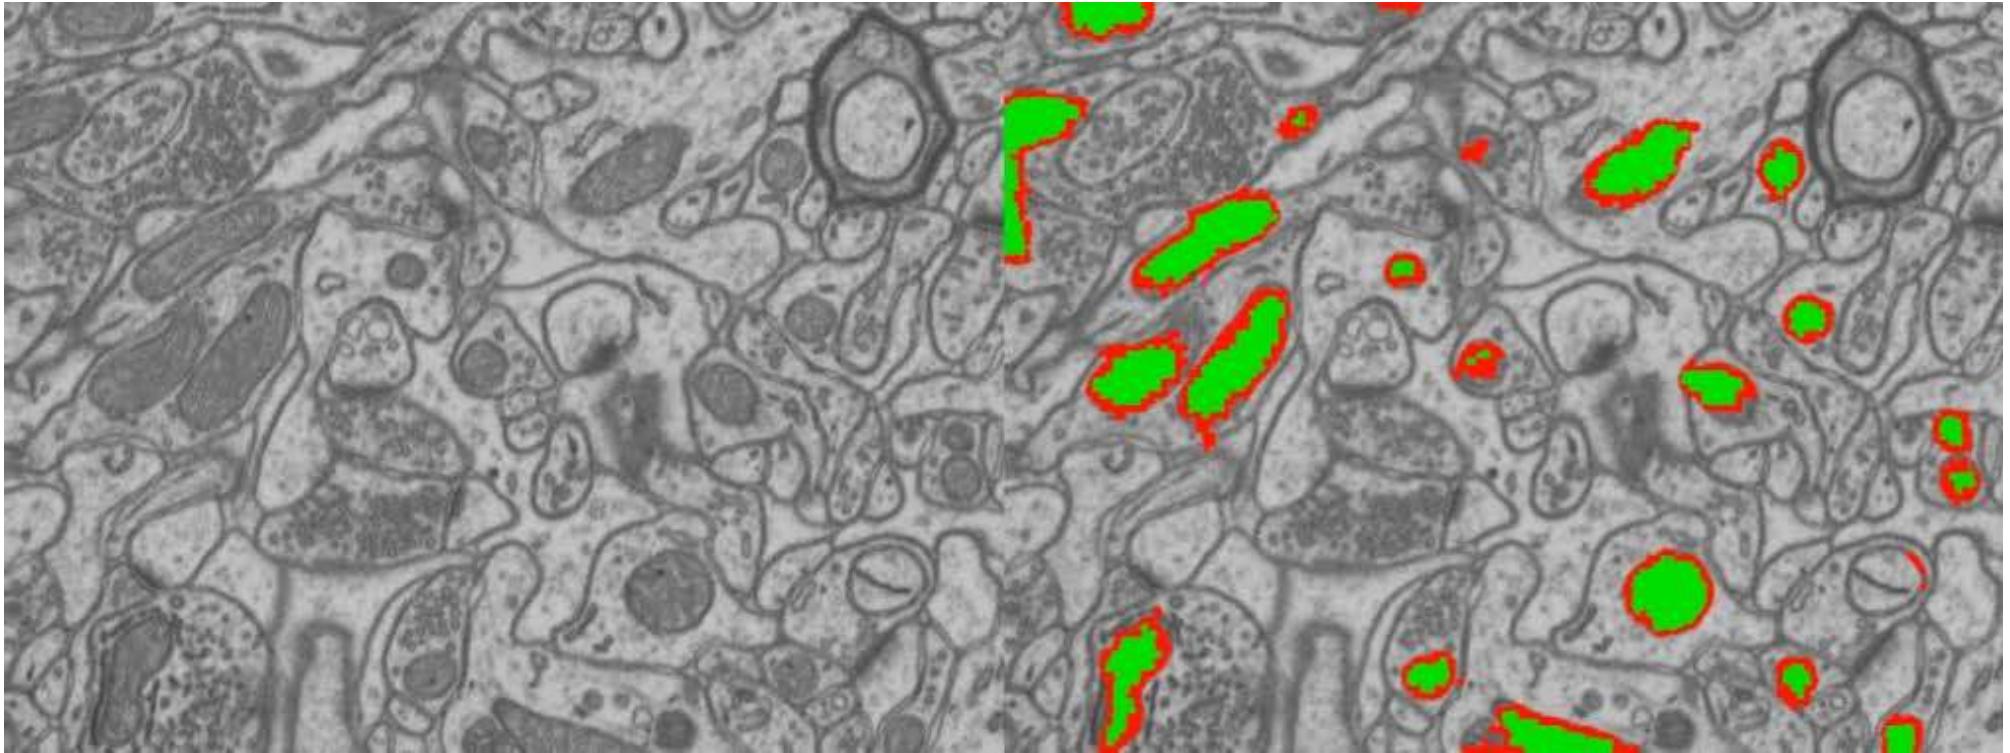
$$E(y|x, \lambda) = \sum_i \underbrace{\psi(y_i|x_i)}_{\text{unary term}} + \lambda \sum_{(i,j) \in \mathcal{E}} \underbrace{\phi(y_i, y_j|x_i, x_j)}_{\text{pairwise term}},$$

with respect to the set of labels $y = [y_1, \dots, y_n]$ given the set of supervoxels $[y_1, \dots, y_n]$, where E is negative log-likelihood of the labels being correct.

3D Mitochondria



Modeling Membranes



Here we use three classes instead of two:

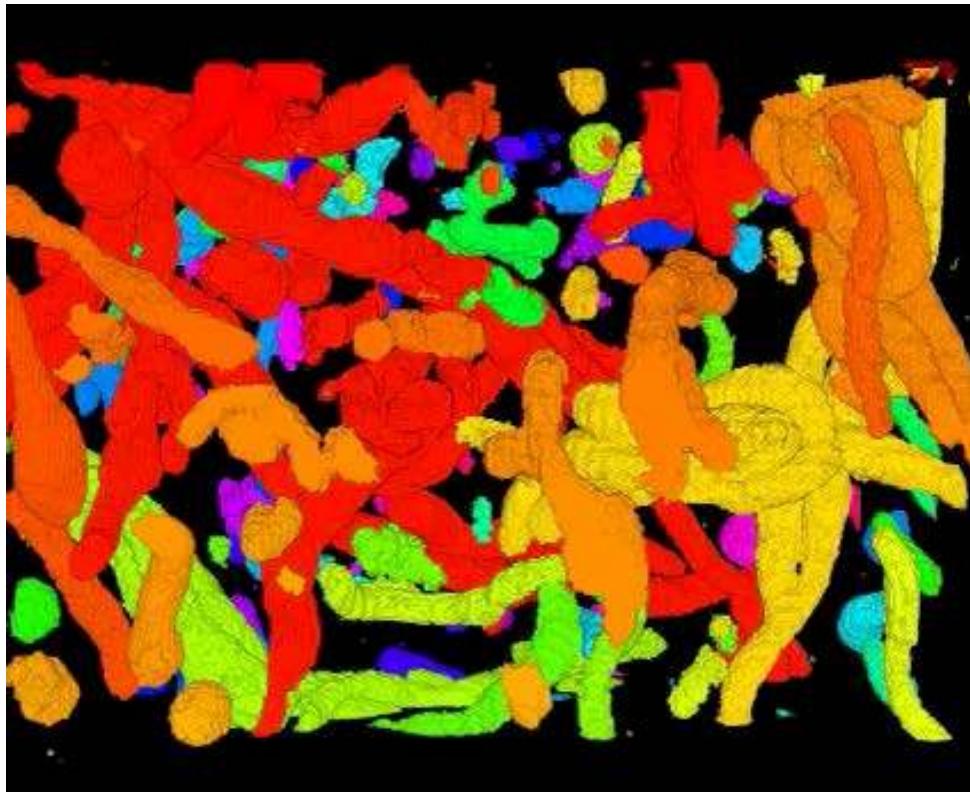
- Inside
- Membrane
- Everything else (Background)

Graph-cut algo works ✓

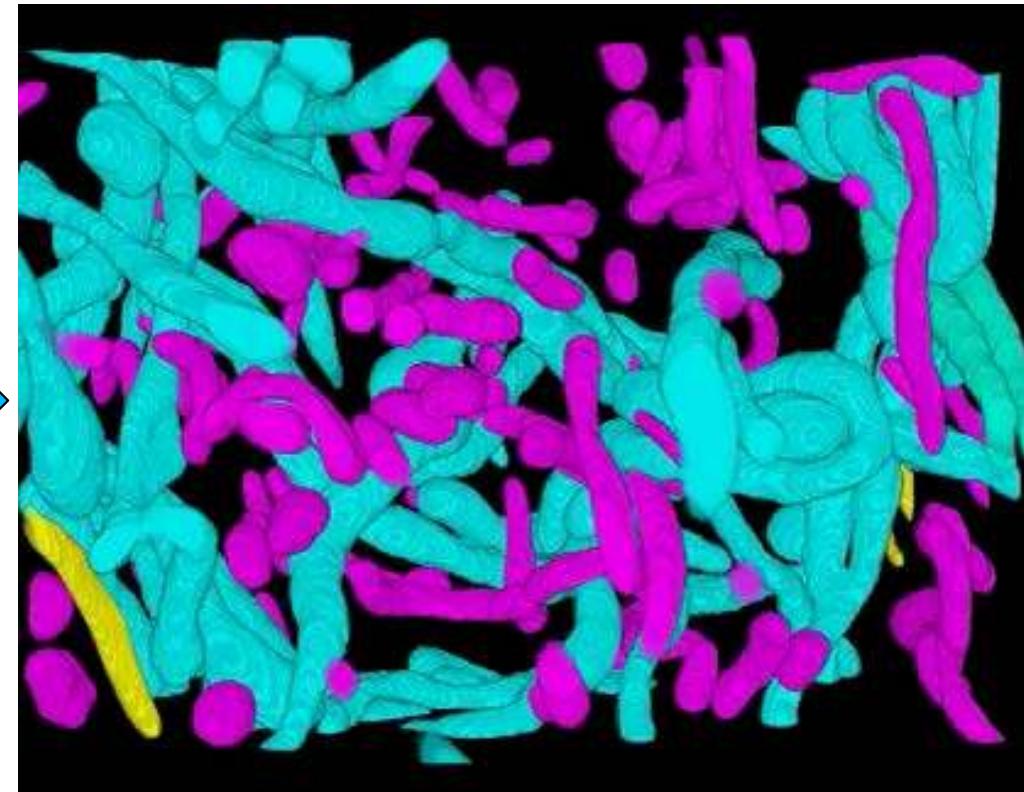
—> Because the inside is fully enclosed by the membranes, we can still find a global optimum.

Speeding up the Analysis Process

3.21 μm \times 3.21 μm \times 1.08 μm : 53 mitochondria



Automated result



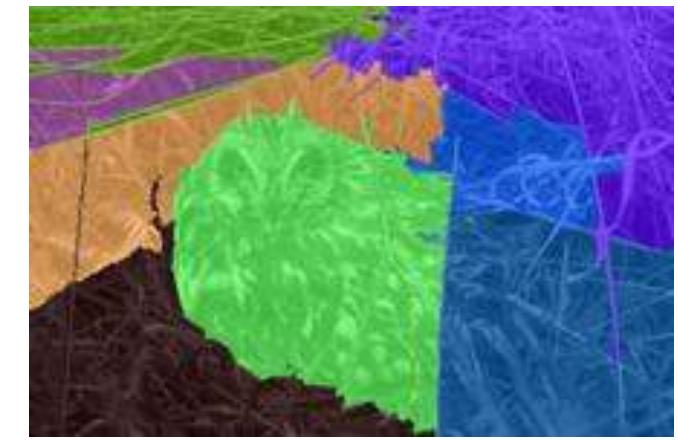
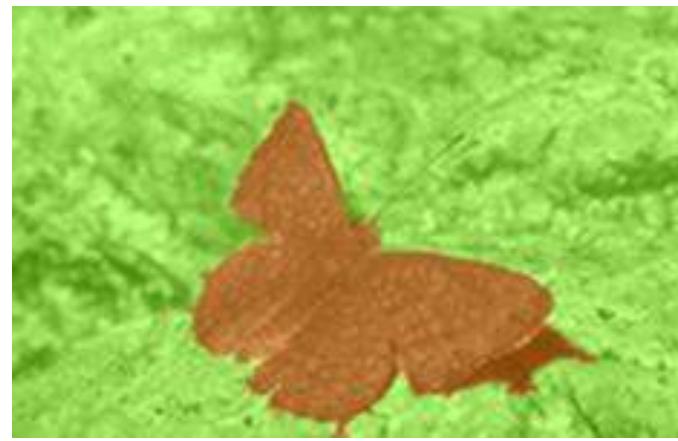
Interactively cleaned-up result

Segmenting mitochondria from volumes of images

- By hand: 6 hours.
 - Semi-automatically: 1.5 hours
- Substantial time saving for the neuroscientists.

Graph-Cut on Ordinary Images

not only in neuroscience



sink and source → should have
semantic meaning

(f_g and b_g)

Galun et al., ICCV'03

Interactive Foreground Extraction

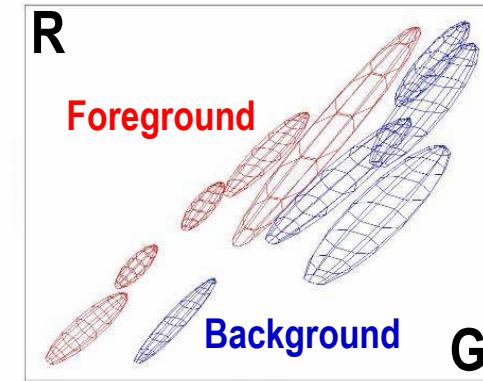
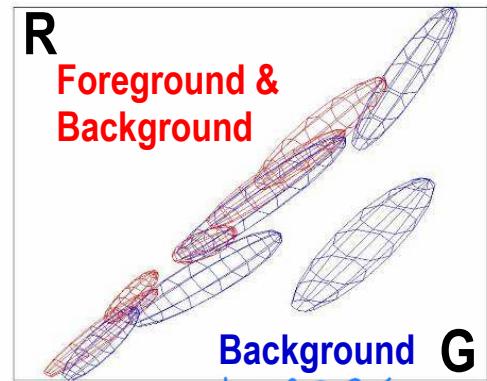


K-means: not good by itself
with combination →
→ gives good results

Automatic segmentation
of flowers



Iterated
graph cut



- K-means to learn color distributions of fg and bg
- Graph cuts to infer the segmentation

Use K-means to assign image pixels either to fg or bg, based on how similar their colors are to the k-means distribution

Relatively Easy Examples



More Difficult Examples

Including cases whose failing

Initial Rectangle

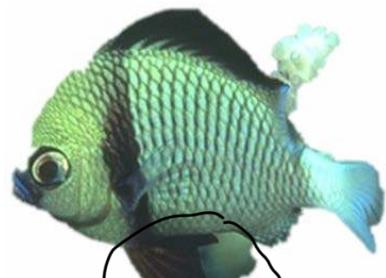


Camouflage & Low Contrast

Troubling!
Fine structure



Initial Result

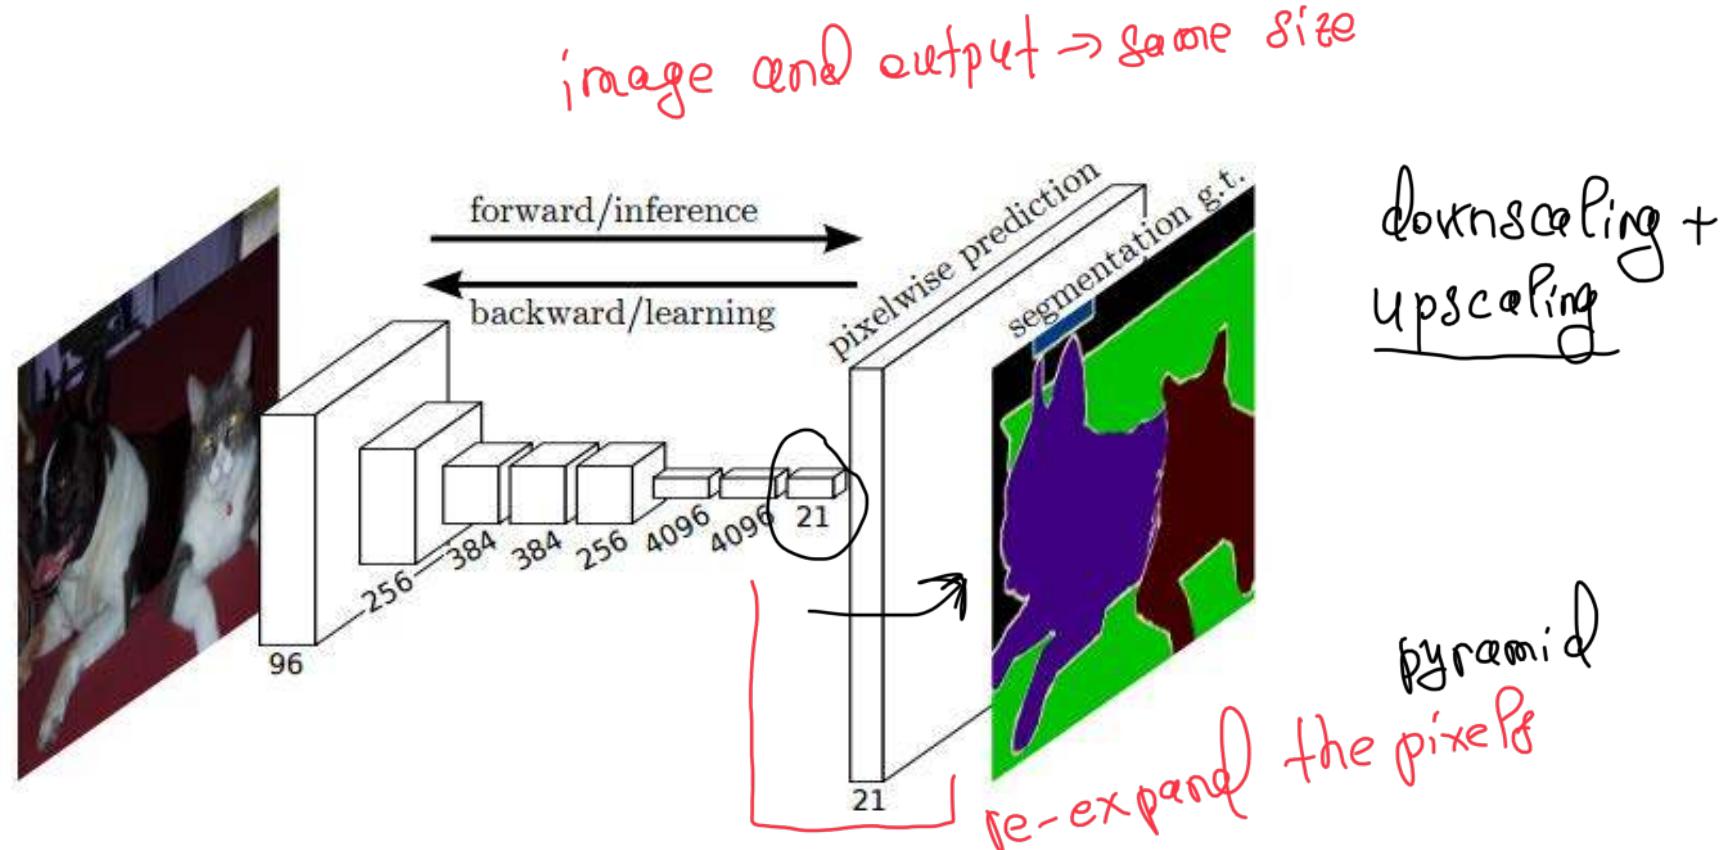


Camouflage
actually

are looking at,
Semantics of where you
No telepathy

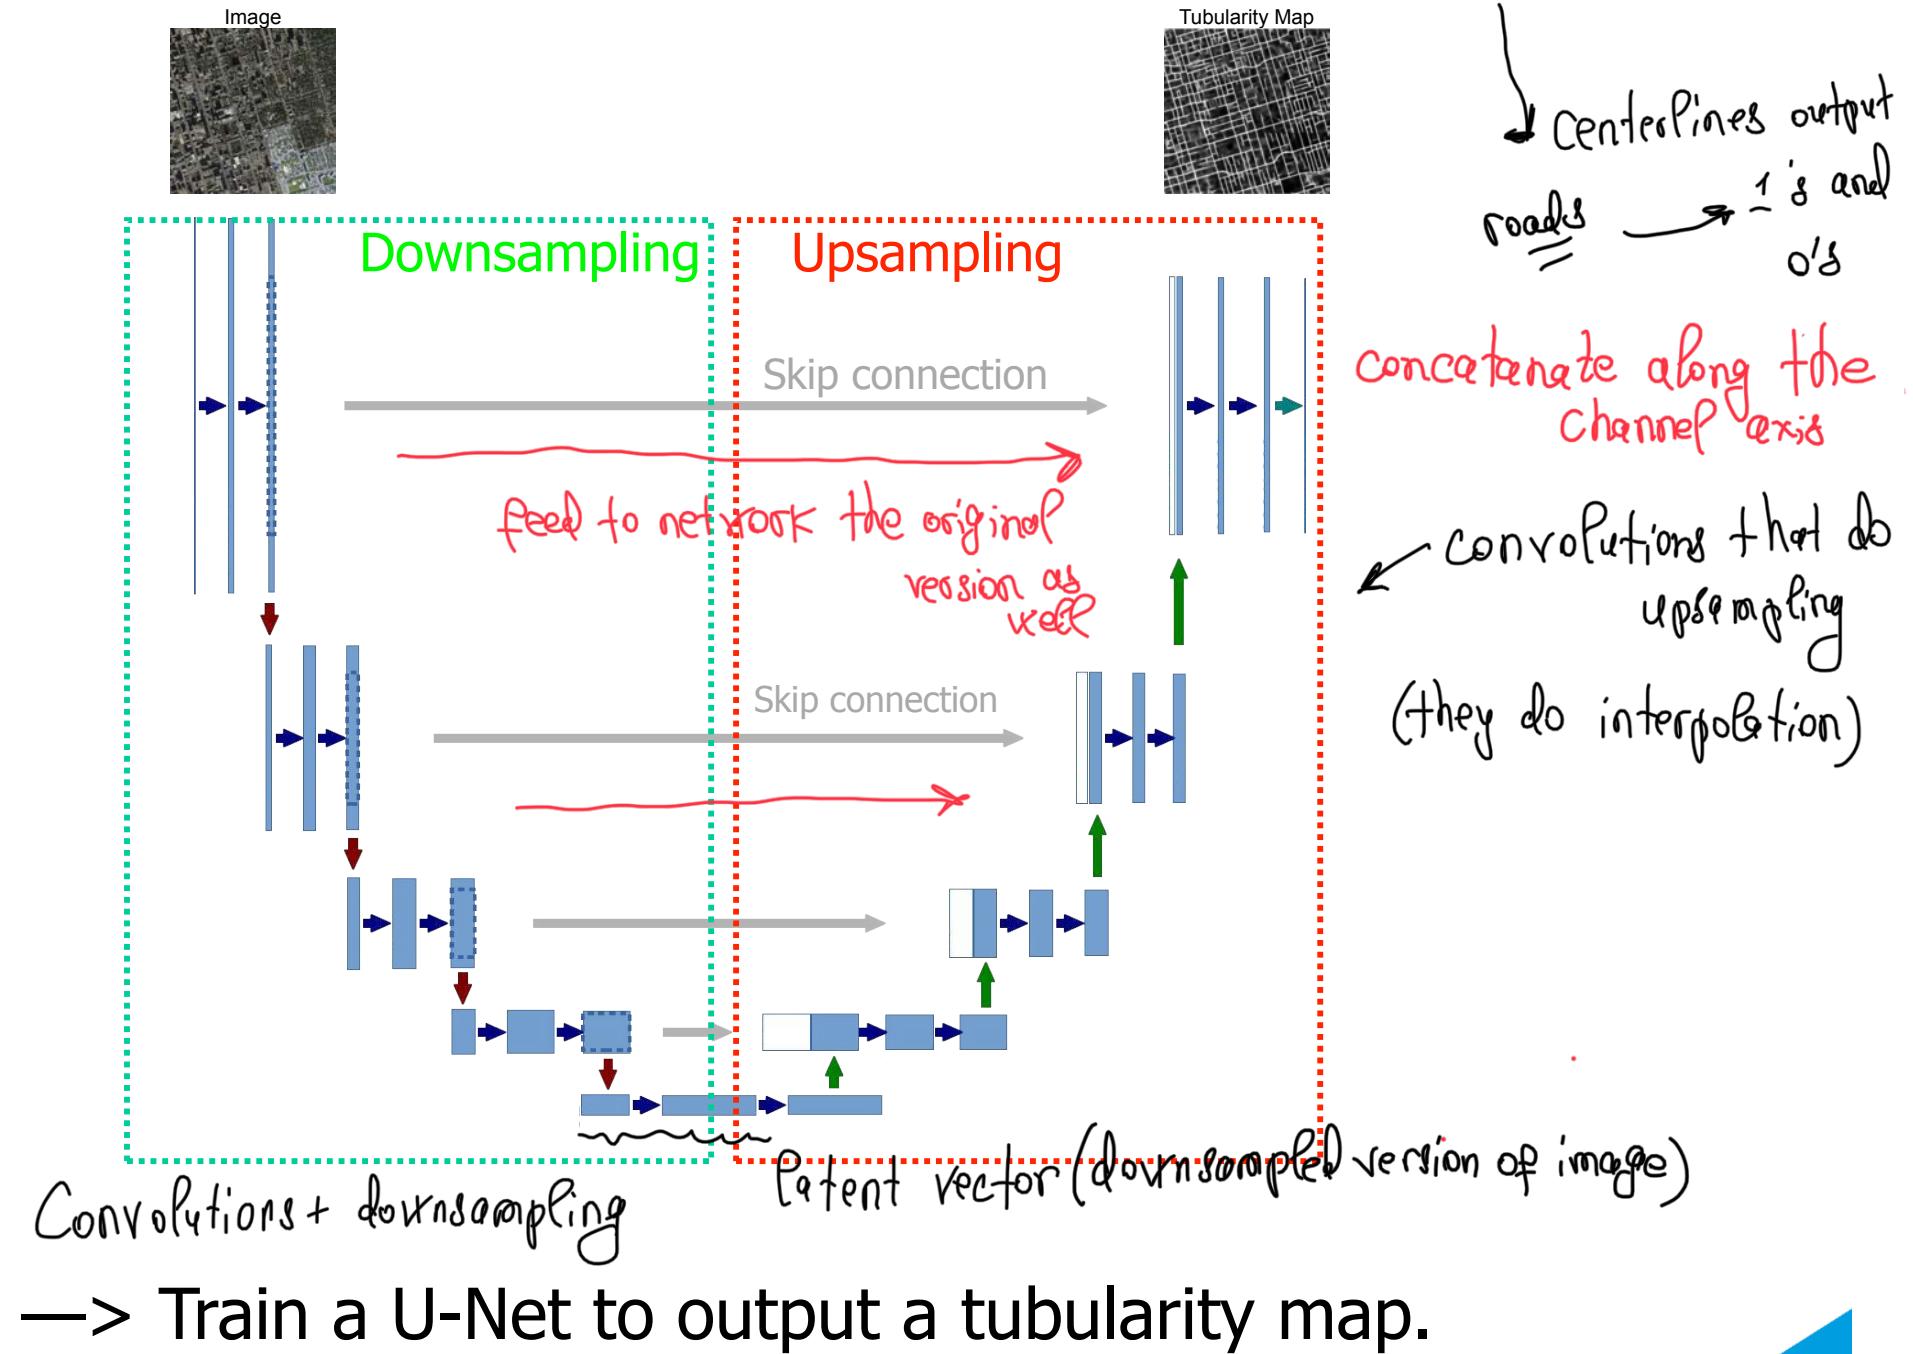


Convolutional Neural Nets



- Connect input layer to output one made of segmentation labels.
- Need layers that both downscale and upscale.
- Connect the lower layers directly to the upper ones. 

Reminder: U-Net for Delineation



U-Net for Segmentation

Train with SGD, or Adam

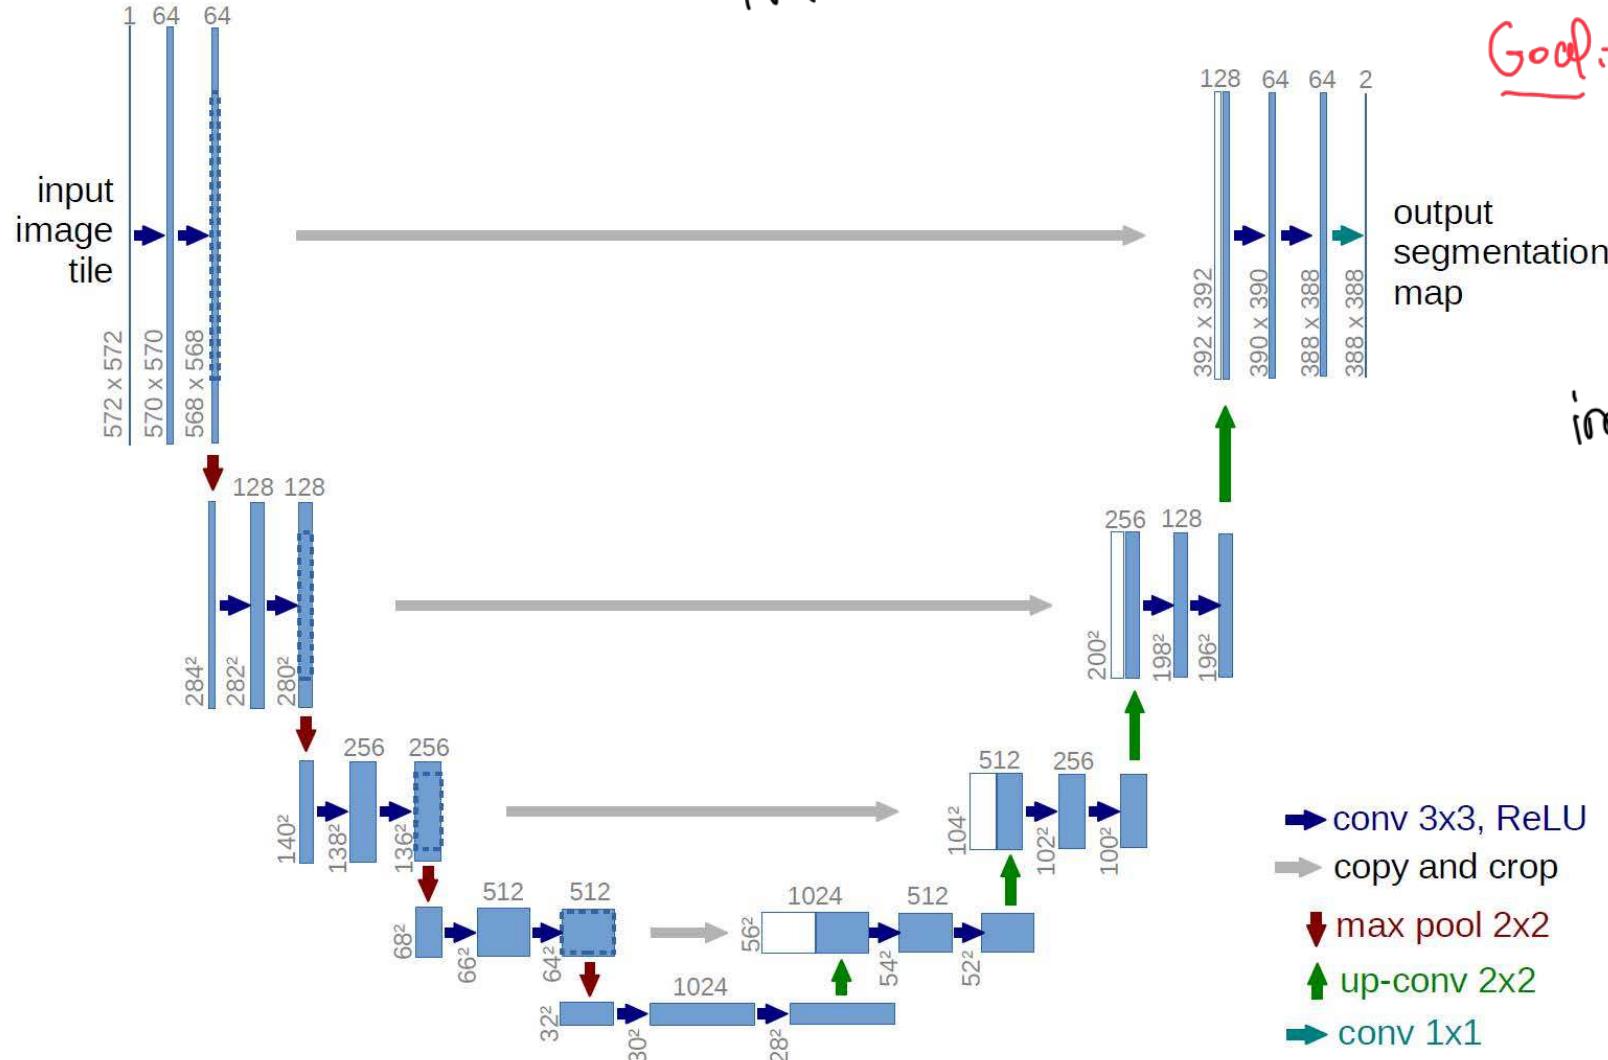
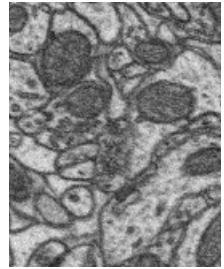


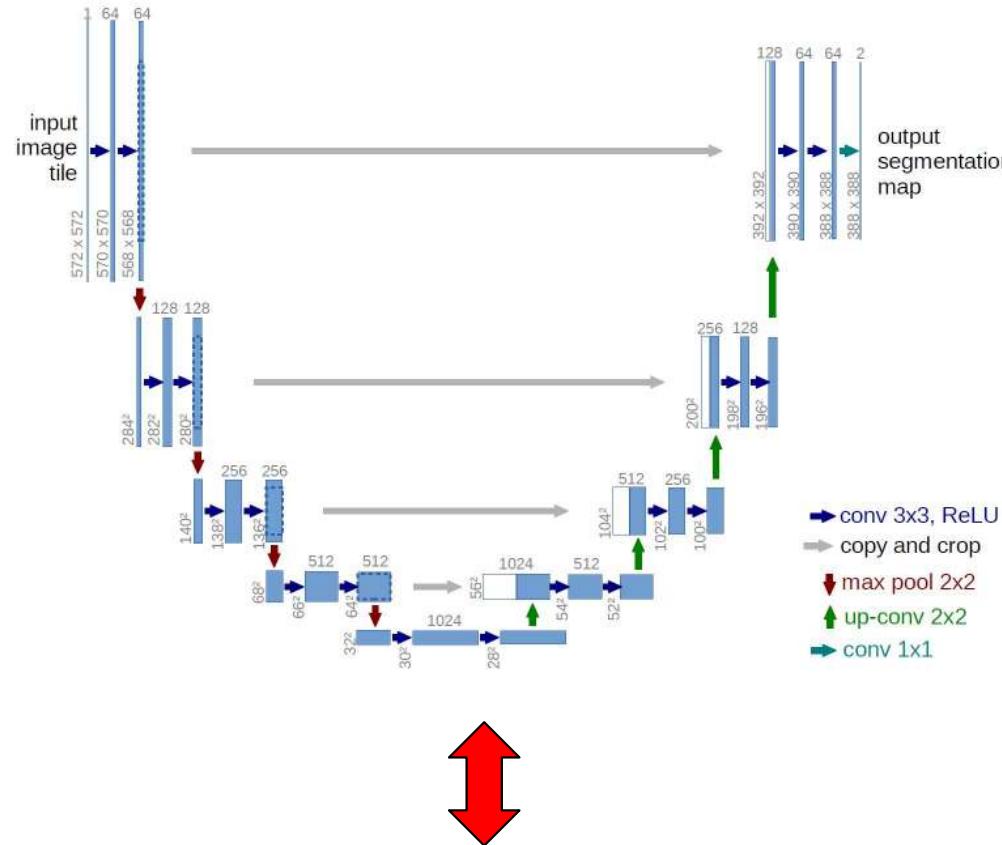
image with zeros and ones

-dria

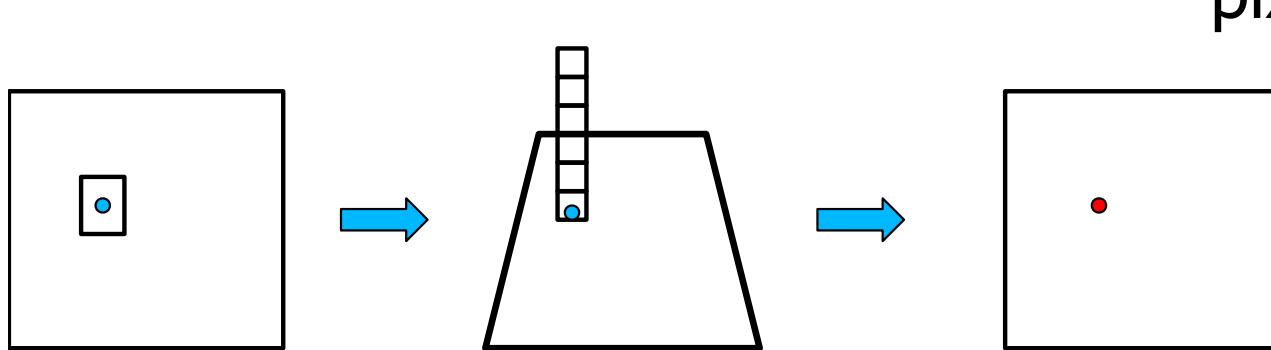
Goal: finding mitochondria

- Same architecture (in more details).
- Train it to produce a segmentation mask instead of a delineation mask.

Potential Interpretation



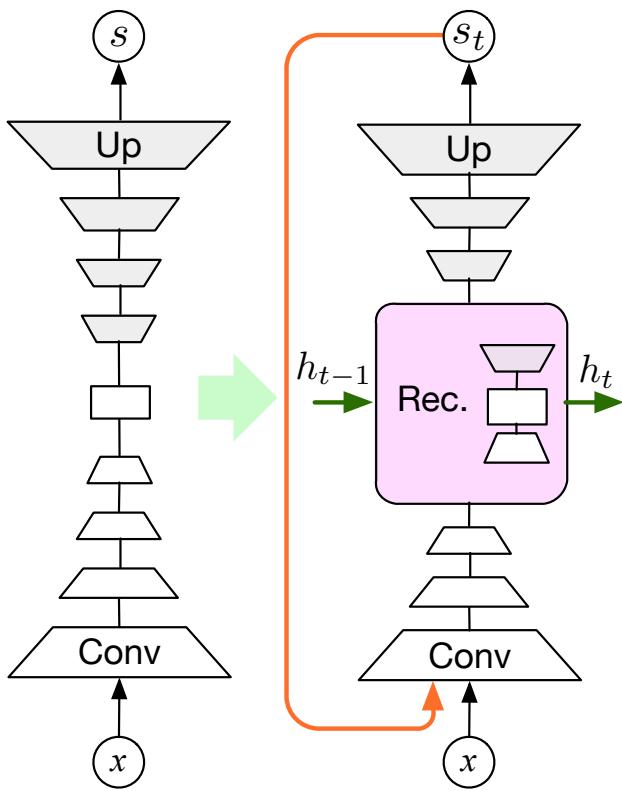
- A key role of the ConvNet is to generate for every output pixel a feature vector containing the output of all the intermediate layers.
- A classifier then assigns a label to the individual pixels.



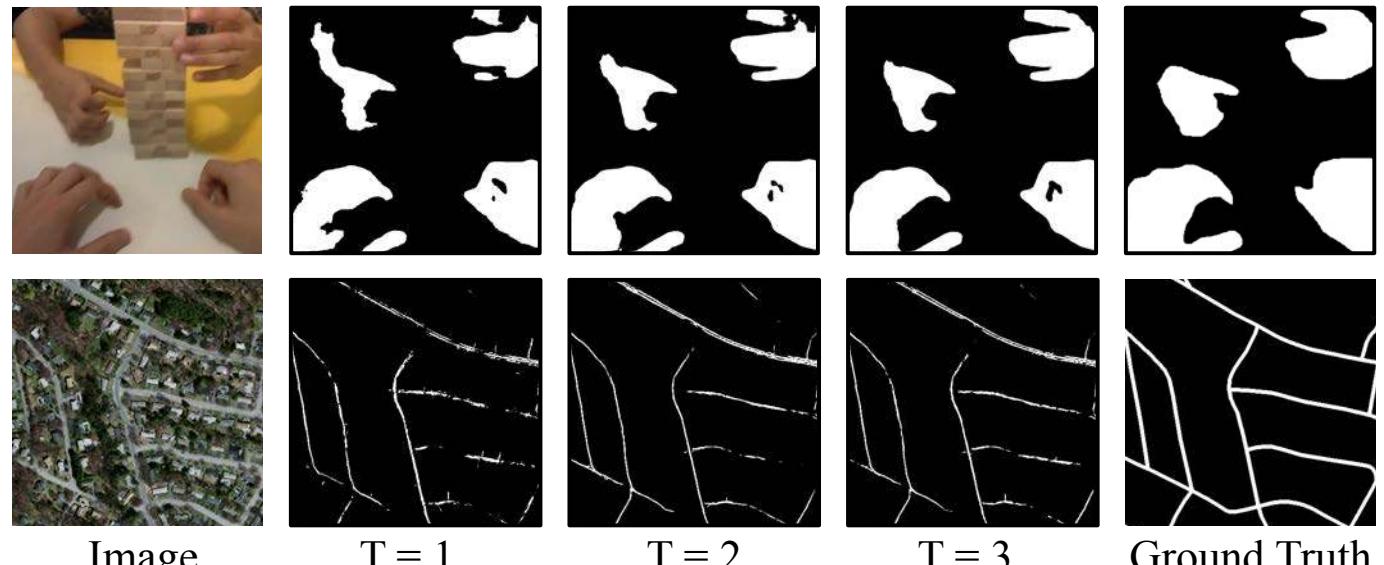
In the Street



Recursive Segmentation



feed recursively to the network
image and prev result

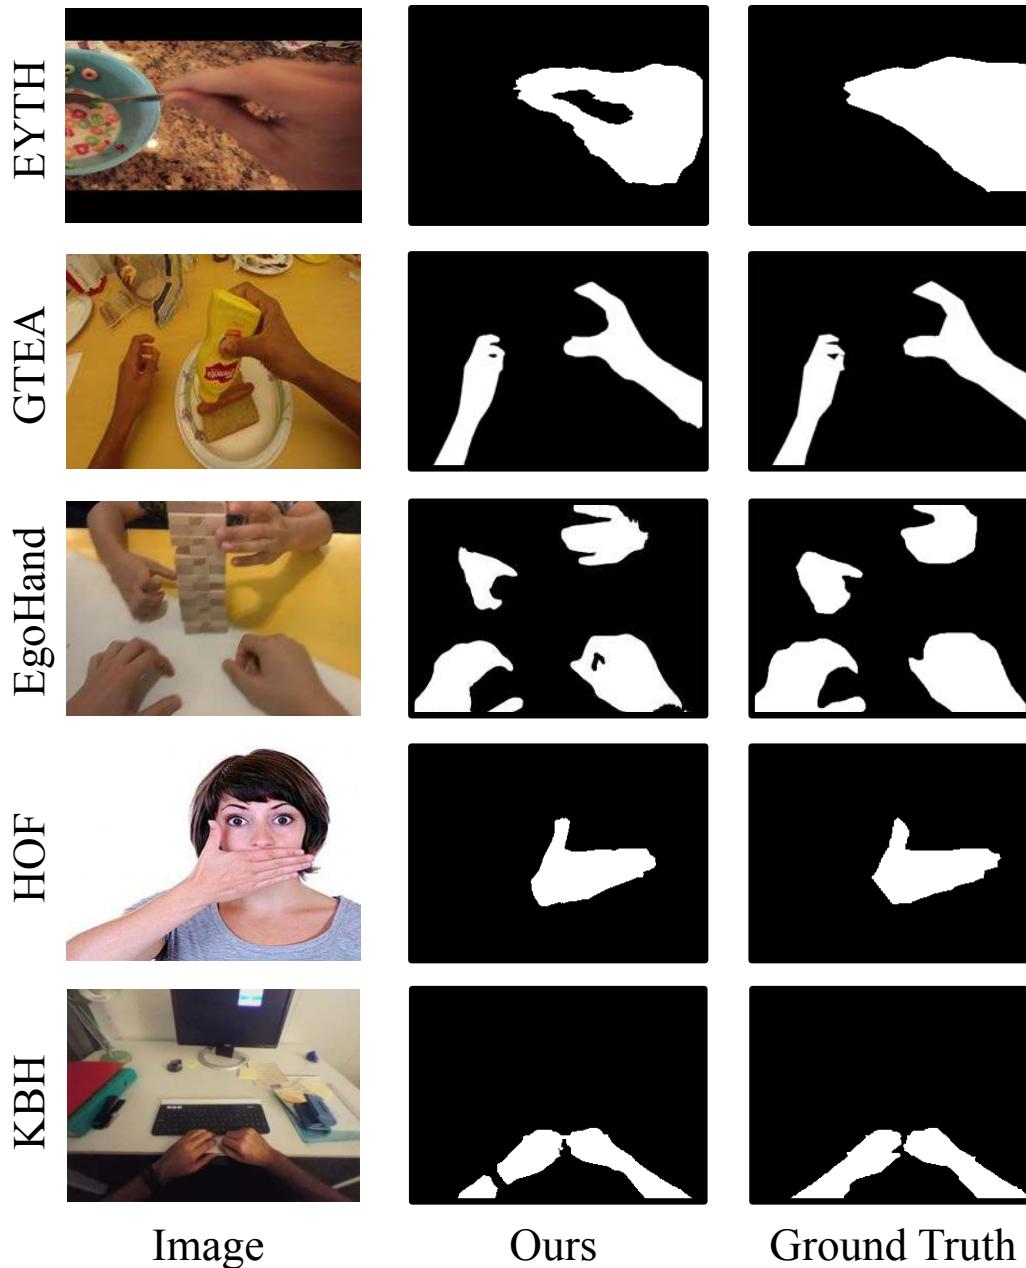


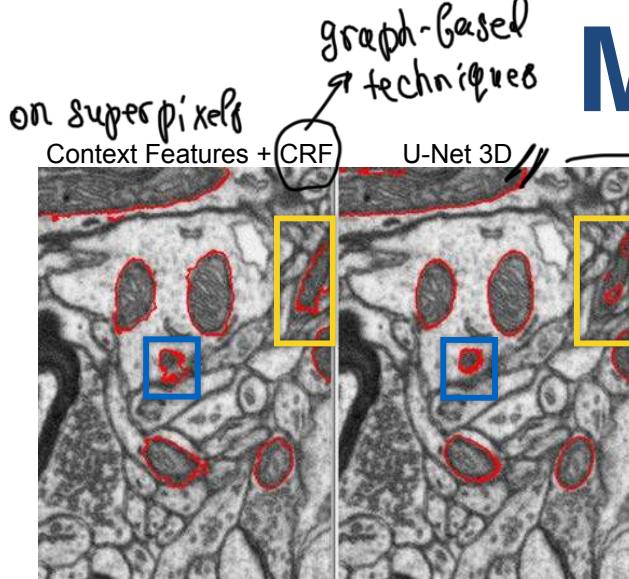
Hand/Road segmentation

U-Net Rec. U-Net

As for delineation, feeding the output back into the network is an effective way to take context into account.

Recursive Hand Segmentation





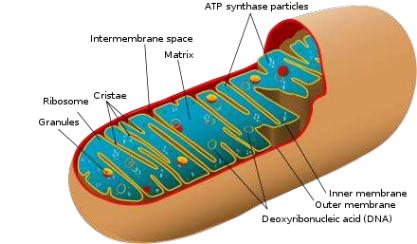
graph-based
techniques

Mitochondria

stacks of img's

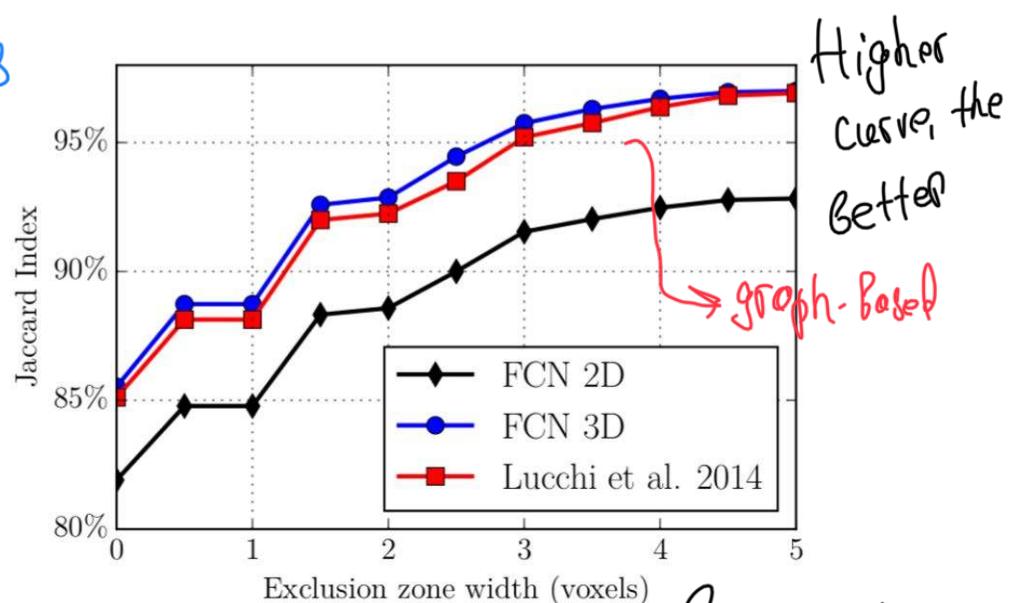
old algorithm
works better
(finds stuff that's
super small)

U-Net
does better
to find boundaries



Striatum Mitochondria

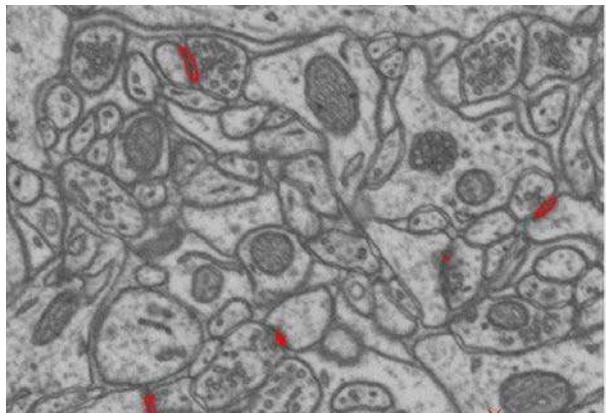
Method	Jaccard Index
Context F. + CRF	84.6%
U-Net 2D	82.4%
U-Net 3D volume	86.1%



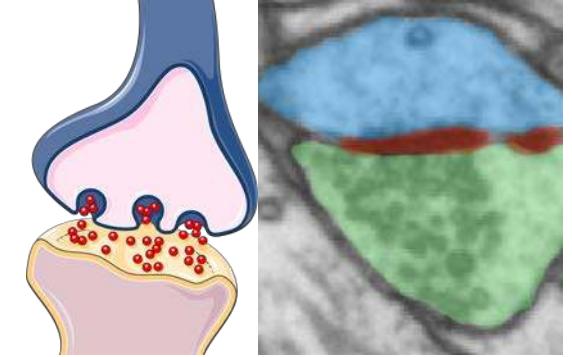
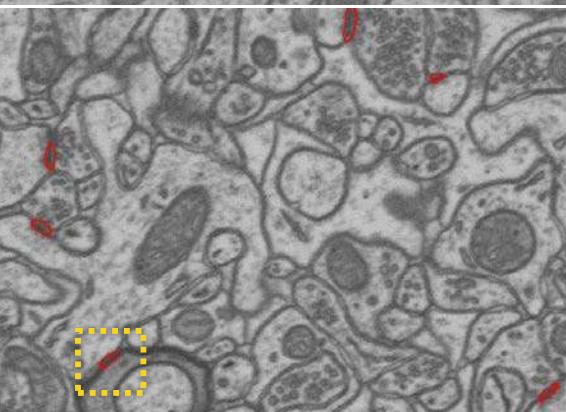
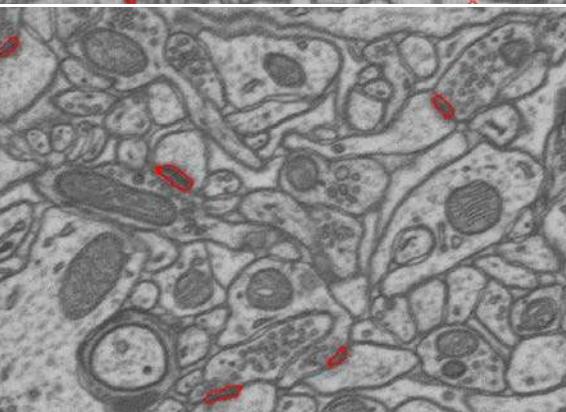
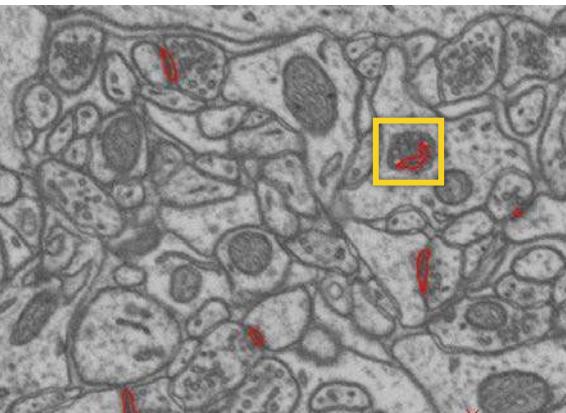
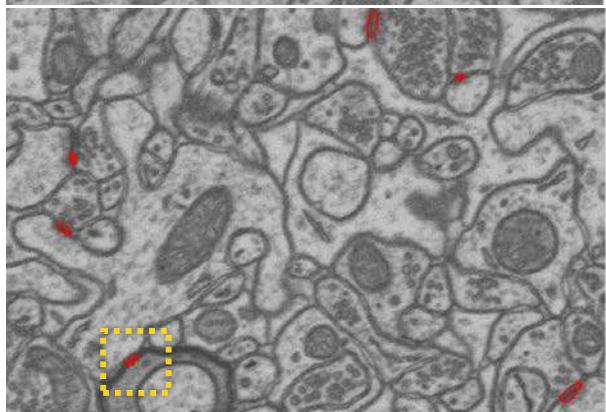
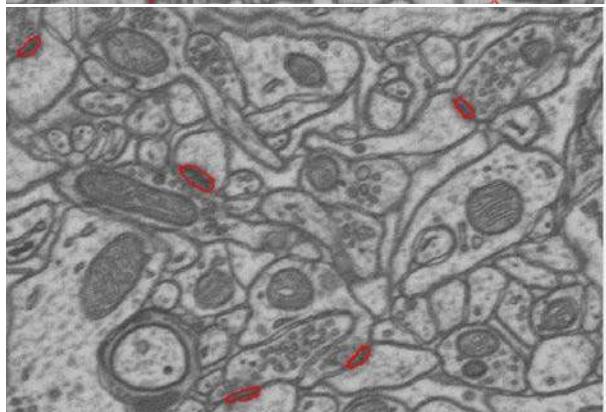
graph-based
gives you spatial information

Synapses

Context Features 3D CRF



U-Net 3D



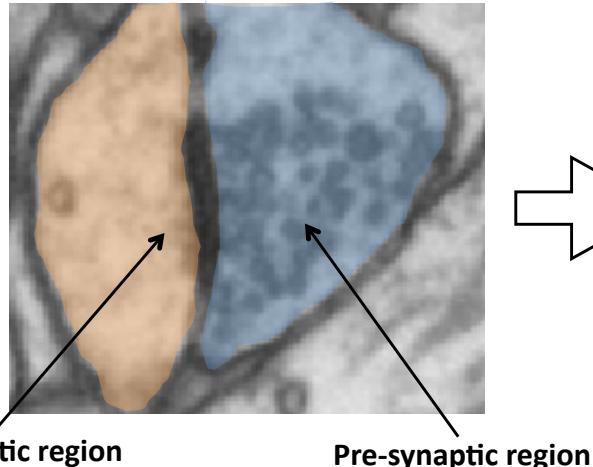
Jaccard Index	Method
66.8%	Context Features 2D
85.2%	Context Features 3D
73.5%	U-Net 2D
77.0%	U-Net 3D

on this specific example

?

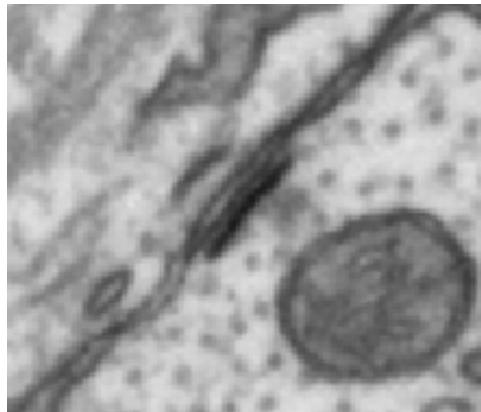
Context-Based Features

Synapse:

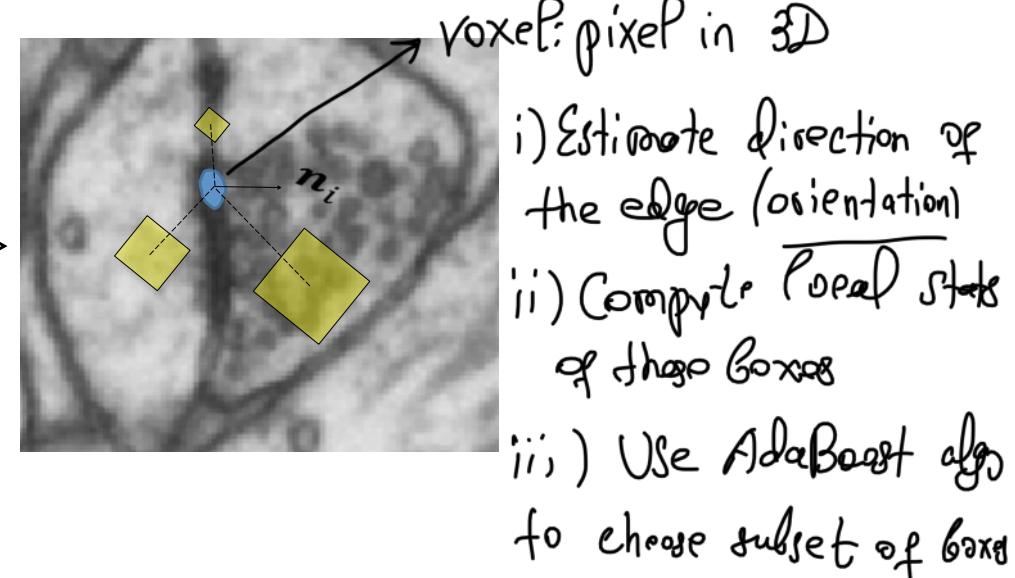


- Dendrite
- No vesicles
- Axon terminal
- Many vesicles

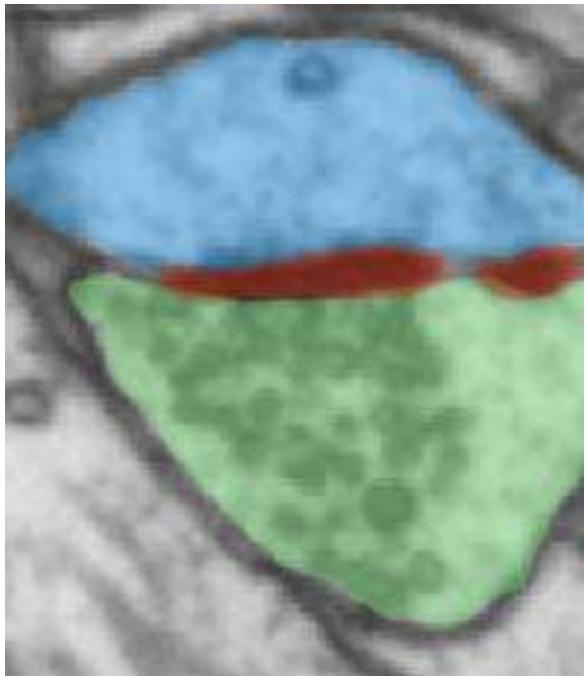
Non-Synapse:



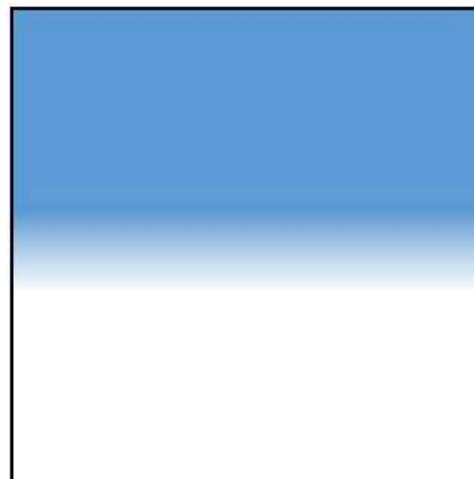
- The old style context-based features explicitly model orientation.
- How can we do that in a deep-learning framework?



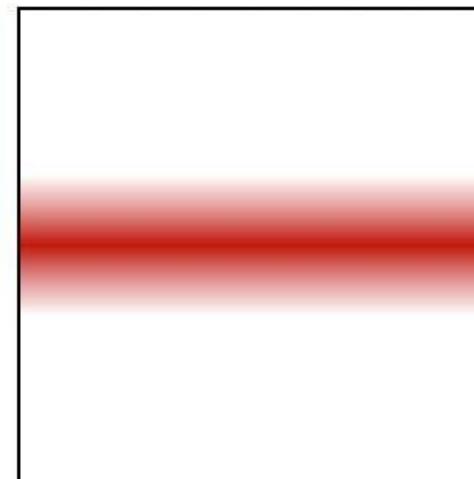
Probabilistic Atlases



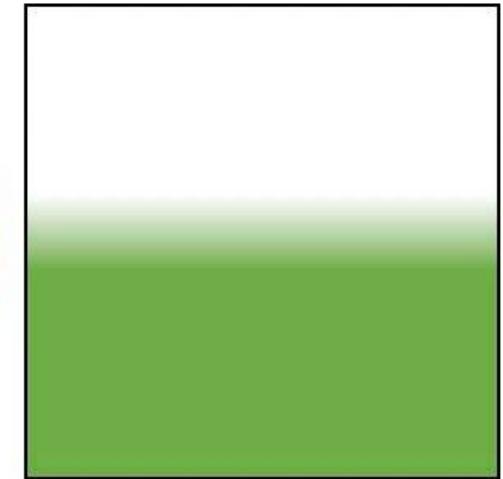
Synapse in canonical orientation



Probability of being
a post-synaptic voxel



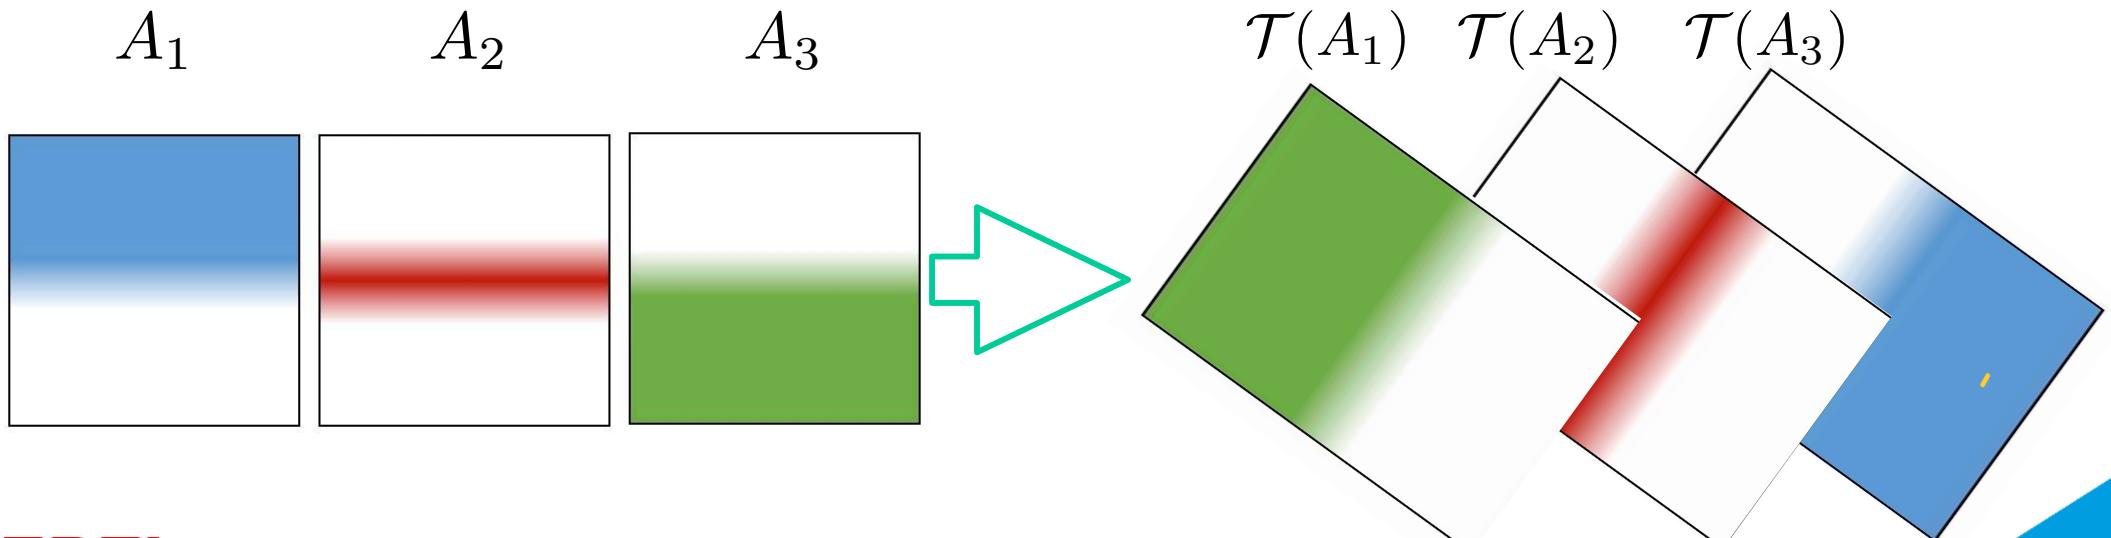
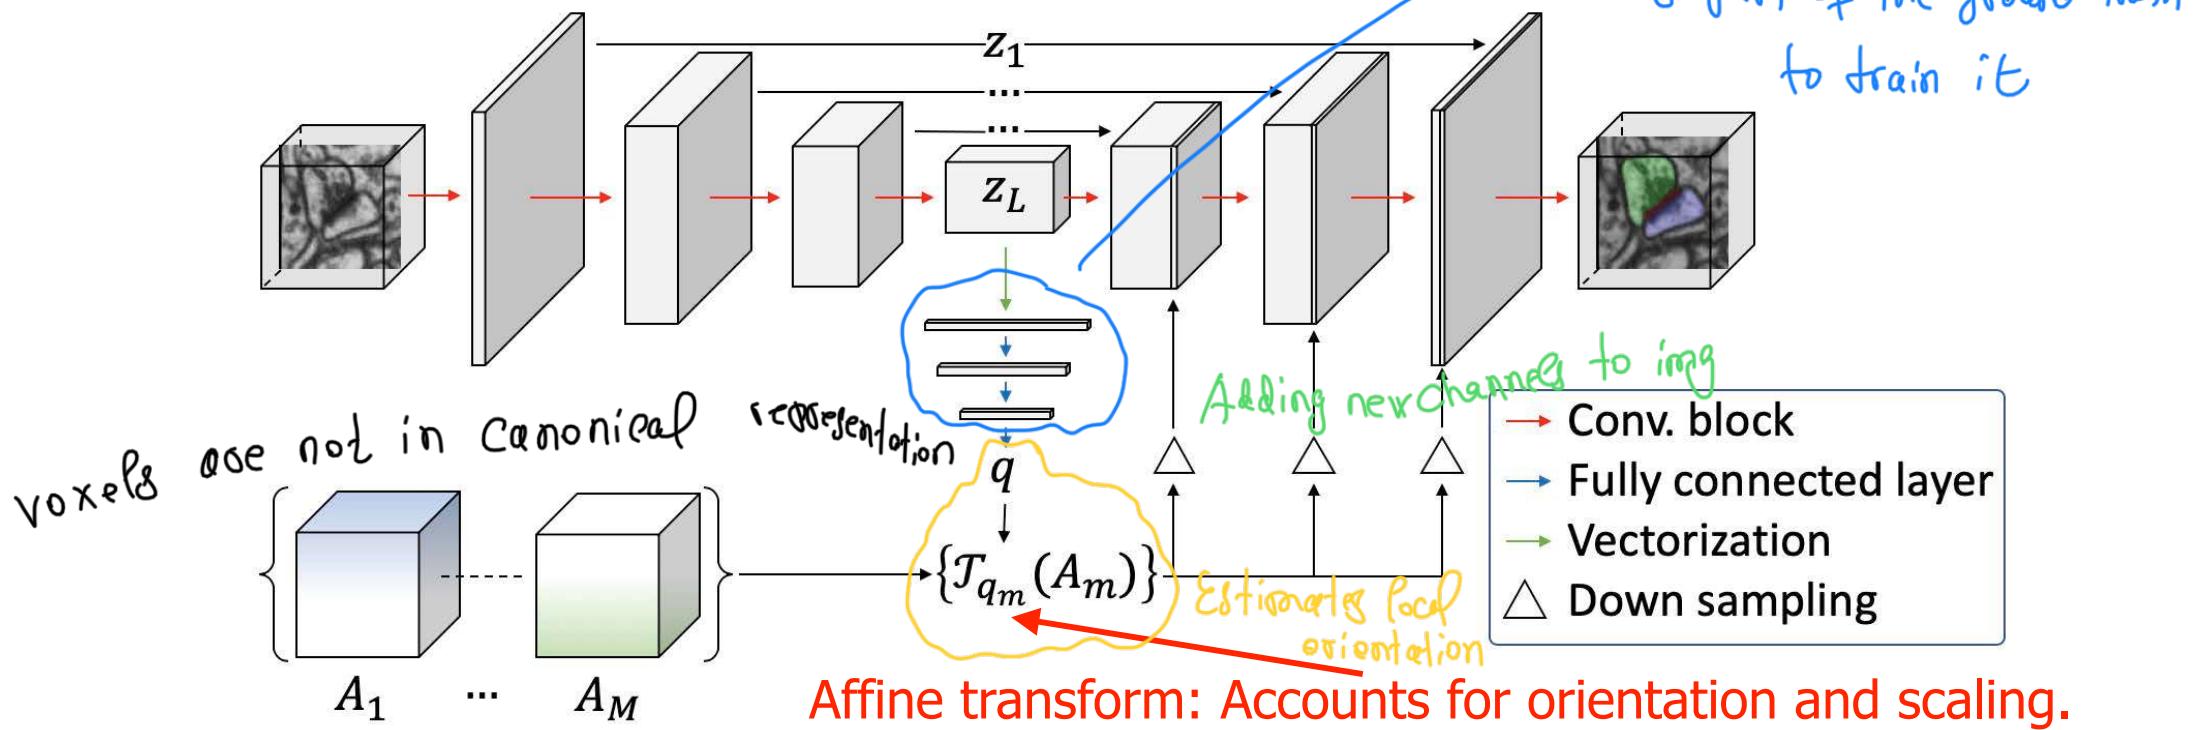
Probability of being
a cleft voxel



Probability of being
a pre-synaptic voxel

P_i likely to be

U-Net + Atlases



3D Synapses

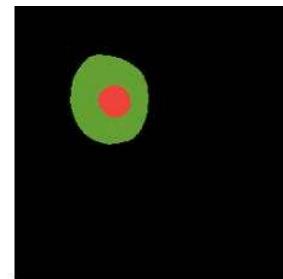
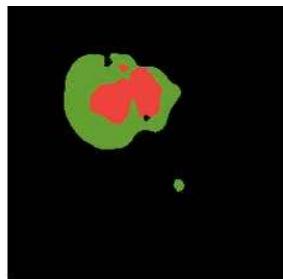
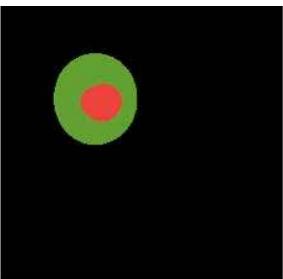
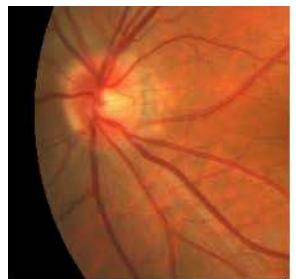
Synaptic Junction Segmentation

Baseline vs PA-Net

Probabilistic - Alphas

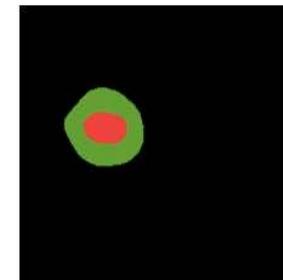
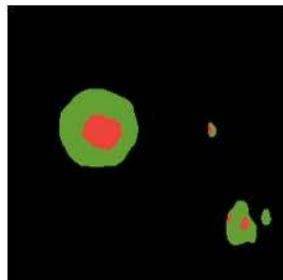
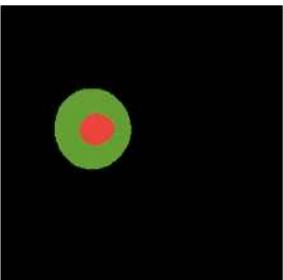
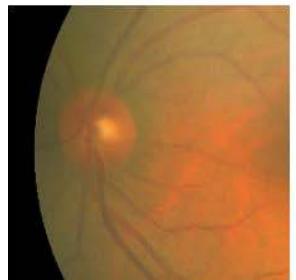
U-Net

Atlases for Optic Disk Segmentation

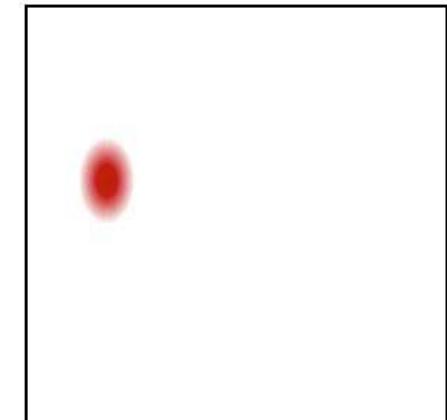
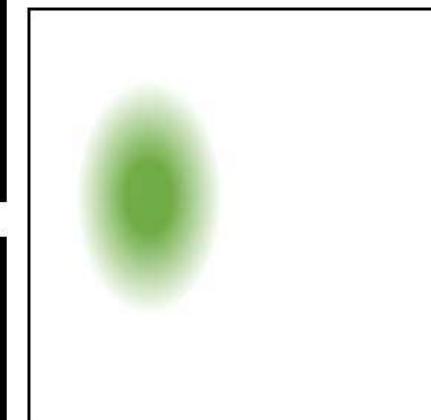


Retina of patients

optic nerve

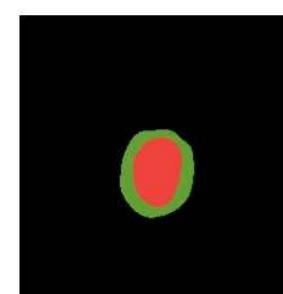
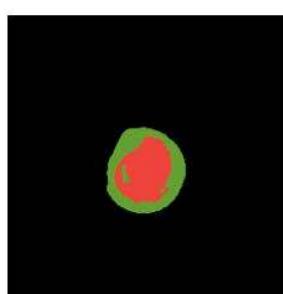
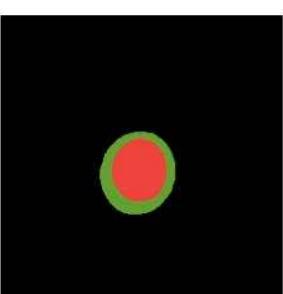


Probabilistic Atlases



Optic disk

Optic nerve



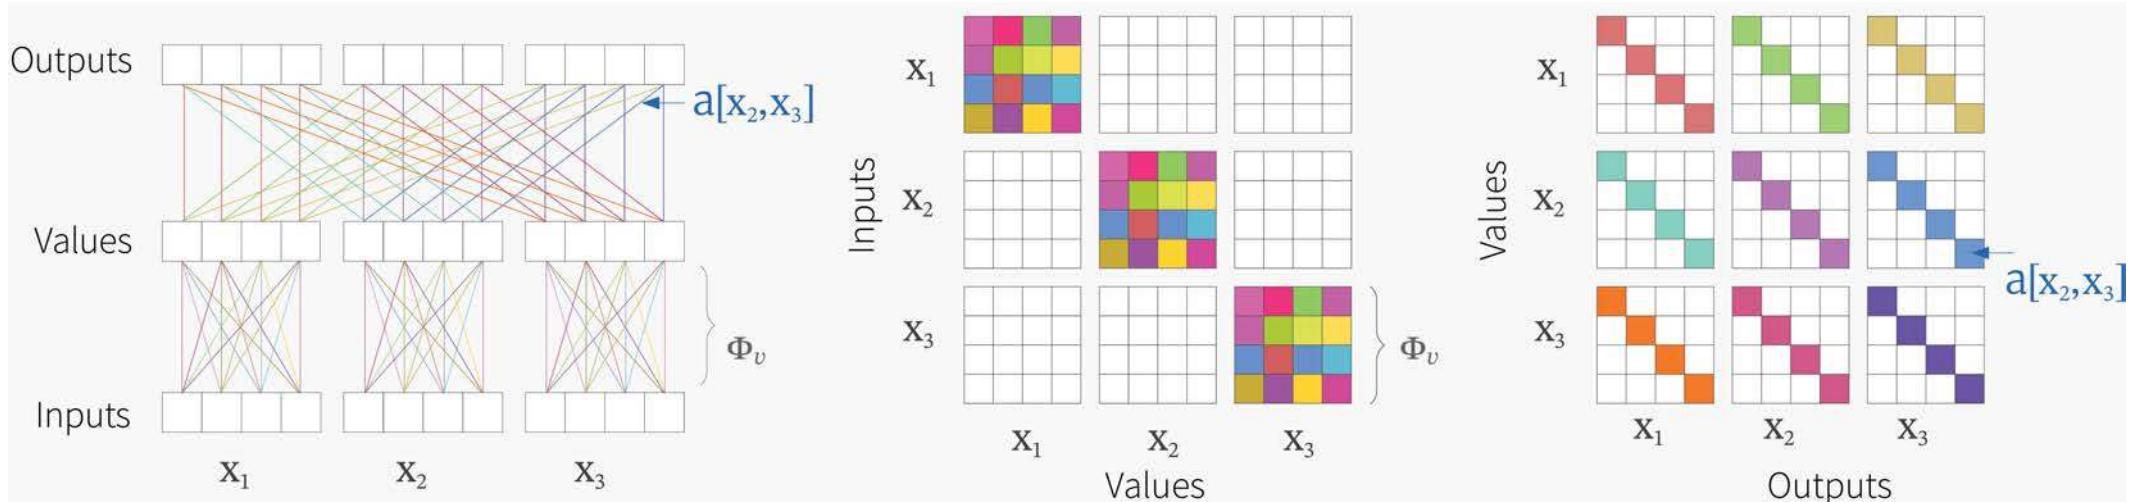
Ground truth

U-Net

PA-Net

fg, outer (inner) disk

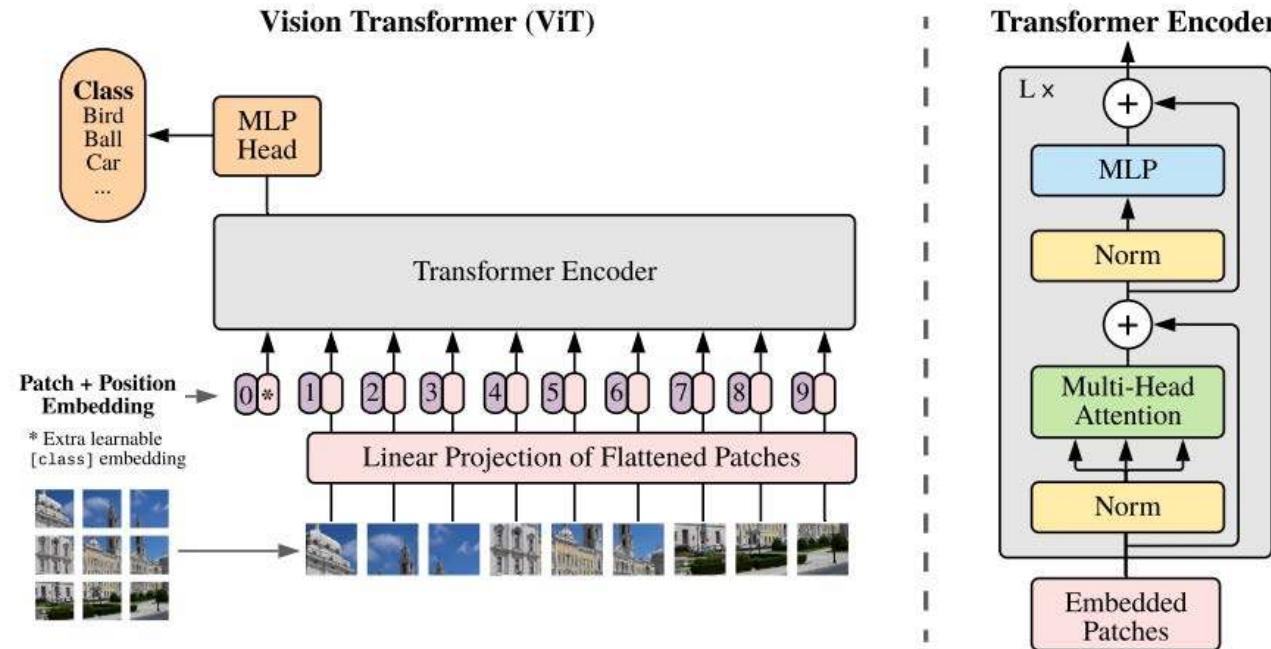
Optional: Transformers in NLP



Given $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_N]$,

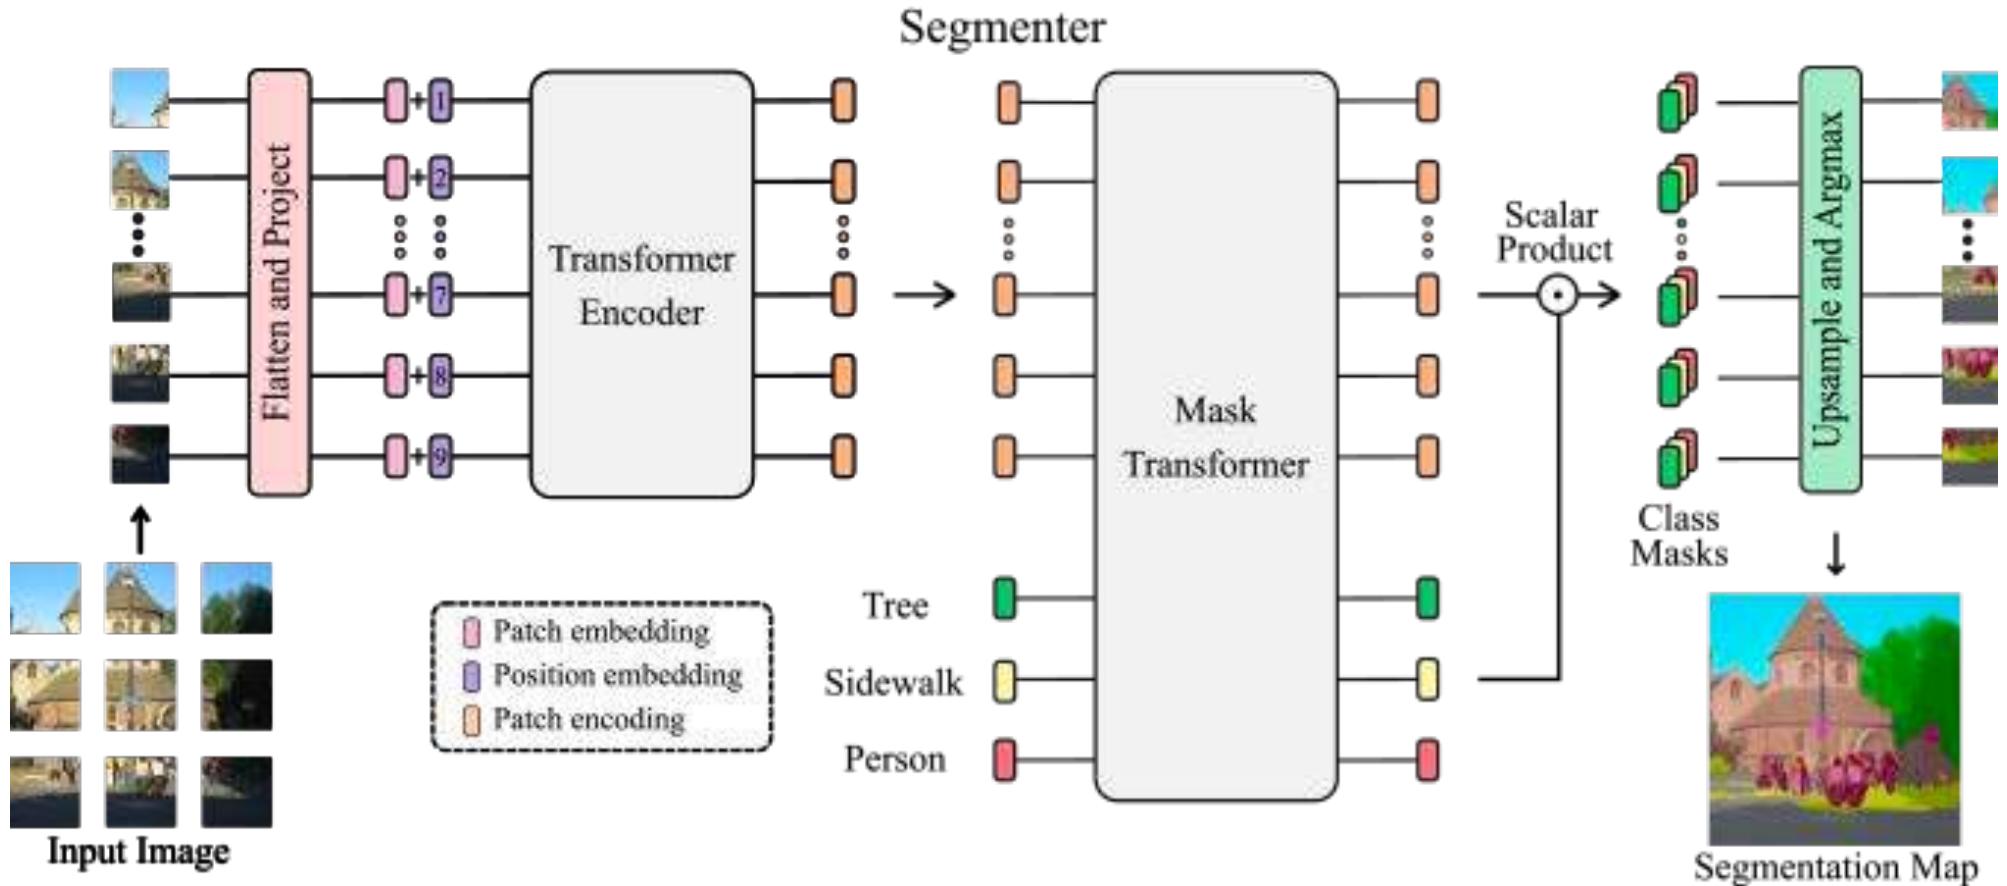
- compute SelfAttention(\mathbf{X}) using far fewer weights than if we used a fully connected network.
- In NLP the \mathbf{x}_n are embedding for words.
- In Vision they can be taken to be image patches.

Vision Transformers for Recognition



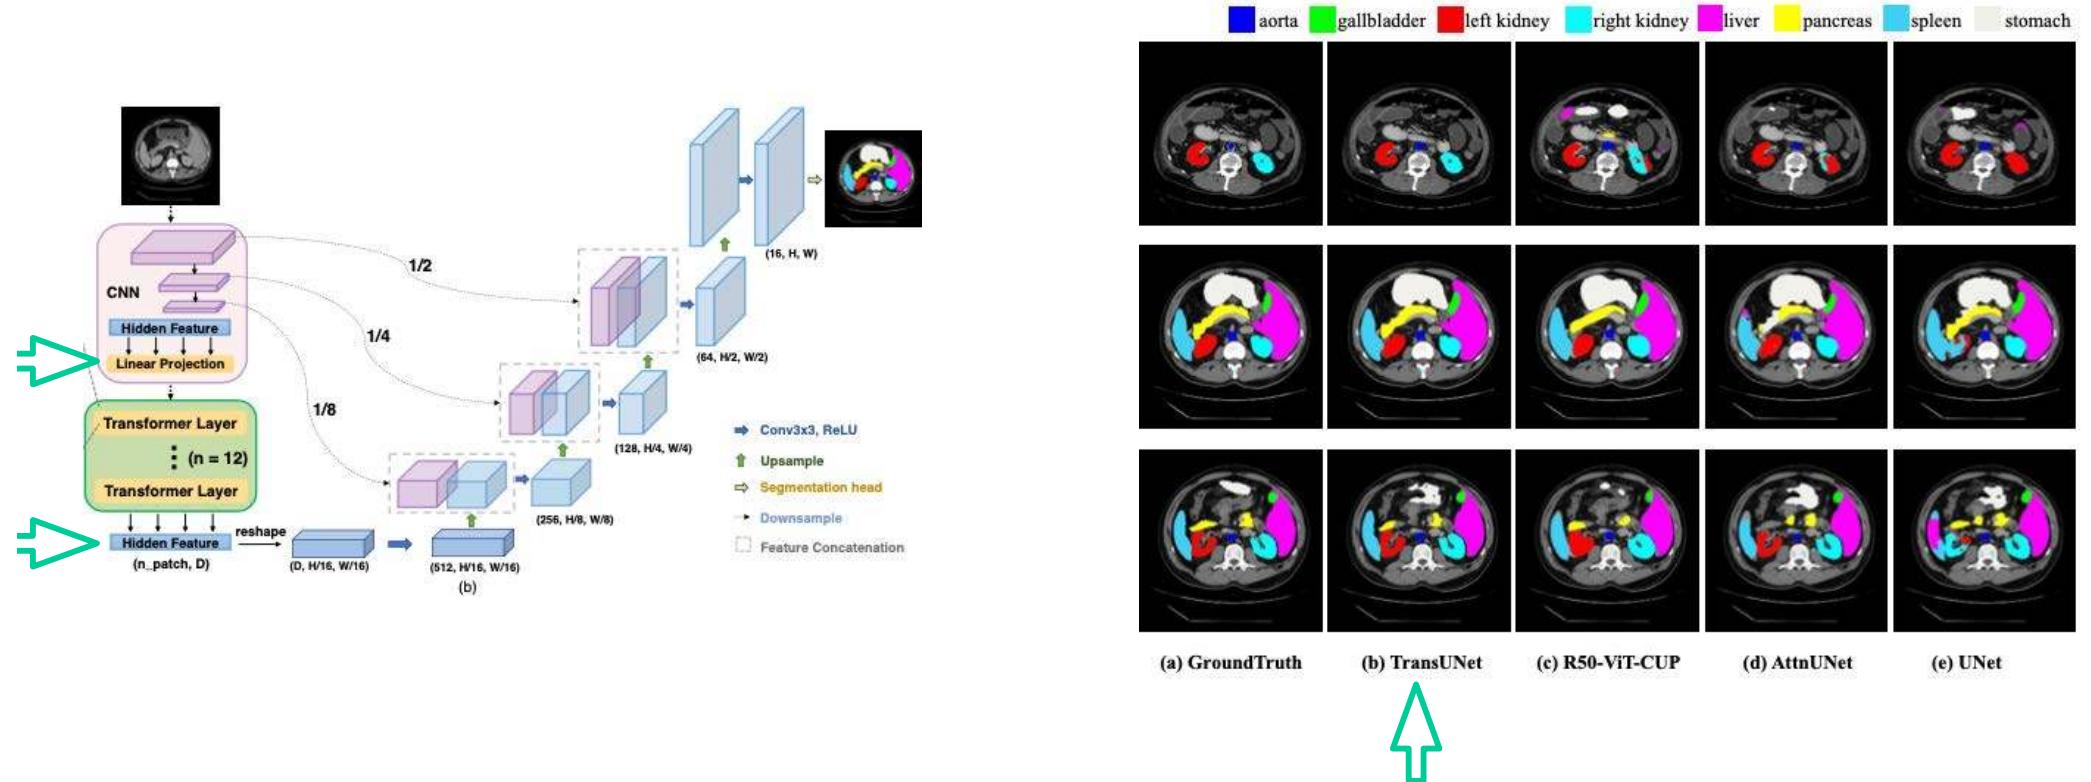
- Break up the images into square patches.
- Transform each path into a feature vector.
- Feed to a transformer architecture.

Vision Transformers for Segmentation



- Replace the “recognition” machinery by a “mask transformer”.
- Pros: Good at modeling long range relationships.
- Cons: Flattening the patches loses some amount of information.

U-NET + Transformers



- A CNN operates at low-resolution and produces a feature vector.
- A transform operates on that feature vector.
- The upsampling is similar to that of U-Net

—> Best of both worlds?

Moral of the Story

- Deep Networks are powerful tools, especially when there is enough training data. *but only for segmentation*
 - However, modeling your problem properly is still needed to achieve the highest possible level of performance. *(e.g. training network to predict orientation)*
 - The old techniques often inform our design choices.
- > Yet another reason why I am still talking about them.

In Short

- Local methods can provide valuable information but are inherently limited.

You need context!

- Domain knowledge, user interaction, and training data can be used to turn this data into usable results:

- Given enough training data, deep nets deliver the best performance today.
if you have limited data
- It can be further enhanced by introducing domain knowledge.
- Given smaller amounts of training data, K-Means and graphical models still have their uses.
old-fashioned techniques
- Same philosophy as for delineation.

What About the Dog?

