

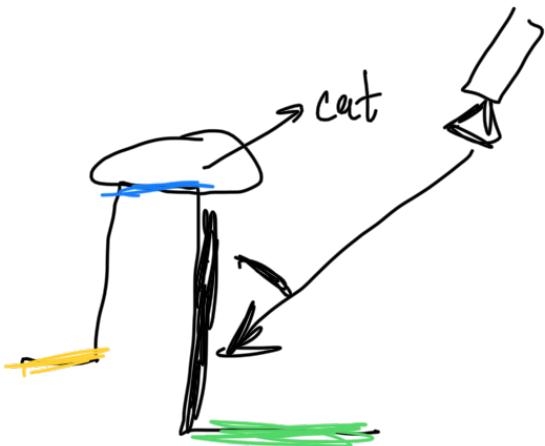
# Shape from X

Capturing shapes from shading

constant albedo, uniform image  
look in

- One image:
  - Shading
  - **Texture**
- Two images or more:
  - Stereo
  - Contours
  - Motion

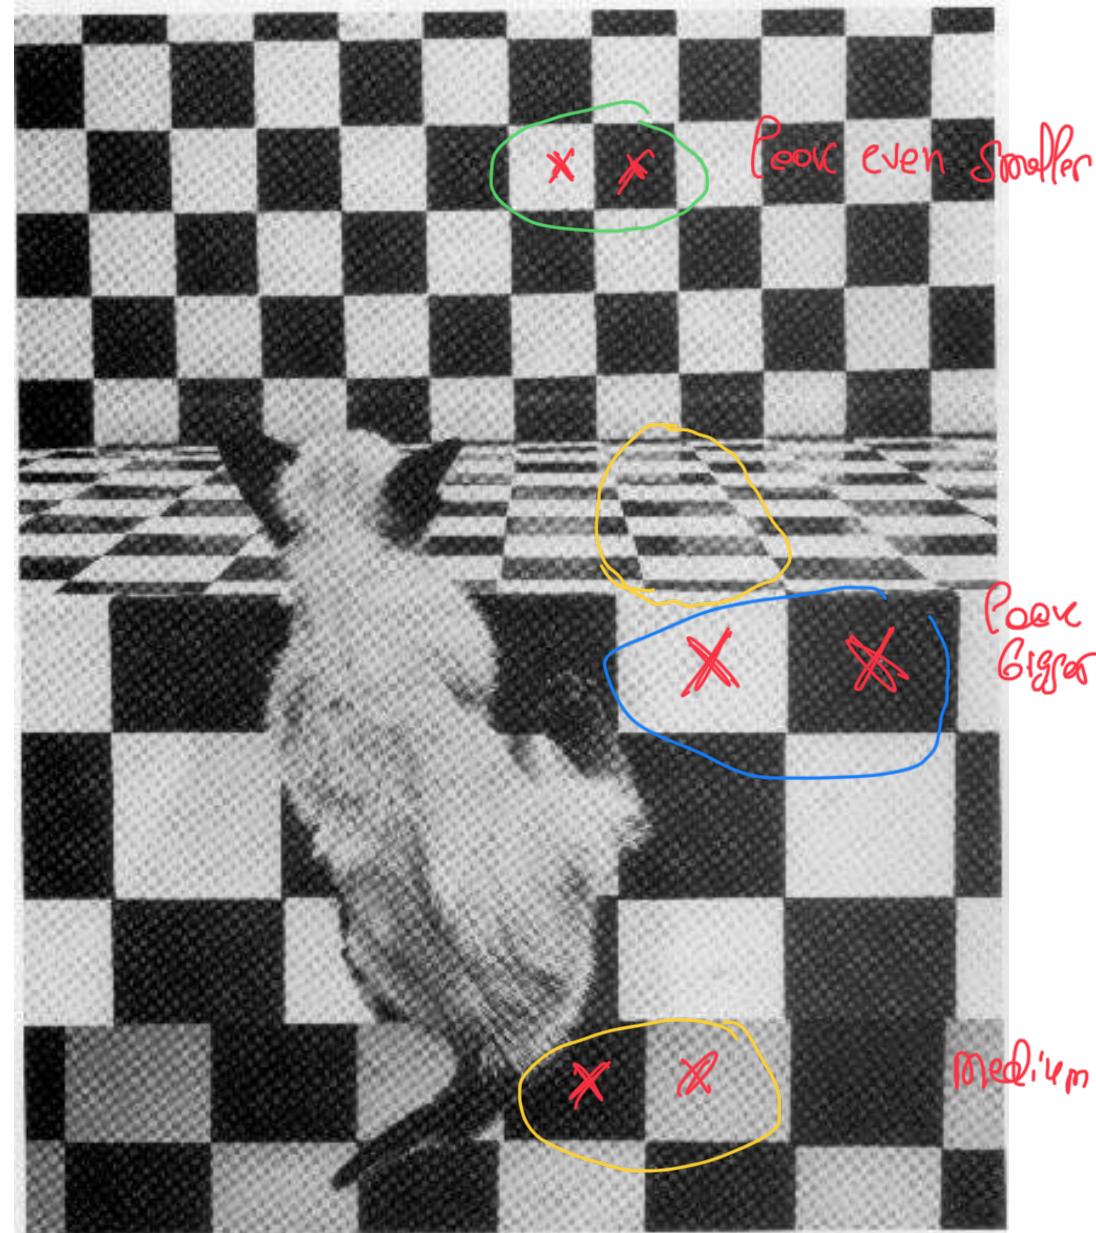
# Shape From X



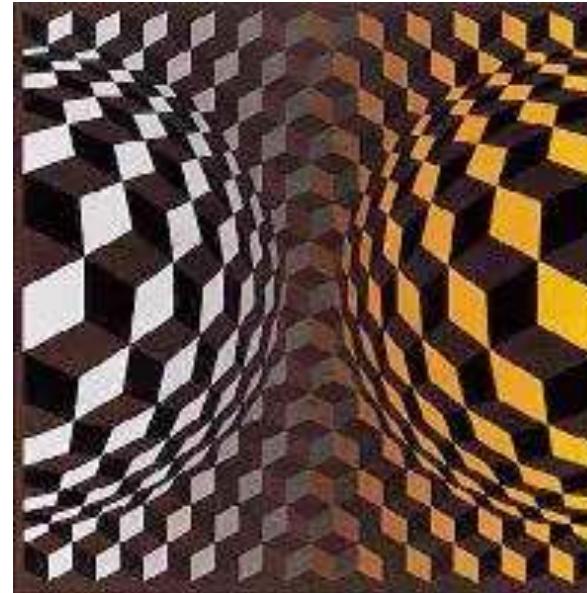
- One image:
  - Shading
  - **Texture**
- Two images or more:
  - Stereo
  - Contours
  - Motion

In perspective, things that are further look smaller

They are not squares in proj  
⇒ distorted,  
foreshortened



# Shape From Texture



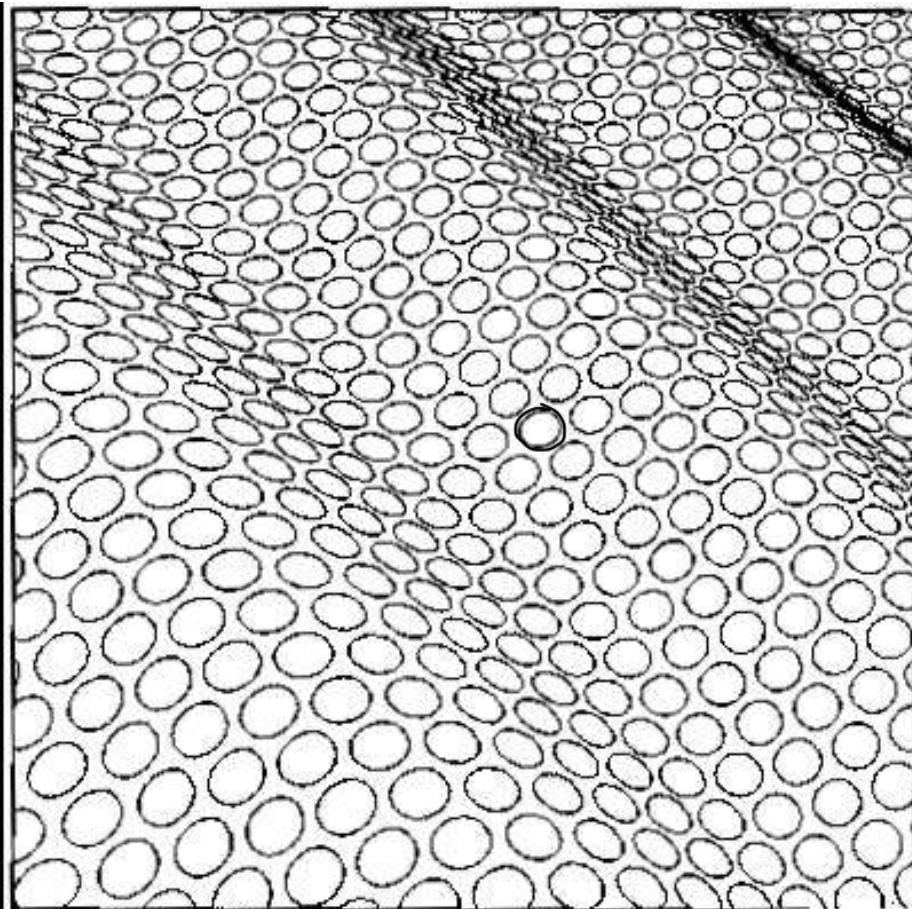
Recover surface orientation or surface shape from image texture:

- Assume texture 'looks the same' at different points on the surface.
- This means that the deformation of the texture is due to the surface curvature.

distance also

# Structural Shape Recovery

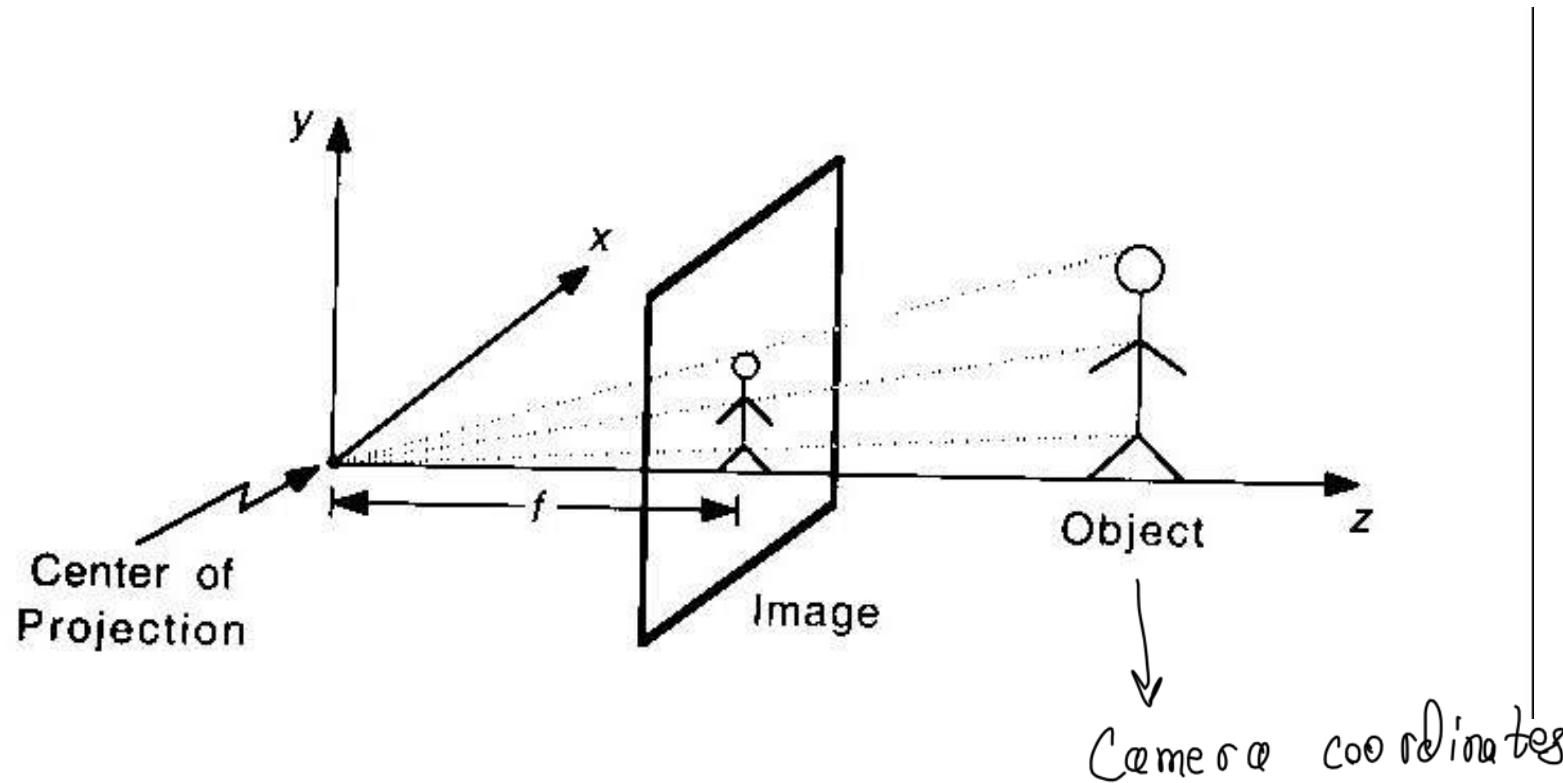
projection deforms them → How they are deformed  
tells us sth about their shapes  
projected circles ✓



**Basic hypothesis:** Texture resides on the surface and has no thickness.

- Computation under:
- Perspective projection
  - Paraperspective projection
  - Orthographic projection

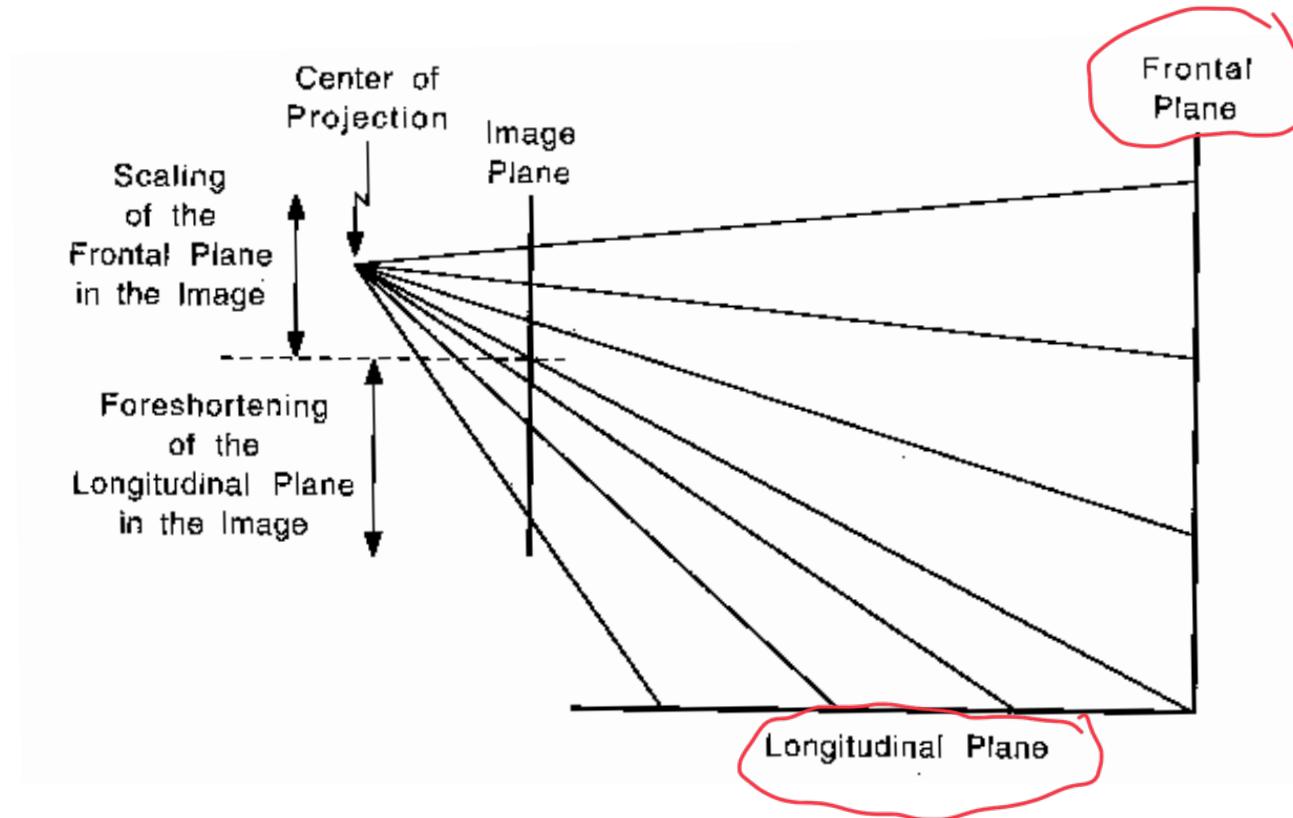
# Reminder: Perspective Projection



$$u = f \frac{x}{z}$$

$$v = f \frac{y}{z}$$

# Perspective Distortion



The perspective projection distortion of the texture

- depends on both depth and surface orientation,  
    tangent surface and position of the camera
- is anisotropic.

# Foreshortening

Distance from Camera → point in the scene

angle at which surface is oriented relative to camera

Depth vs Orientation:

- Infinitesimal vector  $[\Delta x, \Delta y, \Delta z]$  at location  $[x, y, z]$   
image of this vector is (according to dynamics eq'n)

*occurs in depth changes do not*

$$\frac{f}{z} [\Delta x - \frac{x}{z} \Delta z, \Delta y - \frac{y}{z} \Delta z]$$

projection of  $\Delta x, \Delta y, \Delta z$

- Two special cases:

- $\Delta z = 0$  :

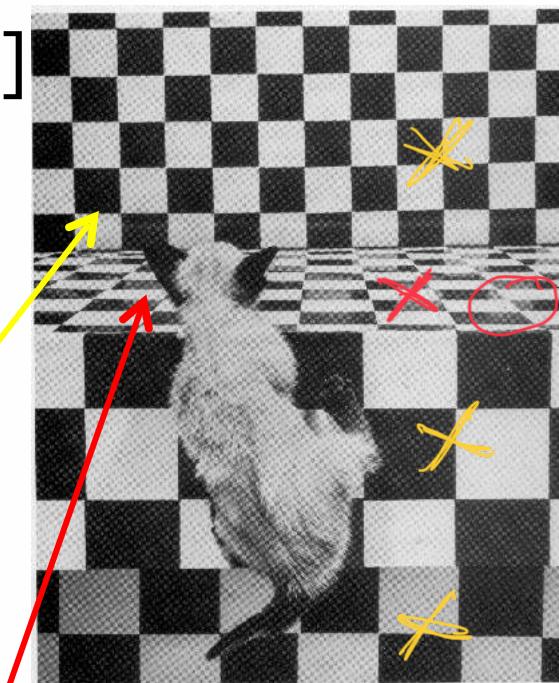
The object is scaled

- $\Delta x = \Delta y = 0$  :

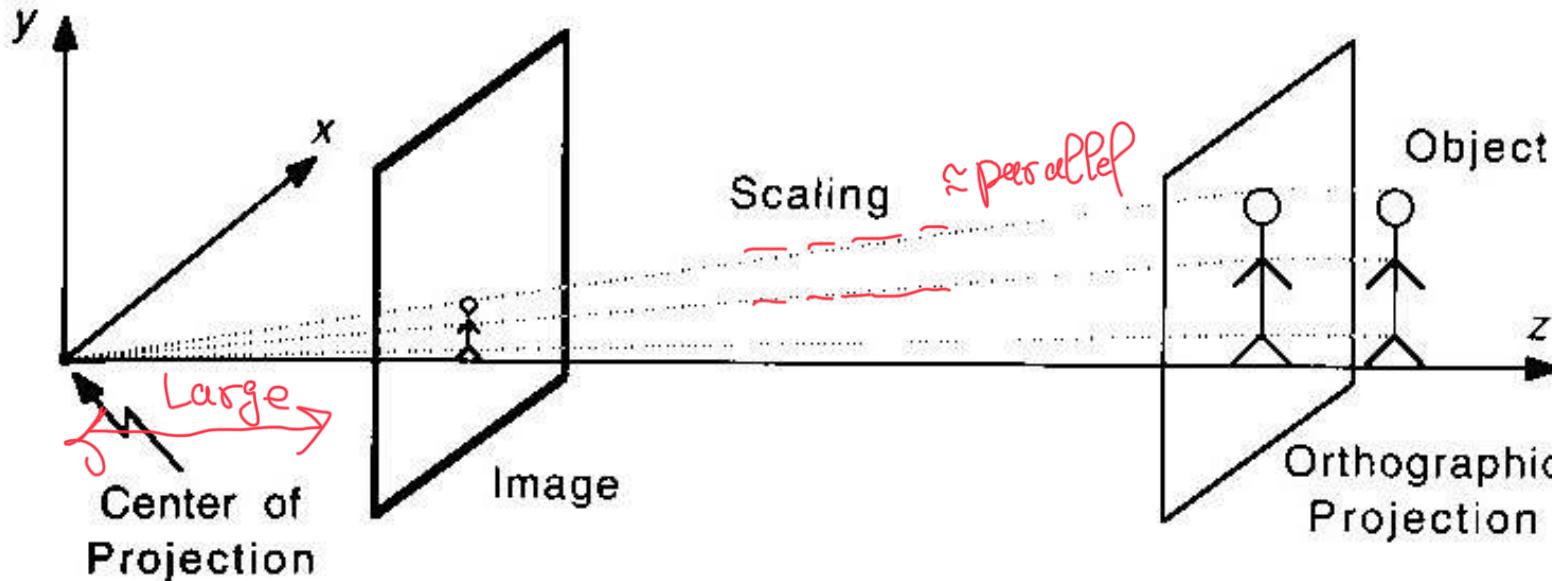
(division by  $z^2$ )

The object is foreshortened

*longitudinal plane*



# Reminder: Orthographic Projection



→ Does not account for the fact that objects appear

Special case of perspective projection: smaller when they are further away

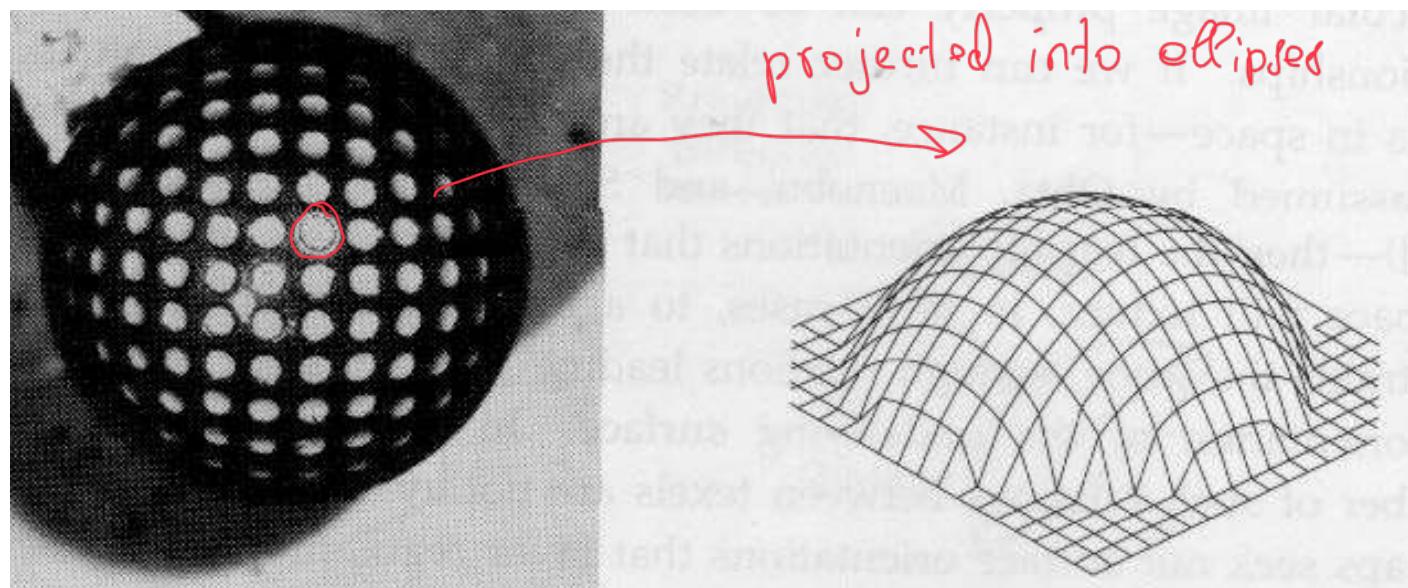
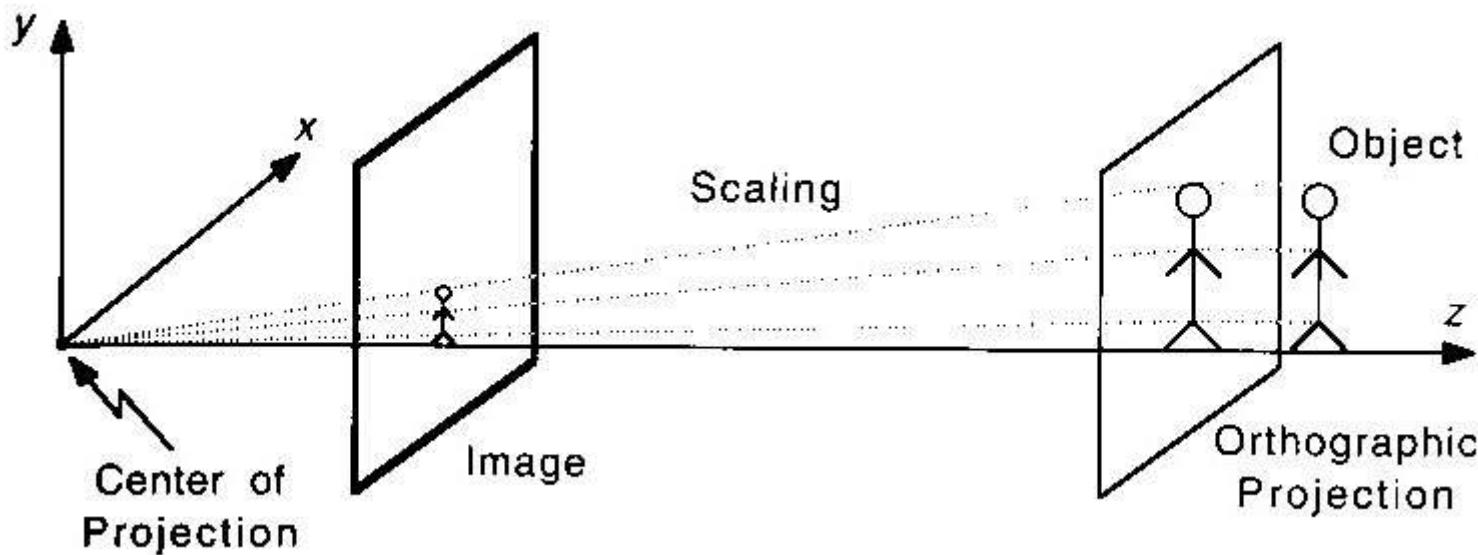
- Large f

- Objects close to the optical axis

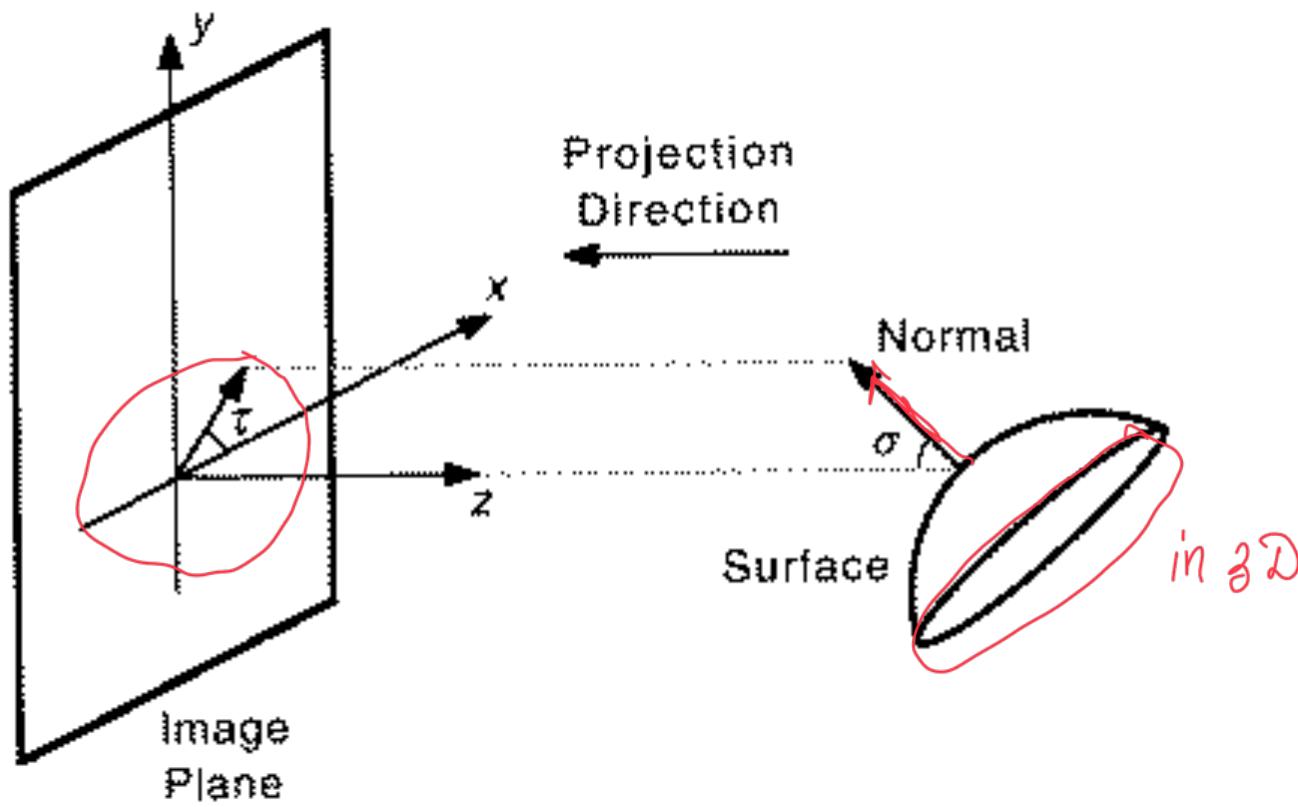
→ Parallel lines mapped into parallel lines.

No foreshortening (distortion that makes an object appear shorter than it's due to angled view)

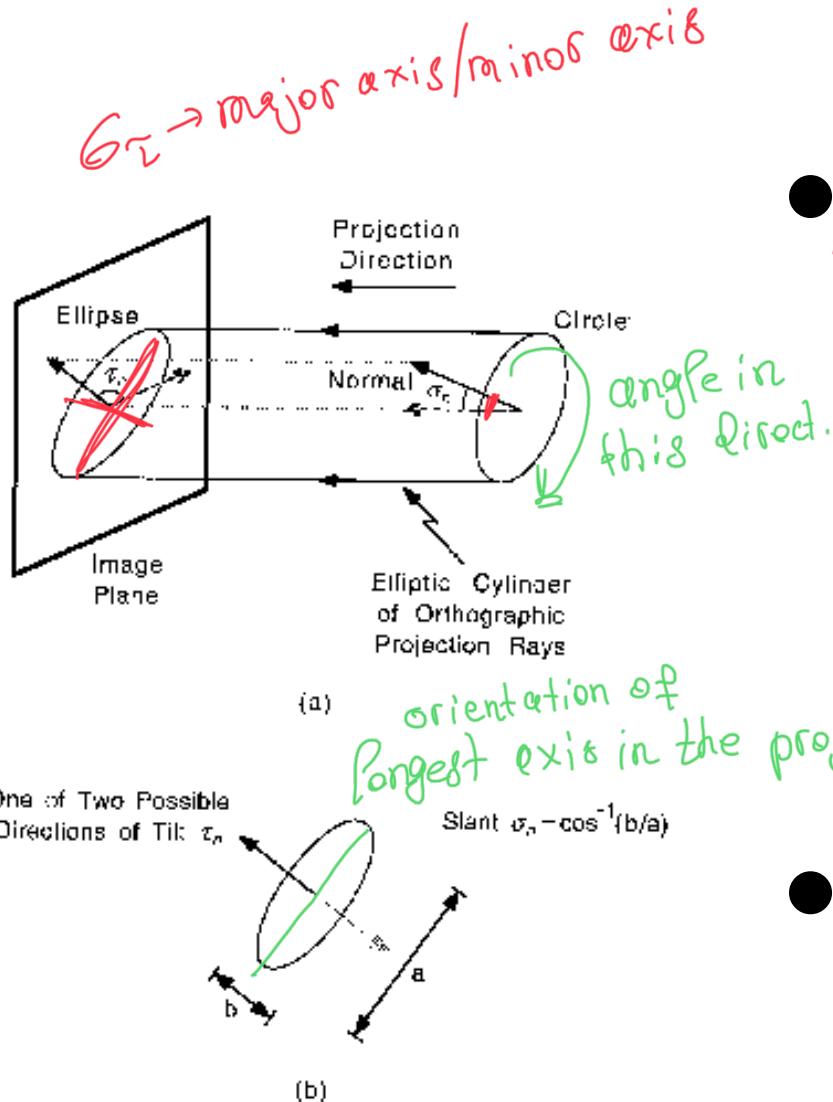
# Orthographic Projection



# Tilt And Slant



# Orthographic Projection



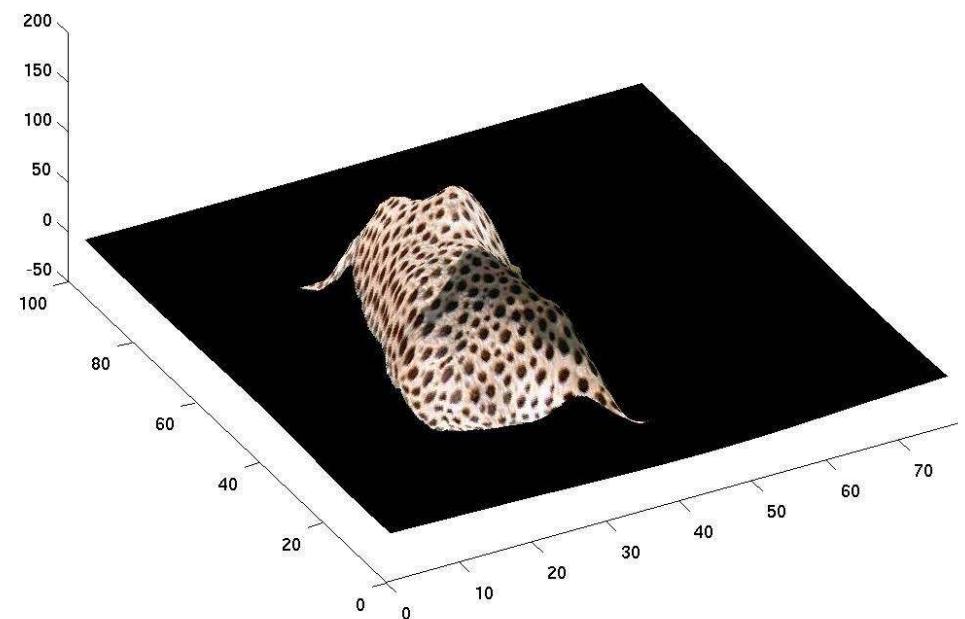
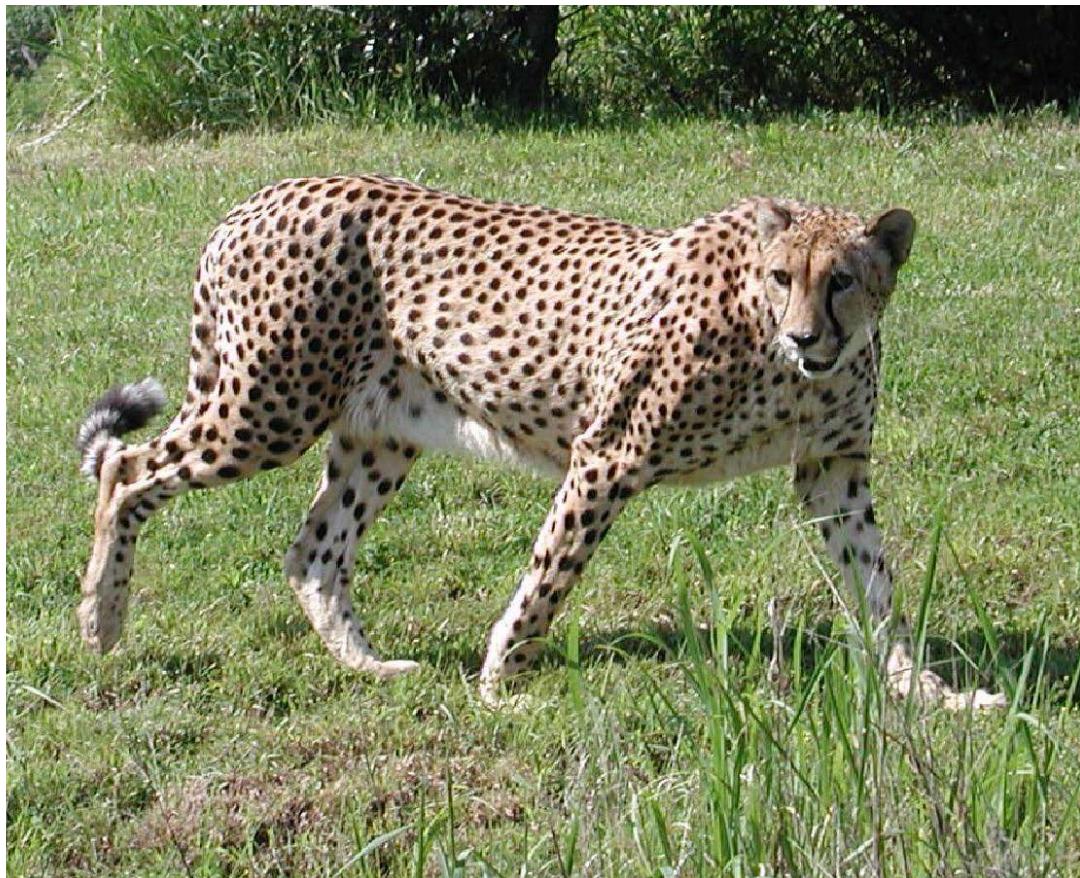
$\alpha=0 \Rightarrow$  projecting into circle

$\alpha \neq 0 \Rightarrow$  more elongated ellipse 'll be'

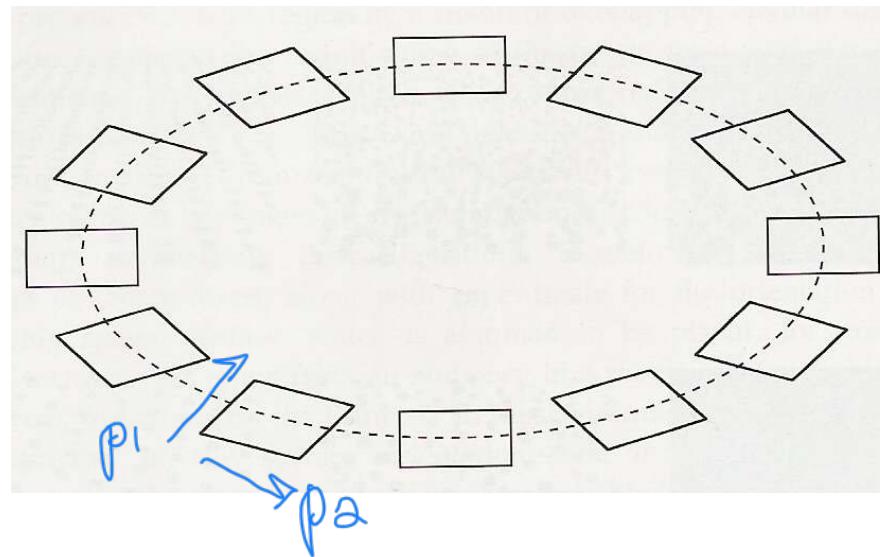
- **Tilt:** Derived from the image direction in which the surface element undergoes maximum compression.

- **Slant:** Derived from the extent of this compression.

# Cheetah



# Perpendicular Lines

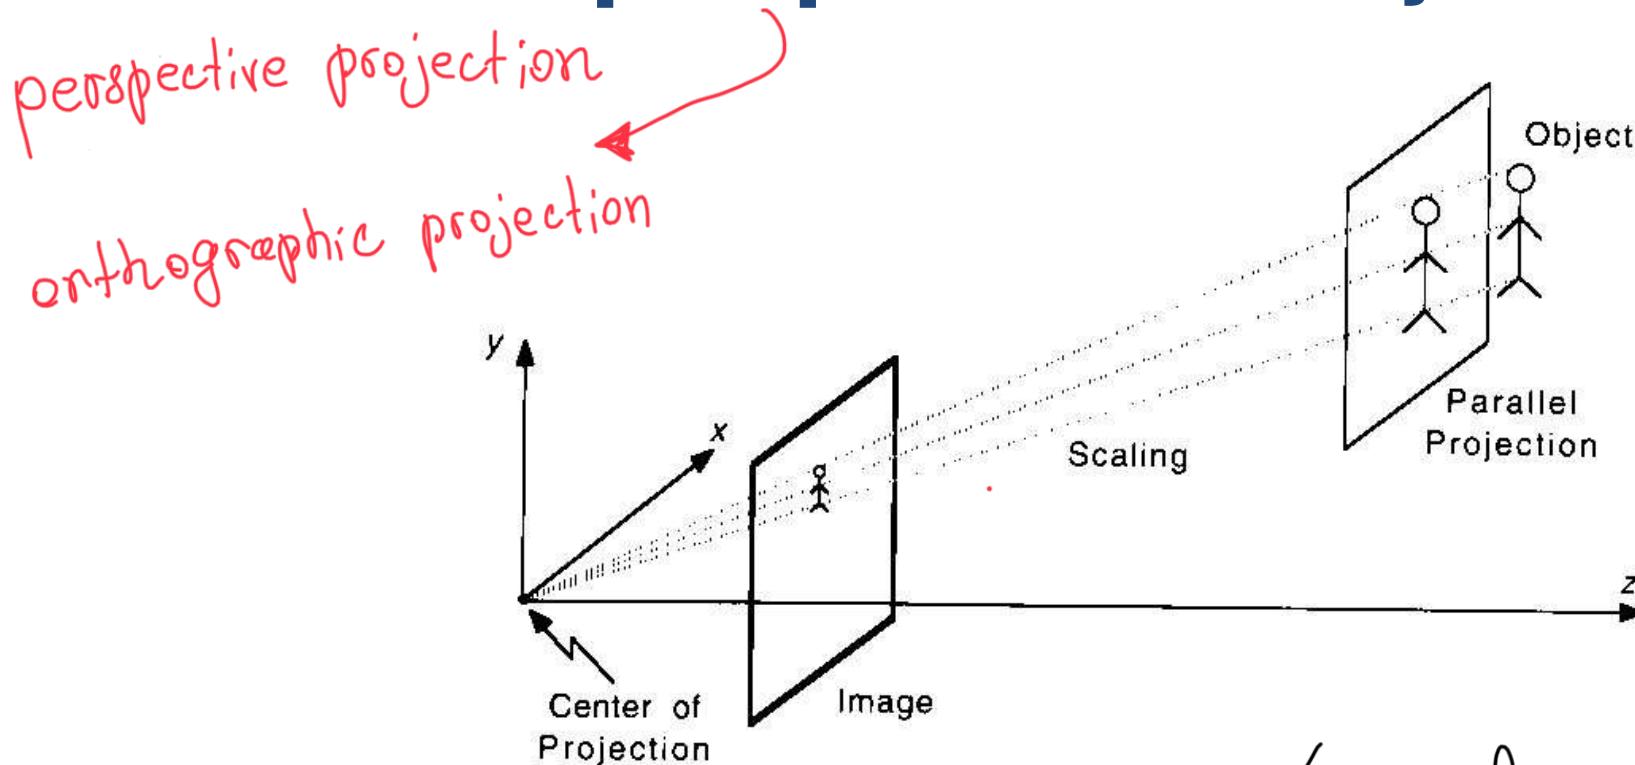


Orthographic projections of squares that are rotated with respect to each other in a plane inclined at  $\omega=60^\circ$  to the image plane.

$$\frac{\|\mathbf{p}_1/l_1 \times \mathbf{p}_2/l_2\|}{2} = \frac{\cos(\omega)}{1 + \cos^2(\omega)}$$

$$\|\mathbf{p}_1/\rho_1\|^2 + \|\mathbf{p}_2/\rho_2\|^2$$

# Paraperspective Projection



(was good for close to optical axis)

Generalization of the orthographic projection:

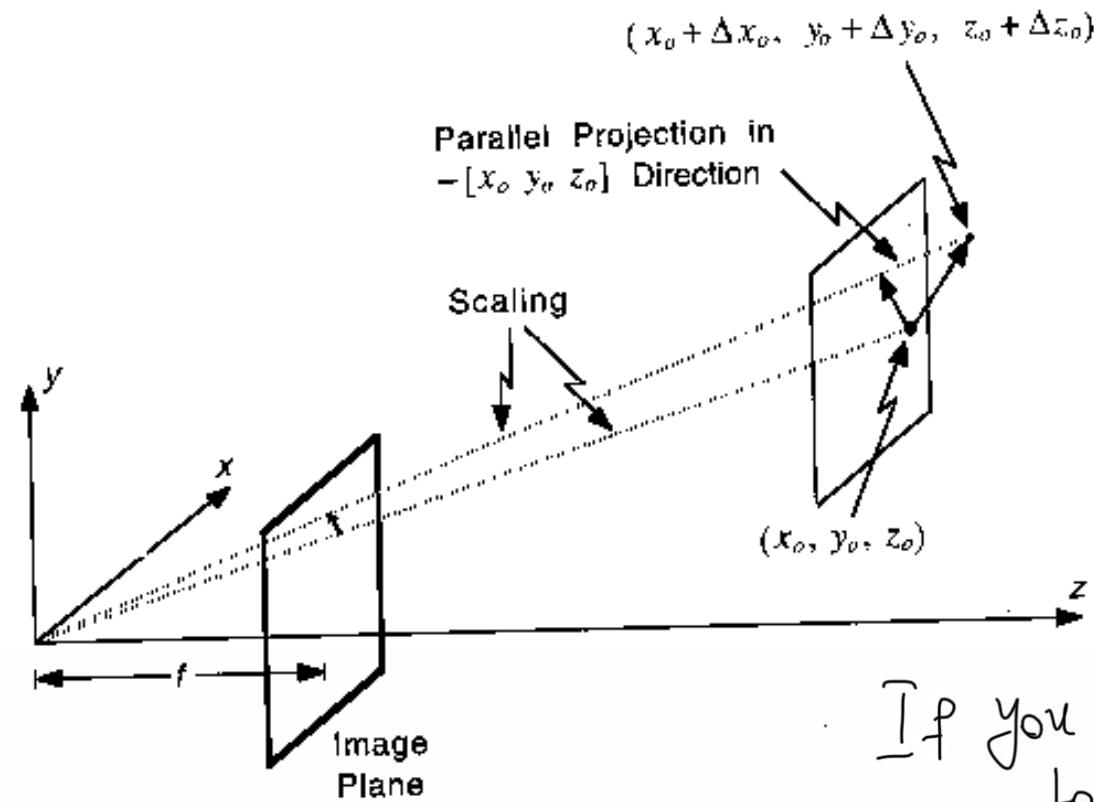
- Object dimensions small wrt distance to the center of projection.

→ Parallel projection followed by scaling

(good for cameras with required with long focal length)

(far away from optical axis)

# Paraperspective Projection



If you compute this for several texture elements

- For planar texels:

Unknown surface normal.

$$A' = -\frac{f^2}{z_0^3} \mathbf{n} \cdot [x_o \ y_o \ z_o] A$$

Projected Area. →  $A'$

True Area.

$\frac{1}{z_0^3}$  Depth

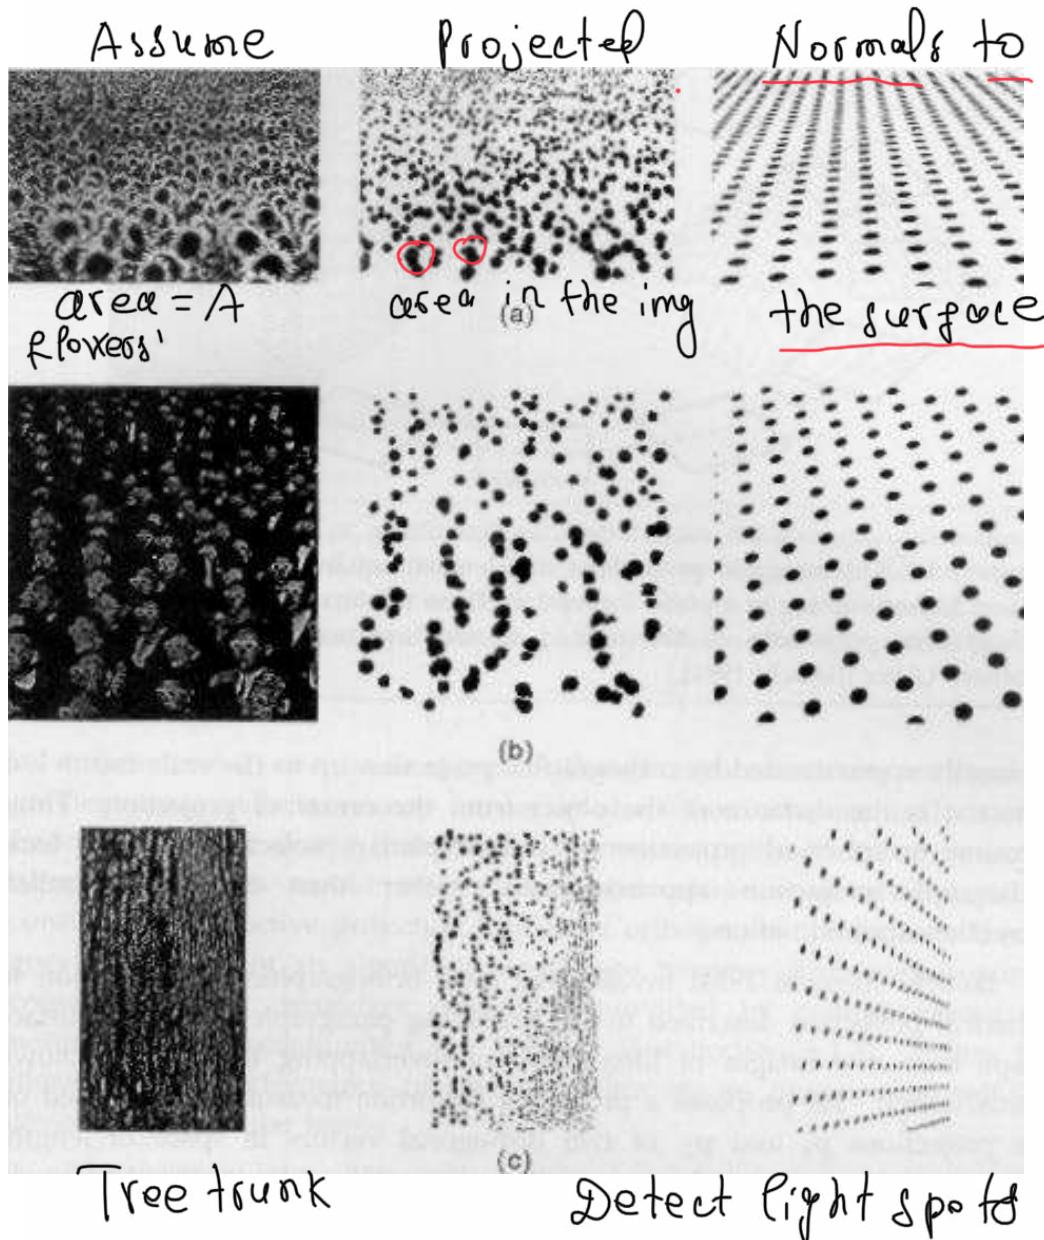
Parse when small when large

Surface faces camera  
Surface angled away from camera

texels appear smaller as their distance from cam ↑

EPFL

# Parapespective Projection



You can write eq's that connect several eq's that connect the area in the world to observed area to the normal

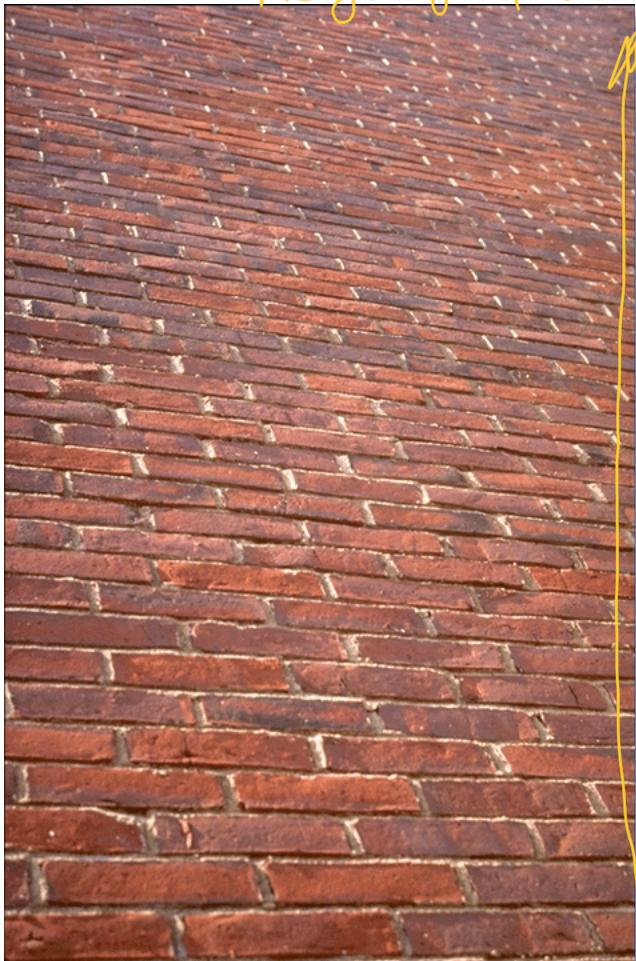
**Texels:**

- Image regions being brighter or darker than their surroundings.
  - Assumed to have the same area in space. Compute normals with the fact of over-constrained task
- Given enough texels, it becomes possible to estimate the normal.

Bottom Pine: If you can detect enough texture elements, then you can estimate shape of surface

# Texture Gradient

Back, things get composed of  
As you go further and further



Textures on the front != textures on the back



So far, we talked about structural texture

# Statistical Shape Recovery

—  
textures



Measure texture density as opposed to texel area, that is, the number of textural primitives per unit surface.

Assuming the texture to be **homogeneous**, we have:

$$\psi \mathbf{n} \propto \mathbf{b}$$

Unknown surface normal.

$$\psi = \begin{bmatrix} u_1 & v_1 & 1 \\ \dots & \dots & \dots \\ u_n & v_n & 1 \end{bmatrix}^t$$

$$\mathbf{b} = [b_1, \dots, b_n]^t$$



density you measure at which  
Image coordinates.

Function of density.

$$\Rightarrow \mathbf{n} = \frac{\psi \mathbf{n}}{\|\psi \mathbf{n}\|}$$

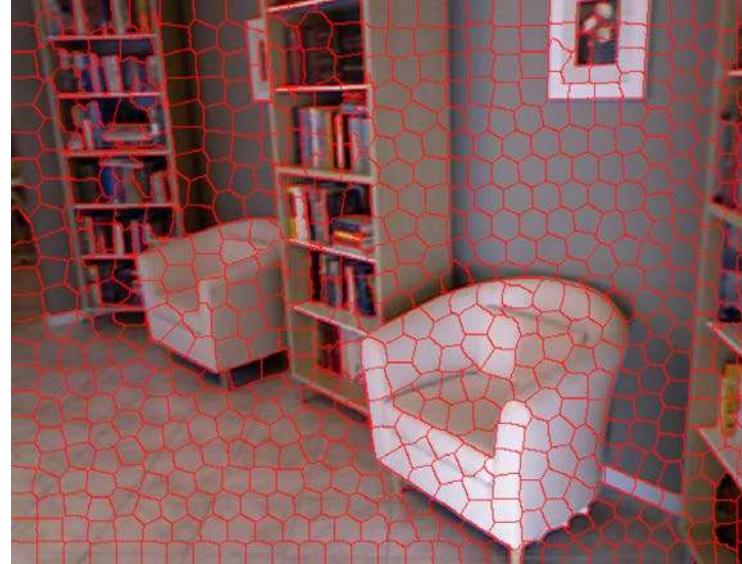
normal

Connect densities observed at different parts of image  
 $\Leftrightarrow$  normal to the plane

# Machine Learning

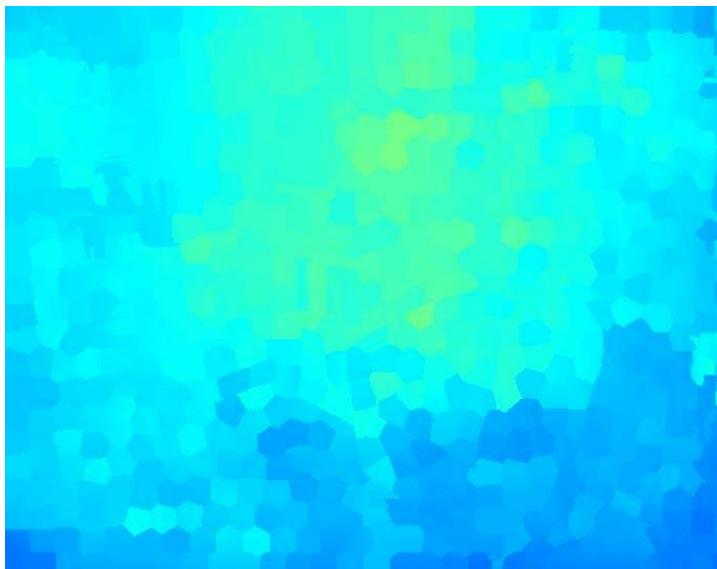


Input Image



Superpixels

Previously, we were assuming to have circles, squares, some textural elements, some textural densities over whole surface with a single normal



Input: Single Image  
Output: Surface Normals

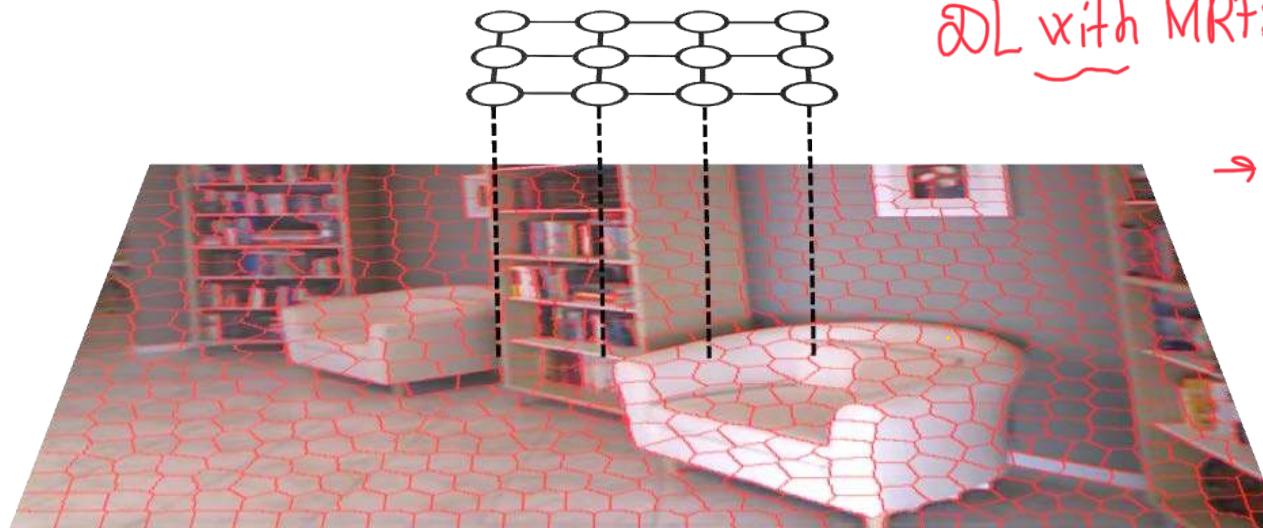
even for non-homogeneous  
textures

Because we tackled w superpixel deeper.  
If we use ML, results will be noisy

# Markov Random Field (MRF)

If you can compute normals  $\Rightarrow$  you can integrate and then compute 3D shape

## Graph with vertices and edges



## Assign values to the nodes to minimize

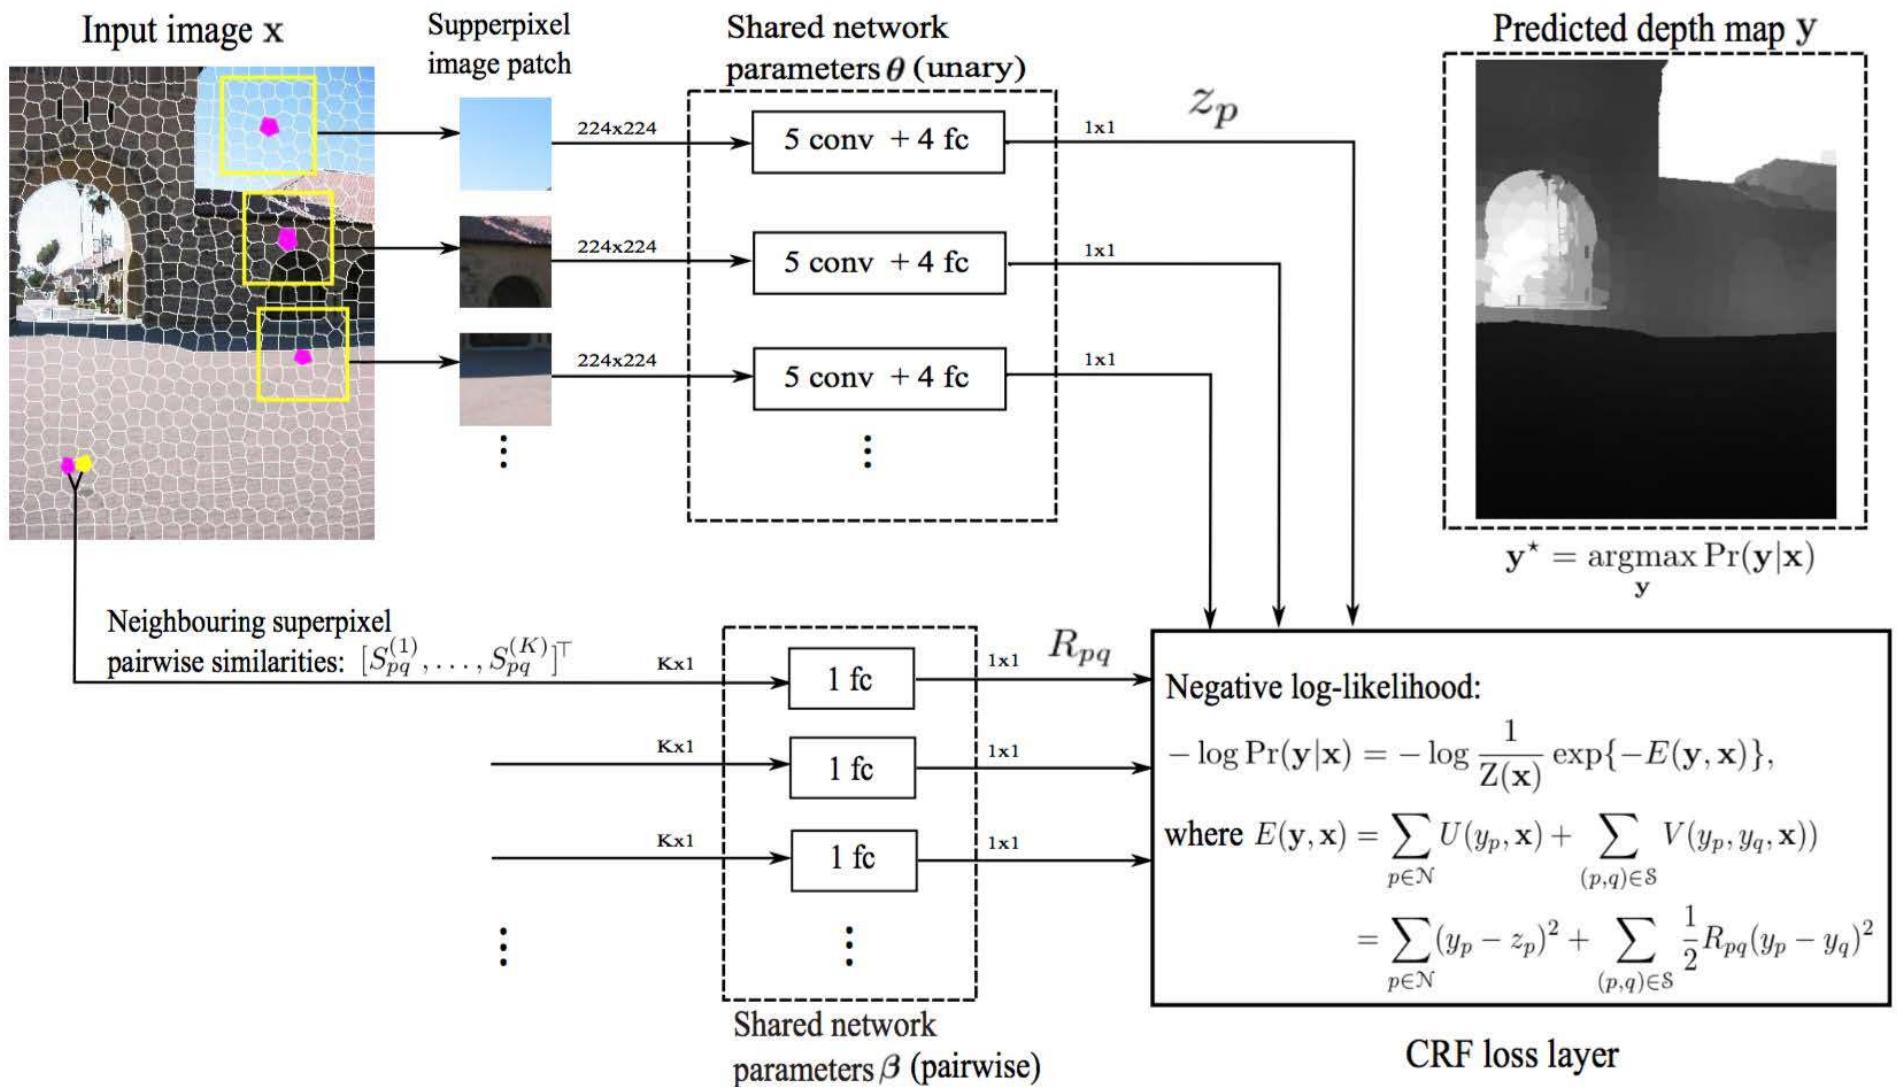
$$E(Y) = \sum_i \varphi(y_i) + \sum_{(i,j)} \psi(y_i, y_j)$$

unary pairwise

Why: Adding some spatial consistency  $\Rightarrow$  Enforces connectivity and smoothness

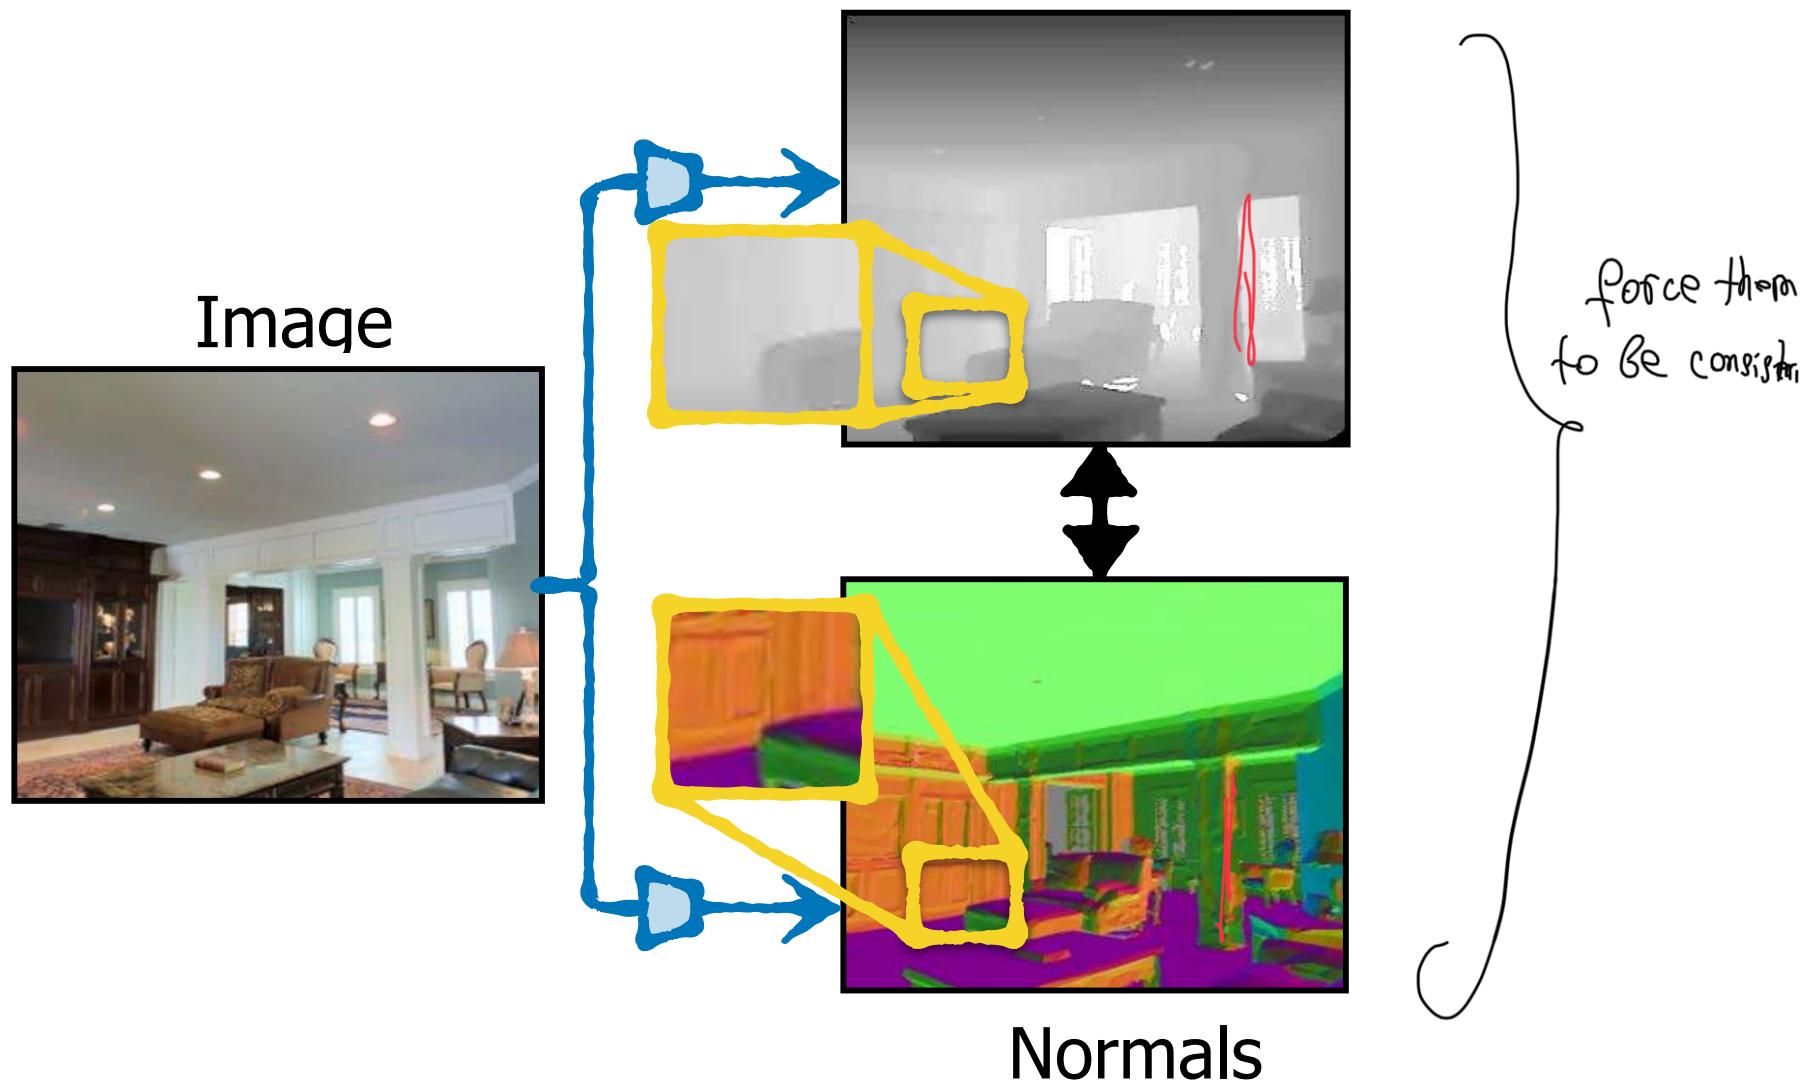
$\rightarrow$  Enforces consistency

# Deep Learning with MRF



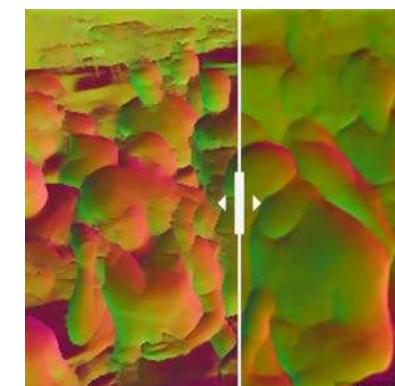
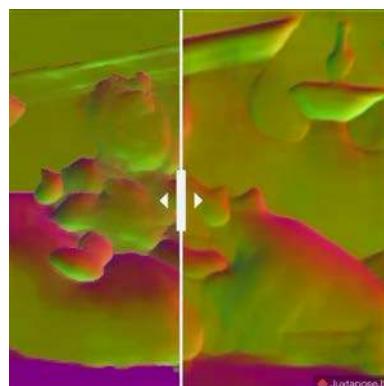
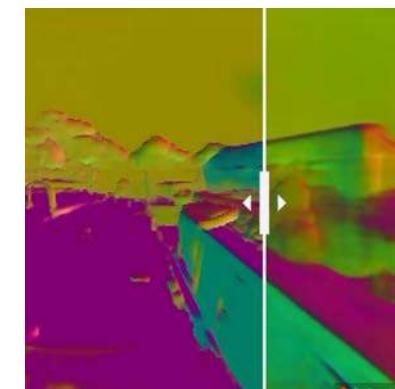
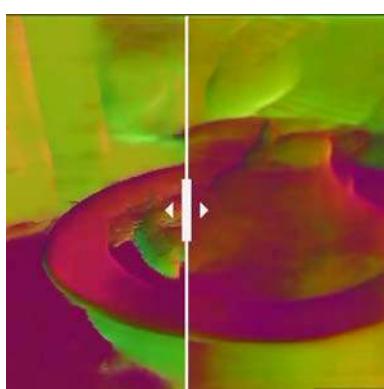
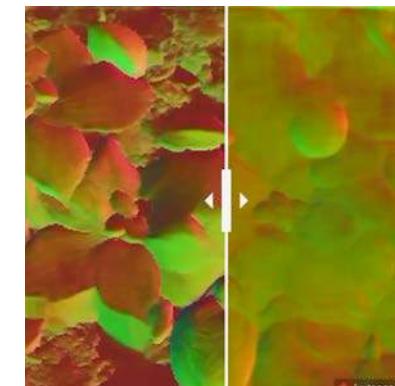
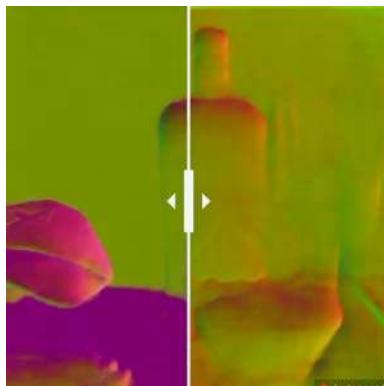
# Enforcing Task Consistency

Across modalities/tasks

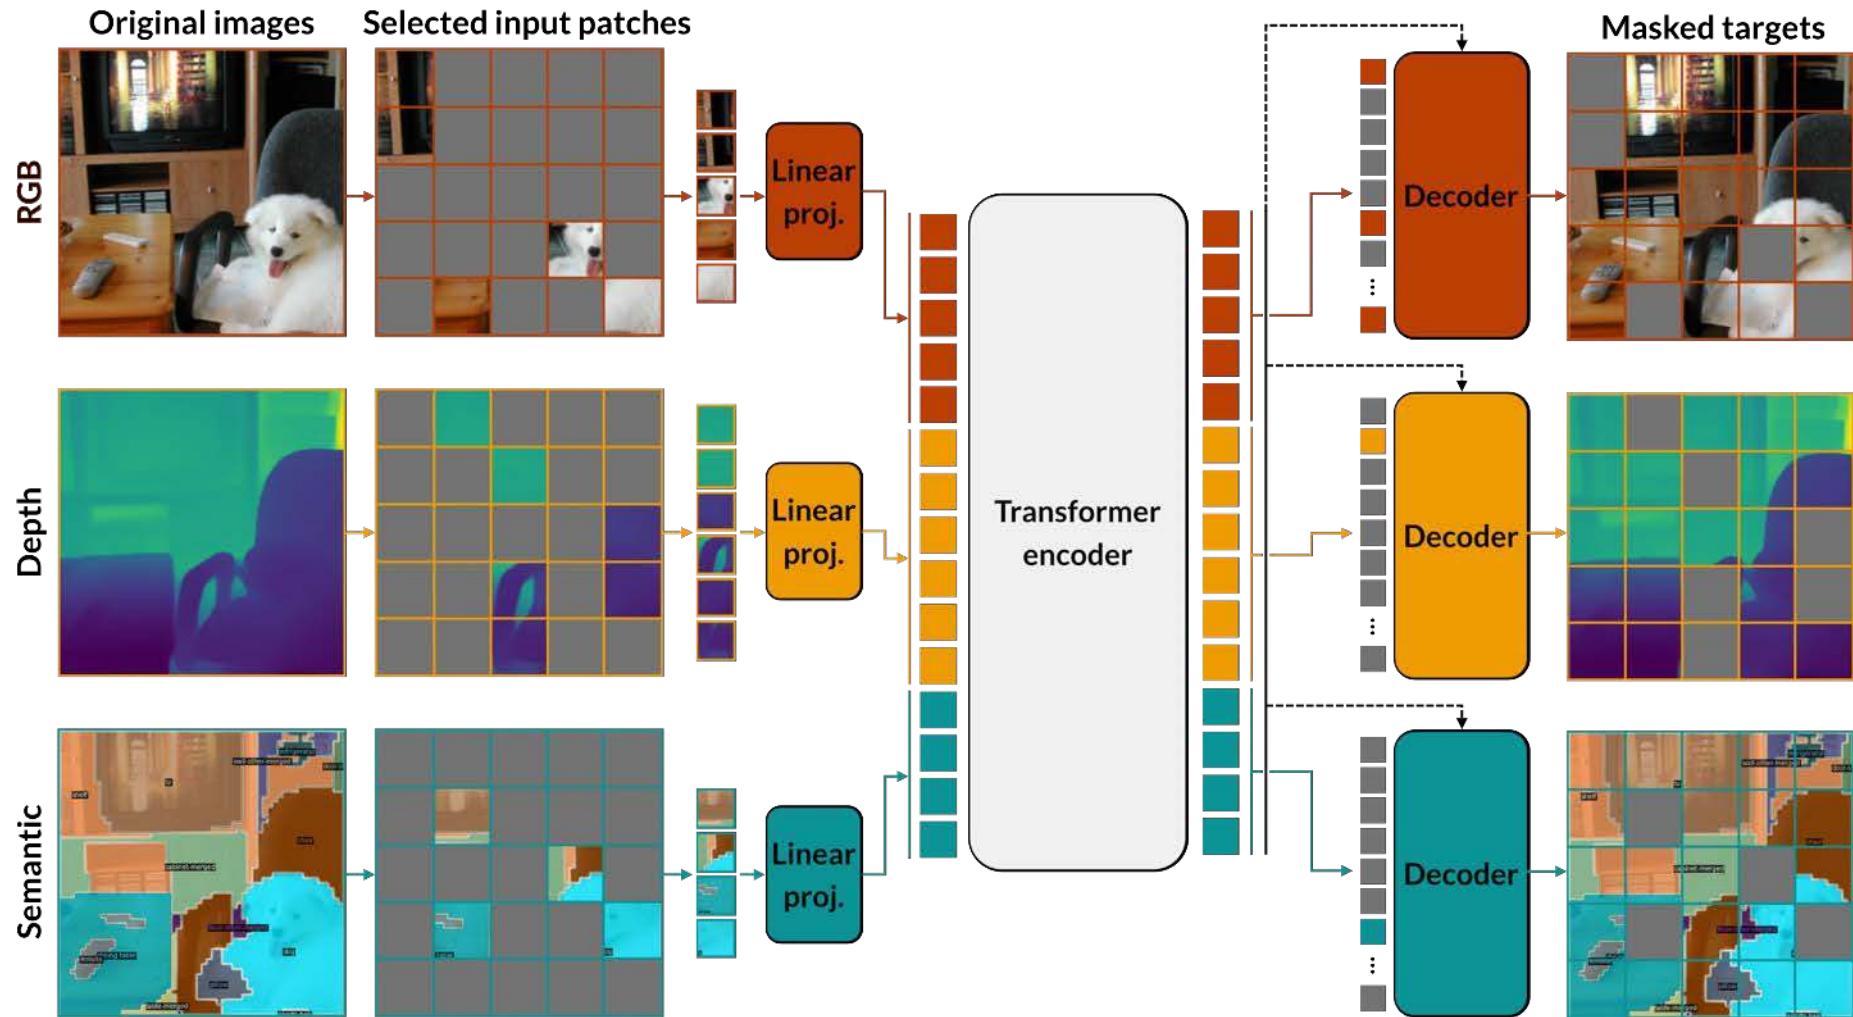


- A network can be trained to predict multiple things.
- Forcing consistency across tasks increases robustness.

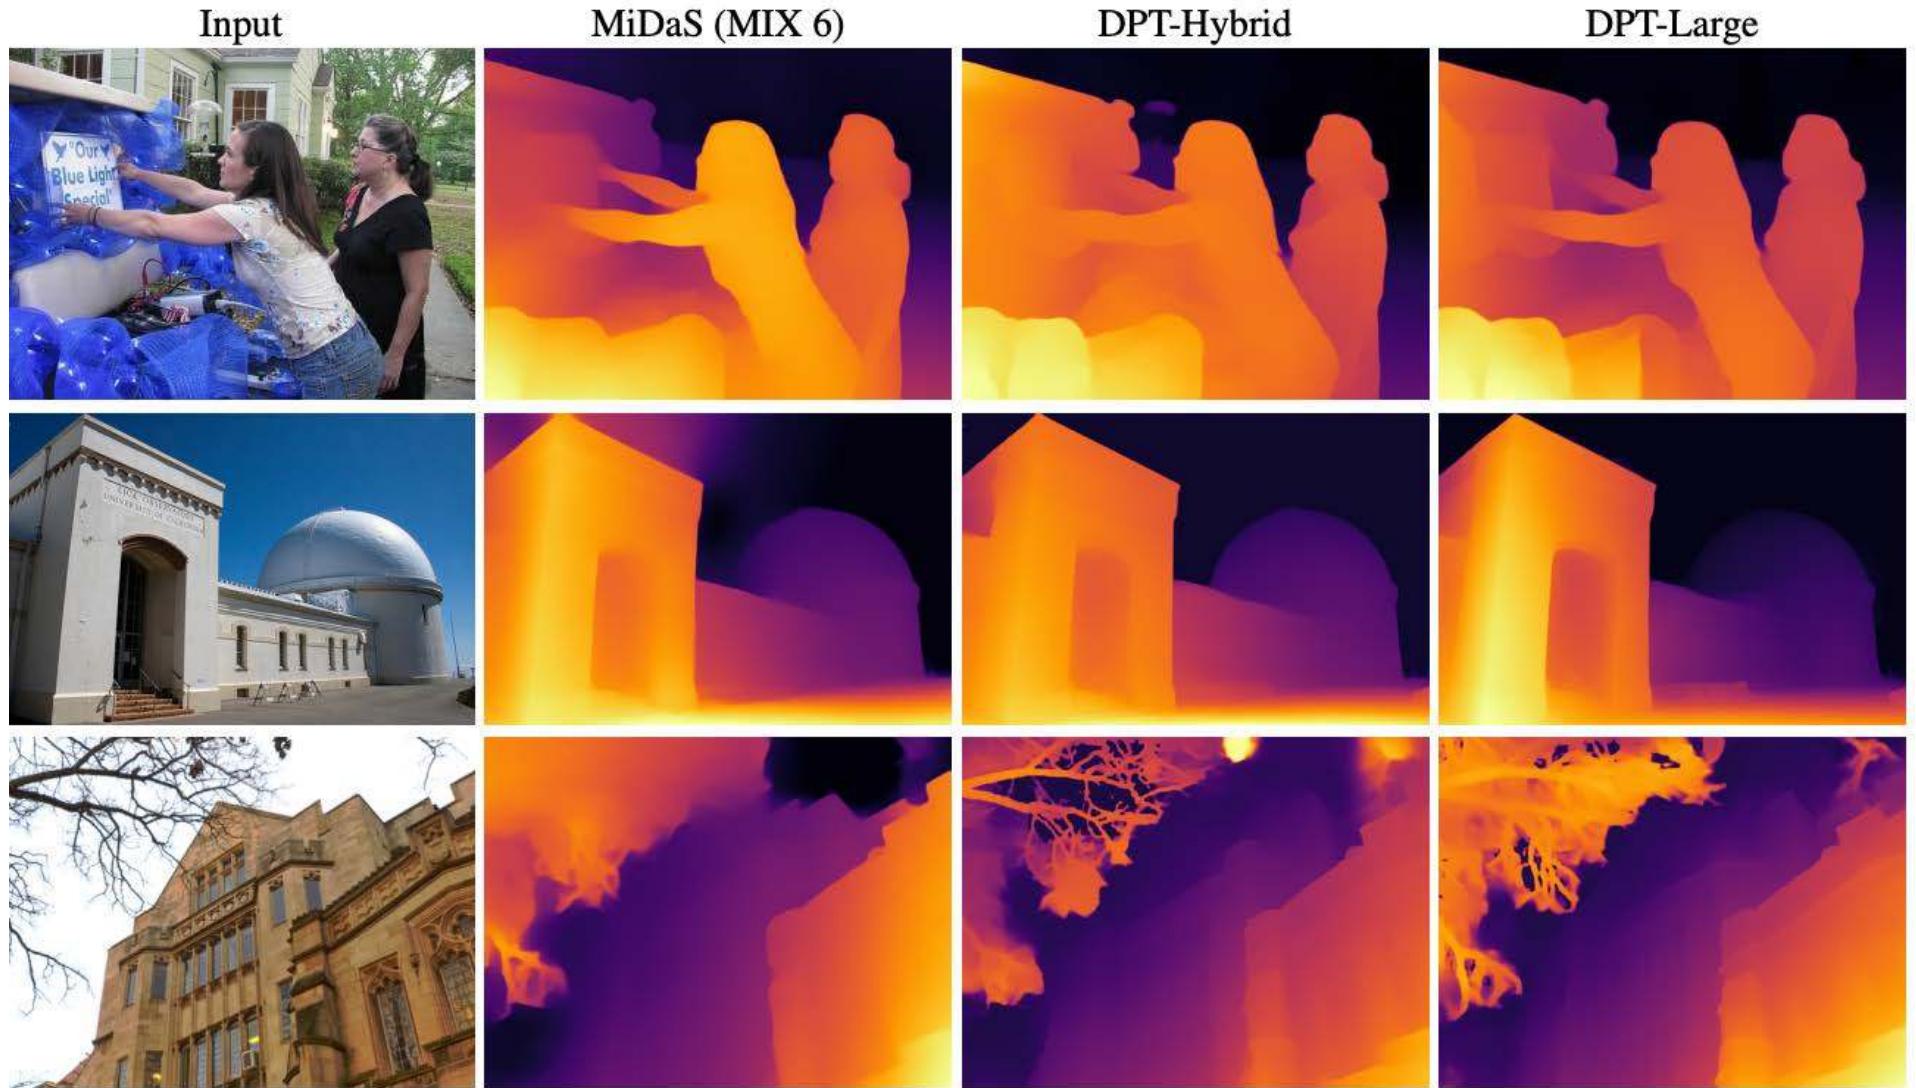
# A Very Diverse Training Database Helps



# .. and so does a Transformer Architecture

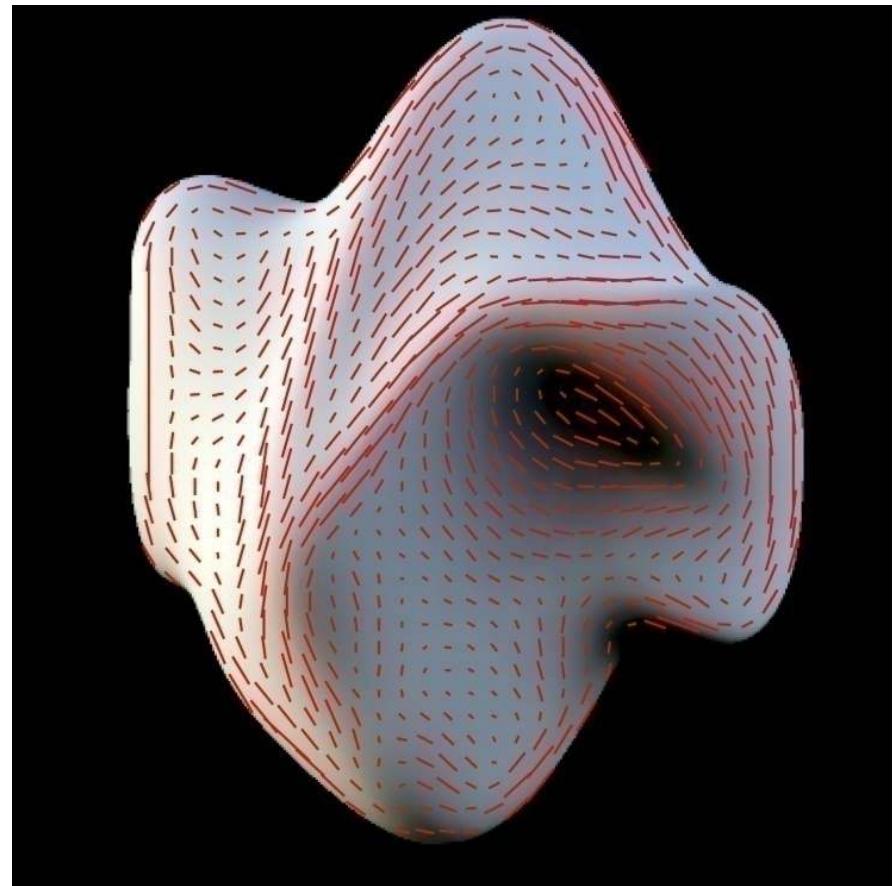
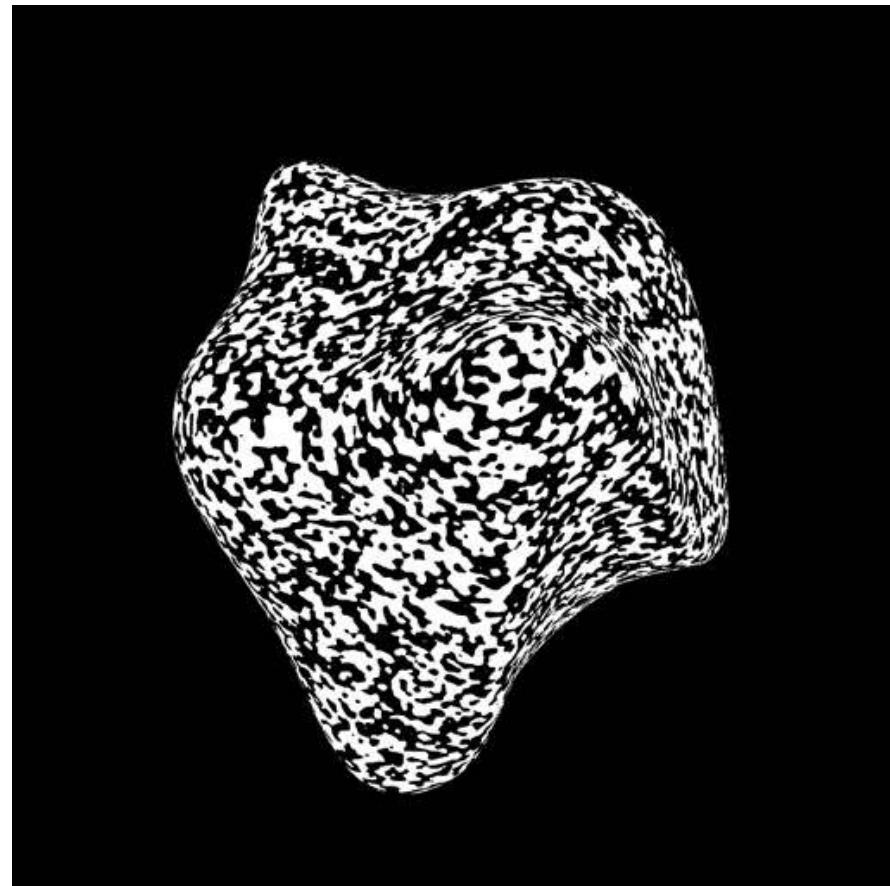


# Using Transformers



- Pros: Good at modeling long range relationships.
- Cons: Flattening the patches loses some amount of information.

# Optional: Illusory Shape Distortion

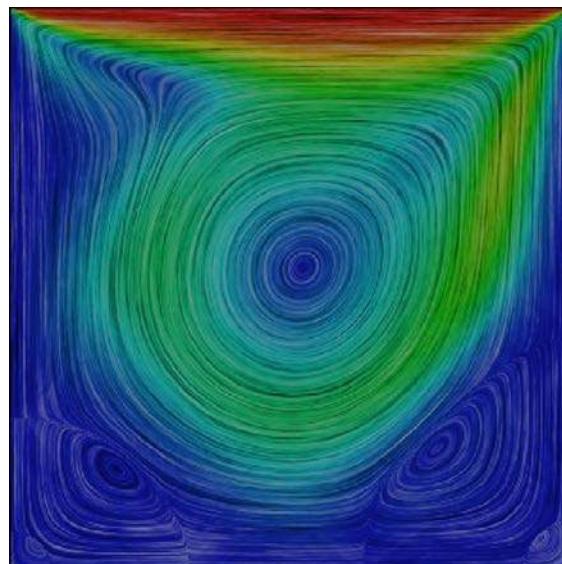


People seem to be sensitive to orientation fields in the cases of both texture and shading.

# Optional: Shape from Smear

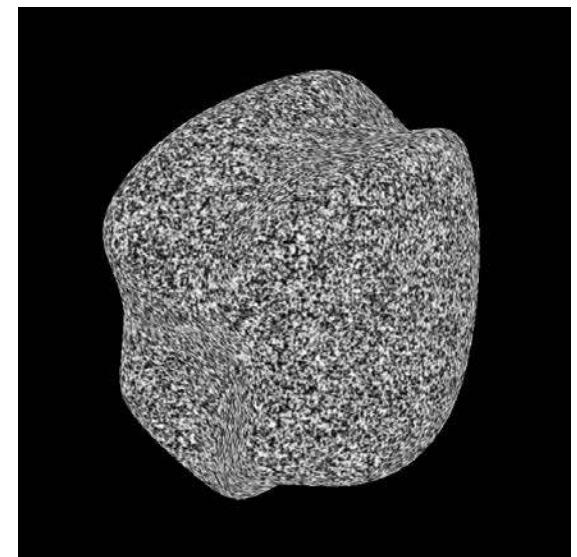
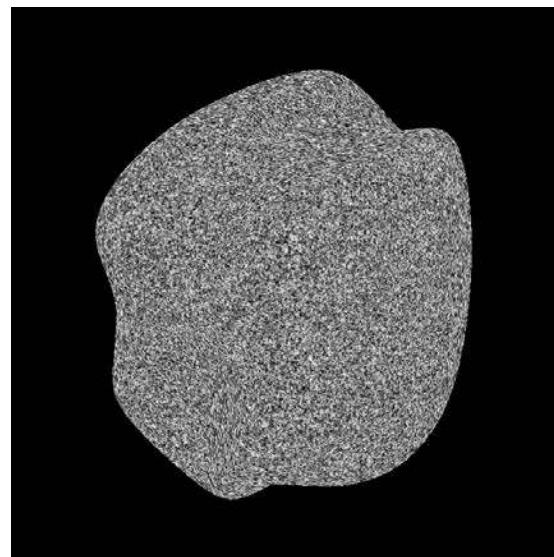
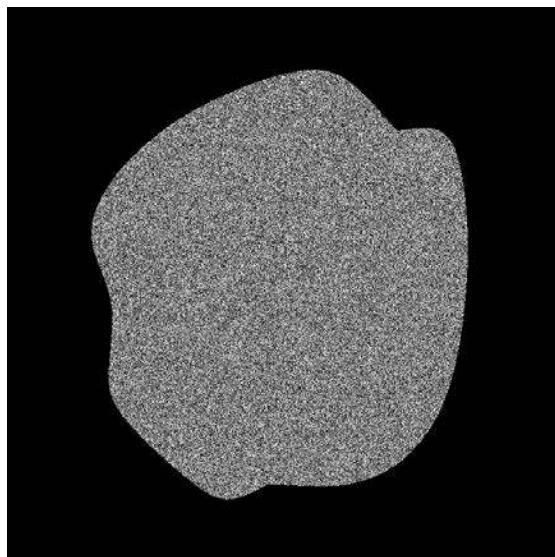
**Hypothesis:** If orientation and scale fields are the key source of information for 3D shape perception, it should be possible to induce a vivid sense of 3D shape by creating 2D patterns with appropriate scale and orientation fields.

**Test:** Use a technique known as Line Integral Convolution to smear the texture along specific orientations and scale appropriately.

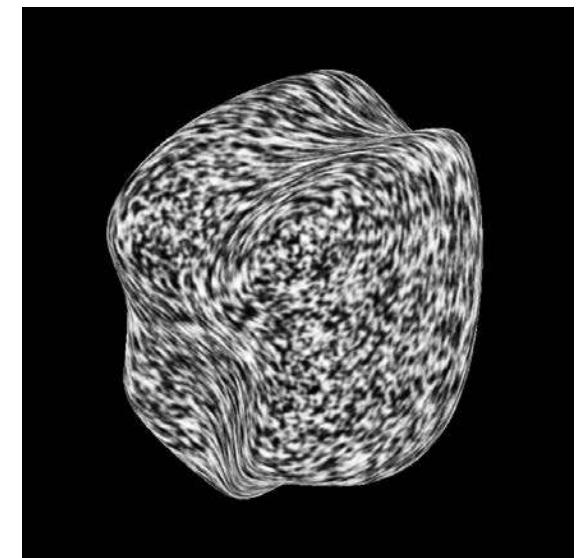
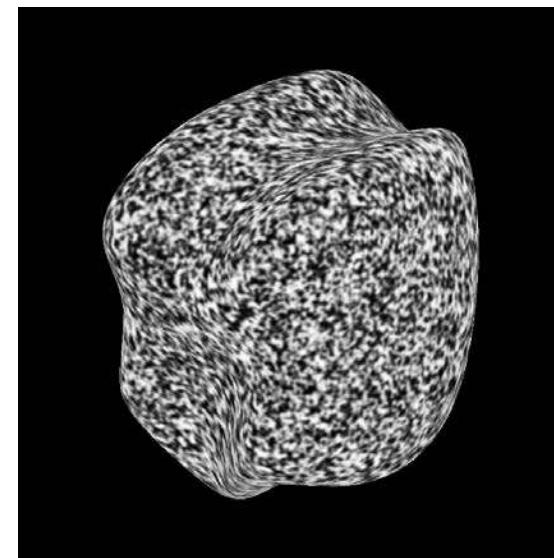
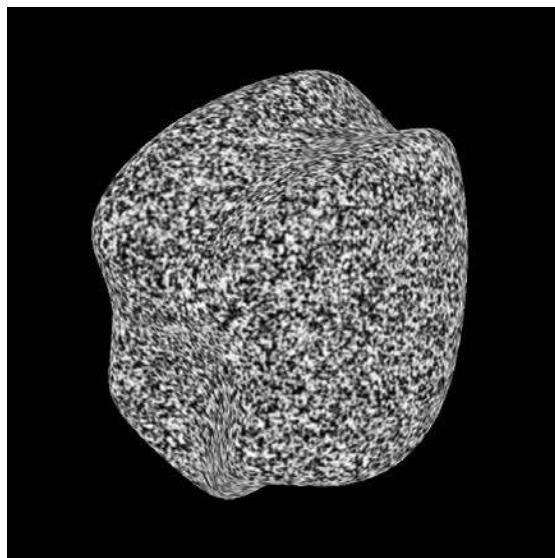


# Optional: Scaling and Smearing

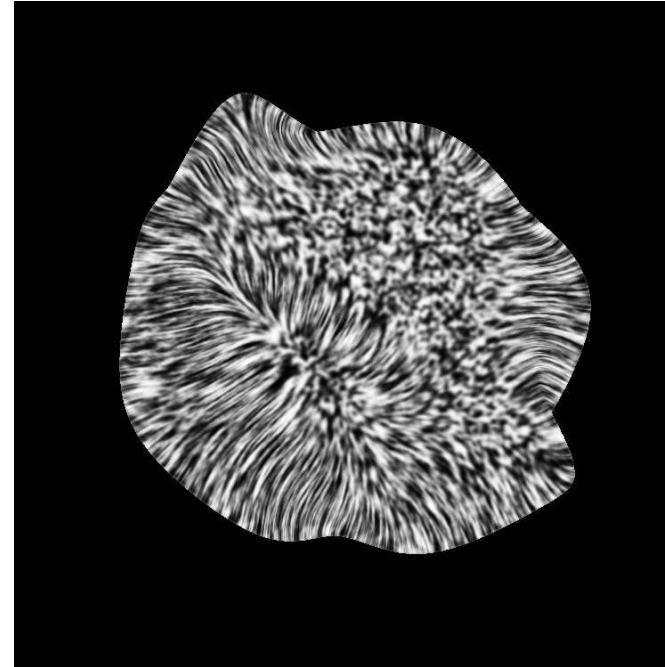
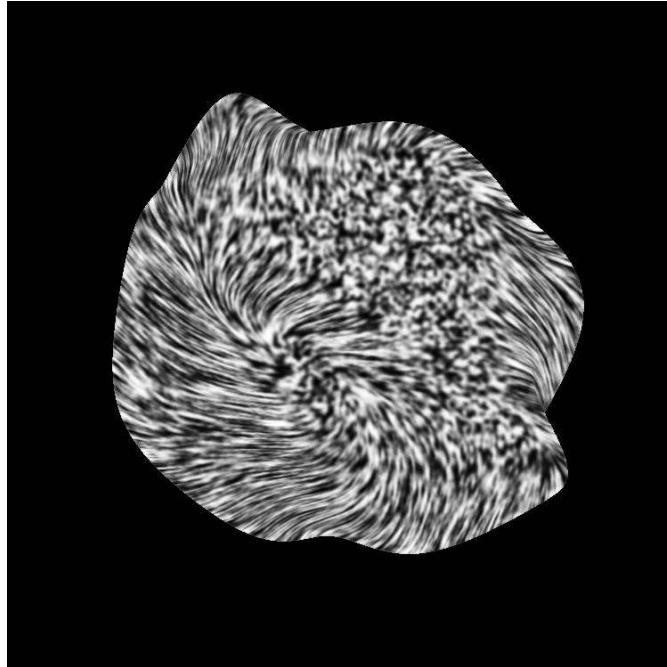
Scaling:



Smearing:



# Optional: Inconsistent Stimulus



The orientation field cannot be integrated

- No depth perception.
- Do we integrate in our heads?
- Is this what the deep nets learn to do?

# Strengths and Limitations

## Strengths:

- Emulates an important human ability.

## Limitations:

- Requires regular texture.
- Involves very strong assumptions.  
Limited set of cases  
trained on specific settings
- Deep learning can be used to weaken them.