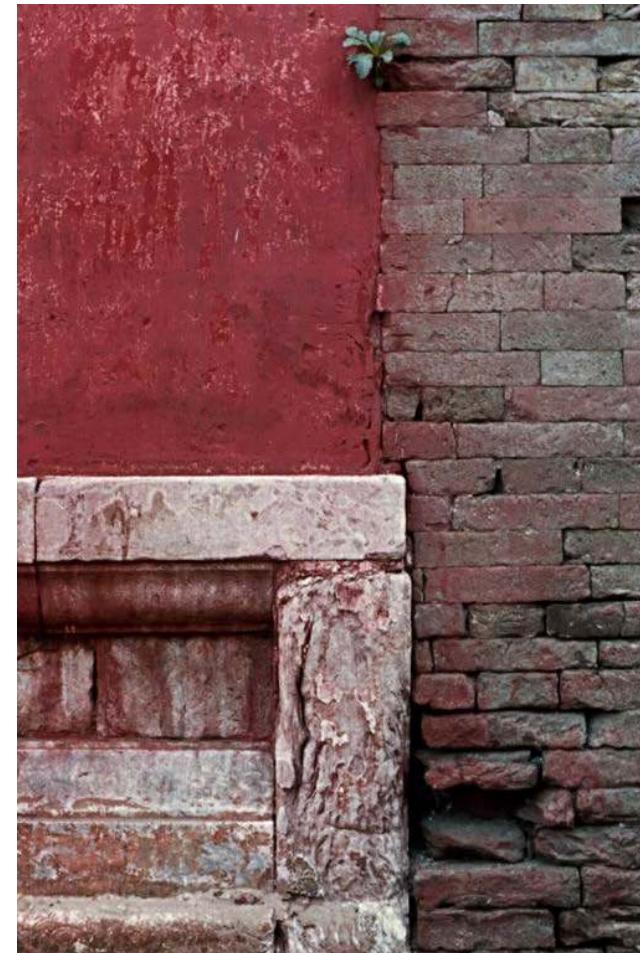


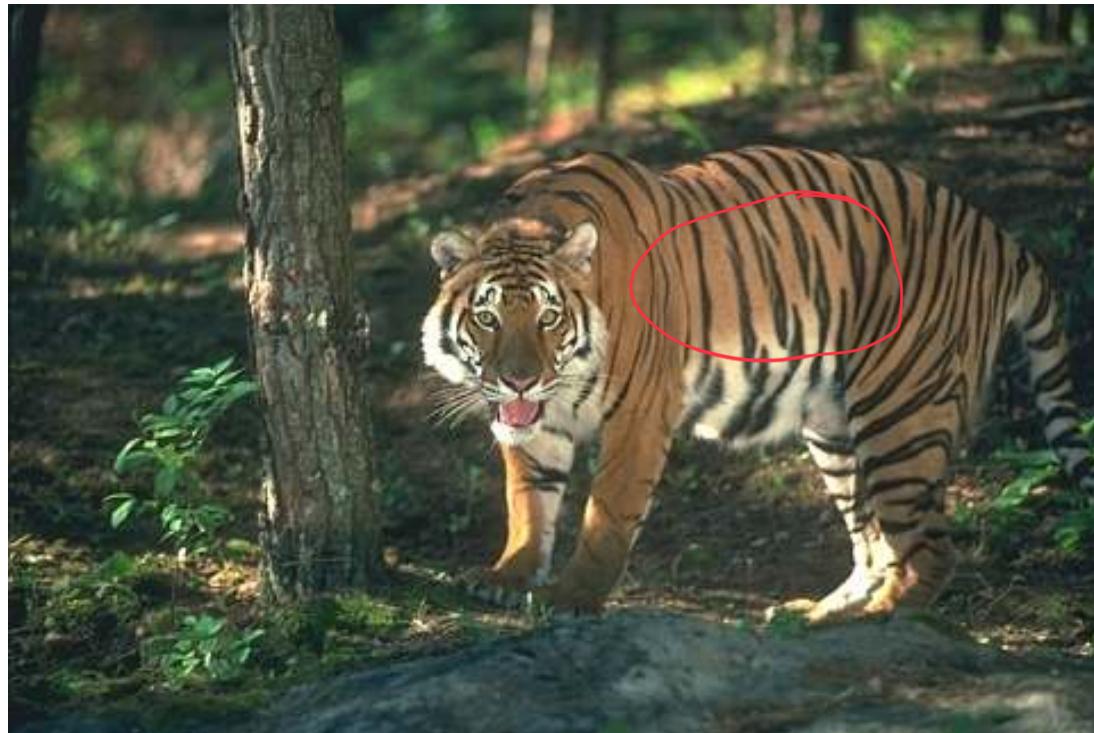
# Texture

→ Local statistics of the image

- What is texture?
- Texture analysis
- Deep Texture



# Reminder: Homogeneous or Not?



↓ What does it mean?

*in some sense*  
What is homogeneous in some parts of these images are the statistical properties, not the actual pixel values.



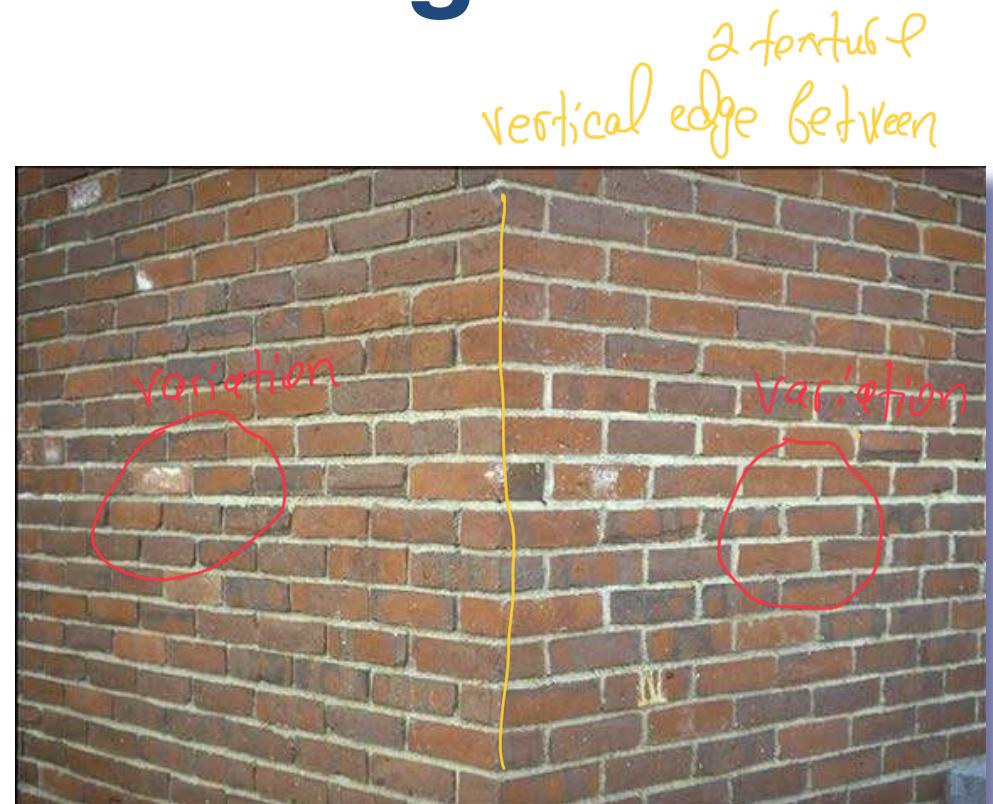
# Texture-Based Segmentation



Ideally, we would like to:

- Assign to individual pixels whose texture is similar the same values to form a textural image.
- Evaluate homogeneity both in the original image and in the textural one.

# Texture-Based Edges



Similarly, we would like to be able to find boundaries between textures.

# What is Texture?



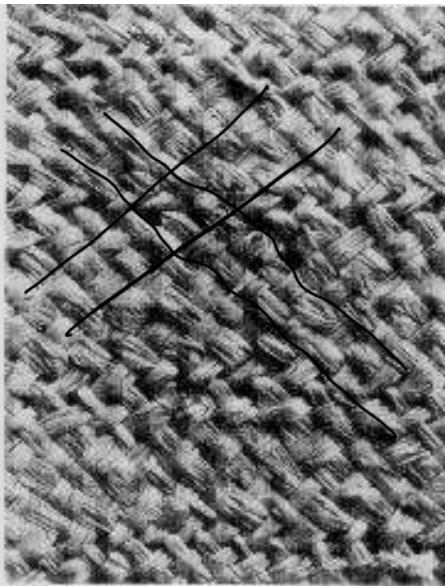
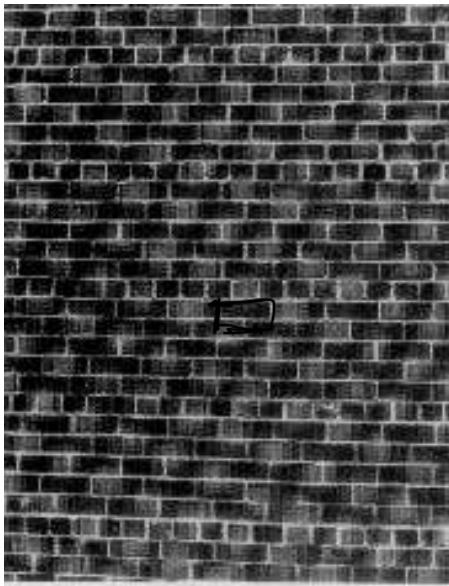
no regular  
shape

Repetition of a basic pattern:

- Structural ← geometry pattern that repeats
  - Statistical ← reasonably similar statistics across whole img (a constant density)
- Non local property, subject to distortions.

texture  
does not 1 pixels

# Structural Textures



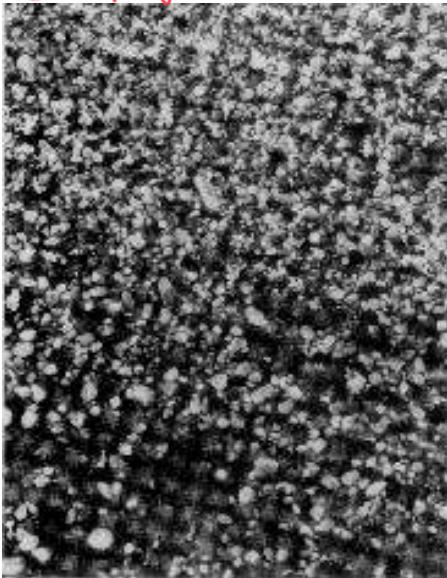
e.g. windows → texels  
not buildings

## Repetitive Texture Elements (Texels)

A texel represents the smallest graphical element in a two-dimensional texture that creates the impression of a textured surface.

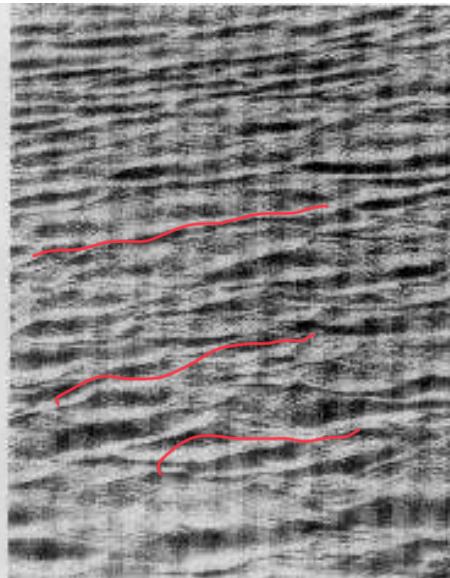
# Statistical Textures

→ constant  
density of pebbles →



statistical property  
that repeats

waves → constant  
distance between



→ statistically consistent  
orange, black stripes → repetitive



no geometric pattern

## Homogeneous Statistical Properties

# Textured vs Smooth

A “featureless” surface can be regarded as the most elementary spatial texture:

- Microstructures define reflectance properties.
- They may be uniform or smoothly varying. *You don't see relief*

→ Texture is a scale dependent phenomenon

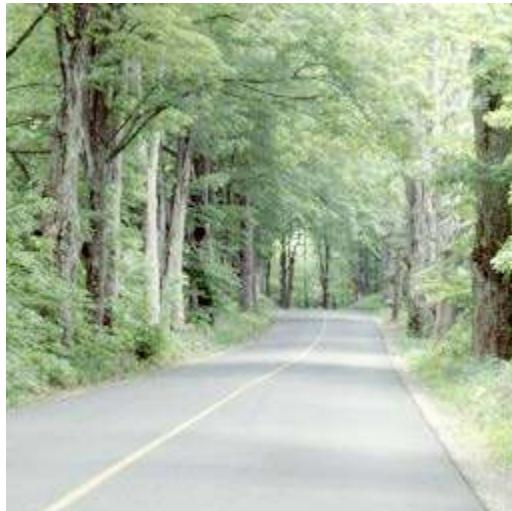
# Scale Dependence



At these two different scales, the texture seems very different.

# Structural vs Statistical

- Segmenting out texels is difficult or impossible in most real images.



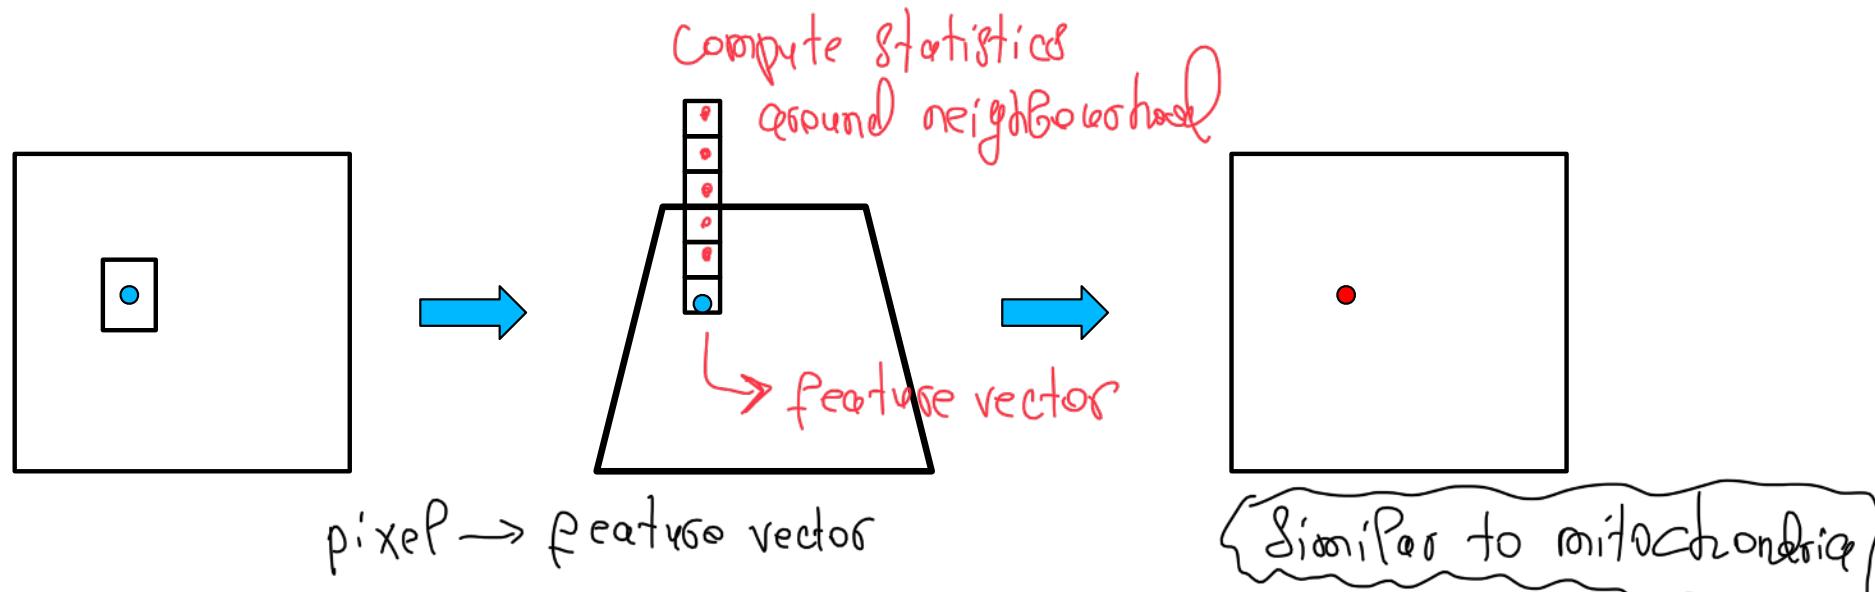
What are the fundamental texture primitives in this image? *elements?*

- Numeric quantities or statistics that describe a texture can be computed from the gray levels or colors alone. (*easier to implement*)  
*image*
- The statistical approach is less intuitive, but more effective in practice.

# Creating Textural Images

There is no such thing of statistics of the pixel

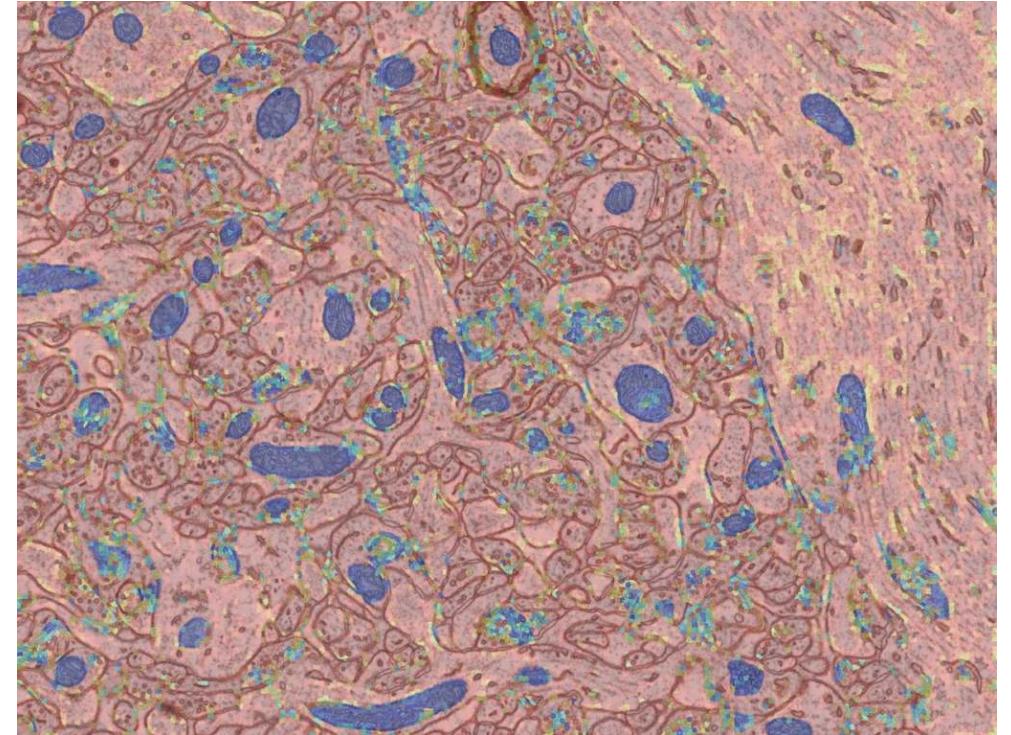
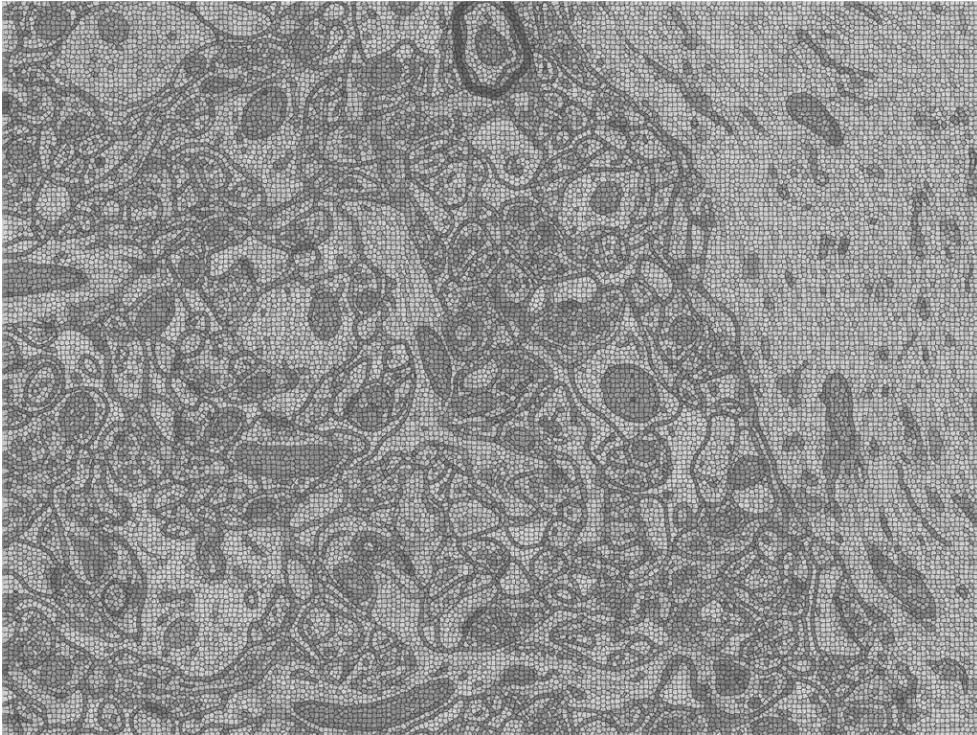
Because texture is non-local, the texture of individual pixels must be estimated using neighborhoods that surround them:



- For each pixel, compute a feature vector using either an image patch or a set of filters.
- Run a classification algorithm to assign a texture value to each pixel.

Convolving img patches

# Reminder: Mitochondria



- Compute image statistics for each superpixel.
- Train a classifier to assign a probability to be  
within a mitochondria.  
→ We used the super pixels to compute local  
statistics.

*to each superpixel  $P$*

# Textural Metrics

## Spectral metrics:

- Texture is characterized by the properties of its Fourier transform.

## Statistical Metrics:

- Texture is a statistical property of the pixels' intensity and color in a region. vrt each other

## Deep Net Metrics:

- They have now mostly superseded the others.
- They encompass the earlier concepts.

Superset of the first 2

# Reminder: Discrete Fourier Transform

$$F(\mu, \nu) = \frac{1}{\sqrt{M * N}} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-2i\pi(\mu x/M + \nu y/N)}$$

*image  $f$*   
*sum of weighted complex exponentials*

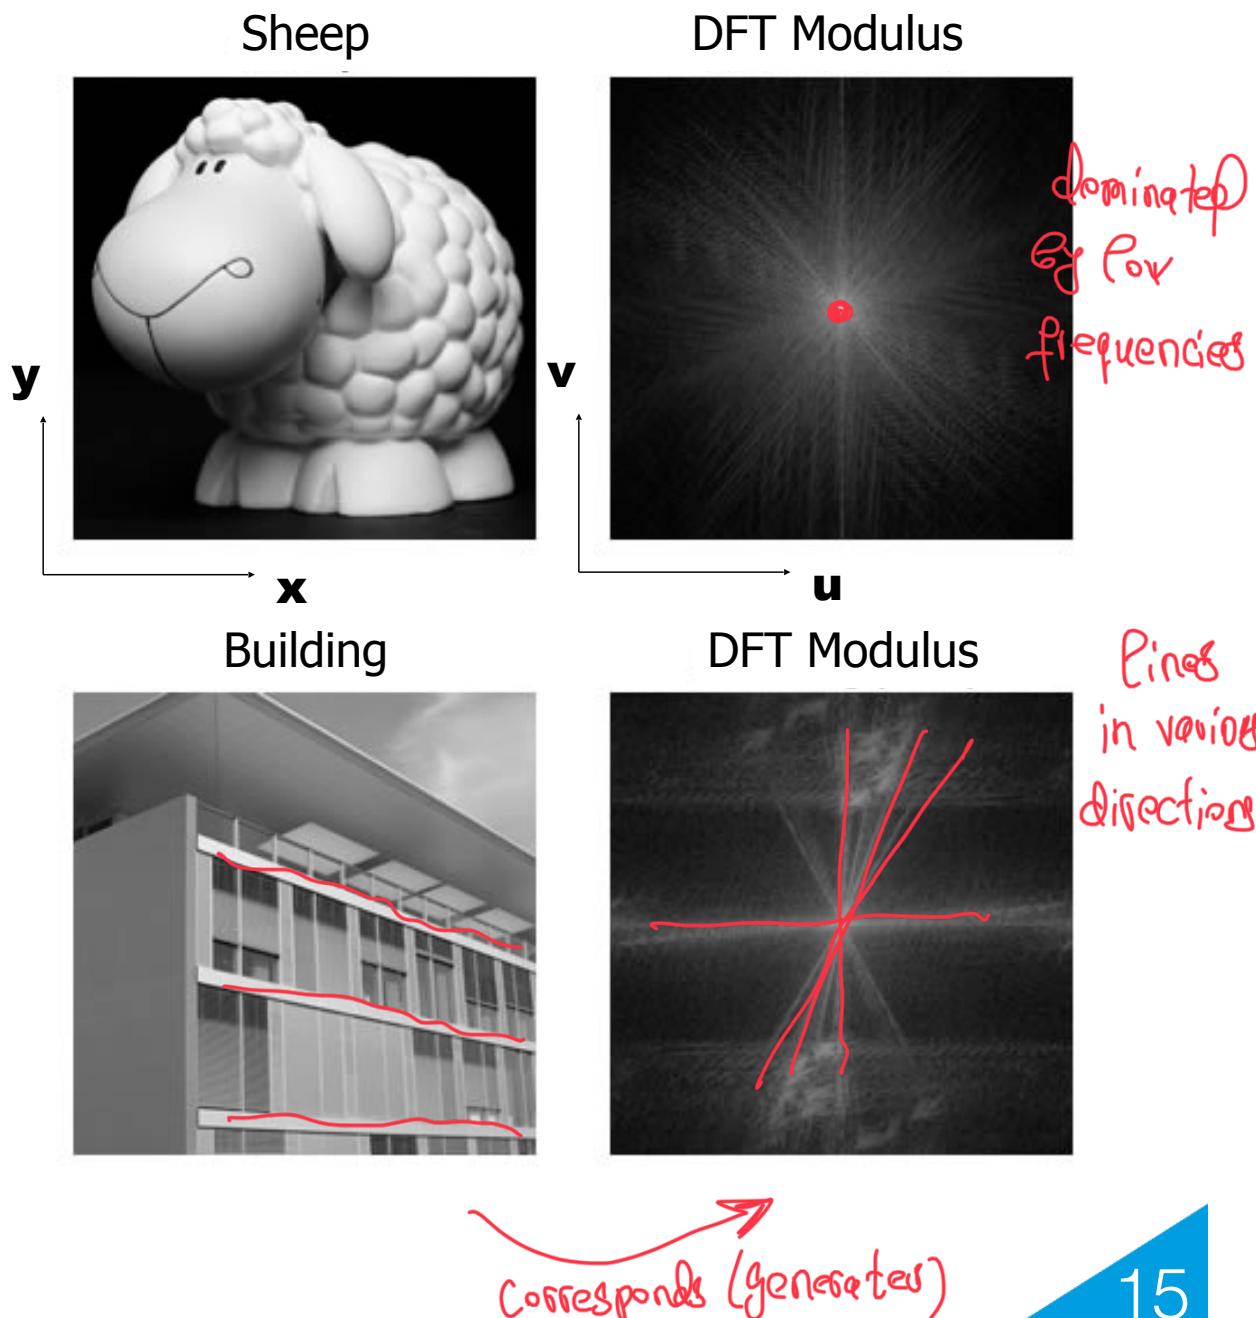
$$f(x, y) = \frac{1}{\sqrt{M * N}} \sum_{\mu=0}^{M-1} \sum_{\nu=0}^{N-1} F(\mu, \nu) e^{+2i\pi(\mu x/M + \nu y/N)}$$

The DFT is the discrete equivalent of the 2D Fourier transform:

- The 2D function  $f$  is written as a sum of sinusoids.
- The DFT of  $f$  convolved with  $g$  is the product of their DFTs.

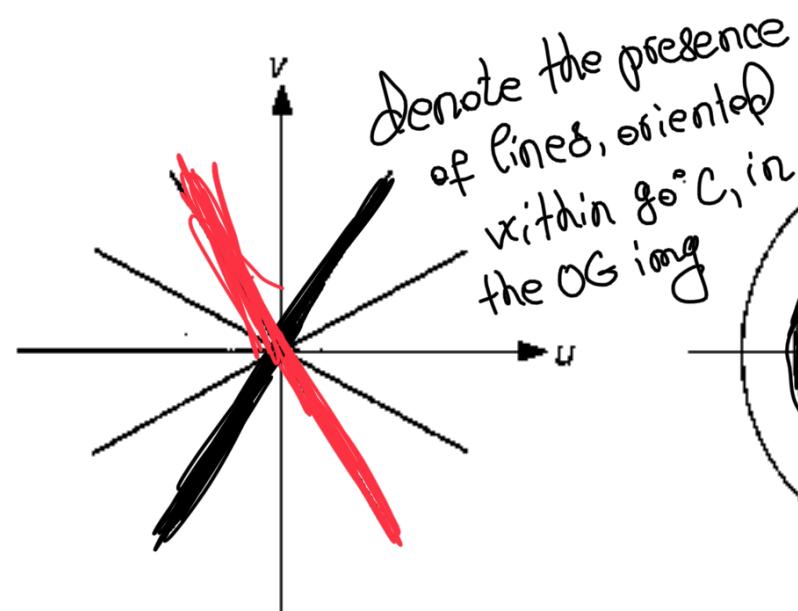
# Spectral Analysis

$F \rightarrow$  complex number

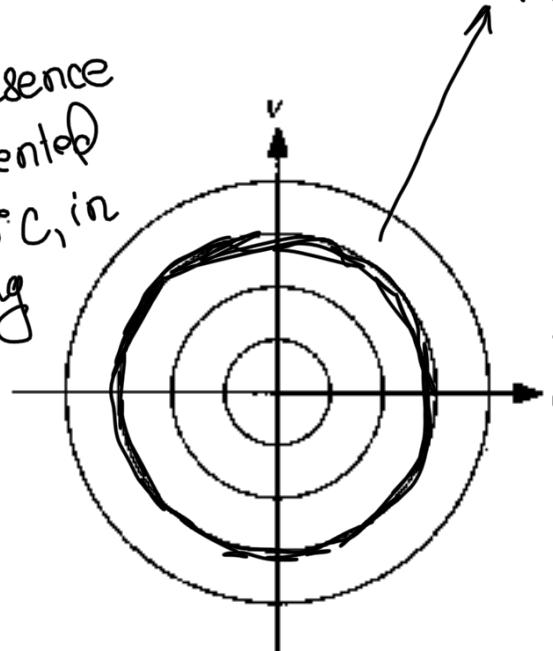


Lines in the DFT modulus images capture the main orientations in the image.

# Texture Analysis

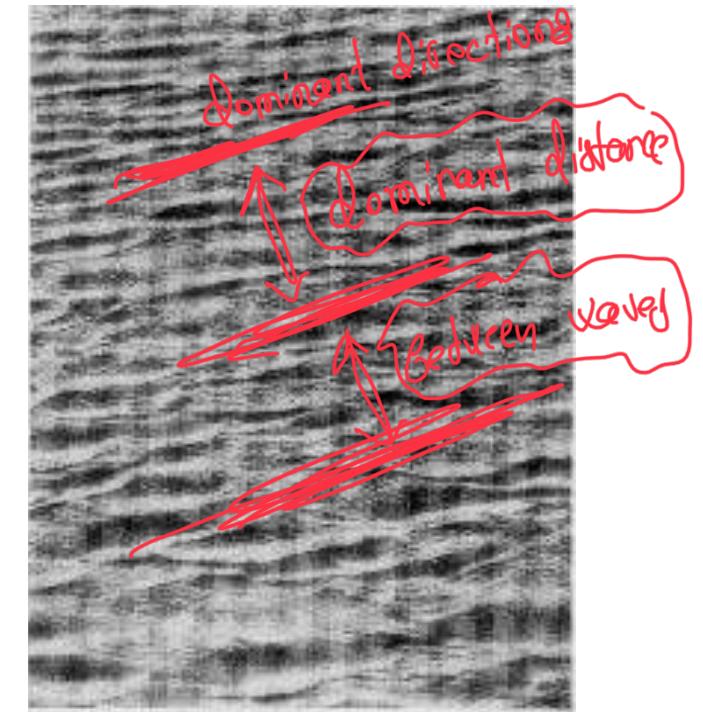


Angular bins  
in  $|DFT|$  image



Radial bins  
in  $|DFT|$  image

responses at particular frequency

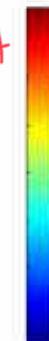
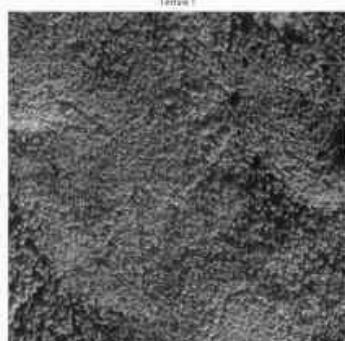


Compute DFT of wavef

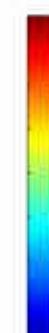
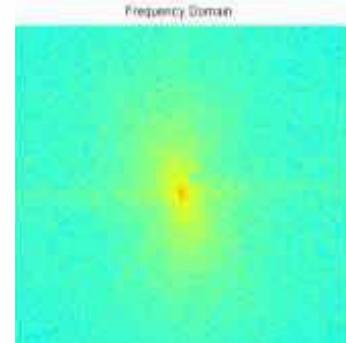
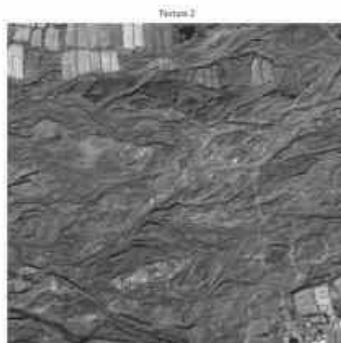
Angular and radial bins in the Fourier domain capture the directionality and fluctuation speed of an image texture, respectively.

# Fourier Texture Classification

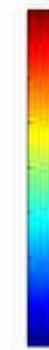
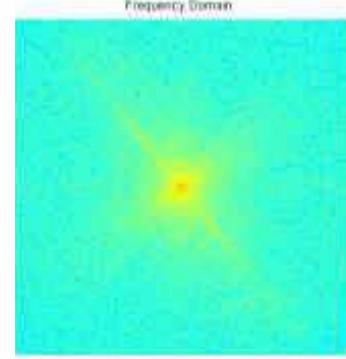
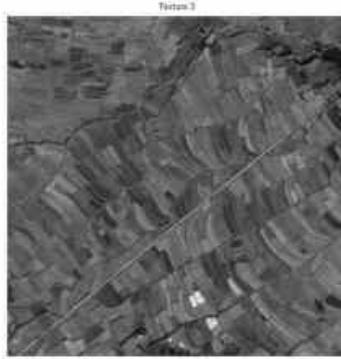
Forest



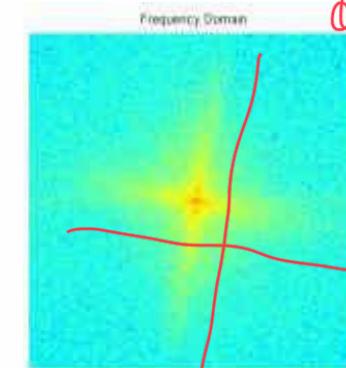
Mud



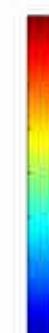
Fields



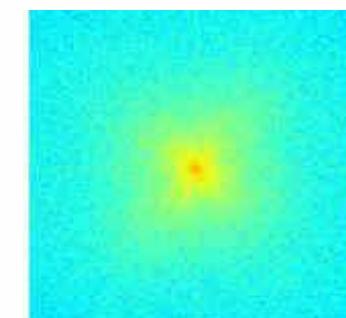
Ponds



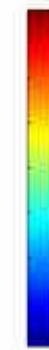
dominant



Village



dominant direction



Water



to differentiate

- For some types of textures, the Fourier spectra are easily distinguishable.
- A classifier can be trained to tell them apart. (Classified classifiers)
- However, one must have the same texture in the whole image patch.

# Limitations

- DFT on small patches is subject to severe boundary effects.
- Only applicable if texture is uniform over large areas.  

- Results can be improved by using wavelets instead, but only up to a point.

→ More local metrics are required.  
  
*for segmentation*

# Statistical Metrics

## **First order gray-level statistics:**

- Statistics of single pixels in terms of histograms.
- Insensitive to neighborhood relationships.

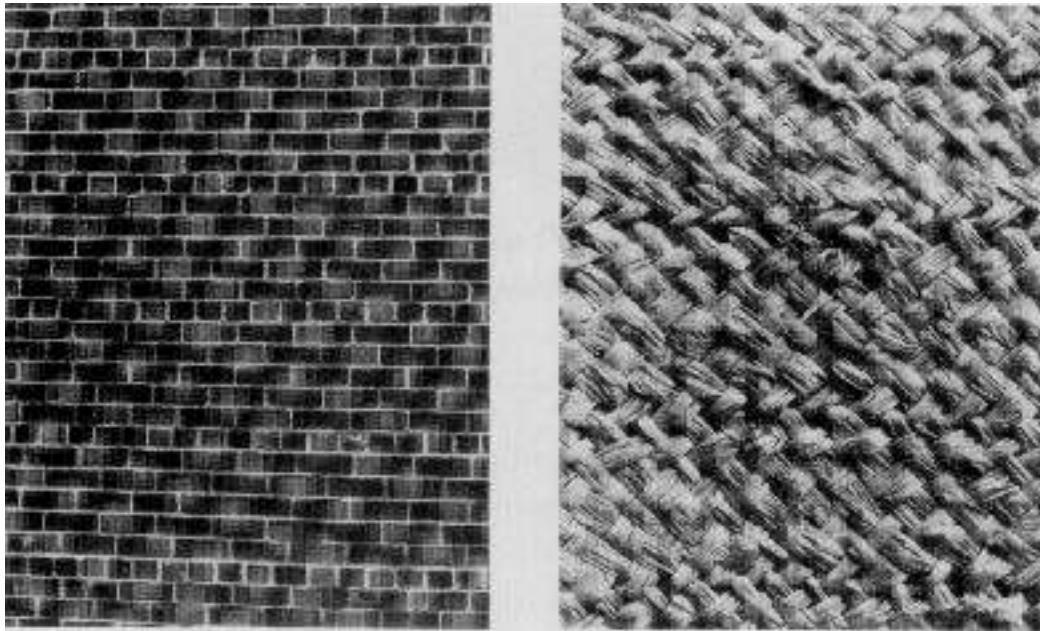
## **Second order gray-level statistics:**

- Statistics of pixel pairs.

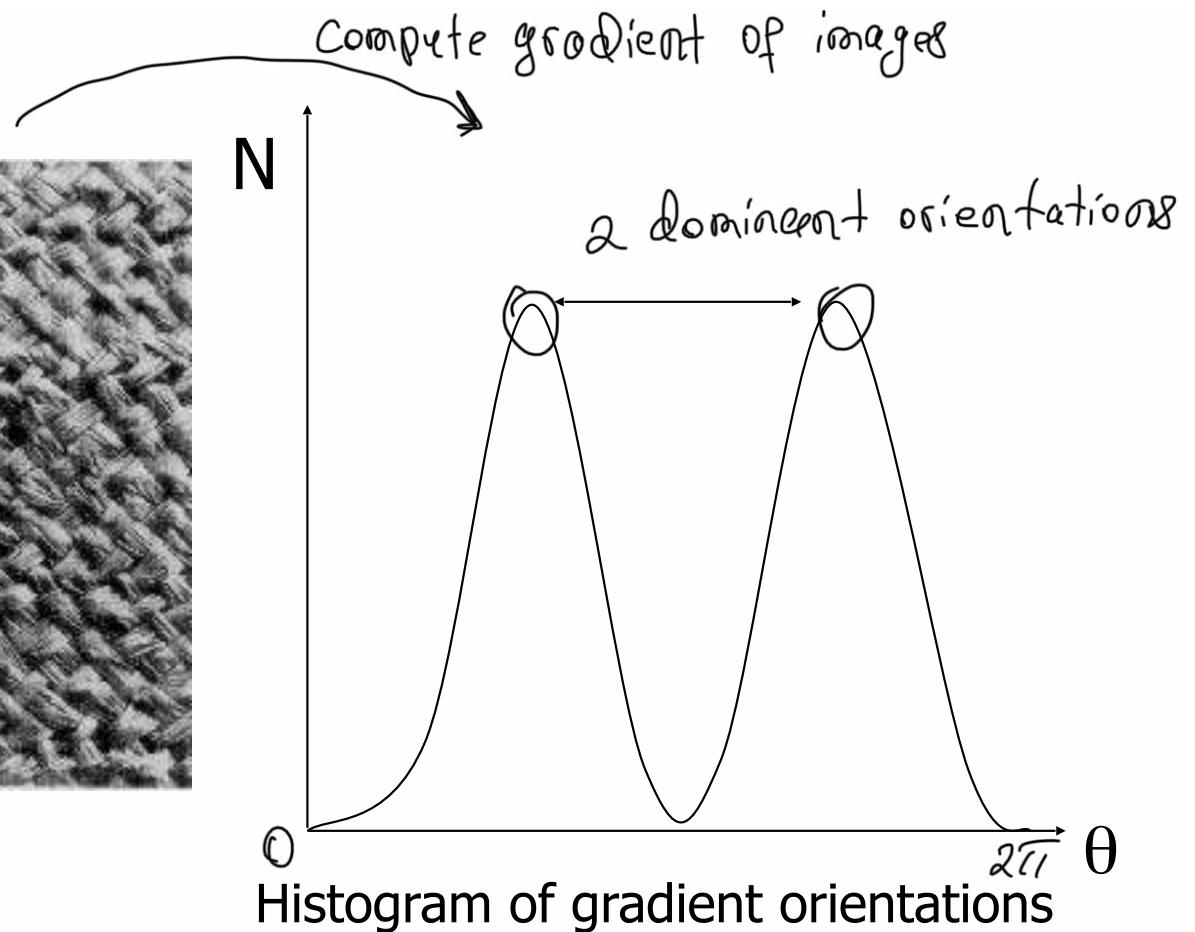
## **Filter-based measures:**

- Statistics of whole neighborhoods.

# Simple First Order Measure



How to recognize texture?



The orientation histogram gives a clue to the orientation of the underlying plane.

recognizable

# More First Order Measures

Edge Density and Direction:

Histogram of intensities

- Edge detection as a first step in texture analysis.
- The number of edge pixels in a fixed-size region tells us how busy that region is.
- The directions of the edges also help characterize the texture

Edgeness per unit area:

- $\{ p : \text{gradient\_magnitude}(p) \geq \text{threshold} \} / N$  where  $N$  is the unit area or region.

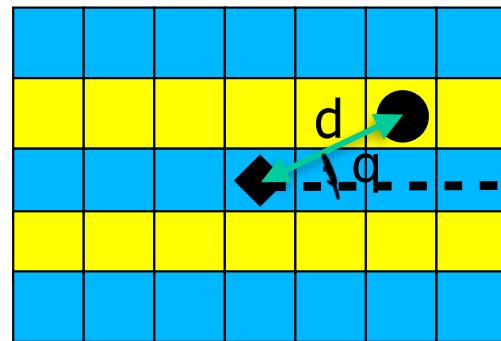
How sharp or pronounced the edges in a texture are

How many pixels ↑

Edge magnitude and direction histograms:

- $\{ H_G, H_t \}$

# Second Order Measures



Histogram of the co-occurrence of particular intensity values in the image.

- Specified in terms of geometric relationships between pixel pairs:
    - Distance
    - Orientation
  - $P(i,j,d,\theta)$  Frequency with which a pixel with value  $j$  occurs at distance  $d$  and orientation  $\theta$  from a pixel with value  $i$ .
- (i)* *(j)*  
of a transition from one gray-level to another

# Simple Example

Small images coded

$0, 1, 2, 3 \rightarrow 2\text{-bits}$

If  $I = \begin{bmatrix} 0 & 1 & 2 & 3 & 1 \\ 2 & 1 & 3 & 1 & 1 \\ 0 & 0 & 2 & 2 & 1 \\ 1 & 2 & 0 & 3 & 1 \\ 0 & 0 & 0 & 0 & 3 \end{bmatrix},$

then  $H = \begin{bmatrix} 4 & 1 & 1 & 2 \\ 0 & 1 & 2 & 1 \\ 1 & 2 & 1 & 1 \\ 0 & 3 & 0 & 0 \end{bmatrix},$  How many times  
; and  $j$  are neighbours

transition from gray level  $P$   
to gray level  $m$   
and  $P(l, m, 1, 0) = \frac{H(l, m)}{20}.$

1 pixel is immediate

neighbors of the other one to the right

# Co-Occurrence Matrix

No need to distinguish between

$$P(\underbrace{m, l}_{\text{transition}}, \Delta i, \Delta j)$$

and

$$P(l, m, \Delta i, \Delta j)$$

→ Co-Occurrence matrix C:

$$C = H + H^T$$

↑  
marked C symmetric

# Second Order Measures

Once you have matrix H, you can find

Contrast:

$$\sum_{i,j=0}^{N-1} P_{i,j} (i-j)^2$$

probability of this happening often  
range for pixels whose colors are different

are black or white  
Maximized: for image where pixels  
Minimized: uniform

Dissimilarity:

$$\sum_{i,j=0}^{N-1} P_{i,j} |i - j|$$

Homogeneity:

$$\sum_{i,j=0}^{N-1} \frac{P_{i,j}}{1 + (i-j)^2}$$

and many more .....

Angular Second Moment ( Energy, Uniformity):

$$\sum_{i,j=0}^{N-1} P_{i,j}^2$$

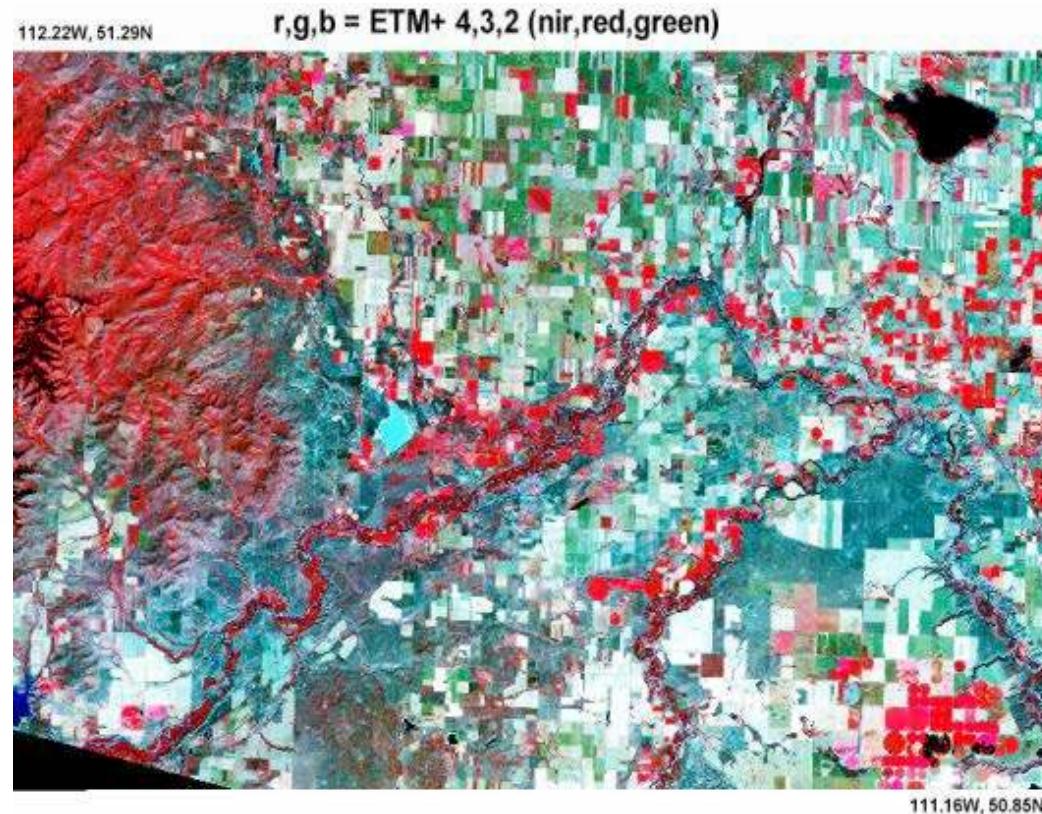
Entropy:

$$\sum_{i,j=0}^{N-1} P_{i,j} (-\ln P_{i,j})$$

measure of disorder  
• 2, 3 colors  $\rightarrow$  entropy low  
• multiple colors  $\rightarrow$  entropy high  
with uniform presence

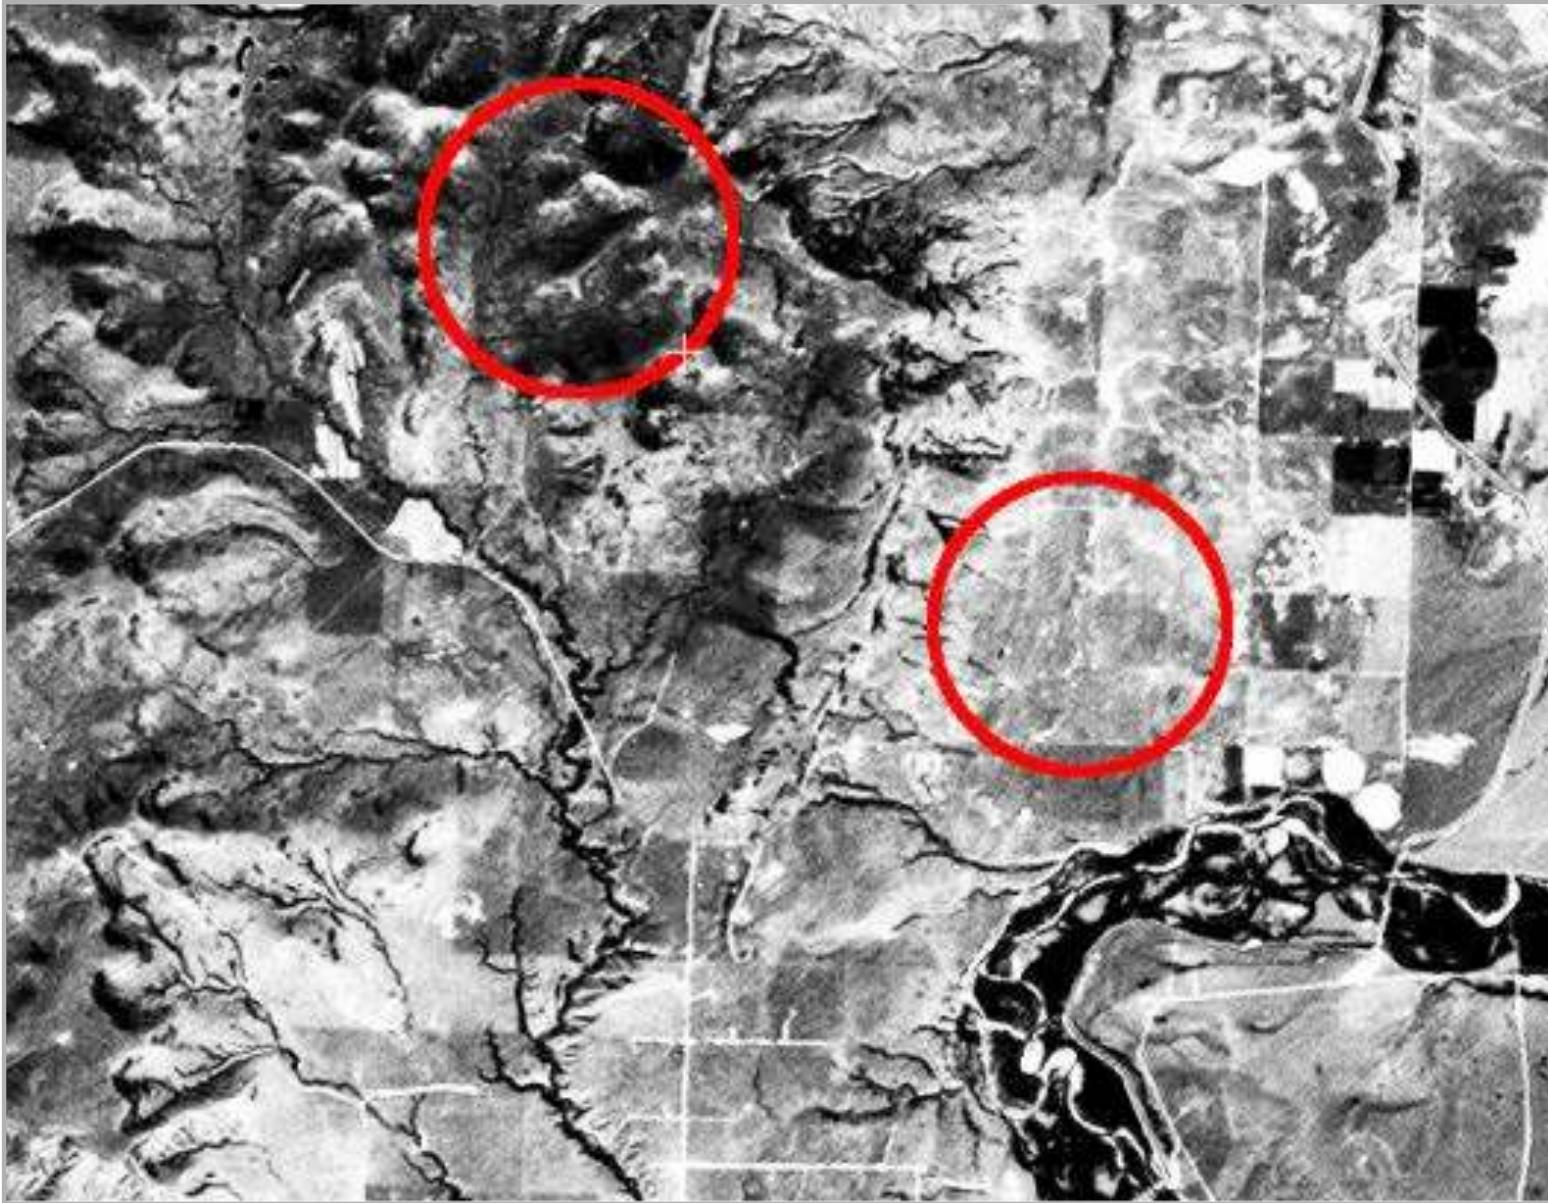
# Landsat Image

Aerial image



The image is excerpted from Path 41, Row 25 of Landsat 7 ETM+, dated 4 September 1999. This is an area in the Rocky Mountain Foothills near Waterton National Park, Alberta. The western edge of the image contains steep slopes and deep valleys. To the east is both grassland and annual crops, mostly grains. The eastern area is bisected by numerous small streams.

# Full Resolution

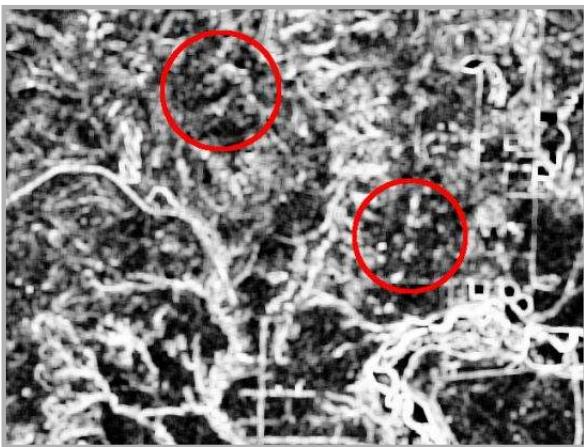


Let us consider two areas, one in the hills, the other in the plain.

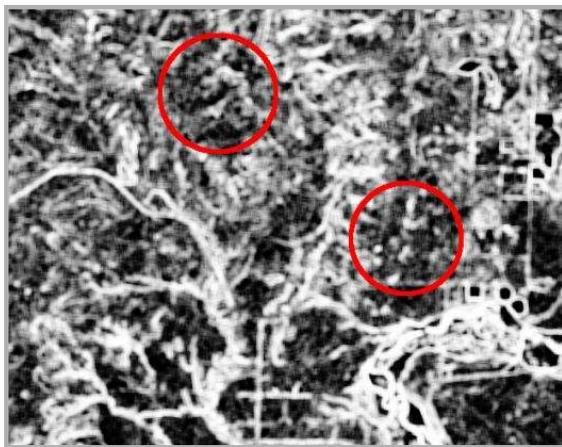
Complete Second-order matrices for those regions  
(co-occurrence matrices)

# Using Second Order Measures

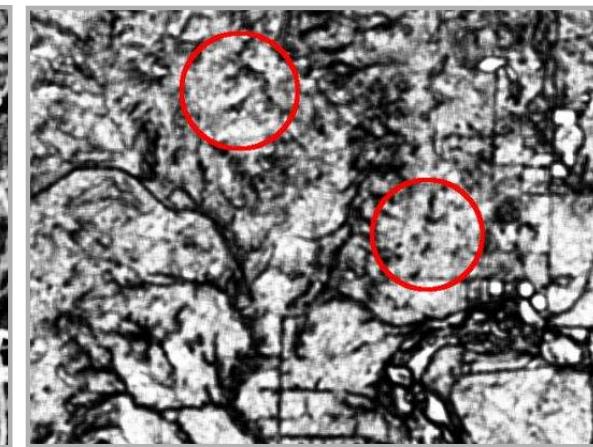
Contrast (for the whole img)



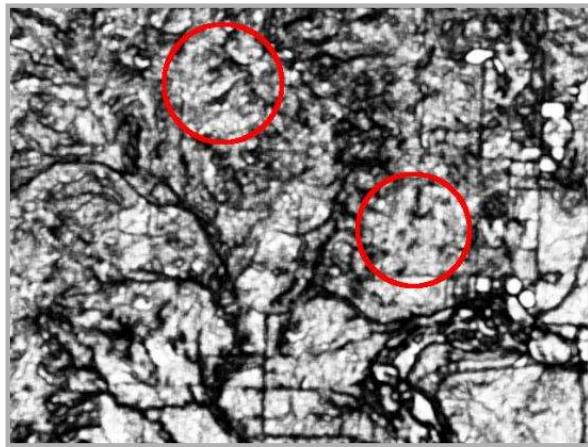
Dissimilarity



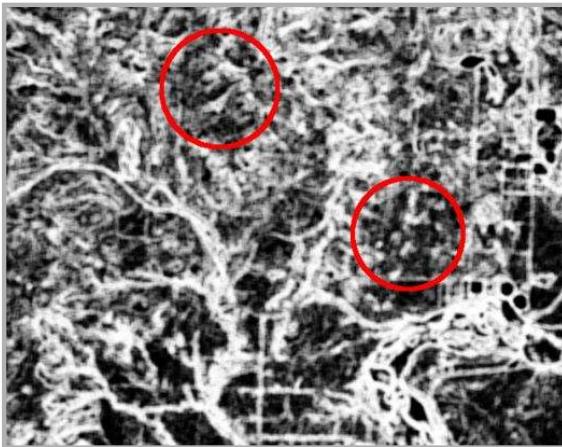
Homogeneity



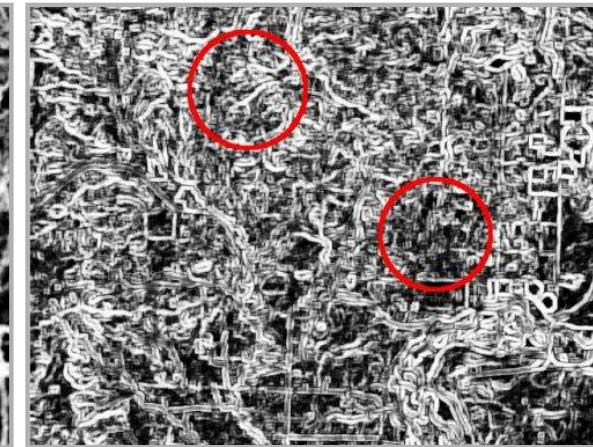
They have different statistics



ASM



Entropy



Correlation

The various measures within the two areas are sufficiently different for a classifier to easily distinguish them.

# Classification

Used to identify eight terrain classes:

- Old residential
- New residential
- Urban
- Lake
- Swamp
- Scrub
- Wood



# Parameter Choices

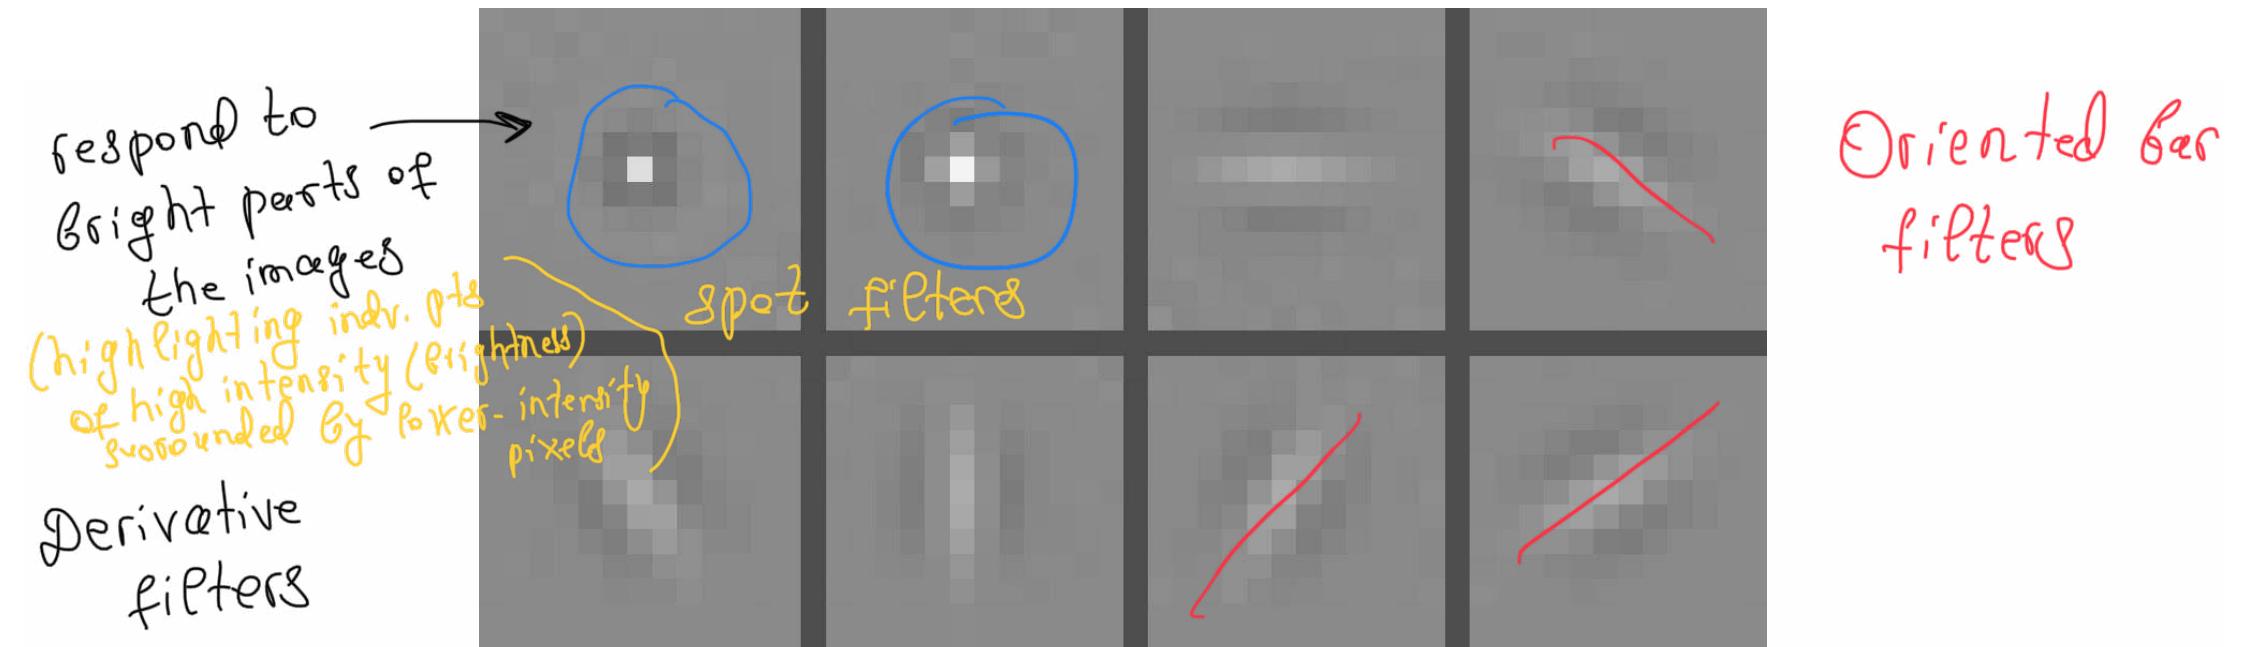
Using co-occurrence matrices requires choosing:

- window size, *→ During the matrix computation*
- direction of offset,
- offset distance,
- what channels to use,
- what measures to use.

How do we choose these parameters?

- Critical question for **all** statistical texture methods.
- Can be addressed using Machine Learning.  
*with good validation set*

# Filter Based Measures

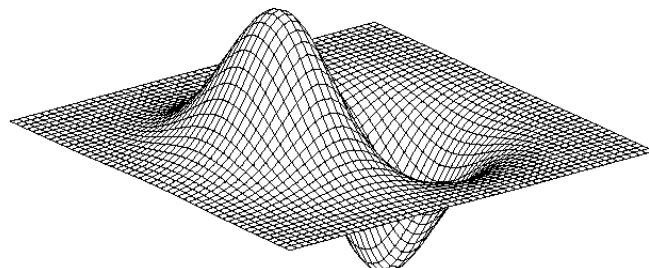


Represent image textures using the responses of a collection of filters. → convolve with block filters

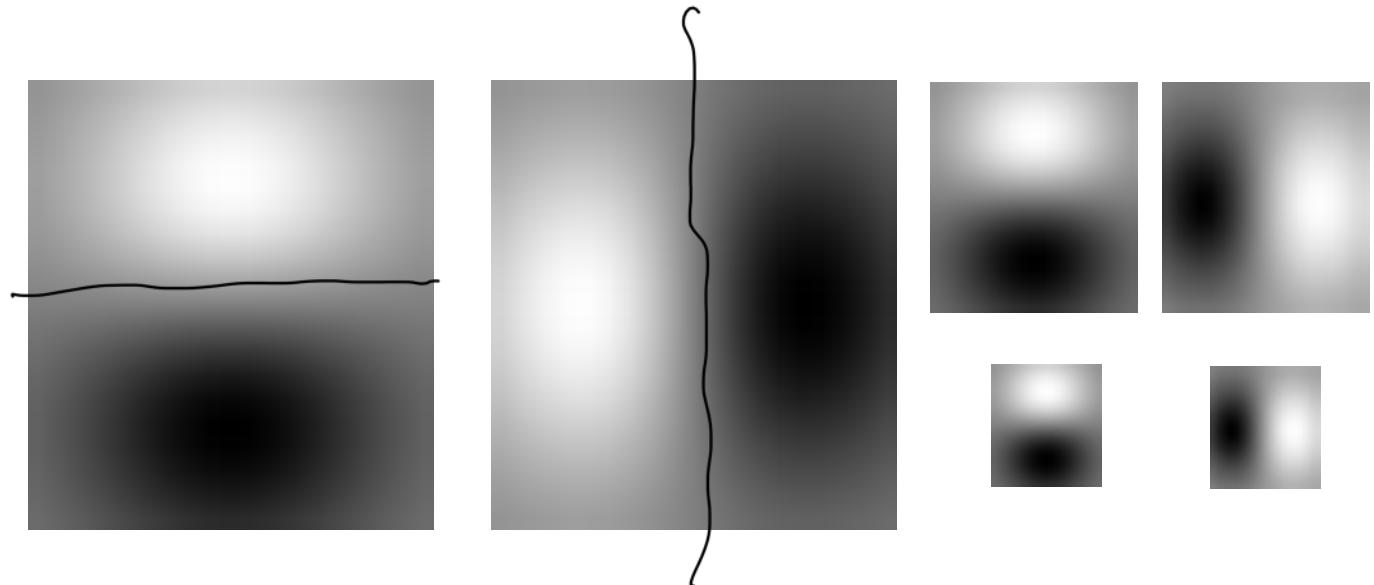
- An appropriate filter bank will extract useful information such as spots and edges. *in various directions*
- Traditionally one or two spot filters and several oriented bar filters. *in various orientations*

# Gaussian Filter Derivatives

(similar to what we did for edges)



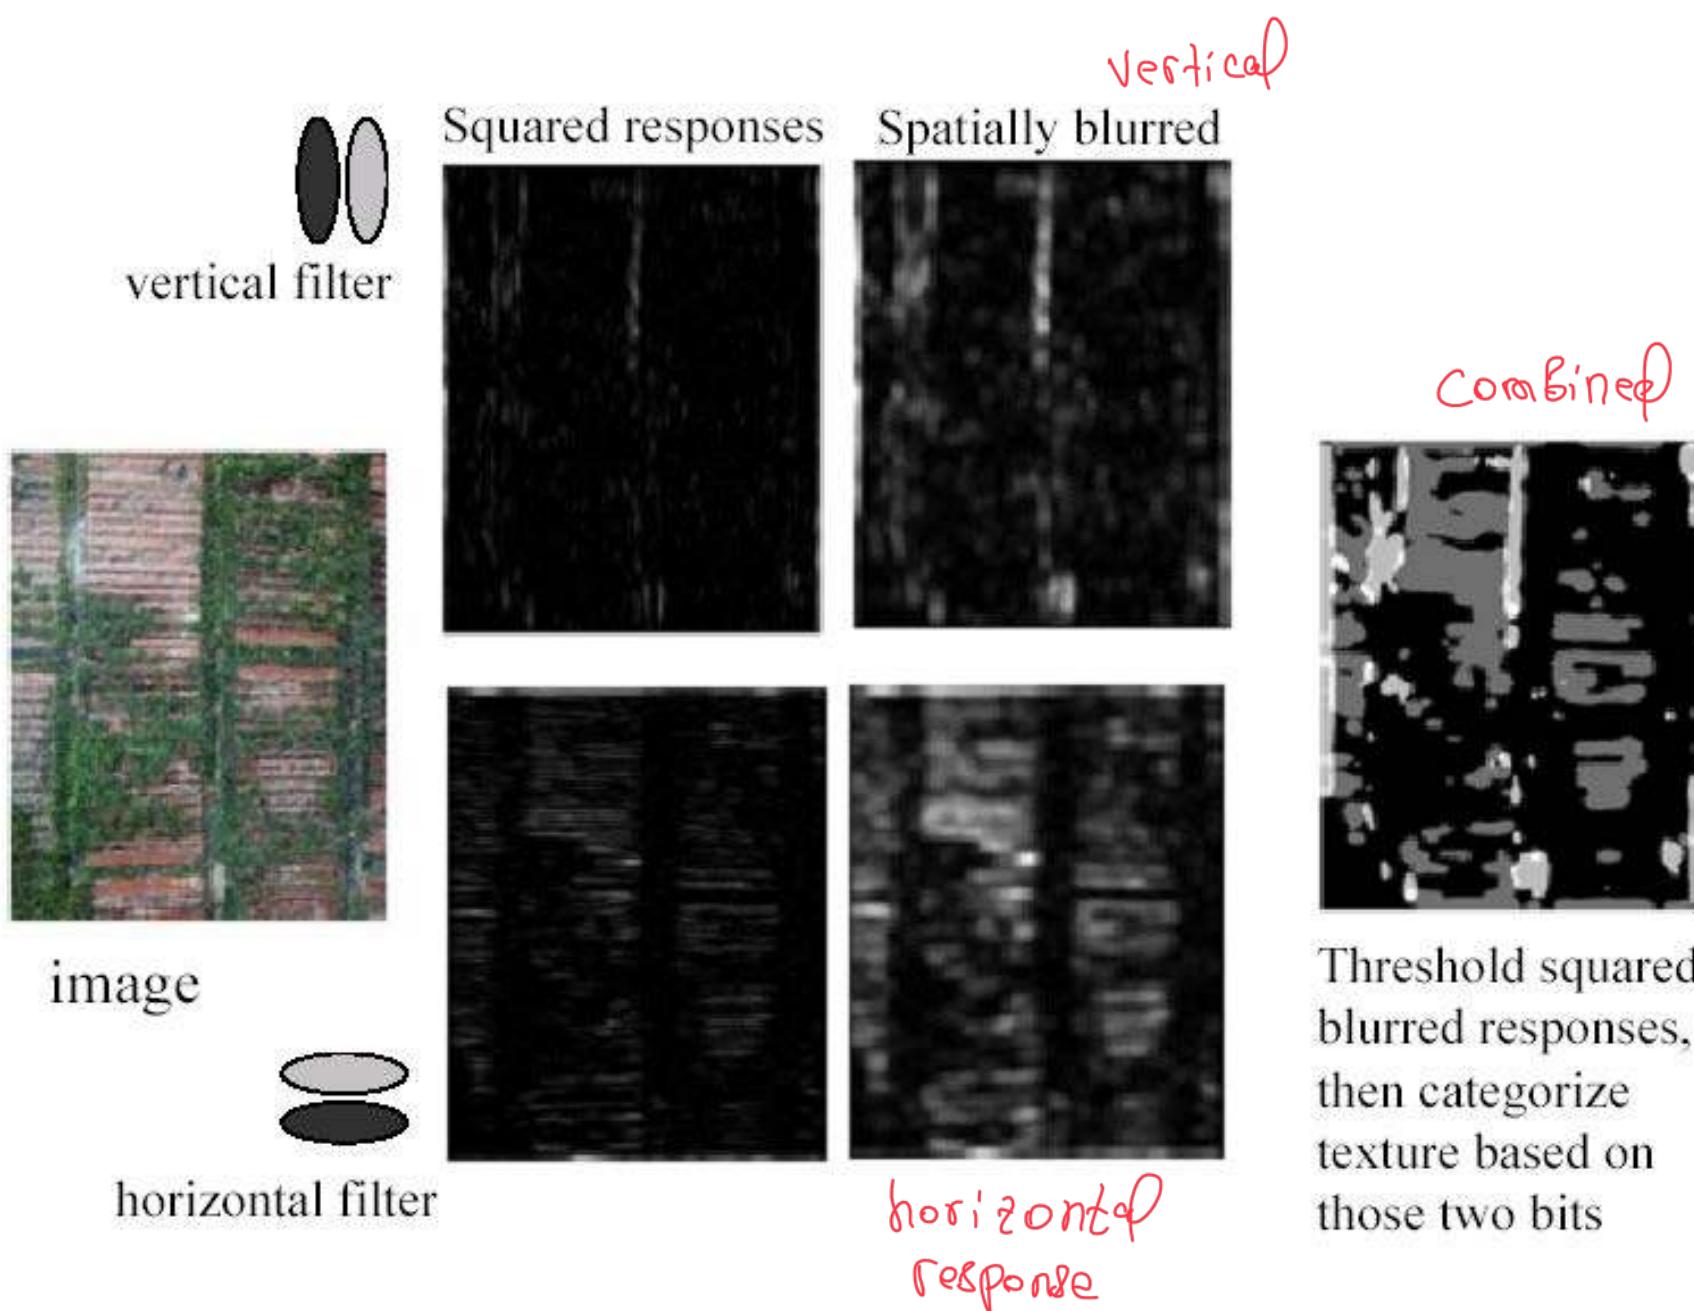
Gaussian Derivative



x and y derivatives at different scales

These filters respond to horizontal and vertical edges.

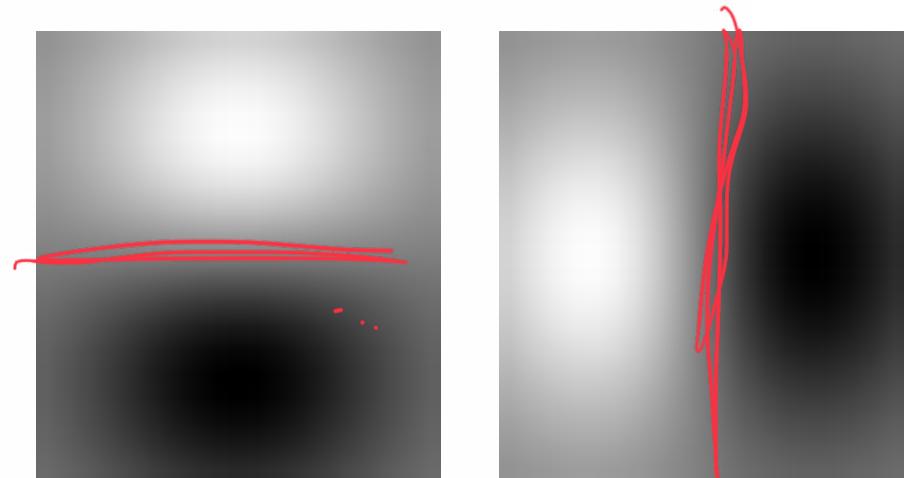
# Horizontal and Vertical Structures



# Oriented Filters

Rotated Gaussian filter

Gaussian function → weighted avg where  
central pixels have the most  
influence and the influence  
decreases for pixels further away

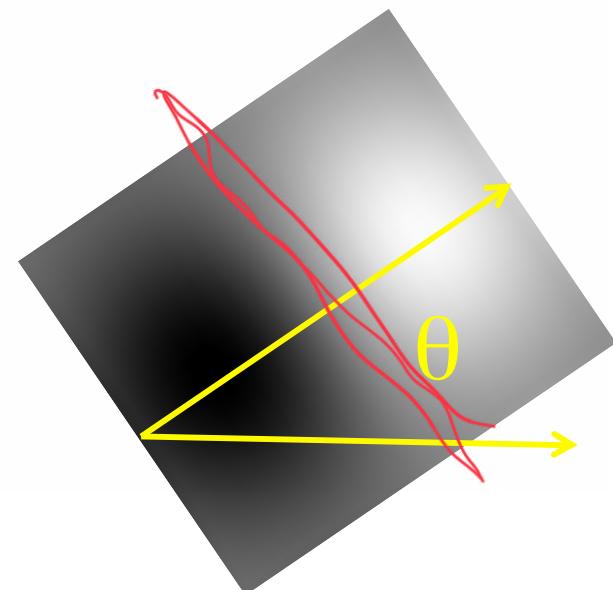


$$\frac{\partial I}{\partial \theta} = \cos(\theta) \frac{\partial I}{\partial x} + \sin(\theta) \frac{\partial I}{\partial y}$$

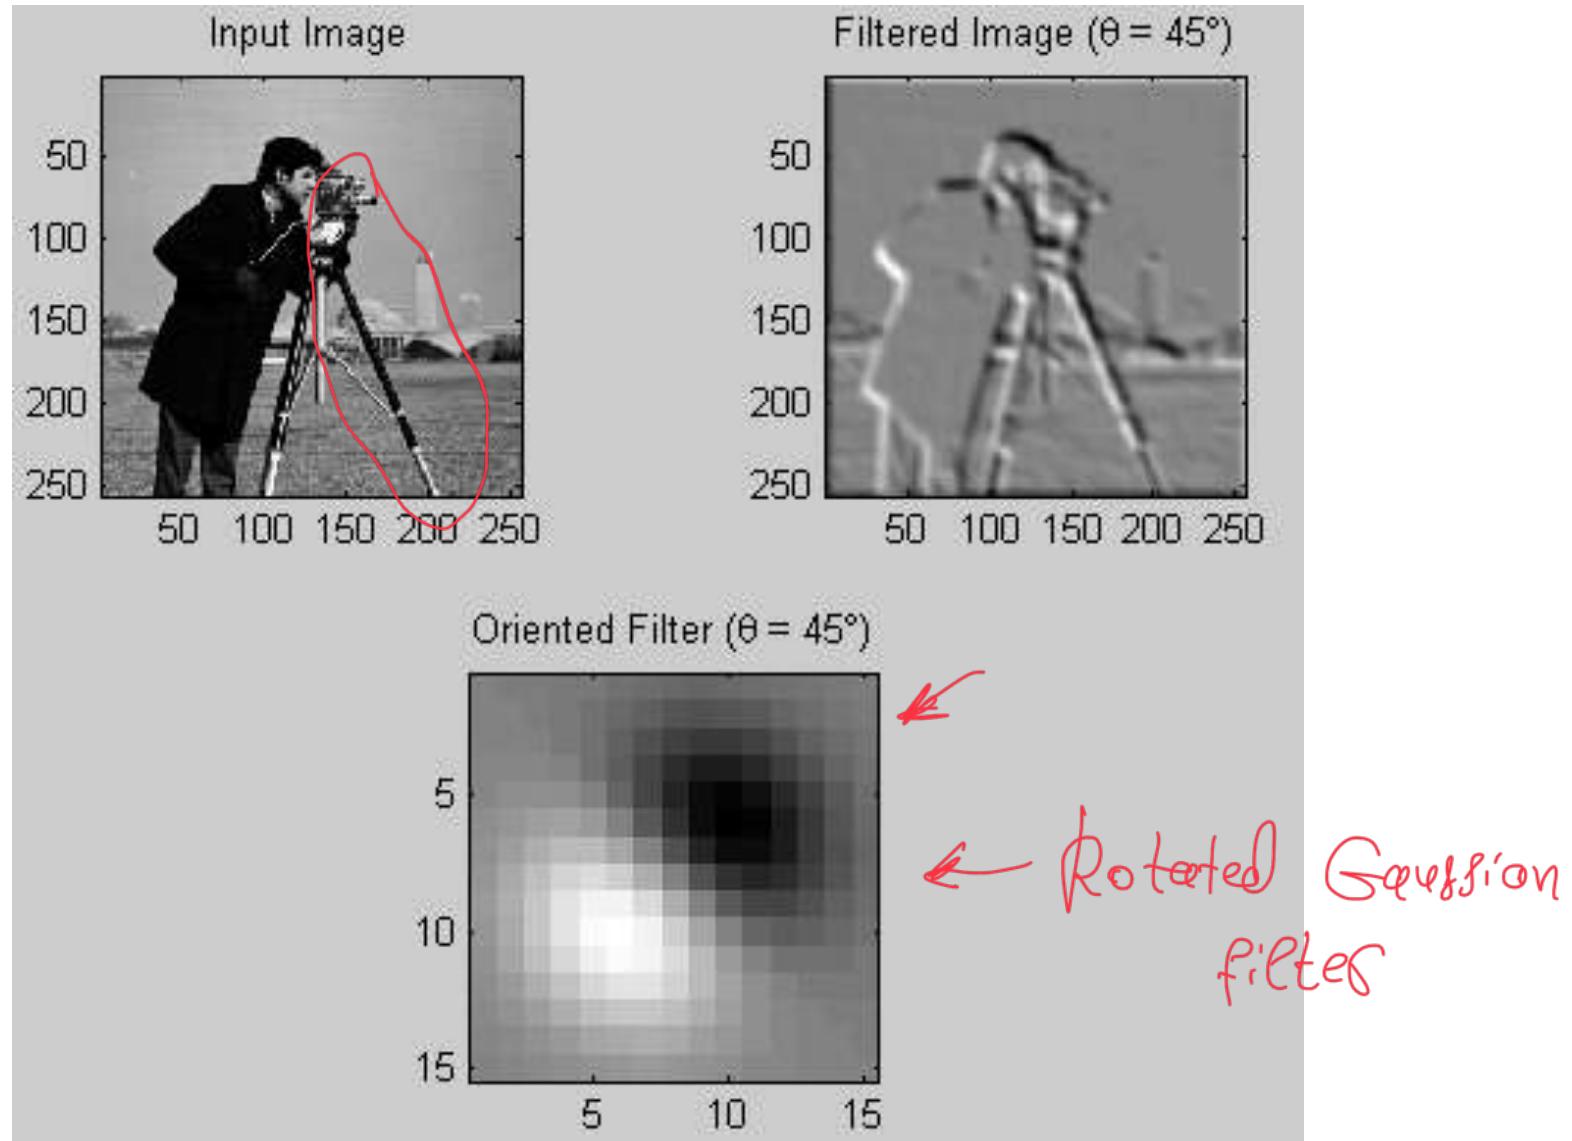
✓

✓

computed with Gaussian

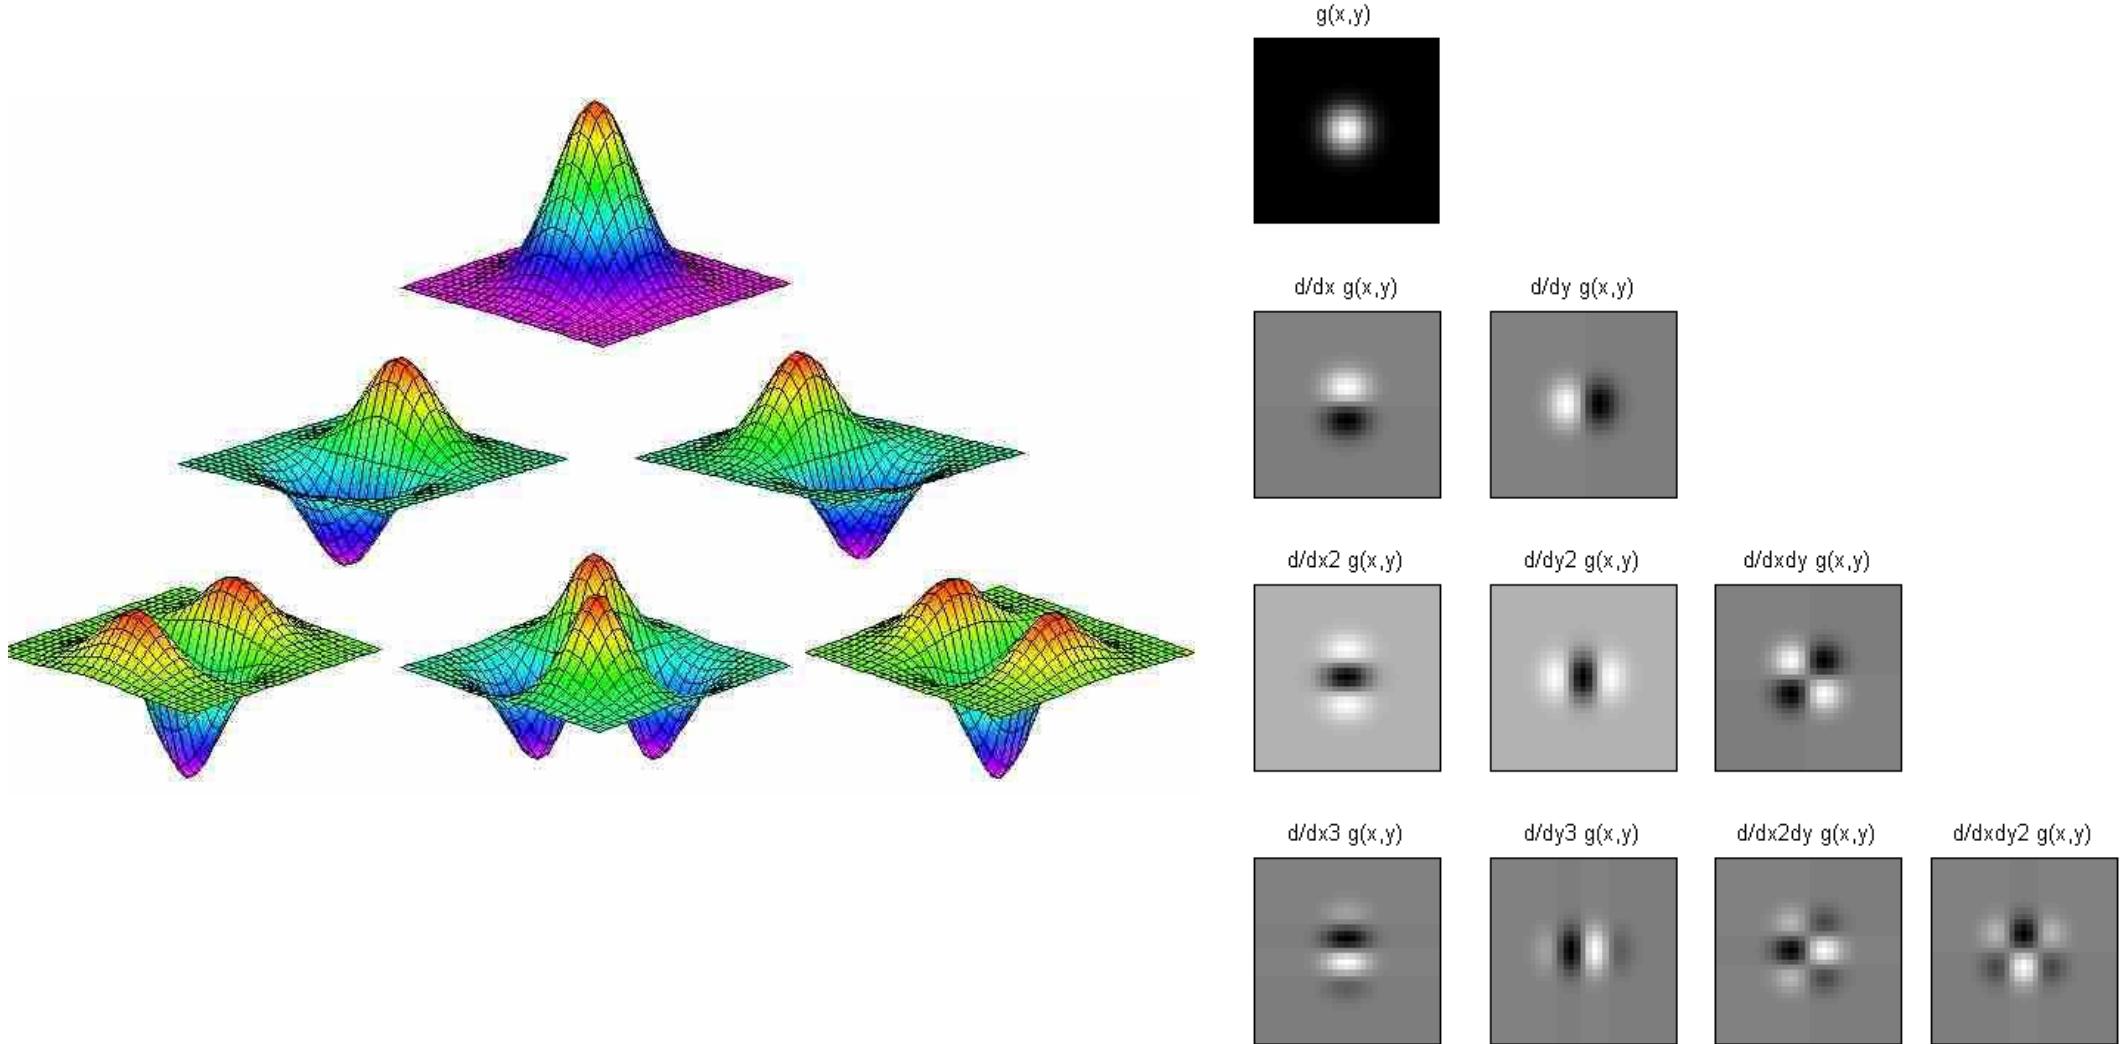


# Directional Gradients



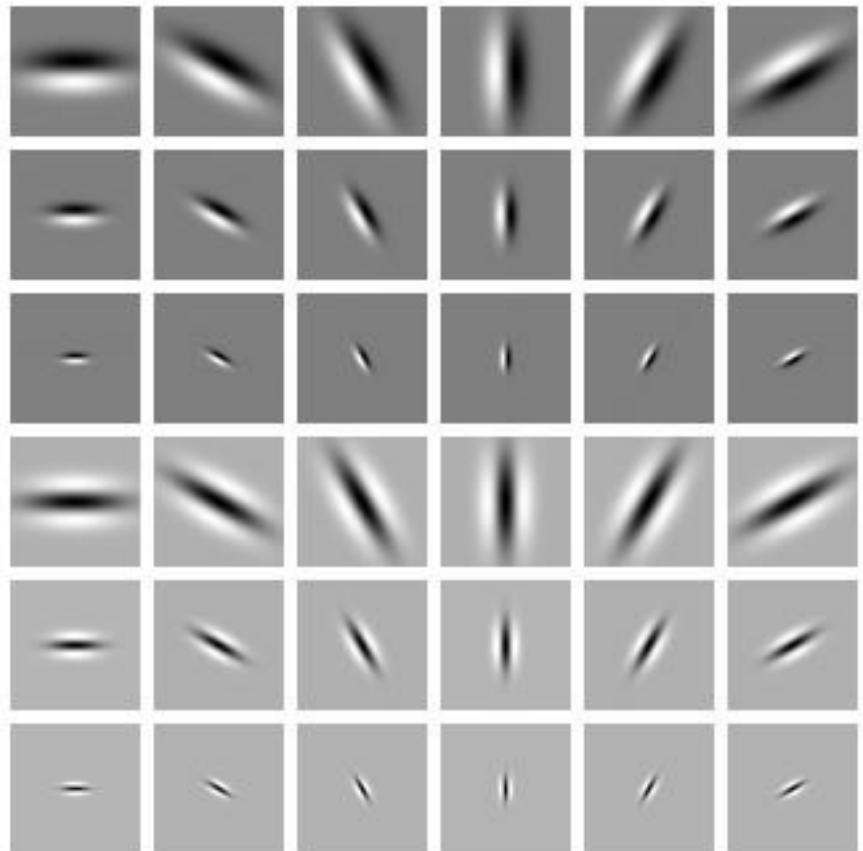
Oriented filters respond to edges in a specific direction.

# Higher Order Derivatives



Higher-order derivatives of the Gaussian filters can be used to compute higher-order image derivatives.

# Filter Bank

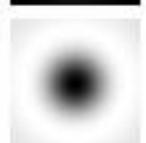


various rotated versions  
of your filters to detect structures  
at various orientations

- Different scales.
- Different orientations.
- Derivatives order 0, 1, 2 ..  
↳ detect changes in intensity at  
different scales



→ 2 spot filters



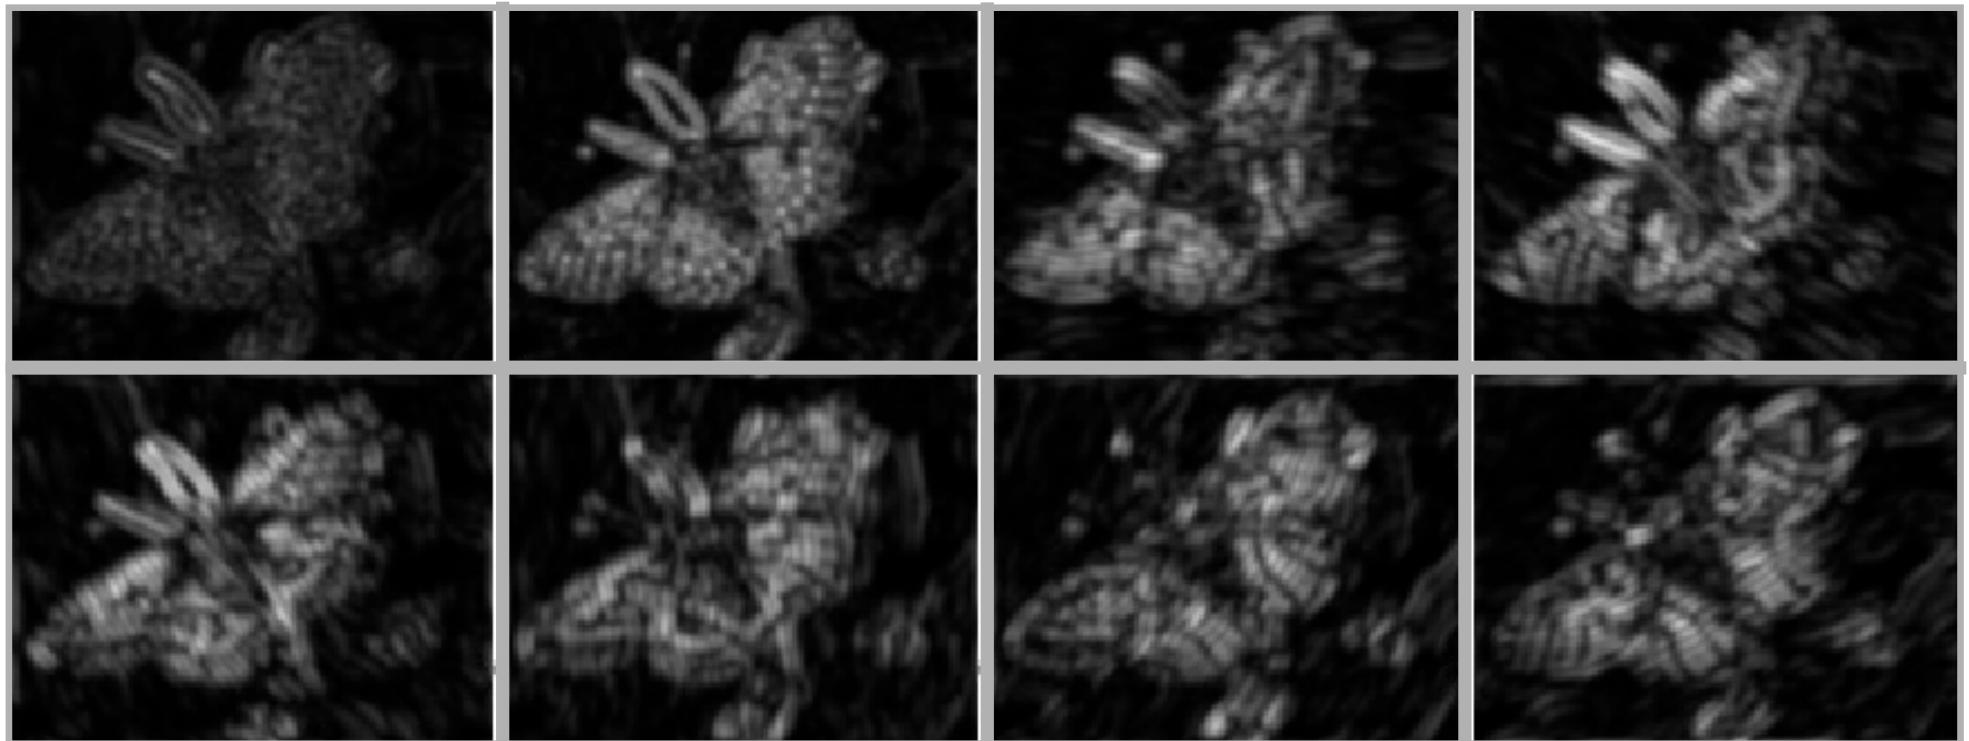
→ For every image pixel, compute a vector of  
responses to each filter.

convolve with each of  
these filters

and then replace that pixel  
with the latter computed vector  
of responses

# Filter Responses: Small Scales

Local Smoothing



Gaussian filters with a small  $\sigma$ . Capture local details.

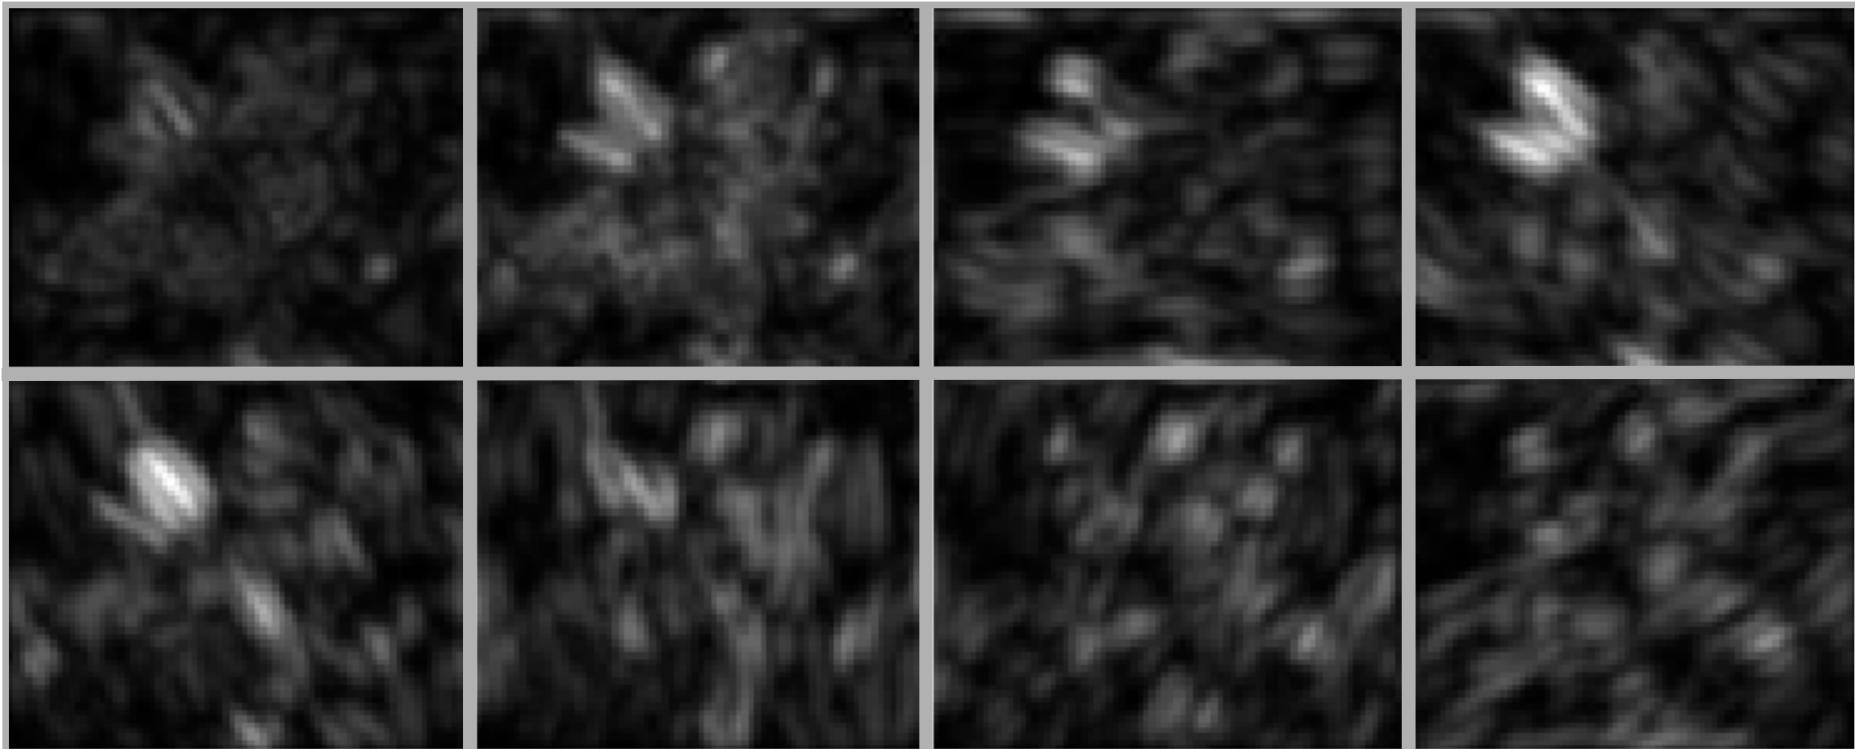
# Filter Responses: Large Scales

Captures more global details in the image:  
global smoothing or blurring



8 responses from small scale  
8 responses from large scale

pixel →  16-dimension to classify different parts of the butterfly



Gaussian filters with a large s. Capture larger details.

# Gabor Filters

→ Smoothing and blurring operator

Gabor filters are the products of a Gaussian filter with oriented sinusoids. They come in pairs, each consisting of a symmetric filter and an anti-symmetric filter:

$$G_{\text{sym}}(x, y) = \underbrace{\cos}_{\text{gives sensitivity to periodic structure (like edges)}}(\underline{k_x}x + k_y y) \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

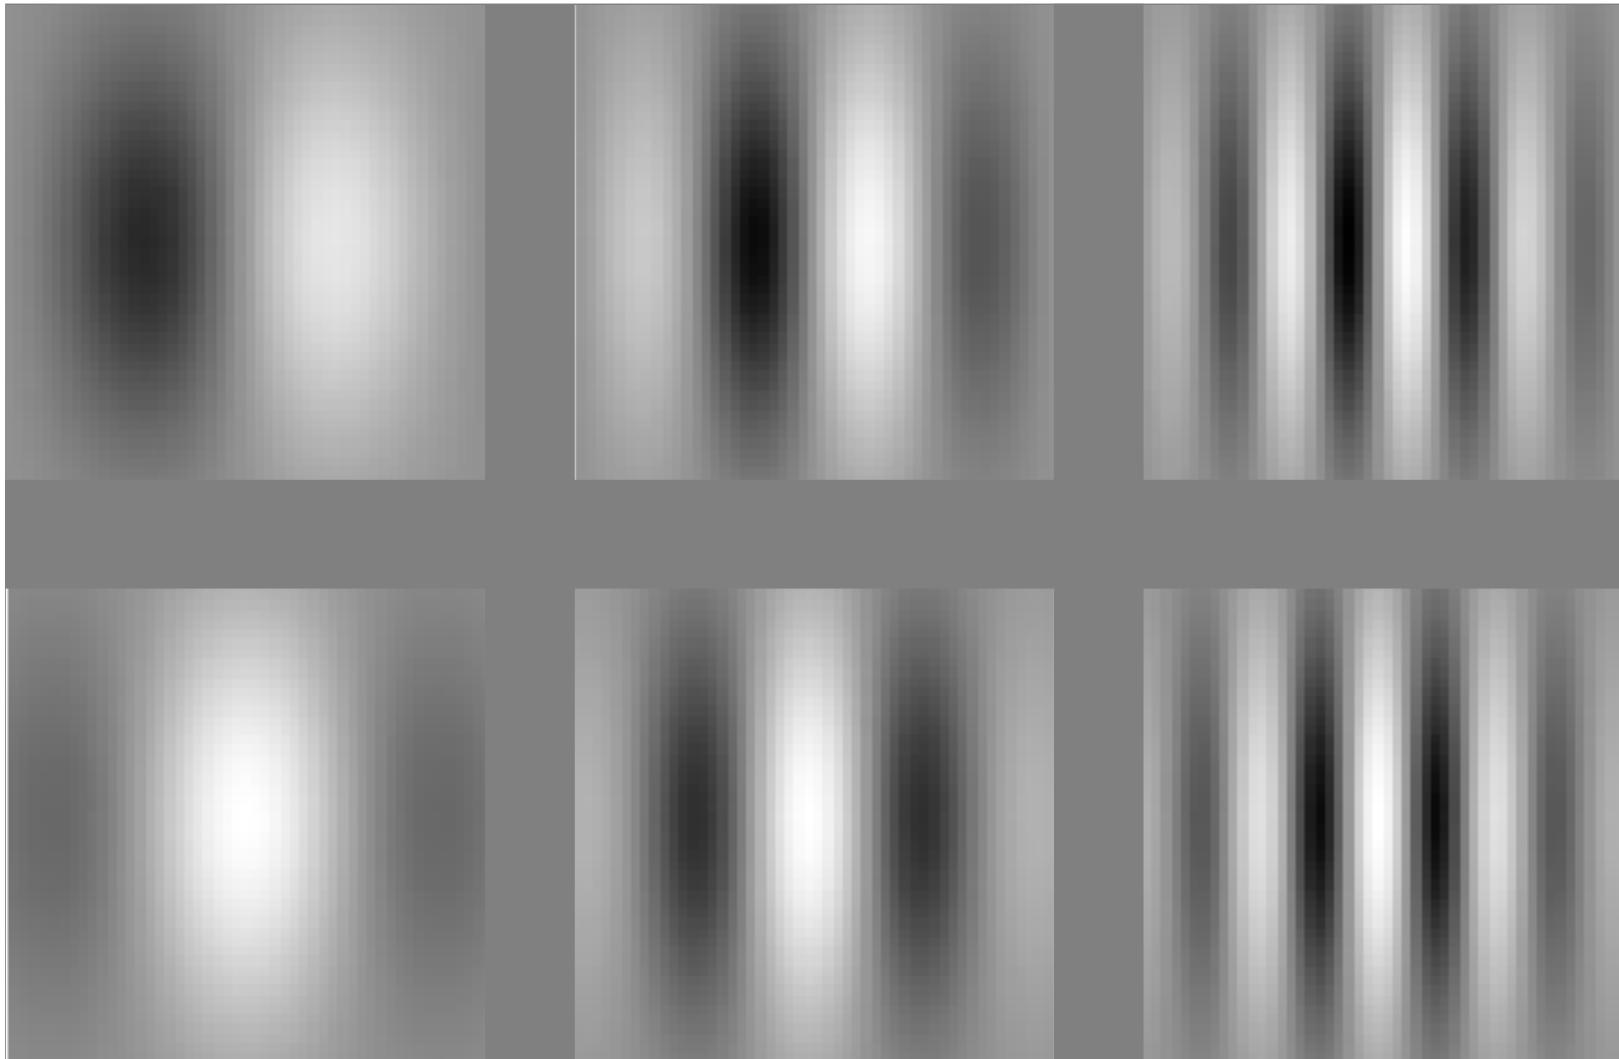
$$G_{\text{asym}}(x, y) = \underbrace{\sin}_{\text{more}}(\underline{k_x}x + k_y y) \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

where  $k_x$  and  $k_y$  determine the spatial frequency and the orientation of the filter and  $\sigma$  determines the scale.

→ A filter bank is formed by varying the frequency, the scale, and the filter orientation of Gaussian filters

# Vertical Derivatives

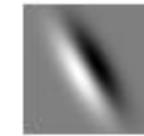
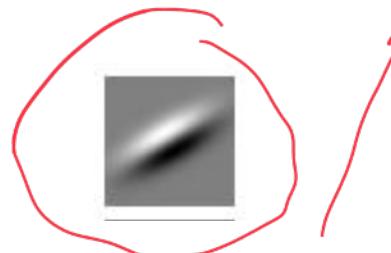
at different  
frequencies



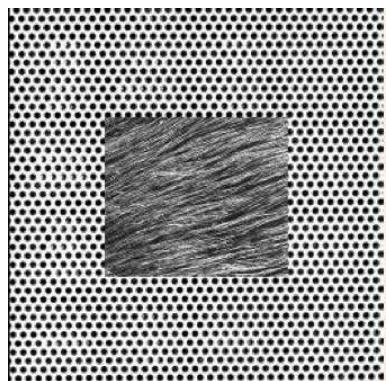
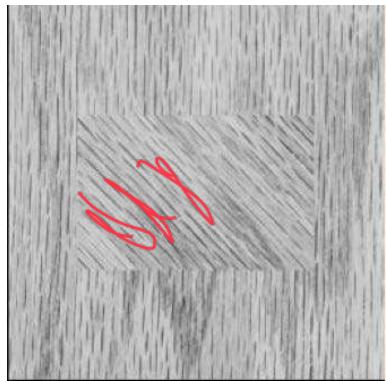
# Gabor Responses

not that strongly

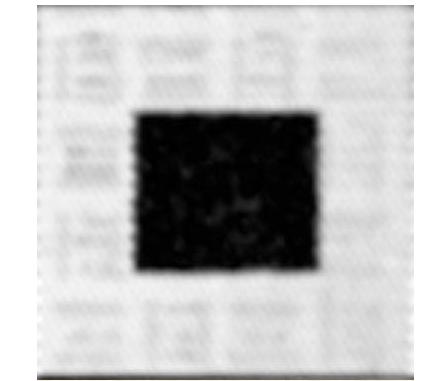
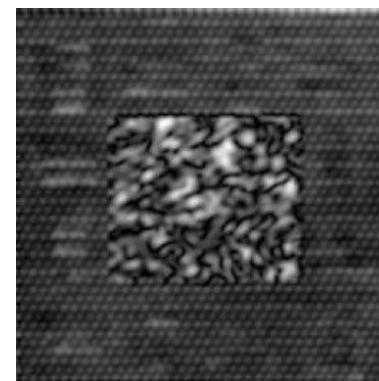
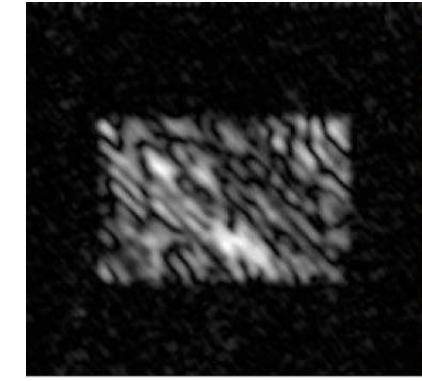
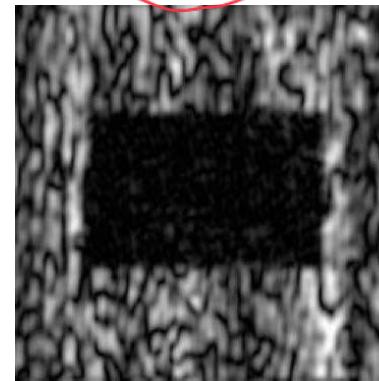
Filters:



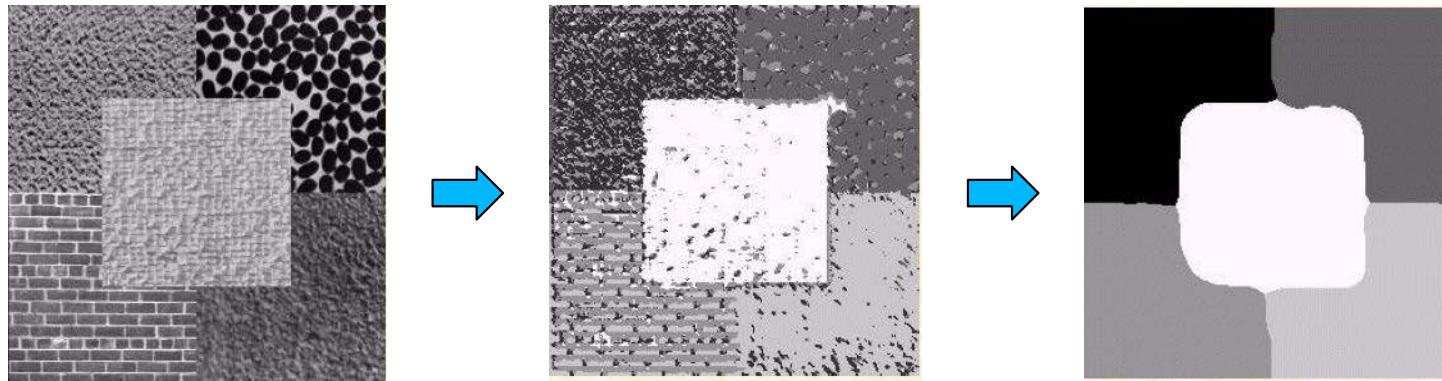
Images:



Responses:

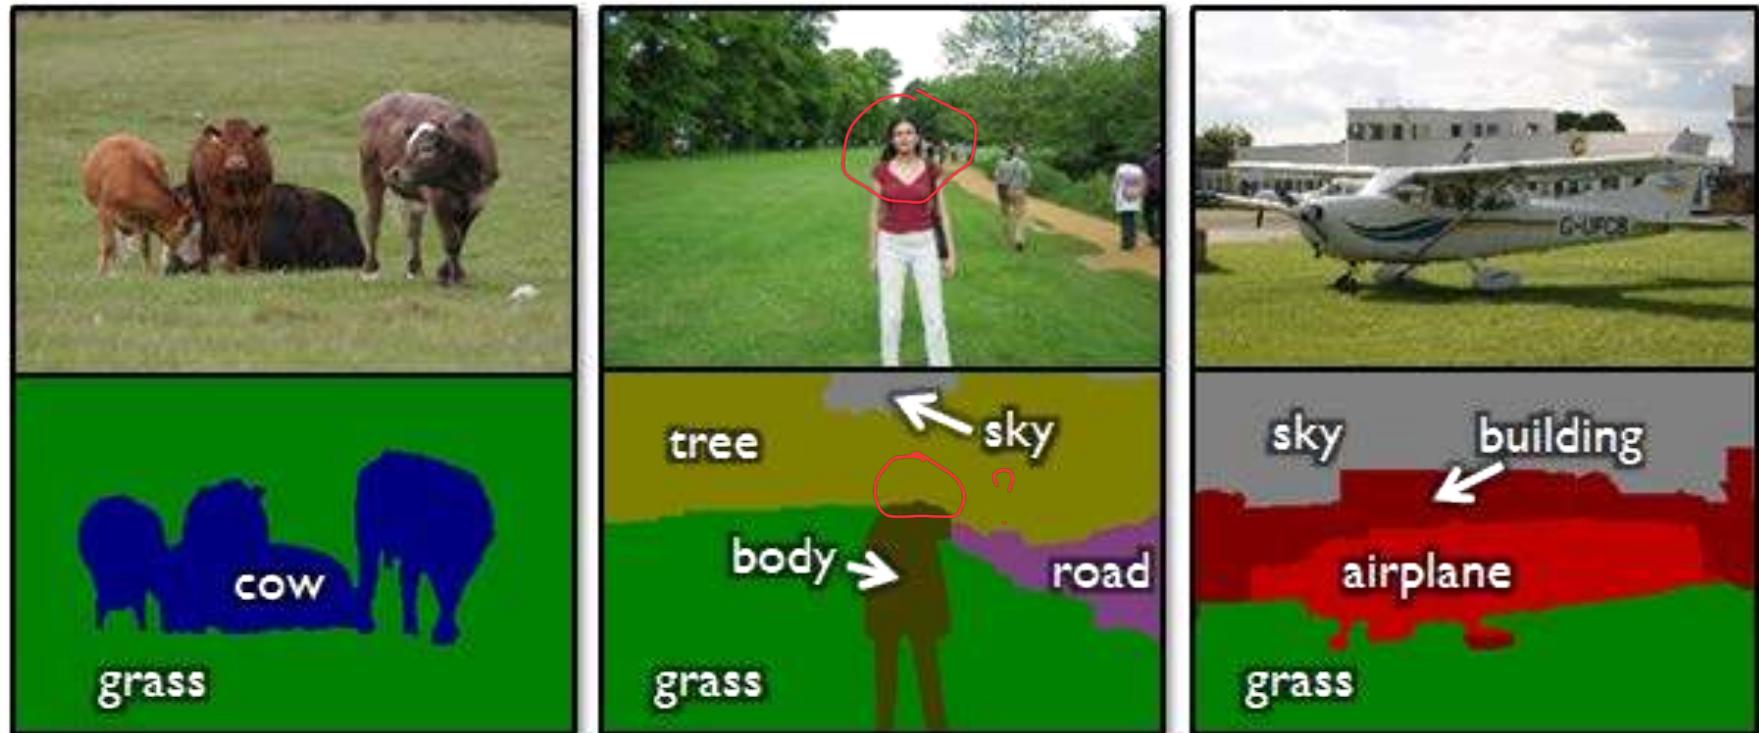


# GABOR FILTER CHARACTERISTICS



- Respond strongly at points in an image where there are components that locally have a particular spatial frequency and orientation.  
*(measure of how often pixel vals change over Space  
pix → pix)*  
→ rate of change of pixel intensity in an img
- In theory, by applying a very large number of Gabor filters at different scales, orientations and spatial frequencies, one can analyze an image into a detailed local description.
- In practice, it is not known how many filters, at what scale, frequencies, and orientations, to use. This tends to be application dependent.

# ML to the Rescue: Texton Boost



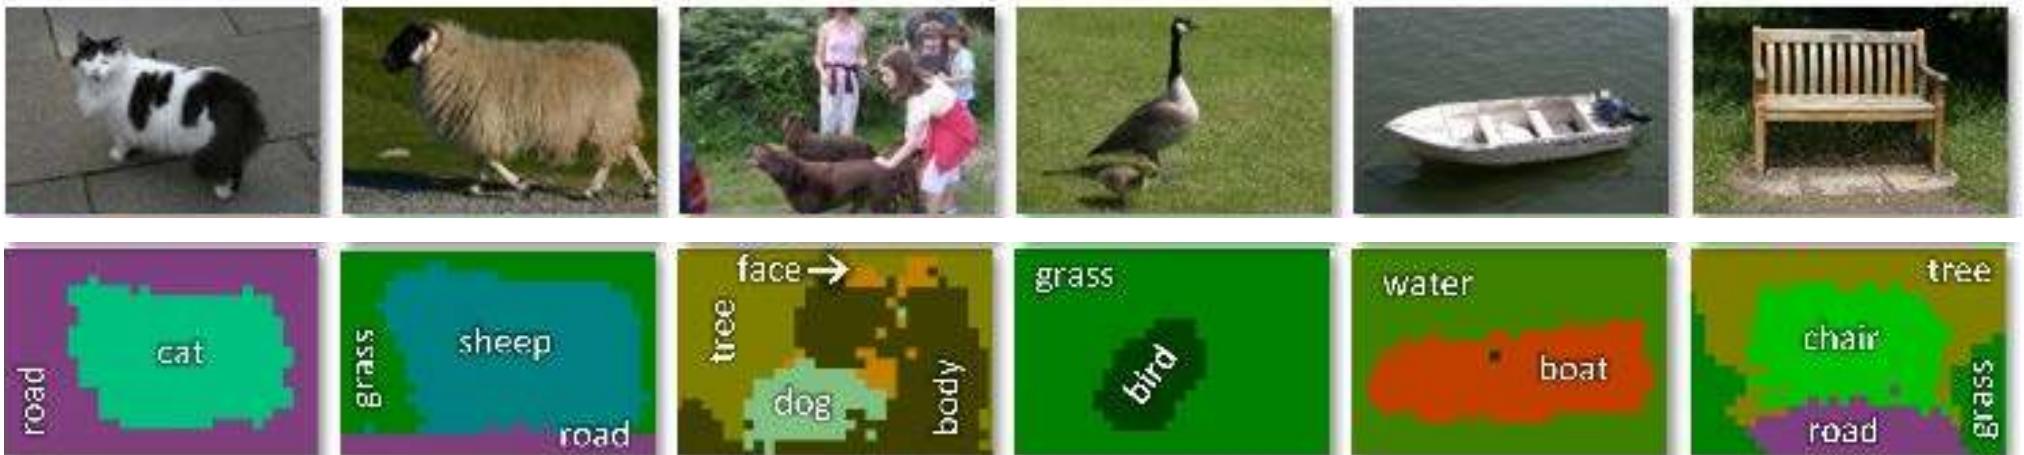
object classes	building	grass	tree	cow	sheep	sky	airplane	water	face	car
bicycle	flower	sign	bird	book	chair	road	cat	dog	body	boat

Use AdaBoost to perform classification on the output of Gabor filters.

find out which filters are useful  
and do segmentation on the filters  
they have ~ texture

Shotton et al., ECCV'06

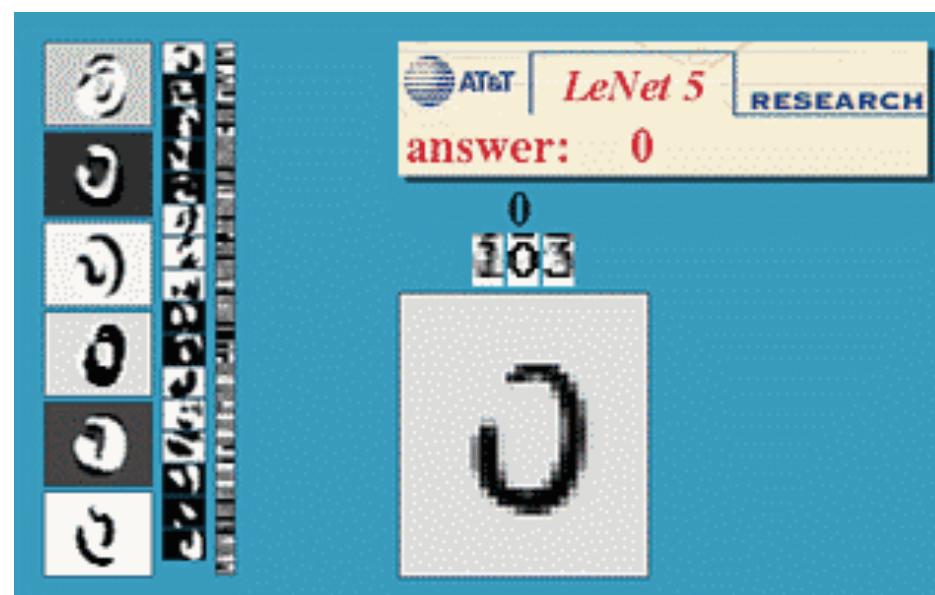
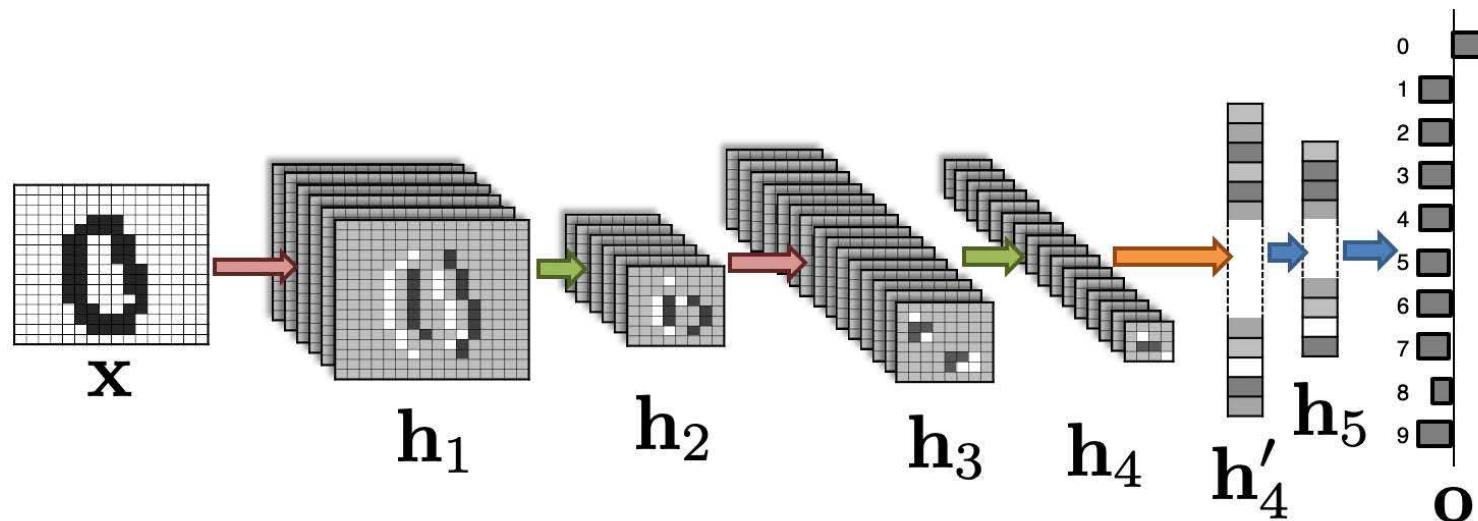
# ML to the Rescue: Texton Forests



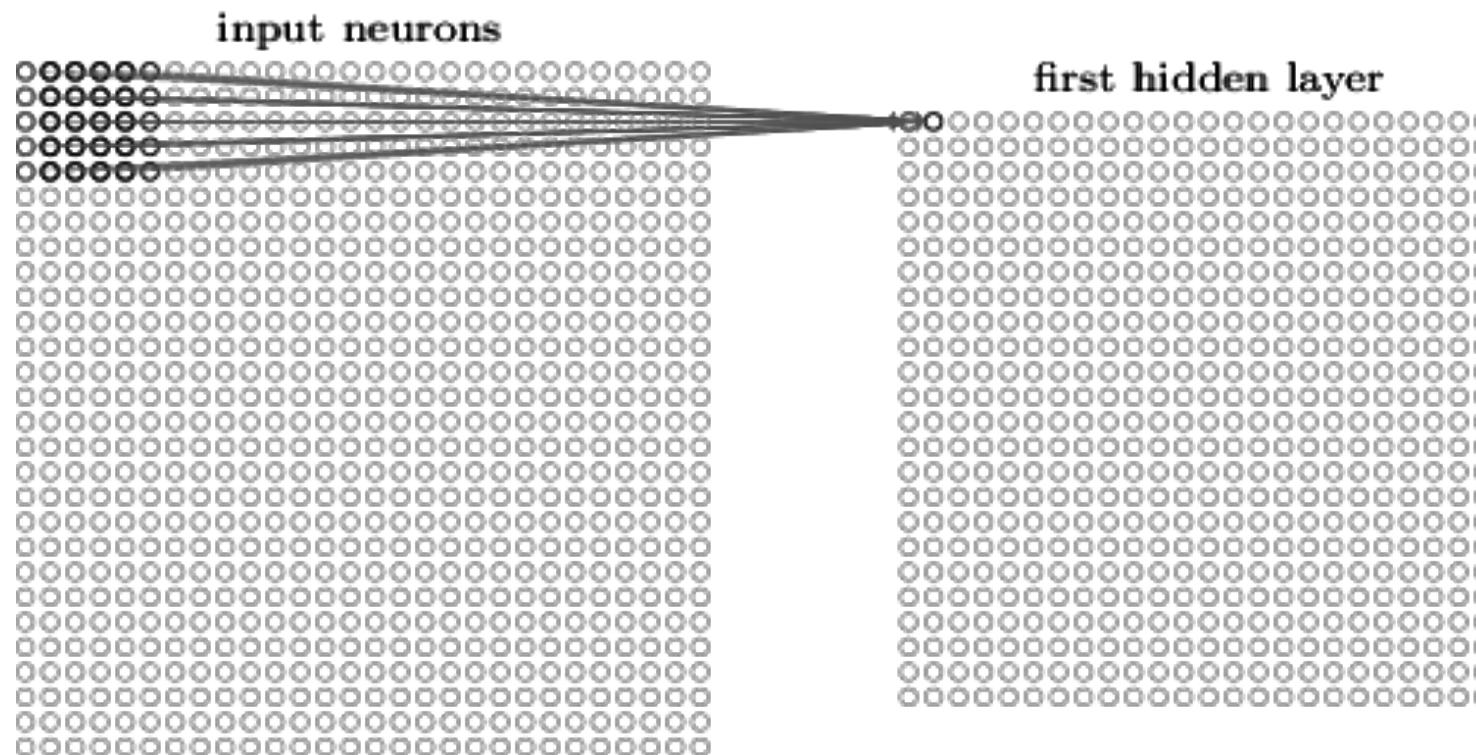
→ Ensemble learning (assembled classifiers)  
methods

- Using Decision Forests to perform classification on the output of Gabor filters works better in this case.
- But what works even better, is .....

# Reminder: ConvNets



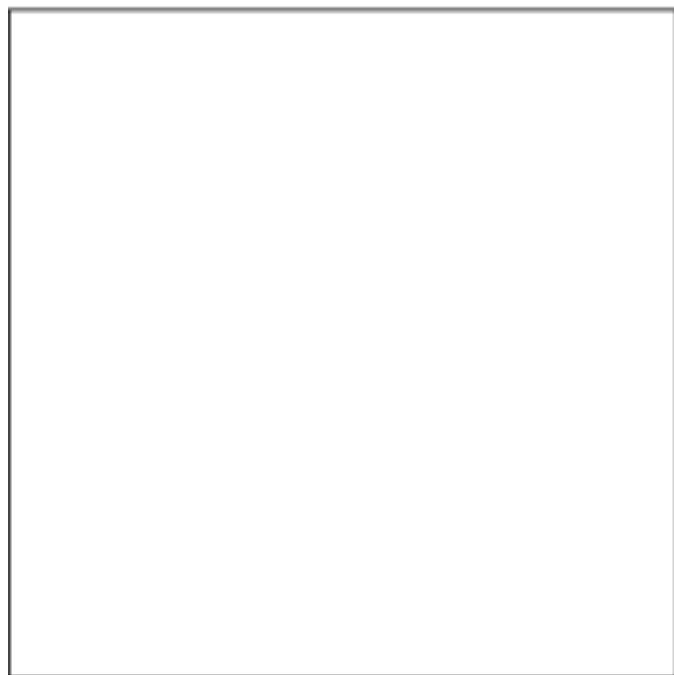
# Reminder: Convolutional Layer



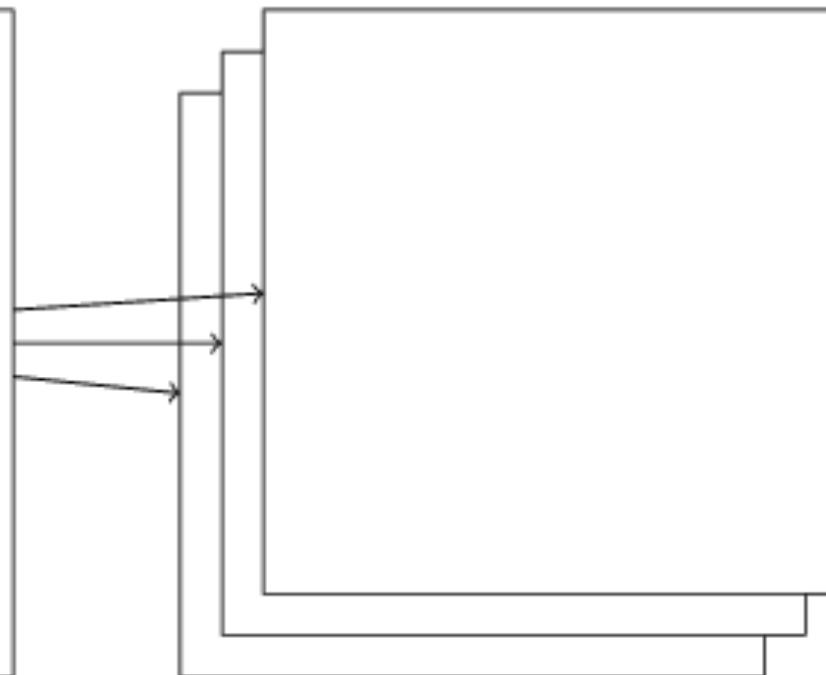
$$\sigma \left( b + \sum_{x=0}^{n_x} \sum_{y=0}^{n_y} w_{x,y} a_{i+x,j+y} \right)$$

# Reminder: Feature Maps

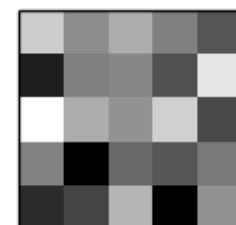
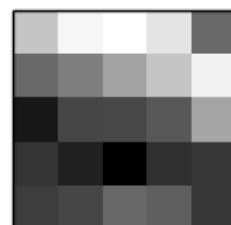
28 × 28 input neurons



first hidden layer:  $3 \times 24 \times 24$  neurons



Filters:

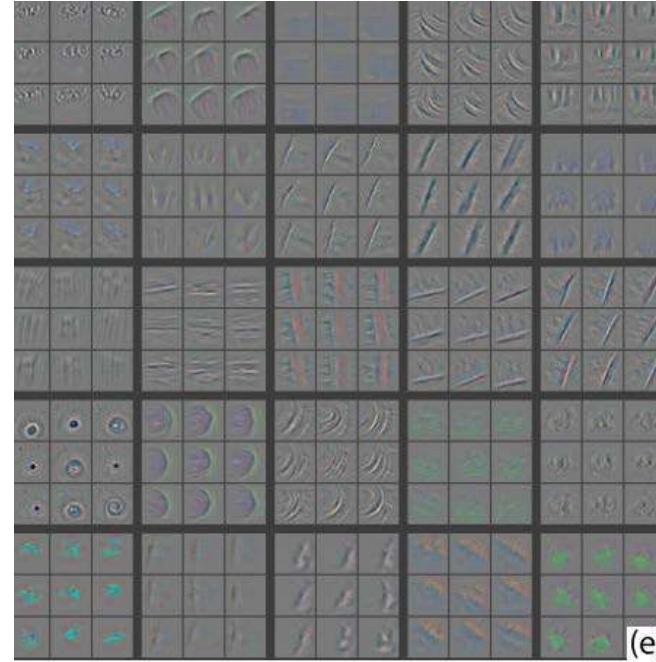
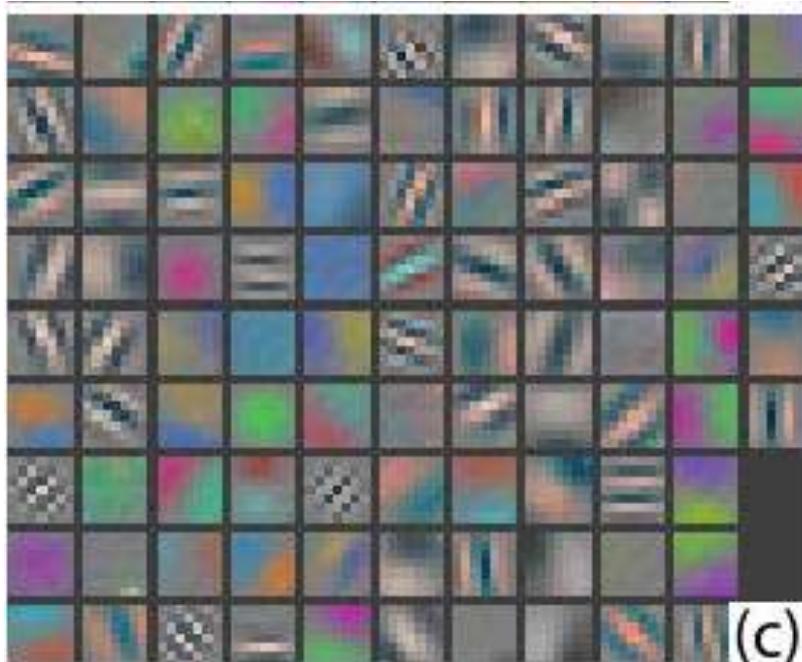


→ It learns the weights

Convolve with different filters and then train network

# Learned Feature Maps

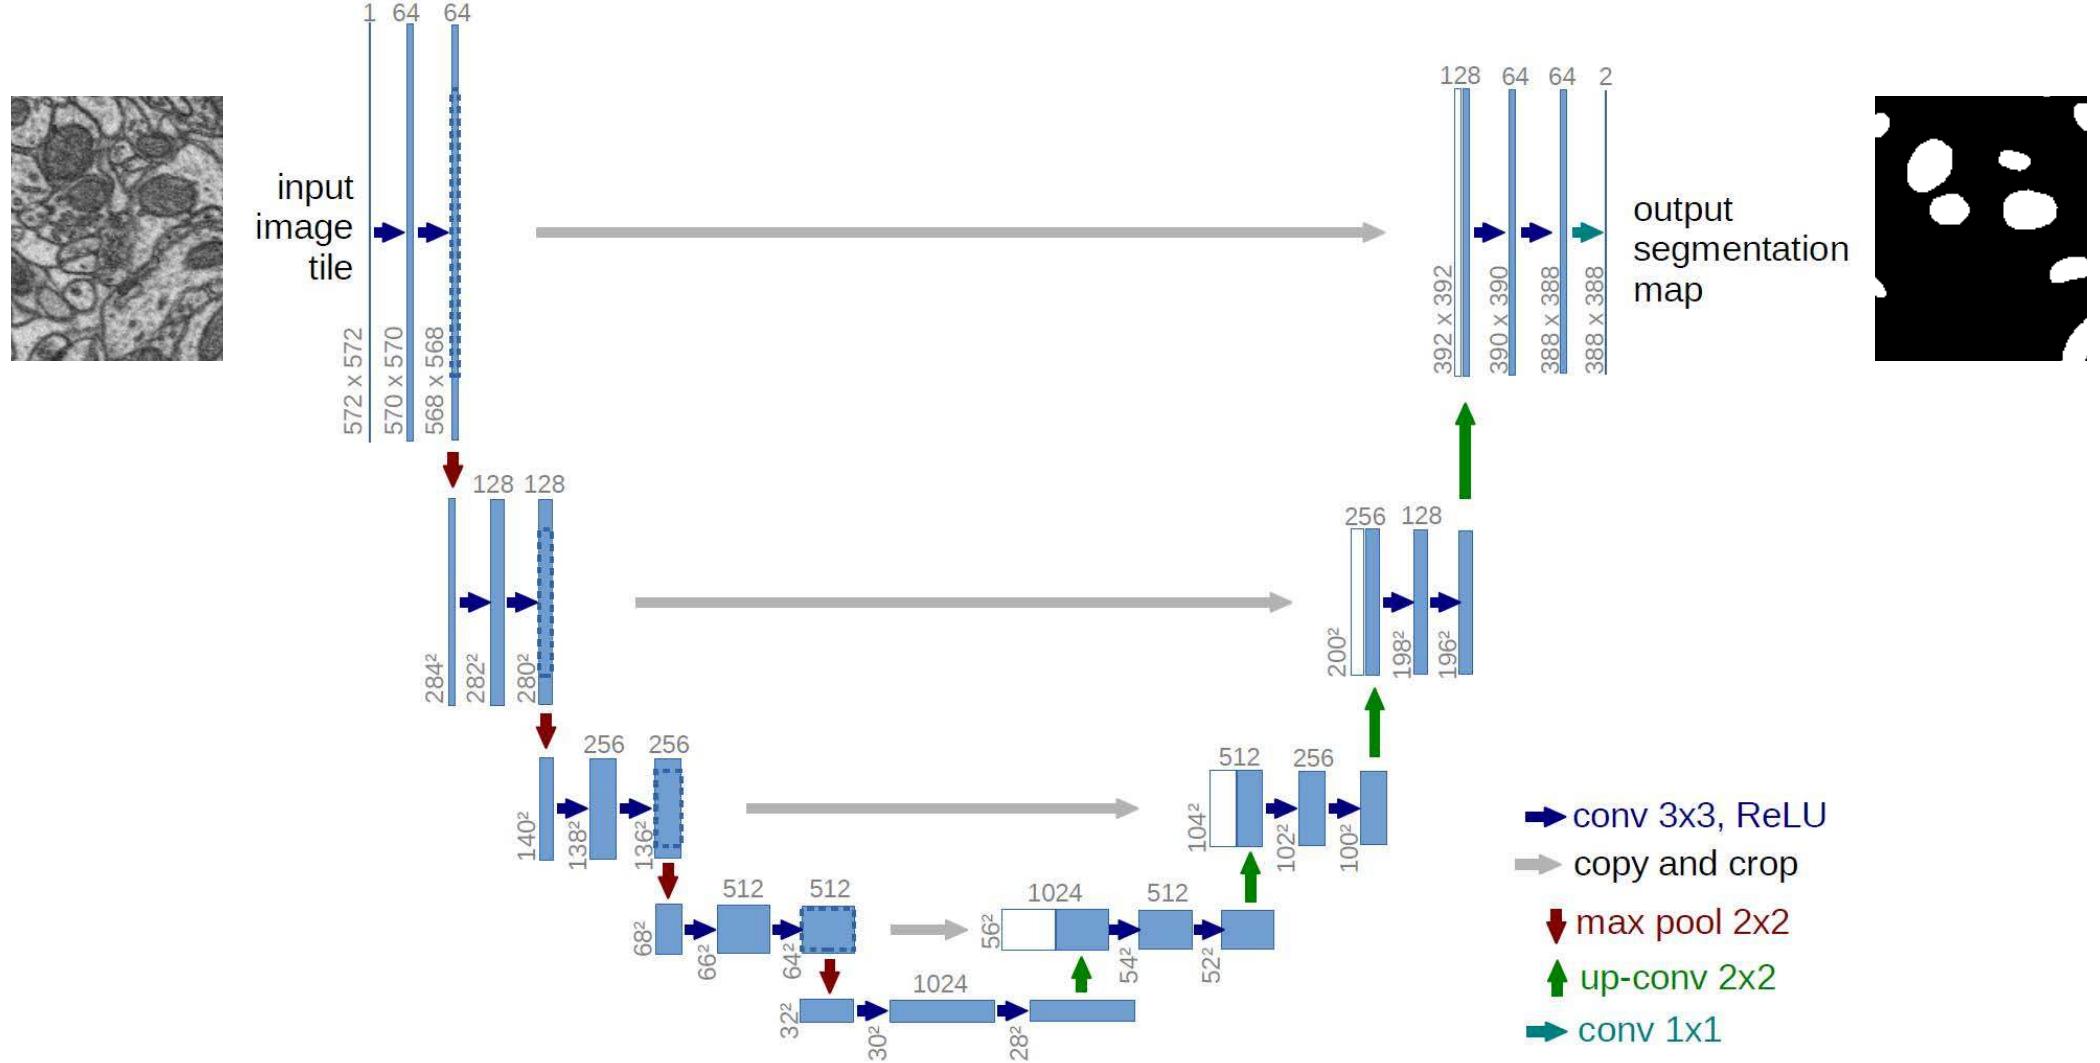
ImageNet



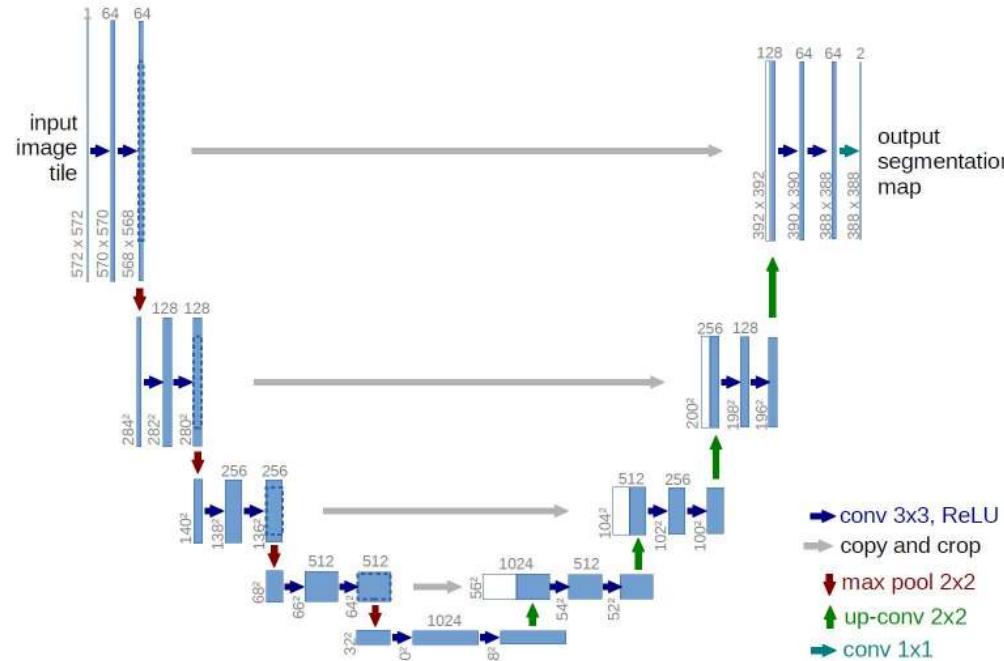
- Some of these convolutional filters look very Gabor like.
- The network requires a large training set to learn an effective filter bank. (*Gabor-like filters*)
- The older techniques still have their place in the absence of such training sets.

manually choose  
Gabor filters

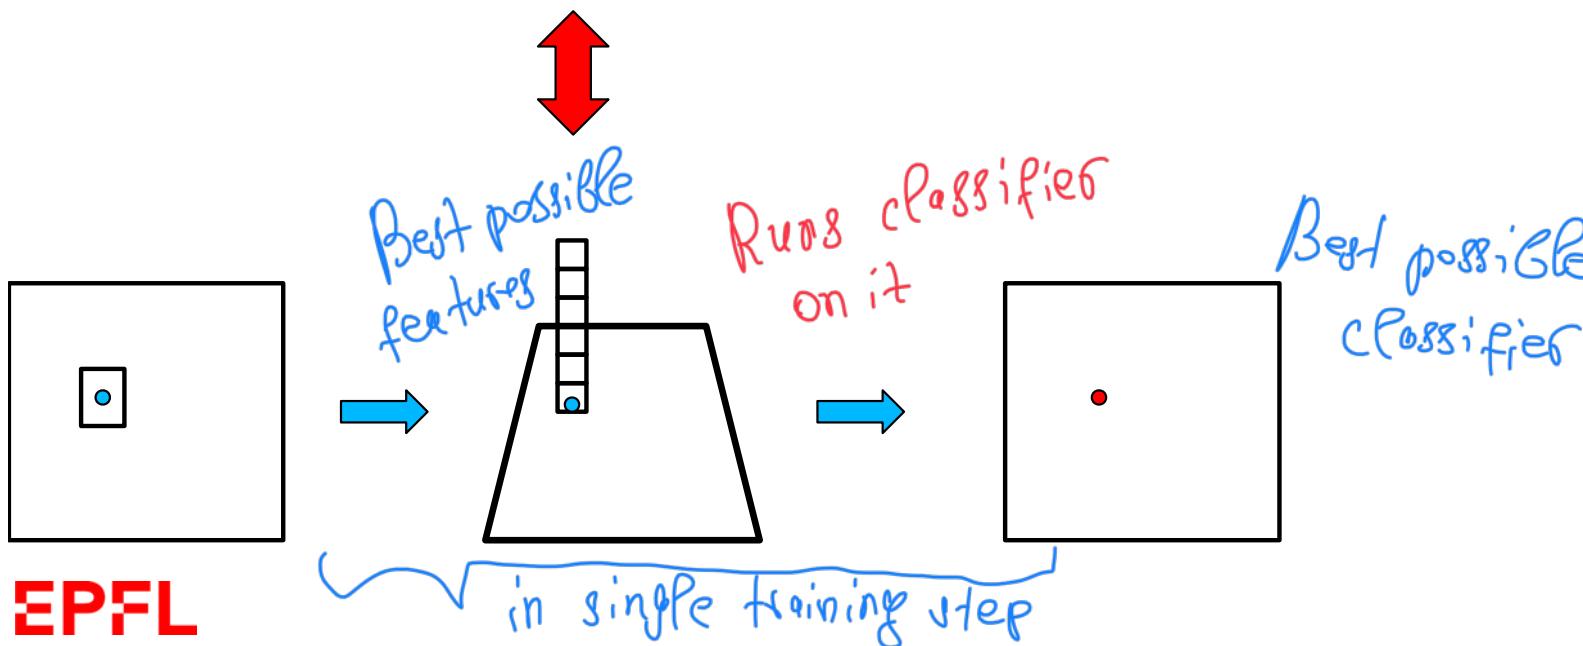
# Reminder: U-Net Architecture



# Reminder: Potential Interpretation



A key role of the ConvNet is to generate for every output pixel a feature vector containing the output of all the intermediate layers.



# In Short

Texture is a key property of objects which is

- Non local → texture of the pixel
  - Non trivial to measure
  - Subject to deformations  
distortions
- Hard to characterize formally and best used in conjunction with effective Machine Learning techniques.
- This seems to be exactly what Convolutional Neural Nets do.